XC32 Compiler for PIC32M

MPLAB XC32 C/C++ Compiler User's Guide for PIC32M MCUs

Notice to Customers

Important:

All documentation becomes dated and this manual is no exception. Microchip tools and documentation are constantly evolving to meet customer needs, so some actual dialogs and/or tool descriptions may differ from those in this document. Please refer to our website (www.microchip.com) to obtain the latest documentation available.

Documents are identified with a "DS" number. This number is located on the bottom of each page, in front of the page number. The numbering convention for the DS number is "DSXXXXXA," where "XXXXX" is the document number and "A" is the revision level of the document.

For the most up-to-date information on development tools, see the MPLAB® X IDE online help. Select the Help menu and then Topics, to open a list of available online help files.

Table of Contents

1. Preface

MPLAB® XC32 C/C++ Compiler for PICM MCUs documentation and support information is discussed in this section.

1.1 Conventions Used in This Guide

The following conventions may appear in this documentation:

Table 1-1. Documentation Conventions

Description	Represents	Examples
Arial font:		
Italic characters	Referenced books	*MPLAB® IDE User's Guide*
	Emphasized text	...is the *only* compiler...
Initial caps	A window	the Output window
	A dialog	the Settings dialog
	A menu selection	select Enable Programmer
Quotes	A field name in a window or dialog	"Save project before build"
Underlined, italic text with right angle bracket	A menu path	*File>Save*
Bold characters	A dialog button	Click **OK**
	A tab	Click the **Power** tab
N'Rnnnn	A number in verilog format, where N is the total number of digits, R is the radix and n is a digit.	4'b0010, 2'hF1
Text in angle brackets < >	A key on the keyboard	Press <Enter>, <F1>
Courier New font:		
Plain Courier New	Sample source code	`#define START`
	Filenames	`autoexec.bat`
	File paths	`c:\mcc18\h`
	Keywords	`_asm, _endasm, static`
	Command-line options	`-Opa+, -Opa-`
	Bit values	`0, 1`
	Constants	`0xFF, 'A'`
Italic Courier New	A variable argument	`file.o`, where `file` can be any valid filename
Square brackets [ ]	Optional arguments	`mcc18 [options] file [options]`
Curly brackets and pipe character: {|}	Choice of mutually exclusive arguments; an OR selection	`errorlevel {0|1}`

..........continued		
Description	**Represents**	**Examples**
Ellipses...	Replaces repeated text	`var_name [, var_name...]`

	Represents code supplied by user	`void main (void)`
		`{ ...`
		
		`}`

1.2 Recommended Reading

The MPLAB® XC32 language toolsuite for PIC32 MCUs consists of a C compilation driver (`xc32-gcc`), a C++ compilation driver (`xc32-g++`), an assembler (`xc32-as`), a linker (`xc32-ld`), and an archiver/librarian (`xc32-ar`). This document describes how to use the MPLAB XC32 C/C++ Compiler. Other useful documents are listed below. The following Microchip documents are available and recommended as supplemental reference resources.

Release Notes (Readme Files)

For the latest information on Microchip tools, read the associated Release Notes (HTML files) included with the software.

MPLAB® XC32 Assembler, Linker and Utilities User's Guide (DS50002186)

A guide to using the 32-bit assembler, object linker, object archiver/librarian and various utilities.

32-Bit Language Tools Libraries (DS50001685)

Lists all library functions provided with the MPLAB XC32 C/C++ Compiler with detailed descriptions of their use.

Dinkum Compleat Libraries

The Dinkum Compleat Libraries are organized into a number of headers – files that you include in your program to declare or define library facilities. A link to the Dinkum libraries is available in the MPLAB X IDE application, on the My MPLAB X IDE tab, References & Featured Links section.

PIC32 Configuration Settings

Lists the Configuration Bit settings for the Microchip PIC32 devices supported by the `#pragma config` of the MPLAB XC32 C/C++ Compiler.

Device-Specific Documentation

The Microchip website contains many documents that describe 32-bit device functions and features. Among these are:

- Individual and family data sheets
- Family reference manuals
- Programmer's reference manuals

C Standards Information

American National Standard for Information Systems – *Programming Language – C*. American National Standards Institute (ANSI), 11 West 42nd. Street, New York, New York, 10036.

This standard specifies the form and establishes the interpretation of programs expressed in the programming language C. Its purpose is to promote portability, reliability, maintainability and efficient execution of C language programs on a variety of computing systems.

C++ Standards Information

Stroustrup, Bjarne, *C++ Programming Language: Special Edition,* 3rd Edition. Addison-Wesley Professional; Indianapolis, Indiana, 46240.

ISO/IEC 14882 C++ Standard. The ISO C++ Standard is standardized by ISO (The International Standards Organization) in collaboration with ANSI (The American National Standards Institute), BSI (The British Standards Institute) and DIN (The German national standards organization).

This standard specifies the form and establishes the interpretation of programs expressed in the programming language C++. Its purpose is to promote portability, reliability, maintainability and efficient execution of C++ language programs on a variety of computing systems.

C Reference Manuals

Harbison, Samuel P. and Steele, Guy L., *C A Reference Manual*, Fourth Edition, Prentice-Hall, Englewood Cliffs, N.J. 07632.

Kernighan, Brian W. and Ritchie, Dennis M., *The C Programming Language*, Second Edition. Prentice Hall, Englewood Cliffs, N.J. 07632.

Kochan, Steven G., *Programming In ANSI C*, Revised Edition. Hayden Books, Indianapolis, Indiana 46268.

Plauger, P.J., *The Standard C Library*, Prentice-Hall, Englewood Cliffs, N.J. 07632.

Van Sickle, Ted., *Programming Microcontrollers in C*, First Edition. LLH Technology Publishing, Eagle Rock, Virginia 24085.

GCC Documents

http://gcc.gnu.org/onlinedocs/

http://sourceware.org/binutils/

2.   Compiler Overview

The MPLAB® XC32 C/C++ Compiler for PIC32M MCUs is defined and described in this section.

2.1   Device Description

The MPLAB XC32 C/C++ Compiler fully supports most Microchip PIC32M and MEC 14 devices.

2.2   Compiler Description and Documentation

The MPLAB XC32 C/C++ Compiler is a full-featured, optimizing compiler that translates standard ANSI C programs into 32-bit device assembly language source. The toolchain supports both the PIC32M microcontroller family. The compiler also supports many command-line options and language extensions that allow full access to the 32-bit device hardware capabilities, and affords fine control of the compiler code generator.

The compiler is a port of the GCC compiler from the Free Software Foundation.

The compiler is available for several popular operating systems, including 32- and 64-bit Windows®, Linux® and Mac OS® X.

The compiler can run in Free or PRO operating mode. The PRO operating mode is a licensed mode and requires an activation key and Internet connectivity to enable it. Free mode is available for unlicensed customers. The basic compiler operation, supported devices and available memory are identical across all modes. The modes only differ in the level of optimization employed by the compiler.

2.2.1   Conventions

Throughout this manual, the term "the compiler" is often used. It can refer to either all or some subset of the collection of applications that form the MPLAB XC32 C/C++ Compiler. Often it is not important to know, for example, whether an action is performed by the parser or code generator application and it is sufficient to say it was performed by "the compiler."

It is also reasonable for "the compiler" to refer to the command-line driver (or just driver) as this is the application that is always executed to invoke the compilation process. The driver for the MPLAB XC32 C/C++ Compiler package is called `xc32-gcc`. The driver for the C/ASM projects is also `xc32-gcc`. The driver for C/C++/ASM projects is `xc32-g++`. The drivers and their options are discussed in 6.7 Driver Option Descriptions. Following this view, "compiler options" should be considered command-line driver options, unless otherwise specified in this manual.

Similarly "compilation" refers to all, or some part of, the steps involved in generating source code into an executable binary image.

2.2.2   ANSI C Standards

The compiler is a fully validated compiler that conforms to the ANSI C standard as defined by the ANSI specification (ANSI x3.159-1989) and described in Kernighan and Ritchie's *The C Programming Language* (second edition). The ANSI standard includes extensions to the original C definition that are now standard features of the language. These extensions enhance portability and offer increased capability. In addition, language extensions for PIC32 MCU embedded-control applications are included.

2.2.3   Optimization

The compiler uses a set of sophisticated optimization passes that employ many advanced techniques for generating efficient, compact code from C/C++ source. The optimization passes include high-level optimizations that are applicable to any C/C++ code, as well as PIC32 MCU-specific optimizations that take advantage of the particular features of the device architecture.

For more on optimizations, see 19. Optimizations.

2.2.4   ANSI Standard Library Support

The compiler is distributed with a complete ANSI C standard library. All library functions have been validated and conform to the ANSI C library standard. The library includes functions for string manipulation, dynamic memory

allocation, data conversion, timekeeping and math functions (trigonometric, exponential and hyperbolic). The standard I/O functions for file handling are also included, and, as distributed, they support full access to the host file system using the command-line simulator. The fully functional source code for the low-level file I/O functions is provided in the compiler distribution, and may be used as a starting point for applications that require this capability.

2.2.5 ISO/IEC C++ Standard

The compiler is distributed with the 2003 Standard C++ Library.

Note: Do not specify an MPLAB XC32 system include directory (for example, `/pic32mx/include/`) in your project properties.The xc32-gcc compilation drivers automatically select the XC libc and their respective include-file directory for you. The xc32-g++ compilation drivers automatically select the Dinkumware libc and their respective include-file directory for you. The Dinkum C libraries can only be used with the C++ compiler. Manually adding a system include file path may disrupt this mechanism and cause the incorrect libc include files to be compiled into your project, causing a conflict between the include files and the library. Note that adding a system include path to your project properties has never been a recommended practice.

2.2.6 Compiler Driver

The compiler includes a powerful command-line driver program. Using the driver program, application programs can be compiled, assembled and linked in a single step.

2.2.7 Documentation

This version of the C compiler is supported under MPLAB X IDE v5.05 or higher is required.

2.3 Compiler and Other Development Tools

The compiler works with many other Microchip tools including:

- MPLAB XC32 assembler and linker - see the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186).
- MPLAB X IDE (v5.05 or higher).
- The MPLAB Simulator.
- All Microchip debug tools and programmers.
- Demo boards and starter kits that support 32-bit devices.

3. Common C Interface

The Common C Interface (CCI) is available with all MPLAB XC C compilers and is designed to enhance code portability between these compilers. For example, CCI-conforming code would make it easier to port from a PIC18 MCU using the MPLAB XC8 C compiler to a PIC32 MCU using the MPLAB XC32 C/C++ Compiler.

The CCI assumes that your source code already conforms to the ANSI Standard. If you intend to use the CCI, it is your responsibility to write code that conforms. Legacy projects will need to be migrated to achieve conformance. A compiler option must also be set to ensure that the operation of the compiler is consistent with the interface when the project is built.

3.1 Background - The Desire for Portable Code

All programmers want to write portable source code.

Portability means that the same source code can be compiled and run in a different execution environment than that for which it was written. Rarely can code be one hundred percent portable, but the more tolerant it is to change, the less time and effort it takes to have it running in a new environment.

Embedded engineers typically think of code portability as being across target devices, but this is only part of the situation. The same code could be compiled for the same target but with a different compiler. Differences between those compilers might lead to the code failing at compile time or runtime, so this must be considered as well.

You can only write code for one target device and only use one brand of compiler; but if there is no regulation of the compiler's operation, simply updating your compiler version can change your code's behavior.

Code must be portable across targets, tools, and time to be truly flexible.

Clearly, this portability cannot be achieved by the programmer alone, since the compiler vendors can base their products on different technologies, implement different features and code syntax, or improve the way their product works. Many a great compiler optimization has broken many an unsuspecting project.

Standards for the C language have been developed to ensure that change is managed and code is more portable. The American National Standards Institute (ANSI) publishes standards for many disciplines, including programming languages. The ANSI C Standard is a universally adopted standard for the C programming language.

3.1.1 The ANSI Standard

The ANSI C Standard has to reconcile two opposing goals: freedom for compilers vendors to target new devices and improve code generation, with the known functional operation of source code for programmers. If both goals can be met, source code can be made portable.

The standard is implemented as a set of rules which detail not only the syntax that a conforming C program must follow, but the semantic rules by which that program will be interpreted. Thus, for a compiler to conform to the standard, it must ensure that a conforming C program functions as described by the standard.

The standard describes *implementation*, the set of tools and the runtime environment on which the code will run. If any of these change, for example, you build for, and run on, a different target device, or if you update the version of the compiler you use to build, then you are using a different implementation.

The standard uses the term behavior to mean the external appearance or action of the program. It has nothing to do with how a program is encoded.

Since the standard is trying to achieve goals that could be construed as conflicting, some specifications appear somewhat vague. For example, the standard states that an `int` type must be able to hold at least a 16-bit value, but it does not go as far as saying what the size of an `int` actually is; and the action of right-shifting a signed integer can produce different results on different implementations; yet, these different results are still ANSI C compliant.

If the standard is too strict, device architectures cannot allow the compiler to conform (see following note). But, if it is too weak, programmers would see wildly differing results within different compilers and architectures, and the standard would lose its effectiveness.

Note: For example, the mid-range PIC® microcontrollers do not have a data stack. Because a compiler targeting this device cannot implement recursion, it (strictly speaking) cannot conform to the ANSI C Standard. This example illustrates a situation in which the standard is too strict for mid-range devices and tools.

The standard organizes source code whose behavior is not fully defined into groups that include the following behaviors:

Implementation-defined behavior	This is unspecified behavior in which each implementation documents how the choice is made.
Unspecified behavior	The standard provides two or more possibilities and imposes no further requirements on which possibility is chosen in any particular instance.
Undefined behavior	This is behavior for which the standard imposes no requirements.

Code that strictly conforms to the standard does not produce output that is dependent on any unspecified, undefined, or implementation-defined behavior. The size of an `int`, which was used as an example earlier, falls into the category of behavior that is defined by implementation. That is to say, the size of an `int` is defined by which compiler is being used, how that compiler is being used, and the device that is being targeted.

All the MPLAB XC compilers conform to the ANSI X3.159-1989 Standard for programming languages (with the exception of the MPLAB XC8 compiler's inability to allow recursion, as mentioned in the footnote). This is commonly called the C89 Standard. Some features from the later standard, C99, are also supported.

For freestanding implementations (or for what are typically call embedded applications), the standard allows non-standard extensions to the language, but obviously does not enforce how they are specified or how they work. When working so closely to the device hardware, a programmer needs a means of specifying device setup and interrupts, as well as utilizing the often complex world of small-device memory architectures. This cannot be offered by the standard in a consistent way.

While the ANSI C Standard provides a mutual understanding for programmers and compiler vendors, programmers need to consider the implementation-defined behavior of their tools and the probability that they may need to use extensions to the C language that are non-standard. Both of these circumstances can have an impact on code portability.

3.1.2 The Common C Interface

The Common C Interface (CCI) supplements the ANSI C Standard and makes it easier for programmers to achieve consistent outcomes on all Microchip devices when using any of the MPLAB XC C compilers.

It delivers the following improvements, all designed with portability in mind.

Refinement of the ANSI C Standard	The CCI documents specific behavior for some code in which actions are implementation-defined behavior under the ANSI C Standard. For example, the result of right-shifting a signed integer is fully defined by the CCI. Note that many implementation-defined items that closely couple with device characteristics, such as the size of an `int`, are not defined by the CCI.
Consistent syntax for non-standard extensions	The CCI non-standard extensions are mostly implemented using keywords with a uniform syntax. They replace keywords, macros and attributes that are the native compiler implementation. The interpretation of the keyword can differ across each compiler, and any arguments to the keywords can be device specific.
Coding guidelines	The CCI can indicate advice on how code should be written so that it can be ported to other devices or compilers. While you may choose not to follow the advice, it will not conform to the CCI.

3.2 Using the CCI

The CCI allows enhanced portability by refining implementation-defined behavior and standardizing the syntax for extensions to the language.

The CCI is something you choose to follow and put into effect, thus it is relevant for new projects, although you can choose to modify existing projects so they conform.

For your project to conform to the CCI, you must do the following things.

- **Enable the CCI**
  Select the MPLAB X IDE widget *Use CCI Syntax* in your project, or use the command-line option that is equivalent.
- **Include <xc.h> in every module**
  Some CCI features are only enabled if this header is seen by the compiler.
- **Ensure ANSI compliance**
  Code that does not conform to the ANSI C Standard does not conform to the CCI.
- **Observe refinements to ANSI by the CCI**
  Some ANSI implementation-defined behavior is defined explicitly by the CCI.
- **Use the CCI extensions to the language**
  Use the CCI extensions rather than the native language extensions.

The next sections detail specific items associated with the CCI. These items are segregated into those that refine the standard, those that deal with the ANSI C Standard extensions, and other miscellaneous compiler options and usage. Guidelines are indicated with these items.

If any implementation-defined behavior or any non-standard extension is not discussed in this document, then it is not part of the CCI. For example, GCC case ranges, label addresses and 24-bit `short long` types are not part of the CCI. Programs which use these features do not conform to the CCI. The compiler may issue a warning or error to indicate a non-CCI feature has been used and the CCI is enabled.

3.3 ANSI Standard Refinement

The following topics describe how the CCI refines the implementation-defined behaviors outlined in the ANSI C Standard.

3.3.1 Source File Encoding

Under the CCI, a source file must be written using characters from the 7-bit ASCII set. Lines can be terminated using a *line feed* (`\n`) or *carriage return* (`\r`) that is immediately followed by a *line feed*. Escaped characters can be used in character constants or string literals to represent extended characters that are not in the basic character set.

Example

The following shows a string constant being defined that uses escaped characters.

```
const char myName[] = "Bj\370rk\n";
```

Differences

All compilers have used this character set.

Migration to the CCI

No action required.

3.3.2 The Prototype for `main`

The prototype for the `main()` function is:

```
int main(void);
```

Example

The following shows an example of how `main()` might be defined:

```
int main(void)
{
    while(1)
        process();
}
```

Differences

The 8-bit compilers used a `void` return type for this function.

Migration to the CCI

Each program has one definition for the `main()` function. Confirm the return type for `main()` in all projects previously compiled for 8-bit targets.

3.3.3   Header File Specification

Header file specifications that use directory separators do not conform to the CCI.

Example

The following example shows two conforming include directives.

```
#include <usb_main.h>
```

```
#include "global.h"
```

Differences

Header file specifications that use directory separators have been allowed in previous versions of all compilers. Compatibility problems arose when Windows-style separators "\" were used and the code was compiled under other host operating systems. Under the CCI, no directory separators should be used.

Migration to the CCI

Any `#include` directives that use directory separators in the header file specifications should be changed. Remove all but the header file name in the directive. Add the directory path to the compiler's include search path or MPLAB X IDE equivalent. This will force the compiler to search the directories specified with this option.

For example, the following code:

```
#include <inc/lcd.h>
```

should be changed to:

```
#include <lcd.h>
```

and the path to the `inc` directory added to the compiler's header search path in your MPLAB X IDE project properties, or on the command-line as follows:

```
-Ilcd
```

3.3.4   Include Search Paths

When you include a header file under the CCI, the file should be discoverable in the paths searched by the compiler that are detailed below.

Header files specified in angle bracket delimiters `< >` should be discoverable in the search paths that are specified by `-I` options (or the equivalent MPLAB X IDE option), or in the standard compiler `include` directories. The `-I` options are searched in the order in which they are specified.

Header files specified in quote characters `"  "` should be discoverable in the current working directory or in the same directories that are searched when the header files are specified in angle bracket delimiters (as above). In the case of an MPLAB X project, the current working directory is the directory in which the C source file is located. If unsuccessful, the search paths should be to the same directories searched when the header file is specified in angle bracket delimiters.

Any other options to specify search paths for header files do not conform to the CCI.

Example

If including a header file, as in the following directive:

```
#include "myGlobals.h"
```

the header file should be locatable in the current working directory, or the paths specified by any `-I` options, or the standard compiler directories. A header file being located elsewhere does not conform to the CCI.

Differences

The compiler operation under the CCI is not changed. This is purely a coding guideline.

Migration to the CCI

Remove any option that specifies header file search paths other than the `-I` option (or the equivalent MPLAB X IDE option), and use the `-I` option in place of this. Ensure the header file can be found in the directories specified in this section.

3.3.5 The Number of Significant Initial Characters in an Identifier

At least the first 255 characters in an identifier (internal and external) are significant. This extends upon the requirement of the ANSI C Standard that states a lower number of significant characters are used to identify an object.

Example

The following example shows two poorly named variables, but names which are considered unique under the CCI.

```
int stateOfPortBWhenTheOperatorHasSelectedAutomaticModeAndMotorIsRunningFast;

int stateOfPortBWhenTheOperatorHasSelectedAutomaticModeAndMotorIsRunningSlow;
```

Differences

Former 8-bit compilers used 31 significant characters by default, but an option allowed this to be extended.

The 16- and 32-bit compilers did not impose a limit on the number of significant characters.

Migration to the CCI

No action required. You can take advantage of the less restrictive naming scheme.

3.3.6 Sizes of Types

The sizes of the basic C types, for example `char`, `int` and `long`, are *not* fully defined by the CCI. These types, by design, reflect the size of registers and other architectural features in the target device. They allow the device to efficiently access objects of this type. The ANSI C Standard does, however, indicate minimum requirements for these types, as specified in `<limits.h>`.

If you need fixed-size types in your project, use the types defined in `<stdint.h>`, for example, `uint8_t` or `int16_t`. These types are consistently defined across all XC compilers, even outside of the CCI.

Essentially, the C language offers a choice of two groups of types: those that offer sizes and formats that are tailored to the device you are using; or those that have a fixed size, regardless of the target.

Example

The following example shows the definition of a variable, `native`, whose size will allow efficient access on the target device; and a variable, `fixed`, whose size is clearly indicated and remains fixed, even though it may not allow efficient access on every device.

```
int native;

int16_t fixed;
```

Differences

This is consistent with previous types implemented by the compiler.

Migration to the CCI

If you require a C type that has a fixed size, regardless of the target device, use one of the types defined by `<stdint.h>`.

3.3.7 Plain `char` Types

The type of a plain `char` is `unsigned char`. It is generally recommended that all definitions for the `char` type explicitly state the signedness of the object.

Example

The following example

```
char foobar;
```

defines an `unsigned char` object called `foobar`.

Differences

The 8-bit compilers have always treated plain `char` as an unsigned type.

The 16- and 32-bit compilers used `signed char` as the default plain `char` type. The `-funsigned-char` option on those compilers changed the default type to be `unsigned char`.

Migration to the CCI

Any definition of an object defined as a plain `char` and using the 16- or 32-bit compilers needs review. Any plain `char` that was intended to be a signed quantity should be replaced with an explicit definition, for example.

```
signed char foobar;
```

You can use the `-funsigned-char` option on MPLAB XC16 and XC32 to change the type of plain `char`, but since this option is not supported on MPLAB XC8, the code is not strictly conforming.

3.3.8 Signed Integer Representation

The value of a signed integer is determined by taking the two's complement of the integer.

Example

The following shows a variable, `test`, that is assigned the value -28 decimal.

```
signed char test = 0xE4;
```

Differences

All compilers have represented signed integers in the way described in this section.

Migration to the CCI

No action required.

3.3.9 Integer Conversion

When converting an integer type to a signed integer of insufficient size, the original value is truncated from the most-significant bit to accommodate the target size.

Example

The following shows an assignment of a value that is truncated.

```
signed char destination;
unsigned int source = 0x12FE;
destination = source;
```

Under the CCI, the value of `destination` after the alignment is -2 (that is, the bit pattern 0xFE).

Differences

All compilers have performed integer conversion in an identical fashion to that described in this section.

Migration to the CCI

No action required.

3.3.10 Bitwise Operations on Signed Values

Bitwise operations on signed values act on the two's complement representation, including the sign bit. See also 3.3.11 Right-Shifting Signed Values.

Example

The following shows an example of a negative quantity involved in a bitwise AND operation.

```
signed char output, input = -13;

output = input & 0x7E;
```

Under the CCI, the value of output after the assignment is 0x72.

Differences

All compilers have performed bitwise operations in an identical fashion to that described in this section.

Migration to the CCI

No action required.

3.3.11 Right-Shifting Signed Values

Right-shifting a signed value will involve sign extension. This will preserve the sign of the original value.

Example

The following example shows a negative quantity involved in a right-shift operation.

```
signed char output, input = -13;

output = input >> 3;
```

Under the CCI, the value of output after the assignment is -2 (that is, the bit pattern 0xFE).

Differences

All compilers have performed right-shifting as described in this section.

Migration to the CCI

No action required.

3.3.12 Conversion of Union Member Accessed Using Member with Different Type

If a union defines several members of different types and you use one member identifier to try to access the contents of another (whether any conversion is applied to the result) is implementation-defined behavior in the standard. In the CCI, no conversion is applied and the bytes of the union object are interpreted as an object of the type of the member being accessed, without regard for alignment or other possible invalid conditions.

Example

The following shows an example of a union defining several members.

```
union {
    signed char code;
    unsigned int data;
    float offset;
} foobar;
```

Code that attempts to extract offset by reading data is not guaranteed to read the correct value.

```
float result;
result = foobbar.data;
```

Differences

All compilers have not converted union members accessed via other members.

Migration to the CCI

No action required.

3.3.13 Default Bit-field `int` Type

The type of a bit-field specified as a plain `int` is identical to that of one defined using `unsigned int`. This is quite different from other objects where the types `int`, `signed` and `signed int` are synonymous. It is recommended that the signedness of the bit-field be explicitly stated in all bit-field definitions.

Example

The following shows an example of a structure tag containing bit-fields that are unsigned integers and with the size specified.

```
struct OUTPUTS {
    int direction :1;
    int parity    :3;
    int value     :4;
};
```

Differences

The 8-bit compilers have previously issued a warning if type `int` was used for bit-fields, but would implement the bit-field with an `unsigned int` type.

The 16- and 32-bit compilers have implemented bit-fields defined using `int` as having a `signed int` type, unless the option `-funsigned-bitfields` was specified.

Migration to the CCI

Any code that defines a bit-field with the plain `int` type should be reviewed. If the intention was for these to be signed quantities, then the type of these should be changed to `signed int`. In the following example:

```
struct WAYPT {
    int log         :3;
    int direction   :4;
};
```

the bit-field type should be changed to `signed int`, as in:

```
struct WAYPT {
    signed int log        :3;
    signed int direction :4;
};
```

3.3.14 Bit-Fields Straddling a Storage Unit Boundary

The standard indicates that implementations can determine whether bit-fields cross a storage unit boundary. In the CCI, bit-fields do not straddle a storage unit boundary; a new storage unit is allocated to the structure, and padding bits fill the gap.

Note that the size of a storage unit differs with each compiler, as this is based on the size of the base data type (for example, `int`) from which the bit-field type is derived. On 8-bit compilers this unit is 8-bits in size; for 16-bit compilers, it is 16 bits; and for 32-bit compilers, it is 32 bits in size.

Example

The following shows a structure containing bit-fields being defined.

```
struct {
        unsigned first  : 6;
        unsigned second :6;
} order;
```

Under the CCI and using MPLAB XC8, the storage allocation unit is byte sized. The bit-field `second` is allocated a new storage unit since there are only 2 bits remaining in the first storage unit in which `first` is allocated. The size of this structure, `order`, is 2 bytes.

Differences

This allocation is identical with that used by all previous compilers.

Migration to the CCI

No action required.

3.3.15 The Allocation Order of Bit-Field

The memory ordering of bit-fields into their storage unit is not specified by the ANSI C Standard. In the CCI, the first bit defined is the least significant bit (LSb) of the storage unit in which it is allocated.

Example

The following shows a structure containing bit-fields being defined.

```
struct {
        unsigned lo  : 1;
        unsigned mid :6;
        unsigned hi  : 1;
} foo;
```

The bit-field `lo` is assigned the least significant bit of the storage unit assigned to the structure `foo`. The bit-field `mid` is assigned the next 6 least significant bits, and `hi`, the most significant bit of that same storage unit byte.

Differences

This is identical with the previous operation of all compilers.

Migration to the CCI

No action required.

3.3.16 The NULL Macro

The NULL macro is defined by `<stddef.h>`; however, its definition is implementation-defined behavior. Under the CCI, the definition of `NULL` is the expression `(0)`.

Example

The following shows a pointer being assigned a null pointer constant via the `NULL` macro.

```
int * ip = NULL;
```

The value of `NULL`, `(0)`, is implicitly converted to the destination type.

Differences

The 32-bit compilers previously assigned `NULL` the expression `((void *)0)`.

Migration to the CCI

No action required.

3.3.17 Floating-Point Sizes

Under the CCI, floating-point types must not be smaller than 32 bits in size.

Example

The following shows the definition for `outY`, which is at least 32-bit in size.

```
float outY;
```

Differences

The 8-bit compilers have allowed the use of 24-bit float and double types.

Migration to the CCI

When using 8-bit compilers, the `float` and `double` type will automatically be made 32 bits in size once the CCI mode is enabled. Review any source code that may have assumed a `float` or `double` type and may have been 24 bits in size.

No migration is required for other compilers.

3.4   ANSI Standard Extensions

The following topics describe how the CCI provides device-specific extensions to the standard.

3.4.1   Generic Header File

A single header file `<xc.h>` must be used to declare all compiler- and device-specific types and SFRs. You *must* include this file into every module to conform with the CCI. Some CCI definitions depend on this header being seen.

Example

The following shows this header file being included, thus allowing conformance with the CCI, as well as allowing access to SFRs.

```
#include <xc.h>
```

Differences

Some 8-bit compilers used `<htc.h>` as the equivalent header. Previous versions of the 16- and 32-bit compilers used a variety of headers to do the same job.

Migration to the CCI

Change:

```
#include <htc.h>
```

previously used in 8-bit compiler code, or family-specific header files, for example, from:

```
#include <p32xxxx.h>
#include <p30fxxxx.h>
#include <p33Fxxxx.h>
#include <p24Fxxxx.h>
#include "p30f6014.h"
```

to:

```
#include <xc.h>
```

3.4.2   Absolute Addressing

Variables and functions can be placed at an absolute address by using the `__at()` construct. Stack-based (`auto` and parameter) variables cannot use the `__at()` specifier.

Example

The following shows two variables and a function being made absolute.

```
int scanMode __at(0x200);
const char keys[] __at(123) = { 'r', 's', 'u', 'd'};
}
```

Differences

The 8-bit compilers have used an `@` symbol to specify an absolute address.

The 16- and 32-bit compilers have used the `address` attribute to specify an object's address.

Migration to the CCI

Avoid making objects and functions absolute if possible.

In MPLAB XC8, change absolute object definitions, for example, from:

```
int scanMode @ 0x200;
```

to:

```
int scanMode __at(0x200);
```

In MPLAB XC16 and XC32, change code, for example, from:

```
int scanMode __attribute__((address(0x200)));
```

to:

```
int scanMode __at(0x200);
```

Caveats

If the `__at()` and `__section()` specifiers are both applied to an object when using MPLAB XC8, the `__section()` specifier is currently ignored.

3.4.3 Far Objects and Functions

The `__far` qualifier can be used to indicate that variables or functions are located in 'far memory'. Exactly what constitutes far memory is dependent on the target device, but it is typically memory that requires more complex code to access. Expressions involving far-qualified objects usually generate slower and larger code.

Use the native keywords discussed in the Differences section to look up information on the semantics of this qualifier.

Some devices may not have such memory implemented, in which case, use of this qualifier is ignored. Stack-based (`auto` and parameter) variables cannot use the `__far` specifier.

Example

The following shows a variable and function qualified using `__far`.

```
__far int serialNo;

__far int ext_getCond(int selector);
```

Differences

The 8-bit compilers have used the qualifier `far` to indicate this meaning. Functions could not be qualified as `far`.

The 16-bit compilers have used the `far` attribute with both variables and functions.

The 32-bit compilers have used the `far` attribute with functions, only.

Migration to the CCI

For 8-bit compilers, change any occurrence of the `far` qualifier, for example, from:

```
far char template[20];
```

to:

```
__far, that is, __far char template[20];
```

In the 16- and 32-bit compilers, change any occurrence of the `far` attribute, for example, from:

```
void bar(void) __attribute__ ((far));

int tblIdx __attribute__ ((far));
```

to:

```
void __far bar(void);

int __far tblIdx;
```

Caveats

None.

3.4.4 Near Objects

The `__near` qualifier can be used to indicate that variables or functions are located in 'near memory'. Exactly what constitutes near memory is dependent on the target device, but it is typically memory that can be accessed with less complex code. Expressions involving near-qualified objects generally are faster and result in smaller code.

Use the native keywords discussed in the Differences section to look up information on the semantics of this qualifier.

Some devices may not have such memory implemented, in which case, use of this qualifier is ignored. Stack-based (`auto` and parameter) variables cannot use the `__near` specifier.

Example

The following shows a variable and function qualified using __near.

```
__near int serialNo;

__near int ext_getCond(int selector);
```

Differences

The 8-bit compilers have used the qualifier `near` to indicate this meaning. Functions could not be qualified as `near`.

The 16-bit compilers have used the `near` attribute with both variables and functions.

The 32-bit compilers have used the `near` attribute for functions, only.

Migration to the CCI

For 8-bit compilers, change any occurrence of the `near` qualifier to __near, for example, from:

```
near char template[20];
```

to:

```
__near char template[20];
```

In 16- and 32-bit compilers, change any occurrence of the `near` attribute to __near, for example, from:

```
void bar(void) __attribute__ ((near));

int tblIdx __attribute__ ((near));
```

to

```
void __near bar(void);

int __near tblIdx;
```

Caveats

None.

3.4.5 Persistent Objects

The __persistent qualifier can be used to indicate that variables should not be cleared by the runtime startup code.

Use the native keywords discussed in the Differences section to look up information on the semantics of this qualifier.

Example

The following shows a variable qualified using __persistent.

```
__persistent int serialNo;
```

Differences

The 8-bit compilers have used the qualifier, `persistent`, to indicate this meaning.

The 16- and 32-bit compilers have used the `persistent` attribute with variables to indicate they were not to be cleared.

Migration to the CCI

With 8-bit compilers, change any occurrence of the `persistent` qualifier to __persistent, for example, from:

```
persistent char template[20];
```

to:

```
__persistent char template[20];
```

For the 16- and 32-bit compilers, change any occurrence of the `persistent` attribute to __persistent, for example, from:

```
int tblIdx __attribute__ ((persistent));
```

to

```
int __persistent tblIdx;
```

Caveats

None.

3.4.6   X and Y Data Objects

The `__xdata` and `__ydata` qualifiers can be used to indicate that variables are located in special memory regions. Exactly what constitutes X and Y memory is dependent on the target device, but it is typically memory that can be accessed independently on separate buses. Such memory is often required for some DSP instructions.

Use the native keywords discussed in the Differences section to look up information on the semantics of these qualifiers.

Some devices may not have such memory implemented; in which case, use of these qualifiers is ignored.

Example

The following shows a variable qualified using `__xdata`, as well as another variable qualified with `__ydata`.

```
__xdata char data[16];

__ydata char coeffs[4];
```

Differences

The 16-bit compilers have used the `xmemory` and `ymemory` space attribute with variables.

Equivalent specifiers have never been defined for any other compiler.

Migration to the CCI

For 16-bit compilers, change any occurrence of the `space` attributes `xmemory` or `ymemory` to `__xdata`, or `__ydata` respectively, for example, from:

```
char __attribute__((space(xmemory)))template[20];
```

to:

```
__xdata char template[20];
```

Caveats

None.

3.4.7   Banked Data Objects

The `__bank(num)` qualifier can be used to indicate that variables are located in a particular data memory bank. The number, `num`, represents the bank number. Exactly what constitutes banked memory is dependent on the target device, but it is typically a subdivision of data memory to allow for assembly instructions with a limited address width field.

Use the native keywords discussed in the Differences section to look up information on the semantics of these qualifiers.

Some devices may not have banked data memory implemented, in which case, use of this qualifier is ignored. The number of data banks implemented will vary from one device to another.

Example

The following shows a variable qualified using `__bank()`.

```
__bank(0) char start;

__bank(5) char stop;
```

Differences

The 8-bit compilers have used the four qualifiers `bank0`, `bank1`, `bank2` and `bank3` to indicate the same, albeit more limited, memory placement.

Equivalent specifiers have never been defined for any other compiler.

Migration to the CCI

For 8-bit compilers, change any occurrence of the `bankx` qualifiers to `__bank()`, for example, from:

```
bank2 int logEntry;
```

to:

```
__bank(2) int logEntry;
```

Caveats

This feature is not yet implemented in MPLAB XC8.

3.4.8   Alignment of Objects

The `__align(alignment)` specifier can be used to indicate that variables must be aligned on a memory address that is a multiple of the alignment specified. The alignment term must be a power of 2. Positive values request that the object's start address be aligned.

Use the native keywords discussed in the Differences section to look up information on the semantics of this specifier.

Example

The following shows variables qualified using `__align()` to ensure they start on an address that is a multiple of 2.

```
__align(2) char coeffs[6];
```

Differences

An alignment feature has never been implemented on 8-bit compilers.

The 16- and 32-bit compilers used the `aligned` attribute with variables.

Migration to the CCI

For 16- and 32-bit compilers, change any occurrence of the `aligned` attribute to `__align`, for example, from:

```
char __attribute__((aligned(4)))mode;
```

to:

```
__align(4) char mode;
```

Caveats

This feature is not yet implemented on MPLAB XC8.

3.4.9   EEPROM Objects

The `__eeprom` qualifier can be used to indicate that variables should be positioned in EEPROM.

Use the native keywords discussed in the Differences section to look up information on the semantics of this qualifier.

Some devices may not implement EEPROM. Use of this qualifier for such devices generates a warning. Stack-based (`auto` and parameter) variables cannot use the `__eeprom` specifier.

Example

The following shows a variable qualified using `__eeprom`.

```
__eeprom int serialNos[4];
```

Differences

The 8-bit compilers have used the qualifier, `eeprom`, to indicate this meaning for some devices.

The 16-bit compilers have used the `space` attribute to allocate variables to the memory space used for EEPROM.

Migration to the CCI

For 8-bit compilers, change any occurrence of the `eeprom` qualifier to `__eeprom`, for example, from:

```
eeprom char title[20];
```

to:

```
__eeprom char title[20];
```

For 16-bit compilers, change any occurrence of the `eedata space` attribute to `__eeprom`, for example, from:

```
int mainSw __attribute__ ((space(eedata)));
```

to:

```
int __eeprom mainSw;
```

Caveats

MPLAB XC8 does not implement the `__eeprom` qualifiers for any PIC18 devices; this qualifier works as expected for other 8-bit devices.

3.4.10 Interrupt Functions

The `__interrupt(type)` specifier can be used to indicate that a function is to act as an interrupt service routine. The `type` is a comma-separated list of keywords that indicate information about the interrupt function.

The current interrupt types are:

<empty>	Implement the default interrupt function.
low_priority	The interrupt function corresponds to the low priority interrupt source.(MPLAB XC8 - PIC18 only)
high_priority	The interrupt function corresponds to the high priority interrupt source.(MPLAB XC8)
save(symbol-list)	Save on entry and restore on exit the listed symbols. (XC16)
irq(irqid)	Specify the interrupt vector associated with this interrupt. (XC16)
altirq(altirqid)	Specify the alternate interrupt vector associated with this interrupt. (XC16)
preprologue(asm)	Specify assembly code to be executed before any compiler-generated interrupt code. (XC16)
shadow	Allow the ISR to utilize the shadow registers for context switching (XC16)
auto_psv	The ISR will set the PSVPAG register and restore it on exit.(XC16)
no_auto_psv	The ISR will not set the PSVPAG register. (XC16)

Use the native keywords discussed in the Differences section to look up information on the semantics of this specifier.

Some devices may not implement interrupts. Use of this qualifier for such devices generates a warning. If the argument to the `__interrupt` specifier does not make sense for the target device, a warning or error is issued by the compiler.

Example

The following shows a function qualified using `__interrupt`.

```
__interrupt(low_priority) void getData(void) {
    if (TMR0IE && TMR0IF) {
        TMR0IF=0;
        ++tick_count;
    }
}
```

Differences

The 8-bit compilers have used the `interrupt` and `low_priority` qualifiers to indicate this meaning for some devices. Interrupt routines were, by default, high priority.

The 16- and 32-bit compilers have used the `interrupt` attribute to define interrupt functions.

Migration to the CCI

For 8-bit compilers, change any occurrence of the `interrupt` qualifier, for example, from:

```
void interrupt myIsr(void)
```

```
void interrupt low_priority myLoIsr(void)
```

to the following, respectively

```
void __interrupt(high_priority) myIsr(void)
```

```
void __interrupt(low_priority) myLoIsr(void)
```

For 16-bit compilers, change any occurrence of the `interrupt` attribute, for example, from:

```
void _attribute_((interrupt(auto_psv,irq(52)))) _T1Interrupt(void);
```

to:

```
void __interrupt(auto_psv,irq(52))) _T1Interrupt(void);
```

For 32-bit compilers, the `__interrupt()` keyword takes two parameters, the vector number and the (optional) IPL value. Change code that uses the `interrupt` attribute, similar to these examples:

```
void __attribute__((vector(0), interrupt(IPL7AUTO), nomips16)) myisr0_7A(void) {}

void __attribute__((vector(1), interrupt(IPL6SRS), nomips16)) myisr1_6SRS(void) {}

/* Determine IPL and context-saving mode at runtime */
void __attribute__((vector(2), interrupt(), nomips16)) myisr2_RUNTIME(void) {}
```

to

```
void __interrupt(0,IPL7AUTO) myisr0_7A(void) {}

void __interrupt(1,IPL6SRS) myisr1_6SRS(void) {}

/* Determine IPL and context-saving mode at runtime */
void __interrupt(2) myisr2_RUNTIME(void) {}
```

Caveats

None.

3.4.11 Packing Objects

The `__pack` specifier can be used to indicate that structures should not use memory gaps to align structure members, or that individual structure members should not be aligned.

Use the native keywords discussed in the Differences section to look up information on the semantics of this specifier.

Some compilers cannot pad structures with alignment gaps for some devices, and use of this specifier for such devices is ignored.

Example

The following shows a structure qualified using `__pack`, as well as a structure where one member has been explicitly packed.

```
__pack struct DATAPOINT {
    unsigned char type;
    int value;
} x-point;
struct LINETYPE {
    unsigned char type;
    __pack int start;
    long total;
} line;
```

Differences

The `__pack` specifier is a new CCI specifier that is with MPLAB XC8. This specifier has no apparent effect since the device memory is byte addressable for all data objects.

The 16- and 32-bit compilers have used the `packed` attribute to indicate that a structure member was not aligned with a memory gap.

Migration to the CCI

No migration is required for MPLAB XC8.

For 16- and 32-bit compilers, change any occurrence of the `packed` attribute, for example, from:

```
struct DOT
{
    char a;
    int x[2] __attribute__ ((packed));
};
```

to

```
struct DOT
{
    char a;
    __pack int x[2];
};
```

Alternatively, you can pack the entire structure, if required.

Caveats

None.

3.4.12 Indicating Antiquated Objects

The `__deprecate` specifier can be used to indicate that an object has limited longevity and should not be used in new designs. It is commonly used by the compiler vendor to indicate that compiler extensions or features can become obsolete, or that better features have been developed and should be used in preference.

Use the native keywords discussed in the Differences section to look up information on the semantics of this specifier.

Example

The following shows a function that uses the `__deprecate` keyword.

```
void __deprecate getValue(int mode)
{
//...
}
```

Differences

No deprecate feature was implemented on 8-bit compilers.

The 16- and 32-bit compilers have used the `deprecated` attribute (note the different spelling) to indicate that objects should be avoided, if possible.

Migration to the CCI

For 16- and 32-bit compilers, change any occurrence of the `deprecated` attribute to `__deprecate`, for example, from:

```
int __attribute__(deprecated) intMask;
```

to:

```
int __deprecate intMask;
```

Caveats

None.

3.4.13 Assigning Objects to Sections

The __section() specifier can be used to indicate that an object should be located in the named section (or psect, using MPLAB XC8 terminology). This is typically used when the object has special and unique linking requirements that cannot be addressed by existing compiler features.

Use the native keywords discussed in the Differences section to look up information on the semantics of this specifier.

Example

The following shows a variable which uses the __section keyword.

```
int __section("comSec") commonFlag;
```

Differences

The 8-bit compilers have previously used the #pragma psect directive to redirect objects to a new section, or psect. However, the __section() specifier is the preferred method to perform this task, even if you are not using the CCI.

The 16- and 32-bit compilers have used the section attribute to indicate a different destination section name. The __section() specifier works in a similar way to the attribute.

Migration to the CCI

For MPLAB XC8, change any occurrence of the #pragma psect directive, such as

```
#pragma psect text%%u=myText
int getMode(int target) {
//...
}
```

to the __section() specifier, as in

```
int __section ("myText") getMode(int target) {
//...
}
```

For 16- and 32-bit compilers, change any occurrence of the section attribute, for example, from:

```
int __attribute__((section("myVars"))) intMask;
```

to:

```
int __section("myVars") intMask;
```

Caveats

None.

3.4.14 Specifying Configuration Bits

The #pragma config directive can be used to program the Configuration bits for a device. The pragma has the form:

```
#pragma config setting = state|value
```

where setting is a configuration setting descriptor (for example, WDT), state is a descriptive value (for example, ON) and value is a numerical value.

Use the native keywords discussed in the Differences section to look up information on the semantics of this directive.

Example

The following shows Configuration bits being specified using this pragma.

```
#pragma config WDT=ON, WDTPS = 0x1A
```

Differences

The 8-bit compilers have used the `__CONFIG()` macro for some targets that did not already have support for the `#pragma config`.

The 16-bit compilers have used a number of macros to specify the configuration settings.

The 32-bit compilers supported the use of `#pragma config`.

Migration to the CCI

For the 8-bit compilers, change any occurrence of the `__CONFIG()` macro, for example,

`__CONFIG(WDTEN & XT & DPROT)`

to the `#pragma config` directive, for example,

`#pragma config WDTE=ON, FOSC=XT, CPD=ON`

No migration is required if the `#pragma config` was already used.

For the 16-bit compilers, change any occurrence of the `_FOSC()` or `_FBORPOR()` macros attribute, for example, from:

`_FOSC(CSW_FSCM_ON & EC_PLL16);`

to:

`#pragma config FCKSMEM = CSW_ON_FSCM_ON, FPR = ECIO_PLL16`

No migration is required for 32-bit code.

Caveats

None.

3.4.15 Manifest Macros

The CCI defines the general form for macros that manifest the compiler and target device characteristics. These macros can be used to conditionally compile alternate source code based on the compiler or the target device.

The macros and macro families are details in the following table.

Table 3-1. Manifest Macros Defined by the CCI

Name	Meaning if defined	Example
`__XC__`	Compiled with an MPLAB XC compiler	`__XC__`
`__CCI__`	Compiler is CCI compliant and CCI enforcement is enabled	`__CCI__`
`__XC##__`	The specific XC compiler used (`##` can be `8`, `16` or `32`)	`__XC#__` where `#` is 8, 16 or 32
`__DEVICEFAMILY__`	The family of the selected target device	`__dsPIC30F__`
`__DEVICENAME__`	The selected target device name	`__18F452__`

Example

The following shows code that is conditionally compiled dependent on the device having EEPROM memory.

```
#ifdef __XC16__
void __interrupt(__auto_psv__) myIsr(void)
#else
void __interrupt(low_priority) myIsr(void)
#endif
```

Differences

Some of these CCI macros are new (for example, `__CCI__`), and others have different names to previous symbols with identical meaning (for example, `__18F452` is now `__18F452__`).

Migration to the CCI

Any code that uses compiler-defined macros needs review. Old macros continue to work as expected, but they are not compliant with the CCI.

Caveats

None.

3.4.16   In-Line Assembly

The `asm()` statement can be used to insert assembly code in-line with C code. The argument is a C string literal that represents a single assembly instruction. Obviously, the instructions contained in the argument are device specific.

Use the native keywords discussed in the Differences section to look up information on the semantics of this statement.

Example

The following shows a `MOVLW` instruction being inserted in-line.

```
asm("MOVLW _foobar");
```

Differences

The 8-bit compilers have used either the `asm()` or `#asm ... #endasm` constructs to insert in-line assembly code.

This is the same syntax used by the 16- and 32-bit compilers.

Migration to the CCI

For 8-bit compilers, change any instance of `#asm ... #endasm` so that each instruction in this `#asm` block is placed in its own `asm()` statement, for example, from:

```
#asm
    MOVLW       20
    MOVWF _i
    CLRF        Ii+1
#endasm
```

to:

```
asm("MOVLW       20");
asm("MOVWF _i");
asm("CLRF        Ii+1");
```

No migration is required for the 16- or 32-bit compilers.

Caveats

None.

3.5   Compiler Features

The following item details the compiler options used to control the CCI.

3.5.1   Enabling the CCI

It is assumed that you are using the MPLAB X IDE to build projects that use the CCI. The widget in the MPLAB X IDE Project Properties to enable CCI conformance is _Use CCI Syntax_ in the Compiler category.

If you are not using this IDE, then the command-line options are `--EXT=cci` for MPLAB XC8 or `-mcci` for MPLAB XC16 and XC32.

Differences

This option has never been implemented previously.

Migration to the CCI

Enable the option.

Caveats

None.

4.   How To's

This section contains help and references for situations that are frequently encountered when building projects for Microchip 32-bit devices. Click the links at the beginning of each section to assist finding the topic relevant to your question. Some topics are indexed in multiple sections.

Start here:

- 4.1 Installing and Activating the Compiler
- 4.2 Invoking the Compiler
- 4.3 Writing Source Code
- 4.4 Getting My Application To Do What I Want
- 4.5 Understanding the Compilation Process
- 4.6 Fixing Code That Does Not Work

4.1   Installing and Activating the Compiler

This section details questions that might arise when installing or activating the compiler.

- 4.1.1 How Do I Install and Activate My Compiler?
- 4.1.2 How Can I Tell if the Compiler has Activated Successfully?
- 4.1.3 Can I Install More Than One Version of the Same Compiler?

4.1.1   How Do I Install and Activate My Compiler?

Installation and activation of the license are performed simultaneously by the XC compiler installer. The guide *Installing and Licensing MPLAB XC C Compilers* (DS50002059) is available on https://www.microchip.com/compilers, under the **Documentation** tab. It provides details on single-user and network licenses, as well as how to activate a compiler for evaluation purposes.

The default install path is different on these operating systems:

Windows:`"C:\Program Files\Microchip\xc32\vx.xx\bin\xclm" -status`

macOS: `macOS: "/Applications/microchip/xc32/vx.xx/bin/xclm" -status`

Linux: `Linux: "/opt/microchip/xc32/vx.xx/bin/xclm" -status`

Where `vx.xx` is is the version of the compiler.

4.1.2   How Can I Tell if the Compiler has Activated Successfully?

If you think the compiler may not have installed correctly or is not working, it is best to verify its operation outside of MPLAB X IDE to isolate possible problems. Try running the compiler from the command line to check for correct operation. You do not actually have to compile code.

From your terminal or command-line prompt, run the license manager xclm with the option `-status`. This option instructs the license manager to print all MPLAB XC licenses installed on your system and exit. So, for example, depending on your operating system, type the following line, replacing the path information with a path that is relevant to your installation.

Windows:`"C:\Program Files\Microchip\xc32\v1.00\bin\xclm" -status`

macOS: `macOS: "/Applications/microchip/xc32/v1.00/bin/xclm" -status`

Linux: `Linux: "/opt/microchip/xc32/v1.00/bin/xclm" -status`

The license manager should run, print all of the MPLAB XC compiler license available on your machine, and quit. Confirm that the your license is listed as activated (for example, Product:swxc32-pro) Note: if it is not activated properly, the compiler will continue to operate, but only in the Free mode. If an error is displayed, or the compiler indicates Free mode, then activation was not successful.

4.1.3 Can I Install More Than One Version of the Same Compiler?

Yes, the compilers and installation process has been designed to allow you to have more than one version of the same compiler installed. For MPLAB X IDE, you can easily swap between version by changing options in the IDE (see 4.2.4 How Can I Select Which Compiler I Want to Build With?.)

Compilers should be installed into a directory whose name is related to the compiler version. This is reflected in the default directory specified by the installer. For example, the MPLAB XC32 compilers v1.00 and v1.10 would typically be placed in separate directories.

```
C:\Program Files\Microchip\xc32\v1.00\
C:\Program Files\Microchip\xc32\v1.10\
```

4.2 Invoking the Compiler

This section discusses how the compiler is run, both on the command-line and from the IDE. It includes information about how to get the compiler to do what you want in terms of options and the build process itself.

4.2.1 How Do I Compile from Within MPLAB X IDE?

4.2.2 How Do I Compile on the Command-Line?

4.2.3 How Do I Compile Using a Make Utility?

4.2.4 How Can I Select Which Compiler I Want to Build With?

4.2.5 How Can I Change the Compiler's Operating Mode?

4.2.6 How Do I Build Libraries?

4.2.7 How Do I Know What Compiler Options Are Available and What They Do?

4.2.8 How Do I Know What the Build Options in MPLAB X IDE Do?

4.2.9 What is Different About an MPLAB X IDE Debug Build?

See also 4.3.3.4 How Do I Stop the Compiler From Using Certain Memory Locations?

See also 4.4.3 What Do I Need to Do When Compiling to Use a Debugger?

See also 4.5.16 How Do I Use Library Files in My Project?

See also 4.5.18 What Optimizations Are Employed by the Compiler?

4.2.1 How Do I Compile from Within MPLAB X IDE?

See the following documentation for information on how to set up a project:

5.4 Project Setup - MPLAB X IDE

4.2.2 How Do I Compile on the Command-Line?

The compiler driver is called `xc32-gcc` for all 32-bit devices; for example, in Windows, it is named `xc32-gcc.exe`. This application should be invoked for all aspects of compilation. It is located in the `bin` directory of the compiler distribution. Avoid running the individual compiler applications (such as the assembler or linker) explicitly. You can compile and link in the one command, even if your project is spread among multiple source files.

The driver is introduced in 6.1 Invoking The Compiler. See 4.2.4 How Can I Select Which Compiler I Want to Build With? to ensure you are running the correct driver if you have more than one installed. The command-line options to the driver are detailed in 6.7 Driver Option Descriptions. The files that can be passed to the driver are listed and described in 6.1.3 Input File Types.

4.2.3 How Do I Compile Using a Make Utility?

When compiling using a make utility (such as `make`), the compilation is usually performed as a two-step process: first generating the intermediate files, and then the final compilation and link step to produce one binary output. This is described in 6.2.2 Multi-Step C Compilation.

4.2.4 How Can I Select Which Compiler I Want to Build With?

The compilation and installation process has been designed to allow you to have more than one compiler installed at the same time For MPLAB X IDE, you can create a project and then build this project with different compilers by simply changing a setting in the project properties.

In MPLAB X IDE, you select which compiler to use when building a project by opening the Project Properties window (`File>Project Properties`) and selecting the Configuration category (`Conf: [default]`). A list of MPLAB XC32 compiler versions is shown in the Compiler Toolchain, on the far right. Select the MPLAB XC32 compiler you require.

Once selected, the controls for that compiler are then shown by selecting the XC32 global options, XC32 Compiler and XC32 Linker categories. These reveal a pane of options on the right; each category has several panes which can be selected from a pull-down menu that is near the top of the pane.

4.2.5 How Can I Change the Compiler's Operating Mode?

The compiler's operating mode (Free, Evaluation, or PRO) is based on its level of optimizations (see 19. Optimizations) which can be specified as a command line option (see 6.7.7 Options for Controlling Optimization.) If you are building under MPLAB X IDE, go to the Project Properties window, click on the compiler name (xc32-gcc for C language projects or xc32-g++ for C++ language projects), and select the Optimization option category to set optimization levels - see 5.4.3 xc32-gcc (32-bit C Compiler).

When building your project, the compiler will emit a warning message if you have selected an option that is not available for your licensed operating mode. The compiler will continue compilation with the option disabled.

4.2.6 How Do I Build Libraries?

When you have functions and data that are commonly used in applications, you can make all the C source and header files available so other developers can copy these into their projects. Alternatively, you can build these modules into object files and package them into library archives, which, along with the accompanying header files, can then be built into an application.

Libraries can be more convenient because there are fewer files to manage. However, libraries do need to be maintained. MPLAB XC32 uses *.a library archives. Be sure to rebuild your library objects when you move your project to a new release of the compiler toolchain.

Using the compiler driver, libraries can begin to be built by listing all the files that are to be included into the library on the command line. None of these files should contain a `main()` function, nor settings for configuration bits or any other such data.

For information on how to create your own libraries, see 6.4.2.2 User-Defined Libraries.

4.2.7 How Do I Know What Compiler Options Are Available and What They Do?

A list of all compiler options can be obtained by using the `--help` option on the command line. Alternatively, all options are listed in 6.7 Driver Option Descriptions in this user's guide. If you are compiling in MPLAB X IDE, see 5.4 Project Setup.

4.2.8 How Do I Know What the Build Options in MPLAB X IDE Do?

Most of the widgets and controls in the MPLAB X IDE Project Properties window, XC32 options, map directly to one command-line driver option or suboption. See 5.4 Project Setup for a list of options and any corresponding command-line options.

4.2.9 What is Different About an MPLAB X IDE Debug Build?

The main difference between a command-line debug build and an MPLAB X IDE debug build is the setting of a preprocessor macro called `__DEBUG` to be 1 when a debug is selected. This macro is not defined if it is not a debug build.

You may make code in your source conditional on this macro using `#ifdef` directives, etc (see 6.7.8 Options for Controlling the Preprocessor) so that you can have your program behave differently when you are still in a development cycle. Some compiler errors are easier to track down after performing a debug build.

In MPLAB X IDE, memory will be reserved for your debugger only when you perform a debug build. See 4.4.3 What Do I Need to Do When Compiling to Use a Debugger?.

4.3 Writing Source Code

This section addresses issues pertaining to the source code you write. It has been subdivided into sections listed below.

4.3.1 C Language Specifics

This section discusses commonly asked source code issues that directly relate to the C language itself.

4.3.1.1 When Should I Cast Expressions?

Expressions can be explicitly cast using the cast operator -- a type in round brackets, for example, (`int`). In all cases, conversion of one type to another must be done with caution and only when absolutely necessary.

Consider the example:

```
unsigned long l;
unsigned short s;

s = l;
```

Here, a `long` type is being assigned to an `int` type and the assignment will truncate the value in `l`. The compiler will automatically perform a type conversion from the type of the expression on the right of the assignment operator (`long`) to the type of the value on the left of the operator (`short`).This is called an implicit type conversion. The compiler typically produces a warning concerning the potential loss of data by the truncation.

A cast to type `short` is not required and should not be used in the above example if a `long` to `short` conversion was intended. The compiler knows the types of both operands and performs the conversion accordingly. If you did use a cast, there is the potential for mistakes if the code is later changed. For example, if you had:

```
s = (short)l;
```

the code works the same way; but if in the future, the type of `s` is changed to a `long`, for example, then you must remember to adjust the cast, or remove it, otherwise the contents of `l` will continue to be truncated by the assignment, which may not be correct. Most importantly, the warning issued by the compiler will not be produced if the cast is in place.

Only use a cast in situations where the types used by the compiler are not the types that you require. For example, consider the result of a division assigned to a floating-point variable:

```
int i, j;
float fl;

fl = i/j;
```

In this case, integer division is performed, then the rounded integer result is converted to a `float` format. So if `i` contained 7 and `j` contained 2, the division yields 3 and this is implicitly converted to a `float` type (3.0) and then assigned to `fl`. If you wanted the division to be performed in a `float` format, then a cast is necessary:

```
fl = (float)i/j;
```

(Casting either `i` or `j` forces the compiler to encode a floating-point division.) The result assigned to `fl` now is 3.5.

An explicit cast can suppress warnings that might otherwise have been produced. This can also be the source of many problems. The more warnings the compiler produces, the better chance you have of finding potential bugs in your code.

4.3.1.2 Can Implicit Type Conversions Change The Expected Results Of My Expressions?

Yes! The compiler will always use integral promotion and there is no way to disable this (see 12.1 Integral Promotion). In addition, the types of operands to binary operators are usually changed so that they have a common type, as specified by the C Standard. Changing the type of an operand can change the value of the final expression, so it is very important that you understand the type C Standard conversion rules that apply when dealing with binary operators. You can manually change the type of an operand by casting; see 4.3.1.1 When Should I Cast Expressions?

4.3.1.3 How Do I Enter Non-English Characters Into My Program?

The ANSI standard and MPLAB XC C do not support extended characters set in character and string literals in the source character set. See 9.8 Constant Types and Formats to see how these characters can be entered using escape sequences.

4.3.1.4 How Can I Use A Variable Defined In Another Source File?

Provided the variable defined in the other source file is not `static` (see 10.2.2 Static Variables) or `auto` (see 10.3 Auto Variable Allocation and Access), adding a declaration for that variable into the current file will allow you to access it. A declaration consists of the keyword `extern` in addition to the type and the name of the variable, as specified in its definition, for example,

```
extern int systemStatus;
```

This is part of the C language. Your favorite C textbook will give you more information.

The position of the declaration in the current file determines the scope of the variable. That is, if you place the declaration inside a function, it will limit the scope of the variable to that function. If you place it outside of a function, it allows access to the variable in all functions for the remainder of the current file.

Often, declarations are placed in header files and then they are `#include`d into the C source code (see 20.2 Preprocessor Directives).

4.3.1.5 How Do I Port My Code To Different Device Architectures?

Porting code to different devices within an architectural family requires a minimum update to application code. However, porting between architectural families can require significant rewrite.

In an attempt to reduce the work to port between architectures, a Common C Interface, or CCI, has been developed. If you use these coding styles, your code will more easily migrate upward. For more on CCI, see 3. Common C Interface.

4.3.2 Device-Specific Features

This section discusses the code that needs to be written to set up or control a feature that is specific to Microchip PIC devices.

- 4.3.2.1 How Do I Set the Configuration Bits?
- 4.3.2.2 How Do I Determine the Cause of Reset?
- 4.3.2.3 How Do I Access SFRS?
- 4.3.3.4 How Do I Stop the Compiler From Using Certain Memory Locations?

See also the following linked information in other sections.

- 4.4.3 What Do I Need to Do When Compiling to Use a Debugger?

4.3.2.1 How Do I Set the Configuration Bits?

These should be set in your code using either a macro or pragma. Earlier versions of MPLAB IDE allowed you to set these bits in a dialog, but MPLAB X IDE requires that they be specified in your source code. Config bits are set in source code using the config pragma. See 8.4 Configuration Bit Access, for more information on the config pragma.

4.3.2.2 How Do I Determine the Cause of Reset?

The bits in the Reset Control (RCON) Register allow you to determine the cause of a Reset. See the data sheet for your target device for a description of the RCON register.

4.3.2.3 How Do I Access SFRS?

The compiler ships with header files, see 8.5 ID Locations, that define variables which are mapped over the top of memory-mapped SFRs. Since these are C variables, they can be used like any other C variable and no new syntax is required to access these registers.

Bits within SFRs can also be accessed. Bit-fields are available in structures which map over the SFR as a whole. See 9.5.2 Bit Fields in Structures.

The name assigned to the variable is usually the same as the name specified in the device data sheet. See 4.3.2.4 How Do I Find The Names Used To Represent SFRs And Bits? if these names are not recognized.

4.3.2.4 How Do I Find The Names Used To Represent SFRs And Bits?

Special function registers, and the bits within them, are accessed via special variables that map over the register (see 4.3.2.3 How Do I Access SFRS?). However, the names of these variables sometimes differ from those indicated in the data sheet for the device you are using.

View the device-specific header file which allows access to these special variables. Begin by searching for the data sheet SFR name. If that is not found, search on what the SFR represents, as comments in the header often spell out what the macros under the comment do.

4.3.3 Memory Allocation

Here are questions relating to how your source code affects memory allocation.

- 4.3.3.1 How Do I Position Variables at an Address I Nominate?
- 4.3.3.2 How Do I Position Functions at an Address I Nominate?
- 4.3.3.3 How Do I Place Variables in Program Memory?
- 4.3.3.4 How Do I Stop the Compiler From Using Certain Memory Locations?
- 4.5.14 Why Are Some Objects Positioned Into Memory That I Reserved?

4.3.3.1 How Do I Position Variables at an Address I Nominate?

The easiest way to do this is to make the variable absolute by using the `address` attribute (see 9.11 Variable Attributes) or the `__at()` CCI construct (see 3.4.2 Absolute Addressing). This means that the address you specify is used in preference to the variable's symbol in generated code. Since you nominate the address, you have full control over where objects are positioned, but you must also ensure that absolute variables do not overlap.

See also 10.3 Auto Variable Allocation and Access for information on moving auto variables, 10.2.1 Non-Auto Variable Allocation for moving non-auto variables and 10.4 Variables in Program Memory for moving program-space variables.

4.3.3.2 How Do I Position Functions at an Address I Nominate?

The easiest way to do this is to make the functions absolute, by using the `address` attribute (see 14.2.1 Function Attributes). This means that the address you specify is used in preference to the function's symbol in generated code. Since you nominate the address, you have full control over where functions are positioned, but must also ensure that absolute functions do not overlap.

4.3.3.3 How Do I Place Variables in Program Memory?

The `const` qualifier implies that the qualified variable is read only. See the `-membedded-data` option in 6.7.1 Options Specific to PIC32M Devices for information about allocating 'const' qualified variables to program memory (Flash). As a consequence of this any variables, except for auto variables or function parameters, qualified `const` are placed in program memory, thus freeing valuable data RAM (see 10.4 Variables in Program Memory). Variables qualified `const` can also be made absolute, so that they can be positioned at an address you nominate.

4.3.3.4 How Do I Stop the Compiler From Using Certain Memory Locations?

Concatenating an address attribute with the `noload` attribute can be used to block out sections of memory. Also, you can use the option `-mreserve`. For more on variable attributes and options, see the following sections in this user's guide:

See the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186) for details on linker scripts.

4.3.4 Variables

This examines questions that relate to the definition and usage of variables and types within a program.

4.3.4.1 Why Are My Floating-point Results Not Quite What I Am Expecting?

First, make sure that if you are watching floating-point variables in MPLAB IDE that the type and size of these match how they are defined. In MPLAB XC32, the float and double types are 32-bit floating-point types by default. The long double type is a 64-bit floating-point type.

The size of the floating point type can be adjusted for `double` types. See 9.4 Floating-Point Data Types.

Since floating-point variables only have a finite number of bits to represent the values they are assigned, they will hold an approximation of their assigned value. A floating-point variable can only hold one of a set of discrete real number values. If you attempt to assign a value that is not in this set, it is rounded to the nearest value. The more bits used by the mantissa in the floating-point variable, the more values can be exactly represented in the set and the average error due to the rounding is reduced.

Whenever floating-point arithmetic is performed, rounding also occurs. This can also lead to results that do not appear to be correct.

4.3.4.2 How Can I Access Individual Bits of a Variable?

There are several ways of doing this. The simplest and most portable way is to define an integer variable and use macros to read, set, or clear the bits within the variable using a mask value and logical operations, such as the following.

```
#define  testbit(var, bit)    (!!(var) & (1 <<(bit)))
#define  setbit(var, bit)     ((var) |= (1 << (bit)))
#define  clrbit(var, bit)     ((var) &= ~(1 << (bit)))
```

These, respectively, test to see if bit number, `bit`, in the integer, `var`, is set; set the corresponding `bit` in `var`; and clear the corresponding `bit` in `var`. Alternatively, a union of an integer variable and a structure with bit-fields (see 9.5.2 Bit Fields in Structures) can be defined, for example,

```
union both {
   unsigned char byte;
   struct {
      unsigned b0:1, b1:1, b2:1, b3:1, b4:1, b5:1, b6:1, b7:1;
   } bitv;
} var;
```

This allows you to access `byte` as a whole (using var.byte), or any bit within that variable independently (using `var.bitv.b0` through `var.bitv.b7`).

4.3.4.3 How Long Can I Make My Variable and Macro Names?

The C Standard indicates that only a number of initial characters in an identifier are significant, but it does not actually state what this number is and it varies from compiler to compiler. For MPLAB XC32, no limit is imposed, but for CCI there is a limit (see 3.3.5 The Number of Significant Initial Characters in an Identifier). CCI Compliant names are more portable across Microchip architectures.

If two identifiers only differ in the non-significant part of the name, they are considered to represent the same object, which will almost certainly lead to code failure.

4.3.5 Functions

This section examines questions that relate to functions.

- 4.3.5.1 What Is the Optimum Size for Functions?
- 4.5.11 How Can I Tell How Big a Function Is?
- 4.5.12 How Do I Know What Resources Are Being Used by Each Function?
- 4.5.13 How Do I Find Out Where Variables and Functions Have Been Positioned?
- 4.3.6.1 How Do I Use Interrupts in C?
- 4.3.5.2 How Do I Stop An Unused Function Being Removed?
- 4.3.5.3 How Do I Make a Function Inline?

4.3.5.1 What Is the Optimum Size for Functions?

Generally speaking, the source code for functions should be kept small as this aids in readability and debugging. It is much easier to describe and debug the operation of a function which performs a small number of tasks. Also smaller-sized functions typically have less side effects, which can be the source of coding errors. On the other hand, in the embedded programming world, a large number of small functions, and the calls necessary to execute them, may result in excessive memory and stack usage. Therefore a compromise is often necessary.

Function size can cause issues with memory paging, as addressed in 14.5 Function Size Limits. The smaller the functions, the easier it is for the linker to allocate them to memory without errors.

4.3.5.2 How Do I Stop An Unused Function Being Removed?

The `__attribute__((keep))` may be applied to a function. The keep attribute will prevent the linker from removing the function with `--gc-sections`, even if it is unused. See the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186) for more information on section garbage collection using the `--gc-sections` option.

4.3.5.3 How Do I Make a Function Inline?

The XC32 compiler does not inline any functions when not optimizing.

By declaring a function inline, you can direct the XC32 compiler to make calls to that function faster. One way XC32 can achieve this is to integrate that function's code into the code for its callers. This makes execution faster by eliminating the function-call overhead; in addition, if any of the actual argument values are constant, their known values may permit simplifications at compile time so that not all of the inline function's code needs to be included. The effect on code size is less predictable; object code may be larger or smaller with function inlining, depending on the particular case.

To declare a function inline, use the inline keyword in its declaration, like this:

```
static inline int
inc (int *a)
{
   return (*a)++;
}
```

When a function is both inline and static, if all calls to the function are integrated into the caller, and the function's address is never used, then the function's own assembler code is never referenced. In this case, XC32 does not actually output assembler code for the function. Some calls cannot be integrated for various reasons (in particular, calls that precede the function's definition cannot be integrated, and neither can recursive calls within the definition). If there is a non-integrated call, then the function is compiled to assembler code as usual. The function must also be compiled as usual if the program refers to its address, because that can't be inlined. Enable optimization level -O1 or greater to enable function inlining.

4.3.6 Interrupts

Interrupt and interrupt service routine questions are discussed in this section.

4.3.6.1 How Do I Use Interrupts in C?

First, be aware of what interrupt hardware is available on your target device. 32-bit devices implement several separate interrupt vector locations and use a priority scheme. For more information, see 15.1 Interrupt Operation.

In C source code, a function can be written to act as the interrupt service routine by using the `interrupt` attribute. Such functions save/restore program context before/after executing the function body code and a different return instruction is used. For more on writing interrupt functions, see 15.2 Writing an Interrupt Service Routine. To populate the interrupt vector table, use the `vector` or `at_vector` attribute. An `__ISR()` macro is provided in the sys/attribs.h header file that simplifies the usage of the `interrupt` and `vector` attributes.

Prior to any interrupt occurring, your program must ensure that peripherals are correctly configured and that interrupts are enabled. For details, see 15.8 Enabling/Disabling Interrupts.

For all other interrupt related tasks, including specifying the interrupt vector, context saving, nesting and other considerations, consult 15. Interrupts.

4.3.7 Assembly Code

This section examines questions that arise when writing assembly code as part of a C project.

4.3.7.1 How Should I Combine Assembly and C Code?

Ideally, any hand-written assembly should be written as separate routines that can be called. This offers some degree of protection from interaction between compiler-generated and hand-written assembly code. Such code can be placed into a separate assembly module that can be added to your project, as specified in 18.1 Mixing Assembly Language and C Variables and Functions.

If necessary, assembly code can be added in-line with C code by using either of two forms of the `asm` instruction; simple or extended. An explanation of these forms, and some examples, are shown in 18.2 Using Inline Assembly Language.

Macros are provided which in-line several simple instructions, as discussed in 18.3 Predefined Macros. More complex in-line assembly that changes register contents and the device state should be used with caution.

See 13.1 Register Usage for those registers used by the compiler.

4.3.7.2 What Do I Need Other Than Instructions in an Assembly Source File?

Assembly code typically needs assembler directives as well as the instructions themselves. The operation of all the directives is described in the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186). Two common directives are discussed below.

All assembly code must be placed in a section, using the .section directive, so it can be manipulated as a whole by the linker and placed in memory. See the "Linker Processing" chapter of the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186) for more information.

Another commonly used directive is `.global` which is used to make symbols accessible across multiple source files. Find more on this directive in the aforementioned user's guide.

4.3.7.3 How Do I Access C Objects from Assembly Code?

Most C objects are accessible from assembly code. There is a mapping between the symbols used in the C source and those used in the assembly code generated from this source. Your assembly should access the assembly-equivalent symbols which are detailed in 18.1 Mixing Assembly Language and C Variables and Functions.

Instruct the assembler that the symbol is defined elsewhere by using the `.global` assembler directive. This is the assembly equivalent of a C declaration, although no type information is present. This directive is not needed and should not be used if the symbol is defined in the same module as your assembly code.

Any C variable accessed from assembly code will be treated as if it were qualified `volatile` (see 9.9.2 Volatile Type Qualifier). Specifying the `volatile` qualifier in C code is preferred as it makes it clear that external code may access the object.

4.3.7.4   How Can I Access SFRs From Within Assembly Code?

The safest way to gain access to SFRs in assembly code is to have symbols defined in your assembly code that equate to the corresponding SFR address. For the XC32 compiler, the xc.h include file can be used from either preprocessed assembly code or C/C++ code.

There is no guarantee that you will be able to access symbols generated by the compilation of C code, even code that accesses the SFR you require.

4.3.7.5   What Things Must I Manage When Writing Assembly Code?

There are several things that you must manage if you are hand-writing assembly code.

- You must place any assembly code you write into a section. See the "Linker Processing" chapter of the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186) for more information.

  Assembly code that is placed in-line with C code will be placed in the same section as the compiler-generated assembly and you should not place this into a separate section.
- You must ensure that any registers you write to in assembly code are not already in use by compiler-generated code. If you write assembly in a separate module, then this is less of an issue as the compiler will, by default, assume that all registers are used by these routines (see 13. Register Usage, registers). No assumptions are made for in-line assembly (see 18.1 Mixing Assembly Language and C Variables and Functions) and you must be careful to save and restore any resources that you use (write) and which are already in use by the surrounding compiler-generated code.

4.4   Getting My Application To Do What I Want

This section provides programming techniques, applications and examples. It also examines questions that relate to making an application perform a specific task.

- 4.4.1 What Can Cause Glitches on Output Ports?
- 4.4.2 How Do I Link Bootloaders and Downloadable Applications?
- 4.4.3 What Do I Need to Do When Compiling to Use a Debugger?
- 4.4.4 How Do I Share Data Between Interrupt and Main-line Code?
- 4.4.5 How Can I Prevent Misuse of My Code?
- 4.4.6 How Do I Use Printf to Send Text to a Peripheral?
- 4.4.7 How Can I Implement a Delay in My Code?
- 4.4.8 How Can I Rotate a Variable?

4.4.1   What Can Cause Glitches on Output Ports?

In most cases, this is caused by using ordinary variables to access port bits or the entire port itself. These variables should be qualified `volatile`. See 9.9.2 Volatile Type Qualifier.

The value stored in a variable mapped over a port (hence the actual value written to the port) directly translates to an electrical signal. It is vital that the values held by these variables only change when the code intends them to, and that they change from their current state to their new value in a single transition. The compiler attempts to write to `volatile` variables in one operation.

4.4.2   How Do I Link Bootloaders and Downloadable Applications?

Exactly how this is done depends on the device you are using and your project requirements, but the general approach when compiling applications that use a bootloader is to allocate discrete program memory space to the bootloader and application so they have their own dedicated memory. In this way the operation of one cannot affect

the other. This will require that either the bootloader or the application is offset in memory. That is, the Reset and interrupt location are offset from address 0 and all program code is offset by the same amount.

Typically the application code is offset, and the bootloader is linked with no offset so that it populates the Reset and interrupt code locations. The bootloader Reset and interrupt code merely contains code which redirects control to the real Reset and interrupt code defined by the application and which is offset.

The contents of the Hex file for the bootloader can be merged with the code of the application by using loadable projects in MPLAB X IDE. (See MPLAB X IDE documentation for details.) This results in a single Hex file that contains the bootloader and application code in the one image. Check for warnings from this application about overlap, which may indicate that memory is in use by both bootloader and the downloadable application.

See the PIC32 Bootloader Application Note (DS01388) on the Microchip website.

4.4.3 What Do I Need to Do When Compiling to Use a Debugger?

You can use debuggers, such as the MPLAB PICkit™ 4 or ICD 4 in-circuit debuggers or the MPLAB REAL ICE™ in-circuit emulator, to debug code built with the MPLAB XC32 compiler. These debuggers use some of the data and program memory of the device for its own use, so it is important that your code does not also use these resources.

The command-line option -g (see 6.7.6 Options for Debugging) is used to tell the compiler to generate debugging information. The compiler can then reserve the memory used by the debugger so that your code will not be placed in these locations.

In the MPLAB X IDE, the appropriate debugger option is specified if you perform a Debug Run. It will not be specified if you perform a regular Run, Build Project, or Clean and Build.

Since some device memory is being reserved for use by the debugger, there is less available for your program and it is possible that your code or data may no longer fit in the device when a debugger is selected. For 32-bit devices, some boot flash memory is required for debugging. In addition, some data memory (RAM) is used by the debug tool and may impact the variable allocation in your application.

The specific memory locations used by the debuggers are attributes of MPLAB X IDE, the debug tool in use and the target device. If you move a project to a new version of the IDE, the resources required may change. For this reason, you should not manually reserve memory for the debugger, or make any assumptions in your code as to what memory is used. A summary of the debugger requirements is available in the MPLAB X IDE help files.

To verify that the resources reserved by the compiler match those required by the debugger, you may view the boot-flash, application-flash and data-memory usage in the map file or memory-usage report.

To create a map file in MPLAB X IDE, open the Project Properties window (*File>Project Properties*) and click on the linker category (xc32-ld). Under "Option Categories," select "Diagnostics." Next to "Generate map file," enter a path and name for the map file. The logical place to put the map file is in the project folder.

Debug Run your code to generate the map file. View in your favorite text viewer.

See also 4.5.14 Why Are Some Objects Positioned Into Memory That I Reserved?.

4.4.4 How Do I Share Data Between Interrupt and Main-line Code?

Variables accessed from both interrupt and main-line code can easily become corrupted or mis-read by the program. The volatile qualifier (see 9.9.2 Volatile Type Qualifier) tells the compiler to avoid performing optimizations on such variables. This will fix some of the issues associated with this problem.

The other issues relates to whether the compiler/device can access the data atomically. With 32-bit PIC devices, this is rarely the case. An atomic access is one where the entire variable is accessed in only one instruction. Such access is uninterruptible. You can determine if a variable is being accessed atomically by looking at the assembler list file (see the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide*, DS50002186, for more information). If the variable is accessed in one instruction, it is atomic. Since the way variables are accessed can vary from statement to statement it is usually best to avoid these issues entirely by disabling interrupts prior to the variable being accessed in main-line code, then re-enable the interrupts afterwards. See 15.8 Enabling/Disabling Interrupts for more information. When writing to Special Function Registers (SFRs), use the SET/CLR/INV registers as described in 8.5 ID Locations.

4.4.5 How Can I Prevent Misuse of My Code?

First, many devices with flash program memory allow all or part of this memory to be write protected. The device configuration bits need to be set correctly for this to take place, so see 8.4 Configuration Bit Access, 3.4.14 Specifying Configuration Bits for CCI, and your device data sheet.

Second, you can prevent third-party code being programmed at unused locations in the program memory by filling these locations with a value rather than leaving them in their default unprogrammed state. You can chose a fill value that corresponds to an instruction or set all the bits so as the values cannot be further modified. (Consider what will happen if your program somehow reaches and starts executing from these filled values. What instruction will be executed?)

Use the `--fill` command to fill unused memory. Find usage information for this command in 6.7.10 Options for Linking.

4.4.6 How Do I Use Printf to Send Text to a Peripheral?

The `printf` function does two things: it formats text based on the format string and placeholders you specify, and sends (prints) this formatted text to a destination (or stream). You may choose the `printf` output go to an LCD, SPI module or USART, for example.

For more on the ANSI C function `printf`, see the *32-bit Language Tool Libraries* manual (DS51685).

To check what is passed to the `printf` function, you may attempt to statically analyze format strings passed to the function by using the `-msmart-io` option (6.7.1 Options Specific to PIC32M Devices). Also you may use the `-Wformat` option to specify a warning when the arguments supplied to the function do not have types appropriate to the format string specified (see 6.7.5 Options for Controlling Warning and Errors).

If you wish to create your own `printf`-type function, you will need to use the attributes `format` and `format_arg` as discussed in 14.2.1 Function Attributes.

4.4.7 How Can I Implement a Delay in My Code?

If an accurate delay is required, or if there are other tasks that can be performed during the delay, then using a timer to generate an interrupt is the best way to proceed.

Microchip does not recommend using a software delay on PIC32 devices as there are many variables that can affect timing such as the configuration of the L1 cache, prefetch cache, & Flash wait states. On PIC32 devices, you may choose to poll the core timer, which increments every two instruction cycles.

4.4.8 How Can I Rotate a Variable?

The C language does not have a rotate operator, but rotations can be performed using the shift and bitwise OR operators. Since the 32-bit devices have a rotate instruction, the compiler will look for code expressions that implement rotates (using shifts and ORs) and use the rotate instruction in the generated output wherever possible.

If you are using CCI, you should consult 3.3.10 Bitwise Operations on Signed Values and 3.3.11 Right-Shifting Signed Values if you will be using signed variables.

For the following example C code:

```
int rotate_left (unsigned a, unsigned s)
{
  return (a << s) | (a >> (32 - s));
}
```

the compiler may generate assembly instructions similar to the following:

```
rotate_left:
    subu      $2,$0,$5
    jr        $31
    ror       $2,$4,$2
```

4.5   Understanding the Compilation Process

This section tells you how to find out what the compiler did during the build process, how it encoded output code, where it placed objects, etc. It also discusses the features that are supported by the compiler.

4.5.1   What's the Difference Between the Free and PRO Modes?

These modes, or editions, mainly differ in the optimizations that are performed when compiling (see 19. Optimizations). Compilers operating in Free mode can compile for all the same devices as supported by the Pro mode. The code compiled in Free or PRO modes can use all the available memory for the selected device. What will be different is the size and speed of the generated compiler output. Free mode output will be less efficient when compared to that produced in Pro mode.

4.5.2   How Can I Make My Code Smaller?

There are a number of ways that this can be done, but results vary from one project to the next. Use the assembly list file to observe the assembly code produced by the compiler to verify that the following tips are relevant to your code. For information on the list file, see the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186).

Use the smallest data types possible as less code is needed to access these. (This also reduces RAM usage.) For examples, a short integer type exists for this compiler. See 9. Supported Data Types and Variables for all data types and sizes.

There are two sizes of floating-point type, as well, and these are discussed in the same section. Replace floating-point variables with integer variables wherever possible. For many applications, scaling a variable's value makes eliminating floating-point operations possible.

Use unsigned types, if possible, instead of signed types, particularly if they are used in expressions with a mix of types and sizes. Try to avoid an operator acting on operands with mixed sizes whenever possible.

Whenever you have a loop or condition code, use a "strong" stop condition, that is, the following:

```
for(i=0; i!=10; i++)
```

is preferable to:

```
for(i=0; i<10; i++)
```

A check for equality (== or !=) is usually more efficient to implement than the weaker < comparison.

In some situations, using a loop counter that decrements to zero is more efficient than one that starts at zero and counts up by the same number of iterations. So you might be able to rewrite the above as:

```
for(i=10; i!=0; i--)
```

Ensure that you enable all the optimizations allowed for the edition of your compiler (see 19. Optimizations). If you have a Pro edition, you can use the -Os option (see 6.7.7 Options for Controlling Optimization) to optimize for size. Otherwise, pick the highest optimization available.

Consider using the a compressed ISA mode such as MIPS16 or microMIPS if it is supported on your device. Use the -mips16 or -mmicromips option for your project to make the compiler default to these modes. Use the mips16 or micromips function attributes to change the mode at the function level. You may also choose the optimized and compressed variants of the libraries in the linker options. Be aware of what optimizations the compiler performs so you can take advantage of them and don't waste your time manually performing optimizations in C code that the compiler already handles, for example, don't turn a multiply-by-4 operation into a shift-by-2 operation as this sort of optimization is already detected.

4.5.3 How Can I Reduce RAM Usage?

Consider using auto variables rather than global or static variables as there is the potential that these may share memory allocated to other auto variables that are not active at the same time. Memory allocation of auto variables is made on a stack, described in 10.3 Auto Variable Allocation and Access.

Rather than pass large objects to, or from, functions, pass pointers which reference these objects. This is particularly true when larger structures are being passed.

Objects that do not need to change throughout the program can be located in program memory using the const qualifier (see 10.4 Variables in Program Memory). This frees up precious RAM, but slows execution.

4.5.4 How Can I Make My Code Faster?

To a large degree, smaller code is faster code, so efforts to reduce code size often decrease execution time. To accomplish this, see 4.5.2 How Can I Make My Code Smaller? and 4.5.6 How Can I Make My Interrupt Routine Faster?. However, there are ways some sequences can be sped up at the expense of increased code size.

Depending on your compiler edition (see 19. Optimizations), you may be able to use the -O3 option (see 6.7.7 Options for Controlling Optimization) to optimize for speed. This will use alternate output in some instances that is faster, but larger.

Generally, the biggest gains to be made in terms of speed of execution come from the algorithm used in a project. Identify which sections of your program need to be fast. Look for loops that might be linearly searching arrays and choose an alternate search method such as a hash table and function. Where results are being recalculated, consider if they can be cached.

4.5.5 How Does the Compiler Place Everything in Memory?

In most situations, assembly instructions and directives associated with both code and data are grouped into sections, and these are then positioned into containers which represent the device memory. To see what sections objects are placed in, use the option `-ai` to view this information in the assembler listing file.

The exception is for absolute variables, which are placed at a specific address when they are defined and which are not placed in a section. For setting absolute variables, use the `address()` attribute specified under 9.11 Variable Attributes.

4.5.6 How Can I Make My Interrupt Routine Faster?

Consider suggestions made in 4.5.2 How Can I Make My Code Smaller? (code size) for any interrupt code. Smaller code is often faster code.

In addition to the code you write in the ISR, there is the code the compiler produces to switch context. This is executed immediately after an interrupt occurs and immediately before the interrupt returns, so must be included in the time taken to process an interrupt. This code is optimal in that only registers used in the ISR will be saved by this code. Thus, the fewer registers used in your ISR will mean potentially less context switch code to be executed.

Generally simpler code will require fewer resources than more complicated expressions. Use the assembly list file to see which registers are being used by the compiler in the interrupt code. For information on the list file, see the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186).

Avoid calling other functions from the ISR. In addition to the extra overhead of the function call, the compiler also saves all general purpose registers that may or may not be used by the called function. Consider having the ISR simply set a flag and return. The flag can then be checked in main-line code to handle the interrupt. This has the advantage of moving the complicated interrupt-processing code out of the ISR so that it no longer contributes to its register usage. Always use the `volatile` qualifier (see 9.9.2 Volatile Type Qualifier for variables shared by the interrupt and main-line code, see 4.4.4 How Do I Share Data Between Interrupt and Main-line Code?.

4.5.7 How Big Can C Variables Be?

This question specifically relates to the size of individual C objects, such as arrays or structures. The total size of all variables is another matter.

To answer this question you need to know in which memory space the variable will be located.

With the default `-membedded-data` option, objects qualified `const` will be located in program memory; other objects will be placed in data memory. Program memory object sizes are discussed in 10.4.1 Size Limitations of const Variables. Objects in data memory are broadly grouped into autos and non-autos and the size limitations of these objects, respectively, are discussed in 10.2.1 Non-Auto Variable Allocation and 10.2.3 Non-Auto Variable Size Limits.

4.5.8 What Optimizations Will Be Applied to My Code?

Code optimizations available depend on the edition of your compiler (see 19. Optimizations). A description of optimization options can be found under 6.7.7 Options for Controlling Optimization.

4.5.9 What Devices are Supported by the Compiler?

New devices are usually added with each compiler release. Check the readme document for a full list of devices supported by a compiler release.

4.5.10 How Do I Know What Code the Compiler Is Producing?

The assembly list file may be set up, using assembler listing file options, to contain a great deal of information about the code, such as the assembly output for almost the entire program, including library routines linked in to your program; section information; symbol listings and more.

The list file may be produced as follows:

- On the command line, create a basic list file using the option:

```
-Wa, -a=projectname.lst
```

- For MPLAB X IDE, right click on your project and select "Properties." In the Project Properties window, click on "xc32-as" under "Categories." From "Option categories," select "Listing file options" and ensure "List to file" is checked.

By default, the assembly list file will have a `.lst` extension.

For information on the list file, see the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186).

4.5.11 How Can I Tell How Big a Function Is?

This size of a function (the amount of assembly code generated for that function) can be determined from the assembly list file. See 4.5.10 How Do I Know What Code the Compiler Is Producing? for more on creating an assembly listing file.

4.5.12 How Do I Know What Resources Are Being Used by Each Function?

In the assembly list file there is information printed for every C function, including library functions. See 4.5.10 How Do I Know What Code the Compiler Is Producing? for more on creating an assembly listing file.

To see information on functions calls, you can view the Call Graph in MPLAB X IDE (*Window>Output>Call Graph*). You must be in debug mode to see this graph. Right click on a function and select "Show Call Graph" to see what calls this function and what it calls.

Auto, parameter and temporary variables used by a function may overlap with those from other functions as these are placed in a compiled stack by the compiler, see 10.3 Auto Variable Allocation and Access.

4.5.13 How Do I Find Out Where Variables and Functions Have Been Positioned?

You can determine where variables and functions have been positioned from either the assembly list file (generated by the assembler) or the map file (generated by the linker). Only global symbols are shown in the map file; all symbols (including locals) are listed in the assembly list file.

There is a mapping between C identifiers and the symbols used in assembly code, which are the symbols shown in both of these files. The symbol associated with a variable is assigned the address of the lowest byte of the variable; for functions it is the address of the first instruction generated for that function.

For more on assembly list files and linker map files, see the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186).

4.5.14 Why Are Some Objects Positioned Into Memory That I Reserved?

Most variables and function are placed into sections that are defined in the linker script. See the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186) for details on linker scripts. However, some variables and function are explicitly placed at an address rather than being linked anywhere in an address range, as described in 4.3.3.1 How Do I Position Variables at an Address I Nominate? and 4.3.3.2 How Do I Position Functions at an Address I Nominate?.

Check the assembly list file to determine the names of sections that hold objects and code. Check the linker options in the map file to see if sections have been linked explicitly or if they are linked anywhere in a class. See the *"MPLAB® XC32 Assembler, Linker and Utilities User's Guide"* (DS50002186) for information on each of these files.

4.5.15 How Do I Know How Much Memory Is Still Available?

A memory usage summary is available from the compiler after compilation (`--report-mem` option), from MPLAB X IDE in the Dashboard window. All of these summaries indicate the amount of memory used and the amount still available, but none indicate whether this memory is one contiguous block or broken into many small chunks. Small blocks of free memory cannot be used for larger objects and so out-of-memory errors may be produced even though the total amount of memory free is apparently sufficient for the objects to be positioned.

Consult the linker map file to determine exactly what memory is still available in each linker class. This file also indicates the largest contiguous block in that class if there are memory page divisions. See the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186) for information on the map file.

4.5.16 How Do I Use Library Files in My Project?

See 4.2.6 How Do I Build Libraries? for information on how you build your own library files. The compiler will automatically include any applicable standard library into the build process when you compile, so you never need to control these files.

To use one or more library files that were built by yourself or a colleague, include them in the list of files being compiled on the command line. The library files can be specified in any position in the file list relative to the source files, but if there is more than one library file, they will be searched in the order specified in the command line.

For example:

```
xc32-gcc -mprocessor=32MZ2048ECH100 main.c int.c mylib.a
```

If you are using MPLAB IDE to build a project, add the library file(s) to the Libraries folder that will shown in your project, in the order in which they should be searched. The IDE will ensure that they are passed to the compiler at the appropriate point in the build sequence.

4.5.17 How Do I Customize the C Runtime Startup Code?

Some applications may require an application-specific version of the C runtime startup code. For instance, you may want to modify the startup code for an application loaded by a bootloader.

To customize the startup code for your application:

1. Start with the default startup code, a copy of which is located in the pic32m-libs.zip file located at:

    ```
    <install-directory>/pic32-libs/
    ```

    Once the file is unzipped, the source code can be found at:

    ```
    pic32m-libs/libpic32/startup/crt0.S
    ```

2. Make a copy of this crt0.S file, rename it, and add it to your project.
3. Change your MPLAB X project to exclude the default startup code by enabling the "Do not link startup code" under XC32 xc32-ld Option categories: Libraries page as shown below. When you build your project, the MPLAB X IDE will build your new application-specific copy of the startup code rather than linking in the default code.

Figure 4-1. Startup Code Properties Setting
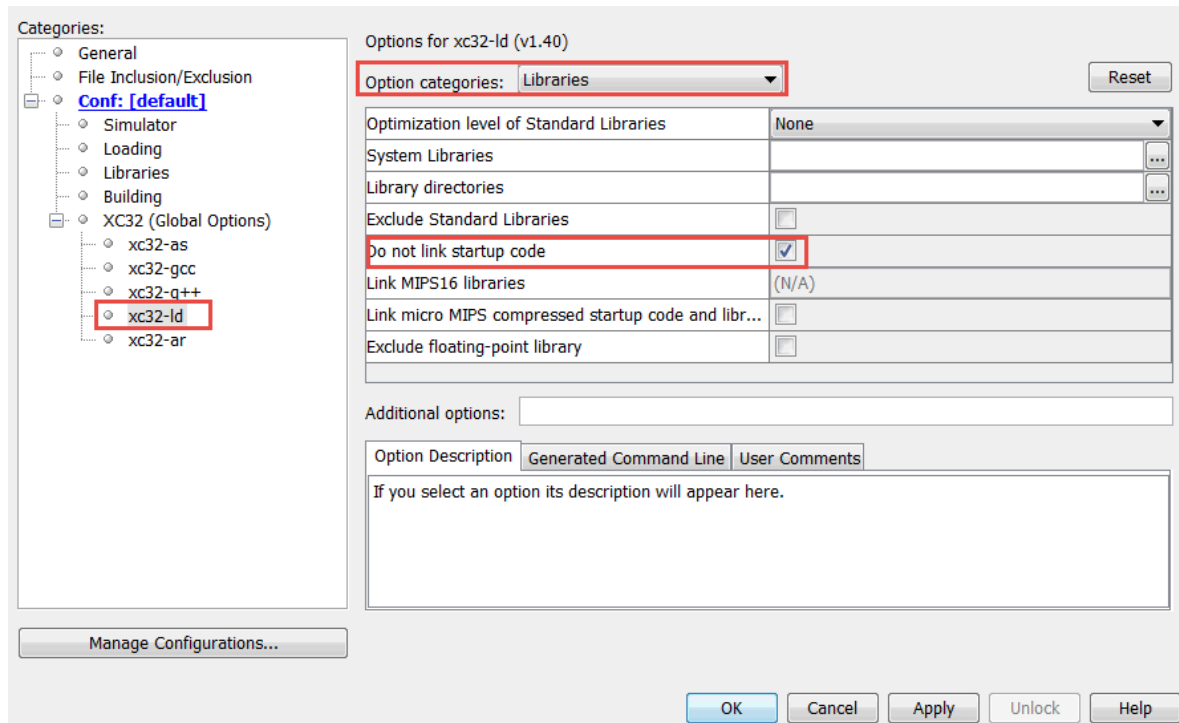
4. You can now edit the assembly code in your new copy of the crt0.S file. The default source code uses macros defined in the device-specific header files, which are included by xc.h. These macros define various device specific behavior. Be sure to take this device-specific code into account when customizing your copy of the code.

Table 4-1. Device-Specific Macros

Device-Specific Macro	Description
`__PIC32_SRS_SET_COUNT`	Defined to the number of Register Sets implemented on the device. The default startup code uses this value to determine how many register sets to use for $GP-register initialization.
`__PIC32_HAS_L1CACHE`	Defined if the device features an L1 cache.
`__PIC32_HAS_MIPS32R2`	Defined if the device supports the MIPS32r2 Instruction Set.
`__PIC32_HAS_MICROMIPS`	Defined if the device supports the microMIPS compressed Instruction Set.
`__PIC32_HAS_DSPR2`	Defined if the device supports the DSPr2 Application-Specific Extension.
`__PIC32_HAS_FPU64`	Defined if the device supports the single- and double-precision hardware Floating Point Unit.
`__PIC32_HAS_SSX`	Defined if the device does *not* require initialization of the bus matrix registers in order to support execution from data memory.
`__PIC32_HAS_MMU_MZ_FIXED`	Defined if the device features a Memory Management Unit that should be pre-initialized to a standard SQI and EBI mapping.
`__PIC32_HAS_INIT_DATA`	Defined if the device requires data initialization by copying from a template located in Flash to RAM.

4.5.18 What Optimizations Are Employed by the Compiler?

Code optimizations available depend on the edition of your compiler (see 19. Optimizations). A description of optimization options can be found under 6.7.7 Options for Controlling Optimization.

4.5.19 Why Do I Get Out-of-Memory Errors When I Select a Debugger?

If you use a hardware tool debugger such as MPLAB PICkit 4 in-circuit debugger or MPLAB ICD 4 in-circuit debugger memory is required for the on-board debug executive.

4.5.20 How Do I Stop My Project's Checksum From Changing?

The checksum that represents your built project (whether this is generated by the MPLAB X IDE or by tools such as Hexmate) is calculated from the generated output of the compiler. Indeed, the algorithms used to obtain the checksum are specifically designed so that even small changes in this output are almost guaranteed to produce a different checksum result. Checksums are not calculated from your project's source code. To ensure that your checksum does not change from build to build, you must ensure that the output of the compiler does not change.

The following actions and situations could cause changes in the compiled output and hence changes in your project's checksum.

- Changing the source code, header files, or library code used by the project between builds.
- Changing the order in which source files or libraries are compiled or linked between builds.
- Having source code that makes using of macros such as `__DATE__` and `__TIME__`, which produce output that is dependent on when the project was built.
- Moving the location of source files between builds, where those files use macros such as `__FILE__`, which produces output that is dependent on where the source file is located.
- Changing the compiler options between builds.
- Changing the compiler version between builds.

Note that the checksum algorithms used by tools such as Hexmate and the MPLAB X IDE can change, which can result in a different checksum for the same compiler output. Such changes are rare, but check the compiler and IDE release notes to see if the tools have been modified.

4.6   Fixing Code That Does Not Work

This section examines issues relating to projects that do not build due to compiler errors, or those that build, but do not work as expected.

4.6.1   How Do I Set Up Warning/Error Messages?

To control message output, see 6.7.5 Options for Controlling Warning and Errors.

4.6.2   How Do I Find the Code that Caused Compiler Errors or Warnings in My Program?

In most instances, where the error is a syntax error relating to the source code, the message produced by the compiler indicates the offending line of code (see 6.6 Compiler Messages). If you are compiling in MPLAB X IDE, then you can double-click the message and have the editor take you to the offending line. But identifying the offending code is not always so easy.

In some instances, the error is reported on the line of code following the line that needs attention. This is because a C statement is allowed to extend over multiple lines of the source file. It is possible that the compiler cannot be able to determine that there is an error until it has started to scan to statement following. So in the following code

```
input = PORTB   // oops - forgot the semicolon
if(input>6)
// ...
```

The missing semicolon on the assignment statement will be flagged on the following line that contains the `if()` statement.

In other cases, the error might come from the assembler, not the code generator. If the assembly code was derived from a C source file, then the compiler will try to indicate the line in the C source file that corresponds to the assembly that is at fault. If the source being compiled is an assembly module, the error directly indicates the line of assembly that triggered the error. In either case, remember that the information in the error relates to some problem is the assembly code, not the C code.

Finally, there are errors that do not relate to any particular line of code at all. An error in a compiler option or a linker error are examples of these. If the program defines too many variables, there is no one particular line of code that is at fault; the program as a whole uses too much data. Note that the name and line number of the last processed file and source can be printed in some situations even though that code is not the direct source of the error.

At the top of each message description, on the right in brackets, is the name of the application that produced this message. Knowing the application that produced the error makes it easier to track down the problem. The compiler application names are indicated in 5. XC32 Toolchain and MPLAB X IDE.

If you need to see the assembly code generated by the compiler, look in the assembly list file. For information on where the linker attempted to position objects, see the map file. See the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186) for information about the list and map files.

4.6.3   How Can I Stop Spurious Warnings From Being Produced?

Warnings indicate situations that could possibly lead to code failure. Always check your code to confirm that it is not a possible source of error. In many situations the code is valid and the warning is superfluous. In this case, you may:

- Inhibit specific warnings by using the `-Wno-` version of the option.
- Inhibit all warnings with the `-w` option.
- In MPLAB X IDE, inhibit warnings in the Project Properties window under each tool category. Also look in the Tool Options window, Embedded button, Suppressible Messages tab.

See 6.7.5 Options for Controlling Warning and Errors for details.

4.6.4   Why Can't I Even Blink an LED?

Even if you have set up the TRIS register and written a value to the port, there are several things that can prevent such a seemingly simple program from working.

- Make sure that the device's configuration registers are set up correctly, as discussed in 8.4 Configuration Bit Access. Make sure that you explicitly specify every bit in these registers and don't just leave them in their default state. All the configuration features are described in your device data sheet. If the configuration bits that specify the oscillator source are wrong, for example, the device clock may not even be running.
- If the internal oscillator is being used, in addition to Configuration bits there may be SFRs you need to initialize to set the oscillator frequency and modes, see 8.5 ID Locations and your device data sheet.
- To ensure that the device is not resetting because of the watchdog time, either turn off the timer in the configuration bits or clear the timer in your code. There are library functions you can use to handle the watchdog timer, described in the 32-bit Language Tool Libraries manual (DS51685). If the device is resetting, it may never reach the lines of code in your program that blink the LED. Turn off any other features that may cause device Reset until your test program is working.
- The device pins used by the port bits are often multiplexed with other peripherals. A pin might be connected to a bit in a port, or it might be an analog input, or it might be the output of a comparator, for example. If the pin connected to your LED is not internally connected to the port you are using, then your LED will never operate as expected. The port function tables in your device data sheets will show other uses for each pin which will help you identify peripherals to investigate.

4.6.5   What Can Cause Corrupted Variables and Code Failure When Using Interrupts?

This is usually caused by having variables used by both interrupt and main-line code. If the compiler optimizes access to a variable or access is interrupted by an interrupt routine, then corruption can occur. See 4.4.4 How Do I Share Data Between Interrupt and Main-line Code? for more information.

5. XC32 Toolchain and MPLAB X IDE

The 32-bit language tools may be used together under MPLAB X IDE to provide GUI development of application code for the PIC32 MCU families of devices. The tools are:

- MPLAB XC32 C/C++ Compiler
- MPLAB XC32 Assembler
- MPLAB XC32 Object Linker
- MPLAB XC32 Object Archiver/Librarian and other 32-bit utilities

5.1 MPLAB X IDE and Tools Installation

In order to use the 32-bit language tools with MPLAB X IDE, you must install:

- MPLAB X IDE, which is available for free on the Microchip website.
- MPLAB XC32 C/C++ Compiler, which includes all of the 32-bit language tools. The compiler is available for free (Free and Evaluation editions) or for purchase (Pro edition) on the Microchip website.

The 32-bit language tools will be installed, by default, in the directory:

- Windows OS - `C:\Program Files\Microchip\xc32\`*x.xx*
- Mac OS - `/Applications/microchip/xc32/`*x.xx*
- Linux OS - `/opt/microchip/xc32/`*x.xx*

where *x.xx* is the version number.

The executables for each tool will be in the `bin` subdirectory:

- C Compiler - `xc32-gcc.exe`
- Assembler - `xc32-as.exe`
- Object Linker - `xc32-ld.exe`
- Object Archiver/Librarian - `xc32-ar.exe`
- Other Utilities - `xc32-`*utility*`.exe`

All device include (header) files are located in the `/pic32mx/include/proc` subdirectory. These files are automatically incorporated when you `#include` the xc.h header file.

Code examples are located in the examples directory.

5.2 MPLAB X IDE Setup

Once MPLAB X IDE is installed on your PC, launch the application and check the settings below to ensure that the 32-bit language tools are properly recognized.

1. From the MPLAB X IDE menu bar, select *Tools>Options* to open the Options dialog. Click on the "Embedded" button and select the "Build Tools" tab.
2. Click on "XC32" under "Toolchain." Ensure that the paths are correct for your installation.
3. Click the **OK** button.

Figure 5-1. XC32 Tool Suite Locations in Windows® OS

5.3 MPLAB X IDE Projects

A project in MPLAB X IDE is a group of files needed to build an application, along with their associations to various build tools. Below is a generic MPLAB X IDE project.

Figure 5-2. Compiler Project Relationships
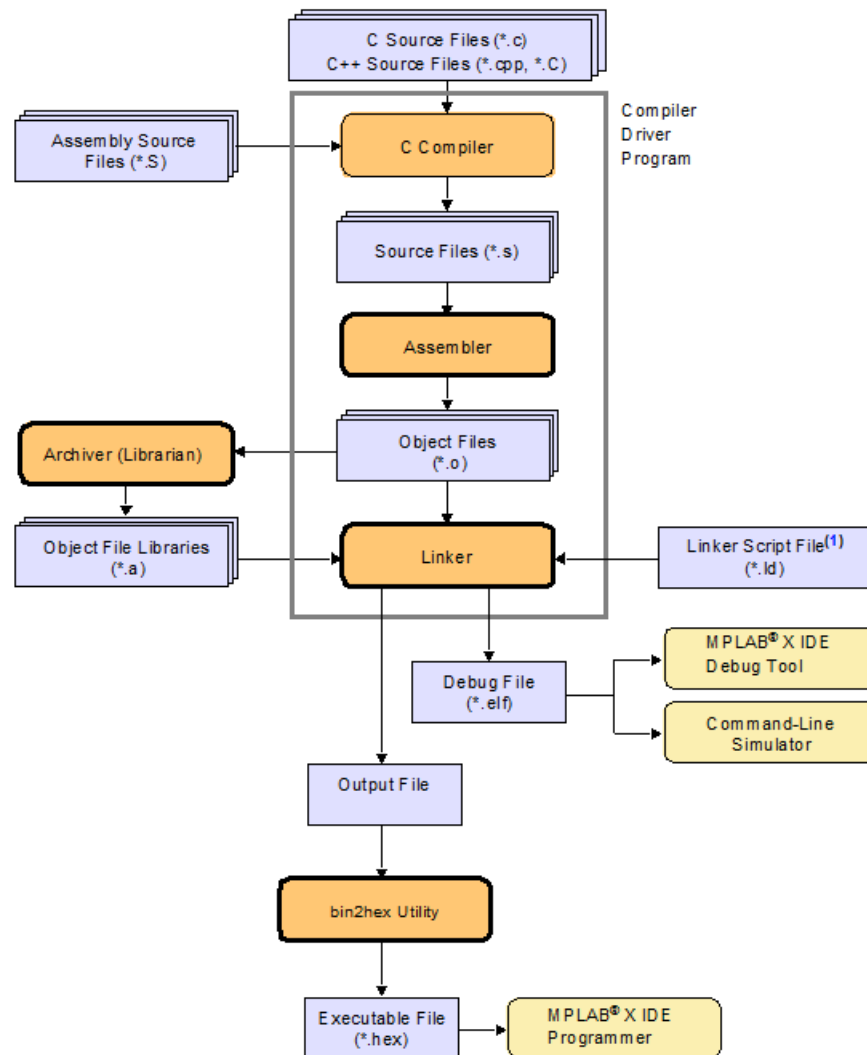
Note 1: The linker can choose the correct linker script file for your project.

In this MPLAB X IDE project, C source files are shown as input to the compiler. The compiler will generate source files for input into the assembler. For more information on the compiler, see the compiler documentation.

Assembly source files are shown as input to the C preprocessor. The resulting source files are input to the assembler. The assembler will generate object files for input into the linker or archiver. For more information on the assembler, see the assembler documentation.

Object files can be archived into a library using the archiver/librarian. For more information on the archiver, see the archiver/librarian documentation.

The object files and any library files, as well as a linker script file (generic linker scripts are added automatically), are used to generate the project output files via the linker. The output file that may be generated by the linker is a debug file (`.elf`) used by the simulator and debug tools which may be input into the bin2hex utility to produce an executable file (`.hex`). For more information on linker script files and using the object linker, see the linker documentation.

For more on projects and related workspaces, see MPLAB X IDE documentation.

5.4   Project Setup

To set up an MPLAB X IDE project for the first time, use the built-in Project Wizard (`File>New Project`). In this wizard, you will be able to select a language toolsuite that uses the 32-bit language tools. For more on the wizard, and MPLAB X IDE projects, see MPLAB X IDE documentation.

Once you have a project set up, you may then set up properties of the tools in MPLAB X IDE.

1. From the MPLAB X IDE menu bar, select _File>Project Properties_ to open a window to set/check project build options.
2. Under "_Conf:[default]_", select a tool from the tool collection to set up.
   - 5.4.1 XC32 (Global Options)
   - 5.4.2 xc32-as (32-bit Assembler)
   - 5.4.3 xc32-gcc (32-bit C Compiler)
   - 5.4.4 xc32-g++ (32-bit C++ Compiler)
   - 5.4.5 xc32-ld (32-Bit Linker)

5.4.1   XC32 (Global Options)

Set up global options for all 32-bit language tools. See also 5.4.6 Options Page Features.

Table 5-1. XC32 (Global Options) All Options Category

Option	Description	Command Line
Use legacy lib	Check this box to link with the recommended Legacy Standard C Library. Uncheck it to link with the non-legacy Standard C Library (HTC).	`-legacy-libc` and `-no-legacy-libc`
Don't delete intermediate files	Don't delete intermediate Files. Place them in the object directory and name them based on the source file.	`-save-temps=obj`
Use Whole-Program and Link-Time Optimizations	When this feature is enabled, the build will be constrained in the following ways: - The per-file build settings will be ignored - The build will no longer be an incremental one (full build only)	`-fwhole-program-flto`
Use GP relative addressing threshold	Put definitions of externally-visible data in a small data section if that data is no bigger than `num` bytes.	`-G num`
Common include dirs	Directory paths entered here will be appended to the already existing include paths of the compiler. Relative paths are from the MPLAB X IDE project directory.	`-I"dir"`

5.4.2   xc32-as (32-bit Assembler)

A subset of command-line options may be specified in MPLAB X IDE. Select a category, and then set up assembler options. For additional options, see MPLAB XC32 Assembler documentation. See also 5.4.6 Options Page Features.

Table 5-2. XC32-AS General Options Category

Option	Description	Command Line
Have symbols in production build	Generate debugging information for source-level debugging in MPLAB X.	`--gdwarf-2`
Keep local symbols	Check to keep local symbols, that is, labels beginning with `.L` (upper case only). Uncheck to discard local symbols.	`--keep-locals`
Exclude floating-point library	Exclude support for floating-point operations reducing code size for applications that do not require floating-point support.	`-mno-float`
Preprocessor macro definitions	Project-specific preprocessor macro defines passed via the compiler's `-D` option.	
Assembler symbols	Define symbol 'sym' to a given 'value'.	`--defsym sym=value`
Preprocessor Include directories	Relative paths are from MPLAB X project directory.	
Assembler Include directories	Relative paths are from MPLAB X project directory. Add a directory to the list of directories the assembler searches for files specified in `.include` directives. You may add as many directories as necessary to include a variety of paths. The current working directory is always searched first and then `-I` directories in the order in which they were specified (left to right) here.	`-I`

Table 5-3. XC32-AS Other Options Category

Option	Description	Command Line

Diagnostics level	Select warnings to display in the Output window.
	- Generate warnings
	- Suppress warnings
	- Fatal warnings
		`--warn`
		`--no-warn`
		`--fatal-warnings`

Include source code	Check for a high-level language listing. High-level listings require that the assembly source code is generated by a compiler, a debugging option like `-g` is given to the compiler, and assembly listings (`-al`) are requested. Uncheck for a regular listing.	`-ah`
Expand macros	Check to expand macros in a listing. Uncheck to collapse macros.	`-am`
Include false conditionals	Check to include false conditionals (`.if`, `.ifdef`) in a listing. Uncheck to omit false conditionals.	`-ac`

..........continued

Option	Description	Command Line
Omit forms processing	Check to turn off all forms processing that would be performed by the listing directives `.psize, .eject, .title` and `.sbttl`. Uncheck to process by listing directives.	`-an`
Include assembly	Check for an assembly listing. This `-a` suboption may be used with other suboptions. Uncheck to exclude an assembly listing.	`-al`
List symbols	Check for a symbol table listing. Uncheck to exclude the symbol table from the listing.	`-as`
Omit debugging directives	Check to omit debugging directives from a listing. This can make the listing cleaner. Uncheck to included debugging directives.	`-ad`

List to file	Use this option if you want the listing for a file. The list file will have the same name as the asm file plus `.lst`.	`-a=$`
		`{CURRENT_QUOTED_IF_SPA`
		`CED_OBJECT_FILE_MINUS_`
		`EXTENSION}.lst`

5.4.3 xc32-gcc (32-bit C Compiler)

Although the MPLAB XC32 C/C++ Compiler works with MPLAB X IDE, it must be acquired separately. The full version may be purchased, or a student (limited-feature) version may be downloaded for free. See the Microchip website (www.microchip.com) for details.

A subset of command-line options may be specified in MPLAB X IDE. Select a category, and then set up compiler options. For additional options, see the *MPLAB X IDE User's Guide* (DS52027/DS50002027), also available on the Microchip website.

See also 5.4.6 Options Page Features.

Table 5-4. XC32-GCC General Category

Option	Description	Command Line
Have symbols in production build	Build for debugging in a production build image.	`-g`
Enable App IO	Support the APPIN/APPOUT debugging feature with REAL ICE	`-mappio-debug`

.........continued		
Option	**Description**	**Command Line**
Isolate each function in a section	This option is often used with the linker's `--gc-sections` option to remove unreferenced functions. Check to place each function into its own section in the output file. The name of the function determines the section's name in the output file. Note: When you specify this option, the assembler and linker may create larger object and executable files and will also be slower. Uncheck to place multiple functions in a section.	`-ffunction-sections`
Place data into its own section	This option is often used with the linker's `--gc-sections` option to remove unreferenced statically-allocated variables. Place each data item into its own section in the output file. The name of the data item determines the name of the section. When you specify this option, the assembler and linker may create larger object and executable files and will also be slower.	`-fdata-sections`
Use indirect calls	Enable full-range calls.	`-mlong-calls`
Generate 16-bit code	By default, generate code for the MIPS16 instruction set, reducing code size.	`-mips16`
Exclude floating-point library	Exclude support for floating-point operations reducing code size for applications that do not require floating-point support.	`-mno-float`

Table 5-5. XC32-GCC Optimization Category

Option	Description	Command Line
Optimization Level	Select an optimization level. Your compiler edition may support only some optimizations. Equivalent to `-On` option, where `n` is an option below: 0 - Do not optimize.The compiler's goal is to reduce the cost of compilation and to make debugging produce the expected results. 1 - Optimize. Optimizing compilation takes somewhat longer, and a lot more host memory for a large function. The compiler tries to reduce code size and execution time. 2 - Optimize even more. The compiler performs nearly all supported optimizations that do not involve a space-speed trade-off. 3 - Optimize yet more favoring speed (superset of O2). s - Optimize yet more favoring size (superset of O2).	`-On`
Unroll loops	This option often increases execution speed at the expense of larger code size. Check to perform the optimization of loop unrolling. This is only done for loops whose number of iterations can be determined at compile time or run time. Uncheck to not unroll loops.	`-funroll-loops`
Omit frame pointer	Check to not keep the Frame Pointer in a register for functions that don't need one. Uncheck to keep the Frame Pointer.	`-fomit-frame-pointer`
Pre-optimization instruction scheduling	Default for optimization level: - Disable - Enable	`-fno-schedule-insns` `-fschedule-insns`
Post-optimization instruction scheduling	Default for optimization level: - Disable - Enable	`-fno-schedule-insns2` `-fschedule-insns2`

Table 5-6. XC32-GCC Preprocessing and Messages Category

Option	Description	Command Line
Preprocessor macros	Project-specific preprocessor macro defines passed via the compiler's -D option.	
Include directories	Search these directories for project-specific include files.	
Make warnings into errors	Check to halt compilation based on warnings as well as errors. Uncheck to halt compilation based on errors only.	-Werror
Additional warnings	Check to enable all warnings. Uncheck to disable warnings.	-Wall
support-ansi	Check to issue all warnings demanded by strict ANSI C. Uncheck to issue all warnings.	-ansi
strict-ansi	Issue all the warnings demanded by strict ISO C and ISO C++; reject all programs that use forbidden extensions, and some other programs that do not follow ISO C and ISO C++.	-pedantic
Use CCI syntax	Enable support for the CCI syntax (see 3. Common C Interface).	-mcci
Use IAR syntax	Enable support for syntax used by other toolchain vendors.	-mext=IAR

5.4.4 xc32-g++ (32-bit C++ Compiler)

A subset of command-line options may be specified in MPLAB X IDE. Select a category, and then set up linker options. For additional options, see MPLAB Object Linker for 32-bit Devices documentation. See also 5.4.6 Options Page Features.

Table 5-7. XC32-G++ C++ Specific Category

Option	Description	Command Line
Generate run time type descriptor information	Enable generation of information about every class with virtual functions for use by the C++ runtime type identification features ('dynamic_cast' and 'typeid'). If you don't use those parts of the language, you can save some space by disabling this option. Note that exception handling uses the same information, but it will generate it as needed. The 'dynamic_cast' operator can still be used for casts that do not require runtime type information, that is, casts to void * or to unambiguous base classes.	-frtti
Enable C++ exception handling	Enable exception handling. Generates extra code needed to propagate exceptions.	-fexceptions

..........continued		
Option	**Description**	**Command Line**
Check that the pointer returned by operator 'new' is non-null	Check that the pointer returned by operator new is non-null before attempting to modify the storage allocated.	`-fcheck-new`
Generate code to check for violation of exception specification	Generate code to check for violation of exception specifications at runtime. Using this option may increase code size in production builds.	`-fenforce-eh-specs`

Table 5-8. XC32-G++ General Category

Option	**Description**	**Command Line**
Have symbols in production build	Build for debugging in a production build image.	`-g`
Enable App IO	Place each function into its own section in the output file if the target supports arbitrary sections. The name of the function or the name of the data item determines the section's name in the output file. This option is useful when combined with the linker's `--gc-sections` option to remove unreferenced functions.	`-ffunction-sections`
Place data into its own section	Place each data item into its own section in the output file if the target supports arbitrary sections. The name of the function or the name of the data item determines the section's name in the output file. This option is useful when combined with the linker's `--gc-sections` option to remove unreferenced variables.	`-fdata-sections`
Use indirect calls	Enable full-range calls.	`-mlong-calls`
Generate 16-bit code	By default, generate code for the MIPS16 instruction set, reducing code size.	`-mips16`
Exclude floating-point library	Exclude support for floating-point operations reducing code size for applications that do not require floating-point support.	`-mno-float`

Table 5-9. XC32-G++ Optimization Category

Option	Description	Command Line
Optimization Level	Select an optimization level. Your compiler edition may support only some optimizations. Equivalent to $-On$ option, where n is an option below: 0 - Do not optimize. The compiler's goal is to reduce the cost of compilation and to make debugging produce the expected results. 1 - Optimize. Optimizing compilation takes somewhat longer, and a lot more host memory for a large function. The compiler tries to reduce code size and execution time. 2 - Optimize even more. The compiler performs nearly all supported optimizations that do not involve a space-speed trade-off. 3 - Optimize yet more favoring speed (superset of O2). s - Optimize yet more favoring size (superset of O2).	`-On`
Unroll loops	Check to perform the optimization of loop unrolling. This is only done for loops whose number of iterations can be determined at compile time or run time. Uncheck to not unroll loops.	`-funroll-loops`
Omit frame pointer	Check to not keep the Frame Pointer in a register for functions that don't need one. Uncheck to keep the Frame Pointer.	`-fomit-frame-pointer`

Pre-optimization instruction scheduling	Default for optimization level:
	- Disable
	
	- Enable
		`-fno-schedule-insns`
		
		`-fschedule-insns`

Post-optimization instruction scheduling	Default for optimization level:
	- Disable
	
	- Enable
		`-fno-schedule-insns2`
		
		`-fschedule-insns2`

Table 5-10. XC32-G++ Preprocessing and Messages Category

Option	Description	Command Line
Preprocessor macros	Project-specific preprocessor macro defines passed via the compiler's $-D$ option.	
Include directories	Search these directories for project-specific include files.	
Make warnings into errors	Check to halt compilation based on warnings as well as errors. Uncheck to halt compilation based on errors only.	`-Werror`
Additional warnings	Check to enable all warnings. Uncheck to disable warnings.	`-Wall`

..........continued		
Option	**Description**	**Command Line**
support-ansi	Check to issue all warnings demanded by strict ANSI C. Uncheck to issue all warnings.	`-ansi`
strict-ansi	Issue all the warnings demanded by strict ISO C and ISO C++; reject all programs that use forbidden extensions, and some other programs that do not follow ISO C and ISO C++.	`-pedantic`
Use CCI syntax	Enable support for the CCI syntax (3. Common C Interface).	`-mcci`
Use IAR syntax	Enable support for syntax used by other toolchain vendors.	`-mext=IAR`

5.4.5   xc32-ld (32-Bit Linker)

A subset of command-line options may be specified in MPLAB X IDE. Select a category, and then set up linker options. For additional options, see MPLAB Object Linker for 32-bit Devices documentation. See also 5.4.6 Options Page Features.

Table 5-11. XC32-LD General Category

Option	**Description**	**Command Line**
Heap Size (bytes)	Specify the size of the heap in bytes. Allocate a run-time heap of size bytes for use by C programs. The heap is allocated from unused data memory. If not enough memory is available, an error is reported.	`--defsym=_min_heap_size=<size>`
Minimum stack size (bytes)	Specify the minimum size of the stack in bytes. By default, the linker allocates all unused data memory for the run-time stack. Alternatively, the programmer may allocate the stack by declaring two global symbols: `__SP_init` and `__SPLIM_init`. Use this option to ensure that at least a minimum sized stack is available. The actual stack size is reported in the link map output file. If the minimum size is not available, an error is reported.	`--defsym=_min_stack_size=<size>`
Allow overlapped sections	Check to not check section addresses for overlaps. Uncheck to check for overlaps.	`--check-sections` `--no-check-sections`
Remove unused sections	Check to not enable garbage collection of unused input sections (on some targets). Uncheck to enable garbage collection.	`--no-gc-sections` `--gc-sections`
Use response file to link	Pass linker options in a file rather than on the command line. On Windows systems, this option allows you to properly link projects with a large number of object files that would normally overrun the command-line length limitation of the Windows OS.	`True`
Additional driver options	Type here any additional driver options not existing in this GUI otherwise. The string you introduce here will be emitted as is in the driver invocation command.	

Table 5-12. XC32-LD Fill Flash Memory Category

Option	Description	Command Line
Which areas to fill	Specify which area of Flash memory to fill. **No Fill** - None (default). **Fill All Unused** - Fill all unused memory. **Provide Range to fill** - Fill a range of memory. Enter a range under "Memory Address Range".	
How to fill it	Specify how to fill Flash memory. **Provide sequence of values** - provide a sequence under the Sequence option. **Constant or incrementing value** - provide a constant, increment/decrement or increment/decrement constant under the same-named option.	
Sequence	When **Provide sequence of values** is selected, enter a sequence. The form is n1, n2, ... where n1 uses C syntax. Example: 0x10, 25, 0x3F, 16.	`--fill=sequence`
Constant	When **Constant or incrementing value** is selected, enter a constant. Specify the constant using C syntax (for example, 0x for hex, 0 for octal). Example: 0x10 is the same as 020 or 16.	`--fill=constant`
Increment/Decrement	When **Constant or incrementing value** is selected, you can select to increment or decrement the initial value of "Constant" on each consecutive address. **No Incrementing** - do not change constant value. **Increment Const** - increment the constant value by the amount specified under the option "Increment/Decrement Constant". **Decrement Const** - decrement the constant value by the amount specified under the option "Increment/Decrement Constant."	
Increment/Decrement Constant	When **Increment Const** or **Decrement Const** is selected, enter a constant increment or decrement value. Specify the constant using C syntax (for example, 0x for hex, 0 for octal). Example: 0x10 is the same as 020 or 16.	`--fill=constant+=incr` `--fill=constant-=decr`
Memory Address Range	When **Provide Range to fill** is selected, enter the range here. Specify range as Start:End where Start and End use C syntax. Example 0x100:0x1FF is the same as 256:511.	`--fill=value@range`

Table 5-13. XC32-LD Libraries Category

Option	Description	Command Line
Optimization level of Standard Libraries	Select an optimization level. Your compiler edition may support only some optimizations. Equivalent to $-On$ option, where n is an option below: 0 - Do not optimize. The compiler's goal is to reduce the cost of compilation and to make debugging produce the expected results. 1 - Optimize. Optimizing compilation takes somewhat longer, and a lot more host memory for a large function. The compiler tries to reduce code size and execution time. 2 - Optimize even more. The compiler performs nearly all supported optimizations that do not involve a space-speed trade-off. 3 - Optimize yet more favoring speed (superset of O2). s - Optimize yet more favoring size (superset of O2).	`-On`
System Libraries	Add libraries to be linked with the project files. You may add more than one.	`--library=name`
Library directories	Add a library directory to the library search path. You may add more than one.	`--library-path="name"`
Exclude Standard Libraries	Check to not use the standard system startup files or libraries when linking. Only use library directories specified on the command line. Uncheck to use the standard system startup files and libraries.	`-nostdlib`
Do no link startup code	Exclude the default startup code because the project provides application-specific startup code.	`-nostartfiles`
Generate 16-bit code	Link the libraries precompiled for the MIPS16 instruction set, reducing code size.	`-mips16`
Exclude floating-point library	Exclude support for floating-point operations reducing code size for applications that do not require floating-point support.	`-mno-float`

Table 5-14. XC32-LD Diagnostics Category

Option	Description	Command Line
Generate map file	Create a map file.	`-Map="file"`
Display memory usage	Check to print memory usage report. Uncheck to not print a report.	`--report-mem`
Generate cross-reference file	Check to create a cross-reference table. Uncheck to not create this table.	`--cref`

..........continued

Option	Description	Command Line
Warn on section realignment	Check to warn if start of section changes due to alignment. Uncheck to not warn.	`--warn-section-align`
Trace Symbols	Add/remove trace symbols.	`--trace-symbol=symbol`

Table 5-15. XC32-LD Symbols and Macros Category

Option	Description	Command Line
Linker symbols	Create a global symbol in the output file containing the absolute address (`expr`). You may use this option as many times as necessary to define multiple symbols in the command line. A limited form of arithmetic is supported for the `expr` in this context: you may give a hexadecimal constant or the name of an existing symbol, or use + and - to add or subtract hexadecimal constants or symbols.	`--defsym=sym`
Preprocessor macro definitions	Add linker macros.	`-Dmacro`

Symbols	Specify symbol information in the output.
	- Keep all
	
	- Strip debugging info
	
	- Strip all symbol info
		`—`
		
		`--strip-debug (-S)`
		
		`--strip-all (-s)`

5.4.6 Options Page Features

The Options section of the Properties page has the following features for all tools:

Table 5-16. Page Features Options

Reset	Reset the page to default values.
Additional options	Enter options in a command-line (non-GUI) format.
Option Description	Click on an option name to see information on the option in this window. Not all options have information in this window.
Generated Command Line	Click on an option name to see the command-line equivalent of the option in this window.

5.5 Project Example

In this example, you will create an MPLAB X IDE project with two C code files.

5.5.1 Run the Project Wizard

In MPLAB X IDE, select *File>New Project* to launch the wizard.

1. **Choose Project:** Select "Microchip Embedded" for the category and "Standalone Project" for the project. Click Next> to continue.
2. **Select Device:** Select the dsPIC30F6014. Click **Next>** to continue.
3. **Select Header:** There is no header for this device so this is skipped.

4. **Select Tool:** Choose a development tool from the list. Tool support for the selected device is shown as a colored circle next to the tool. Mouse over the circle to see the support as text. Click **Next>** to continue.

5. **Select Compiler:** Choose a version of the XC32 toolchain. Click **Next>** to continue.

6. **Select Project Name and Folder:** Enter a project name, such as `MyXC32Project`. Then select a location for the project folder. Click **Finish** to complete the project creation and setup.

Once the Project Wizard has completed, the Project window should contain the project tree. For more on projects, see the MPLAB X IDE documentation.

5.5.2 Set Build Options

Select *File>Project Properties* or right click on the project name and select "Properties" to open the Project Properties dialog.

1. Under "Conf:[default]>XC32 (Global Options)", select "xc32-gcc."
2. Under "Conf:[default]>XC32 (Global Options)", select "xc32-ld."
3. Select "Diagnostics" from the "Option Categories." Then enter a file name to "Generate map file," that is, `example.map`.
4. Click the **OK** button on the bottom of the dialog to accept the build options and close the dialog.

5.5.3 Build the Project

Right-click on the project name, "MyXC32Project," in the project tree and select "Build" from the pop-up menu. The Output window displays the build results.

If the build did not complete successfully, check these items:

1. Review the previous steps in this example. Make sure you have set up the language tools correctly and have all the correct project files and build options.
2. If you modified the sample source code, examine the Build tab of the Output window for syntax errors in the source code. If you find any, click on the error to go to the source code line that contains that error. Correct the error, and then try to build again.

5.5.4 Output Files

View the project output files by opening the files in MPLAB X IDE.

1. Select *File>Open File*. In the Open dialog, find the project directory.
2. Under "Files of type" select "All Files" to see all project files.
3. Select *File>Open File*. In the Open dialog, select "example.map." Click **Open** to view the linker map file in an MPLAB X IDE editor window. For more on this file, see the linker documentation.
4. Select *File>Open File*. In the Open dialog, return to the project directory and then go to the *dist>default>production* directory. Notice that there is only one hex file, "MyXC32Project.X.production.hex." This is the primary output file. Click **Open** to view the hex file in an MPLAB X IDE editor window. For more on this file, see the Utilities documentation.

   There is also another file, "MyXC32Project.X.production.elf." This file contains debug information and is used by debug tools to debug your code. For information on selecting the type of debug file, see 5.4.1 XC32 (Global Options).

5.5.5 Further Development

Usually, your application code will contain errors and not work the first time. Therefore, you will need a debug tool to help you develop your code. Using the output files previously discussed, several debug tools exist that work with MPLAB X IDE to help you do this. You may choose from simulators, in-circuit emulators or in-circuit debuggers, either manufactured by Microchip Technology or third-party developers. Please see the documentation for these tools to learn how they can help you. When debugging, you will use *Debug>Debug Project* to run and debug your code. Please see MPLAB X IDE documentation for more information.

Once you have developed your code, you will want to program it into a device. Again, there are several programmers that work with MPLAB X IDE to help you do this. Please see the documentation for these tools to see how they can help you. When programming, you will use "Make and Program Device Project" button on the debug toolbar. Please see MPLAB X IDE documentation concerning this control.

6.   Command-line Driver

The MPLAB XC32 C Compiler command-line driver, `xc32-gcc`, can be invoked to perform all aspects of compilation, including C code generation, assembly and link steps. Its use is the recommended way to invoke the compiler, as it hides the complexity of all the internal applications and provides a consistent interface for all compilation steps. Even if an IDE is used to assist with compilation, the IDE will ultimately call `xc32-gcc`.

If you are developing a project that contains C++ source code, an alternate driver called `xc32-g++` is supplied and that will link in an alternate set of libraries. Its operation is similar to the `xc32-gcc` driver.

This chapter describes the steps that the driver takes during compilation, the files that the driver can accept and produce, as well as the command-line options that control the compiler's operation.

6.1   Invoking The Compiler

This section explains how to invoke `xc32-gcc` on the command line and discusses the input files that can be passed to the compiler.

Environment variables that can be set to specify certain aspects of the compiler's behavior are also explained.

6.1.1   Driver Command-line Format

The `xc32-gcc` driver can be used to compile and assemble C and assembly source files, as well as link object files and library archives to form a final program image.

The `xc32-g++` driver must instead be used when the module source is written in C++.

The driver has the following basic command format:

```
xc32-gcc [options] files
```

So, for example, to compile and link the C source file `hello.c`, you could use the command:

```
xc32-gcc -mprocessor=32MZ2048ECH100 -O2 -o hello.elf hello.c
```

The format for the C++ driver is similar.

```
xc32-g++ [options] files
```

And this driver is used in a similar way, for example:

```
xc32-g++ -mprocessor=32MZ2048ECH100 -O2 -o hello.elf hello.cpp
```

Throughout this manual, it is assumed that the compiler applications are in your console's search path. See 6.1.2 Environment Variables for information on the environment variable that specifies the search locations. Alternatively, use the full directory path along with the driver name when executing the compiler.

It is customary to declare *options* (identified by a leading dash "-" or double dash "--") before the files' names; however, this is not mandatory.

Command-line options are case sensitive, with their format and description being supplied in 6.7 Driver Option Descriptions. Many of the command-line options accepted by `xc32-gcc` are common to all the MPLAB XC compilers, to allow greater portability between devices and compilers.

The *files* can be any mixture of C/C++ and assembler source files, as well as relocatable object files and archive files. While the order in which these files are listed does not directly affect the operation of the program, it can affect the allocation of code or data. Note, that the order of the archive files will dictate the order in which they are searched, and in some situations, this might affect which modules are linked in to the program.

It is recommended that C-only projects use the `xc32-gcc` driver. If you have C++ source code or a mix of C and C++ source, always use `xc32-g++` driver to ensure the correct libraries are linked in.

6.1.2   Environment Variables

The environment variables in this section are optional, but, if defined, they will be used by the compiler. The compiler driver, or other subprogram, may choose to determine an appropriate value for some of the following environment variables if they are not set. The driver, or other subprogram, takes advantage of internal knowledge about the installation of the compiler. As long as the installation structure remains intact, with all subdirectories and executables remaining in the same relative position, the driver or subprogram will be able to determine a usable value.

The XC32-style variables should be used for new projects; however, the PIC32-style variables may be used for legacy projects.

Table 6-1. Compiler-Related Environment Variables

Variable	Description
`XC32_C_INCLUDE_PATH` or `PIC32_C_INCLUDE_PATH`	This variable's value is a semicolon-separated list of directories, much like your host terminal's `PATH` environment variable. When the compiler searches for header files, it tries the directories listed in the variable, after the directories specified with `-I` but before the standard header file directories. If the environment variable is undefined, the preprocessor chooses an appropriate value based on the standard installation. By default, the following directories are searched for the following include files: `<install-path>\pic32mx\include`
`XC32_COMPILER_PATH` or `PIC32_COMPILER_PATH`	The value of this variable is a semicolon-separated list of directories, much like your host terminal's `PATH` environment variable. The compiler tries the directories thus specified when searching for subprograms, if it can't find the subprograms using `PIC32_EXEC_PREFIX`
`XC32_EXEC_PREFIX` or `PIC32_EXEC_PREFIX`	If this environment variable is set, it specifies a prefix to use in the names of subprograms executed by the compiler. No directory delimiter is added when this prefix is combined with the name of a subprogram, but you can specify a prefix that ends with a slash if you wish. If the compiler cannot find the subprogram using the specified prefix, it tries looking in your `PATH` environment variable. If the environment variable is not set or set to an empty value, the compiler driver chooses an appropriate value based on the standard installation. If the installation has not been modified, this will result in the driver being able to locate the required subprograms. Other prefixes specified with the `-B` command line option take precedence over the user- or driver-defined value of the variable. Under normal circumstances it is best to leave this value undefined and let the driver locate subprograms itself.

`XC32_LIBRARY_PATH` or `PIC32_LIBRARY_PATH`	This variable's value is a semicolon-separated list of directories, much like `PATH`. This variable specifies a list of directories to be passed to the linker. The driver's default evaluation of this variable is:
	
	`<install-path>\lib;`
	
	`<install-path>\pic32mx\lib`

`TMPDIR`	If this variable is set, it specifies the directory to use for temporary files. The compiler uses temporary files to hold the output of one stage of compilation that is to be used as input to the next stage: for example, the output of the preprocessor, which is the input to the compiler proper.

6.1.3   Input File Types

The `xc32-gcc` driver distinguishes source files, intermediate files and library files solely by the file type, or extension. The following case-sensitive extensions, listed below are recognized.

Table 6-2. Input File Types

Extension	File format
.c	C source file
.i	Preprocessed C source file
.cpp	C++ source file
.ii	Preprocessed C++ source file
.s	Assembler source file
.S	Assembly source file requiring preprocessing
.o	Relocatable object code file
.a	Archive (library) file
other	A file to be passed to the linker

There are no compiler restrictions imposed on the names of source files, but be aware of case, name-length, and other restrictions that are imposed by your host operating system.

Avoid using the same base name for assembly and C/C++ source files, even if they are located in different directories. So, for example, if a project contains a C source file called `init.c`, do not also add to the project an assembly source file with the name `init.s`. Avoid, also, having source files with the same base name as the MPLAB X IDE project name.

The terms *source file* and *module* are often used interchangeably, but they refer to the source code at different points in the compilation sequence.

A source file is a file that contains all or part of a program. It may contain C/C++ code, as well as preprocessor directives and commands. Source files are initially passed to the preprocessor by the compiler driver.

A module is the output of the preprocessor for a given source file, after the inclusion of any header files specified by `#include` preprocessor directives, and after the processing and subsequent removal of other preprocessor directives (with the possible exception of some commands for debugging). Thus, a module is usually the amalgamation of a source file and several header files, and it is this output that is passed to the remainder of the compiler applications. A module is also referred to as a *translation unit*.

Like assembly source files, these terms can also be applied to assembly files, which can be preprocessed and can include other header files.

6.2   The C Compilation Sequence

When you compile a project, many internal applications are called by the driver to do the work. This section introduces these internal applications and describes how they relate to the build process, especially when a project consists of multiple source files. This information should be of particular interest if you are using a make system to build projects.

6.2.1   Single-Step C Compilation

Full compilation of one or more C source files, including the link step, can be performed in just one command using the `xc32-gcc` driver.

6.2.1.1   Compiling a Single C File

The following is a simple C program that adds two numbers. To illustrate how to compile and link a program consisting of a single C source file, copy the code into any text editor and save it as a plain text file with the name `ex1.c`.

```
#include <xc.h>

unsigned int
```

```
add(unsigned int a, unsigned int b)
{
  return a + b;
}

int
main(void)
{
  unsigned int x, y, z;
  x = 2;
  y = 5;
  z = add(x, y);

  return 0;
}
```

In the interests of clarity, this code does not specify device configuration bits, nor has any useful purpose.

Compile the program by typing the following command at the prompt in your favorite terminal. For the purpose of this discussion, it is assumed that in your terminal you have changed into the directory containing the source file you just created, and that the compiler is installed in the standard directory location and is in your host's search path.

```
xc32-gcc -mprocessor=32MZ2048ECH100  -o ex1.elf ex1.c
```

This command compiles the `ex1.c` source file for a 32MZ2048ECH100 device and has the output written to `ex1.elf`, which may be loaded into the MPLAB X IDE.

If a hex file is required, for example, to load into a device programmer, then use the following command:

```
xc32-bin2hex ex1.elf
```

This creates an Intel hex file named `ex1.hex`.

The driver will compile the source file, regardless of whether it has changed since the last build command. Development environments (such as MPLAB X IDE) and make utilities must be employed to achieve incremental builds (see 6.2.2 Multi-Step C Compilation).

Unless otherwise specified, an ELF file (this is by default called `a.out`) is produced as the final output.

6.2.1.2   Compiling Multiple C Files

This section demonstrates how to compile and link a project, in a single step, that consists of multiple C source files.

Copy the example code shown into a text file called `add.c`.

```
/* add.c */
#include <xc.h>

unsigned int
add(unsigned int a, unsigned int b)
{
  return a + b;
}
```

And place the following code in another file, `ext.c`.

```
/* ex1.c */
#include <xc.h>

unsigned int add(unsigned int a, unsigned int b);

int
main(void) {
  unsigned int x, y, z;
  x = 2;
  y = 5;
  z = add(x, y);

  return 0;
}
```

In the interests of clarity, this code does not specify device configuration bits, nor has any useful purpose.

Compile both files by typing the following at the prompt:

```
xc32-gcc -mprocessor=32MZ2048ECH100  -o ex1.elf ex1.c add.c
```

This command compiles the modules `ex1.c` and `add.c` in the one step. The compiled modules are linked with the relevant compiler libraries and the executable file `ex1.elf` is created.

6.2.2   Multi-Step C Compilation

A multi-step compilation method can be employed to build projects consisting of one or more C source files. Make utilities can use this feature, taking note of which source files have changed since the last build to speed up compilation. Incremental builds are also be performed by integrated development environments, such as the MPLAB X IDE when selecting the Build Project icon or menu item.

Make utilities typically call the compiler multiple times: once for each source file to generate an intermediate file and once to perform the second stage compilation, which links the intermediate files to form the final output. If only one source file has changed since the last build, the intermediate file corresponding to the unchanged source file need not be regenerated.

For example, the files `ex1.c` and `add.c` are to be compiled using a make utility. The command lines that the make utility could use to compile these files might be something like:

```
xc32-gcc -mprocessor=32MZ2048ECH100 -c ex1.c
xc32-gcc -mprocessor=32MZ2048ECH100 -c add.c
```

```
xc32-gcc -mprocessor=32MZ2048ECH100  -o ex1.elf ex1.o add.o
```

The `-c` option used with the first two commands will compile the specified file into the intermediate file format, but not link. The resultant intermediate files are linked in the final step to create the final output `ex1.elf`. All the files that constitute the project must be present when performing the second stage of compilation.

The above example uses the command-line driver, `xc32-gcc`, to perform the final link step. You can explicitly call the linker application, `pic32m-ld`, but this is not recommended as the commands are complex and when driving the linker application directly, you must specify linker options, not driver options, as shown above.

For more information on using the linker, see the MPLAB® XC32 Assembler, Linker and Utilities User's Guide relevant to your project.

You may also wish to generate intermediate files to construct your own library archive files.

See MPLAB® XC32 Assembler, Linker and Utilities User's Guide relevant to your project for more information on library creation.

6.3   The C++ Compilation Sequences

When you compile a project, many internal applications are called by the driver to do the work. This section introduces these internal applications and describes how they relate to the build process, especially when a project consists of multiple source files. This information should be of particular interest if you are using a make system to build projects.

6.3.1   Single-step C++ Compilation

A single command-line instruction can be used to compile one file or multiple files.

6.3.1.1   Compiling a Single C++ File

The following is a simple C++ program. To illustrate how to compile and link a program consisting of a single C++ source file, copy the code into any text editor and save it as a plain text file with the name `ex2.cpp`.

```
/* ex2.cpp */
#include <xc.h>
#include <iostream>
#include <vector>
```

```
#include <deque>
#include <list>
#include <set>
#include <map>
#include <string>
#include <algorithm>
#include <iterator>
#include <functional>
#include <numeric>
using namespace std;
//Device - Specific Configuration - bit settings
#pragma config FPLLMUL=MUL_20, FPLLIDIV=DIV_2, FPLLODIV=DIV_1
#pragma config FWDTEN=OFF
#pragma config POSCMOD=HS, FNOSC=PRIPLL, FPBDIV=DIV_8
template < class T >
inline void
print_elements(const T & coll, const char *optcstr = "") {
    typename T::const_iterator pos;
    std::cout << optcstr;
    for (pos = coll.begin(); pos != coll.end(); ++pos) {
        std::cout << *pos << ' ';
    }
    std::cout << std::endl;
}
template < class T >
inline void
insert_elements(T & coll, int first, int last) {
    for (int i = first; i <= last; ++i) {
        coll.insert(coll.end(), i);
    }
}
int
main(void) {
    //Direct stdout to UART 1 for use with the simulator
    __XC_UART = 1;
    deque<int>coll;
    insert_elements(coll, 1, 9);
    insert_elements(coll, 1, 9);
    print_elements(coll, "on entry: ");
    // sortelements
    sort(coll.begin(), coll.end());
    print_elements(coll, "sorted: ");
    //sorted reverse
    sort(coll.begin(), coll.end(), greater < int >());
    print_elements(coll, "sorted >: ");
    while (1);
}
```

The first line of the program includes the header file `xc.h`, which provides definitions for all Special Function Registers (SFRs) on the target device. The second file of the program includes the header file, which provides the necessary prototypes for the peripheral library.

Compile the program by typing the following command at the prompt in your favorite terminal. For the purpose of this discussion, it is assumed that in your terminal you have changed into the directory containing the source file you just created, and that the compiler is installed in the standard directory location and is in your host's search path.

```
xc32-g++ -mprocessor=32MX795F512L -Wl,--defsym=_min_heap_size=0xF000 -o ex2.elf ex2.cpp
```

The option `-o ex2.elf` names the output executable file. This elf file may be loaded into MPLAB X IDE.

If a hex file is required, for example, to load into a device programmer, then use the following command

```
xc32-bin2hex ex2.elf
```

This creates an Intel hex file named `ex2.hex`.

6.3.2   Compiling Multiple C and C++ Files

This section demonstrates how to compile and link multiple C and C++ files in a single step.

File 1

```
/* main.cpp */
#include <xc.h>
#include <iostream>
using namespace std;
//Device - Specific Configuration - bit settings
#pragma config FPLLMUL=MUL_20, FPLLIDIV=DIV_2, FPLLODIV=DIV_1
#pragma config FWDTEN=OFF
#pragma config POSCMOD=HS, FNOSC=PRIPLL, FPBDIV=DIV_8
// add() must have C linkage
extern "C" {
    extern unsigned int add(unsigned int a, unsigned int b);
}
int main(void) {
    int myvalue = 6;
    //Direct stdout to UART 1 for use with the simulator
    __XC_UART = 1;
    std::cout << "original value: " << myvalue << endl;
    myvalue = add(myvalue, 3);
    std::cout << "new value:
    while (1);
}
```

File 2

```
/* ex3.c */
unsigned int
add(unsigned int a, unsigned int b)
{
return(a+b);
}
```

Compile both files by typing the following at the prompt:

```
xc32-g++ -mprocessor=32MX795F512L -o ex3.elf main.cpp ex3.c
```

The command compiles the modules `main.cpp` and `ex3.c`. The compiled modules are linked with the compiler libraries for C++ and the executable file `ex3.elf` is created.

Note: Use the xc32-g++ driver (as opposed to the xc32-gcc driver) in order to link the project with the C++ support libraries necessary for the C++ source file in the project.

6.4   Runtime Files

In addition to the C/C++ and assembly source files and user-defined libraries specified on the command line, the compiler can also link into your project compiler-generated source files and pre-compiled library files, whose content falls into the following categories:

- C/C++ standard library routines
- Implicitly called arithmetic library routines
- User-defined library routines
- The runtime start-up code

Note: Some PIC32 target devices allow you to select to boot in either the MIPS32® or microMIPS™ ISA mode via a device configuration bit (BOOTISA). On these devices, if your BOOTISA bit is set to microMIPS mode, pass the `-mmicromips` mode to the xc32-gcc/g++ compilation driver to tell it to link with the microMIPS variant of the runtime start-up code. If your BOOTISA bit is set to MIPS32 mode, pass the `-mno-micromips` option to the compilation driver so that the MIPS32 variant of the runtime start-up code is linked.

6.4.1   Location and Naming Convention

By default, the compiler uses the directory `<install-directory>/lib/gcc/` to store the specific libraries and the directory `<install-directory>/pic32mx/lib` to store the target-specific libraries, based on the target device family.

The target libraries that are distributed with the compiler are built for the corresponding command-line options:

- Size vs. speed (`-Os` vs. `-O3`)
- MIPS16 vs. MIPS32 vs. microMIPS ISA mode (`-mips16` vs. `-mno-mips16` vs. `-mmicromips`)
- Software floating-point vs no floating-point support (`-msoft-float` vs. `-mno-float`)

The following examples provide details on which of the library subdirectories are searched.

If no optimization option has been specified or an optimization of level 2 or lower has been specified, for example, `xc32-gcc -O1 foo.c`, then the libraries in the top-level of the library subdirectories are used.

If `-Os` level optimizations have been chosen, for example, `xc32-g++ -Os foo.cpp`, then the libraries in the `size` directories of the library subdirectories are used.

If `-Os` level optimizations and the MIPS16 instruction set option have both been chosen, for example, `xc32-gcc -Os -mips16 foo.c`, then the libraries in the `size/mips16` directories of the library subdirectories are searched.

6.4.2 Library Files

The names of the C/C++ standard library files appropriate for the selected target device, and other driver options, are determined by the driver.

The target libraries, called multilibs, are built multiple times with a permuted set of options. When the compiler driver is called to compile and link an application, the driver chooses the version of the target library that has been built with the same options.

By default, the 32-bit language tools use the directory `<install-directory>/lib/gcc/` to store the specific libraries and the directory `<install-directory>/pic32mx/lib` to store the target-specific libraries. Both of these directory structures contain subdirectories for each of the multilib combinations specified above.

The target libraries that are distributed with the compiler are built for the corresponding command-line options:

- Size vs speed (`-Os` vs. `-O3`)
- MIPS16 vs MIPS32 vs microMIPS ISA mode (`-mips16` vs. `-mno-mips16` vs `-mmicromips`)
- Software floating-point vs no floating-point support (`-msoft-float` vs `-mno-float`)

The following examples provide details on which of the multilibs subdirectories are chosen.

`xc32-gcc foo.c` `xc32-g++ foo.cpp`	For this example, no command line options have been specified (that is, the default command line options are being used). In this case, the `.` subdirectories are used.
`xc32-gcc -Os foo.c` `xc32-g++ -Os foo.cpp`	For this example, the command line option for optimizing for size has been specified (that is, `-Os` is being used). In this case, the `./size` subdirectories are used.
`xc32-gcc -O2 foo.c` `xc32-g++ -O2 foo.cpp`	For this example, the command line option for optimizing has been specified; however, this command line option optimizes for neither size nor space (that is, `-O2` is being used). In this case, the `.` subdirectories are used.
`xc32-gcc -Os -mips16 foo.c` `xc32-g++ -Os -mips16 foo.cpp`	For this example, the command line options for optimizing for size and for MIPS16 code have been specified (that is, `-Os` and `-mips16` are being used). In this case, the `./size/mips16` subdirectories are used.

6.4.2.1 Standard Libraries

The C/C++ standard libraries contain a standardized collection of functions, such as string, math and input/output routines. The range of these functions are described in 17. Library Routines.

These libraries also contain C/C++ routines that are implicitly called by the output code of the code generator. These are routines that perform tasks such as floating-point operations and that may not directly correspond to a C/C++ function call in the source code.

6.4.2.2 User-Defined Libraries

User-defined libraries may be created and linked in with programs as required. Library files are more easy to manage and may result in faster compilation times, but must be compatible with the target device and options for a particular project. Several versions of a library may need to be created to allow it to be used for different projects.

User-created libraries that should be searched when building a project can be listed on the command line along with the source files.

As with Standard C/C++ library functions, any functions contained in user-defined libraries should have a declaration added to a header file. It is common practice to create one or more header files that are packaged with the library file. These header files can then be included into source code when required.

6.4.3 Peripheral Library Functions

For PIC32MX devices only:

Many of the peripherals of the PIC32 devices are supported by the peripheral library functions provided with the compiler tools. Please refer to the MPLAB® Harmony Libraries for new projects. For legacy support, these PLIB Libraries will be available for download from `http://www.microchip.com/pic32_peripheral_lib`.

For All PIC32M devices:

MPLAB Harmony includes a set of peripheral libraries, drivers and system services that are readily accessible for application development. For access to the peripheral header files, go to the Microchip web site (www.microchip.com), click on the **Design Support** tab and download MPLAB Harmony and MPLAB Code Configurator. The path to the peripheral libraries is:

For Windows: `C:\microchip\harmony\<version>\framework\peripheral`

For Mac/Linux: `~\microchip\harmony\<version>\framework\peripheral`

6.4.4 Runtime Startup Code Location

The runtime startup code is contained in multiple versions of precompiled object file modules, to support architectural differences between device families.

For C programs, there is only one start-up module.

The source code for this module can be found in the `pic32m-libs.zip` file located at `<install-directory>/pic32-libs/`. Once the file is unzipped, the source code can be found at: `pic32m-libs/libpic32/startup/crt0.S`. It is precompiled into the following library location: `<install-directory>/pic32mx/lib/crt0.o`.

For C++ programs, code from five object files link sequentially to create a single initialization routine, which initializes the C++ runtime environment.

The PIC32M precompiled startup object modules are located in the following location: `<install-directory>/pic32mx/lib/`. The files have the names: `cpprt0.o`, `crti.o`, and `crtn.o`. The precompiled startup modules are located in the following location: `<install-directory>/lib/gcc/pic32mx/<gcc-version>/`. The files have the names: `crtbegin.o` and `crtend.o`.

6.4.5 Startup and Initialization

Note: Some PIC32 target devices allow you to select to boot in either the MIPS32® or microMIPS™ ISA mode via a device configuration bit (BOOTISA). On these devices, if your BOOTISA bit is set to microMIPS mode, pass the `-mmicromips` mode to the xc32-gcc/g++ compilation driver to tell it to link with the microMIPS variant of the runtime start-up code. If your BOOTISA bit is set to MIPS32 mode, pass the `-mno-micromips` option to the compilation driver so that the MIPS32 variant of the runtime start-up code is linked.

The runtime startup code performs initialization tasks that must be executed before the `main()` function in the C/C++ program is executed. For information on the tasks performed by this code, see 16. Main, Runtime Start-up and Reset.

The compiler will select the appropriate runtime startup code, based on the selected target device and other compiler options.

- The startup code initializes the L1 cache when available.

- It enables the DSPr2 engine when available.
- It also initializes the Translation Lookaside Buffer (TLB) of the Memory Management Unit (MMU) for the External Bus Interface (EBI) or Serial Quad Interface (SQI) when available. The device-specific linker script creates a table of TLB initialization values that the startup code uses to initialize the TLB at startup.

Important: When your target MCU is configured to use the microMIPS compressed ISA at startup (and for interrupts/exceptions), be sure to pass the `-mmicromips` option to xc32-gcc when linking, and use the `micromips` function attribute on all of your Interrupt Service Routines (ISRs). Using the `-mmicromips` option and the `micromips` attribute ensures that your startup code and ISR code are compiled for the microMIPS ISA when the **BOOTISA Configuration bit is set to micromips**. Likewise, be sure that you linkwith the MIPS32 startup code, and your ISRs are not `micromips` attributed when the `BOOTISA` bit is set to MIPS32.

For C:

There is only one start-up module, which initializes the C runtime environment.

The source code for this is found in the pic32m-libs.zip file located at:

`<install-directory>/pic32-libs/`

Once the file is unzipped, the source code can be found at:

`pic32m-libs/libpic32/startup/crt0.S.`

It is precompiled into the following library location:

`<install-directory>/pic32mx/lib/crt0.o.`

Multilib versions of these modules exist in order to support architectural differences between device families.

For C++:

Code from five object files link sequentially to create a single initialization routine, which initializes the C++ runtime environment.

The PIC32M precompiled startup objects are located in the following location:

`<install-directory>/pic32mx/lib/.`

The files have the following names: `cpprt0.o`, `crti.o`, and `crtn.o`.

The GCC precompiled startup objects are located in the following location:

`<install-directory>/lib/gcc/pic32mx/<gcc-version>/.`

The files have the following names: `crtbegin.o` and `crtend.o`.

Multilib variations of these modules exist in order to support architectural differences between device families and also optimization settings.

For more information about what the code in these start-up modules actual does, see 16.2 Runtime Start-Up Code.

6.5   Compiler Output

There are many files created by the compiler during the compilation. A large number of these are intermediate files and some are deleted after compilation is complete, but many remain and are used for programming the device, or for debugging purposes.

6.5.1   Output Files

The compilation driver can produce output files with the following extensions, which are case-sensitive.

Table 6-3. File Names

Extensions	Definition
file.hex	Executable file
file.elf	ELF debug file
file.o	Object file (intermediate file)
file.s	Assembly code file (intermediate file)
file.i	Preprocessed C file (intermediate file)
file.ii	Preprocessed C++ file (intermediate file)
file.map	Map file

The names of many output files use the same base name as the source file from which they were derived. For example the source file `input.c` will create an object file called `input.o`.

The main output file is an ELF file called `a.elf`, unless you override that name using the `-o` option.

If you are using an IDE, such as MPLAB X IDE, to specify options to the compiler, there is typically a project file that is created for each application. The name of this project is used as the base name for project-wide output files, unless otherwise specified by the user. However check the manual for the IDE you are using for more details.

Note: Throughout this manual, the term *project name* will refer to the name of the project created in the IDE.

The compiler is able to directly produce a number of the output file formats which are used by Microchip development tools.

The default behavior of `xc32-gcc` and `xc32-g++` is to produce an ELF output. To make changes to the file's output or the file names, see 6.7 Driver Option Descriptions.

6.5.2   Diagnostic Files

Two valuable files produced by the compiler are the assembly list file, produced by the assembler, and the map file, produced by the linker.

The assembly list file contains the mapping between the original source code and the generated assembly code. It is useful for information such as how C source was encoded, or how assembly source may have been optimized. It is essential when confirming if compiler-produced code that accesses objects is atomic, and shows the region in which all objects and code are placed.

The option to create a listing file in the assembler is `-a` (or `-Wa,-a` if passed to the driver). There are many variants to this option, which may be found in the "*MPLAB® XC32 Assembler, Linker and Utilities User's Guide*" (DS50002186). To pass the option from the compiler, see 6.7.9 Options for Assembling.

There is one list file produced for each build. There is one assembler listing file for each translation unit. This is a pre-link assembler listing so it will not show final addresses. Thus, if you require a list file for each source file, these files must be compiled separately, see 6.2.2 Multi-Step C Compilation. This is the case if you build using MPLAB IDE. Each list file will be assigned the module name and extension `.lst`.

The map file shows information relating to where objects were positioned in memory. It is useful for confirming that user-defined linker options were correctly processed, and for determining the exact placement of objects and functions.

The option to create a map file in the linker is `-Map file` (or `-Wl,-Map=file`, if passed to the driver), which can be found in the "*MPLAB® XC32 Assembler, Linker and Utilities User's Guide*" (DS50002186). To pass the option from the compiler, see 6.7.10 Options for Linking.

There is one map file produced when you build a project, assuming the linker was executed and ran to completion.

6.6   Compiler Messages

All compiler applications use textual messages to report feedback during the compilation process.

There are several types of messages, described below. The behavior of the compiler when encountering a message of each type is also listed.

Warning Messages	Indicates source code or other situations that can be compiled, but is unusual and might lead to runtime failures of the code. The code or situation that triggered the warning should be investigated; however, compilation of the current module will continue, as will compilation of any remaining modules.
Error Messages	Indicates source code that is illegal or that compilation of code cannot take place. Compilation will be attempted for the remaining source code in the current module (however the cause of the initial error might trigger further errors) and compilation of the other modules in the project will take place, but the project will not be linked.
Fatal Messages	Indicates a situation in which compilation cannot proceed and which forces the compilation process to stop immediately.

For information on options that control compiler output of errors, warnings or comments, see 6.7.4 Options for Controlling the C++ Dialect.

6.7   Driver Option Descriptions

Most aspects of the compilation process can be controlled using options passed to the command-line driver, `xc32-gcc`.

The GCC compiler on which the MPLAB XC32 C Compiler is based provides many options in addition to those discussed in this document. It is recommended that you avoid any option that has not been documented here, especially those that control the generation or optimization of code.

All single letter options are identified by a leading *dash* character, "`-`", for example, `-c`. Some single letter options specify an additional data field which follows the option name immediately and without any *whitespace*, for example, `-Idir`. Options are case sensitive, so `-c` is a different option to `-C`. All options are identified by single or double leading dash character, for example, `-c` or `--version`.

Use the `--help` option to obtain a brief description of accepted options on the command line.

If you are compiling from within the MPLAB X IDE, it will issue explicit options to the compiler that are based on the selections in the project's **Project Properties** dialog. The default project options might be different to the default options used by the compiler when running on the command line, so you should review these to ensue that they are acceptable.

6.7.1   Options Specific to PIC32M Devices

These options are specific to the PIC32M device, not the compiler.

Table 6-4. PIC32M Device-Specific Options

Option	Definition
`-G num`	Put global and static items less than or equal to *num* bytes into the small data or bss section instead of the normal data or bss section. This allows the data to be accessed using a single instruction. All modules should be compiled with the same `-G num` value.
`-mappio-debug`	Enable the APPIN/APPOUT debugging library functions for the Microchip debugger and in-circuit emulator. This feature allows you to use the DBPRINTF and related functions and macros as described in the "*32-bit Language Tool Libraries*" document (DS51685). Enable this option only when using a target PIC32 device that supports the APPIN/APPOUT feature.
`-mcci`	Enables the Microchip Common C Interface compilation mode.

..........continued	
Option	**Definition**
`-mcheck-zero-division` `-mno-check-zero-division`	Trap (do not trap) on integer division by zero. The default is `-mcheck-zero-division`.
`-membedded-data` `-mno-embedded-data`	Allocate variables to the read-only data section first if possible, then next in the small data section if possible, otherwise in data. This gives slightly slower code than the default, but reduces the amount of RAM required when executing, and thus may be preferred for some embedded systems.
`-mframe-header-opt`	Allows the compiler to omit a few instructions for each function that does not use its incoming frame header. This feature usually improves both execution speed and code size.
`-mgen-pie-static`	Generate position-independent code suitable for statically linking into a Position-Independent Executable (PIE). Such code access all constant addresses through a Global Offset Table (GOT). A special ELF loader, running on the target device, resolves the GOT entries and loads the final executable image into memory. Pass this option to the xc32-gcc compilation driver when compiling, assembling, and linking.
`-mno-hi-addr-opt`	The Special Function Register (SFR) Access Efficiency feature adds the address attribute to Peripheral SFRs defined in the processor header file. With this added information, a new compiler optimization reduces the number of registers required to access multiple SFRs within a single function. This also allows the compiler to remove redundant load instructions. Enabled by default at optimization levels -O2, -Os, & -O3. **Note:** If you are building a static library that accesses an SFR and you want that same prebuilt library to work across devices that may have the SFRs located at a different address (for example TMR1 is at different addresses on device A and device B), compile your library with the `-mno-hi-addr-opt` option. This will result in larger code, but the SFR address will be determined at link time.
`-minterlink-compressed`	Generate code that is link compatible with MIPS16 and microMIPS code.
`-mips16` `-mno-mips16`	Generate (do not generate) MIPS16 code. This is only available in the PRO edition. **Note:** On PIC32M devices, bit 0 of the program counter indicates the Instruction Set Architecture (ISA) mode. When this bit is clear, the device is running in MIPS32 mode. When this bit is set, the device is running in either MIPS16 or microMIPS mode, depending on the device core. This means that if you execute a hard-coded jump, bit 0 must be set to the appropriate value for your target function. Hard-coded jumps are most commonly seen when jumping from a bootloader to a bootloaded application.
`-mjals` `-mno-jals`	Generate (do not generate) 'jals' for microMIPS by recognizing that the branch delay slot instruction can be 16 bits. This implies that the function call cannot switch the current mode during the linking stage because we don't have 'jals' that supports 16-bit branch delay slot instructions. You man need to use `-mno-jals` if you have link errors when attempting to link a microMIPS object/library with a MIPS32 object/library.
`-mlong-calls` `-mno-long-calls`	Disable (do not disable) use of the `jal` instruction. Calling functions using `jal` is more efficient but requires the caller and callee to be in the same 256 megabyte segment. This option has no effect on abicalls code. The default is `-mno-long-calls`.
`-mmemcpy` `-mno-memcpy`	Force (do not force) the use of `memcpy()` for non-trivial block moves. The default is `-mno-memcpy`, which allows GCC to inline most constant-sized co-pies.

..........continued

Option	Definition
`-mmicromips` `-mno-micromips`	Generate (do not generate) microMIPS™ instructions. This feature is available only in the PRO edition. When your device is configured to boot to the microMIPS compressed ISA mode (for example,`#pragma config BOOTISA=MICROMIPS`), use the `-mmicromips` option when linking to specify the microMIPS startup code. **Note:** On PIC32M devices, bit 0 of the program counter indicates the Instruction Set Architecture (ISA) mode. When this bit is clear, the device is running in MIPS32 mode. When this bit is set, the device is running in either MIPS16 or microMIPS mode, depending on the device core. This means that if you execute a hard-coded jump, bit 0 must be set to the appropriate value for your target function. Hard-coded jumps are most commonly seen when jumping from a bootloader to a bootloaded application.
`-mno-float`	Do not use software floating-point libraries.
`-mno-peripheral-libs`	`-mno-peripheral-libs` is now the default. `-mperipheral-libs` is optional. By default, the peripheral libraries are linked specified via the device-specific linker script. Do not use the standard peripheral libraries when link-ing.
`-mprocessor`	Selects the device for which to compile. (for example,`-mprocessor=32MX360F512L`)
`-mreserve`	When building a project for debugging in MPLAB X IDE, the IDE passes the `-mreserve` option to the toolchain in order to reserve memory for use by the debug executive. This mechanism replaces the hard-coded reserved memory regions in the linker script.
`-msmart-io=[0|1|2]`	This option attempts to statically analyze format strings passed to `printf`, `scanf` and the `'f'` and `'v'` variations of these functions. Uses of nonfloating-point format arguments will be converted to use an integer-only variation of the library function. For many applications, this feature can reduce program-memory usage. `-msmart-io=0` disables this option, while `-msmart-io=2` causes the compiler to be optimistic and convert function calls with variable or unknown format arguments. `-msmart-io=1` is the default and will convert only when the compiler can prove that floating-point support is not required.

`-mtext="scn-name"`	Places text (program code) to be placed in a section named "`scn-name`", rather than the default .text section. No white space should appear around the '='. This command can be useful when developing a bootloader and you want to map your code in a custom linker script. Example:
	`xc32-gcc bootloader.c`
	
	`-mtext="MySectionName,address(0x9D00a000)"`
	
	`-mprocessor=32MX795F512L`

`-muninit-const-in-rodata` `-mno-uninit-const-in-rodata`	Put uninitialized `const` variables in the read-only data section. This option is only meaningful in conjunction with `-membedded-data`.
`--nofallback`	Require an MPLAB XC32 Pro license and do not fall back to a lesser license.

6.7.2   Options for Controlling the Kind of Output

The following options control the kind of output produced by the compiler.

Table 6-5. Kind-of-Output Control Options

Option	Definition
`-c`	Compile or assemble the source files, but do not link. The default file extension is `.o`.
`-E`	Stop after the preprocessing stage (that is, before running the compiler proper). The default output file is `stdout`.
`-fexceptions`	Enable exception handling. You may need to enable this option when compiling C code that needs to interoperate properly with exception handlers written in C++.
`-o file`	Place the output in `file`.
`-S`	Stop after compilation proper (that is, before invoking the assembler). The default output file extension is `.s`.
`-v`	Print the commands executed during each stage of compilation.

`-x`	You can specify the input language explicitly with the `-x` option:
	`-x language`
	
	Specify explicitly the language for the following input files (rather than letting the compiler choose a default based on the file name suffix). This option applies to all following input files until the next `-x` option. The following values are supported by the compiler:
	
	`c`
	`c++`
	`c-header`
	`cpp-output`
	`assembler`
	`assembler-with-cpp`
	
	`-x none`
	
	Turn off any specification of a language, so that subsequent files are handled according to their file name suffixes. This is the default -behavior but is needed if another `-x` option has been used. For exam-ple:
	
	`xc32-gcc -x assembler foo.asm bar.asm -x none main.c mabonga.s`
	
	Without the `-xnone`, the compiler assumes all the input files are for the assembler.

`--help`	Print a description of the command line options.

6.7.3 Options for Controlling the C Dialect

The following options define the kind of C dialect used by the compiler.

Table 6-6. C Dialect Control Options

Option	Definition
`-ansi`	Support all (and only) ANSI-standard C programs.

..........continued	
Option	**Definition**
`-aux-info filename`	Output to the given filename prototyped declarations for all functions declared and/or defined in a translation unit, including those in header files. This option is silently ignored in any language other than C. Besides declarations, the file indicates, in comments, the origin of each declaration (source file and line), whether the declaration was implicit, prototyped or unprototyped (`I`, `N` for new or `O` for old, respectively, in the first character after the line number and the colon), and whether it came from a declaration or a definition (`C` or `F`, respectively, in the following character). In the case of function definitions, a K&R-style list of arguments followed by their declarations is also provided, inside comments, after the declaration.
`-fcheck-new` / `-fno-check-new` (default)	Check that the pointer returned by operator new is non-null.
`-fenforce-eh-specs` (default) / `-fno-enforce-eh-specs`	Generate/Do not generate code to check for violation of exception specifications at runtime. The `-fno-enforce-eh-specs` option violates the C++ standard, but may be useful for reducing code size in production builds, much like defining 'NDEBUG'. This does not give user code permission to throw exceptions in violation of the exception specifications; the compiler will still optimize based on the specifications, so throwing an unexpected exception will result in undefined behavior.
`-ffreestanding`	Assert that compilation takes place in a freestanding en-vironment. This implies `-fno-builtin`. A freestanding environment is one in which the standard library may not exist, and program start-up may not necessarily be at main. The most obvious example is an OS kernel. This is equivalent to `-fno-hosted`.
`-fno-asm`	Do not recognize `asm`, `inline` or `typeof` as a keyword, so that code can use these words as identifiers. You can use the keywords `__asm__`, `__inline__` and `__typeof__` instead. `-ansi` implies `-fno-asm`.

`-fno-builtin`
 	
`-fno-builtin-function`
	Don't recognize built-in functions that do not begin with `__builtin_` as prefix.

`-fno-exceptions`	Disable C++ exception handling. This option disables the generation of extra code needed to propagate exceptions.
`-fno-rtti`	Enable/Disable runtime type-identification features. The `-fno-rtti` option disables generation of information about every class with virtual functions for use by the C++ runtime type identification features ('`dynamic_cast`' and '`typeid`'). If you don't use those parts of the language, you can save some space by using this flag. Note that exception handling uses the same information, but it will generate it as needed. The '`dynamic_cast`' operator can still be used for casts that do not require runtime type information, that is, casts to `void *` or to unambiguous base classes.
`-fsigned-char`	Let the type `char` be signed, like `signed char`. (This is the default.)

..........continued

Option	Definition

`-fsigned-bitfields`
 	
`-funsigned-bitfields`
 	
`-fno-signed-bitfields`
 	
`-fno-unsigned-bitfields`
	These options control whether a bit field is signed or unsigned, when the declaration does not use either signed or unsigned. By default, such a bit field is signed, unless `-traditional` is used, in which case bit fields are always unsigned.

`-funsigned-char`	Let the type `char` be unsigned, like `unsigned char`.

6.7.4 Options for Controlling the C++ Dialect

The following options define the kind of C++ dialect used by the compiler.

Table 6-7. C++ Dialect Control Options

Option	Definition
`-ansi`	Support all (and only) ANSI-standard C++ programs.
`-aux-info filename`	Output to the given filename prototyped declarations for all functions declared and/or defined in a translation unit, including those in header files. This option is silently ignored in any language other than C++. Besides declarations, the file indicates, in comments, the origin of each declaration (source file and line), whether the declaration was implicit, prototyped or unprototyped (`I`, `N` for new or `O` for old, respectively, in the first character after the line number and the colon), and whether it came from a -declaration or a definition (`C` or `F`, respectively, in the -following character). In the case of function definitions, a K&R-style list of arguments followed by their declarations is also provided, inside comments, after the declaration.
`-ffreestanding`	Assert that compilation takes place in a freestanding en-vironment. This implies `-fno-builtin`. A freestanding environment is one in which the standard library may not exist, and program start-up may not necessarily be at main. The most obvious example is an OS kernel. This is equivalent to `-fno-hosted`.
`-fno-asm`	Do not recognize `asm`, `inline` or `typeof` as a keyword, so that code can use these words as identifiers. You can use the keywords `__asm__`, `__inline__` and `__typeof__` instead. `-ansi` implies `-fno-asm`.

`-fno-builtin`
 	
`-fno-builtin-function`
	Don't recognize built-in functions that do not begin with `__builtin_` as prefix.

`-fsigned-char`	Let the type `char` be signed, like `signed char`. (This is the default.)

`-fsigned-bitfields`
 	
`-funsigned-bitfields`
 	
`-fno-signed-bitfields`
 	
`-fno-unsigned-bitfields`
	These options control whether a bit field is signed or unsigned, when the declaration does not use either signed or unsigned. By default, such a bit field is signed, unless `-traditional` is used, in which case bit fields are always unsigned.

`-funsigned-char`	Let the type `char` be unsigned, like `unsigned char`.

6.7.5 Options for Controlling Warning and Errors

Warnings are diagnostic messages that report constructions that are not inherently erroneous, but that are risky or suggest there may have been an error.

You can request many specific warnings with options beginning -W; for example, -Wimplicit, to request warnings on implicit declarations. Each of these specific warning options also has a negative form beginning -Wno- to turn off warnings; for example, -Wno-implicit. This manual lists only one of the two forms, whichever is not the default.

The following options control the amount and kinds of warnings produced by the compiler.

Table 6-8. Warning and Error Options Implied by All Warnings

Option	Definition
-fsyntax-only	Check the code for syntax, but don't do anything beyond that.
-pedantic	Issue all the warnings demanded by strict ANSI C. Reject all programs that use forbidden extensions.
-pedantic-errors	Like -pedantic, except that errors are produced rather than warnings.
-w	Inhibit all warning messages.
-Wall	This enables all the warnings about constructions that some users consider questionable and that are easy to avoid (or modify to prevent the warning), even in conjunction with macros. Note that some warning flags are not implied by -Wall. Some of them warn about constructions that users generally do not consider questionable, but which occasionally you might wish to check for; others warn about constructions that are necessary or hard to avoid in some cases, and there is no simple way to modify the code to suppress the warning. Some of them are enabled by -Wextra but many of them must be enabled individually.
-Waddress	Warn about suspicious uses of memory addresses. These include using the address of a function in a conditional expression, such as void func(void); if (func), and comparisons against the memory address of a string literal, such as if (x == "abc"). Such uses typically indicate a programmer error: the address of a function always evaluates to true, so their use in a conditional usually indicates that the programmer forgot the parentheses in a function call; and comparisons against string literals result in unspecified behavior and are not portable in C, so they usually indicate that the programmer intended to use strcmp.
-Wchar-subscripts	Warn if an array subscript has type char.
-Wcomment	Warn whenever a comment-start sequence /* appears in a /*comment, or whenever a Backslash-Newline appears in a // comment.
-Wdiv-by-zero	Warn about compile-time integer division by zero. To inhibit the warning messages, use -Wno-div-by-zero. Floating-point division by zero is not warned about, as it can be a legitimate way of obtaining infinities and NaNs. This is the default.
-Wformat	Check calls to printf and scanf, etc., to make sure that the arguments supplied have types appropriate to the format string specified.
-Wimplicit	Equivalent to specifying both -Wimplicit-int and -Wimplicit-function-declaration.
-Wimplicit-function-declaration	Give a warning whenever a function is used before being declared.
-Wimplicit-int	Warn when a declaration does not specify a type.

..........continued

Option	Definition
-Wmain	Warn if the type of `main` is suspicious. `main` should be a function with external linkage, returning `int`, taking either zero, two, or three arguments of appropriate types.

-Wmissing-braces	Warn if an aggregate or union initializer is not fully bracketed. In the following example, the initializer for `a` is not fully bracketed, but that for `b` is fully bracketed.
	`int a[2][2] = { 0, 1, 2, 3 };`
	
	`int b[2][2] = { { 0, 1 }, { 2, 3 } };`

-Wno-multichar	Warn if a multi-character `character` constant is used. Usually, such constants are typographical errors. Since they have implementation-defined values, they should not be used in portable code. The following example illustrates the use of a multi-character `character` constant:
	`char`
	
	`xx(void)`
	
	`{`
	
	`return('xx');`
	
	`}`

-Wparentheses	Warn if parentheses are omitted in certain contexts, such as when there is an assignment in a context where a truth value is expected, or when operators are nested whose precedence people often find confusing.
-Wreturn-type	Warn whenever a function is defined with a return-type that defaults to `int`. Also warn about any `return` statement with no return-value in a function whose return-type is not `void`.
-Wsequence-point	Warn about code that may have undefined semantics because of violations of sequence point rules in the C standard. The C standard defines the order in which expressions in a C program are evaluated in terms of sequence points, which represent a partial ordering between the execution of parts of the program: those executed before the sequence point and those executed after it. These occur after the evaluation of a full expression (one which is not part of a larger expression), after the evaluation of the first operand of a `&&`, `||`, `? :` or `,` (comma) operator, before a function is called (but after the evaluation of its arguments and the expression denoting the called function), and in certain other places. Other than as expressed by the sequence point rules, the order of evaluation of subexpressions of an expression is not specified. All these rules describe only a partial order rather than a total order. For example, if two functions are called within one expression with no sequence point between them, the order in which the functions are called is not specified. However, the standards committee has ruled that function calls do not overlap. It is not specified when between sequence points modifications to the values of objects take effect. Programs whose behavior depends on this have undefined behavior. The C standard specifies that "Between the previous and next sequence point, an object shall have its stored value modified, at most once, by the evaluation of an expression. Furthermore, the prior value shall be read only to determine the value to be stored." If a program breaks these rules, the results on any particular implementation are entirely unpredictable. Examples of code with undefined behavior are `a = a++;`, `a[n] = b[n++]` and `a[i++] = i;`. Some more complicated cases are not diagnosed by this option and it may give an occasional false positive result, but in general it has been found fairly effective at detecting this sort of problem in programs.

..........continued	
Option	**Definition**
-Wswitch	Warn whenever a switch statement has an index of enumeral type and lacks a case for one or more of the named codes of that enumeration. (The presence of a default label prevents this warning.) case labels outside the enumeration range also provoke warnings when this option is used.
-Wsystem-headers	Print warning messages for constructs found in system header files. Warnings from system headers are normally suppressed on the assumption that they usually do not indicate real problems and would only make the compiler output harder to read. Using this command line option tells the compiler to emit warnings from system headers as if they occurred in user code. However, note that using -Wall in conjunction with this option does not warn about unknown pragmas in system headers. For that, -Wunknown-pragmas must also be used.
-Wtrigraphs	Warn if any trigraphs are encountered (assuming they are enabled).
-Wuninitialized	Warn if an automatic variable is used without first being initialized. These warnings are possible only when optimization is enabled, because they require data flow information that is computed only when optimizing. These warnings occur only for variables that are candidates for register allocation. Therefore, they do not occur for a variable that is declared volatile, or whose address is taken, or whose size is other than 1, 2, 4 or 8 bytes. Also, they do not occur for structures, unions or arrays, even when they are in registers. Note that there may be no warning about a variable that is used only to compute a value that itself is never used, because such computations may be deleted by data flow analysis before the warnings are printed.
-Wunknown-pragmas	Warn when a #pragma directive is encountered which is not understood by the compiler. If this command line option is used, warnings will even be issued for unknown pragmas in system header files. This is not the case if the warnings were only enabled by the -Wall command line option.
-Wunused	Warn whenever a variable is unused aside from its declaration, whenever a function is declared static but never defined, whenever a label is declared but not used, and whenever a statement computes a result that is explicitly not used. In order to get a warning about an unused function parameter, both -W and -Wunused must be specified. Casting an expression to void suppresses this warning for an expression. Similarly, the unused attribute suppresses this warning for unused variables, parameters and labels.
-Wunused-function	Warn whenever a static function is declared but not defined or a non-inline static function is unused.
-Wunused-label	Warn whenever a label is declared but not used. To suppress this warning, use the unused attribute.
-Wunused-parameter	Warn whenever a function parameter is unused aside from its declaration. To suppress this warning, use the unused attribute.
-Wunused-variable	Warn whenever a local variable or non-constant static variable is unused aside from its declaration. To suppress this warning, use the unused attribute.
-Wunused-value	Warn whenever a statement computes a result that is explicitly not used. To suppress this warning, cast the expression to void.

The following -W options are not implied by -Wall. Some of them warn about constructions that users generally do not consider questionable, but which you might occasionally wish to check for. Others warn about constructions that

are necessary or hard to avoid in some cases, and there is no simple way to modify the code to suppress the warning.

Table 6-9. Warning and Error Options Not Implied by All Warnings

Option	Definition

`-W`	Print extra warning messages for these events:
	• A nonvolatile automatic variable might be changed by a call to `longjmp`. These warnings are possible only in optimizing compilation. The compiler sees only the calls to `setjmp`. It cannot know where `longjmp` will be called. In fact, a signal handler could call it at any point in the code. As a result, a warning may be generated even when there is in fact no problem, because `longjmp` cannot in fact be called at the place that would cause a problem.
	• A function could exit both via `return value;` and `return;`. Completing the function body without passing any return statement is treated as `return;`.
	• An expression-statement or the left-hand side of a comma expression contains no side effects. To suppress the warning, cast the unused expression to void. For example, an expression such as `x[i,j]` causes a warning, but `x[(void)i,j]` does not.
	• An unsigned value is compared against zero with < or <=.
	• A comparison like `x<=y<=z` appears, This is equivalent to `(x<=y ? 1 : 0) <= z`, which is a different interpretation from that of ordinary mathematical notation.
	• Storage-class specifiers like `static` are not the first things in a declaration. According to the C Standard, this usage is obsolescent.
	• If `-Wall` or `-Wunused` is also specified, warn about unused arguments.
	• A comparison between signed and unsigned values could produce an incorrect result when the signed value is converted to unsigned (but don't warn if `-Wno-sign-compare` is also specified).
	• An aggregate has a partly bracketed initializer. For example, the following code would evoke such a warning, because braces are missing around the initializer for `x.h`:
	`struct s { int f, g; };`
	`struct t { struct s h; int i; };`
	`struct t x = { 1, 2, 3 };`
	• An aggregate has an initializer that does not initialize all members. For example, the following code would cause such a warning, because `x.h` would be implicitly - initialized to zero:
	`struct s { int f, g, h; };`
	`struct s x = { 3, 4 };`

`-Waggregate-return`	Warn if any functions that return structures or unions are defined or called.
`-Wbad-function-cast`	Warn whenever a function call is cast to a non-matching type. For example, warn if `int foof()` is cast to anything `*`.
`-Wcast-align`	Warn whenever a pointer is cast, such that the required alignment of the target is increased. For example, warn if a `char *` is cast to an `int *`.
`-Wcast-qual`	Warn whenever a pointer is cast, so as to remove a type qualifier from the target type. For example, warn if a `const char *` is cast to an ordinary `char *`.

..........continued

Option	Definition
-Wconversion	Warn if a prototype causes a type conversion that is different from what would happen to the same argument in the absence of a prototype. This includes conversions of fixed point to floating and vice versa, and conversions changing the width or signedness of a fixed point argument, except when the same as the default promotion. Also, warn if a negative integer constant expression is implicitly converted to an unsigned type. For example, warn about the assignment `x = -1` if `x` is unsigned. But do not warn about explicit casts like `(unsigned) -1`.
-Werror	Make all warnings into errors.
-Winline	Warn if a function can not be inlined, and either it was declared as inline, or else the `-finline-functions` option was given.
-Wlarger-than-len	Warn whenever an object of larger than `len` bytes is defined.
-Wlong-long -Wno-long-long	Warn if `long long` type is used. This is default. To inhibit the warning messages, use `-Wno-long-long`. Flags `-Wlong-long` and `-Wno-long-long` are taken into account only when `-pedantic` flag is used.
-Wmissing-declarations	Warn if a global function is defined without a previous declaration. Do so even if the definition itself provides a -prototype.
-Wmissing-format-attribute	If `-Wformat` is enabled, also warn about functions that might be candidates for format attributes. Note these are only possible candidates, not absolute ones. This option has no effect unless `-Wformat` is enabled.
-Wmissing-noreturn	Warn about functions that might be candidates for attribute `noreturn`. These are only possible candidates, not absolute ones. Care should be taken to manually verify functions. In fact, do not ever return before adding the `noreturn` attribute, otherwise subtle code generation bugs could be introduced.
-Wmissing-prototypes	Warn if a global function is defined without a previous prototype declaration. This warning is issued even if the definition itself provides a prototype (this option can be used to detect global functions that are not declared in header files).
-Wnested-externs	Warn if an `extern` declaration is encountered within a function.
-Wno-deprecated-declarations	Do not warn about uses of functions, variables and types marked as deprecated by using the `deprecated` attribute.
-Wpadded	Warn if padding is included in a structure, either to align an element of the structure or to align the whole structure.
-Wpointer-arith	Warn about anything that depends on the size of a function type or of `void`. The compiler assigns these types a size of 1, for convenience in calculations with `void *`pointers and pointers to functions.
-Wredundant-decls	Warn if anything is declared more than once in the same scope, even in cases where multiple declaration is valid and changes nothing.
-Wshadow	Warn whenever a local variable shadows another local variable.
-Wsign-compare -Wno-sign-compare	Warn when a comparison between signed and unsigned values could produce an incorrect result when the signed value is converted to unsigned. This warning is also enabled by `-W`. To get the other warnings of `-W` without this warning, use `-W -Wno-sign-compare`.
-Wstrict-prototypes	Warn if a function is declared or defined without specifying the argument types (an old-style function definition is permitted without a warning if preceded by a declaration which specifies the argument types).

..........continued

Option	Definition
-Wtraditional	Warn about certain constructs that behave differently in traditional and ANSI C. • Macro arguments occurring within string constants in the macro body. These would substitute the argument in traditional C, but are part of the constant in ANSI C. • A function declared external in one block and then used after the end of the block. • A switch statement has an operand of type `long`. • A nonstatic function declaration follows a static one. This construct is not accepted by some traditional C compilers.
-Wundef	Warn if an undefined identifier is evaluated in an `#if` -directive.
-Wwrite-strings	Give string constants the type `const char[length]` so that copying the address of one into a non-const `char *` pointer gets a warning. At compile time, these warnings help you find code that you can try to write into a string constant, but only if you have been very careful about using `const` in declarations and prototypes. Otherwise, it's just a nuisance, which is why `-Wall` does not request these warnings.

6.7.6 Options for Debugging

The following options are used for debugging.

Table 6-10. Debugging Options

Option	Definition
-g	Produce debugging information. The compiler supports the use of `-g` with `-O` making it possible to debug optimized code. The shortcuts taken by optimized code may occasionally produce surprising results: • Some declared variables may not exist at all • Flow of control may briefly move unexpectedly • Some statements may not be executed because they compute constant results or their values were already at hand • Some statements may execute in different places because they were moved out of loops Nevertheless it proves possible to debug optimized output. This makes it reasonable to use the optimizer for programs that might have bugs.
-Q	Makes the compiler print out each function name as it is compiled, and print some statistics about each pass when it finishes.
-save-temps -save-temps=cwd	Don't delete intermediate files. Place them in the current directory and name them based on the source file. Thus, compiling `foo.c` with `-c -save-temps` would produce the following files: `foo.i` (preprocessed file) `foo.s` (assembly language file) `foo.o` (object file)

..........continued

Option	Definition

-save-temps=obj	Similar to -save-temps=cwd, but if the -o option is specified, the temporary files are based on the object file. If the -o option is not specified, the -save-temps=obj switch behaves like -save-temps.
	For example:
	
	``
	xc32-gcc -save-temps=obj -c
	          foo.c
	``
	
	``
	xc32-gcc -save-temps=obj -c bar.c -o
	          dir/xbar.o
	``
	
	``
	xc32-gcc -save-temps=obj foobar.c -o
	          dir2/yfoobar
	``
	
	would create foo.i, foo.s, dir/xbar.i, dir/xbar.s, dir2/yfoobar.i, dir2/yfoobar.s, and dir2/yfoobar.o.

6.7.7   Options for Controlling Optimization

The following options control compiler optimizations.

Table 6-11. General Optimization Options

Option	Edition	Definition
-O0	All	Do not optimize (this is the default). Without -O, the compiler's goal is to reduce the cost of compilation and to make debugging produce the expected results. Statements are independent: if you stop the program with a breakpoint between statements, you can then assign a new value to any variable or change the program counter to any other statement in the function and get exactly the results you would expect from the source code. The compiler only allocates variables declared register in registers.
-O -O1	All	Optimization level 1. Optimizing compilation takes somewhat longer and a lot more host memory for a large function. With -O, the compiler tries to reduce code size and execution time. When -O is specified, the compiler turns on -fthread-jumps and -fdefer-pop. The compiler turns on -fomit-frame-pointer.
-O2	PRO	Optimization level 2. The compiler performs nearly all supported optimizations that do not involve a space-speed trade-off. -O2 turns on all optional -optimizations except for loop unrolling (-funroll-loops), function inlining (-finline-functions), and strict aliasing optimizations (-fstrict-aliasing). It also turns on force copy of memory operands (-fforce-mem) and Frame Pointer elimination (-fomit-frame-pointer). As compared to -O, this option increases both compilation time and the performance of the generated code.
-O3	PRO	Optimization level 3. -O3 turns on all optimizations specified by -O2 and also turns on the inline-functions option.
-Os	PRO	Optimize for size. -Os enables all -O2 optimizations that do not typically increase code size. It also performs further optimizations designed to reduce code size.

..........continued

Option	Edition	Definition

`-flto`	PRO	This option runs the standard link-time optimizer. When invoked with source code, the compiler adds an internal bytecode representation of the code to special sections in the object file. When the object files are linked together, all the function bodies are read from these sections and instantiated as if they had been part of the same translation unit. To use the link-time optimizer, specify `-flto` both at compile time and during the final link. For example:
		
		 `xc32-gcc -c -O1 -flto -mprocessor=32MX795F512L foo.c`
		
		 `xc32-gcc -c -O1 -flto -mprocessor=32MX795F512L bar.c`
		
		 `xc32-gcc -o myprog.elf -flto -O3 -mprocessor=32MX795F512L foo.o bar.o`
		
		 Another (simpler) way to enable link-time optimization is,
		
		 `xc32-gcc -o myprog.elf -flto -O3 -mprocessor=32MX795F512L foo.c bar.c`
		
		 Link time optimizations do not require the presence of the whole program to operate. If the program does not require any symbols to be exported, it is possible to combine `-flto` with `-fwhole-program` to allow the interprocedural optimizers to use more aggressive assumptions which may lead to improved optimization opportunities.
		
		 Regarding portability: The bytecode files are versioned and there is a strict version check, so bytecode files generated in one version of XC32 may not work with an older/newer version of the XC32 compiler.

`-ftoplevel -reorder`	All	Allow reordering top-level functions, variables, and asm statements. They may not be output in the same order that they appear in the input file. This option also allows removal of unreferenced static variables. Use this options with optimization level `-O1` or greater.
`-fwhole -program`	PRO	Assume that the current compilation unit represents the whole program being compiled. All public functions and variables, with the exception of main and those merged by attribute `externally_visible`, become static functions and in effect are optimized more aggressively by interprocedural optimizers. While this option is equivalent to proper use of the static keyword for programs consisting of a single file, in combination with option `-flto`, this flag can be used to compile many smaller scale programs since the functions and variables become local for the whole combined compilation unit, not for the single source file itself.

The following options control specific optimizations. The `-O2` option turns on all of these optimizations except `-funroll-loops`, `-funroll-all-loops` and `-fstrict-aliasing`.

You can use the following flags in the rare cases when "fine-tuning" of optimizations to be performed is desired.

Table 6-12. Specific Optimization Options

Option	Definition
`-falign-functions` `-falign-functions=n`	Align the start of functions to the next power-of-two greater than `n`, skipping up to `n` bytes. For instance, `-falign-functions=32` aligns functions to the next 32-byte boundary, but `-falign-functions=24` would align to the next 32-byte boundary only if this can be done by skipping 23 bytes or less. `-fno-align-functions` and `-falign-functions=1` are equivalent and mean that functions are not aligned. The assembler only supports this flag when `n` is a power of two, so `n` is rounded up. If `n` is not specified, use a machine-dependent default.

..........continued

Option	Definition
-falign-labels -falign-labels=n	Align all branch targets to a power-of-two boundary, skipping up to n bytes like -falign-functions. This option can easily make code slower, because it must insert dummy operations for when the branch target is reached in the usual flow of the code. If -falign-loops or -falign-jumps are applicable and are greater than this value, then their values are used instead. If n is not specified, use a machine-dependent default which is very likely to be 1, meaning no alignment.
-falign-loops -falign-loops=n	Align loops to a power-of-two boundary, skipping up to n bytes like -falign-functions. The hope is that the loop is executed many times, which makes up for any execution of the dummy operations. If n is not specified, use a machine-dependent default.
-fcaller-saves	Enable values to be allocated in registers that are -clobbered by function calls, by emitting extra instructions to save and restore the registers around such calls. Such allocation is done only when it seems to result in better code than would otherwise be produced.
-fcse-follow-jumps	In common subexpression elimination, scan through jump instructions when the target of the jump is not reached by any other path. For example, when CSE encounters an if statement with an else clause, CSE follows the jump when the condition tested is false.
-fcse-skip-blocks	This is similar to -fcse-follow-jumps, but causes CSE to follow jumps which conditionally skip over blocks. When CSE encounters a simple if statement with no else clause, -fcse-skip-blocks causes CSE to follow the jump around the body of the if.
-fexpensive-optimizations	Perform a number of minor optimizations that are relatively expensive.
-ffunction-sections -fdata-sections	Place each function or data item into its own section in the output file. The name of the function or the name of the data item determines the section's name in the output file. Only use these options when there are significant benefits for doing so. When you specify these options, the assembler and linker may create larger object and executable files and is also slower.
-fgcse	Perform a global common subexpression elimination pass. This pass also performs global constant and copy propagation.
-fgcse-lm	When -fgcse-lm is enabled, global common subexpression elimination attempts to move loads which are only killed by stores into themselves. This allows a loop containing a load/store sequence to change to a load outside the loop, and a copy/store within the loop.
-fgcse-sm	When -fgcse-sm is enabled, a store motion pass is run after global common subexpression elimination. This pass attempts to move stores out of loops. When used in conjunction with -fgcse-lm, loops containing a load/store sequence can change to a load before the loop and a store after the loop.
-fmove-all-movables	Forces all invariant computations in loops to be moved outside the loop.
-fno-defer-pop	Always pop the arguments to each function call as soon as that function returns. The compiler normally lets arguments accumulate on the stack for several function calls and pops them all at once.

..........continued	
Option	**Definition**

`-fno-peephole`
 	
`-fno-peephole2`
	Disable machine specific peephole optimizations. Peephole optimizations occur at various points during the compilation. `-fno-peephole` disables peephole optimization on machine instructions, while `-fno-peephole2` disables high level peephole optimizations. To disable peephole entirely, use both options.

`-foptimize-register-move`
 	
`-fregmove`
	Attempt to reassign register numbers in move instructions and as operands of other simple instructions in order to maximize the amount of register tying. `-fregmove` and `-foptimize-register-moves` are the same optimization.

`-freduce-all-givs`	Forces all general-induction variables in loops to be strength reduced. These options may generate better or worse code. Results are highly dependent on the structure of loops within the source code.
`-frename-registers`	Attempt to avoid false dependencies in scheduled code by making use of registers left over after register allocation. This optimization most benefits processors with lots of registers. It can, however, make debugging impossible, since variables no longer stay in a "home register."
`-frerun-cse-after-loop`	Rerun common subexpression elimination after loop optimizations has been performed.
`-frerun-loop-opt`	Run the loop optimizer twice.
`-fschedule-insns`	Attempt to reorder instructions to eliminate instruction stalls due to required data being unavailable.
`-fschedule-insns2`	Similar to `-fschedule-insns`, but requests an additional pass of instruction scheduling after register allocation has been done.
`-fstrength-reduce`	Perform the optimizations of loop strength reduction and elimination of iteration variables.

..........continued

Option	Definition

-fstrict-aliasing	Allows the compiler to assume the strictest aliasing rules applicable to the language being compiled. For C, this activates optimizations based on the type of expressions. In particular, an object of one type is assumed never to reside at the same address as an object of a different type, unless the types are almost the same. For example, an `unsigned int` can alias an `int`, but not a `void*` or a `double`. A character type may alias any other type.
	Pay special attention to code like this:
	`union a_union {`
	`int i;`
	`double d;`
	`};`
	
	`int f() {`
	`union a_union t;`
	`t.d = 3.0;`
	`return t.i;`
	`}`
	The practice of reading from a different union member than the one most recently written to (called "type-punning") is common. Even with `-fstrict-aliasing`, type-punning is allowed, provided the memory is accessed through the union type. The code above works as expected. However, this code might not:
	`int f() {`
	`a_union t;`
	`int* ip;`
	`t.d = 3.0;`
	`ip = &t.i;`
	`return *ip;`
	`}`

-fthread-jumps	Perform optimizations where a check is made to see if a jump branches to a location where another comparison subsumed by the first is found. If so, the first branch is redirected to either the destination of the second branch or a point immediately following it, depending on whether the condition is known to be true or false.
-funroll-loops	Perform the optimization of loop unrolling. This is only done for loops whose number of iterations can be determined at compile time or run time. `-funroll-loops` implies both `-fstrength-reduce` and `-frerun-cse-after-loop`.
-funroll-all-loops	Perform the optimization of loop unrolling. This is done for all loops and usually makes programs run more slowly. `-funroll-all-loops` implies `-fstrength-reduce`, as well as `-frerun-cse-after-loop`.
-fuse-caller-save	Allows the compiler to use the caller-save register model. When combined with inter-procedural optimizations, the compiler can generate more efficient code.

Options of the form `-fflag` specify machine-independent flags. Most flags have both positive and negative forms. The negative form of `-ffoo` would be `-fno-foo`. In the table below, only one of the forms is listed (the one that is not the default).

Table 6-13. Machine-Independent Optimization Options

Option	Definition
-fforce-mem	Force memory operands to be copied into registers before doing arithmetic on them. This produces better code by making all memory references potential common subexpressions. When they are not common subexpressions, instruction combination should eliminate the separate register load. The -O2 option turns on this option.
-finline-functions	Integrate all simple functions into their callers. The compiler heuristically decides which functions are simple enough to be worth integrating in this way. If all calls to a given function are integrated, and the function is declared static, then the function is normally not output as assembler code in its own right.
-finline-limit=n	By default, the compiler limits the size of functions that can be inlined. This flag allows the control of this limit for functions that are explicitly marked as inline (that is, marked with the inline keyword). n is the size of functions that can be inlined in number of pseudo instructions (not counting parameter handling). The default value of n is 10000. Increasing this value can result in more inlined code at the cost of compilation time and memory consumption. Decreasing usually makes the compilation faster and less code is inlined (which presumably means slower programs). This option is particularly useful for programs that use inlining. **Note:** Pseudo instruction represents, in this particular context, an abstract measurement of function's size. In no way does it represent a count of assembly instructions and as such, its exact meaning might change from one release of the compiler to an another.
-fkeep-inline-functions	Even if all calls to a given function are integrated, and the function is declared static, output a separate run time callable version of the function. This switch does not affect extern inline functions.
-fkeep-static-consts	Emit variables are declared static const when optimization isn't turned on, even if the variables are not referenced. The compiler enables this option by default. If you want to force the compiler to check if the variable was referenced, regardless of whether or not optimization is turned on, use the -fno-keep-static-consts option.
-fno-function-cse	Do not put function addresses in registers. Make each instruction that calls a constant function contain the function's address explicitly. This option results in less efficient code, but some strange hacks that alter the assembler output may be confused by the optimizations performed when this option is not used.
-fno-inline	Do not pay attention to the inline keyword. Normally this option is used to keep the compiler from expanding any functions inline. If optimization is not enabled, no functions can be expanded inline.
-fomit-frame-pointer	Do not keep the Frame Pointer in a register for functions that don't need one. This avoids the instructions to save, set up and restore Frame Pointers. It also makes an extra register available in many functions.
-foptimize-sibling-calls	Optimize sibling and tail recursive calls.

6.7.8   Options for Controlling the Preprocessor

The following options control the compiler preprocessor.

Table 6-14. Preprocessor Options

Option	Definition
-C	Tell the preprocessor not to discard comments. Used with the -E option.

-dD	Tell the preprocessor to not remove macro definitions into the output, in their proper sequence.
-Dmacro	Define macro *macro* with string 1 as its definition.
-Dmacro=defn	Define macro *macro* as *defn*. All instances of -D on the command line are processed before any -U options.
-dM	Tell the preprocessor to output only a list of the macro -definitions that are in effect at the end of preprocessing. Used with the -E option.
-dN	Like -dD except that the macro arguments and contents are omitted. Only #define name is included in the output.
-fno-show-column	Do not print column numbers in diagnostics. This may be necessary if diagnostics are being scanned by a program that does not understand the column numbers, such as DejaGnu.
-H	Print the name of each header file used, in addition to other normal activities.
-I-	Any directories you specify with -I options before the -I- options are searched only for the case of #include "file". They are not searched for #include <*file*>. If additional directories are specified with -I options after the -I-, these directories are searched for all #include directives. (Ordinarily all -I directories are used this way.) In addition, the -I- option inhibits the use of the current -directory (where the current input file came from) as the first search directory for #include "file". There is no way to override this effect of -I-. With -I. you can specify -searching the directory that was current when the compiler was invoked. That is not exactly the same as what the preprocessor does by default, but it is often satisfactory. -I- does not inhibit the use of the standard system -directories for header files. Thus, -I- and -nostdinc are independent. **Note:** Do not specify an MPLAB XC32 system include directory (for example, /pic32mx/include/) in your project properties. The xc32-gcc and xc32-g++ compilation drivers automatically select the default C libc or the C++ libc and their respective include-file directory for you. Manually adding a system include file path may disrupt this mechanism and cause the incorrect libc include files to be compiled into your project, causing a conflict between the include files and the library. Note that adding a system include path to your project properties has never been a recommended practice.
-Idir	Add the directory *dir* to the head of the list of directories to be searched for header files. This can be used to override a system header file, substituting your own version, since these directories are searched before the system header file directories. If you use more than one -I option, the directories are scanned in left-to-right order. The standard system directories come after.
-idirafter dir	Add the directory dir to the second include path. The directories on the second include path are searched when a header file is not found in any of the directories in the main include path (the one that -I adds to).
-imacros file	Process file as input, discarding the resulting output, before processing the regular input file. Because the output generated from the file is discarded, the only effect of -imacros *file* is to make the macros defined in file available for use in the main input. Any -D and -U options on the command line are always -processed before -imacros *file*, regardless of the order in which they are written. All the -include and -imacros options are processed in the order in which they are written.
-include file	Process file as input before processing the regular input file. In effect, the contents of file are compiled first. Any -D and -U options on the command line are always processed before -include *file*, regardless of the order in which they are written. All the -include and -imacros options are processed in the order in which they are written.

`-M`	Tell the preprocessor to output a rule suitable for `make` describing the dependencies of each object file. For each source file, the preprocessor outputs one make-rule whose target is the object file name for that source file and whose dependencies are all the `#include` header files it uses. This rule may be a single line or may be continued with `\-newline` if it is long. The list of rules is printed on standard output instead of the preprocessed C program. `-M` implies `-E` (see 6.7.2 Options for Controlling the Kind of Output).
`-MD`	Like `-M` but the dependency information is written to a file and compilation continues. The file containing the dependency information is given the same name as the source file with a `.d` extension.
`-MF file`	When used with `-M` or `-MM`, specifies a file in which to write the dependencies. If no `-MF` switch is given, the preprocessor sends the rules to the same place it would have sent preprocessed output. When used with the driver options, `-MD` or `-MMD`, `-MF`, overrides the default dependency output file.
`-MG`	Treat missing header files as generated files and assume they live in the same directory as the source file. If `-MG` is specified, then either `-M` or `-MM` must also be specified. `-MG` is not supported with `-MD` or `-MMD`.
`-MM`	Like `-M` but the output mentions only the user header files included with `#include "file"`. System header files included with `#include <file>` are omitted.
`-MMD`	Like `-MD` except mention only user header files, not system header files.

`-MP`	This option instructs CPP to add a phony target for each dependency other than the main file, causing each to depend on nothing. These dummy rules work around errors `make` gives if you remove header files without updating the make-file to match. This is typical output:
	
	`test.o: test.c test.h`
	
	`test.h:`

`-MQ`	Same as `-MT`, but it quotes any characters which are special to `make`. `-MQ '$(objpfx)foo.o'` gives `$$(objpfx)foo.o: foo.c` The default target is automatically quoted, as if it were given with `-MQ`.
`-MT target`	Change the target of the rule emitted by dependency generation. By default, CPP takes the name of the main input file, including any path, deletes any file suffix such as `.c`, and appends the platform's usual object suffix. The result is the target. An `-MT` option sets the target to be exactly the string you specify. If you want multiple targets, you can specify them as a single argument to `-MT`, or use multiple `-MT` options. For example: `-MT '$(objpfx)foo.o'` might give `$(objpfx)foo.o: foo.c`
`-nostdinc`	Do not search the standard system directories for header files. Only the directories you have specified with `-I` options (and the current directory, if appropriate) are searched. (See 6.7.11 Options for Directory Search) for information on `-I`. By using both `-nostdinc` and `-I-`, the include-file search path can be limited to only those directories explicitly specified.
`-P`	Tell the preprocessor not to generate `#line` directives. Used with the `-E` option (see 6.7.2 Options for Controlling the Kind of Output).
`-trigraphs`	Support ANSI C trigraphs. The `-ansi` option also has this effect.

`-Umacro`	Undefine macro `macro`. `-U` options are evaluated after all `-D` options, but before any `-include` and `-imacros` options.
`-undef`	Do not predefine any nonstandard macros (including architecture flags).

6.7.9   Options for Assembling

The following options control assembler operations.

Table 6-15. Assembly Options

Option	Definition
`-Wa,option`	Pass `option` as an option to the assembler. If `option` contains commas, it is split into multiple options at the commas.

6.7.10   Options for Linking

If any of the options `-c`, `-S` or `-E` are used, the linker is not run and object file names should not be used as arguments.

Table 6-16. Linking Options

Option	Definition
`-fill=<options`	A memory-fill option to be passed on to the linker.
`-Ldir`	Add directory `dir` to the list of directories to be searched for libraries specified by the command line option `-l`.
`-legacy-libc`	Use legacy include files and libraries (C32 v1.12 and before). The format of the include file and libraries changed in C32 v2.00 to match HI-TECH C compiler format.

`-llibrary`	Search the library named `library` when linking.
	The linker searches a standard list of directories for the library, which is actually a file named `liblibrary.a`. The linker then uses this file as if it had been specified precisely by name.
	
	It makes a difference where in the command you write this option. The linker processes libraries and object files in the order they are specified. Thus, `foo.o -lz bar.o` searches library `z` after file `foo.o` but before `bar.o`. If `bar.o` refers to functions in `libz.a`, those functions may not be loaded.
	
	The directories searched include several standard system directories, plus any that you specify with `-L`.
	
	Normally the files found this way are library files (archive files whose members are object files). The linker handles an archive file by scanning through it for members which define symbols that have been referenced but not defined yet. But if the file found is an ordinary object file, it is linked in the usual fashion. The only difference between using an `-l` option (for example,`-lmylib`) and specifying a file name (for example,`libmylib.a`) is that `-l` searches several directories, as specified.
	
	By default the linker is directed to search:
	
	`<install-path>\lib`
	
	for libraries specified with the `-l` option. For a compiler installed into the default location, this would be:
	
	`Program Files\Microchip\mplab32\<version>\lib`
	
	This behavior can be overridden using the environment variables.
	
	See also the `INPUT` and `OPTIONAL` linker script directives.

`-mips16`	Link the MIPS16 ISA variant of the libraries.

..........continued

Option	Definition
-mmicromips	Link the microMIPS compressed ISA variant of the libraries and startup code. The ISA of the startup code must match the setting of the BOOTISA config bit. Therefore, use the -mmicromips link option to link the microMIPS startup code only when you are using #pragma config BOOTISA=MICROMIPS in your source code.
-nodefaultlibs	Do not use the standard system libraries when linking. Only the libraries you specify are passed to the linker. The compiler may generate calls to memcmp, memset and memcpy. These entries are usually resolved by entries in the standard compiler libraries. These entry points should be supplied through some other mechanism when this option is specified.
-nostdlib	Do not use the standard system start-up files or libraries when linking. No start-up files and only the libraries you specify are passed to the linker. The compiler may generate calls to memcmp, memset and -memcpy. These entries are usually resolved by entries in standard compiler libraries. These entry points should be supplied through some other mechanism when this option is specified.
-relaxed-math	This relaxed-compliance math library is enabled with the -relaxed-math xc32-gcc command-line option at link. This library provides alternative floating-point support routines that are faster and smaller than the default math routines, but make some sacrifices in compliance. For instance, it does not do all of the infinity, overflow and NaN checking, etc., of a fully-compliant library. It does not return all of the detailed feedback from that checking. However, this reduced compliance is usually sufficient for most applications.
-s	Remove all symbol table and relocation information from the executable.
-u symbol	Pretend *symbol* is undefined to force linking of library modules to define the symbol. It is legitimate to use -u multiple times with different symbols to force loading of additional library modules.
-Wl,option	Pass *option* as an option to the linker. If *option* contains commas, it is split into multiple options at the commas.
-Xlinker option	Pass *option* as an option to the linker. You can use this to supply system-specific linker options that the compiler does not know how to recognize.

6.7.11 Options for Directory Search

The following options specify to the compiler where to find directories and files to search.

Table 6-17. Directory Search Options

Option	Definition
-Bprefix	This option specifies where to find the executables, libraries, include files and data files of the compiler itself. The compiler driver program runs one or more of the sub-programs xc32-cpp, xc32-as and xc32-ld. It tries *prefix* as a prefix for each program it tries to run. For each sub-program to be run, the compiler driver first tries the -B prefix, if any. Lastly, the driver searches the current PATH environment variable for the subprogram. -B prefixes that effectively specify directory names also apply to libraries in the linker, because the compiler translates these options into -L options for the linker. They also apply to include files in the preprocessor, because the compiler translates these options into -isystem options for the preprocessor. In this case, the compiler appends include to the prefix.

..........continued	
Option	**Definition**
-specs=file	Process file after the compiler reads in the standard `specs` file, in order to override the defaults that the `xc32-gcc` driver program uses when -determining what switches to pass to `xc32-as`, `xc32-ld`, etc. More than one `-specs=file` can be specified on the command line, and they are processed in order, from left to right.

6.7.12 Options for Code Generation Conventions

Options of the form `-fflag` specify machine-independent flags. Most flags have both positive and negative forms. The negative form of `-ffoo` would be `-fno-foo`. In the table below, only one of the forms is listed (the one that is not the default).

Table 6-18. Code Generation Convention Options

Option	**Definition**
-fargument-alias -fargument-noalias -fargument-noalias-global	Specify the possible relationships among parameters and between parameters and global data. `-fargument-alias` specifies that arguments (parameters) may alias each other and may alias global storage. `-fargument-noalias` specifies that arguments do not alias each other, but may alias global storage. `-fargument-noalias-global` specifies that arguments do not alias each other and do not alias global storage. Each language automatically uses whatever option is required by the language standard. You should not need to use these options yourself.
-fcall-saved-reg	Treat the register named *reg* as an allocatable register saved by functions. It may be allocated even for temporaries or variables that live across a call. Functions compiled this way saves and restores the register *reg* if they use it. It is an error to use this flag with the Frame Pointer or Stack Pointer. Use of this flag for other registers that have fixed pervasive roles in the machine's execution model produces disastrous results. A different sort of disaster results from the use of this flag for a register in which function values are returned. This flag should be used consistently through all modules.
-fcall-used-reg	Treat the register named *reg* as an allocatable register that is clobbered by function calls. It may be allocated for temporaries or variables that do not live across a call. Functions compiled this way do not save and restore the register *reg*. It is an error to use this flag with the Frame Pointer or Stack Pointer. Use of this flag for other registers that have fixed pervasive roles in the machine's execution model produces disastrous results. This flag should be used consistently through all modules.
-ffixed-reg	Treat the register named *reg* as a fixed register. Generated code should never refer to it (except perhaps as a Stack Pointer, Frame Pointer or in some other fixed role). *reg* must be the name of a register (for example,`-ffixed-$0`).
-fno-ident	Ignore the `#ident` directive.
-fpack-struct	Pack all structure members together without holes. Usually you would not want to use this option, since it makes the code sub-optimal, and the offsets of structure members won't agree with system libraries.

..........continued	
Option	**Definition**
`-fpcc-struct-return`	Return short `struct` and `union` values in memory like longer ones, rather than in registers. This convention is less efficient, but it has the advantage of allowing capability between 32-bit compiled files and files compiled with other compilers. Short structures and unions are those whose size and alignment match that of an integer type.
`-fshort-enums`	Allocate to an `enum` type only as many bytes as it needs for the declared range of possible values. Specifically, the `enum` type is equivalent to the smallest integer type that has enough room.
`-fverbose-asm` `-fno-verbose-asm`	Put extra commentary information in the generated assembly code to make it more readable. `-fno-verbose-asm`, the default, causes the extra information to be omitted and is useful when comparing two assembler files.
`-fvolatile`	Consider all memory references through pointers to be volatile.
`-fvolatile-global`	Consider all memory references to external and global data items to be volatile. The use of this switch has no effect on static data.
`-fvolatile-static`	Consider all memory references to static data to be volatile.

7.   ANSI C Standard Issues

This compiler conforms to the ANSI X3.159-1989 Standard for programming languages. This is commonly called the C89 Standard. It is referred to as the ANSI C Standard in this manual. Some features from the later standard, C99, are also supported.

7.1   Divergence Fom the ANSI C Standard

There are no divergences from the ANSI C standard.

7.2   Extensions to the ANSI C Standard

C/C++ code for the MPLAB XC32 C/C++ Compiler differs from the ANSI C standard in these areas: keywords, statements and expressions.

7.2.1   Keyword Differences

The new keywords are part of the base GCC implementation and the discussions in the referenced sections are based on the standard GCC documentation, tailored for the specific syntax and semantics of the 32-bit compiler port of GCC.

- Specifying Attributes of Variables – 9.11 Variable Attributes
- Specifying Attributes of Functions – 14.2 Function Attributes and Specifiers
- Inline Functions – 14.9 Inline Functions
- Variables in Specified Registers – 9.11 Variable Attributes
- Complex Numbers – 9.7 Complex Data Types
- Double-Word Integers – 9.3 Integer Data Types
- Referring to a Type with `typeof` – 9.9 Standard Type Qualifiers

7.2.2   Statement Differences

The statement differences are part of the base GCC implementation, and the discussions in the referenced sections are based on the standard GCC documentation, tailored for the specific syntax and semantics of the 32-bit compiler port of GCC.

- Labels as Values – 12.3 Labels as Values
- Conditionals with Omitted Operands – 12.4 Conditional Operator Operands
- Case Ranges – 12.5 Case Ranges

7.2.3   Expression Differences

Expression differences are:

Binary constants – 9.8 Constant Types and Formats.

7.3   Implementation-Defined Behavior

Certain features of the ANSI C standard have implementation-defined behavior. This means that the exact behavior of some C code can vary from compiler to compiler. The exact behavior of the MPLAB XC32 C/C++ Compiler is detailed throughout this documentation, and is fully summarized in 23. Implementation-Defined Behavior.

8.   Device-Related Features

The MPLAB XC32 C/C++ Compiler supports a number of special features and extensions to the C/C++ language which are designed to ease the task of producing ROM-based applications. This chapter documents the special language features which are specific to these devices.

8.1   Device Support

MPLAB XC32 C/C++ Compiler aims to support all PIC32 devices. However, new devices in these families are frequently released.

8.2   Device Header Files

There is one header file that is recommended be included into each source file you write. The file is `<xc.h>` and is a generic file that will include other device-specific header files when you build your project.

Inclusion of this file will allow access to SFRs via special variables, as well as `#define`S which allow the use of conventional register names from within assembly language files.

8.2.1   CP0 Register Definitions Header File

The CP0 register definitions header file (`cp0defs.h`) is a file that contains definitions for the CP0 registers and their fields. In addition, it contains macros for accessing the CP0 registers. These macros are defined in `cp0defs.h` so the application code should `#include <cp0defs.h>`.

The CP0 register definitions header file is located in the `pic32mx/include` directory of your compiler installation directory. The CP0 register definitions header file is automatically included when you include the generic device header file, `xc.h`.

The CP0 register definitions header file was designed to work with either Assembly or C/C++ files. The CP0 register definitions header file is dependent on macros defined within the processor generic header file.

8.3   Stack

The PIC32 devices use what is referred to in this user's guide as a "software stack". This is the typical stack arrangement employed by most computers and is ordinary data memory accessed by a push-and-pop type instruction and a stack pointer register. The term "hardware stack" is used to describe the stack employed by Microchip 8-bit devices, which is only used for storing function return addresses.

The PIC32 devices use a dedicated stack pointer register `sp` (register 29) for use as a software Stack Pointer. All processor stack operations, including function calls, interrupts and exceptions, use the software stack. It points to the next free location on the stack. The stack grows downward, towards lower memory addresses.

By default, the size of the stack is 1024 bytes. The size of the stack can be changed by specifying the size on the linker command line using the `--defsym _min_stack_size` linker command line option. An example of allocating

a stack of 2048 bytes using the command line is:

```
xc32-gcc foo.c -Wl,--defsym,_min_stack_size=2048
```

The run-time stack grows downward from higher addresses to lower addresses. Two working registers are used to manage the stack:

• Register 29 (`sp`) – This is the Stack Pointer. It points to the next free location on

the stack.

• Register 30 (`fp`) – This is the Frame Pointer. It points to the current function's

frame.

No stack overflow detection is supplied.

The C/C++ run-time start-up module initializes the stack pointer during the start-up and initialization sequence, see 16.2.3 Initialize Stack Pointer and Heap.

8.4 Configuration Bit Access

The PIC32 devices have several locations which contain the Configuration bits or fuses. These bits specify fundamental device operation, such as the oscillator mode, watchdog timer, programming mode and code protection. Failure to correctly set these bits may result in code failure, or a non-running device.

The `#pragma config` directive specifies the processor-specific configuration settings (i.e., Configuration bits) to be used by the application. Refer to the "PIC32 Configuration Settings" online help (found under *MPLAB X IDE>Help>Help Contents>XC32 Toolchain*) for more information. (If using the compiler from the command line, this help file is located at the default location at: `Program Files/ Microchip/ <install-dir>/<version>/ docs/PIC32ConfigSet.html`.)

Configuration settings may be specified with multiple `#pragma config` directives. The compiler verifies that the configuration settings specified are valid for the processor for which it is compiling. If a given setting in the Configuration word has not been specified in any `#pragma config` directive, the bits associated with that setting default to the unprogrammed value. Configuration settings should be specified in only a single translation unit (a C/C++ file with all of its include files after preprocessing).

For each Configuration word for which a setting is specified with the `#pragma config` directive, the compiler generates a read-only data section named `.config_address`, where *address* is the hexadecimal representation of the address of the Configuration word. For example, if a configuration setting was specified for the Configuration word located at address `0xBFC02FFC`, a read-only data section named `.config_BFC02FFC` would be created.

8.4.1 Syntax

The following shows the meta syntax notation for the different forms the pragma may take.

```
pragma-config-directive:
     # pragma config setting-list
setting-list:
     setting
   | setting-list, setting
setting:
     setting-name = value-name
```

The setting-name and value-name are device specific and can be determined by using the PIC32ConfigSet.html document located in the installation directory, docs folder.

All `#pragma config` directives should be placed outside of a function definition as they do not define executable code.

PIC32MZ config pragmas include `config_alt`, `config_bf1`, `config_abf1`, `config_bf2`, and `config_abf2` pragmas to support placing configuration bit values in the alternate, boot flash 1, alternate boot flash 1, boot flash 2, and alternate boot flash 2 PIC32MZ memory regions, respectively. (Example: `#pragma config_bf2 FWDTEN=off`)

Integer values for config pragmas can be set using the `config` and `config_region` pragmas. (Examples: `#pragma config_bf2 TSEQ = 1` and `#pragma config USERID = 0x1234u`)

8.4.2 Example

The following example shows how the `#pragma config` directive might be utilized. The example does the following:

- Enables the Watchdog Timer
- Sets the Watchdog Postscaler to 1:128
- Selects the HS Oscillator for the Primary Oscillator

```
#pragma config FWDTEN = ON, WDTPS = PS128
#pragma config POSCMOD = HS
...
int main (void)
{
```

```
   ...
   }
```

8.5   ID Locations

User-defined ID locations are implemented in one Configuration Word. These locations should be programmed using the #pragma config directive. See 8.4 Configuration Bit Access.

Example: `#pragma config USERID=0x1234`.

8.6   Using SFRs From C Code

The Special Function Registers (SFRs) are registers which control aspects of the MCU operation or that of peripheral modules on the device. These registers are memory mapped, which means that they appear at specific addresses in the device memory map. With some registers, the bits within the register control independent features.

Memory-mapped SFRs are accessed by special C variables that are placed at the addresses of the registers and use special attributes. These variables can be accessed like any ordinary C variable so that no special syntax is required to access SFRs.

The SFR variables are predefined in header files and will be accessible once the `<xc.h>` header file (see 8.2 Device Header Files) has been included into your source code. Structures are also defined by these header files to allow access to bits within the SFR.

The names given to the C variables, which map over the registers and bit variables, or bit fields, within the registers are based on the names specified in the device data sheet. The names of the structures that hold the bit fields will typically be those of the corresponding register followed by `bits`. For example, the following shows code that includes the generic header file, clears PORTB as a whole and sets bit 2 of PORTB using the structure/bit field definitions

Note: The symbols `PORTB` and `PORTBbits` refer to the same register and resolve to the same address. Writing to one register will change the values held by both.

```c
#include <xc.h>
int main(void)
{
  PORTBCLR = 0xFFFFu;
  PORTBbits.RB2 = 1;
  PORTBSET= _PORTB_RB2_MASK;
}
```

For use with assembly, the `PORTB` register is declared as: `.extern PORTB`.

To confirm the names that are relevant for the device you are using, check the device specific header file that `<xc.h>` will include for the definitions of each variable. These files will be located in the `pic32mx/include/proc` directory of the compiler and will have a name that represents the device. There is a one-to-one correlation between device and header file names that will be included by `<xc.h>`, for example, when compiling for a PIC32MX360F512L device, the `<xc.h>` header file will include `<proc/p32mx360f512l.h>`. Remember that you do not need to include this chip-specific file into your source code; it is automatically included by `<xc.h>`.

Some of the PIC32 SFRs have associated registers that allow the bits within the SFR to be set, cleared or toggled atomically. For example, the `PORTB` SFR has the write-only registers `PORTBSET`, `PORTBCLR` and `PORTBINV` associated with it. Writing a `'1'` to a bit location in these registers sets, clears or toggles, respectively, the corresponding bit in the `PORTB` SFR. So to set bit 1 in `PORTB`, you can use the following code:

`PORTBSET = 0x2;`

or alternatively, using macros provided in the device header files:

`PORTBSET = _PORTB_RB1_MASK;`

8.6.1 CP0 Register Definitions

When the CP0 register definitions header file is included from an Assembly file, the CP0 registers are defined as:

```
#define _CP0_register_name $register_number, select_number
```

For example, the `IntCtl` register is defined as:

```
#define _CP0_INTCTL $12, 1
```

When the CP0 register definitions header file is included from a C file, the CP0 registers and selects are defined as:

```
#define _CP0_register_name register_number
#define _CP0_register_name_SELECT select_number
```

For example, the `IntCtl` register is defined as:

```
#define _CP0_INTCTL 12
#define _CP0_INTCTL_SELECT 1
```

8.6.2 CP0 Register Field Definitions

When the CP0 register definitions header file is included from either an Assembly or a C/C++ file, three `#defines` exist for each of the CP0 register fields.

`_CP0_register_name_field_name_POSITION` – the starting bit location

`_CP0_register_name_field_name_MASK` – the bits that are part of this field are set

`_CP0_register_name_field_name_LENGTH` – the number of bits that this field occupies

For example, the vector spacing field of the `IntCtl` register has the following defines:

```
#define _CP0_INTCTL_VS_POSITION 0x00000005
#define _CP0_INTCTL_VS_MASK     0x000003E0
#define _CP0_INTCTL_VS_LENGTH   0x00000005
```

8.6.3 CP0 Access Macros

When the CP0 register definitions header file is included from a C file, CP0 access macros are defined. Each CP0 register may have up to six different access macros defined:

`_CP0_GET_register_name ()`	Returns the value for register, `register_name`.
`_CP0_SET_register_name (val)`	Sets the register, `register_name`, to `val`, and returns `void`. Only defined for registers that contain a writable field.
`_CP0_XCH_register_name (val)`	Sets the register, `register_name`, to `val`, and returns the previous register value. Only defined for registers that contain a writable field.
`_CP0_BIS_register_name (set)`	Sets the register, `register_name`, to (`reg |= set`), and returns the previous register value. Only defined for registers that contain writable bit fields.
`_CP0_BIC_register_name (clr)`	Sets the register, `register_name`, to (`reg &= ~clr`), and returns the previous register value. Only defined for registers that contain writable bit fields.
`_CP0_BCS_register_name (clr, set)`	Sets the register, `register_name`, to (`reg = (reg & ~clr)	set`), and returns the previous register value. Only defined for registers that contain writable bit fields.

8.6.4   Address Translation Macros

System code may need to translate between virtual and physical addresses, as well as between kernel segment addresses. The macros are defined in `sys/kmem.h` so the application code should `#include <sys/kmem.h>`. Macros are provided to make these translations easier and to determine the segment an address is in.

`KVA_TO_PA(v)`	Translate a kernel virtual address to a physical address.
`PA_TO_KVA0(pa)`	Translate a physical address to a KSEG0 virtual address.
`PA_TO_KVA1(pa)`	Translate a physical address to a KSEG1 virtual address.
`KVA0_TO_KVA1(v)`	Translate a KSEG0 virtual address to a KSEG1 virtual address.
`KVA1_TO_KVA0(v)`	Translate a KSEG1 virtual address to a KSEG0 virtual address.
`IS_KVA(v)`	Evaluates to 1 if the address is a kernel segment virtual address, zero otherwise.
`IS_KVA0(v)`	Evaluate to 1 if the address is a KSEG0 virtual address, zero otherwise.
`IS_KVA1(v)`	Evaluate to 1 if the address is a KSEG1 virtual address, zero otherwise.
`IS_KVA01(v)`	Evaluate to 1 if the address is either a KSEG0 or a KSEG1 virtual address, zero otherwise.

9. Supported Data Types and Variables

The MPLAB XC32 C/C++ Compiler supports a variety of data types and attributes. These data types and variables are discussed here. For information on where variables are stored in memory, see 10. Memory Allocation and Access.

9.1 Identifiers

A C/C++ variable identifier (the following is also true for function identifiers) is a sequence of letters and digits, where the underscore character "_" counts as a letter. Identifiers cannot start with a digit. Although they may start with an underscore, such identifiers are reserved for the compiler's use and should not be defined by your programs. Such is not the case for assembly domain identifiers, which often begin with an underscore

Identifiers are case sensitive, so `main` is different than `Main`.

All characters are significant in an identifier, although identifiers longer than 31 characters in length are less portable.

9.2 Data Representation

The compiler stores multibyte values in little-endian format. That is, the Least Significant Byte is stored at the lowest address.

For example, the 32-bit value `0x12345678` would be stored at address `0x100` as:

Address	0x100	0x101	0x102	0x103
Data	0x78	0x56	0x34	0x12

9.3 Integer Data Types

Integer values in the compiler are represented in 2's complement and vary in size from 8 to 64 bits. These values are available in compiled code via 9.3.2 limits.h.

Type	Bits	Min	Max
char, signed char	8	-128	127
unsigned char	8	0	255
short, signed short	16	-32768	32767
unsigned short	16	0	65535
int, signed int, long, signed long	32	-2^{31}	$2^{31}-1$
unsigned int, unsigned long	32	0	$2^{32}-1$
long long, signed long long	64	-2^{63}	$2^{63}-1$
unsigned long long	64	0	$2^{64}-1$

9.3.1 Signed and Unsigned Character Types

By default, values of type plain `char` are signed values. This behavior is implementation-defined by the C standard, and some environments (notably, PowerPC) define a plain C/C++ `char` value to be unsigned. The command line option `-funsigned-char` can be used to set the default type to unsigned for a given translation unit.

9.3.2 limits.h

The `limits.h` header file defines the ranges of values which can be represented by the integer types.

Macro Name	Value	Description
CHAR_BIT	8	The size, in bits, of the smallest non-bit field object.
SCHAR_MIN	-128	The minimum value possible for an object of type signed char.
SCHAR_MAX	127	The maximum value possible for an object of type signed char.
UCHAR_MAX	255	The maximum value possible for an object of type unsigned char.
CHAR_MIN	-128 (or 0, see 9.3.1 Signed and Unsigned Character Types)	The minimum value possible for an object of type char.
CHAR_MAX	127 (or 255, see 9.3.1 Signed and Unsigned Character Types)	The maximum value possible for an object of type char.
MB_LEN_MAX	16	The maximum length of multibyte character in any locale.
SHRT__MIN	-32768	The minimum value possible for an object of type short int.
SHRT_MAX	32767	The maximum value possible for an object of type short int.
USHRT_MAX	65535	The maximum value possible for an object of type unsigned short int.
INT_MIN	-2^{31}	The minimum value possible for an object of type int.
INT_MAX	$2^{31}-1$	The maximum value possible for an object of type int.
UINT_MAX	$2^{32}-1$	The maximum value possible for an object of type unsigned int.
LONG_MIN	-2^{31}	The minimum value possible for an object of type long.
LONG_MAX	$2^{31}-1$	The maximum value possible for an object of type long.
ULONG_MAX	$2^{32}-1$	The maximum value possible for an object of type unsigned long.
LLONG_MIN	-2^{63}	The minimum value possible for an object of type long long.
LLONG_MAX	$2^{63}-1$	The maximum value possible for an object of type long long.
ULLONG_MAX	$2^{64}-1$	The maximum value possible for an object of type unsigned long long.

9.4   Floating-Point Data Types

The compiler uses the IEEE-754 floating-point format. Detail regarding the implementation limits is available to a translation unit in `float.h`.

Type	Bits
`float`	32
`double`	32
`long double`	64

Variables may be declared using the `float`, `double` and `long double` keywords, respectively, to hold values of these types. Floating-point types are always signed and the `unsigned` keyword is illegal when specifying a floating-point type. All floating-point values are represented in little endian format with the Least Significant Byte (LSB) at the lower address.

This format is described in the table below, where:

- Sign is the sign bit which indicates if the number is positive or negative.
- For 32-bit floating-point values, the exponent is 8 bits which is stored as excess 127 (that is, an exponent of 0 is stored as 127).
- For 64-bit floating-point values, the exponent is 11 bits which is stored as excess 1023 (that is, an exponent of 0 is stored as 1023).
- Mantissa is the mantissa, which is to the right of the radix point. There is an implied bit to the left of the radix point which is always 1 except for a zero value, where the implied bit is zero. A zero value is indicated by a zero exponent.

The value of this number for 32-bit floating-point values is:

$(-1)^{sign} \times 2^{(exponent-127)} \times 1.\,mantissa$

and for 64-bit values

$(-1)^{sign} \times 2^{(exponent-1023)} \times 1.\,mantissa.$

Here is an example of the IEEE 754 32-bit format shown in the following table. Note that the Most Significant bit of the mantissa column (that is, the bit to the left of the radix point) is the implied bit, which is assumed to be 1 unless the exponent is zero (in which case the float is zero).

Table 9-1. Floating-Point Format Example IEEE 754

Format	Number	Biased Exponent	1.mantissa	Decimal
32-bit	7DA6B69Bh	11111011b (251)	1.01001101011011010011011b (1.302447676659)	2.77000e+37 —

The example in table above can be calculated manually as follows.

The sign bit is zero; the biased exponent is 251, so the exponent is 251-127=124. Take the binary number to the right of the decimal point in the mantissa. Convert this to decimal and divide it by 2^{23} where 23 is the number of bits taken up by the mantissa, to give 0.302447676659. Add 1 to this fraction. The floating-point number is then given by:

$-1^0 \times 2^{124} \times 1.302447676659$

which becomes:

$1 \times 2.126764793256e+37 \times 1.302447676659$

which is approximately equal to:

$2.77000e+37$

Binary floating-point values are sometimes misunderstood. It is important to remember that not every floating-point value can be represented by a finite-sized floating-point number. The size of the exponent in the number dictates the range of values that the number can hold, and the size of the mantissa relates to the spacing of each value that can be represented exactly. Thus the 64-bit floating-point format allows for values with a larger range of values and that can be more accurately represented.

So, for example, if you are using a 32-bit wide floating-point type, it can exactly store the value 95000.0. However, the next highest number it can represent is (approximately) 95000.00781 and it is impossible to represent any value in between these two in such a type as it will be rounded. This implies that C/C++ code which compares floating-point type may not behave as expected. For example:

```
volatile float myFloat;
myFloat = 95000.006;
if(myFloat == 95000.007)      // value will be rounded
    LATA++;                   // this line will be executed!
```

in which the result of the `if()` expression will be true, even though it appears the two values being compared are different.

The characteristics of the floating-point formats are summarized in the following table. The symbols in this table are preprocessor macros which are available after including `<float.h>` in your source code. Two sets of macros are available for `float` and `double` types, where *XXX* represents `FLT` and `DBL`, respectively. So, for example, `FLT_MAX` represents the maximum floating-point value of the `float` type. `DBL_MAX` represents the same values for the `double` type. As the size and format of floating-point data types are not fully specified by the ANSI Standard, these macros allow for more portable code which can check the limits of the range of values held by the type on this implementation.

Table 9-2. Ranges of Floating-Point Type Values

Symbol	Meaning	32-bit Value	64-bit Value
*XXX*_RADIX	Radix of exponent representation	2	2
*XXX*_ROUNDS	Rounding mode for addition	1	
*XXX*_MIN_EXP	Min. *n* such that FLT_RADIX^{n-1} is a normalized float value	-125	-1021
*XXX*_MIN_10_EXP	Min. *n* such that 10^n is a normalized float value	-37	-307
*XXX*_MAX_EXP	Max. *n* such that FLT_RADIX^{n-1} is a normalized float value	128	1024
*XXX*_MAX_10_EXP	Max. *n* such that 10^n is a normalized float value	38	308
*XXX*_MANT_DIG	Number of FLT_RADIX mantissa digits	24	53
*XXX*_EPSILON	The smallest number which added to 1.0 does not yield 1.0	1.1920929e-07	2.2204460492503131e-16

9.5 Structures and Unions

MPLAB XC32 C/C++ Compiler supports `struct` and `union` types. Structures and unions only differ in the memory offset applied to each member.

These types will be at least 1 byte wide. Bit fields are fully supported.

Structures and unions may be passed freely as function arguments and function return values. Pointers to structures and unions are fully supported.

9.5.1   Structure and Union Qualifiers

The MPLAB XC32 C/C++ Compiler supports the use of type qualifiers on structures. When a qualifier is applied to a structure, all of its members will inherit this qualification. In the following example the structure is qualified `const`.

```
const struct {
        int number;
        int *ptr;
} record = { 0x55, &i };
```

In this case, the entire structure will be placed into the program memory and each member will be read-only. Remember that all members are usually initialized if a structure is `const` as they cannot be initialized at runtime.

If the members of the structure were individually qualified `const`, but the structure was not, then the structure would be positioned into RAM, but each member would be read-only. Compare the following structure with the above.

```
struct {
        const int number;
        int * const ptr;
} record = { 0x55, &i};
```

9.5.2   Bit Fields in Structures

MPLAB XC32 C/C++ Compiler fully supports bit fields in structures.

Bit fields are always allocated within 8-bit storage units, even though it is usual to use the type `unsigned int` in the definition. Storage units are aligned on a 32-bit boundary, although this can be changed using the `packed` attribute.

The first bit defined will be the Least Significant bit of the word in which it will be stored. When a bit field is declared, it is allocated within the current 8-bit unit if it will fit; otherwise, a new byte is allocated within the structure. Bit fields can never cross the boundary between 8-bit allocation units. For example, the declaration:

```
struct {
        unsigned        lo : 1;
        unsigned        dummy : 6;
        unsigned        hi : 1;
} foo;
```

will produce a structure occupying 1 byte.

Unnamed bit fields may be declared to pad out unused space between active bits in control registers. For example, if `dummy` is never referenced, the structure above could have been declared as:

```
struct {
        unsigned        lo : 1;
        unsigned           : 6;
        unsigned        hi : 1;
} foo;
```

A structure with bit fields may be initialized by supplying a *comma*-separated list of initial values for each field. For example:

```
struct {
        unsigned        lo  : 1;
        unsigned        mid : 6;
        unsigned        hi  : 1;
} foo = {1, 8, 0};
```

Structures with unnamed bit fields may be initialized. No initial value should be supplied for the unnamed members, for example:

```
struct {
        unsigned        lo  : 1;
        unsigned            : 6;
        unsigned        hi  : 1;
} foo = {1, 0};
```

will initialize the members `lo` and `hi` correctly.

The MPLAB XC compiler supports anonymous unions. These are unions with no identifier and whose members can be accessed without referencing the enclosing union. These unions can be used when placing inside structures. For example:

```
struct {
        union {
        int x;
        double y;
    };
} aaa;

int main(void)
{
    aaa.x = 99;
    // ...}
```

Here, the union is not named and its members accessed as if they are part of the structure. Anonymous unions are not part of any C Standard and so their use limits the portability of any code.

9.6   Pointer Types

There are two basic pointer types supported by the MPLAB XC32 C/C++ Compiler: data pointers and function pointers. Data pointers hold the addresses of variables which can be indirectly read, and possible indirectly written, by the program. Function pointers hold the address of an executable function which can be called indirectly via the pointer.

9.6.1   Combining Type Qualifiers and Pointers

It is helpful to first review the ANSI C/C++ standard conventions for definitions of pointer types.

Pointers can be qualified like any other C/C++ object, but care must be taken when doing so as there are two quantities associated with pointers. The first is the actual pointer itself, which is treated like any ordinary C/C++ variable and has memory reserved for it. The second is the target, or targets, that the pointer references, or to which the pointer points. The general form of a pointer definition looks like the following:

target_type_&_qualifiers * *pointer's_qualifiers pointer's_name;*

Any qualifiers to the right of the `*` (that is, next to the pointer's name) relate to the pointer variable itself. The type and any qualifiers to the left of the `*` relate to the pointer's targets. This makes sense since it is also the `*` operator that dereferences a pointer, which allows you to get from the pointer variable to its current target.

Here are three examples of pointer definitions using the `volatile` qualifier. The fields in the definitions have been highlighted with spacing:

```
volatile int *          vip ;
int           * volatile  ivp ;
volatile int * volatile  vivp ;
```

The first example is a pointer called `vip`. It contains the address of `int` objects that are qualified `volatile`. The pointer itself — the variable that holds the address — is *not* `volatile`; however, the objects that are accessed when the pointer is dereferenced are treated as being `volatile`. In other words, the target objects accessible via the pointer may be externally modified.

The second example is a pointer called `ivp` which also contains the address of `int` objects. In this example, the pointer itself is `volatile`, that is, the address the pointer contains may be externally modified; however, the objects that can be accessed when dereferencing the pointer are not `volatile`.

The last example is of a pointer called `vivp` which is itself qualified `volatile`, and which also holds the address of `volatile` objects.

Bear in mind that one pointer can be assigned the addresses of many objects; for example, a pointer that is a parameter to a function is assigned a new object address every time the function is called. The definition of the pointer must be valid for every target address assigned.

Note: Care must be taken when describing pointers. Is a "const pointer" a pointer that points to `const` objects, or a pointer that is `const` itself? You can talk about "pointers to const" and "const pointers" to help clarify the definition, but such terms may not be universally understood.

9.6.2 Data Pointers

Pointers in the compiler are all 32 bits in size. These can hold an address which can reach all memory locations.

9.6.3 Function Pointers

The MPLAB XC compiler fully supports pointers to functions, which allows functions to be called indirectly. These are often used to call one of several function addresses stored in a user-defined C/C++ array, which acts like a lookup table.

Function pointers are always 32 bits in size and hold the address of the function to be called.

Any attempt to call a function with a function pointer containing NULL will result in an ifetch Bus Error.

9.6.4 Special Pointer Targets

Pointers and integers are not interchangeable. Assigning an integer constant to a pointer will generate a warning to this effect. For example:

```
const char * cp = 0x123; // the compiler will flag this as bad code
```

There is no information in the integer constant, 0x123, relating to the type or size of the destination. This code is also not portable and there is a very good chance of code failure if pointers are assigned integer addresses and dereferenced, particularly for PIC® devices that have more than one memory space.

Always take the address of a C/C++ object when assigning an address to a pointer. If there is no C/C++ object defined at the destination address, then define or declare an object at this address which can be used for this purpose. Make sure the size of the object matches the range of the memory locations that can be accessed.

For example, a checksum for 1000 memory locations starting at address `0xA0001000` is to be generated. A pointer is used to read this data. You may be tempted to write code such as:

```
int * cp;
cp = 0xA0001000;  // what resides at 0xA0001000???
```

and increment the pointer over the data. A much better solution is this:

```
int * cp;
int __attribute__((address(0xA0001000))) inputData [1000];
cp = &inputData;
// cp is incremented over inputData and used to read values there
```

In this case, the compiler can determine the size of the target and the memory space. The array size and type indicates the size of the pointer target.

Take care when comparing (subtracting) pointers. For example:

```
if(cp1 == cp2)
  ; take appropriate action
```

The ANSI C standard only allows pointer comparisons when the two pointer targets are the same object. The address may extend to one element past the end of an array.

Comparisons of pointers to integer constants are even more risky, for example:

```
if(cp1 == 0xA0000100)
  ; take appropriate action
```

A NULL pointer is the one instance where a constant value can be assigned to a pointer and this is handled correctly by the compiler. A NULL pointer is numerically equal to 0 (zero), but this is a special case imposed by the ANSI C standard. Comparisons with the macro NULL are also allowed.

9.7 Complex Data Types

Complex data types are currently not implemented in MPLAB XC32 C/C++ Compiler.

9.8 Constant Types and Formats

A constant is used to represent a numerical value in the source code, for example 123 is a constant. Like any value, a constant must have a C/C++ type. In addition to a constant's type, the actual value can be specified in one of several formats. The format of integral constants specifies their radix. MPLAB XC32 C/C++ supports the ANSI standard radix specifiers as well as ones which enables binary constants to be specified in C code.

The formats used to specify the radices are given in the table below. The letters used to specify binary or hexadecimal radices are case insensitive, as are the letters used to specify the hexadecimal digits.

Table 9-3. Radix Formats

Radix	Format	Example
binary	`0b` *number* or `0B` *number*	0b10011010
octal	`0` *number*	0763
decimal	*number*	129
hexadecimal	`0x` *number* or `0X` *number*	0x2F

Any integral constant will have a type of `int`, `long int` or `long long int`, so that the type can hold the value without overflow. Constants specified in octal or hexadecimal may also be assigned a type of `unsigned int`, `unsigned long int` or `unsigned long long int` if the signed counterparts are too small to hold the value.

The default types of constants may be changed by the addition of a suffix after the digits, for example, `23U`, where `U` is the suffix. The table below shows the possible combination of suffixes and the types that are considered when assigning a type. So, for example, if the suffix `l` is specified and the value is a decimal constant, the compiler will assign the type `long int`, if that type will hold the constant; otherwise, it will assigned `long long int`. If the constant was specified as an octal or hexadecimal constant, then unsigned types are also considered.

Table 9-4. Suffixed and Assigned Types

Suffix	Decimal	Octal or Hexadecimal
`u` or `U`	unsigned int unsigned long int unsigned long long int	unsigned int unsigned long int unsigned long long int
`l` or `L`	long int long long int	long int unsigned long int long long int unsigned long long int
`u` or `U`, and `l` or `L`	unsigned long int unsigned long long int	unsigned long int unsigned long long int
`ll` or `LL`	long long int	long long int unsigned long long int
`u` or `U`, and `ll` or `LL`	unsigned long long int	unsigned long long int

Here is an example of code that may fail because the default type assigned to a constant is not appropriate:

```
unsigned long int result;
unsigned char shifter;

int main(void)
{
    shifter = 40;
    result = 1 << shifter;
    // code that uses result
}
```

The constant `1` will be assigned an `int` type hence the result of the shift operation will be an `int` and the upper bits of the `long` variable, `result`, can never be set, regardless of how much the constant is shifted. In this case, the value 1 shifted left 40 bits will yield the result 0, not 0x10000000000.

The following uses a suffix to change the type of the constant, hence ensure the shift result has an `unsigned long` type.

```
result = 1UL << shifter;
```

Floating-point constants have `double` type unless suffixed by `f` or `F`, in which case it is a `float` constant. The suffixes `l` or `L` specify a `long double` type.

Character constants are enclosed by single quote characters, `'`, for example `'a'`. A character constant has `int` type, although this may be optimized to a `char` type later in the compilation.

Multi-byte character constants are accepted by the compiler but are not supported by the standard libraries.

String constants, or string literals, are enclosed by double quote characters `"  "`, for example `"hello world"`. The type of string constants is `const char *` and the character that make up the string are stored in the program memory, as are all objects qualified `const`.

To comply with the ANSI C standard, the compiler does not support the extended character set in characters or character arrays. Instead, they need to be escaped using the backslash character, as in the following example:

```
const char name[] = "Bj\370rk";
printf("%s's Resum\351", name); \\ prints "Bjørk's Resumé"
```

Assigning a string literal to a pointer to a non-`const char` will generate a warning from the compiler. This code is legal, but the behavior if the pointer attempts to write to the string will fail. For example:

```
char * cp= "one";         // "one" in ROM, produces warning
const char * ccp= "two";  // "two"  in ROM, correct
```

Defining and initializing a non-const array (i.e., not a pointer definition) with a string,

```
char ca[]= "two"; // "two" different to the above
```

is a special case and produces an array in data space which is initialized at start-up with the string `"two"` (copied from program space), whereas a string constant used in other contexts represents an unnamed `const` -qualified array, accessed directly in program space.

The MPLAB XC32 C/C++ Compiler will use the same storage location and label for strings that have identical character sequences. For example, in the code snippet

```
if(strncmp(scp, "hello world", 6) == 0)
    fred = 0;
if(strcmp(scp, "hello world") == 0)
    fred++;
```

the two identical character string greetings will share the same memory locations. The link-time optimization must be enabled to allow this optimization when the strings may be located in different modules.

Two adjacent string constants (that is, two strings separated *only* by white space) are concatenated by the compiler. Thus:

```
const char * cp = "hello"  "world";
```

will assign the pointer with the address of the string `"hello world"`.

9.9   Standard Type Qualifiers

Type qualifiers provide additional information regarding how an object may be used. The MPLAB XC32 C/C++ Compiler supports both ANSI C qualifiers and additional special qualifiers which are useful for embedded applications and which take advantage of the PIC MCU architecture.

9.9.1   Const Type Qualifier

The MPLAB XC32 C/C++ Compiler supports the use of the ANSI type qualifiers `const` and `volatile`.

The `const` type qualifier is used to tell the compiler that an object is read only and will not be modified. If any attempt is made to modify an object declared `const`, the compiler will issue a warning or error.

Usually a `const` object must be initialized when it is declared, as it cannot be assigned a value at any point at runtime. For example:

```
const int  version = 3;
```

will define `version` as being an `int` variable that will be placed in the program memory, will always contain the value 3, and which can never be modified by the program.

Objects qualified const are placed into the program memory unless the `-mno-embedded-data` option is used.

9.9.2   Volatile Type Qualifier

The `volatile` type qualifier is used to tell the compiler that an object cannot be guaranteed to retain its value between successive accesses. This prevents the optimizer from eliminating apparently redundant references to objects declared `volatile` because it may alter the behavior of the program to do so.

Any SFR which can be modified by hardware or which drives hardware is qualified as `volatile` and any variables which may be modified by interrupt routines should use this qualifier as well. For example:

```
extern volatile unsigned int WDTCON __attribute__((section("sfrs")));
```

The `volatile` qualifier does not guarantee that any access will be atomic, but the compiler will try to implement this.

The code produced by the compiler to access `volatile` objects may be different than that to access ordinary variables and typically the code will be longer and slower for `volatile` objects, so only use this qualifier if it is necessary. However, failure to use this qualifier when it is required may lead to code failure.

Another use of the `volatile` keyword is to prevent variables from being removed if they are not used in the C/C++ source. If a non-`volatile` variable is never used, or used in a way that has no effect on the program's function, then it may be removed before code is generated by the compiler.

A C/C++ statement that consists only of a `volatile` variable's name will produce code that reads the variable's memory location and discards the result. For example the entire statement:

```
PORTB;
```

will produce assembly code that reads `PORTB`, but does nothing with this value. This is useful for some peripheral registers that require reading to reset the state of interrupt flags. Normally such a statement is not encoded as it has no effect.

9.10 Compiler-Specific Qualifiers

There are currently no non-standard qualifiers implemented in MPLAB XC32 C/C++ Compiler. Attributes are used to control variables and functions.

9.11 Variable Attributes

The compiler keyword `__attribute__` allows you to specify special attributes of variables or structure fields. This keyword is followed by an attribute specification inside double parentheses.

To specify multiple attributes, separate them by commas within the double parentheses, for example:

`__attribute__ ((aligned (16), packed))`.

Note: It is important to use variable attributes consistently throughout a project. For example, if a variable is defined in file A with the `aligned` attribute, and declared `extern` in file B without `aligned`, then a link error may result.

address (addr)

Specify an absolute virtual address for the variable. This attribute can be used in conjunction with a section attribute.

Note: For a data variable on a target device without an L1 cache, the address is typically in the range [0xA0000000,0xA00FFFFC], as defined in the linker script as the 'kseg1_data_mem' region. For data variables on a target feature an L1 data cache, the address is typically in the range [0x80000000,0x800FFFFC] as defined in the linker script as the 'kseg0_data_mem' region. Take special care to use the correct kseg region for your device or you may end up with more than one variable allocated to the same physical address.

This attribute can be used to start a group of variables at a specific address:

```
int foo __attribute__((section("mysection"),address(0xA0001000)));
int bar __attribute__((section("mysection")));
int baz __attribute__((section("mysection")));
```

Keep in mind that the compiler performs no error checking on the specified address. The section will be located at the specified address regardless of the memory-region ranges listed in the linker script or the actual ranges on the target device. This application code is responsible for ensuring that the address is valid for the target device and application.

Also, be aware that variables attributed with an absolute address are not accessed via GP-relative addressing. This means that they may be more expensive to access than non-address attributed variables.

In addition, to make effective use of absolute sections and the best-fit allocator, standard program-memory and data-memory sections should not be mapped in the linker script. The built-in linker script does not map most standard sections such as the `.text`, `.data`, `.bss`, or `.ramfunc` section. By not mapping these sections in the linker script, we allow these sections to be allocated using the best-fit allocator rather than the sequential allocator. Sections that are unmapped in the linker script can flow around absolute sections whereas sections that are linker-script mapped are grouped together and allocated sequentially, potentially causing conflicts with absolute sections.

Finally, note that "small" data and bss (`.sdata`, `.sbss`, etc.) sections are still mapped in the built-in default linker script. This is because "small" data variables must be grouped together so that they are within range of the more efficient GP-relative addressing mode. *To avoid conflict with these linker-script mapped sections, choose high addresses for your absolute-address variables.*

Note: In almost all cases, you will want to combine the address attribute with the space attribute to indicate code or data with `space(prog)` or `space(data)`, respectively. Also, see the description for the attribute `space(memory-space)`.

aligned (n)

The attributed variable will be aligned on the next `n` byte boundary.

The `aligned` attribute can also be used on a structure member. Such a member will be aligned to the indicated boundary within the structure.

If the alignment value `n` is omitted, the alignment of the variable is set 8 (the largest alignment value for a basic data type).

Note that the `aligned` attribute is used to increase the alignment of a variable, not reduce it. To decrease the alignment value of a variable, use the `packed` attribute.

cleanup (function)

Indicate a function to call when the attributed automatic function scope variable goes out of scope.

The indicated function should take a single parameter, a pointer to a type compatible with the attributed variable, and have `void` return type.

coherent

The coherent variable attribute causes the compiler/linker to place the variable into a unique section that is allocated to the kseg1 region, rather than the kseg0 region (which is the default on L1 cached devices). This means that the variable is accessed through the uncached address.

For devices featuring an L1 data cache, data variables are allocated to the KSEG0 data-memory region (`kseg0_data_mem`), making it accessible through the L1 cache. Likewise, the linker-allocated heap and stack are allocated to the KSEG0 region.

The is a `coherent` variable attribute that allows you to create a DMA buffer allocated to the kseg1_data_mem region:

```
unsigned int __attribute__((coherent)) buffer [1024];
```

When combining the `coherent` attribute with the `address` attribute, be sure to use the default data-memory region address for the device. On devices featuring an L1 data cache, the default data-memory region is kseg0_data_mem:

```
unsigned int __attribute__((coherent,address(0x80001000))) buffer[1024]
```

The `__pic32_alloc_coherent(size_t)` and `__pic32_free_coherent(void*)` functions allocate and free memory from the uncached kseg1_data_mem region. The default stack is allocated to the cached kseg0_data_mem region, but you may want to create an uncached DMA buffer, so you can use these functions to allocate an uncached buffer. These functions call the standard `malloc()`/`free()` functions, but the pointers that they use are translated from kseg0 to kseg1.

```
#include<xc.h>
void jak(void){
    char* buffer = __pic32_alloc_coherent(1024);
    if (buffer){
      /* do somehing */
    }
    else{
      /* handle error */
    }
    if (buffer){
      __pic32_free_coherent(buffer);
    }
}
```

deprecated

deprecated (msg)

When a variable specified as `deprecated` is used, a warning is generated. The optional `msg` argument, which must be a string, will be printed in the warning, if present.

noload

The `noload` attribute causes the variable or function to be placed in a section that has the `noload` attribute set. This attribute tells consumers of the ELF file not to load the contents of the section. This attribute can be useful when you just want to reserve memory for something, but you don't want to clear or initialize memory.

persistent

The `persistent` attribute specifies that the variable should not be initialized or cleared at startup. Use a variable with the `persistent` attribute to store state information that will remain valid after a device Reset. The `persistent`

attribute causes the compiler to place the variable in special .bss-like section that does not get cleared by the default startup code. Because the section is always in data space, this attribute is not compatible with the `space()` attribute.

```
int last_mode __attribute__((persistent));
```

The `persistent` attribute implies the `coherent` attribute. That is, `persistent` attributed variables are accessed via the uncached address.

packed

The attributed variable or structure member will have the smallest possible alignment. That is, no alignment padding storage will be allocated for the declaration. Used in combination with the `aligned` attribute, `packed` can be used to set an arbitrary alignment restriction greater or lesser than the default alignment for the type of the variable or structure member.

section ("section-name")

Place the variable into the named section.

For example,

```
unsigned int dan __attribute__ ((section (".quixote")))
```

Variable `dan` will be placed in section `.quixote`.

The `-fdata-sections` command line option has no effect on variables defined with a `section` attribute unless `unique_section` is also specified.

space(memory-space)

The `space` attribute can be used to direct the compiler to allocate a variable in a specific memory space. Valid memory spaces are `prog` for program memory, `data` for data memory, and `serial_mem` for serial memory such as SPI Flash. The `data` space is the default space for non-const variables.

The `prog`, `data`, and `serial_mem` spaces normally correspond to the `kseg0_prog_mem`, `ksegN_data_mem`, and `serial_mem` memory regions, respectively, as specified in the default device-specific linker scripts.

This attribute also controls how initialized data is handled. The linker generates an entry in the data-initialization template for the default `space(data)`. But, it does not generate an entry for `space(prog)` or `space(serial_mem)`, since the variable is located in non-volatile memory. Typically, this means that `space(data)` applies to variables that will be initialized at runtime startup; while `space(prog)` and `space(serial_mem)` apply to variables that will be programmed by an in-circuit programmer or a bootloader. For example,

```
const unsigned int __attribute__((space(prog))) jack = 10;
const unsigned int __attribute__((space(serial_mem))) zori = 1;
signed int __attribute__((space(data))) oz = 5;
```

unique_section

Place the variable in a uniquely named section, just as if `-fdata-sections` had been specified. If the variable also has a `section` attribute, use that section name as the prefix for generating the unique section name.

For example,

```
int tin __attribute__ ((section (".ofcatfood"), unique_section)
```

Variable `tin` will be placed in section `.ofcatfood`.

unused

Indicate to the compiler that the variable may not be used. The compiler will not issue a warning for this variable if it is not used.

weak

The `weak` attribute causes the declaration to be emitted as a weak symbol. A weak symbol indicates that if a global version of the same symbol is available, that version should be used instead.

When `weak` is applied to a reference to an external symbol, the symbol is not required for linking. For example:

```
extern int __attribute__((weak)) s;
int foo() {
  if (&s) return s;
  return 0; /* possibly some other value */
}
```

In the above program, if `s` is not defined by some other module, the program will still link but `s` will not be given an address. The conditional verifies that `s` has been defined (and returns its value if it has). Otherwise `'0'` is returned. There are many uses for this feature, mostly to provide generic code that can link with an optional library.

10. Memory Allocation and Access

There are two broad groups of RAM-based variables: auto/parameter variables, which are allocated to some form of stack, and global/static variables, which are positioned freely throughout the data memory space. The memory allocation of these two groups is discussed separately in the following sections.

10.1 Address Spaces

Unlike the 8- and 16-bit PIC devices, the PIC32 has a unified programming model. PIC32 devices provide a single 32-bit wide address space for all code, data, peripherals and Configuration bits.

Memory regions within this single address space are designated for different purposes; for example, as memory for instruction code or memory for data. Internally the device uses separate buses* to access the instructions and data in these regions, thus allowing for parallel access. The terms program memory and data memory, which are used on the 8- and 16-bit PIC devices, are still relevant on PIC32 devices, but the smaller parts implement these in different address spaces.

All addresses used by the CPU within the device are virtual addresses. These are mapped to physical addresses by the system control processor (CP0).

* The device can be considered a Harvard architecture in terms of its internal bus arrangement.

10.2 Variables in Data Memory

Most variables are ultimately positioned into the data memory. The exceptions are non-`auto` variables which are qualified as `const`, which are placed in the program memory space, see 9.9.1 Const Type Qualifier.

Due to the fundamentally different way in which auto variables and non-`auto` variables are allocated memory, they are discussed separately. To use the C/C++ language terminology, these two groups of variables are those with automatic storage duration and those with permanent storage duration, respectively.

Note: The terms "local" and "global" are commonly used to describe variables, but are not ones defined by the language standard. The term "local variable" is often taken to mean a variable which has scope inside a function, and "global variable" is one which has scope throughout the entire program. However, the C/C++ language has three common scopes: block, file (that is, internal linkage) and program (that is, external linkage), so using only two terms to describe these can be confusing. For example, a `static` variable defined outside a function has scope only in that file, so it is not globally accessible, but it can be accessed by more than one function inside that file, so it is not local to any one function either. In terms of memory allocation, variables are allocated space based on whether it is an `auto` or not, hence the grouping in the following sections.

10.2.1 Non-Auto Variable Allocation

Non-`auto` variables (those with permanent storage duration) are located by the compiler into any of the available data banks. This is done in a two-stage process: placing each variable into an appropriate section and later linking that section into data memory.

The compiler considers three categories of non-`auto` variable which all relate to the value the variable should contain by the time the program begins. The following sections are used for the categories described.

- `.pbss` - These sections are used to store variables which use the `persistent` attribute, whose values should not be altered by the runtime start-up code. They are not cleared or otherwise modified at start-up.
- `.bss` - These sections (also `.sbss`) contain any uninitialized variables, which are not assigned a value when they are defined, or variables which should be cleared by the runtime start-up code.
- `.data` - These sections (also `.sdata`) contain the RAM image of any initialized variables, which are assigned a non-zero initial value when they are defined and which must have a value copied to them by the runtime start-up code.

Note that the data section used to hold initialized variables is the section that holds the RAM variables themselves. There is a corresponding section (called `.dinit`) that is placed into program memory (so it is non-volatile) and which is used to hold the initial values that are copied to the RAM variables by the runtime start-up code.

10.2.2 Static Variables

All `static` variables have permanent storage duration, even those defined inside a function which are "local static" variables. Local `static` variables only have scope in the function or block in which they are defined, but unlike `auto` variables, their memory is reserved for the entire duration of the program. Thus, they are allocated memory like other non-`auto` variables. Static variables may be accessed by other functions via pointers, since they have permanent duration.

Variables which are `static` are guaranteed to retain their value between calls to a function, unless explicitly modified via a pointer.

Variables which are `static` and initialized have their initial value assigned only once during the program's execution. Thus, they may be preferable over initialized `auto` objects which are assigned a value every time the block they are defined in begins execution. Any initialized static variables are initialized in the same way as other non-`auto` initialized objects by the runtime start-up code, see 6.4.3 Peripheral Library Functions. Static variables are located in the same sections as their non-`static` counterparts.

10.2.3 Non-Auto Variable Size Limits

Arrays of any type (including arrays of aggregate types) are fully supported by the compiler. So too are the structure and union aggregate types, see 9.5 Structures and Unions. There are no theoretical limits as to how large these objects can be made.

10.2.4 Changing the Default Non-Auto Variable Allocation

There are several ways in which non-`auto` variables can be located in locations other than the default.

Variables can be placed in other device memory spaces by the use of qualifiers. For example if you wish to place variables in the program memory space, then the `const` specifier should be used (see 9.9.1 Const Type Qualifier).

If you wish to prevent all variables from using one or more data memory locations so that these locations can be used for some other purpose, it is best to define a variable (or array) using the `address` attribute so that it consumes the memory space, see 9.11 Variable Attributes .

If only a few non-`auto` variables are to be located at specific addresses in data space memory, then the variables can be located using the `address` attribute. This attribute is described in 9.11 Variable Attributes.

10.2.5 Data Memory Allocation Macros

The `sys/attribs.h` header file provides many macros for commonly used attributes in order to enhance code readability.

`__section__(s)`	Apply the `section` attribute with section name `s`.
`__unique_section__`	Apply the `unique_section` attribute.
`__ramfunc__`	Locate the attributed function in the RAM function code section.
`__longramfunc__`	Locate the attributed function in the RAM function code section and apply the `longcall` attribute.
`__longcall__`	Apply the `longcall` attribute.
`__ISR(v,ipl)`	Apply the `interrupt` attribute with priority level `ipl` and the `vector` attribute with vector number `v`.
`__ISR_AT_VECTOR(v,ipl)`	Apply the interrupt attribute with priority level `ipl` and the `at_vector` attribute with vector number `v`. This macro is especially useful on PIC32 devices that feature variable vector offsets.
`__ISR_SINGLE__`	Specifies a function as an Interrupt Service Routine in single-vector mode. This places a jump at the single-vector location to the interrupt handler.

`__ISR_SINGLE_AT_VECTOR__`	Places the entire single-vector interrupt handler at the vector 0 location. When used, ensure that the vector spacing is set to accommodate the size of the handler.

10.3   Auto Variable Allocation and Access

This section discusses allocation of `auto` variables (those with automatic storage duration). This also includes function parameter variables, which behave like `auto` variables, as well as temporary variables defined by the compiler.

The `auto` (short for *automatic*) variables are the default type of local variable. Unless explicitly declared to be `static`, a local variable will be made `auto`. The `auto` keyword may be used if desired.

The `auto` variables, as their name suggests, automatically come into existence when a function is executed, then disappear once the function returns. Since they are not in existence for the entire duration of the program, there is the possibility to reclaim memory they use when the variables are not in existence and allocate it to other variables in the program.

The software stack of the PIC32 is used to store all `auto` variables. Functions are reentrant and each instance of the function has its own area of memory on the stack for its auto and parameter variables, as described below. See 8.3 Stack and 16.2.3 Initialize Stack Pointer and Heap for more information on the stack.

The compiler dedicates General Purpose Register 29 as the software Stack Pointer. All processor stack operations, including function call, interrupts and exceptions use the software stack. The stack grows downward from high addresses to low addresses.

By default, the size of the stack is 1024 bytes. The size of the stack may be changed by specifying the size on the linker command line using the `--defsym_min_stack_size` linker command line option. An example of allocating a stack of 2048 bytes using the command line is:

```
xc32-gcc foo.c -Wl,--defsym,_min_stack_size=2048
```

The run-time stack grows downward from higher addresses to lower addresses (see the figure below). The compiler uses two working registers to manage the stack:

- Register 29 (`sp`) – This is the Stack Pointer. It points to the next free location on the stack.
- Register 30 (`fp`) – This is the Frame Pointer. It points to the current function's frame. Each function, if required, creates a new frame from which automatic and temporary variables are allocated. Compiler optimization may eliminate Stack Pointer references via the Frame Pointer to equivalent references via the Stack Pointer. This optimization allows the Frame Pointer to be used as a General Purpose Register.

Figure 10-1. Stack Frame

The the standard qualifiers const and volatile may both be used with auto variables and these do not affect how they are positioned in memory. This implies that a local const-qualified object is still an auto object and, as such, will be allocated memory on the stack in the data memory, not in the program memory like with non-auto const objects.p

Local Variable Size Limits

There is no theoretical maximum size for auto variables.

10.4   Variables in Program Memory

The only variables that are placed into program memory are those that are not `auto` and which have been qualified `const`. If the `-mno-embedded-data` option is used, then even `const` objects are placed in RAM rather than the program memory. Any `auto` variables qualified `const` are placed on the stack along with other `auto` variables.

Note: For devices with internal Flash, a const-qualified variable combined with the default `-membedded-data` option can be used to place a constant in nonvolatile Flash memory.

Any `const`-qualified (`auto` or non-`auto`) variable will always be read-only and any attempt to write to these in your source code will result in an error being issued by the compiler.

A `const` object is usually defined with initial values, as the program cannot write to these objects at runtime. However this is not a requirement. An uninitialized `const` object is allocated space in the bss section, along with other uninitialized RAM variables, but is still treated as read-only by the compiler.

```
const char IOtype = 'A';  // initialized const object
const char buffer[10];    // I just reserve memory in RAM
```

10.4.1 Size Limitations of `const` Variables

There is no theoretical maximum size for const variables.

10.4.2 Changing the Default Allocation

If you only intend to prevent all variables from using one or more program memory locations so that you can use those locations for some other purpose, you are best reserving the memory using the memory adjust options.

If only a few non-auto `const` variables are to be located at specific addresses in program space memory, then the variables should use the `address` attribute to locate them at the desired location. This attribute is described in 9.11 Variable Attributes .

10.5 Variable in Registers

Allocating variables to registers, rather than to a memory location, can make code more efficient. With MPLAB XC32 C/C++ Compiler, variables may be allocated to registers as part of code optimizations. For optimization levels 1 and higher, the values assigned to variables may cached in a register. During this time, the memory location associated with the variable may not hold a valid value.

The `register` keyword may be used to indicate your preference for the variable to be allocated a register, but this is just a recommendation and may not be honored. The specific register may be indicated as well, but this is not recommended as your register choice may conflict with the needs of the compiler. Using a specific register in your code may cause the compiler to generate less efficient code.

As indicated in 14.6 Function Parameters, parameters may be passed to a function via a register.

Example 10-1. Variables in Registers

```
volatile unsigned int special;
unsigned int example (void)
{
  register unsigned int my_reg __asm__ ("$4");
  my_reg += special;
  return my_reg;
}
```

10.6 Application-Defined Memory Regions

On occasion, an application may require a new memory region that was not defined in the default device-specific linker scripts. One such case may be when using external memory connected to the External Bus Interface (EBI) on a PIC32MZ target device.

One way to handle adding a new memory region would be to create a custom linker script, add the new memory region, and explicitly map your sections to the new region. Another way to add a new memory region would be to add an application-defined memory region to your C/C++ source code. Some applications may even choose to combine a custom linker script and an application-defined memory region.

10.6.1 Advantages of an Application-Defined Memory Region

Using an application-defined memory region in source code can provide a few advantages over adding a memory region to a custom linker script.

- Portability: An application-defined memory region can help reduce the need for a custom linker script. This can be beneficial because you can use the device's default linker script and thereby avoid potential migration issues between XC32 versions, as well as between different PIC32 variants.
- Best-Fit Allocation: Region-attributed variables and functions are handled by the linker's best-fit allocator, as described in the *MPLAB XC32 Assembler, Linker, and Utilities User's Guide* (DS50002186). Sections mapped to a new region in a custom linker script must be explicitly mapped using the `SECTIONS` command and are allocated sequentially.

10.6.2   Advantages of a Linker Script-Defined Memory Region

Using a custom linker script with a new memory region can also provide advantages over an application-defined memory region.

- Update memory mapping at link time: A custom linker script allows you to easily switch between different memory mappings without rebuilding your C/C++ code.
- Standard GNU Binutils: A user migrating from another toolchain based on GNU Binutils may already be familiar and comfortable with creating new memory regions in a custom linker script. XC32 supports this standard mechanism.

10.6.3   Using an Application-Defined Memory Region

To use this feature, work through the following sections.

Define a New Memory Space

The XC32 toolsuite requires information about each memory region. In order for the XC32 linker to be able to properly assign memory, you must specify information about the size of the memory region available and the origin of the memory region.

Define an application-defined memory region, with the origin and the size, using the **`'region'`** pragma as shown below.

```
#pragma region name=name origin=address size=bytes
```

where *name* is a quoted string containing the name of the region, *address* is the starting address of the region, and *bytes* is the size in bytes of the region.

> **Example 10-2.   Defining a New Memory Space**
> ```
> #pragma region name="ext_mem" origin=0xC0000000 size=0x1000
> ```

In this example, we define an application-defined memory region to be used for external memory. We name the region "`ext_mem`" and specify that the starting address is `0xC0000000` and that it has a size of `0x1000` bytes. Consult your PIC32 device data sheet for information about the external-memory interface options and the memory mappings available on your device.

Define Variables within a Region

When you have defined a new memory region, you can then assign a variable to that region. Use the region attribute on a variable to specify that it should be allocated to the specified region. This requires the memory region definition to be present. Given the definition in the previous subsection, you can make the following variable definition:

```
int ext_array[256] __attribute__((region("ext_mem")));
```

`ext_array` will be allocated in the previously declared region **`"ext_mem"`**.

Once the variable has been defined with the region attribute, it may be accessed using normal C syntax.

When called with the `--report-mem` linker command-line option, the linker prints a summary of memory usage to stdout. When the application-defined memory region is used, the length of each section allocated to a region and the total region memory used is displayed.

Define Functions within a Region

You can also use the region attribute to assign individual functions to the region. This requires the memory region definition to be present. Given the definition in the previous subsection, you can make the following function definition:

```
int __attribute__((far, region("ext_mem"))) foo()
{
   ext_array[2] = ext_array[0] + ext_array[1] ;
   return 0;
}
```

Use the region attribute with the far attribute to allocate the function in our example "`ext_mem`" region. In this case, we need the `far` attribute because the address of our "`ext_mem`" region is located outside of the 256 MB segment of our default program-memory region, `kseg0_program_mem`, as defined in our default linker script. Using the `far` attribute tells the compiler to generate a long call.

Initializing Memory Interfaces

When your application-defined memory region corresponds to an external-memory interface such as the Serial Quad Interface (SQI) or External Bus Interface (EBI), you will likely need to configure the interface module. For instance, the EBI module must be configured to understand such things as the type, size, and bus width of each attached device. See the device data sheet and the family reference manual for your target device.

The default XC32 runtime start-up code uses a linker-generated data-initialization template placed in a section named .dinit. (See 16.2.5 Initialize of Clear Variable and RAM Functions Using the Data-Initialization Template). For variables or functions placed in an application-defined memory region, the application must execute any memory-interface configuration code before the runtime start-up code attempts to initialize these variables or functions.

The default runtime start-up code provides an `_on_reset()` weak hook. This routine is called after initializing a minimum 'C' context, but before data initialization. You can provide your memory-interface configuration code in this hook. See 16.3 The On Reset Routine for more information on this important hook.

Example 10-3. Hardware Init Before Data Init

```
/* The _on_reset() function will be called by the default
   runtime start-up code prior to data initialization. */
void _on_reset (void)
{
   /*  Call a function that configures the EBI control
       registers for the target board. */
   configure_ebi_sram();
}
```

On some target devices, your application may also need to enable the Memory Management Unit (MMU) and initialize the Translation Lookaside Buffer (TLB). On many devices, the XC32 toolchain provides a default mapping suitable for the SQI and EBI interfaces. See your target device data sheet for information on default memory mapping that is specific to your target device. For devices where a default SQI and EBI mapping is provided, you can override the default mapping by providing your own `__pic32_tlb_init_ebi_sqi()` function.

The source code for this is found in the pic32m-libs.zip file located at:

`<install-directory>/pic32-libs/`

Once the file is unzipped, the source code can be found at:

`pic32m-libs/libpic32/stubs/pic32_init_tlb_ebi.S.`

The following sections provide example cases using the application-defined memory region feature.

Example 10-4. Case 1

Variables can be placed in external memory by using the region attribute.

```
#pragma region name="ext_mem" origin=0x0xC0000000 size=0x1000

signed int   ea1  __attribute__((region("ext_mem")));
unsigned int ea2  __attribute__((region("ext_mem")));
```

```
signed int   eb1[10]  __attribute__((region("ext_mem"))) = {10,20};
signed long  eb2[10]  __attribute__((region("ext_mem"))) =
                                        {0x987654321, 0x12345678};
```

Example 10-5. Case 2

Functions can be placed in external memory by using the region attribute. Since functions default to `space(prog)`, the function is assumed to be programmed into the region and will not be initialized by the runtime start-up code.

```
#pragma region name="ext_flash" origin=0xC0000000 size=0x1000

int ea1 __attribute__((region("ext_flash"))) ;
int eb1 __attribute__((region("ext_flash"))) = 0x1000 ;
int ec1 __attribute__((region("ext_flash"))) = 0x2000 ;

void __attribute__((region("ext_flash"))) foo()
{
    ea1 = eb1 + ec1 ;
}

void main()
{
    foo();
}
```

Apply the `far` attribute to `foo()`, since it is out of range of the default `kseg0_program_mem` region. Alternatively, use `-mlong-calls` option to compile the above example.

Example 10-6. Case 3

Combine the region attribute with the space(data) attribute to indicate that the function code should be initialized by the runtime start-up code's data initialization template. In this case, the code for the function is contained in the data-initialization template and copied to the memory region at startup.

```
#pragma region name="myebi_sram"  origin=0xC0004000  size=0x100
void __attribute__((far,space(data),region("myebi_sram"))) fn_in_sram()
{ /* Code here */ }
```

Example 10-7. Case 4

Combine the region attribute with the address() attribute to place a variable at an absolute address within the region.

```
#pragma region name="myebi_2"  origin=0xC0001000  size=0x10
unsigned long __attribute__((region("myebi_2"),address(0xC0001000))) paws =
0xAAAABBBB;
```

10.7   Dynamic Memory Allocation

The run-time heap is an uninitialized area of data memory that is used for dynamic memory allocation using the standard C library dynamic memory management functions, `calloc`, `malloc` and `realloc` along with the C++ new operator. Most C++ applications will require a heap.

If you do not use any of these functions, then you do not need to allocate a heap. By default, a heap is not created.

In MPLAB X, you can specify a heap size in the project properties for the xc32-ld linker. MPLAB X will automatically pass the option to the linker when building your project.

If you do want to use dynamic memory allocation, either directly, by calling one of the memory allocation functions, or indirectly, by using a standard C library function that uses one of these functions, then a heap must be created. A heap is created by specifying its size on the linker command line using the `--defsym,min_heap_size` linker command line option. An example of allocating a heap of 512 bytes using the command line is:

```
xc32-gcc foo.c -Wl,--defsym,min_heap_size=512
```

An example of allocating a heap of `0xF000` bytes using the xc32-g++ driver is:

```
xc32-g++ vector.cpp -Wl,--defsym,min_heap_size=0xF000
```

The linker allocates the heap immediately before the stack.

10.8   Memory Models

MPLAB XC32 C/C++ Compiler does not use fixed memory models to alter allocation of variables to memory.

The `-G` option (see 6.7.1 Options Specific to PIC32M Devices), which controls the gp-relative addressing threshold, is similar to the small-data/large-data/scalar-data memory models offered by the Microchip compilers for the 8- and 16-bit architectures. The value specified with this option indicates the maximum size of objects that will be allocated to the small data sections, for example, `sbss`, `sdata`, etc. Variables allocated to the small-data sections require fewer instructions to access than variables allocated to the other data sections. For example:

```
xc32-gcc -G128 -mprocessor=32MX795F512L main.c
```

In this example, data objects up to 128 bytes in size will be located in the efficient small-data or small-bss section.

In general larger `-G` values result in more efficient code. However, gp-relative addressing is limited to 64-KB of small data.

11.   Fixed-Point Arithmetic Support

The MPLAB XC32 C/C++ Compiler supports fixed-point arithmetic. This, according to the N1169 draft of ISO/IEC TR 18037, the ISO C99 technical report on Embedded C. It is available at: http://www.open-std.org/JTC1/SC22/WG14/www/projects#18037.

This chapter describes the implementation-specific details of the types and operations supported by the compiler under this draft standard.

Because of the sensitivity of DSP applications to performance, application developers have historically tended to write functions in assembly. However, the XC32 compiler reduces, and may even eliminate, the need to write assembly code. This chapter describes coding styles and usage tips that can help you to obtain the best optimizations for your DSP application and to take advantage of the DSP-enhanced core that is available on many of the PIC32 MCU families.

Several Microchip PIC32 MCUs feature a DSP-enhanced core with four 64-bit accumulators. The DSP-enhanced core also provides a set of new instructions and a new architectural state, with computational support for fractional data types, SIMD (Single Instruction, Multiple Data), saturation, and other operations commonly used in DSP applications.

Note: Consult the data sheet for your specific target device to determine whether your target device supports the DSP-enhanced core.

11.1   Enabling Fixed-Point Arithmetic Support

Fixed-point arithmetic support is enabled by default in the MPLAB XC32 C/C++ compiler. The compiler automatically enables support for the DSP-enhanced core when you are compiling for an appropriate PIC32 target device as selected by the `-mprocessor` option.

11.2   Data Types

All 12 of the primary fixed-point types and their aliases, described in Section 4.1 "Overview and principles of the fixed-point data types" of the N1169 draft of ISO/IEC TR 18037, are supported. Fixed-point data values contain fractional and optional integral parts. The format of fixed-point data in XC32 are as specified in the table below.

In the formats shown, "s" is the sign bit for a signed type (there is no sign bit for an unsigned type). The period character "." is the specifier that separates the integral part and the fractional part. The numeric digits represent the number of bits in the integral part or in the fractional part.

Table 11-1. Fixed Point Formats

Type	Format	Description
`short _Fract`	s.7 (Q7)	1 bit sign, 7 bits fraction
`unsigned short _Fract`	0.8	8 bits fraction
`_Fract`	s.15 (Q15)	1 bit sign, 15 bits fraction
`unsigned _Fract`	0.16	16 bits fraction
`long _Fract`	s.31 (Q31)	1 bit sign, 31 bits fraction
`unsigned long _Fract`	0.32	32 bits fraction
`long long _Fract`	s.63 (Q63)	1 bit sign, 63 bits fraction
`unsigned long long _Fract`	0.64	64 bits fraction
`short _Accum`	s8.7 (Q8.7)	1 bit sign, 8 bits integer, 7 bits fraction
`unsigned short _Accum`	8.8	8 bits integer, 8 bits fraction

..........continued

Type	Format	Description
`_Accum`	s16.15 (Q16.15)	1 bit sign, 16 bits integer, 15 bits fraction
`unsigned _Accum`	16.16	16 bits integer, 16 bits fraction
`long _Accum`	s32.31 (Q32.31)	1 bit sign, 32 bits integer, 31 bits fraction
`unsigned long _Accum`	32.32	32 bits integer, 32 bits fraction
`long long _Accum`	s32.31 (Q32.31)	1 bit sign, 32 bits integer, 31 bits fraction
`unsigned long long _Accum`	32.32	32 bits integer, 32 bits fraction

The `_Sat`-type specifier, indicating that the values are saturated, may be used with any type as described in the N1169 draft of ISO/IEC TR 18037. For example:

`_Sat short _Fract`; // s.7 (Q7)

`_Sat _Fract`; // s.15 (Q15)

`_Sat long _Fract`; // s.31 (Q31)

`_Sat long long _Fract`; // s.63 (Q63)

`_Sat short _Accum`; // s8.7 (Q8.7)

`_Sat _Accum`; // s16.15 (Q16.15)

`_Sat long _Accum`; // s32.31 (Q32.31)

`_Sat long long _Accum`; // s32.31 (Q32.31)

The representation of unsigned types differ from their corresponding signed types due to the fact that one bit is utilized to store the sign. Signed types saturate at the most negative and positive numbers representable in the corresponding format. Unsigned types saturate at zero and the most positive number representable in the format.

The default behavior of overflow on signed or unsigned types is saturation. The pragmas described in Section 4.1.3 "Rounding and Overflow" of the N1169 draft of ISO/IEC TR 18037 to control the rounding and overflow behavior are not supported.

Use a suffix in a fixed-point literal constant:

- `hr` or `HR` for `short _Fract` and `_Sat short _Fract`
- `r` or `R` for `_Fract` and `_Sat _Fract`
- `lr` or `LR` for `long _Fract` and `_Sat long _Fract`
- `llr` or `LLR` for `long long _Fract` and `_Sat long long _Fract`
- `uhr` or `UHR` for `unsigned short _Fract` and `_Sat unsigned short _Fract`
- `ur` or `UR` for `unsigned _Fract` and `_Sat unsigned _Fract`
- `ulr` or `ULR` for `unsigned long _Fract` and `_Sat unsigned long _Fract`
- `ullr` or `ULLR` for `unsigned long long _Fract` and `_Sat unsigned long long _Fract`
- `hk` or `HK` for `short _Accum` and `_Sat short _Accum`
- `k` or `K` for `_Accum` and `_Sat _Accum`
- `lk` or `LK` for `long _Accum` and `_Sat long _Accum`
- `llk` or `LLK` for `long long _Accum` and `_Sat long long _Accum`
- `uhk` or `UHK` for `unsigned short _Accum` and `_Sat unsigned short _Accum`
- `uk` or `UK` for `unsigned _Accum` and `_Sat unsigned _Accum`
- `ulk` or `ULK` for `unsigned long _Accum` and `_Sat unsigned long _Accum`

- `ullk` or `ULLK` for `unsigned long long _Accum` and `_Sat unsigned long long _Accum`

11.3   External Definitions

The MPLAB XC32 C compiler provides an include file, stdfix.h, which defines various pre-processor macros related to fixed-point support.

Fixed-point conversion specifiers for formatted I/O, as described in Section 4.1.9 "Formatted I/O functions for fixed-point arguments" of the N1169 draft of ISO/IEC TR 18037, are not supported in the current MPLAB XC32 standard C libraries. Fixed-point variables may be displayed via `(s)printf` by casting them to the appropriate floating-point representation (float for `_Fract`, double for `long _Fract` and `_Accum`) and then displaying the value in that format. To scan a fixed-point number via `(s)scanf`, first scan it as the appropriate float or double floating-point number and then cast the value obtained to the desired fixed-point type.

The fixed-point functions described in Section 4.1.7 of N1169 are not provided in the current MPLAB XC32 standard C libraries.

Fixed-point constants, with suffixes of `k` (`K`) and `r` (`R`), as described in Section 4.1.5 of N1169, are supported by the MPLAB XC32 C compiler.

Usage example:

```
#include <stdfix.h>
void main () {
  int i;
  fract a[5] = {0.5,0.4,0.2,0.0,-0.1};
  fract b[5] = {0.1,0.8,0.6,0.5,-0.1};
  accum dp = 0;
  /* Calculate dot product of a[] and b[] */
  for (i=0; i<5; i++) {
      dp += a[i] * b[i];
  }
}
```

Integer representation of Q15 and Q31

The Q15 data type can be represented by the 16-bit integer data type (`short`) and the Q31 data type can be represented by the 32-bit integer data type (`int`). These types are necessary when using the compiler's DSP built-in functions (see 25. Built-In Functions). Typedefs are useful for Q15 and Q31 as follows:

```
typedef short q15;
typedef int q31;
```

The four 64-bit accumulators in the DSP-enhanced core can be represented by the "long long" data type.

```
typedef long long a64;
```

To initialize Q15 variables, multiply the fractional value (for example, 0.1234) by `0x1.0p15`. To initialize Q31 variables, programmers can multiply the fractional value by `0x1.0p31`.

```
Ex: /* Q15 Example  */
typedef short q15;
q15 a = 0.1234 * 0x1.0p15;

/* ----------------------------------------------------- */
Ex: /* Q31 Example  */
typedef int q31;
q31 b = 0.2468 * 0x1.0p31;
```

11.4   SIMD Variables

The 8-bit unsigned integer data and Q15 fractional data are packed in a single 32-bit register, and the new instructions operate simultaneously on the multiple data in the register in Single Instruction, Multiple Data (SIMD) fashion. This feature provides computation parallelism for increased application performance.

You can directly take advantage of SIMD parallelism by declaring SIMD data types as described here. In addition, the compiler may automatically take advantage of SIMD parallelism using auto-vectorization as described in the Auto-vectorization section below.

To declare SIMD data types, typedefs with special `vector_size` attributes are required. For example,

```
typedef signed char v4i8 __attribute__ ((vector_size(4)));
typedef signed char v4q7 __attribute__ ((vector_size(4)));
typedef short v2i16 __attribute__ ((vector_size(4)));
typedef short v2q15 __attribute__ ((vector_size(4)));
```

`v4i8`, `v4q7`, `v2i16`, and `v2q15` are SIMD data types that consist of 4, 4, 2, and 2 elements of 8 bits, 8 bits, 16 bits, and 16 bits respectively in a single variable/register.

SIMD data types are powerful and can be applied to fixed-point data types as well. For example:

```
typedef _Sat unsigned short _Fract
  sat_v4uqq __attribute__ ((vector_size(4)));
typedef _Sat unsigned _Fract
  sat_v2uhq __attribute__ ((vector_size(4)));
typedef _Sat unsigned short _Accum
  sat_v2uha __attribute__ ((vector_size(4)));
typedef _Sat _Fract
  sat_v2hq __attribute__ ((vector_size(4)));
typedef _Sat short _Accum
  sat_v2ha __attribute__ ((vector_size(4)));
```

To initialize SIMD variables is similar to initializing aggregate data. The following examples show how to initialize SIMD variables.

Example:

```
/* v4i8 Example */
v4i8 a = {1, 2, 3, 4};
v4i8 b;
b = (v4i8) {5, 6, 7, 8};
/* ---------------------------------------------------- */
Ex: /* v2q15 Example */
v2q15 a = {0x0fcb, 0x3a75};
v2q15 b;
b = (v2q15) {0.1234 * 32768.0, 0.4567 * 32768.0};
```

Data is stored from the right-to-left location of a register. For the example of v4i8 a = {1, 2, 3, 4}, the register stores 1, 2, 3, and 4 from the right-to-left location as shown below.

a[3]=4	a[2]=3	a[1]=2	a[0]=1

Bit 31   Bit 0

Most arithmetic operations will simply work on the SIMD operands in the register irrespective of their right-to-left location within the register. However, you must be aware of such instructions that directly refer to the left or right portions of a register. For example, `MAQ_SA.W.PHL`.

11.5 Accessing Elements in SIMD Variables

The use of SIMD variables enables operations on multiple data in parallel. However, in certain situations, programmers need to access elements inside a SIMD variable. This can be done by using a union type that unites a SIMD type and an array of a basic type as follows.

```
typedef union
{
  v4i8 a;
  unsigned char b[4];
} v4i8_union;
typedef short q15;
typedef union
```

```
  {
    v2q15 a;
    q15 b[2];
  } v2q15_union;
```

As shown in the figure above for a `v4i8` variable, `b[0]` is used to access the first element in the variable. The element `b[0]` is right-most position. The following examples show how to extract or assign elements.

Example:

```
/* v4i8 Example */
v4i8 i;
unsigned char j, k, l, m;
v4i8_union temp;
/* Assume we want to extract from i.  */
temp.a = i;
j = temp.b[0];
k = temp.b[1];
l = temp.b[2];
m = temp.b[3];
/* Assume we want to assign j, k, l, m to i.  */
temp.b[0] = j;
temp.b[1] = k;
temp.b[2] = l;
temp.b[3] = m;
i = temp.a;
/* ------------------------------------------------------- */
```

Example:

```
/* v2q15 Example  */
v2q15 i;
q15 j, k;
v2q15_union temp;
/* Assume we want to extract from i.  */
temp.a = i;
j = temp.b[0];
k = temp.b[1];
/* Assume we want to assign j, k to i.  */
temp.b[0] = j;
temp.b[1] = k;
i = temp.a;
```

Using SIMD data types is a very powerful technique. Programmers can enjoy the performance improvement from SIMD data types by calling the DSP built-in functions (see 25. Built-In Functions) and/or using generic C operators. For SIMD data types, the compiler can map C operators (e.g., +, -, *, /) to hardware instructions directly, so long as the selected target PIC32 MCU features the DSP-enhanced core.

Note: In many cases, optimization level -O1 or greater may be required to optimize the code to use the SIMD instruction.

Here are some examples:

```
typedef signed char v4i8 __attribute__ ((vector_size(4)));
v4i8 a, b, c;
c = a + b; // compiler generates addu.qb
c = a - b; // compiler generates subu.qb
/* ------------------------------------------------------- */
typedef short v2q15 __attribute__ ((vector_size(4)));
v2q15 d, e, f;
f = d + e; // compiler generates addq.ph
f = d - e; // compiler generates subq.ph
/* ------------------------------------------------------- */
typedef short v2i16 __attribute__ ((vector_size(4)));
v2i16 x, y, z;
z = x * y; // compiler generates mul.ph
/* ------------------------------------------------------- */
typedef _Sat _Fract sat_v2hq __attribute__ ((vector_size(4)));
sat_v2hq a, b, c;
c = a + b; // compiler generates addq_s.ph
```

```
c = a - b; // compiler generates subq_s.ph
c = a * b; // compiler generates mulq_rs.ph
```

Note: When char or short data elements are packed into SIMD data types, the first data must be aligned to 32 bits; otherwise, the unaligned memory accesses may generate general exceptions or decrease performance.

11.6   Array Alignment and Data Layout

The compiler provides a mechanism to specify the alignment of variables by using `__attribute__ ((aligned (bytes)))`. The alignment is important to loading or storing SIMD variables: "`v4i8`" and "`v2q15`". If an array is aligned to a 4-byte boundary, that is, word-aligned, the compiler can load or store four 8-bit data for `v4i8` variables (or two 16-bit data for `v2q15` variables) at a time using the load word class of instructions. The following example shows that when a char array A is aligned to a 4-byte boundary, we can cast this array to a `v4i8` array and load four items to a `v4i8` variable at a time by using the `lwx` instruction. However, if this char array A is not aligned to a 4-byte boundary, executing the following code will result in an address exception due to a mis-aligned load.

Example:

```
/* v4i8 Example */
   char A[128] __attribute__ ((aligned (4)));
   v4i8 test (int i)
   {
     v4i8 a;
     v4i8 *myA = (v4i8 *)A;
     a = myA[i];
     return a;
   }
# Illustrative generated assembly with optimizations
test:
   lui     $2,%hi(A)
   sll     $4,$4,2
   addiu   $2,$2,%lo(A)
   lwx      $2,$2($4)
   j       $31
```

After SIMD data is loaded from memory into a register, ensure that the SIMD variables in the register are ready for use without requiring any rearrangement of the data. To avoid such data rearrangement which can reduce the benefit of parallelism, design your array with an efficient data layout that is favorable for SIMD calculations.

11.7   C Operators

11.7.1   Operations on Fixed-Point Variables

Support for fixed-point types includes:

*   prefix and postfix increment and decrement operators (++, --)
*   unary arithmetic operators (+, -, !)
*   binary arithmetic operators (+, -, *, /)
*   binary shift operators (<<, >>)
*   relational operators (<, <=, >=, >)
*   equality operators (==, !=)
*   assignment operators (+=, -=, *=, /=, <<=, >>=)
*   conversions to and from integer, floating-point, or fixed-point types

This is an example of a multiplication operator.

Example:

```
#include <stdfix.h>
sat fract test (sat fract a, sat fract b)
{
    return (a * b);
```

```
    }
# Illustrative generated assembly with optimizations
test:
    mulq_rs.ph    $2,$4,$5
    j    $31
```

11.8   Operations on SIMD Variables

Some specific C operators can be applied to SIMD variables. They are +, -, *, /, unary minus, ^, |, &, ~. The DSP-enhanced core provides SIMD addition and subtraction instructions for v4i8 and v2q15, allowing the XC32 to generate appropriate instructions for addition and subtraction of v4i8 and v2q15 variables. For other operators, the compiler synthesizes a sequence of instructions. The examples here show compiler-generated SIMD instructions when the appropriate operator is applied to SIMD variables.

Example:

```
/* v4i8 Addition */
v4i8 test (v4i8 a, v4i8 b)
{
  return a + b;
}
# Illustrative generated assembly with optimizations
test:
  addu.qb $2, $4, $5
  j        $31
# ------------------------------------------------------------
```

Example:

```
/* v4i8 Subtraction */
v4i8 test (v4i8 a, v4i8 b)
{
  return a - b;
}
# Illustrative generated assembly with optimizations
test:
  subu.qb $2, $4, $5
  j $31
# ------------------------------------------------------------
```

Example:

```
/* v2q15 Addition */
v2q15 test (v2q15 a, v2q15 b)
{
  return a + b;
}
# Illustrative generated assembly with optimizations
test:
  addq.ph $2, $4, $5
  j $31
# ------------------------------------------------------------
```

Example:

```
/* v2q15 Subtraction */
v2q15 test (v2q15 a, v2q15 b)
{
  return a - b;
}
# Illustrative generated assembly with optimizations
test:
  subq.ph $2, $4, $5
  j        $31
```

In situations where your application requires special integer and fractional calculations and the compiler cannot generate them automatically, you can use the DSP built-in functions (see 25. Built-In Functions).

11.9 DSP Built-In Functions

Built-in functions are very similar to standard function calls in syntax. Your application passes parameters to a built-in function, and the built-in function returns the result to variables. The difference between a built-in function and a standard function is that the compiler directly can map the built-in function to a specific instruction sequence for better performance. The DSP built-in functions are listed in (see 25. Built-In Functions).

11.10 DSP Control Register

The DSP-enhanced core includes a DSP control register that has six fields: CCOND (condition code bits), OUFLAG (overflow/underflow bits), EFI (extract fail indicator bit), C (carry bit), SCOUNT (size count bits) and POS (position bits). The compiler treats the SCOUNT and POS fields as global variables, such that instructions that modify SCOUNT or POS are never optimized away. These instructions include WRDSP, EXTPDP, EXTPDPV, and MTHLIP. A function call that jumps to a function containing WRDSP, EXTPDP, EXTPDPV, or MTHLIP is also never deleted by the compiler.

For correctness, you must assume that a function call clobbers all fields of the DSP control register. That is, do not depend on the values in CCOND, OUFLAG, EFI or C across a function-call boundary. Re-initialize the values of CCOND, OUFLAG, EFI or C before using them. Note that because SCOUNT and POS fields are treated as global variables, the values of SCOUNT and POS are always valid across function-call boundaries and can be used without re-initialization.

The following example shows possibly incorrect code. The first built-in function "`__builtin_mips_addsc`" sets the carry bit (C) in the DSP control register, and the second built-in function "`__builtin_mips_addwc`" reads the carry bit (C) from the DSP control register. However, a function call "func" inserted between "`__builtin_mips_addsc`" and "`__builtin_mips_addwc`" may change the carry bit to affect the correct result of "`__builtin_mips_addwc`".

```
Incorrect Ex:
int test (int a, int b, int c, int d)
{
  __builtin_mips_addsc (a, b);
  func();   // may clobber the carry bit
  return __builtin_mips_addwc (c, d);
}
```

The previous example may be corrected by moving "func" before the first built-in function or after the second built-in function as follows. p

```
Corrected Ex:
int test (int a, int b, int c, int d)
{
  func();    // may affect the carry bit
  __builtin_mips_addsc (a, b);
  return __builtin_mips_addwc (c, d);
}
/* ------------------------------------------------------- */
int test (int a, int b, int c, int d)
{
  int i;
  __builtin_mips_addsc (a, b);
  i = __builtin_mips_addwc (c, d);
  func();    // may affect the carry bit
  return i;
}
```

11.11 Using Accumulators

To access only HI or LO of an accumulator, use a union type as follows.

```
typedef union
{
  long long a;   // One 64-bit accumulator
```

```
  int b[2];      // 32-bit HI and LO
} a64_union;
```

To access HI, use b[1]. To access LO, use b[0].

Example:

```
int test (long long a, v2q15 b, v2q15 c)
{
  a64_union temp;
  temp.a = __builtin_mips_dpaq_s_w_ph (a, b, c);
  return temp.b[0];  // access LO.
}
```

11.12 Mixed-Mode Operations

11.12.1 Multiply "32-bit int * 32-bit int = 64-bit long long int"
Multiply "32-bit int * 32-bit int = 64-bit long long int"

To multiply 32 bits by 32 bits to obtain a 64-bit result, cast the 32-bit integer to a 64-bit integer (long long) and then perform the multiplication operation as follows.

Example:

```
long long test (int a, int b)
{
  return (long long) a * b;
  // Same as (long long) a * (long long) b
  // NOT the same as (long long) (a * b)
}
```

We can then access the highest 32-bit result from HI as follows.

Example:

```
typedef union
{
  long long a;   // One 64-bit accumulator
  int b[2];      // 32-bit HI and LO
} a64_union;
int test (int a, int b)
{
  a64_union temp;
  temp.a = (long long) a * b;
  return temp.b[1];  // Access the HI 32 bits
}
# Illustrative generated assembly with optimizations
test:
  mult    $4,$5
  mfhi    $2
  j       $31
```

11.12.2 Multiply and Add "32-bit int * 32-bit int + 64-bit long long = 64-bit long long int"
To perform multiplication and addition, cast the 32-bit integer to 64-bit (long long) and then perform multiplication and addition as follows.

Example:

```
long long test (int a, int b, long long c)
{
  return c + (long long) a * b;
}
# Illustrative generated assembly with optimizations
test:
  mtlo $6
  mthi $7
  madd $4, $5
```

```
  mflo $2
  mfhi $3
  j $31
```

11.13 Auto-Vectorization to SIMD

The compiler supports auto-vectorization for loops at optimization level -O3. The advantage of auto-vectorization is that the compiler can recognize scalar variables (which can be integer, fixed-point, or floating-point types) in order to utilize SIMD (Single Instruction, Multiple Data) instructions automatically. In the ideal case, when auto-vectorization is used, there is no need to use SIMD variables explicitly.

Example:

```
/* add8.c */
unsigned char a[32], b[32], c[32];
void add8() {
  int i;
  for (i = 0; i < 32; i++)
  {
    c[i] = a[i] + b[i];
  }
}
# Illustrative generated assembly code
add8:
  lui    v0,0x0
  addiu  v0,v0,0
  lui    a0,0x0
  addiu  a0,a0,0
  lui    v1,0x0
  addiu  v1,v1,0
  addiu  a3,v0,32
  lw     a2,0(a0)
  lw     a1,0(v0)
  addiu  v0,v0,4
  addiu  a0,a0,4
  addu.qb a1,a2,a1
  addiu  v1,v1,4
  bne    v0,a3,1c <add8+0x1c>
  sw     a1,-4(v1)
  jr     ra
```

In `add8.c`, elements in two arrays of `unsigned char` are added together. The compiler automatically generates the code for `addu.qb` to add four elements at a time.

For existing C code, try auto-vectorization at the -O3 optimization level without any modifications to see if the compiler can auto vectorize the loops. In some cases, if the loop body is too complex, the compiler will not be able to auto-vectorize the loop; in this case, you may choose to restructure and simplify the loop.

11.14 FIR Filter Example Project

The DSP_Intrinsics example project shows the advantages of a FIR filter implementation based on DSP Built-in Functions compared to traditional C code, without any DSP optimizations. First, the filter is implemented for Q15 input data represented as arrays of short variables:

```
short coeff[BUFSIZE] __attribute__ ((aligned(4)));
short delay[BUFSIZE] __attribute__ ((aligned(4)));
```

The traditional C code version for Q15 inputs doesn't use SIMD variables or DSP Built-in Functions. Instead, it implements the Q15 x Q15 multiplication by checking the saturation condition (both operands are 0x8000 or -1) and then left shifting the integer multiplication result by 1 bit before adding it to the accumulator.

```
long long FIR_Filter_Traditional_16(short *delay, short *coeff, int buflen)
{
    int i;
    short x, y;
```

```
    // 64-bit accumulator for result
    long long ac0 = 0;

    for (i = 0; i < buflen; i++) {
        x = coeff[i];
        y = delay[i];

        // check saturation condition
        if ((unsigned short)x == 0x8000 && (unsigned short)y == 0x8000) {
            ac0 += 0x7fffffff;
        } else {
            // multiply (Q15 x Q15) needs left shift
            // result is added to accumulator variable
            ac0 += ((x * y) << 1);
        }
    }
    return ac0;
}
```

For this C implementation the compiler generates inefficient assembly code. It does not have enough information to auto-vectorize the loop.

```
# Illustrative generated assembly code
FIR_Filter_Traditional_16:
    …
    mul     $24,$13,$15
    sll     $25,$24,1
    addu    $12,$2,$25
    sra     $9,$25,31
    sltu    $11,$12,$2
    addu    $3,$3,$9
    move    $2,$12
    addu    $3,$11,$3
```

In the DSP Intrinsics approach of the filter the input buffers are casted to the v2q15 SIMD vector type defined above in section SIMD Variables. Inside the loop, the "__builtin_mips_dpaq_s_w_ph" DSP built-in function is called. The result of the call is stored in the accumulator variable, of type Q32.31 (64-bit), represented as integer type a64 (see definition in section Integer representation of Q15 and Q31).

```
a64 FIR_Filter_Intrinsics_16(short *delay, short *coeff, int buflen)
{
    int i;
    v2q15 *my_delay = (v2q15 *)delay;
    v2q15 *my_coeffs = (v2q15 *)coeff;
    // 64-bit accumulator for result
    a64 ac0 = 0;
    for (i = 0; i < buflen/2; i++) {
        ac0 = __builtin_mips_dpaq_s_w_ph (ac0,
my_delay[i],
my_coeffs[i]);
    }
    return ac0;
}
```

This function generates "dpaq_s.w.ph" assembly DSP instructions that apply the "Dot Product with Accumulate" operation on two sets of Q15 values. The result is stored in one of the four 64-bit accumulators in the DSP-enhanced core.

```
# Illustrative generated assembly code
FIR_Filter_Intrinsics_16:
    …
    mtlo    $0
    mthi    $0
    addiu   $8,$7,-1
    li      $6,1
    andi    $10,$8,0x7
    dpaq_s.w.ph     $ac0,$2,$3
    addiu   $3,$4,4
    beq     $6,$7,.L60
    addiu   $2,$5,4
```

The project targets the PIC32MZ2048EFM144 device. The tools used are MPLAB X IDE v3.10, MPLAB XC32 v1.40 compiler and PIC32 MZ EF Starter Kit/Simulator. Optimization level is set at "-O3" with "-funroll-loops" option enabled. Using an internal timer to count the ticks during calls to the 2 functions operating on the same Q15 data buffers reveals that the Intrinsics version is approximately **4.52 times faster** than the traditional C version.

Example timing output for input data buffers of size 2048:

```
16-bit without DSP Intrinsics: timer ticks 1180
16-bit with DSP Intrinsics: timer ticks: 261
```

The project includes filter implementations for buffers of 32-bit integer data types using similar approaches as for the previous 16-bit versions.

First implementation uses multiply and add operators. The 32-bit is casted to 64-bit before multiply as it was described previously in 11.12.2 Multiply and Add "32-bit int * 32-bit int + 64-bit long long = 64-bit long long int":

```
for (i = 0; i < buflen; i++) {
    acc += (long long)data[i] * coeff[i];
}

# Illustrative generated assembly code
FIR_Filter_32:
    …
    lwx    $24,$2($4)
    lwx    $25,$2($5)
    addiu    $2,$2,4
    madd    $ac0,$24,$25
```

The Intrinsics variant uses the "__builtin_mips_madd" to operate the multiply add.

```
for (i = 0; i < buflen; i++) {
    acc = __builtin_mips_madd (acc, data[i], coeff[i]);
}

# Illustrative generated assembly code
FIR_Filter_Intrinsics_32:
    …
    lw    $17,0($3)
    lw    $24,0($2)
    addiu    $7,$7,1
    addiu    $3,$3,4
    addiu    $2,$2,4
    madd    $ac0,$17,$24
```

The tick counts for the 32-bit implementations are very similar. In both cases the compiler is now generating the "MADD" DSP instructions.

Example timing output for input data buffers of size 2048:

```
32-bit without DSP Intrinsics: timer ticks 520
32-bit with DSP Intrinsics: timer ticks 484
```

12.   Operators and Statements

The MPLAB XC32 C/C++ Compiler supports all ANSI operators. The exact results of some of these are implementation-defined. Implementation-defined behavior is fully documented in 23. Implementation-Defined Behavior. The following sections illustrate code operations that are often misunderstood, as well as additional operations that the compiler is capable of performing.

12.1   Integral Promotion

When there is more than one operand to an operator, they typically must be of exactly the same type. The compiler will automatically convert the operands, if necessary, so they do have the same type. The conversion is to a "larger" type so there is no loss of information; however, the change in type can cause different code behavior to what is sometimes expected. These form the standard type conversions.

Prior to these type conversions, some operands are unconditionally converted to a larger type, even if both operands to an operator have the same type. This conversion is called *integral promotion* and is part of Standard C behavior. The MPLAB XC32 C/C++ Compiler performs these integral promotions where required, and there are no options that can control or disable this operation. If you are not aware that the type has changed, the results of some expressions are not what would normally be expected.

Integral promotion is the implicit conversion of enumerated types, `signed` or `unsigned` varieties of `char`, `short int` or bit field types to either `signed int` or `unsigned int`. If the result of the conversion can be represented by an `signed int`, then that is the destination type, otherwise the conversion is to `unsigned int`.

Consider the following example:

```
unsigned char count, a=0, b=50;
if(a - b < 10)
   count++;
```

The `unsigned char` result of `a - b` is 206 (which is not less than 10), but both `a` and `b` are converted to `signed int` via integral promotion before the subtraction takes place. The result of the subtraction with these data types is -50 (which *is* less than 10) and hence the body of the `if()` statement is executed.

If the result of the subtraction is to be an `unsigned` quantity, then apply a cast. For example:

```
if((unsigned int)(a - b) < 10)
   count++;
```

The comparison is then done using `unsigned int`, in this case, and the body of the `if()` would not be executed.

Another problem that frequently occurs is with the bitwise compliment operator, `~`. This operator toggles each bit within a value. Consider the following code:

```
unsigned char count, c;
c = 0x55;
if( ~c == 0xAA)
   count++;
```

If `c` contains the value 0x55, it often assumed that `~c` will produce 0xAA, however the result is 0xFFFFFFAA and so the comparison in the above example would fail. The compiler may be able to issue a mismatched comparison error to this effect in some circumstances. Again, a cast could be used to change this behavior.

The consequence of integral promotion as illustrated above is that operations are not performed with `char` -type operands, but with `int` -type operands. However there are circumstances when the result of an operation is identical regardless of whether the operands are of type `char` or `int`. In these cases, the MPLAB XC32 C/C++ Compiler will not perform the integral promotion so as to increase the code efficiency. Consider the following example:

```
unsigned char a, b, c;
a = b + c;
```

Strictly speaking, this statement requires that the values of b and c should be promoted to `unsigned int`, the addition performed, the result of the addition cast to the type of a, and then the assignment can take place. Even if the result of the `unsigned int` addition of the promoted values of b and c was different to the result of the `unsigned char` addition of these values without promotion, after the `unsigned int` result was converted back to `unsigned char`, the final result would be the same. If an 8-bit addition is more efficient than a 32-bit addition, the compiler will encode the former.

If, in the above example, the type of a was `unsigned int`, then integral promotion would have to be performed to comply with the ANSI C standard.

12.2   Type References

Another way to refer to the type of an expression is with the `typeof` keyword. This is a non-standard extension to the language. Using this feature reduces your code portability.

The syntax for using this keyword looks like `sizeof`, but the construct acts semantically like a type name defined with `typedef`.

There are two ways of writing the argument to `typeof`: with an expression or with a type. Here is an example with an expression:

```
typeof (x[0](1))
```

This assumes that x is an array of functions; the type described is that of the values of the functions.

Here is an example with a `typename` as the argument:

```
typeof (int *)
```

Here the type described is a pointer to `int`.

If you are writing a header file that must work when included in ANSI C programs, write _ _typeof_ _ instead of `typeof`.

A `typeof` construct can be used anywhere a `typedef` name could be used. For example, you can use it in a declaration, in a cast, or inside of `sizeof` or `typeof`.

- This declares y with the type of what x points to:
  ```
  typeof (*x) y;
  ```
- This declares y as an array of such values:
  ```
  typeof (*x) y[4];
  ```
- This declares y as an array of pointers to characters:
  ```
  typeof (typeof (char *)[4]) y;
  ```

  It is equivalent to the following traditional C declaration:

  ```
  char *y[4];
  ```

To see the meaning of the declaration using `typeof`, and why it might be a useful way to write, let's rewrite it with these macros:

```
#define pointer(T) typeof(T *)
#define array(T, N) typeof(T [N])
```

Now the declaration can be rewritten this way:

```
array (pointer (char), 4) y;
```

Thus, `array (pointer (char), 4)` is the type of arrays of four pointers to `char`.

12.3   Labels as Values

You can get the address of a label defined in the current function (or a containing function) with the unary operator '`&&`'. This is a non-standard extension to the language. Using this feature reduces your code portability.

The value returned has type `void *`. This value is a constant and can be used wherever a constant of that type is valid. For example:

```
void *ptr;
...
ptr = &&foo;
```

To use these values, you need to be able to jump to one. This is done with the computed goto statement, `goto *exp;`. For example:

```
goto *ptr;
```

Any expression of type `void *` is allowed.

One way of using these constants is in initializing a static array that will serve as a jump table:

```
static void *array[] = { &&foo, &&bar, &&hack };
```

Then you can select a label with indexing, like this:

```
goto *array[i];
```

Note: This does not check whether the subscript is in bounds. (Array indexing in C never does.)

Such an array of label values serves a purpose much like that of the `switch` statement. The `switch` statement is cleaner and therefore preferable to an array.

Another use of label values is in an interpreter for threaded code. The labels within the interpreter function can be stored in the threaded code for fast dispatching.

This mechanism can be misused to jump to code in a different function. The compiler cannot prevent this from happening, so care must be taken to ensure that target addresses are valid for the current function.

12.4   Conditional Operator Operands

The middle operand in a conditional expression may be omitted. Then if the first operand is nonzero, its value is the value of the conditional expression. This is a non-standard extension to the language. Using this feature reduces your code portability.

Therefore, the expression:

```
x ? : y
```

has the value of `x` if that is nonzero; otherwise, the value of `y`.

This example is perfectly equivalent to:

```
x ? x : y
```

In this simple case, the ability to omit the middle operand is not especially useful. When it becomes useful is when the first operand does, or may (if it is a macro argument), contain a side effect. Then repeating the operand in the middle would perform the side effect twice. Omitting the middle operand uses the value already computed without the undesirable effects of recomputing it.

12.5   Case Ranges

You can specify a range of consecutive values in a single case label, like this:

```
case low ... high:
```

This has the same effect as the proper number of individual case labels, one for each integer value from `low` to `high`, inclusive. This is a non-standard extension to the language. Using this feature reduces your code portability.

This feature is especially useful for ranges of ASCII character codes:

```
case 'A' ... 'Z':
```

Be careful: Write spaces around the..., otherwise it may be parsed incorrectly when you use it with integer values. For example, write this:

```
case 1 ... 5:
```

rather than this:

```
case 1...5:
```

13. Register Usage

This chapter examines registers used by the compiler to generate assembly from C/C++ source code.

13.1 Register Usage

The assembly generated from C/C++ source code by the compiler will use certain registers that are present on the PIC MCU device. The compiler assumes that nothing other than code it generates can alter the contents of these registers, but an extended assembly language format can be used to indicate to the compiler registers used in assembly code so that code can be adjusted accordingly.

13.2 Register Conventions

The 32 general purpose registers contained in the PIC32 are shown in the table below. Some of these registers are assigned a dedicated task by the compiler. The name used in assembly code and the usage is indicated.

Table 13-1. Register Conventions

Register Number	Software Name	Use
$0	zero	Always 0 when read.
$1	at	Assembler temporary variable. Do not use the $at register from source code unless you fully understand the implications.
$2-$3	v0-v1	Return value from functions.
$4-$7	a0-a3	Used for passing arguments to functions.
$8-$15	t0-t7	Temporary registers used by compiler for expression evaluation. Values not saved across function calls.
$16-$23	s0-s7	Temporary registers whose values are saved across function calls.
$24-$25	t8-t9	Temporary registers used by compiler for expression evaluation. Values not saved across function calls.
$26-$27	k0-k1	Reserved for interrupt/trap handler.
$28	gp	Global Pointer.
$29	sp	Stack Pointer.
$30	fp or s8	Frame Pointer if needed. Additional temporary saved register if not.
$31	ra	Return address for functions.

The PIC32MZ family uses the microMIPS compressed instruction-set architecture. You can use the `micromips` function attribute to compile the function for the microMIPS compressed mode. This compressed ISA generally results in a ~30% reduction in overall application code size at the expense of ~2% in performance. The microcontroller can switch between the MIPS32 and microMIPS modes on a function call. Consult your device data sheet to determine if your target device supports the microMIPS ISA.

Example function:

```
#include <xc.h>
void
__attribute__((micromips))
peanut (void)
{
  // function code here
}
```

Important: Standard function calls can switch between MIPS32 and microMIPS modes. However, when calling a MIPS32 library function from a microMIPS function, the compiler may generate a compressed `jals` instruction to call the library function. A `jals` instruction cannot change modes to MIPS32 and upon linking, you may receive an error, "Unsupported jump between ISA modes; consider recompiling with interlinking enabled." In that case, add the `-mno-jals` option to the Alternative Options field in your project properties for xc32-gcc, so it is passed to the compiler.

14.   Functions

The following sections describe how function definitions are written, and specifically how they can be customized to suit your application. The conventions used for parameters and return values, as well as the assembly call sequences are also discussed.

14.1   Writing Functions

Functions may be written in the usual way in accordance with the C/C++ language.

The only specifier that has any effect on function is `static`. Interrupt functions are defined with the use of the `interrupt` attribute, see 14.2 Function Attributes and Specifiers.

A function defined using the `static` specifier only affects the scope of the function, that is, limits the places in the source code where the function may be called. Functions that are `static` may only be directly called from code in the file in which the function is defined. The equivalent symbol used in assembly code to represent the function may change if the function is `static`, see 10.2.2 Static Variables. This specifier does not change the way the function is encoded.

14.2   Function Attributes and Specifiers

14.2.1   Function Attributes

address(addr)

The `address` attribute specifies an absolute virtual address for the function. Be sure to specify the `address` attribute using an appropriate virtual address for the target device. The address is typically in the range [0x9D000000,0x9D0FFFFC], as defined in the linker script as the 'kseg0_program_mem' memory region. For example,

`__attribute__((address(0x9D008000))) void bar (void);`

The compiler performs no error checking on the address. The section containing the function will be located at the specified address regardless of the memory-regions specified in the linker script or the actual memory ranges on the target device. The application code must ensure that the address is valid for the target device.

To make effective use of absolute sections and the new best-fit allocator, standard program-memory and data-memory sections should not be mapped in the linker script. The built-in linker script does not map most standard sections such as the `.text`, `.data`, `.bss`, or `.ramfunc` sections. By not mapping these sections in the linker script, we allow these sections to be allocated using the best-fit allocator rather than the sequential allocator. Sections that are unmapped in the linker script can flow around absolute sections, whereas sections that are linker-script mapped are grouped together and allocated sequentially, potentially causing conflicts with absolute sections.

alias ("symbol")

Indicates that the function is an alias for another symbol. For example:

```
void foo (void) { /* stuff */ }
__attribute__ ((alias("foo"))) void bar (void);
```

Symbol bar is considered to be an alias for the symbol `foo`.

always_inline

If the function is declared `inline`, always inline the function, even if no optimization level was specified.

at_vector

Place the body of the function at the indicated exception vector address.

See 15. Interrupts and 15.4 Exception Handlers.

const

If a pure function determines its return value exclusively from its parameters (that is, does not examine any global variables), it may be declared `const`, allowing for even more aggressive optimization. Note that a function which de-references a pointer argument is not `const` since the pointer de-reference uses a value which is not a parameter, even though the pointer itself is a parameter.

`deprecated`

`deprecated (msg)`

When a function specified as `deprecated` is used, a warning is generated. The optional *msg* argument, which must be a string, will be printed in the warning if present. The `deprecated` attribute may also be used for variables and types.

`far`

Always invoke the function by first loading its address into a register and then using the contents of that register. This allows calling a function located beyond the 28-bit addressing range of the direct `CALL` instruction.

`format (type, format_index, first_to_check)`

The `format` attribute indicates that the function takes a `printf`, `scanf`, `strftime`, or `strfmon` style format string and arguments and that the compiler should type check those arguments against the format string, just as it does for the standard library functions.

The `type` parameter is one of `printf`, `scanf`, `strftime` or `strfmon` (optionally with surrounding double underscores, for example, `__printf__`) and determines how the format string will be interpreted.

The `format_index` parameter specifies which function parameter is the format string. Function parameters are numbered from the left-most parameter, starting from 1.

The `first_to_check` parameter specifies which parameter is the first to check against the format string. If `first_to_check` is zero, type checking is not performed and the compiler only checks the format string for consistency (for example, `vfprintf`).

`format_arg (index)`

The `format_arg` attribute specifies that a function manipulates a `printf` style format string and that the compiler should check the format string for consistency. The function attribute which is a format string is identified by `index`.

`function_replacement_prologue`

This feature allows the application to redirect one function to another implementation at runtime without replacing the existing function. This is achieved by changing the method of invoking functions through a function replacement table instead of from a linker-resolved address. Initially the function address in the table points to the location of the entry point of the original function's prologue. The application can then replace the table entry with the new function address. Now, during the program execution, the control will pass to the new function address and returns to the caller function. This feature adds a new `function_replacement_prologue` function attribute. To redirect the function, modify the corresponding Function Replacement Table entry at runtime.

Example C code:

```
int a, b, c, d;
int __attribute__((function_replacement_prologue)) foo (void)
{
   a = b + c;
   return (a);
}
int main()
{
   d = foo();
   return 0;
}
```

Example generated assembly code:

```
# Function Replacement Table entries, located in data memory
   .section .fixtable, data
fixtable.foo:
```

```
    .word       cont.foo            # By default, populate the table with the address
                                    # of the original implementation. Redirect to
                                    # another implementation by overwriting this
                                    # location with the address of the new implementation.

    .section .text, code
    .globl foo
    .ent foo
    .type foo, @function
foo:
    # Begin Function Replacement Table Prologue
    lui $25,%hi(fixtable.foo)       # Load address from .fixtable above
    lw $25,%lo(fixtable.foo)($25)
    j $25                           # Jump to address loaded from table
    nop
cont.foo:
    # End Function Replacement Table Prologue
    addiu $sp,$sp,-8
    sw $fp,4($sp)
    move $fp,$sp
    lw $3,%gp_rel(b)($28)
    ...........
    j $31
    nop
```

interrupt (priority)

Generate prologue and epilogue code for the function as an interrupt handler function. See 15. Interrupts . The argument specified the interrupt priority level using the symbols IPL*n*SOFT, IPL*n*SRS, or IPL*n*AUTO where *n* represents the 7 levels of priority and SOFT|SRS|AUTO specifies the context saving mode.

keep

The `__attribute__((keep))` may be applied to a function. The keep attribute will prevent the linker from removing the function with `--gc-sections`, even if it is unused.

```
void __attribute__((keep)) foo(void);
```

longcall

Functionally equivalent to `far`.

malloc

Any non-Null Pointer return value from the indicated function will not alias any other pointer which is live at the point when the function returns. This allows the compiler to improve optimization.

micromips

Generate code for the function in the compressed microMIPS instruction set.

mips16

Generate code for the function in the MIPS16 instruction set.

naked

Generate no prologue or epilogue code for the function.

near

Always invoke the function with an absolute `CALL` instruction, even when the `-mlong-calls` command line option is specified.

no_fpu

The `no_fpu` attribute specifies that the interrupt function should not preserve the Floating-Point Unit (FPU) context. The attribute should always be used in conjunction with the `interrupt (priority)` attribute. In addition to suppressing the FPU context saving code, this attribute causes the compiler to disable the FPU by clearing the CU1 bit of the CP0 Status register. If your interrupt service routine uses a floating-point operation while the FPU is disabled, the device will take a general exception. The interrupt service routine restores the original value of the status register before returning from the interrupt.

noinline

The function will never be considered for inlining.

noload

Causes the variable or function to be placed in a section that has the `noload` attribute set. It tells consumers of the ELF files not to load the contents of the section. This attribute can be useful when you just want to reserve memory for something, but you don't want to clear or initialize memory.

void bar() _ *attribute_* ((noload))

nomips16

Always generate code for the function in the MIPS32[®] instruction set, even when compiling the translation unit with the `-mips16` command line option.

nonnull (index, ...)

Indicate to the compiler that one or more pointer arguments to the function must be non-null. If the compiler determines that a Null Pointer is passed as a value to a non-null argument, and the `-Wnonnull` command line option was specified, a warning diagnostic is issued.

If no arguments are given to the `nonnull` attribute, all pointer arguments of the function are marked as non-null.

noreturn

Indicate to the compiler that the function will never return. In some situations, this can allow the compiler to generate more efficient code in the calling function since optimizations can be performed without regard to behavior if the function ever did return. Functions declared as `noreturn` should always have a return type of void.

optimize

You can now use the `optimize` attribute to specify different optimization options for various functions within a source file. Arguments can either be numbers or strings. Numbers are assumed to be an optimization level. Strings that begin with O are assumed to be an optimization option. This feature can be used for instance to have frequently executed functions compiled with more aggressive optimization options that produce faster and larger code, while other functions can be called with less aggressive options.

```
int __attribute__((optimize("-O3"))) pandora (void)
{
  if (maya > axton) return 1;
  return 0;
}
```

pure

If a function has no side effects other than its return value, and the return value is dependent only on parameters and/or (nonvolatile) global variables, the compiler can perform more aggressive optimizations around invocations of that function. Such functions can be indicated with the `pure` attribute.

ramfunc

Treat the function as if it was in data memory. Allocate the function at the highest appropriately aligned address for executable code. Note that due to `ramfunc` alignment and placement requirements, the `address` attribute should not be used with the `ramfunc` attribute. The presence of the `ramfunc` section causes the linker to emit the symbols necessary for the crt0.S start-up code to initialize the bus matrix appropriately for executing code out of data memory.

Use this attribute along with the `far/longcall` attribute and the `section` attribute. For example:

```
  __attribute__((ramfunc,section(".ramfunc"),far,unique_section))
  unsigned int myramfunct (void_
    { /* code */ }
```

A macro in the sys/attribs.h header file makes the `ramfunc` attribute simple to use:

```
#include <sys/attribs.h>
__longramfunc__ unsigned int myramfunct (void)
{ /* code */ }
```

`section("name")`

Place the function into the named section.

For example:

```
void __attribute__ ((section (".wilma"))) baz () {return;}
```

Function `baz` will be placed in section `.wilma`.

The `-ffunction-sections` command line option has no effect on functions defined with a `section` attribute.

`unique_section`

Place the function in a uniquely named section, as if `-ffunction-sections` had been specified. If the function also has a `section` attribute, use that section name as the prefix for generating the unique section name.

For example:

```
void __attribute__ ((section (".fred"), unique_section) foo (void) {return;}
```

Function `foo` will be placed in section `.fred.foo`.

`unused`

Indicate to the compiler that the function may not be used. The compiler will not issue a warning for this function if it is not used.

`used`

Indicate to the compiler that the function is always used and code must be generated for the function even if the compiler cannot see a reference to the function. For example, if inline assembly is the only reference to a static function.

`vector (num)`

Generate a branch instruction at the indicated exception vector which targets the function. See 15. Interrupts and 15.4 Exception Handlers.

`warn_unused_result`

A warning will be issued if the return value of the indicated function is unused by a caller.

`weak`

A weak symbol indicates that if another version of the same symbol is available, that version should be used instead. For example, this is useful when a library function is implemented such that it can be overridden by a user written function.

14.3 Allocation of Function Code

Code associated with C/C++ functions is normally always placed in the program Flash memory of the target device.

Functions may be located in and executed from RAM rather than Flash by using the `__ramfunc__` and `__longramfunc__` macros.

Functions specified as a RAM function will be copied to RAM by the start-up code and all calls to those functions will reference the RAM location. Functions located in RAM will be in a different 512MB memory segment than functions located in program memory, so the `longcall` attribute should be applied to any RAM function, which will be called from a function not in RAM. The `__longramfunc__` macro will apply the `longcall` attribute as well as place the function in RAM.

```
#include <sys/attribs.h>
/* function 'foo' will be placed in RAM */
```

```
void __ramfunc__ foo (void)
{
}

/* function 'bar' will be placed in RAM and will be invoked
   using the full 32 bit address */
void __longramfunc__ bar (void)
{
}
```

Note: Specifying `__longramfunc__` is functionally equivalent to specifying both `__ramfunc__` and `__longcall__`.

14.4 Changing the Default Function Allocation

The assembly code associated with a C/C++ function can be placed at an absolute address. This can be accomplished by using the `address` attribute and specifying the virtual address of the function, see 9.11 Variable Attributes.

Functions can also be placed at specific positions by placing them in a user-defined section and then linking this section at an appropriate address, see 9.11 Variable Attributes.

14.5 Function Size Limits

There are no theoretical limits as to how large functions can be made.

14.6 Function Parameters

MPLAB XC uses a fixed convention to pass arguments to a function. The method used to pass the arguments depends on the size and number of -arguments involved.

Note: The names "argument" and "parameter" are often used interchangeably, but typically an argument is the actual value that is passed to the function and a parameter is the variable defined by the function to store the argument.

The Stack Pointer is always aligned on an 8-byte boundary.

- All integer types smaller than a 32-bit integer are first converted to a 32-bit value. The first four 32 bits of arguments are passed via registers `a0-a3` (see the table below for how many registers are required for each data type).
- Although some arguments may be passed in registers, space is still allocated on the stack for all arguments to be passed to a function (see the figure below). Application code should not assume that the current argument value is on the stack, even when space is allocated.
- When calling a function:
    - Registers `a0-a3` are used for passing arguments to functions. Values in these registers are not preserved across function calls.
    - Registers `t0-t7` and `t8-t9` are caller saved registers. The calling function must push these values onto the stack for the registers' values to be saved.
    - Registers `s0-s7` are called saved registers. The function being called must save any of these registers it modifies.
    - Register `s8` is a saved register if the optimizer eliminates its use as the Frame Pointer. `s8` is a reserved register otherwise.
    - Register `ra` contains the return address of a function call.

Table 14-1. Registers Required

Data Type	Number of Registers Required
char	1

..........continued

Data Type	Number of Registers Required
short	1
int	1
long	1
long long	2
float	1
double	1
long double	2
structure	Up to 4, depending on the size of the struct.

Figure 14-1. Passing Arguments

Example 1:

```
int add (int, int)
a= add (5, 10);
```

SP + 4	undefined		a0	5
SP	undefined		a1	10

Example 2:

```
void foo (long double, long double)
call= foo (10.5, 20.1);
```

SP + 12	undefined		a0	10.5
SP + 8			a1	
SP + 4	undefined		a2	20.1
SP			a3	

Example 3:

```
void calculate (long double, long double, int)
calculate (50.3, 100.0, .10);
```

	.10			
SP + 16				
SP + 12	undefined		a0	50.3
SP + 8			a1	
SP + 4	undefined		a2	100.0
SP			a3	

14.7 Function Return Values

Function return values are returned in registers.

Integral or pointer value are placed in register `v0`. All floating-point values, regardless of precision, are returned in floating-point register `$f0`.

If a function needs to return an actual structure or union – not a pointer to such an object – the called function copies this object to an area of memory that is reserved by the caller. The caller passes the address of this memory area in

register $4 when the function is called. The function also returns a pointer to the returned object in register `v0`. Having the caller supply the return object's space allows re-entrance.

14.8   Calling Functions

By default, functions are called using the direct form of the call (`jal`) instruction. This allows calls to destinations within a 256 MB segment. This operation can be changed through the use of attributes applied to functions or command-line options so that a longer, but unrestricted, call is made.

The `-mlong-calls` option, see 6.7.1 Options Specific to PIC32M Devices, forces a register form of the call to be employed by default. Generated code is longer, but calls are not limited in terms of the destination address.

The attributes `longcall` or `far` can be used with a function definition to always enforce the longer call sequence for that function. The `near` attribute can be used with a function so that calls to it use the shorter direct call, even if the `-mlong-calls` option is in force.

14.9   Inline Functions

Enter a short description of your concept here (optional).

By declaring a function `inline`, you can direct the compiler to integrate that function's code into the code for its callers. This usually makes execution faster by eliminating the function-call overhead. In addition, if any of the actual argument values are constant, their known values may permit simplifications at compile time, so that not all of the inline function's code needs to be included. The effect on code size is less predictable. Machine code may be larger or smaller with inline functions, depending on the particular case.

Note: Function inlining will only take place when the function's definition is visible (not just the prototype). In order to have a function inlined into more than one source file, the function definition may be placed into a header file that is included by each of the source files.

To declare a function inline, use the `inline` keyword in its declaration, like this:

```
inline int
inc (int *a)
{
   (*a)++;
}
```

(If you are using the `-traditional` option or the `-ansi` option, write `__inline__` instead of `inline`.) You can also make all "simple enough" functions inline with the command-line option `-finline-functions`. The compiler heuristically decides which functions are simple enough to be worth integrating in this way, based on an -estimate of the function's size.

Note: The inline keyword will only be recognized with `-finline` or optimizations enabled.

Certain usages in a function definition can make it unsuitable for inline substitution. Among these usages are: use of `varargs`, use of `alloca`, use of variable-sized data, use of computed `goto` and use of nonlocal `goto`. Using the command-line option `-Winline` will warn when a function marked `inline` could not be substituted, and will give the reason for the failure.

In compiler syntax, the `inline` keyword does not affect the linkage of the function.

When a function is both `inline` and `static`, if all calls to the function are integrated into the caller and the function's address is never used, then the function's own assembler code is never referenced. In this case, the compiler does not actually output assembler code for the function, unless you specify the command-line option `-fkeep-inline-functions`. Some calls cannot be integrated for various reasons (in particular, calls that precede the function's definition cannot be integrated and neither can recursive calls within the definition). If there is a non-integrated call, then the function is compiled to assembler code as usual. The function must also be compiled as usual if the program refers to its address, because that can't be inlined. The compiler will only eliminate `inline` functions if they are declared to be `static` and if the function definition precedes all uses of the function.

When an `inline` function is not `static`, then the compiler must assume that there may be calls from other source files. Since a global symbol can be defined only once in any program, the function must not be defined in the other source files, so the calls therein cannot be integrated. Therefore, a non-`static` inline function is always compiled on its own in the usual fashion.

If you specify both `inline` and `extern` in the function definition, then the definition is used only for inlining. In no case is the function compiled on its own, not even if you refer to its address explicitly. Such an address becomes an external reference, as if you had only declared the function and had not defined it.

This combination of `inline` and `extern` has a similar effect to a macro. Put a function definition in a header file with these keywords and put another copy of the definition (lacking `inline` and `extern`) in a library file. The definition in the header file will cause most calls to the function to be inlined. If any uses of the function remain, they will refer to the single copy in the library.

15. Interrupts

Interrupt processing is an important aspect of most microcontroller applications. Interrupts may be used to synchronize software operations with events that occur in real time. When interrupts occur, the normal flow of software execution is suspended, and special functions are invoked to process the event. At the completion of interrupt processing, previous context information is restored and normal execution resumes.

PIC32 devices support multiple interrupts, from both internal and external sources. The devices allow high-priority interrupts to override any lower priority interrupts that may be in progress.

The compiler provides full support for interrupt processing in C/C++ or inline assembly code. This section presents an overview of interrupt processing.

15.1 Interrupt Operation

The compiler incorporates features allowing interrupts to be fully handled from C/C++ code. Interrupt functions are often called interrupt handlers or Interrupt Service Routines (ISRs).

Each interrupt source typically has a control bit in an SFR which can disable that interrupt source. Check your device data sheet for full information how your device handles interrupts.

Interrupt code is the name given to any code that executes as a result of an interrupt occurring. Interrupt code completes at the point where the corresponding return from interrupt instruction is executed. This contrasts with *main-line code*, which, for a freestanding application, is usually the main part of the program that executes after Reset.

15.2 Writing an Interrupt Service Routine

An interrupt handler function is different than an ordinary function in that it handles the context save and restore to ensure that upon return from interrupt, the program context is maintained. A different code sequence is used to return from these functions as well.

Several attributes can be used to ensure that the compiler generates the correct code for an ISR. Macros are provided so that this is easier to accomplish, see the following sections.

There are several actions that the compiler needs to take to generate an interrupt service routine. The compiler has to be told to use an alternate form of return code. The function also needs to be linked to the interrupt vector. All ISRs must use either the MIPS32®r2 or the microMIPS™ ISA modes. Apply the 'nomips16' function attribute to each interrupt function.

Note: For devices that support multiple Instruction Set Architecture (ISA) modes, there may be a configuration bit that determines which mode the device uses for an exception/interrupt. If your device is configured to use the microMIPS ISA on interrupt, be sure to apply the `micromips` function attribute to your interrupt function. Consult your the data sheet for your target device to determine if it supports handling exceptions and interrupts in an alternate ISA mode.

An interrupt function must be declared as type `void` and may not have parameters. This is the only function prototype that makes sense for an interrupt function since they are never directly called in the source code.

Interrupt functions must not be called directly from C/C++ code (due to the different return instruction that is used), but they themselves may call other functions, both user-defined and library functions, but be aware that this may use additional registers which will need to be saved and restored by the context switch code.

A function is marked as an interrupt handler function, or ISR, via either the `interrupt` attribute or the interrupt pragma*. While each method is functionally equivalent to the other, the `interrupt` attribute is more commonly used and therefore the recommended method. The interrupt is specified as handling interrupts of a specific priority level or for operating in single vector mode.

For all interrupt vectors without specific handlers, a default interrupt handler will be installed. The default interrupt handler is supplied by the libpic32.a library and will cause a debug breakpoint and reset the device. An application may override the default handler and provide an application-specific default interrupt handler by declaring an interrupt function with the name `_DefaultInterrupt`.

Note: * Pre-processor macros are not expanded in pragma directives.

15.2.1   Interrupt Attribute

```
__attribute__((interrupt([IPLn[SRS|SOFT|AUTO]])))
```

Where *n* is in the range of 0..7, inclusive.

Use the `interrupt` attribute to indicate that the specified function is an interrupt handler. The compiler generates function entry and exit sequences suitable for use in an interrupt handler when this attribute is present. The generated code preserves context by either using a shadow register set (SRS) or using generated software instructions (SOFT) to push context onto the stack. See the example below for an `interrupt` attribute.

Note: Some PIC32 target devices allow the exception/interrupt code to be in either the MIPS32® or microMIPS™ ISA mode via a device configuration bit (BOOTISA). On these devices, if your BOOTISA bit is set to microMIPS mode, add the 'micromips' attribute to your interrupt function. If your BOOTISA bit is set to MIPS32 mode, add the 'nomicromips' attribute to your interrupt function. See your device data sheet for more information on this configuration bit.

> **Example 15-1. Interrupt Attribute**
> ```
> void __attribute__((interrupt(IPL7SRS))) bambam (void);
> ```

Many PIC32 devices allow us to specify, via configuration-bit settings, which interrupt priority level will use the shadow register set (for example, `#pragma config` FSRSSEL=PRIORITY_7). Refer to the device data sheet to determine if your PIC32 target device supports this feature. This means we must specify which context-saving mechanism to use for each interrupt handler. The compiler will generate interrupt function prologue and epilogue code utilizing shadow register context saving for the `IPLnSRS` Interrupt Priority Level (IPL) specifier. It will use software context saving for the `IPLnSOFT` IPL specifier.

Other PIC32 variants may have 8 register sets (1 standard set and 7 shadow register sets) meaning that there are enough shadow register sets for every interrupt priority level. Therefore, you should use the IPLnSRS IPL specifier for every interrupt service routine on these device variants.

Note: Application code is responsible for applying the correct IPL specifier value to each ISR. The interrupt source's priority level must match the ISR's IPL value (for example, IPLnSRS) or the interrupt will not be handled correctly. Mismatching priority levels may result in critical runtime problems such as a stack overflow that overwrites data memory. This can include corruption of memory reserved for use by the Debug Executive, causing the debug tool to behave erratically.

The compiler also supports an `IPLnAUTO` IPL specifier that uses the run-time value in `SRSCTL` to determine whether it should use software or SRS context-saving code. The compiler defaults to using `IPLnAUTO` when the IPL specifier is omitted from the `interrupt()` attribute.

For devices that do not support a shadow register set for interrupt context saving, use IPLnSOFT for all interrupt handlers.

Note: `SRS` has the shortest latency and `SOFT` has a longer latency due to registers saved on the stack. `AUTO` adds a few cycles to test if `SRS` or `SOFT` should be used.

For `IPL7(SRS | SOFT | AUTO)`, the compiler assumes that nothing can interrupt priority 7. This means that there is no reason to save `EPC` or `SRSCtl` and that global disabling of interrupts is unnecessary.

The `IPLnSAVEALL` interrupt priority specifier can be used with the interrupt attribute. Use this specifier in place of `IPLnSOFT` to force software context saving of all software-saved general registers even if they are not used within the Interrupt Service Routine (ISR). This attribute can be useful for some RTOS implementations.

The `keep_interrupts_masked` attribute can be used to modify the behavior of an interrupt handler. The attribute keeps interrupts masked for the whole function. Without this attribute, the XC32 compiler re-enables interrupts for as much of the function as it can. By keeping interrupts masked, support for nested interrupts is disables. Users can re-enable them as necessary in their own code.

The attribute `keep_interrupts_masked` can be combined with the `interrupt` attribute. This attribute causes the Interrupt Service Routine (ISR) prologue code to not re-enable interrupts. Application code may then choose whether and when to re-enable interrupts in the ISR.

15.2.2   Interrupt Pragma

Note: The interrupt pragma is provided only for compatibility when porting code from other compilers. The `interrupt` function attribute is the preferred and more common way to write an interrupt service routine.

```
# pragma interrupt function-name IPLn[AUTO|SOFT|SRS] [vector [@]vector-number [, vector-
number-list]]
# pragma interrupt function-name  single [vector [@] 0
```

Where *n* is in the range of 0..7, inclusive.

The `IPLn[AUTO|SOFT|SRS]` IPL specifier may be all uppercase or all lowercase.

The function definition for a handler function indicated by an interrupt pragma must follow in the same translation unit as the pragma itself.

The `interrupt` attribute will also indicate that a function definition is an interrupt handler. It is functionally equivalent to the interrupt pragma.

For example, the definitions of `foo` below both indicate that it is an interrupt handler function for an interrupt of priority 4 that uses software context saving.

```
#pragma interrupt foo IPL4SOFT
void foo (void)
```

is functionally equivalent to

```
void __attribute__  ((interrupt(IPL4SOFT))) foo (void)
```

15.2.3   __ISR Macros

The `<sys/attribs.h>` header file provides macros intended to simplify the application of attributes to interrupt functions. There are also vector macros defined in the processor header files. (See the appropriate header file in the compiler's `/pic32mx/include/proc` directory.)

Note: Some PIC32 target devices allow the exception/interrupt code to be in either the MIPS32® or microMIPS™ ISA mode via a device configuration bit (BOOTISA). On these devices, if your BOOTISA bit is set to microMIPS mode, add the 'micromips' attribute to your interrupt function. If your BOOTISA bit is set to MIPS32 mode, add the 'nomicromips' attribute to your interrupt function. See your device data sheet for more information on this configuration bit.

__ISR(V, IPL)

Use the `__ISR(v, IPL)` macro to assign the vector-number location and associate it with the specified IPL. This will place a jump to the interrupt handler at the associated vector location. This macro also applies the `nomips16` attribute since PIC32 devices require that interrupt handlers must use the MIPS32 instruction set.

The following example creates an interrupt handler function for the core timer interrupt that has an interrupt priority level of two. The compiler places a dispatch function at the associated vector location. To reach this function, the core timer interrupt flag and enable bits must be set, and the interrupt priority should be set to a level of two. The compiler generates software context-saving code for this handler function.

> **Example 15-2. Core Timer Vector, IPL2SOFT**
>
> ```
> #include <xc.h>
> #include <sys/attribs.h>
> void __ISR(_CORE_TIMER_VECTOR, IPL2SOFT) CoreTimerHandler(void);
> ```

The example below creates an interrupt handler function for the core software interrupt 0 that has an interrupt priority level of three. The compiler places a dispatch function at the associated vector location. To reach this function, the core software interrupt flag and enable bits must be set, and the interrupt priority should be set to a level of three. The device configuration fuses must assign Shadow Register Set 1 to interrupt priority level three. The compiler generates code that assumes that register context will be saved in SRS1.

Example 15-3. Core Software 0 Vector, IPL3SRS

```
#include <xc.h>
#include <sys/attribs.h>
void __ISR(_CORE_SOFTWARE_0_VECTOR,IPL3SRS) CoreSoftwareInt0Handler(void);
```

The example below creates an interrupt handler function for the core software interrupt 1 that has an interrupt priority level of zero. The compiler places a dispatch function at the associated vector location. To reach this function, the core software interrupt 1 flag and enable bits must be set, and the interrupt priority should be set to a level of zero. The compiler generates code that determines at run time whether software context saving is required.

Example 15-4. Core Software 1 Vector, IPL0AUTO

```
#include <xc.h>
#include <sys/attribs.h>
void __ISR(_CORE_SOFTWARE_1_VECTOR, IPL0AUTO) CoreSoftwareInt1Handler(void);
```

The next example is functionally equivalent to Example 14-4. Because the IPL specifier is omitted, the compiler assumes IPL0AUTO.

Example 15-5. Core Software 1 Vector, Default

```
#include <xc.h>
#include <sys/attribs.h>
void __ISR(_CORE_SOFTWARE_1_VECTOR) _CoreSoftwareInt1Handler(void);
```

__ISR_AT_VECTOR(v, IPL)

Use the __ISR_AT_VECTOR(v, IPL) to place the entire interrupt handler at the vector location and associate it with the software-assigned interrupt priority. Application code is responsible for making sure that the vector spacing is set to accommodate the size of the handler. This macro also applies the nomips16 attribute since ISR functions are required to be MIPS32.

The following example creates an interrupt handler function for the core timer interrupt that has an interrupt priority level of two. The compiler places the entire interrupt handler at the vector location. It does not use a dispatch function. To reach this function, the core timer interrupt flag and enable bits must be set, and the interrupt priority should be set to a level of two. The compiler generates software context-saving code for this handler function.

Example 15-6. Core Timer Vector, IPL2SOFT

```
#include <xc.h>
#include <sys/attribs.h>
void __ISR_AT_VECTOR(_CORE_TIMER_VECTOR, IPL2SOFT) CoreTimerHandler(void);
```

INTERRUPT-VECTOR MACROS

Each processor-support header file provides a macro for each interrupt-vector number (for example, /pic32mx/include/proc/p32mx360f512l.h. See the appropriate header file in the compiler install directory). When used in conjunction with the __ISR() macro provided by the sys\attribs.h header file, these macros help make an Interrupt Service Routine easier to write and maintain.

The example below creates an interrupt handler function for the Timer 1 interrupt that has an interrupt priority level of seven. The compiler places a dispatch function at the vector location associated with macro _TIMER_1_VECTOR as defined in the device-specific header file. To reach this function, the Timer 1 interrupt flag and enable bits must be set, and the interrupt priority should be set to a level of seven. For devices that allow assignment of shadow registers to specific IPL values, the device Configuration bit settings must assign Shadow Register Set 1 to interrupt priority level seven. The compiler generates code that assumes that register context will be saved in SRS1.

Example 15-7. Interrupt-Vector with Handler

```
#include <xc.h>
#include <sys/attribs.h>
void __ISR (_TIMER_1_VECTOR, IPL7SRS) Timer1Handler (void);
```

15.3   Associating a Handler Function with an Exception Vector

For PIC32 devices, each interrupt source is mapped to an exception vector, as specified in the device data sheet. For devices with a constant vector spacing, a default of four words of space are reserved at each vector address for a dispatch to the handler function for that exception source. For devices with a variable vector spacing, the default linker script adjusts each vector's spacing for the size of the designated interrupt function.

An interrupt handler function can be associated with an interrupt vector either as the target of a dispatch function located at the exception vector address, or as being located directly at the exception vector address. A single handler function can be the target of multiple dispatch functions.

The association of a handler function to one or more exception vector addresses is specified via a vector attribute on the function declaration. For compatibility purposes, you may also associate a handler function to a vector address using a clause of the interrupt pragma, a separate vector pragma, or a vector attribute on the function declaration.

15.3.1   Vector Attribute

A handler function can be associated with one or more exception vector addresses via an attribute. The `at_vector` attribute indicates that the handler function should itself be placed at the exception vector address. The `vector` attribute indicates that a dispatch function should be created at the exception vector address(es) which will transfer control to the handler function.

For example, the following declaration specifies that function `foo` will be created as an interrupt handler function of priority four. `foo` will be located at the address of exception vector 54.

```
void __attribute__ ((interrupt(IPL4SOFT))) __attribute__ ((at_vector(54))) foo void);
```

The following declaration specifies that function `foo` will be created as an interrupt handler function of priority four. Define dispatch functions targeting `foo` at exception vector addresses 52 and 53.

```
void __attribute__ ((interrupt(IPL4SOFT))) __attribute__ ((vector(53, 52))) foo void)
```

Handler functions that are linked directly to the vector will be executed faster. Although the vector spacing can be adjusted, there is limited space between vectors and linking a substantial handler function directly at a vector may cause it to overlap the higher vector locations, preventing their use. In such situations, using a dispatch function is a safer option.

The newer devices family features variable offsets for vector spacing. The compiler and linker work together to treat the OFF*nnn* SFRs as initialized data so that they are initialized at startup. This means there is no need for application code to initialize the OFF*nnn* SFRs. This also means that it is often more efficient to place the ISR within the vector table rather than using a dispatch function.

Example 15-8. Example Interrupt Service Routine

```
#include <xc.h>
#include <sys/attribs.h>
void
__ISR_AT_VECTOR(_CORE_TIMER_VECTOR, IPL7SRS)
CoreTimerHandler(void)
{
    // ISR code here
}
```

15.3.2 Interrupt-Pragma Vector Clause

Note: The interrupt pragma and its vector clause are provided only for compatibility when porting code from other compilers. The vector function attribute is the preferred way to associate a handler function to an exception vector address.

The interrupt pragma has an optional `vector` clause following the priority specifier.

```
# pragma interrupt function-nameIPL-specifier [vector [@]vector-number [, vector-number-list]]
```

A dispatch function targeting the specified handler function will be created at the exception vector address for the specified vector numbers. If the first vector number is specified with a preceding "@" symbol, the handler function itself will be located there directly.

For example, the following pragma specifies that function `foo` will be created as an interrupt handler function of priority four. `foo` will be located at the address of exception vector 54. A dispatch function targeting `foo` will be created at exception vector address 34.

```
#pragma interrupt foo IPL4AUTO vector @54, 34
```

The following pragma specifies that function `bar` will be created as an interrupt handler function of priority five. `bar` will be located in general purpose program memory (.text section). A dispatch function targeting `bar` will be created at exception vector address 23.

```
#pragma interrupt bar IPL5SOFT vector 23
```

15.3.3 Vector Pragma

Note: The vector pragma is provided only for compatibility when porting code from other compilers. The vector function attribute is the preferred way to associate a handler function to an exception vector address.

The `vector` pragma creates one or more dispatch functions targeting the indicated function. For target functions specified with the `interrupt` pragma, this functions as if the vector clause had been used. The target function of a `vector` pragma can be any function, including external functions implemented in assembly or by other means.

```
# pragma vector function-name vector vector-number [, vector-number-list]
```

The following pragma defines a dispatch function targeting `foo` at exception vector address 54.

```
#pragma vector foo 54
```

15.4 Exception Handlers

The PIC32 devices also have exception vectors for non-interrupt exceptions. These exceptions are grouped into bootstrap exceptions and general exceptions.

15.4.1 Bootstrap Exception

A Reset exception is any exception which occurs while bootstrap code is running ($Status_{BEV}$=1). All Reset exceptions are vectored to `0xBFC00380`.

At this location, the 32-bit toolchain places a branch instruction targeting a function named `_bootstrap_exception_handler()`. In the standard library, a default weak version of this function is provided, which merely causes a software Reset. When compiling for in-circuit debugging or emulation, the default implementation of `_bootstrap_exception_handler` will first cause a software breakpoint and then a software Reset. If the user application provides an implementation of `_bootstrap_exception_handler()`, that implementation will be used instead.

15.4.2 General Exception

A general exception is any non-interrupt exception which occurs during program execution outside of bootstrap code ($Status_{BEV}$=0). General exceptions are vectored to offset `0x180` from `EBase`.

At this location, the 32-bit toolchain places a branch instruction targeting a function named `_general_exception_context()`. The provided implementation of this function saves context, calls an

application handler function, restores context and performs a return from the exception instruction. The context saved is the `hi` and `lo` registers, and all General Purpose Registers except `s0-s8`, which are defined to be preserved by all called functions and so are not necessary to actively save here again.

```
void _general_exception_handler (void);
```

A weak default implementation of `_general_exception_handler()` is provided in the standard library which merely causes a software Reset. When compiling for in-circuit debugging or emulation, the default implementation of `_general_exception_handler` will first cause a software breakpoint and then a software Reset. If the user application provides an implementation of `_general_exception_handler()`, that implementation will be used instead.

15.4.3   Simple TLB Refill Exception

During an instruction fetch or data access, a TLB refill exception occurs when no TLB entry matches a reference to a mapped address space and the EXL bit is 0 in the Status register. Note that this is distinct from the case in which an entry matches, but has the valid bit off. In that case, a TLB Invalid exception occurs.

```
void _simple_tlb_refill_exception_handler(void);
```

A weak default implementation of `_simple_tlb_refill_exception_handler()` is provided which merely causes a software Reset.

When compiling for in-circuit debugging or emulation, the default implementation of `_simple_tlb_refill_exception_handler` will first cause a software breakpoint and then a software Reset.

15.4.4   Cache Error Exception

A cache-error exception occurs when an instruction or data reference detects a cache tag or data error. This exception is not maskable. To avoid disturbing the error in the cache array the exception vector is to an unmapped, uncached address. This exception is precise.

```
void _cache_err_exception_handler(void);
```

A weak default implementation of `_cache_err_exception_handler()` is provided which merely causes a software Reset. When compiling for in-circuit debugging or emulation, the default implementation of `_cache_err_exception_handler` will first cause a software breakpoint and then a software Reset.

15.5   Interrupt Service Routine Context Switching

The standard calling convention for C/C++ functions will already preserve `zero`, `s0-s7`, `gp`, `sp`, and `fp`. `k0` and `k1` are used by the compiler to access and preserve non-GPR context, but are always accessed atomically (that is, in sequences with global interrupts disabled), so they need not be preserved actively. A handler function will actively preserve the `a0-a3`, `t0-t9`, `v0`, `v1` and `ra` registers in addition to the standard registers.

An interrupt handler function will also actively save and restore processor status registers that are utilized by the handler function. Specifically, the `EPC`, `SR`, `hi` and `lo` registers are preserved as context. All available DSP accumulators are preserved as necessary.

In addition, if a DSP accumulator register is preserved, the DSP Control register is also preserved.

Handler functions may use a shadow register set to preserve the General Purpose Registers, enabling lower latency entry into the application code of the handler function. On some devices, the shadow register set is assigned to an interrupt priority level (IPL) using the device Configuration bit settings (for example, `#pragma config FSRSSEL=PRIORITY_6`). While on other devices, the shadow register set may be hard wired to IPL7. Consult the target device's data sheet for more information on the shadow register set.

15.5.1   Context Restoration

Any objects saved by software are automatically restored by software before the interrupt function returns. The order of restoration is the reverse to that used when context is saved.

15.6 Latency

There are two elements that affect the number of cycles between the time the interrupt source occurs and the execution of the first instruction of your ISR code. These are:

- **Processor Servicing of Interrupt** – The amount of time it takes the processor to recognize the interrupt and branch to the first address of the interrupt vector. To determine this value, refer to the processor data sheet for the specific processor and interrupt source being used.
- **ISR Code** – The compiler saves the registers that were used by the ISR. Moreover, if the ISR calls an ordinary function, then the compiler will save all the working registers, even if they are not all used explicitly in the ISR itself. This must be done, because the compiler cannot know, in general, which resources are used by the called function.

15.7 Nesting Interrupts

Interrupts may be nested. The interrupt priority scheme implemented in the PIC32 architecture allows you to specify which interrupt sources may be interruptible by others. See your device data sheet for explicit details on interrupt operation.

The compiler Interrupt Service Routine prologue code automatically re-enables interrupts by default.

15.8 Enabling/Disabling Interrupts

The following builtin functions inspect or manipulate the current CPU interrupt state:

```
unsigned int __builtin_get_isr_state(void)
void __builtin_set_isr_state(unsigned int)
unsigned int __builtin_disable_interrupts(void)
void __builtin_enable_interrupts(void)
```

15.9 ISR Considerations

There are a few things to consider when writing an interrupt service routine.

As with all compilers, limiting the number of registers used by the interrupt function, or any functions called by the interrupt function, may result in less context switch code being generated and executed by the compiler, see 15.6 Latency. Keeping interrupt functions small and simple will help you achieve this.

When interrupt latency is a concern, avoid calling other functions from your ISR. You may be able to replace a function call with a volatile flag that is handled by your application's main control loop.

16.   Main, Runtime Start-up and Reset

When creating C/C++ code, there are elements that are required to ensure proper program operation: a `main` function must be present; start-up code will be needed to initialize and clear variables and setup registers and the processor; and Reset conditions will need to be handled.

16.1   The Main Function

The identifier `main` is special. It must be used as the name of a function that will be the first function to execute in a program. You must always have one and only one function called `main` in your programs. Code associated with `main`, however, is not the first code to execute after Reset. Additional code provided by the compiler and known as the runtime start-up code is executed first and is responsible for transferring control to the `main()` function.

16.2   Runtime Start-Up Code

A C/C++ program requires certain objects to be initialized and the processor to be in a particular state before it can begin execution of its function `main()`. It is the job of the runtime start-up code to perform these tasks. The runtime start-up code is executed before `main()`, but if you require any special initialization to be performed immediately after Reset, you should use the On Reset feature described in **16.3 The On Reset Routine**.

The PIC32 start-up code must perform the following:

1.   Switch to the selected instruction set (ISA) mode.
2.   Jump to NMI handler (_nmi_handler) if an NMI occurred.
3.   Initialize stack pointer and heap.
4.   Initialize global pointer in all register sets available on the selected target device.
5.   Call the application-provided "on reset" procedure (`_on_reset`).
6.   Call the `__pic32_init_cache` procedure to initialize the L1 cache on target devices that feature an L1 cache.
7.   Call the `__pic32_tlb_init_ebi_sqi` procedure to initialize the TLB on -target devices that use pre-mapped EBI and SQI external memory regions.
8.   Clear uninitialized small bss sections.
9.   Initialize data using the linker-generated data-initialization template.
10.   If the target device that features a bus matrix and the application uses a RAM function, initialize the bus matrix for execution from data memory.
11.   Initialize the CP0 registers.
12.   Enable the DSPr2 engine on target devices that feature it. Also, enable and configure the Floating-Point Unit (FPU) on devices that feature it.
13.   Call the "On Bootstrap" procedure (`_on_bootstrap`).
14.   Change the location of exception vectors.
15.   For C++, call the C++ initialization code to invoke all constructors for file-scope static storage objects.
16.   Call `main()`.

The following provisions are made regarding the run-time model:

- Kernel mode only
- KSEG1 only
- RAM functions are attributed with `__ramfunc__` or `__longramfunc__`, (defined in `sys/attribs.h`), meaning that all RAM functions end up in the `.ramfunc` section and the function is `ramfunc` attributed.

16.2.1   Switch to the Selected Instruction Set (ISA) Mode

Some PIC32 MCUs support both the MIPS32 and microMIPS Instruction Set Architecture (ISA) modes. The microMIPS instruction set delivers the same functionality as the MIPS32 ISA, with the added benefit of smaller code size.

Devices that support both the MIPS32 and microMIPS ISA modes use the BOOTISA configuration bit in a device Configuration Word to determine the ISA mode on boot. The device can be configured to boot to either the MIPS32 or the microMIPS ISA mode. See the target-device data sheet for more information on the BOOTISA bit.

The microMIPS ISA supplies assembler-source code compatibility with MIPS32 instead of binary compatibility. Because of this, the XC32 toolchain provides a copy of the runtime start-up code compiled for the MIPS32 ISA as well as a copy compiled for the microMIPS ISA. The toolchain determines which copy to link based on the presence of the -mmicromips command-line option. In the MPLAB X IDE project properties, select *xc32-ld >Option category: Libraries> "Link microMIPS compressed startup code and libraries"* to get the -mmicromips option.

For added flexibility, the default start-up code attempts to ensure that the linked Precompiled mode matches the current ISA mode at runtime. To enable this, a binary code sequence is required that can be run in either instruction set and change code paths, depending on the instruction set that is being used.

The following binary sequence achieves this goal:

```
0x1000wxyz  // where w,x,y,z represent hexadecimal digits
0x00000000
```

For the MIPS32 instruction set, this binary sequence is interpreted as:

```
// branch to location of more MIPS32 instructions
BEQ $0, $0, wxyz
NOP
```

For the microMIPS instruction set, this binary sequence is interpreted as:

```
ADDI32 $0, $0, wxyz // do nothing
NOP                 // fall through to more microMIPS instructions
```

In the default runtime startup-code, we place this binary sequence at the _reset symbol, which is then located at the reset vector by the default linker script. We follow this binary sequence with a `jal _startup` to jump to the remainder of the startup code.

This sequence is included only for devices that support both the MIPS32 and microMIPS ISA modes.

On PIC32M devices, bit 0 of the address indicates the ISA mode. When this bit is clear, the device is running in MIPS32 Mode. When this bit is set, the device is running in either MIPS16 or microMIPS mode, depending on the device core. This means that if you execute a hard-coded jump, bit 0 must be set to the appropriate value for your target function. Hard-coded jumps are most commonly seen when jumping from a bootloader to a bootloaded application.

16.2.2 Jump to NMI Handler (`_nmi_handler`) if an NMI Occurred

If a Non-Maskable Interrupt (NMI) caused entry to the Reset vector, which is located at 0xBFC00000 on PIC32M MIPS cores, the startup code's _reset function jumps to an NMI Handler procedure named _nmi_handler. A weak version of the NMI handler procedure is provided that performs an `ERET`.

To override the default NMI Handler with an application-specific handler, use an assembly-code .S file to create a routine named __nmi_handler. This routine *must* be written in assembly code because the startup code calls this routine before the C runtime environment is initialized. The __nmi_handler routine must either use only the k0, k1 CPU registers or it must save context before using other registers.

16.2.3 Initialize Stack Pointer and Heap

The Stack Pointer (sp) register must be initialized in the start-up code. To enable the start-up code to initialize the sp register, the linker must initialize a variable which points to the end of KSEG0/KSEG1 data memory.

Note: The end of data memory is different based on whether RAM functions exist. If RAM functions exist, then part of the DRM must be configured for the kernel program to contain the RAM functions, and the Stack Pointer is located one word prior to the beginning of the DRM kernel program boundary address. If RAM functions do not exist, then the Stack Pointer is located at the true end of DRM.

The linker allocates the stack to KSEG0 on devices featuring an L1 data cache. It allocates the stack to KSEG1 on devices that do not have an L1 cache.

This variable is named _stack. The user can change the minimum amount of stack space allocated by providing the command line option --defsym _min_stack_size=N to the linker. _min_stack_size is provided by the linker script with a default value of 1024. On a similar note, the user may wish to utilize a heap with their application. While the start-up code does not need to initialize the heap, the standard C libraries (sbrk) must be made aware of the heap location and its size. The linker creates a variable to identify the beginning of the heap. The location of the heap is the end of the utilized KSEG0/KSEG1 data memory.

The linker allocates the heap to KSEG0 on devices that have an L1 cache. It allocates the heap to KSEG1 on devices that do not have an L1 cache.

This variable is named _heap. A user can change the minimum amount of heap space allocated by providing the command line option --defsym _min_heap_size=M to the linker. If the heap is used when the heap size is set to zero, the behavior is the same as when the heap usage exceeds the minimum heap size. Namely, it overflows into the space allocated for the stack.

The heap and the stack use the unallocated KSEG0/KSEG1 data memory, with the heap starting from a low address in KSEG0/KSEG1 data memory, and growing upwards towards the stack while the stack starts at a higher address in KSEG1 data memory and grows downwards towards the heap. The linker attempts to allocate the heap and stack together in the largest gap of memory available in the KSEG0/KSEG1 data memory region. If enough space is not available based on the minimum amount of heap size and stack size requested, the linker issues an error.

Figure 16-1. Stack and Heap Layout

Figure 16-2. Stack and Heap Layout with RAM Functions

The linker must then group all of the above input sections together. This grouping is handled by the default linker script. The run-time start-up code must initialize the gp register to point to the "middle" of this output section. To enable the start-up code to initialize the gp register, the linker script must initialize a variable which is 32 KB from the start of the output section containing the "small" variables and constants. This variable is named _gp (to match core linker scripts). Besides being initialized in the standard GPR set, the Global Pointer must also be initialized in the register shadow set.

16.2.4   Initialize Global Pointer

The compiler toolchain supports Global Pointer (`gp`) relative addressing. Loads and stores to data residing within 32KB of either side of the address stored in the `gp` register can be performed in a single instruction using the `gp` register as the base register. Without the Global Pointer, loading data from a static memory area takes two instructions – one to load the Most Significant bits of the 32-bit constant address computed by the compiler/linker and one to do the data load.

To utilize `gp`-relative addressing, the compiler and assembler must group all of the "small" variables and constants into one of the following sections:

`.lit4.`	`lit8`
`.sdata.`	`sbss`
`.sdata.*`	`sbss.*`
`.gnu.linkonce.s.*`	`.gnu.linkonce.sb.*`

The linker must then group all of the above input sections together. This grouping is handled by the default linker script. The run-time start-up code must initialize the `gp` register to point to the "middle" of this output section. To enable the start-up code to initialize the `gp` register, the linker script must initialize a variable which is 32 KB from the start of the output section containing the "small" variables and constants. This variable is named `_gp` (to match core linker scripts).

Some PIC32 MCUs have more than one register set. The additional register sets can be used as interrupt shadow register sets. The Global Pointer must be initialized in each of the register sets. The default start-up code does this by looping through each of the register sets.

In the loop, the CP0 SRSCtl register's PSS field must be set to the shadow set in which to initialize the global pointer. In the source code, we start with the highest register set, as defined by the PIC32_SRS_SET_COUNT macro, and work down to zero. By initializing the global pointer in the previous set as iterate through the register sets, we initialize the register in each of the sets on the device.

16.2.5   Initialize of Clear Variable and RAM Functions Using the Data-Initialization Template

Those non-`auto` objects which are not initialized must be cleared before execution of the program begins. This task is also performed by the runtime start-up code.

Uninitialized variables are those which are not `auto` objects and which are not assigned a value in their definition, for example `output` in the following example:

```
int output;
int main(void) { ...
```

Such uninitialized objects will only require space to be reserved in RAM where they will reside and be accessed during program execution (runtime).

There are two uninitialized data sections—`.sbss` and `.bss`. The `.sbss` section is a data segment containing uninitialized variables less than or equal to n bytes where n is determined by the `-Gn` command line option. The `.bss` section is a data segment containing uninitialized variables not included in `.sbss`.

Another task of the runtime start-up code is to ensure that any initialized variables contain their initial value before the program begins execution. Initialized variables are those which are not `auto` objects and which are assigned an initial value in their -definition, for example `input` in the following example:

```
int input = 88;
int main(void) { ...
```

Such initialized objects have two components: their initial value (0x0088 in the above example) stored in program memory (that is, placed in the HEX file), and space for the -variable reserved in RAM, which will reside and be accessed during program execution -(runtime).

The runtime start-up code will copy all the blocks of initial values from program memory to RAM so the variables will contain the correct values before `main` is executed.

Since `auto` objects are dynamically created, they require code to be positioned in the function in which they are defined to perform their initialization. It is possible that the initial value of an `auto` object may change on each instance of the function and so the initial values cannot be stored in program memory and copied. As a result, initialized `auto` objects are not considered by the runtime start-up code, but are instead initialized by assembly code in each function output.

Variables whose contents should be preserved over a Reset should be qualified with the `persistent` attribute, see 9.9 Standard Type Qualifiers. Such variables are linked at a different area of memory and are not altered by the runtime start-up code in any way.

Four initialized data sections exist: `.sdata`, `.data`, `.lit4`, and `.lit8`. The `.sdata` section is a data segment containing initialized variables less than or equal to n bytes, where n is determined by the `-Gn` command line option. The `.data` section is a data segment containing initialized variables not included in `.sdata`. The `.lit4` and `.lit8` sections contain constants, (usually floating-point) which the assembler stores in memory rather than in the instruction stream.

Note: Initialized `auto` variables can impact code performance, particularly if the objects are large in size. Consider using global or `static` objects instead.

In order to clear or initialize these sections, the linker creates a data-initialization template, which is loaded into an output section named `.dinit`. The linker creates this special `.dinit` section, allocated in program memory, to hold the template for the run-time initialization of data. The C/C++ start-up module, `crt0.o`, interprets this template and copies initial data values into initialized data sections. This includes sections containing `ramfunc` attributed functions. Other data sections (such as `.bss`) are cleared before the `main()` function is called. The persistent data section (`.pbss`) is not affected. When the application's main program takes control, all variables and RAM functions in data memory have been initialized.

The data initialization template contains one record for each output section in data memory. The template is terminated by a null instruction word. The format of a data initialization record is:

```
/* data init record */
struct data_record {
  char *dst;            /* destination address  */
  unsigned int len;     /* length in bytes      */
  unsigned int format;  /* format code          */
  char dat[0];          /* variable-length data */
};
```

The first element of the record is a pointer to the section in data memory. The second and third elements are the section length and format code, respectively. The last element is an optional array of data bytes. For bss-type sections, no data bytes are required.

Currently supported format codes are:

- 0 – Fill the output section with zeros
- 1 – Copy each byte of data from the data array

16.2.6 Initialize Bus Matrix Registers

On some of the PIC32 MCUs, the bus matrix registers (`BMXDKPBA`, `BMXDUDBA`, `BMXDUPBA`) must be initialized by the start-up code if any RAM functions exist. The startup code leaves these registers in their state when RAM functions do not exist in the projects. The linker collects all RAM functions and allocates them to a section of data memory that is aligned on a 2K-alignment boundary. To determine whether any RAM functions exist in the application, the linker provides a variable that contains the beginning address of this section. This variable is named `_ramfunc_begin`.

In addition, the linker provides a 2K-aligned variable required for the boundary register (BMXDKPBA). The variable is named `_bmxdkpba_address`. The linker also provides two variables that contains the addresses for the bus matrix register. These variables are named `_bmxdkpba_address`, `_bmxdudba_address`, and `_bmxdupba_address`.

The linker ensures that RAM functions are aligned to a 2K-alignment boundary as is required by the BMXDKPBA register.

On other PIC32 devices, no special bus initialization is required to execute RAM functions.

Figure 16-3. Bus Matrix Initialization
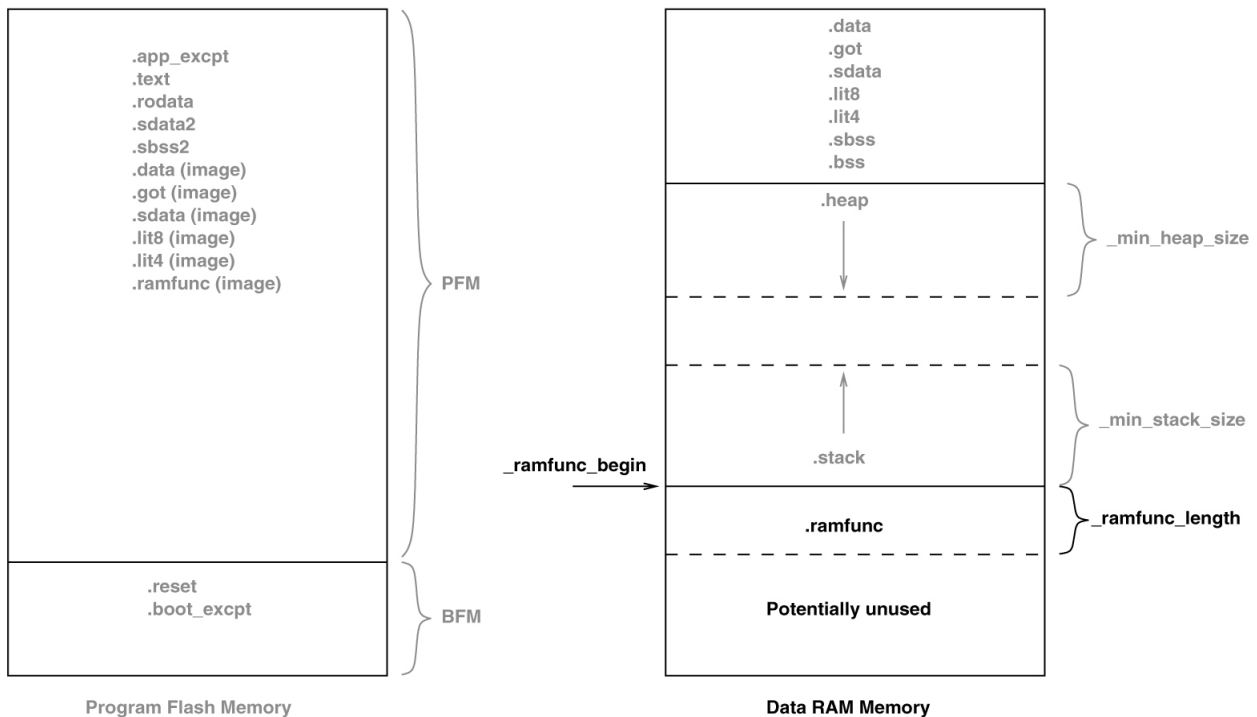
Initialize CP0 Registers

The `CP0` registers are initialized in the following order:

1. `Count` register
2. `Compare` register
3. `EBase` register
4. `IntCtl` register
5. `Cause` register
6. `Status` register

Hardware Enable Register (`HWREna` – CP0 Register 7, Select 0)

This register contains a bit mask that determines which hardware registers are accessible via the `RDHWR` instruction. Privileged software may determine which of the hardware registers are accessible by the `RDHWR` instruction. In doing so, a register may be virtualized at the cost of handling a Reserved Instruction Exception, interpreting the instruction, and returning the virtualized value. For example, if it is not desirable to provide direct access to the `Count` register, access to the register may be individually dis-abled, and the return value can be virtualized by the operating system.

No initialization is performed on this register in the PIC32 start-up code.

Bad Virtual Address Register (`BadVAddr` – CP0 Register 8, Select 0)

This register is a read-only register that captures the most recent virtual address that caused an Address Error exception (`AdEL` or `AdES`).No initialization is performed on this register in the PIC32 start-up code.

Count Register (`Count` – CP0 Register 9, Select 0)

This register acts as a timer, incrementing at a constant rate, whether or not an instruction is executed, retired, or any forward progress is made through the pipeline. The counter increments every other clock if the `DC` bit in the `Cause` register is '0'. The `Count` register can be written for functional or diagnostic purposes, including at Reset or to synchronize processors. By writing the `CountDM` bit in the `Debug` register, it is possible to control whether the `Count` register continues incrementing while the processor is in Debug mode. This register is cleared in the default PIC32 start-up code.

Status Register (`Status` – CP0 Register 12, Select 0)

This register is a read/write register that contains the operating mode, Interrupt Enabling, and the diagnostic states of the processor. Fields of this register combine to create operating modes for the processor.

The following settings are initialized by the default PIC32 start-up code (`0b000000000x0xx0? 00000000000000000`):

- Access to Coprocessor 0 not allowed in User mode (`CU0 = 0`)
- User mode uses configured endianess (`RE = 0`)
- No change to exception vectors location (`BEV = no change`)
- No change to flag bits that indicate reason for entry to the Reset exception vector (`SR, NMI = no change`)
- Interrupt masks are cleared to disable any pending interrupt requests (`IM7..IM2 = 0, IM1..IM0 = 0`)
- Interrupt priority level is 0 (`IPL = 0`)
- Base mode is Kernel mode (`UM = 0`)
- Error level is normal (`ERL = 0`)
- Exception level is normal (`EXL = 0`)
- Interrupts are disabled (`IE = 0`)

The DSPr2 engine is enabled on target devices featuring the DSPr2 engine (MX = 1).

The IEEE 754 compliant Floating-Point Unit is enabled for target devices that support the FPU. The FPU is configured in the FR64 mode, which defines 32 64-bit float-ing-point general registers (FPRs) with all formats supported in each register (`CU1=1`) and (`FR=1`).

Interrupt Control Register (`IntCtl` – CP0 Register 12, Select 1)

This register controls the expanded interrupt capability added in Release 2 of the Architecture, including vectored interrupts and support for an external interrupt controller.

This register contains the vector spacing for interrupt handling. The vector spacing portion of this register (bits 9..5) is initialized with the value of the `_vector_spacing` symbol by the PIC32 start-up code. All other bits are set to '1'.

Shadow Register Control Register (`SRSCtl` – CP0 Register 12, Select 2)

This register controls the operation of the GPR shadow sets in the processor. The default startup code uses the SRSCtl register when it initializes the Global Pointer register in all register sets. However, it restores the original SRSCtl value after the GP register is initialized.

Shadow Register Map Register (`SRSMap` – CP0 Register 12, Select 3)

This register contains eight 4-bit fields that provide the mapping from a vector number to the shadow set number to use when servicing such an interrupt. The values from this register are not used for a non-interrupt exception, or a non-vectored interrupt (`CauseIV = 0` or `IntCtlVS = 0`). In such cases, the shadow set number comes from `SRSCtlESS`. If `SRSCtlHSS` is zero, the results of a software read or write of this register are UNPREDICTABLE. The operation of the processor is UNDEFINED if a value is written to any field in this register that is greater than the value of `SRSCtlHSS`. The `SRSMap` register contains the shadow register set numbers for vector numbers 7..0. The same shadow set number can be established for multiple interrupt vectors, creating a many-to-one mapping from a vector to a single shadow register set number.

No initialization is performed on this register in the PIC32 start-up code.

Cause Register (`Cause` – CP0 Register 13, Select 0)

This register primarily describes the cause of the most recent exception. In addition, fields also control software interrupt requests and the vector through which interrupts are dispatched. With the exception of the `DC`, `IV`, and `IP1..IP0` fields, all fields in the `Cause` register are read-only. Release 2 of the Architecture added optional support for an External Interrupt Controller (EIC) interrupt mode, in which `IP7..IP2` are interpreted as the Requested Interrupt Priority Level (`RIPL`).

The following settings are initialized by the PIC32 start-up code:

- Enable counting of `Count` register (`DC = no change`)
- Use the special exception vector (`16#200`) (`IV = 1`)
- Disable software interrupt requests (`IP1..IP0 = 0`)

Exception Program Counter (`EPC` – CP0 Register 14, Select 0)

This register is a read/write register that contains the address at which processing resumes after an exception has been serviced. All bits of the `EPC` register are significant and must be writable. For synchronous (precise) exceptions, the `EPC` contains one of the following:

- The virtual address of the instruction that was the direct cause of the exception
- The virtual address of the immediately preceding branch or jump instruction, when the exception causing instruction is a branch delay slot and the `Branch Delay` bit in the `Cause` register is set.

On new exceptions, the processor does not write to the `EPC` register when the `EXL` bit in the `Status` register is set; however, the register can still be written via the `MTC0` instruction.

No initialization is performed on this register in the PIC32 start-up code.

Processor Identification Register (`PRId` – CP0 Register 15, Select 0)

This register is a 32-bit read-only register that contains information identifying the manufacturer, manufacturer options, processor identification, and revision level of the processor.

No initialization is performed on this register in the PIC32 start-up code.

Exception Base Register (`EBase` – CP0 Register 15, Select 1)

This register is a read/write register containing the base address of the exception vectors used when `Status`BEV equals 0, and a read-only CPU number value that may be used by software to distinguish different processors in a multi-processor system. The `EBase` register provides the ability for software to identify the specific processor within a multi-processor system, and allows the exception vectors for each processor to be different, especially in systems composed of heterogeneous processors. Bits 31..12 of the `EBase` register are concatenated with zeros to form the base of the exception vectors when `Status`BEV is 0. The exception vector base address comes from fixed defaults when `Status`BEV is 1, or for any EJTAG Debug exception. The Reset state of bits 31..12 of the `EBase` register initialize the exception base register to `16#8000-0000`, providing backward compatibility with Release 1 implementations. Bits 31..30 of the `EBase` register are fixed with the value `2#10` to force the exception base address to be in KSEG0 or KSEG1 unmapped virtual address segments.

If the value of the exception base register is to be changed, this must be done with `Status`BEV equal 1. The operation of the processor is UNDEFINED if the Exception Base field is written with a different value when `Status`BEV is 0.

Combining bits 31..30 with the Exception Base field allows the base address of the exception vectors to be placed at any 4K byte page boundary. If vectored interrupts are used, a vector offset greater than 4K byte can be generated. In this case, bit 12 of the Exception Base field must be zero. The operation of the processor is UNDEFINED if software writes bit 12 of the Exception Base field with a 1 and enables the use of a vectored interrupt whose offset is greater than 4K bytes from the exception base address.

This register is initialized with the value of the `_ebase_address` symbol by the PIC32 start-up code. `_ebase_address` is provided by the linker script with a default value of the start of KSEG1 program memory. The user can change this value by providing the command line option `--defsym _ebase_address=`*A* to the linker.

Config Register (`Config` – CP0 Register 16, Select 0)

This register specifies various configuration and capabilities information. Most of the fields in the `Config` register are initialized by hardware during the Reset exception pr-ocess, or are constant.

No initialization is performed on this register in the PIC32 start-up code.

Config1 Register (`Config1` – CP0 Register 16, Select 1)

This register is an adjunct to the `Config` register and encodes additional information about the capabilities present on the core. All fields in the `Config1` register are read-only.

No initialization is performed on this register in the PIC32 start-up code.

Config2 Register (`Config2` – CP0 Register 16, Select 2)

This register is an adjunct to the `Config` register and is reserved to encode additional capabilities information. `Config2` is allocated for showing the configuration of level 2/3 caches. These fields are reset to 0 because L2/L3 caches are not supported on the core. All fields in the `Config2` register are read-only.

No initialization is performed on this register in the PIC32 start-up code.

Config3 Register (`Config3` – CP0 Register 16, Select 3)

This register encodes additional capabilities. All fields in the `Config3` register are read-only.

No initialization is performed on this register in the PIC32 start-up code.

Debug Register (`Debug` – CP0 Register 23, Select 0)

This register is used to control the debug exception and provide information about the cause of the debug exception, and when re-entering at the debug exception vector due to a normal exception in Debug mode. The read-only information bits are updated every time the debug exception is taken, or when a normal exception is taken when already in Debug mode. Only the `DM` bit and the `EJTAGver` field are valid when read from non-Debug mode. The values of all other bits and fields are UNPREDICTABLE. Operation of the processor is UNDEFINED if the `Debug` register is written from non-Debug mode.

No initialization is performed on this register in the PIC32 start-up code.

Trace Control Register (`TraceControl` – CP0 Register 23, Select 1)

This register provides control and status information. The `TraceControl` register is only implemented if the EJTAG Trace capability is present.

No initialization is performed on this register in the PIC32 start-up code.

Trace Control 2 Register (`TraceControl2` – CP0 Register 23, Select 2)

This register provides additional control and status information. The `TraceControl2` register is only implemented if the EJTAG Trace capability is present.

No initialization is performed on this register in the PIC32 start-up code.

User Trace Data Register (`UserTraceData` – CP0 Register 23, Select 3)

When this register is written to, a trace record is written indicating a type 1 or type 2 user format. This type is based on the `UT` bit in the `TraceControl` register. This register cannot be written in consecutive cycles. The trace output data is UNPREDICTABLE if this register is written in consecutive cycles. The `UserTraceData` register is only implemented if the EJTAG Trace capability is present.

No initialization is performed on this register in the PIC32 start-up code.

TraceBPC Register (`TraceBPC` – CP0 Register 23, Select 4)

This register is used to control start and stop of tracing using an EJTAG hardware breakpoint. The hardware breakpoint would then be set as a triggered source and optionally also as a Debug exception breakpoint. The `TraceBPC` register is only impl-emented if both the hardware breakpoints and the EJTAG Trace cap are present.

No initialization is performed on this register in the PIC32 start-up code.

Debug2 Register (`Debug2` – CP0 Register 23, Select 5)

This register holds additional information about complex breakpoint exceptions. The `Debug2` register is only implemented if complex hardware breakpoints are present.

No initialization is performed on this register in the PIC32 start-up code.

Debug Exception Program Counter (`DEPC` – CP0 Register 24, Select 0)

This register is a read/write register that contains the address at which processing resumes after a debug exception or Debug mode exception has been serviced. For synchronous (precise) debug and Debug mode exceptions, the `DEPC` contains either:

- The virtual address of the instruction that was the direct cause of the debug exception, or
- The virtual address of the immediately preceding branch or jump instruction, when the debug exception causing instruction is in a branch delay slot, and the Debug Branch Delay (`DBD`) bit in the `Debug` register is set.

For asynchronous debug exceptions (debug interrupt, complex break), the `DEPC` contains the virtual address of the instruction where execution should resume after the debug handler code is executed.

No initialization is performed on this register in the PIC32 start-up code.

Error Exception Program Counter (`ErrorEPC` – CP0 Register 30, Select 0)

This register is a read/write register, similar to the `EPC` register, except that it is used on error exceptions. All bits of the `ErrorEPC` are significant and must be writable. It is also used to store the program counter on Reset, Soft Reset, and Non-Maskable Interrupt (NMI) exceptions. The `ErrorEPC` register contains the virtual address at which instruction processing can resume after servicing an error. This address can be:

- The virtual address of the instruction that caused the exception, or
- The virtual address of the immediately preceding branch or jump instruction when the error causing instruction is a branch delay slot.

Unlike the `EPC` register, there is no corresponding branch delay slot indication for the `ErrorEPC` register.

No initialization is performed on this register in the PIC32 start-up code.

Debug Exception Save Register (`DeSave` – CP0 Register 31, Select 0)

This register is a read/write register that functions as a simple memory location. This register is used by the debug exception handler to save one of the GPRs that is then used to save the rest of the context to a pre-determined memory area (such as in the EJTAG Probe). This register allows the safe debugging of exception handlers and other types of code where the existence of a valid stack for context saving cannot be assumed.

No initialization is performed on this register in the PIC32 start-up code.

16.2.7 Call "On Bootstrap" Procedure

A procedure is called after initializing the CP0 registers. This procedure allows your application to perform actions during bootstrap (that is, while `StatusBEV` is set) and before entering into the main routine. An empty weak version of this procedure (`_on_bootstrap`) is provided with the start-up code. This procedure may be used for performing hardware initialization and/or for initializing the environment required by an RTOS.

16.2.8 Change Location of Exception Vectors

Immediately before executing any application code, the `StatusBEV` is cleared to change the location of the exception vectors from the bootstrap location to the normal location.

16.2.9 Call the C++ Initialization Code

Invoke all constructors for C++ file-scope static-storage objects. The startup code must call the constructors last because the low-level initialization must be done before executing application code.

16.2.10 Call Main

The last thing that the start-up code performs is a call to the main routine. If the user returns from main, the start-up code goes into an infinite loop. When you are compiling for use with a debugger in MPLAB X IDE with the `-mdebugger` option, this loop contains a software breakpoint.

16.2.11 Symbols Required by Start-Up Code and C/C++ Library

This section details the symbols that are required by the start-up code and C/C++ library. Currently the default device-specific linker script defines these symbols. If an application provides a custom linker script, the user must ensure that all of the following symbols are provided in order for the start-up code and C library to function:

Symbol Name	Description
`_bmxdkpba_address`	The address to place into the `BMXDKPBA` register if `_ramfunc_length` is greater than 0.
`_bmxdudba_address`	The address to place into the `BMXDUDBA` register if `_ramfunc_length` is greater than 0.

Symbol Name	Description
..........continued	
_bmxdupba_address	The address to place into the BMXDUPBA register if _ramfunc_length is greater than 0.
_ebase_address	The initialization value for the ExceptionBase field of the EBASE register. The ExceptionBase is the base address for the exception vectors, adjustable to a resolution of 4 Kbytes. The default device-specific linker scripts provided with the XC32 toolchain provide a default location for the ExceptionBase.
_end	The end of data allocation.
_gp	Points to the "middle" of the small variables region. By convention this is 0x8000 bytes from the first location used for small variables.
_heap	The starting location of the heap in DRM.
_ramfunc_begin	The starting location of the RAM functions. This should be located at a 2K boundary as it is used to initialize the BMXDKPBA register.
_ramfunc_length	The length of the .ramfunc section.
_stack	The starting location of the stack in DRM. Remember that the stack grows from the bottom of data memory so this symbol should point to the bottom of the section allocated for the stack.
_vector_spacing	The initialization value for the vector spacing field in the IntCtl register.

16.2.12 Exceptions

In addition, two weak general exception handlers are provided that can be overridden by the application — one to handle exceptions when StatusBEV is 1 (_bootstrap_exception_handler), and one to handle exceptions when StatusBEV is 0 (_general_exception_handler). Both the weak Reset exception handler and the weak general exception handler provided with the start-up code causes a software Reset. The start-up code arranges for a jump to the bootstrap exception handler to be located at 0xBFC00380, and a jump to the general exception handler to be located at EBASE + 0x180.

Both handlers must be attributed with the nomips16 [for example, __attribute__ ((nomips16))], since the start-up code jumps to these functions.

When the BOOTISA configuration bit is set for exceptions to be in the microMIPS mode, both handlers must be attributed with the micromips attribute [for example, __attribute__ ((micromips))].

Figure 16-4. Exceptions
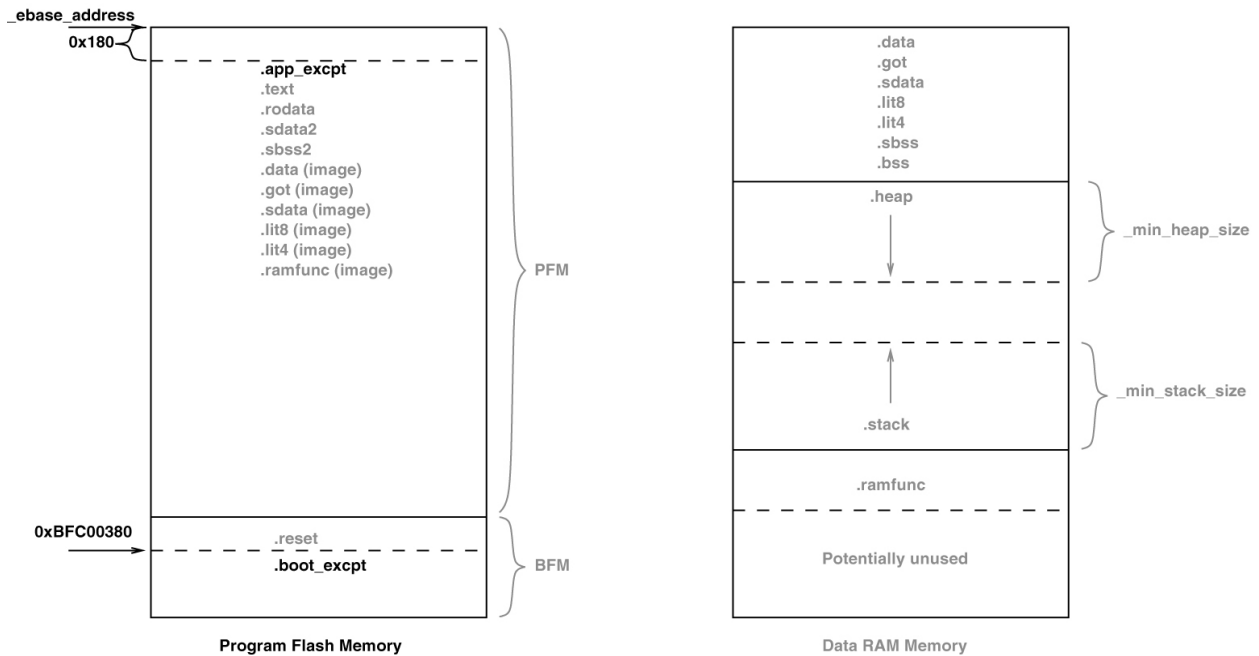
16.3 The On Reset Routine

Some hardware configurations require special initialization, often within the first few instruction cycles after Reset. To achieve this, there is a hook provided via the on Reset routine.

This routine is called after initializing a minimum 'C' context. An empty weak version of this procedure (`_on_reset`) is provided with the start-up code. A stub for this routine can be found in `pic32m-libs/libpic32/stubs` in the installation directory of your compiler.

Special consideration needs to be taken by the user if this procedure is written in 'C'. Most importantly, statically allocated variables are not initialized (with either the specified initializer or a zero as required for uninitialized variables). The stack pointer has been initialized when this routine is called.

16.3.1 Clearing Objects

The runtime start-up code will clear all memory locations occupied by uninitialized variables so they will contain zero before `main()` is executed.

Variables whose contents should be preserved over a Reset should use the `persistent` attribute, see 9.9 Standard Type Qualifiers for more information. Such variables are linked in a different area of memory and are not altered by the runtime start-up code in any way.

17. Library Routines

17.1 Using Library Routines

Library functions or routines (and any associated variables) will be automatically linked into a program once they have been referenced in your source code. The use of a function from one library file will not include any other functions from that library. Only used library functions will be linked into the program output and consume memory.

Your program will require declarations for any functions or symbols used from libraries. These are contained in the standard C header (`.h`) files. Header files are not library files and the two files types should not be confused. Library files contain precompiled code, typically functions and variable definitions; the header files provide declarations (as opposed to definitions) for functions, variables and types in the library files, as well as other preprocessor macros.

```c
#include <math.h>    // declare function prototype for sqrt

int main(void)
{
  double i;

  // sqrt referenced; sqrt will be linked in from library file
  i = sqrt(23.5);
}
```

A comprehensive and readily accessible set of peripheral libraries, drivers, and system services is available from MPLAB® Harmony. Please see www.microchip.com/mplab/mplab-harmony for details.

18.   Mixing C/C++ and Assembly Language

Assembly language code can be mixed with C/C++ code using two different techniques: writing assembly code and placing it into a separate assembler module, or including it as in-line assembly in a C/C++ module. This section describes how to use assembly language and C/C++ modules together. It gives examples of using C/C++ variables and functions in assembly code, and examples of using assembly language variables and functions in C/C++.

The more assembly code a project contains, the more difficult and time consuming its maintenance will be. As the project is developed, the compiler may work in different ways as some optimizations look at the entire program. The assembly code is more likely to fail if the compiler is updated due to differences in the way the updated compiler may work. These factors do not affect code written in C/C++.

Note: If assembly must be added, it is preferable to write this as self-contained routine in a separate assembly module rather than in-lining it in C code.

18.1   Mixing Assembly Language and C Variables and Functions

The following guidelines indicate how to interface separate assembly language modules with C modules.

- Follow the register conventions described in 13.2 Register Conventions. In particular, registers $4-$7 are used for parameter passing. An assembly -language function will receive parameters, and should pass arguments to called functions, in these registers.
- Table 13-1 describes which registers must be saved across non-interrupt function calls.
- Interrupt functions must preserve all registers. Unlike a normal function call, an interrupt may occur at any point during the execution of a program. When returning to the normal program, all registers must be as they were before the interrupt occurred.
- Variables or functions declared within a separate assembly file that will be referenced by any C source file should be declared as global using the assembler directive .global. Undeclared symbols used in assembly files will be treated as externally defined.

The following example shows how to use variables and functions in both assembly language and C regardless of where they were originally defined.

The file ex1.c defines foo and cVariable to be used in the assembly language file. The C file also shows how to call an assembly function, asmFunction, and how to access the assembly defined variable, asmVariable.

Example 18-1. Mixing C and Assembly

```
/*
** file: ex1.S
*/
#include <xc.h>

    /* define which section (for example "text")
     * does this portion of code resides in. Typically,
     * all your code will reside in .text section as
     * shown below.
     */
    .text

    /*  This is important for an assembly programmer. This
     * directive tells the assembler that don't optimize
     * the order of the instructions as well as don't insert
     * 'nop' instructions after jumps and branches.
     */
    .set noreorder

/************************************************************************
 * asmFunction(int bits)
 * This function clears the specified bites in IOPORT A.
 ************************************************************/
.global asmFunction
.ent asmFunction
asmFunction:
    /* function prologue - save registers used in this function
```

```
     * on stack and adjust stack-pointer
     */
    addiu   sp, sp, -4
    sw      s0, 0(sp)

    la      s0, LATACLR
    sw      a0, 0(s0)       /* clear specified bits */

    la      s0, PORTA
    lw      s1, 0(s0)
    la      s0, cVariable
    sw      s1, 0(s0)       /* keep a copy */

    /* function epilogue - restore registers used in this function
     * from stack and adjust stack-pointer
     */
    lw      s0, 0(sp)
    addiu   sp, sp,

    addu    s1, ra, zero
    jal     foo
    nop
    addu    ra, s1, zero
    nop
    /* return to caller */
    jr      ra
    nop
.end asmFunction

    .bss
    .global asmVariable
    .align 2
asmVariable: .space 4
```

The file `ex1.S` defines `asmFunction` and `asmVariable` as required for use in a linked application. The assembly file also shows how to call a C function, `foo`, and how to access a C defined variable, `cVariable`.

```
/*
 * file: ex2.c
 */
#include <xc.h>
extern void asmFunction(int bits);
extern unsigned int asmVariable;
volatile unsigned int cVariable = 0;
volatile unsigned int jak = 0;

int
main(void) {
    TRISA = 0;
    LATA = 0xC6FFul;
    asmFunction(0xA55Au);
    while (1) {
        asmVariable++;
    }
}

void
foo(void) {
    jak++;
}
```

In the C file, `ex2.c`, external references to symbols declared in an assembly file are declared using the standard `extern` keyword; note that `asmFunction` is a `void` function and is declared accordingly.

In the assembly file, `ex1.S`, the symbols `asmFunction` and `asmVariable` are made globally visible through the use of the `.global` assembler directive and can be accessed by any other source file.

18.2 Using Inline Assembly Language

Within a C/C++ function, the `asm` statement may be used to insert a line of assembly-language code into the assembly language that the compiler generates. Inline -assembly has two forms: simple and extended.

In the **simple** form, the assembler instruction is written using the syntax:

```
asm ("instruction");
```

where `instruction` is a valid assembly-language construct. If you are writing inline assembly in ANSI C programs, write `__asm__` instead of `asm`.

Note: Only a single string can be passed to the simple form of inline assembly.

In an **extended** assembler instruction using `asm`, the operands of the instruction are specified using C/C++ expressions. The extended syntax is:

```
asm("template" [ : [ "constraint"(output-operand) [ , ... ] ]
              [ : [ "constraint"(input-operand) [ , ... ] ]
                  [ "clobber" [ , ... ] ]
              ]
          ]);
```

You must specify an assembler instruction `template`, plus an operand `constraint` string for each operand. The `template` specifies the instruction mnemonic, and optionally placeholders for the operands. The `constraint` strings specify operand constraints, for example, that an operand must be in a register (the usual case), or that an operand must be an immediate value.

Constraint letters and modifiers supported by the compiler are listed in the following tables.

Table 18-1. Register Constraint Letters Supported by the Compiler

Letter	Constraint
c	A register suitable for use in an indirect jump
d	An address register. This is equivalent to `@code{r}` unless generating MIPS16 code
ka	Registers that can be used as the target of multiply-accumulate instructions
l	The `@code{lo}` register. Use this register to store values that are no bigger than a word
x	The concatenated `@code{hi}` and `@code{lo}` registers. Use this register to store double-word values

Table 18-2. Integer Constraint Letters Supported by the Compiler

Letter	Constraint
I	A signed 32-bit constant (for arithmetic instructions)
J	Integer zero
K	An unsigned 32-bit constant (for logic instructions)
L	A signed 32-bit constant in which the lower 32 bits are zero. Such constants can be loaded using `@code{lui}`
M	A constant that cannot be loaded using `@code{lui}`, `@code{addiu}` or `@code{ori}`
N	A constant in the range -65535 to -1 (inclusive)
O	A signed 15-bit constant

..........continued	
Letter	**Constraint**
P	A constant in the range 1 to 65535 (inclusive)

Table 18-3. General Constraint Letters Supported by the Compiler

Letter	**Constraint**
R	An address that can be used in a non-macro load or store.

Table 18-4. Constraint Modifiers Supported by the Compiler

Letter	**Constraint**
=	Means that this operand is write-only for this instruction: the previous value is discarded and replaced by output data
+	Means that this operand is both read and written by the instruction
&	Means that this operand is an `earlyclobber` operand, which is modified before the instruction is finished using the input operands. Therefore, this operand may not lie in a register that is used as an input operand or as part of any memory address
d	Second register for operand number *n*, that is, `%d`*n*`..`
q	Fourth register for operand number **n**, that is, `%q`*n*`..`
t	Third register for operand number *n*, that is, `%t`*n*`..`

18.2.1   Inline Examples

Insert Bit Field

This example demonstrates how to use the `INS` instruction to insert a bit field into a 32-bit wide variable. This function-like macro uses inline assembly to emit the `INS` instruction, which is not commonly generated from C/C++ code.

```
/* MIPS32r2 insert bits */
#define _ins(tgt,val,pos,sz) __extension__({               \
    unsigned int __t = (tgt), __v = (val);                 \
    __asm__ ("ins %0,%z1,%2,%3"            /* template */ \
        : "+d" (__t)                       /* output   */ \
        : "dJ" (__v), "I" (pos), "I" (sz)); /* input    */ \
    __t;                                                   \
})
```

Here `__v`, `pos`, and `sz` are input operands. The `__v` operand is constrained to be of type 'd' (an address register) or 'J' (integer zero). The `pos` and `sz` operands are constrained to be of type 'I' (a signed 32-bit constant).

The `__t` output operand is constrained to be of type 'd' (an address register). The '+' modifier means that this operand is both read and written by the instruction and so the operand is both an input and an output.

The following example shows this macro in use.

```
unsigned int result;
void example (void)
{
    unsigned int insertval = 0x12;
    result = 0xAAAAAAAAu;
    result = _ins(result, insertval, 4, 8);
    /* result is now 0xAAAAA12A */
}
```

For this example, the compiler may generate assembly code similar to the following.

```
    li      $2,-1431699456          # 0xaaaa0000
    ori     $2,$2,0xaaaa         # 0xaaaa0000 | 0xaaaa

    li      $3,18           # 0x12
    ins     $2,$3,4,8               # inline assembly

    lui     $3,%hi(result)       # assign the result back
    j       $31             # return
    sw      $2,%lo(result)($3)
```

Multiple Assembler Instructions

This example demonstrates how to use the WSBH and ROTR instructions together for a byte swap. The WSBH instruction is a 32-bit byte swap within each of the two halfwords. The ROTR instruction is a rotate right by immediate. This function-like macro uses inline assembly to create a "byte-swap word" using instructions that are not commonly generated from C/C++ code.

The following shows the definition of the function-like macro, _bswapw.

```
/* MIPS32r2 byte-swap word */
#define _bswapw(x) __extension__({          \
    unsigned int __x = (x), __v;            \
    __asm__ ("wsbh %0,%1;\n\t"              \
            "rotr %0,16" /* template */ \
            : "=d" (__v)  /* output */   \
            : "d" (__x))  /* input*/ ;   \
    __v;                                    \
})
```

Here __x is the C expression for the input operand. The operand is constrained to be of type 'd', which denotes an address register.

The C expression __v is the output operand. This operand is also constrained to be of type 'd'. The '=' means that this operand is write-only for this instruction: the previous value is discarded and replaced by output data.

The function-like macro is shown in the following example assigning to result the content of value, swapped.

```
unsigned int result;
int example (void)
{
    unsigned int value = 0x12345678u;
    result = _bswapw(value);
    /* result == 0x78563412 */
}
```

The compiler may generate assembly code similar to the following for this example:

```
    li      $2,305397760            # 0x12340000
    addiu   $2,$2,22136             # 0x12340000 + 0x5678
    wsbh    $2,$2                   # From inline asm
    rotr    $2,16                   # From inline asm
    lui     $2,%hi(result)          # assign back to result
    j       $31                     # return
    sw      $3,%lo(result)($2)
```

18.2.2 Equivalent Assembly Symbols

C/C++ symbols can be accessed directly with no modification in extended assembly code.

18.3 Predefined Macros

Several predefined macros are available once you include <xc.h>. The exact operation of these macros is dependent on the instruction set employed. The following table shows general purpose predefined macros and their operation.

Table 18-5. Predefined Macros in XC.H

Macro	Description
_bcc0(rn, sel, clr)	For the CP0 register specified by rn and sel, clear bits corresponding to those bits in clr which are non-zero.
_bcsc0(rn, sel, clr, set)	For the CP0 register specified by rn and sel, clear bits corresponding to those bits in clr which are non-zero, and set bits corresponding to those bits in set which are non-zero.
_bsc0(rn, sel, set)	For the CP0 register specified by rn and sel, clear bits corresponding to those bits in clr which are non-zero.
_bswapw(x)	See <xc.h> file. Byte-swap word.
_cache(op,addr)	Do an operation to a cache line. See the device documentation for details on the available operations.
_clo(x)	Count leading ones in x.
_clz(x)	Count leading zeros in x.
_ctz(x)	Count trailing zeros in x.
_dclo(x)	Simulate 64-bit count leading ones in x.
_dclz(x)	Simulate 64-bit count leading zeros in x.
_dctz(x	Simulate 64-bit count trailing zeros in x.
_ehb()	Insert Execution Hazard Barrier instruction.
_ext(x,pos,sz)	See <xc.h> file. Extract bitfield from a 32-bit variable.
_get_byte(addr, errp)	Return the least significant byte of addr.
_get_dword(addr, errp)	Return the least significant 64-bit word of addr.
_get_half(addr, errp)	Return the least significant 16-bit word of addr.
_get_word(addr, errp)	Return the least significant 32-bit word of addr.
_ins(tgt,val,pos,sz)	See <xc.h> file. Insert bits.
_jr_hb()	See <xc.h> file. Jump register with hazard barrier.
_mfc0(rn, sel)	See <xc.h> file. Move a value from a coprocessor 0 register.
_mtc0(rn, sel, v)	See <xc.h> file. Move a value to a coprocessor 0 register.
_mxc0(rn, sel, v)	See <xc.h> file. Exchange a value with a value in a coprocessor 0 register.
_nop()	Insert a No Operation instruction.
_prefetch(hint,x)	Prefetch instruction for memory reference optimization. An application that knows in advance it may need data can arrange for it to be brought into cache. 'hint' defines which sort of prefetch this is.
_put_byte(addr, v)	Write the least significant byte of addr with v.
_put_dword(addr, v)	Write the least significant 64-bit word of addr with v.
_put_half(addr, v)	Write the least significant 16-bit word of addr with v.
_put_word(addr, v)	Write the least significant 32-bit word of addr with v.

..........continued	
Macro	**Description**
`_rdpgpr(regno)`	See `<xc.h>` file. Read register from previous register set.
`_sync()`	Insert Synchronize Shared Memory instruction.
`_synci(addr)`	Synchronize the I-cache with the D-cache; Run instruction for each cache-line-sized block after writing instructions but before executing them.
`_wait()`	Insert instruction to enter Standby mode.
`_wrpgpr(regno, val)`	See `<xc.h>` file. Write to a register in the previous register set.
`_wsbh(x)`	See `<xc.h>` file. 32-bit byte-swap within each of the two halfword.
`__XC32_PART_SUPPORT_UPDATE`	This macro expands to the letter corresponding to the part-support update version being used. The value is based upon the major and minor version numbers of the current release. For example, part-support update version v1.42(B) will have `#define __XC32_PART_SUPPORT_UPDATE B`
`__XC32_PART_SUPPORT_VERSION`	This macro expands to a numeric value corresponding to the part-support update version being used. The value is based upon the major and minor version numbers of the current release. For example, part-support update version v1.42(B) will have `#define __XC32_PART_SUPPORT_UPDATE 1420`.

19. Optimizations

Different MPLAB XC32 C/C++ Compiler editions support different levels of optimization. Some editions are free to download and others must be purchased. Visit https://www.microchip.com/mplab/compilers for more information on C and C++ licenses.

The compiler editions are:

Edition	Cost	Description
Professional (PRO)	Yes	Implemented with the highest optimizations and performance levels.
Free	No	Implemented with the most code optimizations restrictions.
Evaluation (EVAL)	No	PRO edition enabled for 60 days and then reverts to Free edition.

Setting Optimization Levels

Different optimizations may be set ranging from no optimization to full optimization, depending on your compiler edition. When debugging code, you may wish to not optimize your code to ensure expected program flow.

For details on compiler options used to set optimizations, see 6.7.7 Options for Controlling Optimization.

20. Preprocessing

All C/C++ source files are preprocessed before compilation. Assembly source files that use the .S extension (upper case) are also preprocessed. A large number of options control the operation of the preprocessor and preprocessed code, see 6.7.8 Options for Controlling the Preprocessor.

20.1 C/C++ Language Comments

A C/C++ comment is ignored by the compiler and can be used to provide information to someone reading the source code. They should be used freely.

Comments may be added by enclosing the desired characters within `/*` and `*/`. The comment can run over multiple lines, but comments cannot be nested. Comments can be placed anywhere in C/C++ code, even in the middle of expressions, but cannot be placed in character constants or string literals.

Since comments cannot be nested, it may be desirable to use the `#if` preprocessor directive to comment out code that already contains comments, for example:

```
#if 0
    result = read();  /* TODO: Jim, check this function is right */
#endif
```

Single-line, C++ style comments may also be specified. Any characters following // to the end of the line are taken to be a comment and will be ignored by the compiler, as shown below:

```
result = read(); // TODO: Jim, check this function is right
```

20.2 Preprocessor Directives

MPLAB XC32 C/C++ Compiler accepts all the standard preprocessor directives, which are listed in the following table.

Table 20-1. Preprocessor Directives

Directive	Meaning	Example
#	Preprocessor null directive, do nothing.	#
#assert	Generate error if condition false.	#assert SIZE > 10

#define	Define preprocessor macro.	`#define SIZE 5`
		`#define FLAG`
		`#define add(a,b) ((a)+(b))`

#elif	Short for #else #if.	see #ifdef
#else	Conditionally include source lines.	see #if
#endif	Terminate conditional source inclusion.	see #if
#error	Generate an error message.	#error Size too big

#if	Include source lines if constant expression true.	`#if SIZE < 10`
		`  c = process(10)`
		`#else`
		`  skip();`
		`#endif`

..........continued		
Directive	**Meaning**	**Example**

`#ifdef`	Include source lines if preprocessor symbol defined.	`#ifdef FLAG`
		`  do_loop();`
		`#elif SIZE == 5`
		`  skip_loop();`
		`#endif`

`#ifndef`	Include source lines if preprocessor symbol not defined.	`#ifndef FLAG`
		`  jump();`
		`#endif`

`#include`	Include text file into source.	`#include <stdio.h>` `#include "project.h"`
`#line`	Specify line number and file name for listing.	`#line 3 final`
`#nn`	(Where *nn* is a number) short for `#line` *nn*.	`#20`
`#pragma`	Compiler specific options.	Refer to 20.3 Pragma Directives
`#undef`	Undefines preprocessor symbol.	`#undef FLAG`
`#warning`	Generate a warning message.	`#warning Length not set`

Macro expansion using arguments can use the `#` character to convert an argument to a string, and the `##` sequence to concatenate arguments. If two expressions are being concatenated, consider using two macros in case either expression requires substitution itself, so for example,

```
#define          paste1(a,b)              a##b
#define          paste(a,b)               paste1(a,b)
```

lets you use the `paste` macro to concatenate two expressions that themselves may require further expansion. The replacement token is rescanned for more macro identifiers, but remember that once a particular macro identifier has been expanded, it will not be expanded again if it appears after concatenation.

The type and conversion of numeric values in the preprocessor domain is the same as in the C domain. Preprocessor values do not have a type, but acquire one as soon as they are converted by the preprocessor. Expressions may overflow their allocated type in the same way that C expressions may overflow.

Overflow may be avoided by using a constant suffix. For example, an `L` after the number indicates it should be interpreted as a long once converted.

So, for example:

```
#define MAX 100000*100000
```

and

```
#define MAX 100000*100000L
```

(note the `L` suffix) will define the values `0x540be400` and `0x2540be400`, respectively.

20.3   Pragma Directives

There are certain compile-time directives that can be used to modify the behavior of the compiler. These are implemented through the use of the ANSI standard `#pragma` facility. Any pragma which is not understood by the compiler will be ignored.

The format of a pragma is:

```
#pragma keyword options
```

where `keyword` is one of a set of keywords, some of which are followed by certain `options`. A description of the keywords is given below.

#pragma interrupt

Mark a function as an interrupt handler. The prologue and epilogue code for the function will perform more extensive context preservation. Note that the `interrupt` attribute (rather than this pragma) is the recommended mechanism for marking a function as an interrupt handler. The interrupt pragma is provided for compatibility with other compilers. See 15. Interrupts and 15.4 Exception Handlers.

#pragma vector

Generate a branch instruction at the indicated exception vector that targets the function. Note that the `vector` attribute (rather than this pragma) is the recommended mechanism for generating an exception/interrupt vector. See 15. Interrupts and 15.4 Exception Handlers.

#pragma config

The `#pragma config` directive specifies the processor-specific configuration settings (that is, Configuration bits) to be used by the application. See 8.4 Configuration Bit Access.

20.4 Predefined Macros

These predefined macros are available for use with the compiler.

20.4.1 32-Bit C/C++ Compiler Macros

The compiler defines a number of macros, most with the prefix "`_MCHP_`", which characterize the various target specific options, the target processor and other aspects of the host environment.

`_MCHP_SZINT`	32 or 64, depending on command line options to set the size of an integer (`-mint32 -mint64`).
`_MCHP_SZLONG`	32 or 64, depending on command line options to set the size of an integer (`-mlong32 -mlong64`).
`_MCHP_SZPTR`	32 always since all pointers are 32 bits.
`__mchp_no_float`	Defined if `-mno-float` specified.
`__NO_FLOAT`	Defined if `-mno-float` specified.
`__SOFT_FLOAT`	Defined if `-mno-float` not specified and the specified device does not feature a hardware Floating-Point Unit (FPU). Indicates that floating-point operations are supported via library calls.
`__HARD_FLOAT`	Defined if `-mno-float` and `-msoft-float` are not specified and the specified device features a hardware Floating-Point Unit (FPU). Indicates that floating-point operations utilize the FPU.
`__PIC__` `__pic__`	The translation unit is being compiled for position independent code.
`__PIC32MX` `__PIC32MX__`	Defined when a PIC32MX device is specified with the `-mprocessor` option.
`__PIC32MZ`	Defined when a PIC32MZ device is specified with the `-mprocessor` option.

`__PIC32_FEATURE_SET__`	The compiler predefines a macro based on the features available for the selected device. These macros are intended to be used when writing code to take advantage of features available on newer devices while maintaining compatibility with older devices. Examples: PIC32MX795F512L would use: `__PIC32_FEATURE_SET__ == 795`, and PIC32MX340F128H would use: `__PIC32_FEATURE_SET__ == 340`. Examples: PIC32MZ2048ECH100 would use: `__PIC32_FEATURE_SET "EC" /* PIC32MZ2048ECH100 */` `__PIC32_FEATURE_SET0 69 /* PIC32MZ2048ECH100 */` `__PIC32_FEATURE_SET1 67 /* PIC32MZ2048ECH100 */`
`PIC32MX`	Defined if `-ansi` is not specified.
`__LANGUAGE_ASSEMBLY` `__LANGUAGE_ASSEMBLY__` `_LANGUAGE_ASSEMBLY`	Defined if compiling a pre-processed assembly file (.S files).
`LANGUAGE_ASSEMBLY`	Defined if compiling a pre-processed assembly file (.S files) and `-ansi` is not specified.
`__LANGUAGE_C` `__LANGUAGE_C__` `_LANGUAGE_C`	Defined if compiling a C file.
`LANGUAGE_C`	Defined if compiling a C file and `-ansi` is not specified.
`__LANGUAGE_C_PLUS_PLUS` `__cplusplus` `_LANGUAGE_C_PLUS_PLUS__`	Defined if compiling a C++ file.
`__EXCEPTIONS`	Defined if X++ exceptions are enabled.
`__GXX_RTTI`	Defined if runtime type information is enabled.
`__processor__`	Where "processor" is the capitalized argument to the `-mprocessor` option. for example, `-mprocessor=32mx12f3456` will define `__32MX12F3456__`
`__XC`	Always defined to indicate that this is a Microchip XC compiler.
`__XC32`	Always defined to indicate this the XC32 compiler.
`__VERSION__`	The `__VERSION__` macro expands to a string constant describing the compiler in use. Do not rely on its contents having any particular form, but it should contain at least the release number. Use the `__XC32_VERSION` macro for a numeric version number.
`__XC32_VERSION` or `__C32_VERSION__`	The C compiler defines the constant `__XC32_VERSION`, giving a numeric value to the version identifier. This macro can be used to construct applications that take advantage of new compiler features while still remaining backward compatible with older versions. The value is based upon the major and minor version numbers of the current release. For example, release version 1.03 will have a `__XC32_VERSION` definition of 1030. This macro can be used, in conjunction with standard preprocessor comparison statements, to conditionally include/exclude various code constructs.

`__mips_dsp 1`
`__mips_dspr2 1`
 	
`__mips_dsp_rev 2`
	The C compiler defines these constants when the selected target device supports the DSPr2 engine.

`__mips_micromips 1`	The compiler defines these constants when we are building for the microMIPS compressed ISA as the default using the `-mmicromips` option.
`__mips_soft_float 1`	The compiler defines this constant when we are compiling for software floating-point operations.

See also the device-specific include files (`pic32mx/include/proc/p32*.h`) for other macros that can be used to determine the features available on the selected device. You will find these macros near the end of the header file.

20.4.2 SDE Compatibility Macros

The MIPS® SDE (Software Development Environment) defines a number of macros, most with the prefix "_MIPS_", which characterize various target specific options, some determined by command line options (for example, `-mint64`). Where applicable, these macros will be defined by the compiler in order to ease porting applications and middleware from the SDE to the compiler.

`_MIPS_SZINT`	
`_MIPS_SZLONG`	
`_MIPS_SZPTR`	
`__mips_no_float`	

`__mips__`
`_mips`
 	
`_MIPS_ARCH_PIC32MX`
 	
`_MIPS_TUNE_PIC32MX`
 	
`_R3000`
 	
`__R3000`
 	
`__R3000__`
 	
`__mips_soft_float`
 	
`__MIPSEL`
 	
`__MIPSEL__`
 	
`_MIPSEL`
	Always defined.

`R3000` `MIPSEL`	Defined if `-ansi` is not specified.
`_mips_fpr`	Defined as 32.
`__mips16`	Defined if `-mips16` specified.
`__mips`	Defined as 32.
`__mips_isa_rev`	Defined as 2.
`_MIPS_ISA`	Defined as `_MIPS_ISA_MIPS32`.

20.4.3 Processor-Specific Macros

The proc/p*.h header files define a number of processor-specific macros. To use these macros, `#include <xc.h>` in your source file. This list is not comprehensive.

`__PIC32_HAS_L1CACHE`	Defined if and only if the target device supports an L1 cache
`__PIC32_HAS_MIPS32R2`	Defined if and only if the target device supports the MIPS32r2 Instruction Set
`__PIC32_HAS_MICROMIPS`	Defined if and only if the target device supports the microMIPS32 Instruction Set
`__PIC32_HAS_MIPS16`	Defined if and only if the target device supports the MIPS16 Instruction Set
`__PIC32_HAS_DSPR2`	Defined if and only if the target device supports the DSP-enhanced core
`__PIC32_HAS_FPU64`	Defined if and only if the target device supports the 64-bit Hardware Floating-Point Unit. Also check `__mips_hard_float` to determine if the compiler is set to compile for the FPU.
`_<Interrupt-Source>_VECTOR`	Defined to the vector number for the interrupt source and intended to be used with the vector function attribute. for example, `#define _CORE_TIMER_VECTOR 0`
`_<Interrupt-Source>_IRQ`	Defined to the IRQ number for the interrupt source. for example, `#define _CORE_SOFTWARE_0_IRQ 1`
`_<SFR>_<Bitfield>_POSITION`	Defined to the position of a bitfield within a special function register (SFR). for example, `#define _WDTCON_ON_POSITION 0x0000000F`
`_<SFR>_<Bitfield>_MASK`	Defined to a mask of the bitfield within a special function register (SFR). for example, `#define _WDTCON_ON_MASK 0x00008000`
`_<SFR>_<Bitfield>_LENGTH`	Defined to the length of the bitfield within a special function register (SFR). for example, `#define _WDTCON_ON_LENGTH 0x00000001`
`<Special-Function-Register>`	Defined to the Special Function Register name. for example, `#define T5CON T5CON` This macro allows preprocessor testing for the existence of an SFR before using it. For example: ```#if defined(T5CON)\n T5CONbits.ON = 1;\n#endif```

21. Linking Programs

See the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186) for more detailed information on the linker.

The compiler will automatically invoke the linker unless the compiler has been requested to stop after producing an intermediate file.

Linker scripts are used to specify the available memory regions and where sections should be positioned in those regions.

The linker creates a map file which details the memory assigned to sections. The map file is the best place to look for memory information.

21.1 Replacing Library Symbols

Unlike with the Microchip MPLAB XC8 compiler, not all library functions can be replaced with user-defined routines using MPLAB XC32 C/C++ Compiler. Only weak library functions (see) can be replaced in this way. For those that are weak, any function you write in your code will replace an identically named function in the library files.

21.2 Linker-Defined Symbols

The 32-bit linker defines several symbols that can be used in your C code development. Please see the *MPLAB® XC32 Assembler, Linker and Utilities User's Guide* (DS50002186) for more information.

The linker defines the symbols `_ramfunc_begin` and `_bmxdkpba_address`, which represent the starting address in RAM where ram functions will be accessed, and the corresponding address in the program memory from which the functions will be copied. They are used by the default runtime start-up code to initialize the bus matrix if ram functions exist in the project, see 14.3 Allocation of Function Code.

The linker also defines the symbol `_stack`, which is used by the runtime start-up code to initialize the stack pointer. This symbol represents the starting address for the software stack.

All the above symbols are rarely required for most programs, but may assist you if you are writing your own runtime start-up code.

21.3 Default Linker Script

For PIC32MX Devices Only:

If no linker script is specified on the command line, the linker will use an internal version known as the built-in default linker script. The default linker script has section mapping that is appropriate for all PIC32 MCUs. It uses an INCLUDE directive to include the device-specific memory regions.

The default linker script is appropriate for most PIC32 MCU applications. Only applications with specific memory-allocation needs will require an application-specific linker script. The default linker script can be examined by invoking the linker with the `--verbose` option:

```
xc32-ld --verbose
```

In a normal tool-suite installation, a copy of the default linker script is located at `\pic32mx\lib\ldscripts \elf32pic32mx.x`. Note that this file is only a copy of the default linker script. The script that the linker uses is internal to the linker.

The device-specific portion of the linker script is located in `\pic32mx\pic32mx\lib\proc\32MXGENERIC \procdefs.ld`, where device is the device value specified to the `-mprocessor` compilation-driver (xc32-gcc) option.

For PIC32MZ and Later Devices:

Single-file linker script for PIC32MZ and later devices: The linker script for PIC32MZ devices are contained within a single file (for example, `pic32mx/lib/proc/32MZ2048ECH100/p32MZ2048ECH100.ld`). This eliminates the

dependency on two files (`elf32pic32mx.x` and `procdefs.ld`) used by the older linker-script model. Like before, the xc32-gcc compilation driver will pass the device-specific linker script to the linker when building with `-mprocessor=device` option.

The default linker script contains the following categories of information:

Note: All addresses specified in the linker scripts should be specified as virtual addresses, not physical addresses.

21.3.1 Output Format and Entry Points

The first several lines of the default linker script define the output format and the entry point for the application.

```
OUTPUT_FORMAT("elf32-tradlittlemips")

OUTPUT_ARCH(pic32mx)

ENTRY(_reset)
```

The `OUTPUT_FORMAT` line selects the object file format for the output file. The output object file format generated by the 32-bit language tools is a traditional, little-endian, MIPS, ELF32 format.

The `OUTPUT_ARCH` line selects the specific machine architecture for the output file. The output files generated by the 32-bit language tools contain information that identifies the file was generated for the PIC32 architecture.

The `ENTRY` line selects the entry point of the application. This is the symbol identifying the location of the first instruction to execute. The 32-bit language tools begins execution at the instruction identified by the `_reset` label.

21.3.2 Default Values for Minimum Stack and Heap Sizes

The next section of the default linker script provides default values for the minimum stack and heap sizes.

```
/*
 * Provide for a minimum stack and heap size
 * - _min_stack_size - represents the minimum space that must
 *                     be made available for the stack. Can
 *                     be overridden from the command line
 *                     using the linker's --defsym option.
 * - _min_heap_size  - represents the minimum space that must
 *                     be made available for the heap. Can
 *                     be overridden from the command line
 *                     using the linker's --defsym option.
 */
EXTERN (_min_stack_size _min_heap_size)
PROVIDE(_min_stack_size = 0x400) ;
PROVIDE(_min_heap_size = 0) ;
```

The `EXTERN` line ensures that the rest of the linker script has access to the default values of `_min_stack_size` and `_min_heap_size` assuming that the user does not override these values using the linker's `--defsym` command line option.

The two `PROVIDE` lines ensure that a default value is provided for both `_min_stack_size` and `_min_heap_size`. The default value for the minimum stack size is `1024` bytes (`0x400`). The default value for the minimum heap size is `0` bytes.

21.3.3   Processor Definitions Include File for PIC32MX Family

The next line in the default linker script pulls in information specific to the processor.

```
INCLUDE procdefs.ld
```

The file `procdefs.ld` is included in the linker script at this point. The file is searched for in the current directory and in any directory specified with the `-L` command line option. The compiler shell ensures that the correct directory is passed to the linker with the `-L` command line option based on the processor selected with the `-mprocessor` command line option.

The processor definitions linker script contains the following pieces of information:

- 21.3.3.1 Inclusion of Processor-Specific Object File(s)
- 21.3.3.3 Base Exception Vector Address and Vector Spacing Symbols
- 21.3.3.4 Memory Address Equates
- 21.3.3.5 Memory Regions
- 21.3.3.6 Configuration Words Input/Output Section Map

21.3.3.1   Inclusion of Processor-Specific Object File(s)

This section of the processor definitions linker script ensures that the processor-specific object file(s) get included in the link.

```
/*************************************************************

 * Processor-specific object file. Contains SFR
       definitions.

*************************************************************/

INPUT("processor.o")
```

The `INPUT` line specifies that `processor.o` should be included in the link as if this file were named on the command line. The linker attempts to find this file in the current directory. If it is not found, the linker searches through the library search paths (that is, the paths specified with the `-L` command line option).

21.3.3.2   Optional Inclusion of Processor-Specific Peripheral Libraries

Note: The legacy peripheral libraries are deprecated and are replaced by the MPLAB Harmony libraries, installed separately.

This section of the processor definitions linker script ensures that the processor-specific peripheral libraries get included, but only if the files exist.

```
/****************************************************************

 * Processor-specific peripheral libraries are
       optional

****************************************************************/

OPTIONAL("libmchp_peripheral.a")

OPTIONAL("libmchp_peripheral_32MX795F512L.a")
```

The OPTIONAL lines specify that `libmchp_peripheral.a` and `libmchp_peripheral_32MX795F512L.a` should be included in the link as if the files were named on the command line. The linker attempts to find these files in the current directory. If they are not found in the current directory, the linker searches through the library search paths. If they are not found in the library search paths, the link process continues without error. The linker will error only when a symbol from the -peripheral library is required but not found elsewhere.

21.3.3.3   Base Exception Vector Address and Vector Spacing Symbols

This section of the processor definitions linker script defines values for the base exception vector address and vector spacing.

```
/*************************************************************
```

```
* For interrupt vector handling
**************************************************************/

_vector_spacing = 0x00000001;

_ebase_address = 0x9FC01000;
```

The first line defines a value of 1 for `_vector_spacing`. The available memory for exceptions only supports a vector spacing of 1. The second line defines the location of the base exception vector address (`EBASE`).

On some devices, the base exception vector address is located in the KSEG0 boot segment. On other devices, the size of the KSEG0 boot segment is not sufficient for the vector table, so the base exception vector address is located in the KSEG0 program segment. In general, devices with at least 12 KB in the KSEG0 boot segment use the boot flash for the exception vector table. Devices with less than 12 KB in the KSEG0 boot segment use the KSEG0 program segment for the exception vector table. Be sure to check the `procdefs.ld` include file for the default address for your target device.

21.3.3.4 Memory Address Equates

This section of the processor definitions linker script provides information about certain memory addresses required by the default linker script.

```
/**************************************************************

* Memory Address Equates

**************************************************************/
```

```
_RESET_ADDR              =
        0xBFC00000;
```

```
_BEV_EXCPT_ADDR              =
        0xBFC00380;
```

```
_DBG_EXCPT_ADDR              =
        0xBFC00480;
```

```
_DBG_CODE_ADDR              =
        0xBFC02000;
```

```
_GEN_EXCPT_ADDR              =
        _ebase_address + 0x180;
```

The `_RESET_ADDR` defines the processor's Reset address. This is the virtual begin address of the IFM boot section in Kernel mode.

The `_BEV_EXCPT_ADDR` defines the address that the processor jumps to when an exception is encountered and $Status_{BEV} = 1$.

The `_DBG_EXCPT_ADDR` defines the address that the processor jumps to when a debug exception is encountered.

The `_DBG_CODE_ADDR` defines the address that is the start address of the debug executive. Note that this address may vary depending on the size of the `KSEG0` boot segment on your target device.

The `_GEN_EXCPT_ADDR` defines the address that the processor jumps to when an exception is encountered and $Status_{BEV} = 0$.

21.3.3.5 Memory Regions

This section of the processor definitions linker script provides information about the memory regions that are available on the device.

```
/**************************************************************

* Memory Regions
```

```
 *

  * Memory regions without attributes cannot be
        used for

  * orphaned sections. Only sections specifically
        assigned to

  * these regions can be allocated into these
        regions.

 ************************************************************/

MEMORY

{

   kseg0_program_mem (rx) : ORIGIN = 0x9D000000,
        LENGTH = 0x8000

kseg0_boot_mem : ORIGIN = 0x9FC00490, LENGTH = 0x970

exception_mem : ORIGIN = 0x9FC01000, LENGTH = 0x1000

kseg1_boot_mem : ORIGIN = 0xBFC00000, LENGTH = 0x490

debug_exec_mem : ORIGIN = 0xBFC02000, LENGTH = 0xFF0

   config3              : ORIGIN = 0xBFC02FF0, LENGTH =
       0x4

   config2              : ORIGIN = 0xBFC02FF4, LENGTH =
       0x4

   config1              : ORIGIN = 0xBFC02FF8, LENGTH =
       0x4

   config0              : ORIGIN = 0xBFC02FFC, LENGTH =
       0x4

kseg1_data_mem : ORIGIN = 0xA0000000, LENGTH = 0x2000

   sfrs                 : ORIGIN = 0xBF800000, LENGTH =
        0x10000

}
```

Note: L1 cache devices use `kseg1_data_mem`.

Eleven memory regions are defined with an associated start address and length:

- Program memory region (`kseg0_program_mem`) for application code
- Boot memory regions (`kseg0_boot_mem` and `kseg1_boot_mem`)
- Exception memory region (`exception_mem`)
- Debug executive memory region (`debug_exec_mem`)
- Configuration memory regions (`config3`, `config2`, `config1`, and `config0`)
- Data memory region (`kseg1_data_mem`)
- SFR memory region (`sfrs`)

The default linker script uses these names to locate sections into the correct regions. Sections which are non-standard become orphaned sections. The attributes of the memory regions are used to locate these orphaned sections. The attributes (`rx`) specify that read-only sections or executable sections can be located into the program

memory regions. Similarly, the attributes (`w!x`) specify that sections that are not read-only and not executable can be located in the data memory region. Since no attributes are specified for the boot memory region, the configuration memory regions, or the SFR memory region, only specified sections may be located in these regions (that is, orphaned sections may not be located in the boot memory regions, the exception memory region, the configuration memory regions, the debug executive memory region, or the SFR memory region).

21.3.3.6 Configuration Words Input/Output Section Map

The last section in the processor definitions linker script is the input/output section map for Configuration Words. This section map is additive to the Input/Output Section Map found in the default linker script (see 21.3.4 Input/Output Section Map). It defines how input sections for Configuration Words are mapped to output sections for Configuration Words. Note that input sections are portions of an application that are defined in source code, while output sections are created by the linker. Generally, several input sections may be combined into a single output section. All output sections are specified within a `SECTIONS` command in the linker script.

For each Configuration Word that exists on the specific processor, a distinct output section named `.config_address` exists where address is the location of the Configuration Word in memory. Each of these sections contains the data created by the `#pragma config` directive (see 20.3 Pragma Directives) for that Configuration Word. Each section is assigned to their respective memory region (`con-fig`*n*).

```
SECTIONS
{
.config_BFC02FF0 : {
*(.config_BFC02FF0)
} > config3
.config_BFC02FF4 : {
*(.config_BFC02FF4)
} > config2
.config_BFC02FF8 : {
*(.config_BFC02FF8)
} > config1
.config_BFC02FFC : {
*(.config_BFC02FFC)
} > config0
}
```

21.3.4 Input/Output Section Map

The last section in the default linker script is the input/output section map. The section map is the heart of the linker script. It defines how input sections are mapped to output sections. Note that input sections are portions of an application that are defined in source code, while output sections are created by the linker. Generally, several input sections may be combined into a single output section. All output sections are specified within a `SECTIONS` command in the linker script.

21.3.4.1 .Config_<Address> Sections

These sections map the Configuration Words to their corresponding absolute addresses on the target device. The compiler's config pragma generates the input sections using this naming convention, and the linker script then maps the compiler-generated input section to the output section mapped to the corresponding absolute address.

21.3.4.2 .Reset Section

This section contains the code that is executed when the processor performs a Reset. This section is located at the Reset address (`_RESET_ADDR`), as specified in the processor definitions linker script and is assigned to the boot

memory region (`kseg1_boot_mem`). The `.reset` output section also contains the C start-up code from
the `.reset.startup` input section.

```
.reset _RESET_ADDR :
{
  KEEP(*(.reset))
  KEEP(*(.reset.startup))
} > kseg1_boot_mem
```

21.3.4.3 .bev_excpt Section

This section contains the handler for exceptions that occur when $Status_{BEV} = 1$. This section is located at the `BEV` exception address (`_BEV_EXCPT_ADDR`) as specified in the processor definitions linker script and is assigned to the boot memory region (`kseg1_boot_mem`).

```
(kseg1_boot_mem).
.bev_excpt _BEV_EXCPT_ADDR :
{
  (*(.bev_handler))
} > kseg1_boot_mem
```

21.3.4.4 .dbg_excpt Section

This section reserves space for the debug exception vector. This section is only allocated if the symbol `_DEBUGGER` has been defined. (This symbol is defined if the `-mde-bugger` command line option is specified to the shell.) This section is located at the debug exception address (`_DBG_EXCPT_ADDR`) as specified in the processor definitions linker script and is assigned to the boot memory region (`kseg1_boot_mem`). The section is marked as `NOLOAD` as it is only intended to ensure that application code cannot be placed at locations reserved for the debug executive.

```
.dbg_excpt _DBG_EXCPT_ADDR (NOLOAD) :
{
  . += (DEFINED (_DEBUGGER) ? 0x8 : 0x0);
} > kseg1_boot_mem
```

21.3.4.5 .dbg_code Section

This section reserves space for the debug exception handler. This section is only allocated if the symbol `_DEBUGGER` has been defined. (This symbol is defined if the `-mde-bugger` command line option is specified to the shell.) This section is located at the debug code address (`_DBG_CODE_ADDR`) as specified in the processor definitions linker script and is assigned to the debug executive memory region (`debug_exec_mem`). The section is marked as `NOLOAD` because it is only intended to ensure that application code cannot be placed at locations reserved for the debug executive.

```
.dbg_code _DBG_CODE_ADDR (NOLOAD) :
{
  . += (DEFINED (_DEBUGGER) ? 0xFF0 : 0x0);
} > debug_exec_mem
```

21.3.4.6 .app_excpt Section

This section contains the handler for exceptions that occur when $StatusBEV = 0$. This section is located at the general exception address (`_GEN_EXCPT_ADDR`) as specified in the processor definitions linker script and is assigned to the exception memory region (`exception_mem`).

```
.app_excpt _GEN_EXCPT_ADDR :
{
  KEEP(*(.gen_handler))
} > exception_mem
```

21.3.4.7 .vector_0 .. .vector_63 Sections (PIC32MX Interrupt Vector Tables)

PIC32MX devices use an Interrupt Vector Controller that provides 64 interrupt vectors with uniform, user-configurable spacing.

For these devices, each vector in the table is created as an output section located at an absolute address based on values of the `_ebase_address` and `_vector_spacing` symbols. There is one `.vector_n` output section for each of the 64 vectors in the table.

These sections contain the handler for each of the interrupt vectors. These sections are located at the correct vectored addresses using the formula:

```
_ebase_address + 0x200 + (_vector_spacing << 5) * n
```

where `n` is the respective vector number.

Each of the sections is followed by an assert that ensures the code located at the vector does not exceed the vector spacing specified.

```
.vector_n _ebase_address + 0x200 + (_vector_spacing << 5) * n :
  {
    KEEP(*(.vector_n))
  } > exception_mem
ASSERT (SIZEOF(.vector_n) < (_vector_spacing << 5), "function at exception vector n too
large")
```

21.3.4.8 .vectors Section

Some PIC32 families feature variable offsets for vector spacing. This feature allows the interrupt vector spacing to be configured according to application needs. A specific interrupt vector offset can be set for each vector using its associated `OFFxxx` register. For details on the interrupt vector-table variable offset feature, refer to the "*PIC32 - Family Reference Manual*" (DS61108) and also the data sheet for your specific PIC32 MCU.

The application source code is responsible for creating a `.vector_n` input section for each interrupt vector. The C/C++ compiler creates this section when either the `-vector(n)` or the `at_vector(n)` attribute is applied to the interrupt service routine. In assembly code, use the `.section` directive to create a new named section.

The device-specific linker script creates a single output section named `.vectors` that groups all of the individual `.vector_n` input sections from the project. The start of the interrupt-vector table is mapped to the address (`_ebase_address + 0x200`). The default value of the `_ebase_address` symbol is also provided in the linker script.

For each vector, the linker script also creates a symbol named `__vector_offset_n`, whose value is the offset of the vector address from the `_ebase_address` address.

```
PROVIDE(_ebase_address = 0x9D000000);
SECTIONS {
 /* Interrupt vector table with vector offsets */
 .vectors _ebase_address + 0x200 :
 {
 /* Symbol __vector_offset_n points to .vector_n if it exists,
  * otherwise points to  the default handler. The
  * vector_offset_init.o module then provides a .data section
  * containing values used to initialize the vector-offset SFRs
  * in the crt0 startup code.
  */
   __vector_offset_0 = (DEFINED(__vector_dispatch_0) ? (. - _ebase_address) :
 __vector_offset_default); KEEP(*(.vector_0))
   __vector_offset_1 = (DEFINED(__vector_dispatch_1) ? (. - _ebase_address) :
 __vector_offset_default); KEEP(*(.vector_1))
   __vector_offset_2 = (DEFINED(__vector_dispatch_2) ? (. - _ebase_address) :
 __vector_offset_default); KEEP(*(.vector_2))
 /* ... */
   __vector_offset_190 = (DEFINED(__vector_dispatch_190) ? (. - _ebase_address) :
 __vector_offset_default); KEEP(*(.vector_190))
 }
}
```

The vector-offset initialization module (`vector_offset_init.o`) uses the `__vector_offset_n` symbols defined in the default linker script. The value of each symbol is the offset of the vector's address from the ebase register's address. The vector-offset initialization module, uses the symbol value to create a `.data` section using the address of the corresponding `OFFxxx` special function register. This means that the standard linker-generated data-initialization template contains the values used to initialize the `OFFxxx` registers.

With these `.data` sections added to the project and the linker-generated data-initialization template, the standard runtime startup code initializes the `OFFxxx` special function registers as regular initialized data. No special code is required in the startup code to initialize the `OFFxxx` registers.

21.3.4.9 .text Section

The standard executable code sections are no longer mapped to the `.text` output section. However, a few special executable sections are still mapped here as shown below. This section is assigned to the program memory region (`kseg0_program_mem`) and has a fill value of `NOP (0)`.

The built-in linker script no longer maps standard `.text` executable code input sections. By not mapping these sections in the linker script, we allow these sections to be allocated using the best-fit allocator rather than the sequential allocator. Sections that are unmapped in the linker script can flow around absolute sections specified in code, whereas sections that are linker-script mapped are grouped together and allocated sequentially. This potentially causes conflicts with absolute sections (using the `address` function attribute).

```
.text ORIGIN(kseg0_program_mem) :
{
*(.stub .gnu.linkonce.t.*)
  KEEP (*(.text.*personality*))
  *(.gnu.warning)
  *(.mips16.fn.*)
  *(.mips16.call.*)
} > kseg0_program_mem =0
```

21.3.4.10 C++ Initilization Sections

The sections `.init`, `.preinit_array`, `.init_array`, `.fini_array`, `.ctors`, and `.dtors` are all used for the construction and destruction of file-scope static-storage C++ objects.

```
/* Global-namespace object initialization */
.init    :
{
  KEEP (*crti.o(.init))
  KEEP (*crtbegin.o(.init))
  KEEP (*(EXCLUDE_FILE (*crtend.o *crtend?.o *crtn.o ).init))
  KEEP (*crtend.o(.init))
  KEEP (*crtn.o(.init))
  . = ALIGN(4) ;
} >kseg0_program_mem
.fini    :
{
  KEEP (*(.fini))
  . = ALIGN(4) ;
} >kseg0_program_mem
.preinit_array    :
{
  PROVIDE_HIDDEN (__preinit_array_start = .);
  KEEP (*(.preinit_array))
  PROVIDE_HIDDEN (__preinit_array_end = .);
  . = ALIGN(4) ;
} >kseg0_program_mem
.init_array    :
{
  PROVIDE_HIDDEN (__init_array_start = .);
  KEEP (*(SORT(.init_array.*)))
  KEEP (*(.init_array))
  PROVIDE_HIDDEN (__init_array_end = .);
  . = ALIGN(4) ;
} >kseg0_program_mem
.fini_array    :
{
  PROVIDE_HIDDEN (__fini_array_start = .);
  KEEP (*(SORT(.fini_array.*)))
  KEEP (*(.fini_array))
  PROVIDE_HIDDEN (__fini_array_end = .);
  . = ALIGN(4) ;
} >kseg0_program_mem
.ctors    :
{
   /* XC32 uses crtbegin.o to find the start of
      the constructors, so we make sure it is
      first.  Because this is a wildcard, it
      doesn't matter if the user does not
      actually link against crtbegin.o; the
      linker won't look for a file to match a
      wildcard.  The wildcard also means that it
```

```
      doesn't matter which directory crtbegin.o
      is in.  */
   KEEP (*crtbegin.o(.ctors))
   KEEP (*crtbegin?.o(.ctors))
   /* We don't want to include the .ctor section from
      the crtend.o file until after the sorted ctors.
      The .ctor section from the crtend file contains the
      end of ctors marker and it must be last */
   KEEP (*(EXCLUDE_FILE (*crtend.o *crtend?.o ) .ctors))
   KEEP (*(SORT(.ctors.*)))
   KEEP (*(.ctors))
   . = ALIGN(4) ;
 } >kseg0_program_mem
 .dtors   :
 {
   KEEP (*crtbegin.o(.dtors))
   KEEP (*crtbegin?.o(.dtors))
   KEEP (*(EXCLUDE_FILE (*crtend.o *crtend?.o ) .dtors))
   KEEP (*(SORT(.dtors.*)))
   KEEP (*(.dtors))
   . = ALIGN(4) ;
 } >kseg0_program_mem
The order of the input sections within each output section is significant.
```

Note: The order of the input sections within each output section is significant.

21.3.4.11 .rodata Section

Standard read-only sections are not mapped in the linker script. A few special read-only sections are still mapped in the linker script, but most sections are unmapped, allowing them to be handled by the best fit allocator. This section is assigned to the program memory region (`kseg0_program_mem`).

```
.rodata   :
{
  *(.gnu.linkonce.r.*)
  *(.rodata1)
} > kseg0_program_mem
```

21.3.4.12 .sdata2 Section

This section collects the small initialized constant global and static data from all of the application's input files. Because of the constant nature of the data, this section is also a read-only section. This section is assigned to the program memory region (`kseg0_program_mem`).

```
/*
 * Small initialized constant global and static data can be
 * placed in the .sdata2 section. This is different from
 * .sdata, which contains small initialized non-constant
 * global and static data.
 */
.sdata2   :
{
  *(.sdata2 .sdata2.* .gnu.linkonce.s2.*)
} > kseg0_program_mem
```

21.3.4.13 .sbss2 Section

This section collects the small uninitialized constant global and static data from all of the application's input files. Because of the constant nature of the data, this section is also a read-only section. This section is assigned to the program memory region (`kseg0_program_mem`).

```
/*
 * Uninitialized constant global and static data (that is,
 * variables which will always be zero). Again, this is
 * different from .sbss, which contains small non-initialized,
 * non-constant global and static data.
 */
.sbss2   :
{
  *(.sbss2 .sbss2.* .gnu.linkonce.sb2.*)
} > kseg0_program_mem
```

21.3.4.14 .dbg_data Section

This section reserves space for the data required by the debug exception handler. This section is only allocated if the symbol _DEBUGGER has been defined. (This symbol is defined if the -mdebugger command line option is specified to the shell.) This section is assigned to the data memory region (kseg1_data_mem). The section is marked as NOLOAD as it is only intended to ensure that application data cannot be placed at locations reserved for the debug executive.

```
.dbg_data (NOLOAD) :
{
  . += (DEFINED (_DEBUGGER) ? 0x200 : 0x0);
} > kseg1_data_mem
```

21.3.4.15 .data Section

The linker generates a data-initialization template that the C start-up code uses to initialize variables.

21.3.4.16 .got Section

This section collects the global offset table from all of the application's input files. This section is assigned to the data memory region (kseg1_data_mem) with a load address located in the program memory region (kseg0_program_mem). A symbol is defined to represent the location of the Global Pointer (_gp).

```
_gp = ALIGN(16) + 0x7FF0 ;
.got    :
{
   *(.got.plt) *(.got)
} > kseg1_data_mem AT> kseg0_program_mem
```

21.3.4.17 .sdata Section

This section collects the small initialized data from all of the application's input files. This section is assigned to the data memory region (kseg1_data_mem) with a load address located in the program memory region (kseg0_program_mem). Symbols are defined to represent the virtual begin (_sdata_begin) and end (_sdata_end) addresses of this section.

```
/*
 * We want the small data sections together, so
 * single-instruction offsets can access them all, and
 * initialized data all before uninitialized, so
 * we can shorten the on-disk segment size.
 */
.sdata    :
{
  _sdata_begin = . ;
  *(.sdata .sdata.* .gnu.linkonce.s.*)
  _sdata_end = . ;
} > kseg1_data_mem AT> kseg0_program_mem
```

21.3.4.18 .lit8 Section

This section collects the 8-byte constants (usually floating-point) which the assembler decides to store in memory rather than in the instruction stream from all of the application's input files. This section is assigned to the data memory region (kseg1_data_mem) with a load address located in the program memory region (kseg0_program_mem).

```
.lit8           :
{
  *(.lit8)
} > kseg1_data_mem AT> kseg0_program_mem
```

21.3.4.19 .lit4 Section

This section collects the 4-byte constants (usually floating-point) which the assembler decides to store in memory rather than in the instruction stream from all of the application's input files. This section is assigned to the data memory region (kseg1_data_mem) with a load address located in the program memory region

(`kseg0_program_mem`). A symbol is defined to represent the virtual end address of the initialized data
(`_data_end`).

```
  .lit4           :
  {
    *(.lit4)
  } > kseg1_data_mem AT> kseg0_program_mem
  _data_end = . ;
```

21.3.4.20 .sbss Section

This section collects the small uninitialized data from all of the application's input files. This section is assigned to the
data memory region (`kseg1_data_mem`). A symbol is defined to represent the virtual begin address of uninitialized
data (`_bss_begin`). Symbols are also defined to represent the virtual begin (`_sbss_begin`) and end (`_sbss_end`)
addresses of this section.

```
  _bss_begin = . ;
  .sbss    :
  {
    _sbss_begin = . ;
    *(.dynsbss)
    *(.sbss .sbss.* .gnu.linkonce.sb.*)
    *(.scommon)
    _sbss_end = . ;
  } > kseg1_data_mem
```

21.3.4.21 .bss Section

This section collects the uninitialized data from all of the application's input files. This section is assigned to the data
memory region (`kseg1_data_mem`). A symbol is defined to represent the virtual end address of uninitialized data
(`_bss_end`). A symbol is also to represent the virtual end address of data memory (`_end`).

```
  .bss    :
  {
    *(.dynbss)
    *(.bss .bss.* .gnu.linkonce.b.*)
    *(COMMON)
    /*
     * Align here to ensure that the .bss section occupies
     * space up to _end.  Align after .bss to ensure correct
     * alignment even if the .bss section disappears because
     * there are no input sections.
     */
    . = ALIGN(32 / 8) ;
  } > kseg1_data_mem
  . = ALIGN(32 / 8) ;
  _end = . ;
  _bss_end = . ;
```

21.3.4.22 .heap Section

The linker now dynamically reserves an area of memory for the heap. The `.heap` section is no longer mapped in the
linker script. The linker finds the largest unused gap of memory after all other sections are allocated and uses that
gap for both the heap and the stack. The minimum amount of space reserved for the heap is determined by the
symbol `_min_heap_size`.

21.3.4.23 .stack Section

The linker now dynamically reserves an area of memory for the stack. The `.stack` section is no longer mapped in
the linker script. The linker finds the largest unused gap of memory after all other sections are allocated and uses that
gap for both the heap and the stack. The minimum amount of space reserved for the stack is determined by the
symbol `_min_stack_size`.

21.3.4.24 .ramfunc Section

The linker now dynamically collects the '`ramfunc`' attributed and "`.ramfunc`" named sections and allocates them
sequentially in an appropriate range of memory. The first `ramfunc` attributed function is placed at the highest
appropriately aligned address.

The presence of a `ramfunc` section causes the linker to emit the symbols necessary for the crt0.S start-up code to initialize the PIC32 bus matrix appropriately.

```
/*
 * RAM functions go at the end of our stack and heap allocation.
 * Alignment of 2K required by the boundary register (BMXDKPBA).
 *
 * RAM functions are now allocated by the linker. The linker generates
 * _ramfunc_begin and _bmxdkpba_address symbols depending on the
 * location of RAM functions.
 */

_bmxdudba_address = LENGTH(kseg1_data_mem) ;
_bmxdupba_address = LENGTH(kseg1_data_mem) ;
```

21.3.4.25 Stack Location

A symbol is defined to represent the location of the Stack Pointer (`_stack`). As described previously, the heap and the stack are now allocated to the largest available gap of memory after other sections have been allocated.

For PIC32 devices with more than 64K of data memory, GP relative addressing mode should not be used. To avoid conflict of using GP-relative addressing to the linker generated symbols, allocate the symbols in section "`_linkergenerated`": extern unsigned int __attribute__((section("_linkergenerated"))) _splim;

21.3.4.26 Debug Sections

The debug sections contain DWARF2 debugging information. They are not loaded into program Flash.

```
/* Stabs debugging sections.  */
.stab          0 : { *(.stab) }
.stabstr       0 : { *(.stabstr) }
.stab.excl     0 : { *(.stab.excl) }
.stab.exclstr  0 : { *(.stab.exclstr) }
.stab.index    0 : { *(.stab.index) }
.stab.indexstr 0 : { *(.stab.indexstr) }
.comment       0 : { *(.comment) }
/* DWARF debug sections.
   Symbols in the DWARF debugging sections are relative to the beginning
   of the section so we begin them at 0.  */
/* DWARF 1 */
.debug         0 : { *(.debug) }
.line          0 : { *(.line) }
/* GNU DWARF 1 extensions */
.debug_srcinfo  0 : { *(.debug_srcinfo) }
.debug_sfnames  0 : { *(.debug_sfnames) }
/* DWARF 1.1 and DWARF 2 */
.debug_aranges  0 : { *(.debug_aranges) }
.debug_pubnames 0 : { *(.debug_pubnames) }
/* DWARF 2 */
.debug_info     0 : { *(.debug_info .gnu.linkonce.wi.*) }
.debug_abbrev   0 : { *(.debug_abbrev) }
.debug_line     0 : { *(.debug_line) }
.debug_frame    0 : { *(.debug_frame) }
.debug_str      0 : { *(.debug_str) }
.debug_loc      0 : { *(.debug_loc) }
.debug_macinfo  0 : { *(.debug_macinfo) }
/* SGI/MIPS DWARF 2 extensions */
.debug_weaknames 0 : { *(.debug_weaknames) }
.debug_funcnames 0 : { *(.debug_funcnames) }
.debug_typenames 0 : { *(.debug_typenames) }
.debug_varnames  0 : { *(.debug_varnames) }
.debug_pubtypes 0 : { *(.debug_pubtypes) }
.debug_ranges   0 : { *(.debug_ranges) }
/DISCARD/ : { *(.rel.dyn) }
.gnu.attributes 0 : { KEEP (*(.gnu.attributes)) }
/DISCARD/ : { *(.note.GNU-stack) }
/DISCARD/ : { *(.note.GNU-stack) *(.gnu_debuglink) *(.gnu.lto_*) *(.discard) }
```

21.3.4.27 Variables Allocated to L1 Cached Memory

For devices featuring an L1 data cache, data variables are now allocated to the KSEG0 data-memory region (`kseg0_data_mem`) making it accessible through the L1 cache. Likewise, the linker-allocated heap and stack are

allocated to the KSEG0 region. The startup code initializes the L1 cache using symbols defining the base addresses in the linker script.

Example:

```
EXTERN  (__pic32_init_cache_program_base_addr)
PROVIDE (__pic32_init_cache_program_base_addr = 0x9D000000) ;
EXTERN  (__pic32_init_cache_data_base_addr)
PROVIDE (__pic32_init_cache_data_base_addr = 0x80000000) ;
```

21.3.4.28 .tlb_init_values Section

Some PIC32 devices feature a Memory Management Unit (MMU) with a Translation Lookaside Buffer (TLB). On some of these devices, the data sheet describes specific ranges of KSEG2/KSEG3 regions as dedicated to the Serial Quad Interface (SQI) and/or the External Bus Interface (EBI).

For these devices, the default startup code calls a module that initializes the TLB for this dedicated memory mapping. The TLB initialization module, `pic32_init_tlb_ebi_sqi.o`, uses the table created by this section to initialize the TLB. For more information on this format, see the copy of the `pic32_init_tlb_ebi_sqi.S` source file located in your `pic32m-libs/libpic32/stubs` directory.

22.   Embedded Compiler Compatibility Mode

All three MPLAB XC C compilers can be placed into a compatibility mode. In this mode, they are syntactically compatible with the non-standard C language extensions used by other non-Microchip embedded compiler vendors. This compatibility allows C source code written for other compilers to be compiled with minimum modification when using the MPLAB XC compilers.

Since very different device architectures may be targeted by other compilers, the semantics of the non-standard extensions may be different to that in the MPLAB XC compilers. This document indicates when the original C code may need to be reviewed.

22.1   Compiling in Compatibility Mode

An option is used to enable vendor-specific syntax compatibility. When using MPLAB XC8, this option is `--ext=vendor;` when using MPLAB XC16 or MPLAB XC32, the option is `-mext=vendor`. The argument vendor is a key that is used to represent the syntax. See the following table for a list of all keys usable with the MPLAB XC compilers.

Table 22-1. Vendor Keys

Vendor key	Syntax	XC8 Support	XC16 Support	XC32 Support
`cci`	Common Compiler Interface	Yes	Yes	Yes

The Common Compiler Interface is a language standard that is common to all Microchip MPLAB XC compilers. The non-standard extensions associated with this syntax are already described in 3. Common C Interface and are not repeated here.

22.2   Syntax Compatibility

The goal of this syntax compatibility feature is to ease the migration process when porting source code from other C compilers to the native MPLAB XC compiler syntax.

Many non-standard extensions are not required when compiling for Microchip devices and, for these, there are no equivalent extensions offered by MPLAB XC compilers. These extensions are then simply ignored by the MPLAB XC compilers, although a warning message is usually produced to ensure that you are aware of the different compiler behavior. You should confirm that your project will still operate correctly with these features disabled.

Other non-standard extensions are not compatible with Microchip devices. Errors will be generated by the MPLAB XC compiler if these extensions are not removed from the source code. You should review the ramifications of removing the extension and decide whether changes are required to other source code in your project.

The following table indicates the various levels of compatibility used in the tables that are presented throughout this guide.

Table 22-2. Level of Support Indicators

Level	Explanation
support	The syntax is accepted in the specified compatibility mode, and its meaning will mimic its meaning when it is used with the original compiler.
support (no args)	In the case of pragmas, the base pragma is supported in the specified compatibility mode, but the arguments are ignored.
native support	The syntax is equivalent to that which is already accepted by the MPLAB XC compiler, and the semantics are compatible. You can use this feature without a vendor compatibility mode having been enabled.

..........continued

Level	Explanation
ignore	The syntax is accepted in the specified compatibility mode, but the implied action is not required or performed. The extension is ignored and a warning will be issued by the compiler.
error	The syntax is not accepted in the specified compatibility mode. An error will be issued and compilation will be terminated.

Note that even if a C feature is supported by an MPLAB XC compiler, addresses, register names, assembly instructions, or any other device-specific argument is unlikely to be valid when compiling for a Microchip device. Always review code which uses these items in conjunction with the data sheet of your target Microchip device.

22.3 Data Type

Some compilers allow use of the boolean type, `bool`, as well as associated values `true` and `false`, as specified by the C99 ANSI Standard. This type and these values may be used by all MPLAB XC compilers when in compatibility mode (see note), as shown in the table below.

As indicated by the ANSI Standard, the `<stdbool.h>` header must be included for this feature to work as expected when it is used with MPLAB XC compilers.

Table 22-3. Support for C99 Bool Type

IAR Compatibility Mode			
Type	XC8	XC16	XC32
`bool`	support	support	support

Do not confuse the boolean type, `bool`, and the integer type, `bit`, implemented by MPLAB XC8.

Note: Not all C99 features have been adopted by all Microchip MPLAB XC compilers.

22.4 Operator

The @ operator may be used with other compilers to indicate the desired memory location of an object. As Table A-4 indicates, support for this syntax in MPLAB C is limited to MPLAB XC8 only.

Any address specified with another device is unlikely to be correct on a new architecture. Review the address in conjunction with the data sheet for your target Microchip device.

Using @ in a compatibility mode with MPLAB XC8 will work correctly, but will generate a warning. To prevent this warning from appearing again, use the reviewed address with the MPLAB C `__at()` specifier instead.

For MPLAB XC16/32, consider using the `address` attribute.

Table 22-4. Support for Non-Standard Operator

IAR Compatibility Mode			
Operator	XC8	XC16	XC32
@	native support	error	error

22.5 Extended Keywords

Non-standard extensions often specify how objects are defined or accessed. Keywords are usually used to indicate the feature. The non-standard C keywords corresponding to other compilers are listed in the following table, as well as the level of compatibility offered by MPLAB XC compilers. The table notes offer more information about extensions.

Table 22-5. Support for Non-Standard Keywords

IAR Compatibility Mode			
Keyword	XC8	XC16	XC32
`__section_begin`	ignore	support	support
`__section_end`	ignore	support	support
`__section_size`	ignore	support	support
`__segment_begin`	ignore	support	support
`__segment_end`	ignore	support	support
`__segment_size`	ignore	support	support
`__sfb`	ignore	support	support
`__sfe`	ignore	support	support
`__sfs`	ignore	support	support
`__asm or asm`[1]	support[2]	native support	native support
`__arm`	ignore	ignore	ignore
`__big_endian`	error	error	error
`__fiq`	support	error	error
`__intrinsic`	ignore	ignore	ignore
`__interwork`	ignore	ignore	ignore
`__irq`	support	error	error
`__little_endian`[3]	ignore	ignore	ignore
`__nested`	ignore	ignore	ignore
`__no_init`	support	support	support
`__noreturn`	ignore	support	support
`__ramfunc`	ignore	ignore	support[4]
`__packed`	ignore[5]	support	support
`__root`	ignore	support	support
`__swi`	ignore	ignore	ignore
`__task`	ignore	support	support
`__weak`	ignore	support	support
`__thumb`	ignore	ignore	ignore
`__farfunc`	ignore	ignore	ignore
`__huge`	ignore	ignore	ignore
`__nearfunc`	ignore	ignore	ignore
`__inline`	support	native support	native support

Note:

1. All assembly code specified by this construct is device-specific and will need review when porting to any Microchip device.
2. The keyword, asm, is supported natively by MPLAB XC8, but this compiler only supports the `__asm` keyword in IAR compatibility mode.
3. This is the default (and only) endianism used by all MPLAB XC compilers.
4. When used with MPLAB XC32, this must be used with the `__longcall__` macro for full compatibility.
5. Although this keyword is ignored, by default, all structures are packed when using MPLAB XC8, so there is no loss of functionality.

22.6 Intrinsic Functions

Intrinsic functions can be used to perform common tasks in the source code. The MPLAB XC compilers' support for the intrinsic functions offered by other compilers is shown in the following table.

Table 22-6. Support for Non-Standard Intrinsic Functions

IAR Compatibility Mode			
Function	XC8	XC16	XC32
`__disable_fiq`[1]	support	ignore	ignore
`__disable_interrupt`	support	support	support
`__disable_irq`[1]	support	ignore	ignore
`__enable_fiq`[1]	support	ignore	ignore
`__enable_interrupt`	support	support	support
`__enable_irq`[1]	support	ignore	ignore
`__get_interrupt_state`	ignore	support	support
`__set_interrupt_state`	ignore	support	support

Note: **1.** These intrinsic functions map to macros which disable or enable the global interrupt enable bit on 8-bit PIC® devices.

The header file `<xc.h>` must be included for supported functions to operate correctly.

22.7 Pragmas

Pragmas may be used by a compiler to control code generation. Any compiler will ignore an unknown pragma, but many pragmas implemented by another compiler have also been implemented by the MPLAB XC compilers in compatibility mode. The table below shows the pragmas and the level of support when using each of the MPLAB XC compilers.

Many of these pragmas take arguments. Even if a pragma is supported by an MPLAB XC compiler, this support may not apply to all of the pragma's arguments. This is indicated in the table.

Table 22-7. Support for Non-Standard Pragmas

IAR Compatibility Mode			
Pragma	XC8	XC16	XC32
`bitfields`	ignore	ignore	ignore
`data_alignment`	ignore	support	support

..........continued			
IAR Compatibility Mode			
Pragma	**XC8**	**XC16**	**XC32**
diag_default	ignore	ignore	ignore
diag_error	ignore	ignore	ignore
diag_remark	ignore	ignore	ignore
diag_suppress	ignore	ignore	ignore
diag_warning	ignore	ignore	ignore
include_alias	ignore	ignore	ignore
inline	support (no args)	support (no args)	support (no args)
language	ignore	ignore	ignore
location	ignore	support	support
message	support	native support	native support
object_attribute	ignore	ignore	ignore
optimize	ignore	native support	native support
pack	ignore	native support	native support
__printf_args	support	support	support
required	ignore	support	support
rtmodel	ignore	ignore	ignore
__scanf__args	ignore	support	support
section	ignore	support	support
segment	ignore	support	support
swi_number	ignore	ignore	ignore
type_attribute	ignore	ignore	ignore
weak	ignore	native support	native support

23.   Implementation-Defined Behavior

This section discusses the choices for implementation defined behavior in compiler.

23.1   Overview

ISO C requires a conforming implementation to document the choices for behaviors defined in the standard as "implementation-defined." The following sections list all such areas, the choices made for the compiler, and the corresponding section number from the ISO/IEC 9899:1999 standard.

23.2   Translation

ISO Standard:	"How a diagnostic is identified (3.10, 5.1.1.3)."
Implementation:	All output to `stderr` is a diagnostic.
ISO Standard:	"Whether each nonempty sequence of white-space characters other than new-line is retained or replaced by one space character in translation phase 3 (5.1.1.2)."
Implementation:	Each sequence of whitespace is replaced by a single character.

23.3   Environment

ISO Standard:	"The name and type of the function called at program start-up in a freestanding environment (5.1.2.1)."
Implementation:	`int main (void);`
ISO Standard:	"The effect of program termination in a freestanding environment (5.1.2.1)."
Implementation:	An infinite loop (branch to self) instruction will be executed.
ISO Standard:	"An alternative manner in which the main function may be defined (5.1.2.2.1)."
Implementation:	`int main (void);`
ISO Standard:	"The values given to the strings pointed to by the `argv` argument to `main` (5.1.2.2.1)."
Implementation:	No arguments are passed to main. Reference to `argc` or `argv` is undefined.
ISO Standard:	"What constitutes an interactive device (5.1.2.3)."
Implementation:	Application defined.
ISO Standard:	"Signals for which the equivalent of `signal(sig, SIG_IGN);` is executed at program start-up (7.14.1.1)."
Implementation:	Signals are application defined.
ISO Standard:	"The form of the status returned to the host environment to indicate unsuccessful termination when the `SIGABRT` signal is raised and not caught (7.20.4.1)."
Implementation:	The host environment is application defined.
ISO Standard:	"The forms of the status returned to the host environment by the `exit` function to report successful and unsuccessful termination (7.20.4.3)."
Implementation:	The host environment is application defined.

ISO Standard:	"The status returned to the host environment by the exit function if the value of its argument is other than zero, `EXIT_SUCCESS`, or `EXIT_FAILURE` (7.20.4.3)."
Implementation:	The host environment is application defined.
ISO Standard:	"The set of environment names and the method for altering the environment list used by the `getenv` function (7.20.4.4)."
Implementation:	The host environment is application defined.
ISO Standard:	"The manner of execution of the string by the system function (7.20.4.5)."
Implementation:	The host environment is application defined.

23.4   Identifiers

ISO Standard:	"Which additional multibyte characters may appear in identifiers and their correspondence to universal character names (6.4.2)."
Implementation:	No.
ISO Standard:	"The number of significant initial characters in an identifier (5.2.4.1, 6.4.2)."
Implementation:	All characters are significant.

23.5   Characters

ISO Standard:	"The number of bits in a byte (C90 3.4, C99 3.6)."
Implementation:	8.
ISO Standard:	"The values of the members of the execution character set (C90 and C99 5.2.1)."
ISO Standard:	"The unique value of the member of the execution character set produced for each of the standard alphabetic escape sequences (C90 and C99 5.2.2)."
Implementation:	The execution character set is ASCII.
ISO Standard:	"The value of a char object into which has been stored any character other than a member of the basic execution character set (C90 6.1.2.5, C99 6.2.5)."
Implementation:	The value of the char object is the 8-bit binary representation of the character in the source character set. That is, no translation is done.
ISO Standard:	"Which of signed char or unsigned char has the same range, representation, and behavior as "plain" char (C90 6.1.2.5, C90 6.2.1.1, C99 6.2.5, C99 6.3.1.1)."
Implementation:	By default on PIC32M, signed char is functionally equivalent to plain char.
ISO Standard:	"The mapping of members of the source character set (in character constants and string literals) to members of the execution character set (C90 6.1.3.4, C99 6.4.4.4, C90 and C99 5.1.1.2)."
Implementation:	The binary representation of the source character set is preserved to the execution character set.
ISO Standard:	"The value of an integer character constant containing more than one character or containing a character or escape sequence that does not map to a single-byte execution character (C90 6.1.3.4, C99 6.4.4.4)."

Implementation:	The compiler determines the value for a multi-character character constant one character at a time. The previous value is shifted left by eight, and the bit pattern of the next character is masked in. The final result is of type `int`. If the result is larger than can be represented by an `int`, a warning diagnostic is issued and the value truncated to `int` size.
ISO Standard:	"The value of a wide character constant containing more than one multibyte character, or containing a multibyte character or escape sequence not represented in the extended execution character set (C90 6.1.3.4, C99 6.4.4.4)."
Implementation:	See previous.
ISO Standard:	"The current locale used to convert a wide character constant consisting of a single multibyte character that maps to a member of the extended execution character set into a corresponding wide character code (C90 6.1.3.4, C99 6.4.4.4)."
Implementation:	LC_ALL
ISO Standard:	"The current locale used to convert a wide string literal into corresponding wide character codes (C90 6.1.4, C99 6.4.5)."
Implementation:	LC_ALL
ISO Standard:	"The value of a string literal containing a multibyte character or escape sequence not represented in the execution character set (C90 6.1.4, C99 6.4.5)."
Implementation:	The binary representation of the characters is preserved from the source character set.

23.6   Integers

ISO Standard:	"Any extended integer types that exist in the implementation (C99 6.2.5)."
Implementation:	There are no extended integer types.
ISO Standard:	"Whether signed integer types are represented using sign and magnitude, two's complement, or one's complement and whether the extraordinary value is a trap representation or an ordinary value (C99 6.2.6.2)."
Implementation:	All integer types are represented as two's complement, and all bit patterns are ordinary values.

23.7   Floating-Point

ISO Standard:	"The accuracy of the floating-point operations and of the library functions in <math.h> and <complex.h> that return floating-point results (C90 and C99 5.2.4.2.2)."
Implementation:	The accuracy is unknown.
ISO Standard:	"The accuracy of the conversions between floating-point internal representations and string representations performed by the library functions in <stdio.h>, <stdlib.h>, and <wchar.h> (C90 and C99 5.2.4.2.2)."
Implementation:	The accuracy is unknown.
ISO Standard:	"The rounding behaviors characterized by non-standard values of FLT_ROUNDS (C90 and C99 5.2.4.2.2)."

Implementation:	No such values are used.
ISO Standard:	"The evaluation methods characterized by non-standard negative values of FLT_EVAL_METHOD (C90 and C99 5.2.4.2.2)."
Implementation:	No such values are used.
ISO Standard:	"The direction of rounding when an integer is converted to a floating-point number that cannot exactly represent the original value (C90 6.2.1.3, C99 6.3.1.4)."
Implementation:	C99 Annex F is followed.
ISO Standard:	"The direction of rounding when a floating-point number is converted to a narrower floating-point number (C90 6.2.1.4, 6.3.1.5)."
Implementation:	C99 Annex F is followed.
ISO Standard:	"How the nearest representable value or the larger or smaller representable value immediately adjacent to the nearest representable value is chosen for certain floating constants (C90 6.1.3.1, C99 6.4.4.2)."
Implementation:	C99 Annex F is followed.
ISO Standard:	"Whether and how floating expressions are contracted when not disallowed by the FP_CONTRACT pragma (C99 6.5)."
Implementation:	The pragma is not implemented.
ISO Standard:	"The default state for the FENV_ACCESS pragma (C99 7.6.1)."
Implementation:	This pragma is not implemented.
ISO Standard:	"Additional floating-point exceptions, rounding modes, environments, and classifications, and their macro names (C99 7.6, 7.12)."
Implementation:	None supported.
ISO Standard:	"The default state for the FP_CONTRACT pragma (C99 7.12.2)."
Implementation:	This pragma is not implemented.
ISO Standard:	"Whether the "inexact" floating-point exception can be raised when the rounded result actually does equal the mathematical result in an IEC 60559 conformant implementation (C99 F.9)."
Implementation:	Unknown.
ISO Standard:	"Whether the "underflow" (and "inexact") floating-point exception can be raised when a result is tiny but not inexact in an IEC 60559 conformant implementation (C99 F.9)."
Implementation:	Unknown.

23.8   Arrays and Pointers

ISO Standard:	"The result of converting a pointer to an integer or vice versa (C90 6.3.4, C99 6.3.2.3)."

Implementation:	A cast from an integer to a pointer or vice versa results uses the binary representation of the source type, reinterpreted as appropriate for the destination type. If the source type is larger than the destination type, the Most Significant bits are discarded. When casting from a pointer to an integer, if the source type is smaller than the destination type, the result is sign extended. When casting from an integer to a pointer, if the source type is smaller than the destination type, the result is extended based on the signedness of the source type.
ISO Standard:	"The size of the result of subtracting two pointers to elements of the same array (C90 6.3.6, C99 6.5.6)."
Implementation:	32-bit signed integer.

23.9 Hints

ISO Standard:	"The extent to which suggestions made by using the register storage-class specifier are effective (C90 6.5.1, C99 6.7.1)."
Implementation:	The register storage class specifier generally has no effect.
ISO Standard:	"The extent to which suggestions made by using the inline function specifier are effective (C99 6.7.4)."
Implementation:	If `-fno-inline` or `-O0` are specified, no functions will be inlined, even if specified with the inline specifier. Otherwise, the function may or may not be `inlined` dependent on the optimization heuristics of the compiler.

23.10 Structures, Unions, Enumerations, and Bit Fields

ISO Standard:	"A member of a union object is accessed using a member of a different type (C90 6.3.2.3)."
Implementation:	The corresponding bytes of the union object are interpreted as an object of the type of the member being accessed without regard for alignment or other possible invalid conditions.
ISO Standard:	"Whether a "plain" `int` bit field is treated as a `signed int` bit field or as an `unsigned int` bit field (C90 6.5.2, C90 6.5.2.1, C99 6.7.2, C99 6.7.2.1)."
Implementation:	By default on PIC32M, a plain `int` bit field is treated as a signed integer. This behavior can be altered by use of the `-funsigned-bitfields` command line option.
ISO Standard:	"Allowable bit field types other than `_Bool`, `signed int`, and `unsigned int` (C99 6.7.2.1)."
Implementation:	No other types are supported.
ISO Standard:	"Whether a bit field can straddle a storage unit boundary (C90 6.5.2.1, C99 6.7.2.1)."
Implementation:	No.
ISO Standard:	"The order of allocation of bit fields within a unit (C90 6.5.2.1, C99 6.7.2.1)."
Implementation:	Bit fields are allocated left to right.
ISO Standard:	"The alignment of non-bit field members of structures (C90 6.5.2.1, C99 6.7.2.1)."

Implementation:	Each member is located to the lowest available offset allowable according to the alignment restrictions of the member type.
ISO Standard:	"The integer type compatible with each enumerated type (C90 6.5.2.2, C99 6.7.2.2)."
Implementation:	If the enumeration values are all non-negative, the type is `unsigned int`, else it is `int`. The `-fshort-enums` command line option can change this.

23.11 Qualifiers

ISO Standard:	"What constitutes an access to an object that has volatile-qualified type (C90 6.5.3, C99 6.7.3)."
Implementation:	Any expression which uses the value of or stores a value to a volatile object is considered an access to that object. There is no guarantee that such an access is atomic.If an expression contains a reference to a volatile object but neither uses the value nor stores to the object, the expression is considered an access to the volatile object or not depending on the type of the object. If the object is of scalar type, an aggregate type with a single member of scalar type, or a union with members of (only) scalar type, the expression is considered an access to the volatile object. Otherwise, the expression is evaluated for its side effects but is not considered an access to the volatile object. For example: `volatile int a;` `a; /* access to 'a' since 'a' is scalar */`

23.12 Declarators

ISO Standard:	"The maximum number of declarators that may modify an arithmetic, structure or union type (C90 6.5.4)."
Implementation:	No limit.

23.13 Statements

ISO Standard:	"The maximum number of case values in a switch statement (C90 6.6.4.2)."
Implementation:	No limit.

23.14 Pre-Processing Directives

ISO Standard:	"How sequences in both forms of header names are mapped to headers or external source file names (C90 6.1.7, C99 6.4.7)."
Implementation:	The character sequence between the delimiters is considered to be a string which is a file name for the host environment.
ISO Standard:	"Whether the value of a character constant in a constant expression that controls conditional inclusion matches the value of the same character constant in the execution character set (C90 6.8.1, C99 6.10.1)."
Implementation:	Yes.

ISO Standard:	"Whether the value of a single-character `character` constant in a constant expression that controls conditional inclusion may have a negative value (C90 6.8.1, C99 6.10.1)."
Implementation:	Yes.
ISO Standard:	"The places that are searched for an included `<  >` delimited header, and how the places are specified or the header is identified (C90 6.8.2, C99 6.10.2)."
Implementation:	`<install directory>/lib/gcc/pic32mx/3.4.4/include<install directory>/pic32mx/include`
ISO Standard:	"How the named source file is searched for in an included `" "` delimited header (C90 6.8.2, C99 6.10.2)."
Implementation:	The compiler first searches for the named file in the directory containing the including file, the directories specified by the `-iquote` command line option (if any), then the directories which are searched for a `<  >` delimited header.
ISO Standard:	"The method by which preprocessing tokens are combined into a header name (C90 6.8.2, C99 6.10.2)."
Implementation:	All tokens, including whitespace, are considered part of the header file name. Macro expansion is not performed on tokens inside the delimiters.
ISO Standard:	"The nesting limit for `#include` processing (C90 6.8.2, C99 6.10.2)."
Implementation:	No limit.
ISO Standard:	"The behavior on each recognized non-`STDC` `#pragma` directive (C90 6.8.6, C99 6.10.6)."
Implementation:	See 9.11 Variable Attributes.
ISO Standard:	"The definitions for `__DATE__` and `__TIME__` when respectively, the date and time of translation are not available (C90 6.8.8, C99 6.10.8)."
Implementation:	The date and time of translation are always available.

23.15   Library Functions

ISO Standard:	"The Null Pointer constant to which the macro NULL expands (C90 7.1.6, C99 7.17)."
Implementation:	`(void *)0`
ISO Standard:	"Any library facilities available to a freestanding program, other than the minimal set required by clause 4 (5.1.2.1)."
Implementation:	See the *32-Bit Language Tools Libraries* (DS51685).
ISO Standard:	"The format of the diagnostic printed by the `assert` macro (7.2.1.1)."
Implementation:	"Failed assertion '*message*' at line *line* of '*filename*'.\n"
ISO Standard:	"The default state for the `FENV_ACCESS` pragma (7.6.1)."
Implementation:	Unimplemented.
ISO Standard:	"The representation of floating-point exception flags stored by the `fegetexceptflag` function (7.6.2.2)."
Implementation:	Unimplemented.
ISO Standard:	"Whether the `feraiseexcept` function raises the inexact exception in addition to the overflow or underflow exception (7.6.2.3)."

Implementation:	Unimplemented.
ISO Standard:	"Floating environment macros other than FE_DFL_ENV that can be used as the argument to the fesetenv or feupdateenv function (7.6.4.3, 7.6.4.4)."
Implementation:	Unimplemented.
ISO Standard:	"Strings other than "C" and "" that may be passed as the second argument to the setlocale function (7.11.1.1)."
Implementation:	None.
ISO Standard:	"The types defined for float_t and double_t when the value of the FLT_EVAL_METHOD macro is less than 0 or greater than 2 (7.12)."
Implementation:	Unimplemented.
ISO Standard:	"The infinity to which the INFINITY macro expands, if any (7.12)."
Implementation:	Unimplemented.
ISO Standard:	"The quiet NaN to which the NAN macro expands, when it is defined (7.12)."
Implementation:	Unimplemented.
ISO Standard:	"Domain errors for the mathematics functions, other than those required by this International Standard (7.12.1)."
Implementation:	None.
ISO Standard:	"The values returned by the mathematics functions, and whether errno is set to the value of the macro EDOM, on domain errors (7.12.1)."
Implementation:	errno is set to EDOM on domain errors.
ISO Standard:	"Whether the mathematics functions set errno to the value of the macro ERANGE on overflow and/or underflow range errors (7.12.1)."
Implementation:	Yes.
ISO Standard:	"The default state for the FP_CONTRACT pragma (7.12.2)
Implementation:	Unimplemented.
ISO Standard:	"Whether a domain error occurs or zero is returned when the fmod function has a second argument of zero (7.12.10.1)."
Implementation:	NaN is returned.
ISO Standard:	"The base-2 logarithm of the modulus used by the remquo function in reducing the quotient (7.12.10.3)."
Implementation:	Unimplemented.
ISO Standard:	"The set of signals, their semantics, and their default handling (7.14)."
Implementation:	The default handling of signals is to always return failure. Actual signal handling is application defined.
ISO Standard:	"If the equivalent of signal(sig, SIG_DFL); is not executed prior to the call of a signal handler, the blocking of the signal that is performed (7.14.1.1)."
Implementation:	Application defined.
ISO Standard:	"Whether the equivalent of signal(sig, SIG_DFL); is executed prior to the call of a signal handler for the signal SIGILL (7.14.1.1)."
Implementation:	Application defined.

ISO Standard:	"Signal values other than `SIGFPE`, `SIGILL`, and `SIGSEGV` that correspond to a computational exception (7.14.1.1)."
Implementation:	Application defined.
ISO Standard:	"Whether the last line of a text stream requires a terminating new-line character (7.19.2)."
Implementation:	Yes.
ISO Standard:	"Whether space characters that are written out to a text stream immediately before a new-line character appear when read in (7.19.2)."
Implementation:	Yes.
ISO Standard:	"The number of null characters that may be appended to data written to a binary stream (7.19.2)."
Implementation:	No null characters are appended to a binary stream.
ISO Standard:	"Whether the file position indicator of an append-mode stream is initially positioned at the beginning or end of the file (7.19.3)."
Implementation:	Application defined. The system level function `open` is called with the `O_APPEND` flag.
ISO Standard:	"Whether a write on a text stream causes the associated file to be truncated beyond that point (7.19.3)."
Implementation:	Application defined.
ISO Standard:	"The characteristics of file buffering (7.19.3)."
ISO Standard:	"Whether a zero-length file actually exists (7.19.3)."
Implementation:	Application defined.
ISO Standard:	"The rules for composing valid file names (7.19.3)."
Implementation:	Application defined.
ISO Standard:	"Whether the same file can be open multiple times (7.19.3)."
Implementation:	Application defined.
ISO Standard:	"The nature and choice of encodings used for multibyte characters in files (7.19.3)."
Implementation:	Encodings are the same for each file.
ISO Standard:	"The effect of the `remove` function on an open file (7.19.4.1)."
Implementation:	Application defined. The system function `unlink` is called.
ISO Standard:	"The effect if a file with the new name exists prior to a call to the `rename` function (7.19.4.2)."
Implementation:	Application defined. The system function `link` is called to create the new file name, then `unlink` is called to remove the old file name. Typically, `link` will fail if the new file name already exists.
ISO Standard:	"Whether an open temporary file is removed upon abnormal program termination (7.19.4.3)."
Implementation:	No.
ISO Standard:	"What happens when the `tmpnam` function is called more than `TMP_MAX` times (7.19.4.4)."
Implementation:	Temporary names will wrap around and be reused.

ISO Standard:	"Which changes of mode are permitted (if any), and under what circumstances (7.19.5.4)."
Implementation:	The file is closed via the system level `close` function and re-opened with the `open` function with the new mode. No additional restriction beyond those of the application defined `open` and `close` functions are imposed.
ISO Standard:	"The style used to print an infinity or NaN, and the meaning of the *n-char-sequence* if that style is printed for a NaN (7.19.6.1, 7.24.2.1)."
Implementation:	No char sequence is printed.NaN is printed as "NaN".Infinity is printed as "[-/+]Inf".
ISO Standard:	"The output for `%p` conversion in the `fprintf` or `fwprintf` function (7.19.6.1, 7.24.2.1)."
Implementation:	Functionally equivalent to `%x`.
ISO Standard:	"The interpretation of a – character that is neither the first nor the last character, nor the second where a ^ character is the first, in the scanlist for `%[` conversion in the `fscanf` or `fwscanf` function (7.19.6.2, 7.24.2.1)."
Implementation:	Unknown
ISO Standard:	"The set of sequences matched by the `%p` conversion in the `fscanf` or `fwscanf` function (7.19.6.2, 7.24.2.2)."
Implementation:	The same set of sequences matched by `%x`.
ISO Standard:	"The interpretation of the input item corresponding to a `%p` conversion in the `fscanf` or `fwscanf` function (7.19.6.2, 7.24.2.2)."
Implementation:	If the result is not a valid pointer, the behavior is undefined.
ISO Standard:	"The value to which the macro `errno` is set by the `fgetpos`, `fsetpos`, or `ftell` functions on failure (7.19.9.1, 7.19.9.3, 7.19.9.4)."
Implementation:	If the result exceeds `LONG_MAX`, `errno` is set to `ERANGE`.Other errors are application defined according to the application definition of the `lseek` function.
ISO Standard:	"The meaning of the *n-char-sequence* in a string converted by the `strtod`, `strtof`, `strtold`, `wcstod`, `wcstof`, or `wcstold` function (7.20.1.3, 7.24.4.1.1)."
Implementation:	No meaning is attached to the sequence.
ISO Standard:	"Whether or not the `strtod`, `strtof`, `strtold`, `wcstod`, `wcstof`, or `wcstold` function sets `errno` to `ERANGE` when underflow occurs (7.20.1.3, 7.24.4.1.1)."
Implementation:	Yes.
ISO Standard:	"Whether the `calloc`, `malloc`, and `realloc` functions return a Null Pointer or a pointer to an allocated object when the size requested is zero (7.20.3)."
Implementation:	A pointer to a statically allocated object is returned.
ISO Standard:	"Whether open output streams are flushed, open streams are closed, or temporary files are removed when the `abort` function is called (7.20.4.1)."
Implementation:	No.
ISO Standard:	"The termination status returned to the host environment by the `abort` function (7.20.4.1)."
Implementation:	By default, there is no host environment.

ISO Standard:	"The value returned by the `system` function when its argument is not a Null Pointer (7.20.4.5)."
Implementation:	Application defined.
ISO Standard:	"The local time zone and Daylight Saving Time (7.23.1)."
Implementation:	Application defined.
ISO Standard:	"The era for the `clock` function (7.23.2.1)."
Implementation:	Application defined.
ISO Standard:	"The positive value for `tm_isdst` in a normalized `tmx` structure (7.23.2.6)."
Implementation:	1
ISO Standard:	"The replacement string for the `%Z` specifier to the `strftime`, `strfxtime`, `wcsftime`, and `wcsfxtime` functions in the "`C`" locale (7.23.3.5, 7.23.3.6, 7.24.5.1, 7.24.5.2)."
Implementation:	Unimplemented.
ISO Standard:	"Whether or when the trigonometric, hyperbolic, base-*e* exponential, base-*e* logarithmic, error, and log gamma functions raise the inexact exception in an IEC 60559 conformant implementation (F.9)."
Implementation:	No.
ISO Standard:	"Whether the inexact exception may be raised when the rounded result actually does equal the mathematical result in an IEC 60559 conformant implementation (F.9)."
Implementation:	No.
ISO Standard:	"Whether the underflow (and inexact) exception may be raised when a result is tiny but not inexact in an IEC 60559 conformant implementation (F.9)."
Implementation:	No.
ISO Standard:	"Whether the functions honor the Rounding Direction mode (F.9)."
Implementation:	The Rounding mode is not forced.

23.16 Architecture

ISO Standard:	"The values or expressions assigned to the macros specified in the headers `<float.h>`, `<limits.h>`, and `<stdint.h>` (C90 and C99 5.2.4.2, C99 7.18.2, 7.18.3)."
Implementation:	See 9.3.2 limits.h.
ISO Standard:	"The number, order, and encoding of bytes in any object (when not explicitly specified in the standard) (C99 6.2.6.1)."
Implementation:	Little endian, populated from Least Significant Byte first. See 9.2 Data Representation.
ISO Standard:	"The value of the result of the `size of` operator (C90 6.3.3.4, C99 6.5.3.4)."
Implementation:	See 9.2 Data Representation.

24. Deprecated Features

The features described below are considered to be obsolete and have been replaced with more advanced functionality. Projects which depend on deprecated features will work properly with versions of the language tools cited. The use of a deprecated feature will result in a warning; programmers are encouraged to revise their projects in order to eliminate any dependency on deprecated features. Support for these features may be removed entirely in future versions of the language tools.

24.1 Variables in Specified Registers

The compiler allows you to put a few global variables into specified hardware registers.

Note: Using too many registers may impair the ability of the 32-bit compiler to compile. It is not recommended that registers be placed into fixed registers.

You can also specify the register in which an ordinary register variable should be allocate.

- Global register variables reserve registers throughout the program. This may be useful in programs such as programming language interpreters which have a couple of global variables that are accessed very often.
- Local register variables in specific registers do not reserve the registers. The compiler's data flow analysis is capable of determining where the specified registers contain live values, and where they are available for other uses. Stores into local register variables may be deleted when they appear to be unused. References to local register variables may be deleted, moved or simplified.

These local variables are sometimes convenient for use with the extended inline assembly (see 18. Mixing C/C++ and Assembly Language), if you want to write one output of the assembler instruction directly into a particular register. (This will work provided the register you specify fits the constraints specified for that operand in the inline assembly statement).

24.2 Defining Global Register Variables

You can define a global register variable like this:

```
register int *foo asm ("t0");
```

Here `t0` is the name of the register which should be used. Choose a register that is normally saved and restored by function calls, so that library routines will not clobber it.

Defining a global register variable in a certain register reserves that register entirely for this use, at least within the current compilation. The register will not be allocated for any other purpose in the functions in the current compilation. The register will not be saved and restored by these functions. Stores into this register are never deleted even if they would appear to be dead, but references may be deleted, moved or simplified.

It is not safe to access the global register variables from signal handlers, or from more than one thread of control, because the system library routines may temporarily use the register for other things (unless you recompile them especially for the task at hand).

It is not safe for one function that uses a global register variable to call another such function `foo` by way of a third function `lose` that was compiled without knowledge of this variable (that is, in a source file in which the variable wasn't declared). This is because `lose` might save the register and put some other value there. For example, you can't expect a global register variable to be available in the comparison-function that you pass to `qsort`, since `qsort` might have put something else in that register. This problem can be avoided by recompiling `qsort` with the same global register variable definition.

If you want to recompile `qsort` or other source files that do not actually use your global register variable, so that they will not use that register for any other purpose, then it suffices to specify the compiler command-line option -ffixed-reg. You need not actually add a global register declaration to their source code.

A function that can alter the value of a global register variable cannot safely be called from a function compiled without this variable, because it could clobber the value the caller expects to find there on return. Therefore, the

function that is the entry point into the part of the program that uses the global register variable must explicitly save and restore the value that belongs to its caller.

The library function `longjmp` will restore to each global register variable the value it had at the time of the `setjmp`.

All global register variable declarations must precede all function definitions. If such a declaration appears after function definitions, the register may be used for other purposes in the preceding functions.

Global register variables may not have initial values because an executable file has no means to supply initial contents for a register.

24.3   Specifying Registers for Local Variables

You can define a local register variable with a specified register like this:

```
register int *foo asm ("t0");
```

Here `t0` is the name of the register that should be used. Note that this is the same syntax used for defining global register variables, but for a local variable it would appear within a function.

Defining such a register variable does not reserve the register; it remains available for other uses in places where flow control determines the variable's value is not live. Using this feature may leave the compiler too few available registers to compile certain functions.

This option does not ensure that the compiler will generate code that has this variable in the register you specify at all times. You may not code an explicit reference to this register in an `asm` statement and assume it will always refer to this variable.

Assignments to local register variables may be deleted when they appear to be unused. References to local register variables may be deleted, moved or simplified.

25.   Built-In Functions

This appendix lists the built-in functions that are specific to MPLAB XC32 C/C++ Compiler.

Built-in functions give the C programmer access to assembler operators or machine instructions that are currently only accessible using inline assembly, but are sufficiently useful that they are applicable to a broad range of applications. Built-in functions are coded in C source files syntactically like function calls, but they are compiled to assembly code that directly implements the function, and do not involve function calls or library routines.

There are a number of reasons why providing built-in functions is preferable to requiring programmers to use inline assembly. They include the following:

Providing built-in functions for specific purposes simplifies coding.

Certain optimizations are disabled when inline assembly is used. This is not the case for built-in functions.

For machine instructions that use dedicated registers, coding inline assembly while avoiding register allocation errors can require considerable care. The built-in functions make this process simpler as you do not need to be concerned with the particular register requirements for each individual machine instruction.

26. Built-In Function Descriptions

This section describes the programmer interface to the compiler built-in functions. Since the functions are "built in", there are no header files associated with them. Similarly, there are no command-line switches associated with the built-in functions – they are always available. The built-in function names are chosen such that they belong to the compiler's namespace (they all have the prefix `__builtin_`), so they will not -conflict with function or variable names in the programmer's namespace.

Built-In Function List

- `__builtin_bcc0(rn,sel,clr)`
- `__builtin_bsc0(rn,sel,set)`
- `__builtin_bcsc0(rn,sel,clr,set)`
- `__builtin_clz(x)`
- `__builtin_ctz(x)`
- `__builtin_mips_cache(op,addr)`
- `__builtin_mxc0(rn,sel,val)`
- `__builtin_set_isr_state(unsigned int)`
- `__builtin_software_breakpoint(void)`
- `unsigned long __builtin_section_begin(quoted-section-name)`
- `unsigned long __builtin_section_end(quoted-section-name)`
- `unsigned long __builtin_section_size(quoted-section-name)`
- `unsigned int __builtin_get_isr_state(void)`

26.1 __builtin_bcc0(rn,sel,clr)

Description

Clear the non-zero bits from `CLR` in CP0 register `REG,SEL`.

Prototype

```
unsigned int __builtin_bcc0(rn,sel,clr);
```

Argument

```
rn : cp0 register number
sel : cp0 select number
clr : 32-bit mask to clear
```

Return Value

```
unsigned int : new value
```

Assembler Operator/Machine Instruction

```
mfc0 t, cs
nor d, s, t
and d, s, j
mtc0 s, cd
```

Error Messages

None.

26.2 __builtin_bsc0(rn,sel,set)

Description

Set the non-zero bits from set in CP0 register `rn,sel`.

Prototype

```
unsigned int __builtin_bsc0(rn,sel,set);
```

Argument

```
rn : cp0 register number
sel : cp0 select number
set : 32-bit mask to set
```

Return Value

```
unsigned int : new value
```

Assembler Operator/ Machine Instruction

```
mfc0 t, cs
or d, s, t
mtc0 s, cd
```

Error Messages

None.

26.3   __builtin_bcsc0(rn,sel,clr,set)

Description

Clear the non-zero bits from CLR and set non-zero bits and set non-zero bits from set in CP0 register rn,sel.

Prototype

```
unsigned int __builtin_bcsc0(rn,sel,clr,set);
```

Argument

```
rn : cp0 register number
sel : cp0 select number
clr : 32-bit mask to clear
set : 32-bit mask to set
```

Return Value

```
unsigned int : new value
```

Assembler Operator/ Machine Instruction

```
mfc0 t, cs
nor d, s, t
and d, s, j
or d, s, t
mtc0 s, cd
```

Error Messages

None.

26.4   __builtin_clz(x)

Description

Count leading (high-order) zero bits in x considered as a 32-bit word.

Prototype

```
unsigned int __builtin_clz(x);
```

Argument

```
x : 32-bit word to count
```

Return Value

```
unsigned int : number of leading zero bits
```

Assembler Operator/ Machine Instruction

```
clz d, s
```

Error Messages

None.

26.5 __builtin_ctz(x)

Description

Count trailing (low-order) zero bits in x considered as a 32-bit word.

Prototype

```
unsigned int __builtin_ctz(x);
```

Argument

```
x : 32-bit word to count
```

Return Value

```
unsigned int : number of trailing zero bits
```

Assembler Operator/ Machine Instruction

```
subu d, s, t
and d, s, t
clz d, s
li d, j
subu d, s, j
```

Error Messages

None.

26.6 __builtin_mips_cache(op,addr)

Description

Perform an operation on a cache line indicated by address. See the PIC32 Family Reference Manual for valid operation values.

Prototype

```
void __builtin_mips_cache(op,addr);
```

Argument

None.

Return Value

```
void
```

Assembler Operator/ Machine Instruction

```
cache op, addr
```

Error Message

None.

26.7 __builtin_mxc0(rn,sel,val)

Description

Exchange (Swap) `val` and CP0 register `rn,sel`.

Prototype

```
unsigned int __builtin_mxc0(rn,sel,val);
```

Argument

```
rn : cp0 register number
sel : cp0 select number
val : 32-bit value to move to the register
```

Return Value

```
unsigned int : previous value in specified CP0 register
```

Assembler Operator/ Machine Instruction

```
mfc0 t, cs
mtc0 s, cd
```

Error Messages

None.

26.8 __builtin_set_isr_state(unsigned int)

Description

Set the Interrupt Priority Level and Interrupt Enable bits using a value obtained from
`__builtin_get_isr_state()`.

Prototype

```
void __builtin_set_isr_state(unsigned int);
```

Argument

An unsigned integer value obtained from `__builtin_get_isr_state()`.

Return Value

None.

Assembler Operator/ Machine Instruction

```
di
ehb
mfc0   $2, $12, 0
ins    $2,$3,10,3
srl    $3,$3,3
ins    $2,$3,0,1
mtc0   $2, $12, 0
ehb
```

Error Messages

None.

26.9 __builtin_software_breakpoint(void)

Description

Insert a software breakpoint.

Note that the `__conditional_software_breakpoint()` macro defined in assert.h provides a lightweight variant of `assert(exp)` that causes only a software breakpoint when the assertion fails rather than printing a message. This macro is disabled if, at the moment of including `<assert.h>`, a macro with the name `NDEBUG` has already been defined of if a macro with the name `__DEBUG` has not been defined. For example:

```
__conditional_software_breakpoint(myPtr!=NULL);
```

Prototype

```
void__builtin_software_breakpoint(void)
```

Argument

None.

Return Value

None.

Assembler Operator/ Machine Instruction

```
sdbbp 0
```

Error Messages

None.

26.10 unsigned long__builtin_section_begin(quoted-section-name)

Remember: This builtin gets run-time information about section addresses and sizes.

Description

Return the beginning address of the quoted section name.

Prototype

```
unsigned long __builtin_section_begin(quoted-section-name);
```

Argument

`quoted-section-name` The name of the section.

Return Value

The address of the section.

Assembler Operator/ Machine Instruction

```
.startof.
```

Error Messages

An "undefined reference" error message will be displayed if the quoted section name does not exist in the link.

26.11 unsigned long __builtin_section_end(quoted-section-name)

Remember: This builtin gets run-time information about section addresses and sizes.

Description

Return the end address of the quoted section name + 1.

Prototype

```
unsigned long __builtin_section_end(quoted-section-name);
```

Argument

`quoted-section-name` The name of the section.

Return Value

The end address of the section + 1.

Assembler Operator/ Machine Instruction

`.endof.`

Error Messages

An "undefined reference" error message will be displayed if the quoted section name does not exist in the link.

26.12 unsigned long __builtin_section_size(quoted-section-name)

👉 **Remember:** This builtin gets run-time information about section addresses and sizes.

Description

Return the size in bytes of the named quoted section.

Prototype

```
unsigned long __builtin_section_size(quoted-section-name);
```

Argument

`quoted-section-name` The name of the section.

Return Value

The size in bytes of the named section.

Assembler Operator/ Machine Instruction

`.sizeof.`

Error Messages

An "undefined reference" error message will be displayed if the quoted section name does not exist in the link.

26.13 unsigned int __builtin_get_isr_state(void)

👉 **Remember:** This builtin inspects or manipulates the current CPU interrupt state.

Description

Get the current Interrupt Priority Level and Interrupt Enable bits.

Prototype

```
unsigned int __builtin_get_isr_state(void);
```

Argument

None.

Return Value

The current IPL and interrupt enable bits in a packed format. This value is to be used with the __builtin_set_isr_state() function.

Assembler Operator/ Machine Instruction

```
mfc0    $3, $12, 0
srl     $2,$3,10
ins     $2,$3,3,1
andi    $2,$2,0xf
sw      $2,0($fp)
```

Error Messages

None.

27. Built-In DSP Functions

Many PIC32 MCUs support a DSP engine, including instructions that are designed to improve the performance of DSP and media applications. The DSPr2 engine provides instructions that operate on packed 8-bit/16-bit integer data, Q7, Q15 and Q31 fractional data.

The XC32 C compiler supports these DSP operations using both the generic vector extensions and a collection of built-in functions. Both kinds of support are enabled automatically when you select a DSP device with the -mprocessor option.

The SCOUNT and POS bits of the DSP control register are global. The WRDSP, EXTPDP, EXTPDPV and MTHLIP instructions modify the SCOUNT and POS bits. During optimization, the compiler will not delete these instructions and it will not delete calls to functions containing these instructions.

At present, the XC32 C compiler provides support for only operations on 32-bit vectors. The vector type associated with 8-bit integer data is usually called v4i8, the vector type associated with Q7 is usually called v4q7, the vector type associated with 16-bit integer data is usually called v2i16, and the vector type associated with Q15 is usually called v2q15. They can be defined in C as follows:

```
typedef signed char v4i8 __attribute__ ((vector_size(4)));
typedef signed char v4q7 __attribute__ ((vector_size(4)));
typedef short v2i16 __attribute__ ((vector_size(4)));
typedef short v2q15 __attribute__ ((vector_size(4)));
```

The v4i8, v4q7, v2i16 and v2q15 values are initialized in the same way as aggregates. For example:

```
v4i8 a = {1, 2, 3, 4};
v4i8 b;
b = (v4i8) {5, 6, 7, 8};

v2q15 c = {0x0fcb, 0x3a75};
v2q15 d;
d = (v2q15) {0.1234 * 0x1.0p15, 0.4567 * 0x1.0p15};
```

Note:

1. The first value is the least significant and the last value is the most significant. The code above will set the lowest byte of a to 1.

2. Q7, Q15 and Q31 values must be initialized with their integer representation. As shown in this example, the integer representation of a Q7 value can be obtained by multiplying the fractional value by 0x1.0p7. The equivalent for Q15 values is to multiply by 0x1.0p15. The equivalent for Q31 values is to multiply by 0x1.0p31.

The table below lists the v4i8 and v2q15 operations for which hardware support exists. The a and b are v4i8 values, and c and d are v2q15 values.

C Code	Instruction
a + b	addu.qb
c + d	addq.ph
a - b	subu.qb
c - d	subq.ph

The table below lists the v2i16 operation for which hardware support exists. The e and f are v2i16 values.

C Code	Instruction
e * f	mul.ph

It is easier to describe the DSP built-in functions if the types are defined beforehand.

```
typedef int q31;
typedef int i32;
```

```
typedef unsigned int ui32;
typedef long long a64;
```

The q31 and i32 operations are actually the same as `int`, but q31 is used to indicate a Q31 fractional value and i32 to indicate a 32-bit integer value. Similarly, a64 is the same as `long long`, but a64 is used to indicate values that will be placed in one of the four DSP accumulators ($ac0, $ac1, $ac2 or $ac3).

Also, some built-in functions prefer or require immediate numbers as parameters, because the corresponding DSP instructions accept both immediate numbers and register operands, or accept immediate numbers only. The immediate parameters are listed as follows.

```
imm0_3: 0 to 3.
imm0_7: 0 to 7.
imm0_15: 0 to 15.
imm0_31: 0 to 31.
imm0_63: 0 to 63.
imm0_255: 0 to 255.
imm_n32_31: -32 to 31.
imm_n512_511: -512 to 511.
```

The following built-in functions map directly to a particular DSP instruction. Please refer to the PIC32 DSP documentation for details on what each instruction does. In the table below, the function provides a way to generate the DSP instruction in your code. For example, in v2q15 __builtin_mips_**addq_ph** (v2q15, v2q15), the **addq_ph** is the actual DSP instruction.

Table 27-1. Map Directly to DSP Instruction

v2q15 __builtin_mips_absq_s_ph (v2q15)
q31 __builtin_mips_absq_s_w (q31)
v2q15 __builtin_mips_addq_ph (v2q15, v2q15)
v2q15 __builtin_mips_addq_s_ph (v2q15, v2q15)
q31 __builtin_mips_addq_s_w (q31, q31)
i32 __builtin_mips_addsc (i32, i32)
v4i8 __builtin_mips_addu_qb (v4i8, v4i8)
v4i8 __builtin_mips_addu_s_qb (v4i8, v4i8)
i32 __builtin_mips_addwc (i32, i32)
i32 __builtin_mips_bitrev (i32)
i32 __builtin_mips_bposge32 (void)
void __builtin_mips_cmp_eq_ph (v2q15, v2q15)
void __builtin_mips_cmp_le_ph (v2q15, v2q15)
void __builtin_mips_cmp_lt_ph (v2q15, v2q15)
i32 __builtin_mips_cmpgu_eq_qb (v4i8, v4i8)
i32 __builtin_mips_cmpgu_le_qb (v4i8, v4i8)
i32 __builtin_mips_cmpgu_lt_qb (v4i8, v4i8)
void __builtin_mips_cmpu_eq_qb (v4i8, v4i8)
void __builtin_mips_cmpu_le_qb (v4i8, v4i8)
void __builtin_mips_cmpu_lt_qb (v4i8, v4i8)
a64 __builtin_mips_dpaq_s_w_ph (a64, v2q15, v2q15)
a64 __builtin_mips_dpaq_sa_l_w (a64, q31, q31)

```
a64 __builtin_mips_dpau_h_qbl (a64, v4i8, v4i8)
```

```
a64 __builtin_mips_dpau_h_qbr (a64, v4i8, v4i8)
```

```
a64 __builtin_mips_dpsq_s_w_ph (a64, v2q15, v2q15)
```

```
a64 __builtin_mips_dpsq_sa_l_w (a64, q31, q31)
```

```
a64 __builtin_mips_dpsu_h_qbl (a64, v4i8, v4i8)
```

```
a64 __builtin_mips_dpsu_h_qbr (a64, v4i8, v4i8)
```

```
i32 __builtin_mips_extp (a64, i32)
```

```
i32 __builtin_mips_extp (a64, imm0_31)
```

```
i32 __builtin_mips_extpdp (a64, i32)
```

```
i32 __builtin_mips_extpdp (a64, imm0_31)
```

```
i32 __builtin_mips_extr_r_w (a64, i32)
```

```
i32 __builtin_mips_extr_r_w (a64, imm0_31)
```

```
i32 __builtin_mips_extr_rs_w (a64, i32)
```

```
i32 __builtin_mips_extr_rs_w (a64, imm0_31)
```

```
i32 __builtin_mips_extr_s_h (a64, i32)
```

```
i32 __builtin_mips_extr_s_h (a64, imm0_31)
```

```
i32 __builtin_mips_extr_w (a64, i32)
```

```
i32 __builtin_mips_extr_w (a64, imm0_31)
```

```
i32 __builtin_mips_insv (i32, i32)
```

```
i32 __builtin_mips_lbux (void *, i32)
```

```
i32 __builtin_mips_lhx (void *, i32)
```

```
i32 __builtin_mips_lwx (void *, i32)
```

```
a64 __builtin_mips_maq_s_w_phl (a64, v2q15, v2q15)
```

```
a64 __builtin_mips_maq_s_w_phr (a64, v2q15, v2q15)
```

```
a64 __builtin_mips_maq_sa_w_phl (a64, v2q15, v2q15)
```

```
a64 __builtin_mips_maq_sa_w_phr (a64, v2q15, v2q15)
```

```
i32 __builtin_mips_modsub (i32, i32)
```

```
a64 __builtin_mips_mthlip (a64, i32)
```

```
q31 __builtin_mips_muleq_s_w_phl (v2q15, v2q15)
```

```
q31 __builtin_mips_muleq_s_w_phr (v2q15, v2q15)
```

```
v2q15 __builtin_mips_muleu_s_ph_qbl (v4i8, v2q15)
```

```
v2q15 __builtin_mips_muleu_s_ph_qbr (v4i8, v2q15)
```

```
v2q15 __builtin_mips_mulq_rs_ph (v2q15, v2q15)
```

```
a64 __builtin_mips_mulsaq_s_w_ph (a64, v2q15, v2q15)
```

```
v2q15 __builtin_mips_packrl_ph (v2q15, v2q15)
```

```
v2q15 __builtin_mips_pick_ph (v2q15, v2q15)
```

```
v4i8 __builtin_mips_pick_qb (v4i8, v4i8)
```

`q31 __builtin_mips_preceq_w_phl (v2q15)`
`q31 __builtin_mips_preceq_w_phr (v2q15)`
`v2q15 __builtin_mips_precequ_ph_qbl (v4i8)`
`v2q15 __builtin_mips_precequ_ph_qbla (v4i8)`
`v2q15 __builtin_mips_precequ_ph_qbr (v4i8)`
`v2q15 __builtin_mips_precequ_ph_qbra (v4i8)`
`v2q15 __builtin_mips_preceu_ph_qbl (v4i8)`
`v2q15 __builtin_mips_preceu_ph_qbla (v4i8)`
`v2q15 __builtin_mips_preceu_ph_qbr (v4i8)`
`v2q15 __builtin_mips_preceu_ph_qbra (v4i8)`
`v2q15 __builtin_mips_precrq_ph_w (q31, q31)`
`v4i8 __builtin_mips_precrq_qb_ph (v2q15, v2q15)`
`v2q15 __builtin_mips_precrq_rs_ph_w (q31, q31)`
`v4i8 __builtin_mips_precrqu_s_qb_ph (v2q15, v2q15)`
`i32 __builtin_mips_raddu_w_qb (v4i8)`
`i32 __builtin_mips_rddsp (imm0_63)`
`v2q15 __builtin_mips_repl_ph (i32)`
`v2q15 __builtin_mips_repl_ph (imm_n512_511)`
`v4i8 __builtin_mips_repl_qb (i32)`
`v4i8 __builtin_mips_repl_qb (imm0_255)`
`a64 __builtin_mips_shilo (a64, i32)`
`a64 __builtin_mips_shilo (a64, imm_n32_31)`
`v2q15 __builtin_mips_shll_ph (v2q15, i32)`
`v2q15 __builtin_mips_shll_ph (v2q15, imm0_15)`
`v4i8 __builtin_mips_shll_qb (v4i8, i32)`
`v4i8 __builtin_mips_shll_qb (v4i8, imm0_7)`
`v2q15 __builtin_mips_shll_s_ph (v2q15, i32)`
`v2q15 __builtin_mips_shll_s_ph (v2q15, imm0_15)`
`q31 __builtin_mips_shll_s_w (q31, i32)`
`q31 __builtin_mips_shll_s_w (q31, imm0_31)`
`v2q15 __builtin_mips_shra_ph (v2q15, i32)`
`v2q15 __builtin_mips_shra_ph (v2q15, imm0_15)`
`v2q15 __builtin_mips_shra_r_ph (v2q15, i32)`
`v2q15 __builtin_mips_shra_r_ph (v2q15, imm0_15)`
`q31 __builtin_mips_shra_r_w (q31, i32)`
`q31 __builtin_mips_shra_r_w (q31, imm0_31)`
`v4i8 __builtin_mips_shrl_qb (v4i8, i32)`

v4i8 __builtin_mips_shrl_qb (v4i8, imm0_7)
v2q15 __builtin_mips_subq_ph (v2q15, v2q15)
v2q15 __builtin_mips_subq_s_ph (v2q15, v2q15)
q31 __builtin_mips_subq_s_w (q31, q31)
v4i8 __builtin_mips_subu_qb (v4i8, v4i8)
v4i8 __builtin_mips_subu_s_qb (v4i8, v4i8)
void __builtin_mips_wrdsp (i32, imm0_63)

The following built-in functions map directly to a particular MIPS DSP REV 2 instruction. Please refer to the PIC32 DSP documentation for details on what each instruction does.

Table 27-2. Map Directly to MIPS DSP Instruction

v4q7 __builtin_mips_absq_s_qb (v4q7);
v2q15 __builtin_mips_addqh_ph (v2q15, v2q15);
v2q15 __builtin_mips_addqh_r_ph (v2q15, v2q15);
q31 __builtin_mips_addqh_r_w (q31, q31);
q31 __builtin_mips_addqh_w (q31, q31);
v2i16 __builtin_mips_addu_ph (v2i16, v2i16);
v2i16 __builtin_mips_addu_s_ph (v2i16, v2i16);
v4i8 __builtin_mips_adduh_qb (v4i8, v4i8);
v4i8 __builtin_mips_adduh_r_qb (v4i8, v4i8);
i32 __builtin_mips_append (i32, i32, imm0_31);
i32 __builtin_mips_balign (i32, i32, imm0_3);
i32 __builtin_mips_cmpgdu_eq_qb (v4i8, v4i8);
i32 __builtin_mips_cmpgdu_le_qb (v4i8, v4i8);
i32 __builtin_mips_cmpgdu_lt_qb (v4i8, v4i8);
a64 __builtin_mips_dpa_w_ph (a64, v2i16, v2i16);
a64 __builtin_mips_dpaqx_s_w_ph (a64, v2q15, v2q15);
a64 __builtin_mips_dpaqx_sa_w_ph (a64, v2q15, v2q15);
a64 __builtin_mips_dpax_w_ph (a64, v2i16, v2i16);
a64 __builtin_mips_dps_w_ph (a64, v2i16, v2i16);
a64 __builtin_mips_dpsqx_s_w_ph (a64, v2q15, v2q15);
a64 __builtin_mips_dpsqx_sa_w_ph (a64, v2q15, v2q15);
a64 __builtin_mips_dpsx_w_ph (a64, v2i16, v2i16);
a64 __builtin_mips_madd (a64, i32, i32);
a64 __builtin_mips_maddu (a64, ui32, ui32);
a64 __builtin_mips_msub (a64, i32, i32);
a64 __builtin_mips_msubu (a64, ui32, ui32);

```
v2i16 __builtin_mips_mul_ph (v2i16, v2i16);
```

```
v2i16 __builtin_mips_mul_s_ph (v2i16, v2i16);
```

```
q31 __builtin_mips_mulq_rs_w (q31, q31);
```

```
v2q15 __builtin_mips_mulq_s_ph (v2q15, v2q15);
```

```
q31 __builtin_mips_mulq_s_w (q31, q31);
```

```
a64 __builtin_mips_mulsa_w_ph (a64, v2i16, v2i16);
```

```
a64 __builtin_mips_mult (i32, i32);
```

```
a64 __builtin_mips_multu (ui32, ui32);
```

```
v4i8 __builtin_mips_precr_qb_ph (v2i16, v2i16);
```

```
v2i16 __builtin_mips_precr_sra_ph_w (i32, i32, imm0_31);
```

```
v2i16 __builtin_mips_precr_sra_r_ph_w (i32, i32, imm0_31);
```

```
i32 __builtin_mips_prepend (i32, i32, imm0_31);
```

```
v4i8 __builtin_mips_shra_qb (v4i8, i32);
```

```
v4i8 __builtin_mips_shra_qb (v4i8, imm0_7);
```

```
v4i8 __builtin_mips_shra_r_qb (v4i8, i32);
```

```
v4i8 __builtin_mips_shra_r_qb (v4i8, imm0_7);
```

```
v2i16 __builtin_mips_shrl_ph (v2i16, i32);
```

```
v2i16 __builtin_mips_shrl_ph (v2i16, imm0_15);
```

```
v2q15 __builtin_mips_subqh_ph (v2q15, v2q15);
```

```
v2q15 __builtin_mips_subqh_r_ph (v2q15, v2q15);
```

```
q31 __builtin_mips_subqh_r_w (q31, q31);
```

```
q31 __builtin_mips_subqh_w (q31, q31);
```

```
v2i16 __builtin_mips_subu_ph (v2i16, v2i16);
```

```
v2i16 __builtin_mips_subu_s_ph (v2i16, v2i16);
```

```
v4i8 __builtin_mips_subuh_qb (v4i8, v4i8);
```

```
v4i8 __builtin_mips_subuh_r_qb (v4i8, v4i8);
```

28. ASCII Character Set

Table 28-1. ASCII Character Set

Least Significant Character	Most Significant Character								
	Hex	0	1	2	3	4	5	6	7
	0	NUL	DLE	Space	0	@	P	'	p
	1	SOH	DC1	!	1	A	Q	a	q
	2	STX	DC2	"	2	B	R	b	r
	3	ETX	DC3	#	3	C	S	c	s
	4	EOT	DC4	$	4	D	T	d	t
	5	ENQ	NAK	%	5	E	U	e	u
	6	ACK	SYN	&	6	F	V	f	v
	7	Bell	ETB	'	7	G	W	g	w
	8	BS	CAN	(	8	H	X	h	x
	9	HT	EM	)	9	I	Y	i	y
	A	LF	SUB	*	:	J	Z	j	z
	B	VT	ESC	+	;	K	[	k	{
	C	FF	FS	,	<	L	\	l	|
	D	CR	GS	-	=	M	]	m	}
	E	SO	RS	.	>	N	^	n	~
	F	SI	US	/	?	O	_	o	DEL

29. Document Revision History

Revision A (August 2018)

Initial revision of the document.

Revision B (April 2020)

- The document was re-formatted throughout, including paragraph numbering and minor editing changes.
- Added new topic 4.5.20 How Do I Stop My Project's Checksum From Changing? to the How To's section.
- Corrected description of option `-fenforce-eh-specs` in Table 5-7.
- Revised 6. Command-line Driver section extensively.
- Updated Example 10-6 and Example 10-7 in 10.6.3 Using an Application-Defined Memory Region.
- Added examples for `keep` and `noload` in the 14.2.1 Function Attributes.
- Adding more information to 16.2.2 Jump to NMI Handler (_nmi_handler) if an NMI Occurred.

The Microchip Website

Microchip provides online support via our website at http://www.microchip.com/. This website is used to make files and information easily available to customers. Some of the content available includes:

- **Product Support** – Data sheets and errata, application notes and sample programs, design resources, user's guides and hardware support documents, latest software releases and archived software
- **General Technical Support** – Frequently Asked Questions (FAQs), technical support requests, online discussion groups, Microchip design partner program member listing
- **Business of Microchip** – Product selector and ordering guides, latest Microchip press releases, listing of seminars and events, listings of Microchip sales offices, distributors and factory representatives

Product Change Notification Service

Microchip's product change notification service helps keep customers current on Microchip products. Subscribers will receive email notification whenever there are changes, updates, revisions or errata related to a specified product family or development tool of interest.

To register, go to http://www.microchip.com/pcn and follow the registration instructions.

Customer Support

Users of Microchip products can receive assistance through several channels:

- Distributor or Representative
- Local Sales Office
- Embedded Solutions Engineer (ESE)
- Technical Support

Customers should contact their distributor, representative or ESE for support. Local sales offices are also available to help customers. A listing of sales offices and locations is included in this document.

Technical support is available through the web site at: http://www.microchip.com/support

Microchip Devices Code Protection Feature

Note the following details of the code protection feature on Microchip devices:

- Microchip products meet the specification contained in their particular Microchip Data Sheet.
- Microchip believes that its family of products is one of the most secure families of its kind on the market today, when used in the intended manner and under normal conditions.
- There are dishonest and possibly illegal methods used to breach the code protection feature. All of these methods, to our knowledge, require using the Microchip products in a manner outside the operating specifications contained in Microchip's Data Sheets. Most likely, the person doing so is engaged in theft of intellectual property.
- Microchip is willing to work with the customer who is concerned about the integrity of their code.
- Neither Microchip nor any other semiconductor manufacturer can guarantee the security of their code. Code protection does not mean that we are guaranteeing the product as "unbreakable."

Code protection is constantly evolving. We at Microchip are committed to continuously improving the code protection features of our products. Attempts to break Microchip's code protection feature may be a violation of the Digital Millennium Copyright Act. If such acts allow unauthorized access to your software or other copyrighted work, you may have a right to sue for relief under that Act.

Legal Notice

Information contained in this publication regarding device applications and the like is provided only for your convenience and may be superseded by updates. It is your responsibility to ensure that your application meets with

your specifications. MICROCHIP MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WHETHER EXPRESS OR IMPLIED, WRITTEN OR ORAL, STATUTORY OR OTHERWISE, RELATED TO THE INFORMATION, INCLUDING BUT NOT LIMITED TO ITS CONDITION, QUALITY, PERFORMANCE, MERCHANTABILITY OR FITNESS FOR PURPOSE. Microchip disclaims all liability arising from this information and its use. Use of Microchip devices in life support and/or safety applications is entirely at the buyer's risk, and the buyer agrees to defend, indemnify and hold harmless Microchip from any and all damages, claims, suits, or expenses resulting from such use. No licenses are conveyed, implicitly or otherwise, under any Microchip intellectual property rights unless otherwise stated.

Trademarks

The Microchip name and logo, the Microchip logo, Adaptec, AnyRate, AVR, AVR logo, AVR Freaks, BesTime, BitCloud, chipKIT, chipKIT logo, CryptoMemory, CryptoRF, dsPIC, FlashFlex, flexPWR, HELDO, IGLOO, JukeBlox, KeeLoq, Kleer, LANCheck, LinkMD, maXStylus, maXTouch, MediaLB, megaAVR, Microsemi, Microsemi logo, MOST, MOST logo, MPLAB, OptoLyzer, PackeTime, PIC, picoPower, PICSTART, PIC32 logo, PolarFire, Prochip Designer, QTouch, SAM-BA, SenGenuity, SpyNIC, SST, SST Logo, SuperFlash, Symmetricom, SyncServer, Tachyon, TempTrackr, TimeSource, tinyAVR, UNI/O, Vectron, and XMEGA are registered trademarks of Microchip Technology Incorporated in the U.S.A. and other countries.

APT, ClockWorks, The Embedded Control Solutions Company, EtherSynch, FlashTec, Hyper Speed Control, HyperLight Load, IntelliMOS, Libero, motorBench, mTouch, Powermite 3, Precision Edge, ProASIC, ProASIC Plus, ProASIC Plus logo, Quiet-Wire, SmartFusion, SyncWorld, Temux, TimeCesium, TimeHub, TimePictra, TimeProvider, Vite, WinPath, and ZL are registered trademarks of Microchip Technology Incorporated in the U.S.A.

Adjacent Key Suppression, AKS, Analog-for-the-Digital Age, Any Capacitor, AnyIn, AnyOut, BlueSky, BodyCom, CodeGuard, CryptoAuthentication, CryptoAutomotive, CryptoCompanion, CryptoController, dsPICDEM, dsPICDEM.net, Dynamic Average Matching, DAM, ECAN, EtherGREEN, In-Circuit Serial Programming, ICSP, INICnet, Inter-Chip Connectivity, JitterBlocker, KleerNet, KleerNet logo, memBrain, Mindi, MiWi, MPASM, MPF, MPLAB Certified logo, MPLIB, MPLINK, MultiTRAK, NetDetach, Omniscient Code Generation, PICDEM, PICDEM.net, PICkit, PICtail, PowerSmart, PureSilicon, QMatrix, REAL ICE, Ripple Blocker, SAM-ICE, Serial Quad I/O, SMART-I.S., SQI, SuperSwitcher, SuperSwitcher II, Total Endurance, TSHARC, USBCheck, VariSense, ViewSpan, WiperLock, Wireless DNA, and ZENA are trademarks of Microchip Technology Incorporated in the U.S.A. and other countries.

SQTP is a service mark of Microchip Technology Incorporated in the U.S.A.

The Adaptec logo, Frequency on Demand, Silicon Storage Technology, and Symmcom are registered trademarks of Microchip Technology Inc. in other countries.

GestIC is a registered trademark of Microchip Technology Germany II GmbH & Co. KG, a subsidiary of Microchip Technology Inc., in other countries.

All other trademarks mentioned herein are property of their respective companies.

© 2020, Microchip Technology Incorporated, Printed in the U.S.A., All Rights Reserved.

ISBN: 978-1-5224-6006-0

Quality Management System

For information regarding Microchip's Quality Management Systems, please visit http://www.microchip.com/quality.

Worldwide Sales and Service

AMERICAS	ASIA/PACIFIC	ASIA/PACIFIC	EUROPE
Corporate Office 2355 West Chandler Blvd. Chandler, AZ 85224-6199 Tel: 480-792-7200 Fax: 480-792-7277 Technical Support: http://www.microchip.com/support Web Address: http://www.microchip.com	**Australia - Sydney** Tel: 61-2-9868-6733 **China - Beijing** Tel: 86-10-8569-7000 **China - Chengdu** Tel: 86-28-8665-5511 **China - Chongqing** Tel: 86-23-8980-9588 **China - Dongguan** Tel: 86-769-8702-9880	**India - Bangalore** Tel: 91-80-3090-4444 **India - New Delhi** Tel: 91-11-4160-8631 **India - Pune** Tel: 91-20-4121-0141 **Japan - Osaka** Tel: 81-6-6152-7160 **Japan - Tokyo** Tel: 81-3-6880- 3770	**Austria - Wels** Tel: 43-7242-2244-39 Fax: 43-7242-2244-393 **Denmark - Copenhagen** Tel: 45-4450-2828 Fax: 45-4485-2829 **Finland - Espoo** Tel: 358-9-4520-820 **France - Paris** Tel: 33-1-69-53-63-20 Fax: 33-1-69-30-90-79
Atlanta Duluth, GA Tel: 678-957-9614 Fax: 678-957-1455	**China - Guangzhou** Tel: 86-20-8755-8029 **China - Hangzhou** Tel: 86-571-8792-8115	**Korea - Daegu** Tel: 82-53-744-4301 **Korea - Seoul** Tel: 82-2-554-7200	**Germany - Garching** Tel: 49-8931-9700 **Germany - Haan** Tel: 49-2129-3766400
Austin, TX Tel: 512-257-3370	**China - Hong Kong SAR** Tel: 852-2943-5100	**Malaysia - Kuala Lumpur** Tel: 60-3-7651-7906	**Germany - Heilbronn** Tel: 49-7131-72400
Boston Westborough, MA Tel: 774-760-0087 Fax: 774-760-0088	**China - Nanjing** Tel: 86-25-8473-2460 **China - Qingdao** Tel: 86-532-8502-7355	**Malaysia - Penang** Tel: 60-4-227-8870 **Philippines - Manila** Tel: 63-2-634-9065	**Germany - Karlsruhe** Tel: 49-721-625370 **Germany - Munich** Tel: 49-89-627-144-0
Chicago Itasca, IL Tel: 630-285-0071 Fax: 630-285-0075	**China - Shanghai** Tel: 86-21-3326-8000 **China - Shenyang** Tel: 86-24-2334-2829	**Singapore** Tel: 65-6334-8870 **Taiwan - Hsin Chu** Tel: 886-3-577-8366	Fax: 49-89-627-144-44 **Germany - Rosenheim** Tel: 49-8031-354-560
Dallas Addison, TX Tel: 972-818-7423 Fax: 972-818-2924	**China - Shenzhen** Tel: 86-755-8864-2200 **China - Suzhou** Tel: 86-186-6233-1526	**Taiwan - Kaohsiung** Tel: 886-7-213-7830 **Taiwan - Taipei** Tel: 886-2-2508-8600	**Israel - Ra'anana** Tel: 972-9-744-7705 **Italy - Milan** Tel: 39-0331-742611 Fax: 39-0331-466781
Detroit Novi, MI Tel: 248-848-4000	**China - Wuhan** Tel: 86-27-5980-5300	**Thailand - Bangkok** Tel: 66-2-694-1351	**Italy - Padova** Tel: 39-049-7625286
Houston, TX Tel: 281-894-5983	**China - Xian** Tel: 86-29-8833-7252	**Vietnam - Ho Chi Minh** Tel: 84-28-5448-2100	**Netherlands - Drunen** Tel: 31-416-690399 Fax: 31-416-690340
Indianapolis Noblesville, IN Tel: 317-773-8323 Fax: 317-773-5453 Tel: 317-536-2380	**China - Xiamen** Tel: 86-592-2388138 **China - Zhuhai** Tel: 86-756-3210040		**Norway - Trondheim** Tel: 47-72884388 **Poland - Warsaw** Tel: 48-22-3325737
Los Angeles Mission Viejo, CA Tel: 949-462-9523 Fax: 949-462-9608 Tel: 951-273-7800			**Romania - Bucharest** Tel: 40-21-407-87-50 **Spain - Madrid** Tel: 34-91-708-08-90 Fax: 34-91-708-08-91
Raleigh, NC Tel: 919-844-7510 **New York, NY** Tel: 631-435-6000			**Sweden - Gothenberg** Tel: 46-31-704-60-40 **Sweden - Stockholm** Tel: 46-8-5090-4654
San Jose, CA Tel: 408-735-9110 Tel: 408-436-4270			**UK - Wokingham** Tel: 44-118-921-5800 Fax: 44-118-921-5820
Canada - Toronto Tel: 905-695-1980 Fax: 905-695-2078			