



Section 50. CPU for Devices with MIPS32[®] microAptiv[™] and M-Class Cores

This section of the manual contains the following topics:

50.1	Introduction	50-2
50.2	Architecture Overview	50-4
50.3	PIC32 CPU Details	50-8
50.4	Special Considerations When Writing to CP0 Registers	50-13
50.5	MIPS32 Architecture	50-14
50.6	CPU Bus	50-15
50.7	Internal System Busses	50-15
50.8	Set/Clear/Invert	50-16
50.9	ALU Status Bits	50-16
50.10	Interrupt and Exception Mechanism	50-17
50.11	Programming Model	50-17
50.12	Floating Point Unit (FPU)	50-24
50.13	Coprocessor 0 (CP0) Registers	50-42
50.14	Coprocessor 1 (CP1) Registers	50-121
50.15	microMIPS Execution	50-132
50.16	MCU ASE Extension	50-132
50.17	MIPS DSP ASE Extension	50-133
50.18	Memory Model (MCU only)	50-133
50.19	Memory Management (MPU only)	50-135
50.20	L1 Caches (MPU only)	50-141
50.21	CPU Instructions	50-145
50.22	MIPS DSP ASE Instructions	50-151
50.23	CPU Initialization	50-153
50.24	Effects of a Reset	50-154
50.25	Related Application Notes	50-155
50.26	Revision History	50-156

Note: This family reference manual section is meant to serve as a complement to device data sheets. Depending on the device variant, this manual section may not apply to all PIC32 devices.

Please consult the note at the beginning of the “**CPU**” chapter in the current device data sheet to check whether this document supports the device you are using.

Device data sheets and family reference manual sections are available for download from the Microchip Worldwide Web site at: <http://www.microchip.com>

50.1 INTRODUCTION

Depending on the device family, PIC32 devices are a complex System-on-Chip (SoC), which are based on the microAptiv™ Microprocessor core or the M-Class Microprocessor core from Imagination Technologies Ltd. This document provides an overview of the CPU system architecture and features of PIC32 microcontrollers that feature these microprocessor cores.

The microAptiv Microprocessor core is a superset of the MIPS® M14KE™ and M14KEc™ Microprocessor cores. These cores are state of the art, 32-bit, low-power, RISC processor cores with the enhanced MIPS32® Release 2 Instruction Set Architecture (ISA).

The M-Class Microprocessor core is a superset of the microAptiv™ Microprocessor core. This 32-bit, low-power, RISC processor core uses the enhanced MIPS32® Release 5 Instruction Set Architecture (ISA).

Visit the Imagination Technologies Ltd. website (www.imgtec.com) to learn more about the microprocessor cores.

Depending on the core configuration, one of two options, MCU or MPU, are used, as shown in [Table 50-1](#).

Table 50-1: microAptiv and M-Class Microprocessor Core Configurations

MCU Features	MPU Features
Split-bus architecture	Unified bus architecture
Integrated DSP ASE (see Note 1)	Integrated DSP ASE (see Note 1)
Integrated MCU™ ASE	Integrated MCU ASE
microMIPS™ code compression	microMIPS code compression
FMT-based MMU	TLB-based MMU
Two shadow register sets	Eight shadow register sets
EJTAG TAP controller	EJTAG TAP controller
Performance counters	Performance counters
Hardware Trace (iFlowtrace®)	Hardware Trace (iFlowtrace)
	Level One (L1) CPU cache

Note 1: This feature is not available on all devices, refer to the “**CPU**” chapter of the specific device data sheet to determine availability.

The primary difference between the MCU and MPU is the presence of an L1 cache and TLB-based MMU on the MPU. These features are used to facilitate PIC32 designs that use operating systems to manage virtual memory.

50.1.1 Key Features Common to All PIC32 Devices with the microAptiv Microprocessor Core

The following key features are common to all PIC32 devices that are based on the microAptiv Microprocessor core:

- microMIPS variable-length instruction mode for compact code
- Vectored interrupt controller with up to 256 interrupt sources
- Atomic bit manipulations on peripheral registers (Single cycle)
- High-speed Microchip ICD port with hardware-based non-intrusive data monitoring and application data streaming functions
- EJTAG debug port allows extensive third party debug, programming and test tools support
- Instruction controlled power management modes
- Five-stage pipelined instruction execution
- Internal code protection to help protect intellectual property
- Arithmetic saturation and overflow handling support
- Zero cycle overhead saturation and rounding operations
- Atomic read-modify-write memory-to-memory instructions
- MAC instructions with up to 4 accumulators
- Native fractional data type (Q15, Q31) with rounding support
- Digital Signal Processing (DSP) Application-Specific Extension (ASE) Revision 2, which adds DSP capabilities with support for powerful data processing operations
- Multiply/Divide unit with a maximum issue rate of one 32 x 32 multiply per clock

50.1.2 Key Features Common to All PIC32 Devices with the M-Class Microprocessor Core

In addition to the features described for devices with the microAptiv core, the following key features are common to all PIC32 devices that are based on the M-Class Microprocessor core:

- Implements the latest MIPS Release 5 Architecture, which includes IP protection and reliability for industrial controllers, Internet of Things (IoT), wearables, wireless communications, automotive, and storage
- Floating Point Unit (FPU)

50.1.3 Related MIPS Documentation

Related MIPS documentation is available for download from the related Imagination Technologies Ltd. product page. Please note that a login may be required to access these documents.

Documentation for the microAptiv core is available for download at:

<http://www.imgtec.com/mips/aptiv/microaptiv.asp>

Documentation for the M-Class core is available for download at:

<http://www.imgtec.com/mips/warrior/mclass.asp>

50.2 ARCHITECTURE OVERVIEW

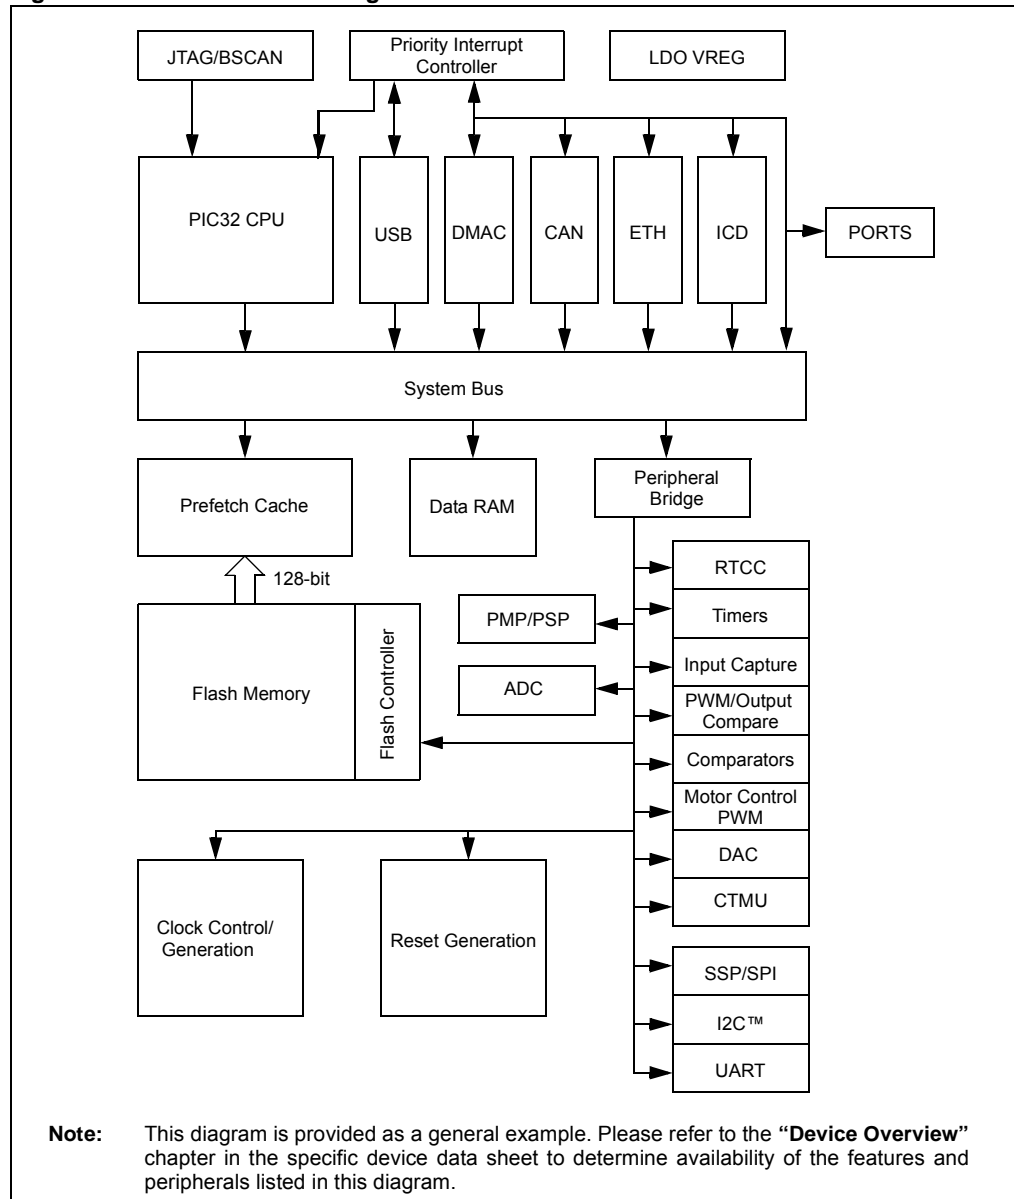
The PIC32 family of devices are complex systems-on-a-chip that contain many features. Included in all processors of the PIC32 family is a high-performance RISC CPU, which can be programmed in 32-bit and 16-bit modes, and even mixed modes.

Devices with the M-Class core include a Floating Point Unit (FPU) that implements the MIPS Release 5 Instruction Set Architecture for floating point computation. The FPU implementation supports the ANSI/IEEE Standard 754 (IEEE Standard for Binary Floating-Point Arithmetic) for single- and double-precision data formats.

Note: Refer to the “CPU” chapter in the specific device data sheet to determine availability of the FPU module in your device.

PIC32 devices contain a high-performance Interrupt Controller, DMA controller, USB controller, in-circuit debugger, a high-performance switching matrix for high-speed data accesses to the peripherals, and on-chip data RAM memory, which holds data and programs. The optional prefetch cache and prefetch buffer for the Flash memory, which hides the latency of the Flash, provides zero Wait state equivalent performance.

Figure 50-1: PIC32 Block Diagram



The peripherals of a PIC32 device connect to the CPU through a System Bus and a series of internal busses. The main peripheral bus connects most of the peripheral units to the System Bus through one or more peripheral bridges.

The PIC32 CPU performs operations under program control. Instructions are fetched by the CPU and are synchronously decoded and executed. Instructions exist in either Program Flash memory or Data RAM memory. In addition, PIC32 devices with the microAptiv and M-Class core incorporate the MIPS DSP Application-Specific Extension Revision 2 that provides digital signal processing (DSP) capabilities with support for a number of powerful data processing operations.

The PIC32 CPU is based on a load/store architecture and performs most operations on a set of internal registers. Specific load and store instructions are used to move data between these internal registers and the outside world.

Figure 50-2: microAptiv™ Microprocessor Core Block Diagram

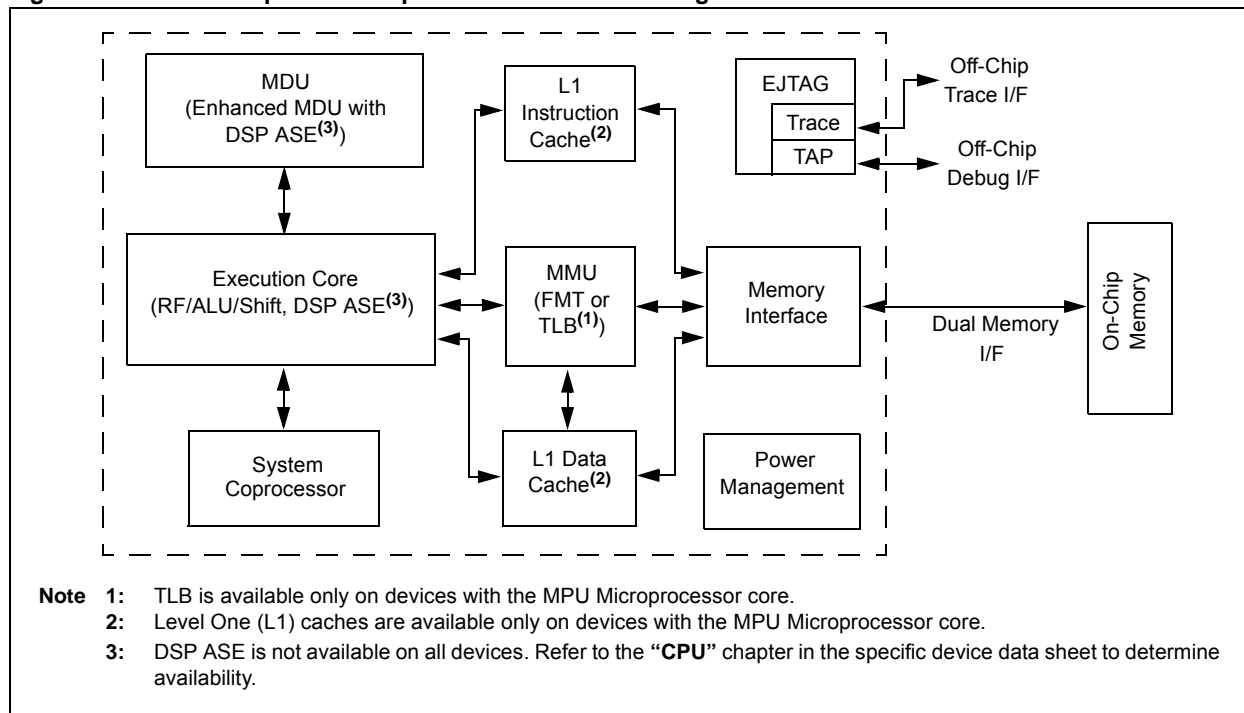
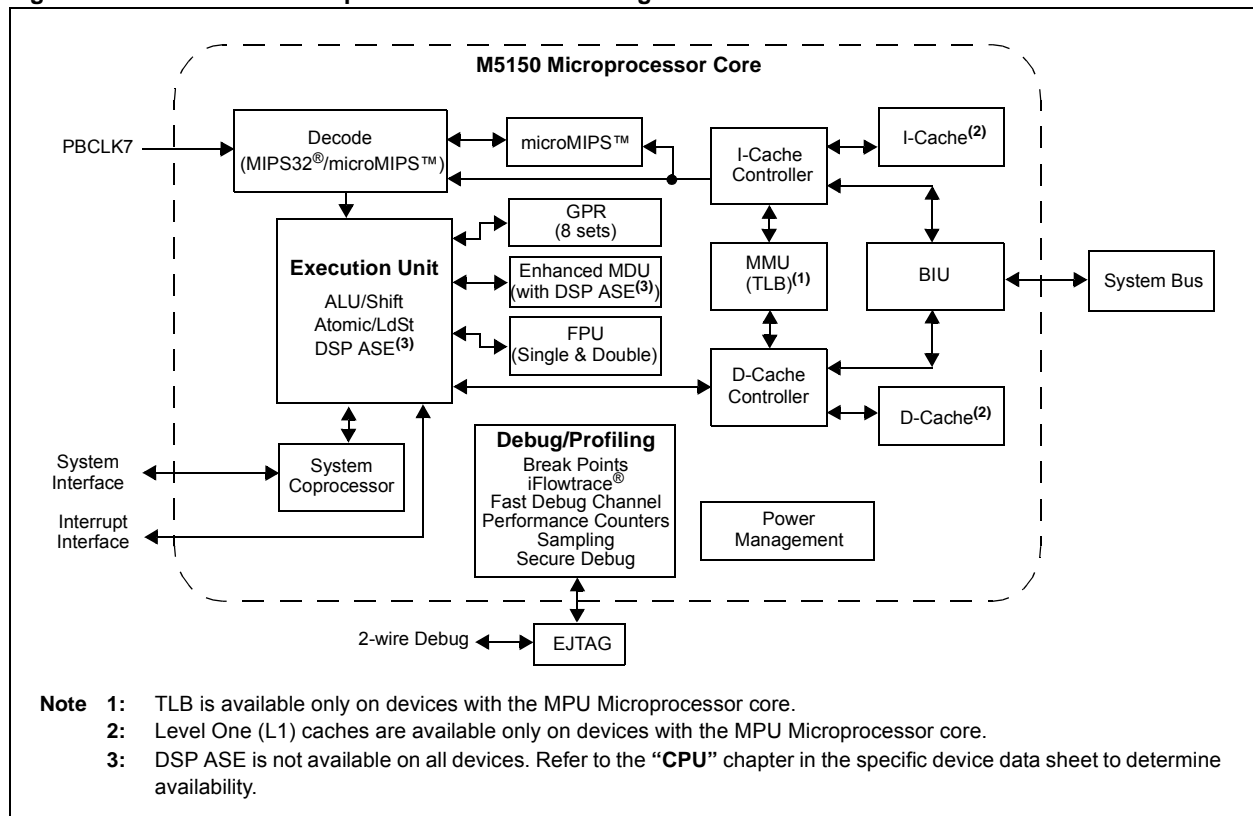


Figure 50-3: M-Class Microprocessor Core Block Diagram



50.2.1 Busses

All PIC32 devices use a System Bus to connect the CPU and other bus masters to memory and other target devices. The System Bus controls and arbitrates accesses between different bus masters and targets. The type of System Bus and the bus architecture in a specific PIC32 device is dependent on which the microAptiv or M-Class CPU core is used.

PIC32 devices based on the MCU Microprocessor core use a split-bus CPU architecture. In this architecture, there are separate busses for instruction fetch and data load/store operations. Both the instruction, or I-side bus, and the data, or D-side bus, are connected to the System Bus. The System Bus allows simultaneous accesses between different bus masters accessing different targets, and uses an arbitration algorithm to serialize accesses from different masters to the same target. Since the CPU has two different data paths to the System Bus, the CPU is effectively two different bus masters to the system. When running from Flash memory, load and store operations to SRAM and the internal peripherals will occur in parallel to instruction fetches from Flash memory.

PIC32 devices based on the MPU core use a unified bus CPU architecture along with a multi-layer (crossbar) System Bus. In this architecture, the CPU has a single interface to the System Bus. The System Bus uses dedicated links to provide multiple independent data paths between bus initiators and targets. This allows for concurrent data transactions on the bus.

Note: Please refer to the “**Memory Organization**” chapter in the specific device data sheet and **Section 3. “Memory Organization”** (DS60001115) of the “*PIC32 Family Reference Manual*” for a description of the System Bus for a specific device.

In addition to the CPU, and depending on the device variant, there are other bus masters in PIC32 devices:

- DMA controller
- In-Circuit Debugger (ICD)
- USB controller
- CAN controller
- Ethernet controller

50.2.2 Core Timer

The PIC32 architecture includes a core timer that is available to application programs. This timer is implemented in the form of two coprocessor registers: the Count register, and the Compare register. The Count register is incremented every two system clock (SYSCLK) cycles. The incrementing of Count can be optionally suspended during Debug mode. The Compare register is used to cause a timer interrupt if desired. An interrupt is generated when the Compare register matches the Count register. An interrupt is taken only if it is enabled in the interrupt controller.

For more information on the core timer, refer to **50.13 “Coprocessor 0 (CP0) Registers”** and **Section 8. “Interrupts.”** (DS60001108) of the “*PIC32 Family Reference Manual*”.

GS: `SYS_CLK_FREQ` is defined in system_config.h, so, we can define `CORE_TIMER_FREQ` as follows:

```
#include "system_config.h"
#define CORE_TIMER_FREQ_HZ (SYS_CLK_FREQ/2)
```

50.3 PIC32 CPU DETAILS

50.3.1 Pipeline Stages

The pipeline consists of five stages:

- Instruction (I) Stage
- Execution (E) Stage
- Memory (M) Stage
- Align (A) Stage
- Writeback (W) Stage

50.3.1.1 I STAGE – INSTRUCTION FETCH

During I stage:

- An instruction is fetched from the instruction SRAM
- microMIPS instructions are converted into instructions that are similar to MIPS32 instructions

50.3.1.2 E STAGE – EXECUTION

During E stage:

- Operands are fetched from the register file
- Operands from the M and A stage are bypassed to this stage
- The Arithmetic Logic Unit (ALU) begins the arithmetic or logical operation for register-to-register instructions
- The ALU calculates the data virtual address for load and store instructions and the MMU performs the fixed virtual-to-physical address translation
- The ALU determines whether the branch condition is true and calculates the virtual branch target address for branch instructions
- Instruction logic selects an instruction address and the MMU performs the fixed virtual-to-physical address translation
- All multiply divide operations begin in this stage

50.3.1.3 M STAGE – MEMORY FETCH

During M stage:

- The arithmetic or logic ALU operation completes
- The data SRAM access is performed for load and store instructions
- A 16 x 16 or 32 x 16 MUL operation completes in the array and stalls for one clock in the M stage to complete the carry-propagate-add in the M stage
- A 32 x 32 MUL operation stalls for two clocks in the M stage to complete the second cycle of the array and the carry-propagate-add in the M stage
- Multiply and divide calculations proceed in the MDU. If the calculation completes before the IU moves the instruction past the M stage, the MDU holds the result in a temporary register until the IU moves the instructions to the A stage (and it is consequently known that it will not be killed).

50.3.1.4 A STAGE – ALIGN

During A stage:

- A separate aligner aligns loaded data with its word boundary
- A MUL operation makes the result available for writeback. The actual register writeback is performed in the W stage
- From this stage, load data or a result from the MDU are available in the E stage for bypassing

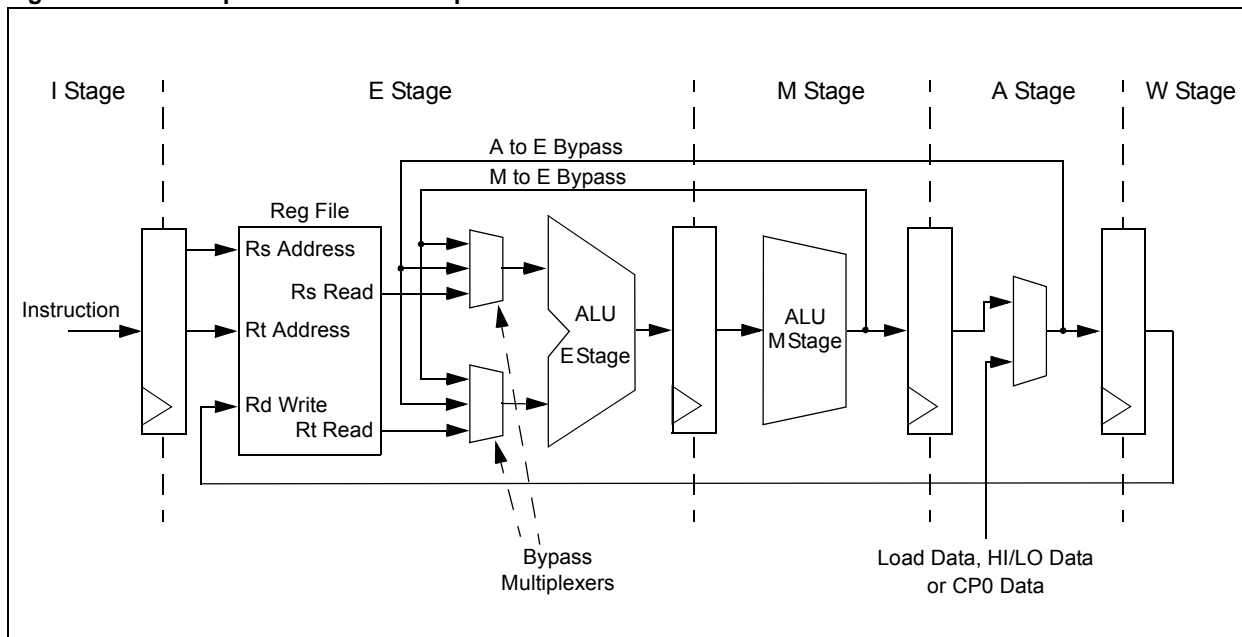
50.3.1.5 W STAGE – WRITEBACK

During W stage:

For register-to-register or load instructions, the result is written back to the register file.

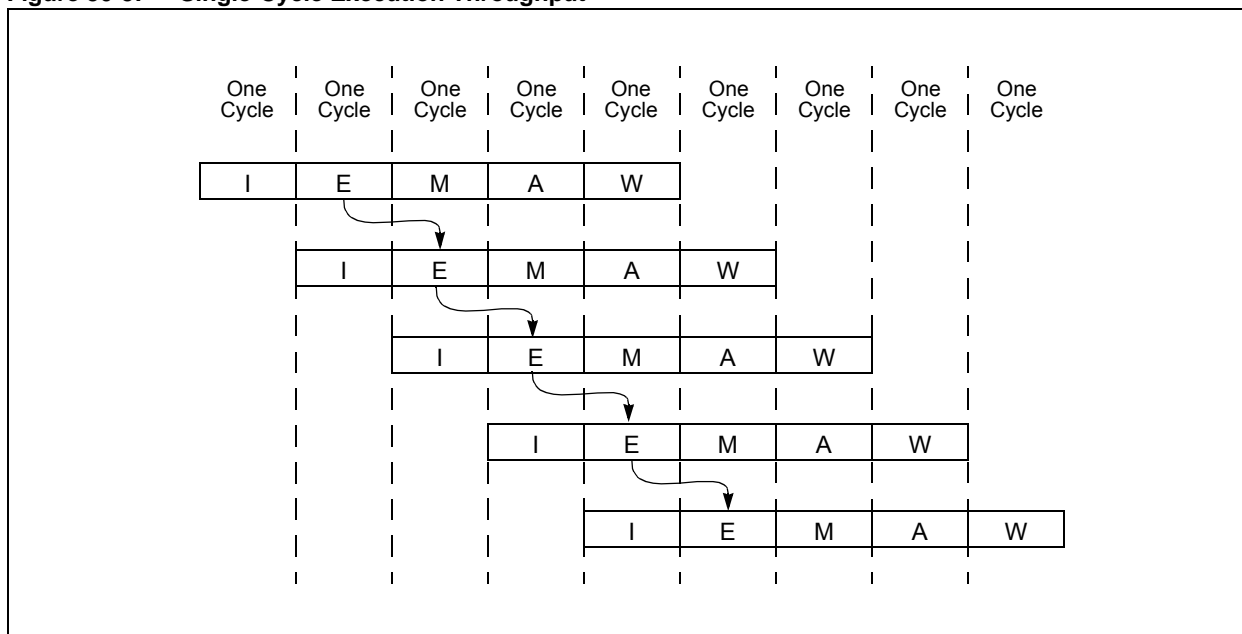
The microprocessor cores implement a “bypass” mechanism that allows the result of an operation to be sent directly to the instruction that needs it without having to write the result to the register, and then read it back.

Figure 50-4: Simplified PIC32 CPU Pipeline



The results of using instruction pipelining in the PIC32 core is a fast, single-cycle instruction execution environment.

Figure 50-5: Single-Cycle Execution Throughput



50.3.2 Execution Unit

The PIC32 Execution Unit is responsible for carrying out the processing of most of the instructions of the MIPS instruction set. The Execution Unit provides single-cycle throughput for most instructions by means of pipelined execution. Pipelined execution, sometimes referred to as “pipelining”, is where complex operations are broken into smaller pieces called stages. Operation stages are executed over multiple clock cycles.

The Execution Unit contains the following features:

- 32-bit adder used for calculating the data address
- Address unit for calculating the next instruction address
- Logic for branch determination and branch target address calculation
- Load aligner
- Bypass multiplexers used to avoid stalls when executing instructions streams where data producing instructions are followed closely by consumers of their results
- Leading Zero/One detect unit for implementing the `CLZ` and `CLO` instructions
- Arithmetic Logic Unit (ALU) for performing bit-wise logical operations
- Shifter and Store Aligner

50.3.3 Multiply/Divide Unit (MDU)

The Multiply/Divide unit (MDU) performs multiply and divide operations. The MDU consists of a 32 x 16 multiplier, result-accumulation registers (HI and LO), multiply and divide state machines, and all multiplexers and control logic required to perform these functions. The high-performance, pipelined MDU supports execution of a 16 x 16 or 32 x 16 multiply operation every clock cycle; 32 x 32 multiply operations can be issued every other clock cycle. Appropriate interlocks are implemented to stall the issue of back-to-back 32 x 32 multiply operations. Divide operations are implemented with a simple 1 bit per clock iterative algorithm and require 35 clock cycles in the worst case to complete. Any attempt to issue a subsequent MDU instruction while a divide is still active causes a pipeline stall until the divide operation is completed.

The microprocessor cores implement an additional multiply instruction, `MUL`, which specifies that lower 32-bits of the multiply result be placed in the register file instead of the HI/LO register pair. By avoiding the explicit move from the LO (`MFLO`) instruction, which required when using the LO register, and by supporting multiple destination registers, the throughput of multiply-intensive operations is increased. Two instructions, multiply-add (`MADD/MADDU`) and multiply-subtract (`MSUB/MSUBU`), are used to perform the multiply-add and multiply-subtract operations. The `MADD` instruction multiplies two numbers, and then adds the product to the current contents of the HI and LO registers. Similarly, the `MSUB` instruction multiplies two operands, and then subtracts the product from the HI and LO registers. The `MADD/MADDU` and `MSUB/MSUBU` operations are commonly used in Digital Signal Processor (DSP) algorithms.

The MDU is a separate pipeline for integer multiply and divide operations and DSP ASE multiply instructions (see **Note**). This pipeline operates in parallel with the integer unit (IU) pipeline and does not stall when the IU pipeline stalls. This allows the long-running MDU operations to be partially masked by system stalls and/or other integer unit instructions. The MDU supports execution of one 32 x 32 multiply or multiply-accumulate operation every clock cycle. The 32-bit divide operation executes in 12-38 clock cycles. The MDU also implements various shift instructions operating on the HI/LO register and multiply instructions as defined in the DSP ASE.

Note: DSP ASE is not available on all devices. Refer to the “CPU” chapter of the specific device data sheet to determine availability
--

50.3.4 Shadow Register Sets

The PIC32 processor implements one or more copies of the General Purpose Registers (GPR) for use by high-priority interrupts. The extra banks of registers are known as shadow register sets. When a high-priority interrupt occurs, the processor automatically switches to a shadow register set without software intervention. This reduces overhead in the interrupt handler and reduces effective latency.

The shadow register sets are controlled by registers located in the System Coprocessor (CP0) as well as the interrupt controller hardware located outside of the CPU core.

For more information on shadow register sets, refer to **Section 8. “Interrupts”** (DS60001108) of the “PIC32 Family Reference Manual”.

50.3.5 Pipeline Interlock Handling

Smooth pipeline flow is interrupted when an instruction in a pipeline stage cannot advance due to a data dependency or a similar external condition. Pipeline interruptions are handled entirely in hardware. These dependencies are referred to as “interlocks”. At each cycle, interlock conditions are checked for all active instructions. An instruction that depends on the result of a previous instruction is an example of an interlock condition.

In general, MIPS processors support two types of hardware interlocks:

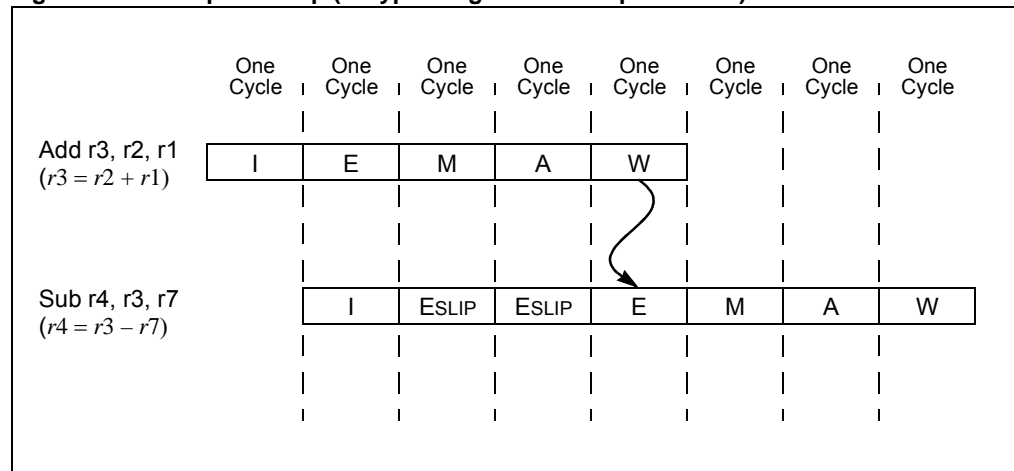
- Stalls – These interlocks are resolved by halting the entire pipeline. All instructions, currently executing in each pipeline stage, are affected by a stall.
- Slips – These interlocks allow one part of the pipeline to advance while another part of the pipeline is held static

In the PIC32 processor core, all interlocks are handled as slips. These slips are minimized by grabbing results from other pipeline stages by using a method called register bypassing, which is described in [50.3.6 “Register Bypassing”](#).

Note: To illustrate the concept of a pipeline slip, the example in [Figure 50-6](#) shows would happen if the PIC32 core did not implement register bypassing.

As shown in [Figure 50-6](#), the sub instruction has a source operand dependency on register r3 with the previous add instruction. The sub instruction slips by two clocks waiting until the result of the add is written back to register r3. This slipping does not occur on the PIC32 family of processors.

Figure 50-6: Pipeline Slip (If Bypassing Was Not Implemented)



50.3.6 Register Bypassing

As mentioned previously, the PIC32 processor implements a mechanism called register bypassing that helps reduce pipeline slips during execution. When an instruction is in the E stage of the pipeline, the operands must be available for that instruction to continue. If an instruction has a source operand that is computed from another instruction in the execution pipeline, register bypassing allows a shortcut to get the source operands directly from the pipeline. An instruction in the E stage can retrieve a source operand from another instruction that is executing in either the M stage or the A stage of the pipeline. As seen in Figure 50-7, a sequence of three instructions with interdependencies does not slip at all during execution. This example uses both A to E, and M to E register bypassing. Figure 50-8 shows the operation of a load instruction utilizing A to E bypassing. Since the result of load instructions are not available until the A pipeline stage, M to E bypassing is not needed.

The performance benefit of register bypassing is that instruction throughput is increased to the rate of one instruction per clock for ALU operations, even in the presence of register dependencies.

Figure 50-7: IU Pipeline M to E Bypass

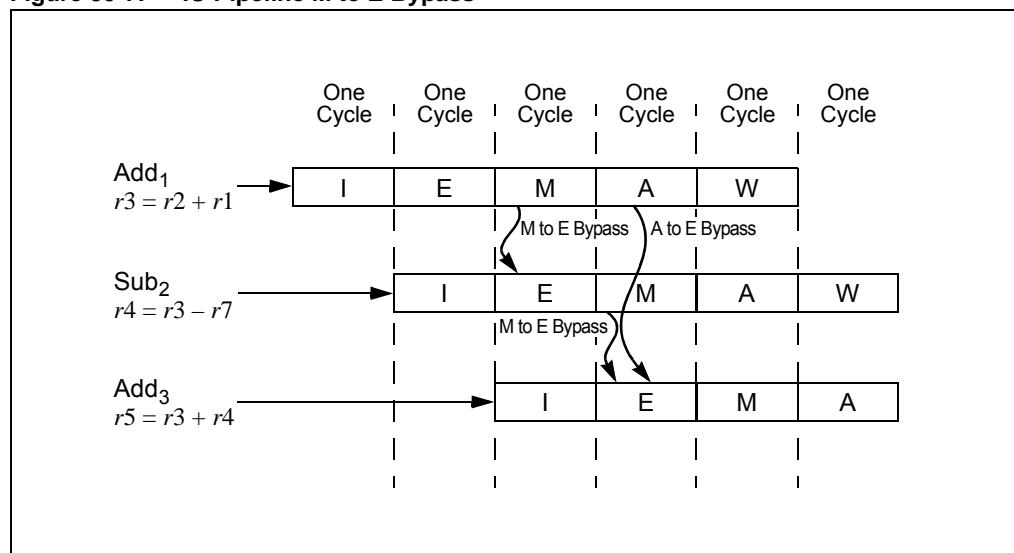
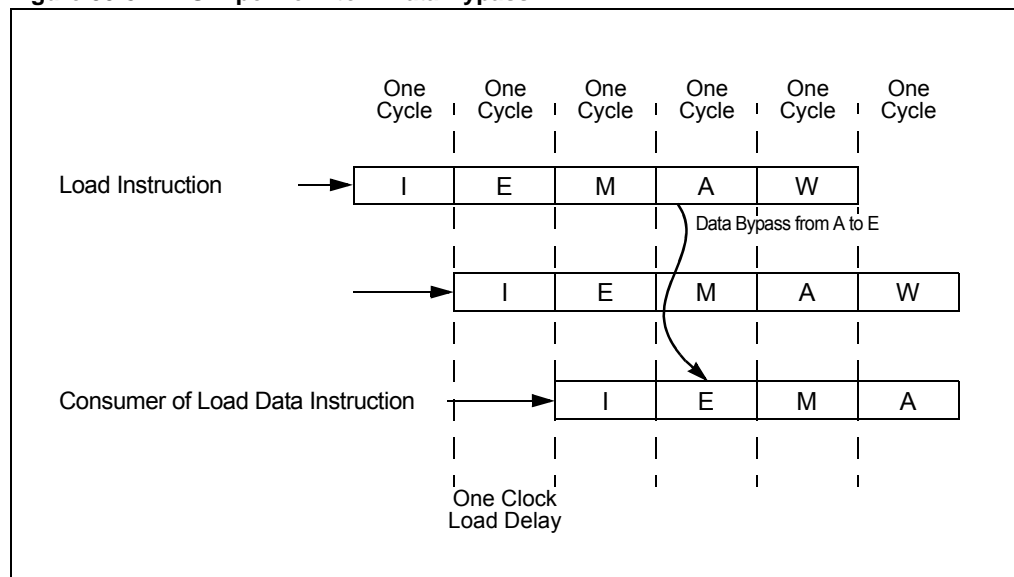


Figure 50-8: IU Pipeline A to E Data Bypass



50.4 SPECIAL CONSIDERATIONS WHEN WRITING TO CP0 REGISTERS

In general, the PIC32 core ensures that instructions are executed following a fully sequential program model. Each instruction in the program sees the results of the previous instruction. There are some deviations to this model. These deviations are referred to as “hazards”.

In privileged software, there are two types of hazards:

- Execution
- Instruction

50.4.1 Execution Hazards

Execution hazards are those created by the execution of one instruction, and seen by the execution of another instruction. [Table 50-2](#) lists the execution hazards.

Table 50-2: Execution Hazards

Created by	Seen by	Hazard On	Spacing (Instructions)
LL	MFC0	LLAddr	1
MTC0	Coprocessor instruction execution depends on the new value of the CU0 bit (Status<28>)	CU0 bit (Status<28>)	1
MTC0	ERET	EPC, DEPC, ErrorEPC	1
MTC0, EI, DI	Interrupted Instruction	IE bit (Status<0>)	1
MTC0	Interrupted Instruction	IP1 and IP0 bits (Cause<1> and <0>)	3
MTC0	TLBR, TLBWI, TLBWR	EntryHi	1
MTC0	TLBP, Load/Store affected by new state	ASID<7:0> bits (EntryHi<7:0>)	1
MTC0	TLBWI, TLBWR	Index	1
MTC0	RDPGPR, WRPGPR	PSS<3:0> bits (SRSCtl<9:6>)	1
MTC0	Instruction is not seeing a core timer interrupt	Compare update that clears the core timer Interrupt	4
MTC0	Instruction affected by change	Any other CP0 register	2

50.4.2 Instruction Hazards

Instruction hazards are those created by the execution of one instruction, and seen by the instruction fetch of another instruction. [Table 50-3](#) lists the instruction hazards.

Table 50-3: Instruction Hazards

Created by	Seen by	Hazard On	Spacing (Instructions)
TLBWR, TLBWI	Instruction fetch using new TLB entry	TLB entry	3
MTC0	Instruction fetch seeing the new value (including a change to ERL followed by an instruction fetch from the useg segment)	Status	
MTC0	Instruction fetch seeing the new value	ASID<7:0> bits (EntryHi<7:0>)	3
MTC0	Instruction fetch seeing the new value	WatchHi and WatchLo	1
MTC0	Interrupted instruction	IP1 and IP0 bits (Cause<1> and <0>)	2
Instruction stream write via cache	Instruction fetch seeing the new instruction stream	Cache entries	3
Instruction stream write via store	Instruction fetch seeing the new instruction stream	Cache entries	System dependent

50.5 MIPS32 ARCHITECTURE

The MIPS32 architecture is based on a fixed-length, regularly encoded instruction set and uses a load/store data model. The architecture is streamlined to support optimized execution of high-level languages. Arithmetic and logic operations use a three-operand format, allowing compilers to optimize complex expressions formulation. Availability of 32 general-purpose registers enables compilers to further optimize code generation for performance by keeping frequently accessed data in registers.

For more information and documentation, refer to the MIPS32 Architecture web page at:

<http://www.imgtec.com/mips/architectures/mips32.asp>

50.5.1 Architecture Release 2

PIC32 devices with the microAptiv core utilize Release 2 of the MIPS32 processor architecture, and implement the following features:

- Vectored interrupts using an external-to-core interrupt controller, which provides the ability to vector interrupts directly to a handler for that interrupt
- Programmable exception vector base, which allows the base address of the exception vectors to be moved for exceptions that occur when `Status_BEV` is '0'. This feature enables any system to place the exception vectors in memory that is appropriate to the system environment.
- Atomic interrupt enable/disable, which includes two added instructions to atomically enable or disable interrupts, and return the previous value of the Status register
- The ability to disable the Count register for highly power-sensitive applications
- GPR shadow registers, which provide the addition of GPR shadow registers and the ability to bind these registers to a vectored interrupt or exception
- Field, Rotate, and Shuffle instructions, which add additional capability in processing bit fields in registers
- Explicit hazard management, which provides a set of instructions to explicitly manage hazards, in place of the cycle-based SSNOP method of dealing with hazards

50.5.2 Architecture Release 5

PIC32 devices with the M-Class core utilize all of the features of Release 2, as well as the following Release 5 features:

- User mode access through the UFR bit in the Config5 Register (CP0 Register 16, Select 5)
- Additional MCU ASE instructions: `ASET` and `ACLR` for setting and clearing atomic 8-bit memory locations

50.6 CPU BUS

The PIC32 devices use two different CPU bus architectures, either split-bus or data/instruction bus, depending on which CPU core is implemented.

50.6.1 Split-bus Architecture

PIC32 devices based on the MCU Microprocessor core have a Split-bus architecture, with two distinct busses to provide parallel instruction and data operations. Load and store operations occur simultaneously as instruction fetches. The two busses are known as the I-side bus, which is used for feeding instructions into the CPU, and the D-side bus, which is used for data transfers.

In the split-bus architecture, the CPU fetches instructions during the I-pipeline stage. A fetch is issued to the I-side bus and is handled by the System Bus. Depending on the address, the System Bus will do one of the following:

- Forward the fetch request to the Prefetch Cache unit (if available)
- Forward the fetch request to the DRM unit, or
- Cause an exception

Instruction fetches always use the I-side bus independent of the addresses being fetched.

The D-side bus processes all load and store operations executed by the CPU. When a load or store instruction is executed, the request is routed to the System Bus by the D-side bus. This operation occurs during the M-pipeline stage and is routed to one of several targets:

- Data RAM
- Prefetch Cache/Flash memory
- Fast Peripheral Bus (Interrupt controller, DMA, Debug unit, USB, Ethernet, GPIO ports)
- General Peripheral Bus (UART, SPI, Flash Controller, EPMP/EPSP, TRCC Timers, Input Capture, PWM/Output Compare, ADC, Dual Compare, I²C, Clock SIB, and Reset SIB)

50.6.2 Data/Instruction Architecture

PIC32 devices based on the MPU core have a unified Data or Instruction bus connected to the System Bus. This architecture uses a multi-layer System Bus to provide multiple simultaneous data transactions between bus initiators and targets.

50.7 INTERNAL SYSTEM BUSSES

The internal busses of the PIC32 processor connect the peripherals to the System Bus. The System Bus routes bus accesses from different initiators to a set of targets utilizing several data paths throughout the device to help eliminate performance bottlenecks.

Some of the paths that the System Bus uses serve a dedicated purpose, while others are shared between several targets.


The data RAM and Flash memory read paths are dedicated paths, allowing low-latency access to the memory resources without being delayed by peripheral bus activity. The high-bandwidth peripherals are placed on a high-speed bus. These include the Interrupt controller, Debug unit, DMA engine, the USB Host/Peripheral unit, and other high-bandwidth peripherals (i.e., CAN, Ethernet engines).

Peripherals that do not require high-bandwidth are located on a separate peripheral bus to save power.

50.8 SET/CLEAR/INVERT

To provide single-cycle bit operations on peripherals, the registers in the peripheral units can be accessed in three different ways depending on peripheral addresses. Each register has four different addresses. Although the four different addresses appear as different registers, they are really just four different methods to address the same physical register.

Figure 50-9: Four Addresses for a Single Physical Register



The base register address provides normal Read/Write access, while the other three provide special write-only functions.

- Normal access
- Set bit atomic RMW access
- Clear bit atomic RMW access
- Invert bit atomic RMW access

Peripheral reads must occur from the base address of each peripheral register. Reading from a Set/Clear/Invert address has an undefined meaning, and may be different for each peripheral.

Writing to the base address writes an entire value to the peripheral register. All bits are written. For example, assume a register contains 0xAAAA5555 before a write of 0x000000FF. After the write, the register will contain 0x000000FF (assuming that all bits are R/W bits).

Writing to the Set address for any peripheral register causes only the bits written as '1's to be set in the destination register. For example, assume that a register contains 0xAAAA5555 before a write of 0x000000FF to the set register address. After the write to the Set register address, the value of the peripheral register will contain 0xAAAA55FF.

Writing to the Clear address for any peripheral register causes only the bits written as '1's to be cleared to '0's in the destination register. For example, assume that a register contains 0xAAAA5555 before a write of 0x000000FF to the Clear register address. After the write to the Clear register address, the value of the peripheral register will contain 0xAAAA5500.

Writing to the Invert address for any peripheral register causes only the bits written as '1's to be inverted, or toggled, in the destination register. For example, assume that a register contains 0xAAAA5555 before a write of 0x000000FF to the invert register address. After the write to the Invert register, the value of the peripheral register will contain 0xAAAA55AA.

50.9 ALU STATUS BITS

Unlike most other PIC microcontrollers, the PIC32 processor does not use Status register flags. Condition flags are used on many processors to help perform decision making operations during program execution. Flags are set based on the results of comparison operations or some arithmetic operations. Conditional branch instructions on these machines then make decisions based on the values of the single set of condition codes.

Instead, the PIC32 processor uses instructions that perform a comparison and stores a flag or value into a General Purpose Register. A conditional branch is then executed with this general purpose register used as an operand.

50.10 INTERRUPT AND EXCEPTION MECHANISM

Note: In this section, the terms “precise” and “imprecise” are used to describe exceptions. A precise exception is one in which the EPC (COP0, Register 14, Select 0) can be used to identify the instruction that caused the exception. For imprecise exceptions, the instruction that caused the exception cannot be identified. Most exceptions are precise. Bus error exceptions may be imprecise.

The PIC32 family of processors implement an efficient and flexible interrupt and exception handling mechanism. Interrupts and exceptions both behave similarly in that the current instruction flow is changed temporarily to execute special procedures to handle an interrupt or exception. The difference between the two is that interrupts are usually a result of normal operation, and exceptions are a result of error conditions such as bus errors.

When an interrupt or exception occurs, the processor does the following:

1. The PC of the next instruction to execute after the handler returns is saved into a coprocessor register.
2. The Cause register is updated to reflect the reason for exception or interrupt.
3. The Status register EXL or ERL bit is set to cause Kernel mode execution.
4. Handler PC is calculated from Ebase and OFFSET values.
5. Automated Interrupt Epilogue can save some of the COP0 state in the stack and automatically update some of the COP0 registers in preparation for interrupt handling.
6. Processor starts execution from new PC.

This is a simplified overview of the interrupt and exception mechanism. Refer to the “**CPU Exceptions and Interrupt Controller**” chapter in the specific device data sheet for details.

50.11 PROGRAMMING MODEL

The PIC32 family of processors is designed to be used with a high-level language such as the C programming language. It supports several data types and uses simple but flexible addressing modes needed for a high-level language. There are 32 General Purpose Registers and two special registers for multiplying and dividing.

There are three different formats for the machine language instructions on the PIC32 processor:

- Immediate or I-type CPU instructions
- Jump or J-type CPU instructions, and
- Registered or R-type CPU instructions

Most operations are performed in registers. The register type CPU instructions have three operands; two source operands and a destination operand.

Having three operands and a large register set allows assembly language programmers and compilers to use the CPU resources efficiently. This creates faster and smaller programs by allowing intermediate results to stay in registers rather than constantly moving data to and from memory.

The immediate format instructions have an immediate operand, a source operand and a destination operand. The jump instructions have a 26-bit relative instruction offset field that is used to calculate the jump destination.

50.11.1 CPU Instruction Formats

A CPU instruction is a single 32-bit aligned word. The CPU instruction formats are:

- Immediate (see [Figure 50-10](#))
- Jump (see [Figure 50-11](#))
- Register (see [Figure 50-12](#))

[Table 50-4](#) describes the fields used in these instructions.

Table 50-4: CPU Instruction Format Fields

Field	Description
opcode	6-bit primary operation code.
rd	5-bit specifier for the destination register.
rs	5-bit specifier for the source register.
rt	5-bit specifier for the target (source/destination) register or used to specify functions within the primary opcode REGIMM.
immediate	16-bit signed immediate used for logical operands, arithmetic signed operands, load/store address byte offsets, and PC-relative branch signed instruction displacement.
instr_index	26-bit index shifted left two bits to supply the low-order 28 bits of the jump target address.
sa	5-bit shift amount.
function	6-bit function field used to specify functions within the primary opcode SPECIAL.

Figure 50-10: Immediate (I-Type) CPU Instruction Format




Figure 50-11: Jump (J-Type) CPU Instruction Format





Figure 50-12: Register (R-Type) CPU Instruction Format



50.11.2 CPU Registers

The PIC32 architecture defines the following CPU registers:

- Thirty-two 32-bit General Purpose Registers (GPRs)
- The standard MIPS32 architecture defines one pair of HI/LO accumulator registers (AC0). The cores in PIC32 devices include the DSP ASE (see **Note**), which provides three additional pairs of HI/LO accumulator registers (AC1, AC2, and AC3). These registers improve the parallelization of independent accumulation routines. DSP instructions that target the accumulators use two instruction bits to specify the destination accumulator.
- A special purpose program counter (PC), which is affected only indirectly by certain instructions; it is not an architecturally visible register.

Note: DSP ASE is not available on all devices. Please consult the “CPU” chapter of the specific device data sheet to determine availability

50.11.2.1 CPU GENERAL PURPOSE REGISTERS

Two of the CPU General Purpose Registers have assigned functions:

- r0 – This register is hard-wired to a value of ‘0’, and can be used as the target register for any instruction the result of which will be discarded. r0 can also be used as a source when a ‘0’ value is needed.
- r31 – This is the destination register used by JAL, BLTZAL, BLTZALL, BGEZAL, and BGEZALL, without being explicitly specified in the instruction word; otherwise, r31 is used as a normal register.

The remaining registers are available for general purpose use.

50.11.2.2 REGISTER CONVENTIONS

Although most of the registers in the PIC32 architecture are designated as General Purpose Registers, as shown in [Table 50-5](#), there are some recommended uses of the registers for correct software operation with high-level languages such as the Microchip MPLAB® XC32 C/C++ compiler.

Table 50-5: Register Conventions

CPU Register	Symbolic Register	Usage
r0	zero	Always ‘0’ (see Note 1)
r1	at	Assembler Temporary
r2 - r3	v0-v1	Function Return Values
r4 - r7	a0-a3	Function Arguments
r8 - r15	t0-t7	Temporary – Caller does not need to preserve contents
r16 - r23	s0-s7	Saved Temporary – Caller must preserve contents
r24 - r25	t8-t9	Temporary – Caller does not need to preserve contents
r26 - r27	k0-k1	Kernel temporary – Used for interrupt and exception handling
r28	gp	Global Pointer – Used for fast-access common data
r29	sp	Stack Pointer – Software stack
r30	s8 or fp	Saved Temporary – Caller must preserve contents <i>OR</i> Frame Pointer – Pointer to procedure frame on stack
r31	ra	Return Address (see Note 1)

Note 1: Hardware enforced, not just convention.

50.11.2.3 CPU SPECIAL PURPOSE REGISTERS

The CPU contains these special purpose registers:

- PC – Program Counter register
- AC0 through AC3 – 64-bit Accumulator register pairs (HI/LO):
 - HI/LO – Multiply and divide register pair (high and low result):
 - During a multiply operation, the HI and LO registers store the product of integer multiply
 - During a multiply-add or multiply-subtract operation, the HI and LO registers store the result of the integer multiply-add or multiply-subtract
 - During a division, the HI and LO registers store the quotient (in LO) and remainder (in HI) of integer divide
 - During a multiply-accumulate, the HI and LO registers store the accumulated result of the operation

Figure 50-13 shows the layout of the CPU registers.

Figure 50-13: CPU Registers

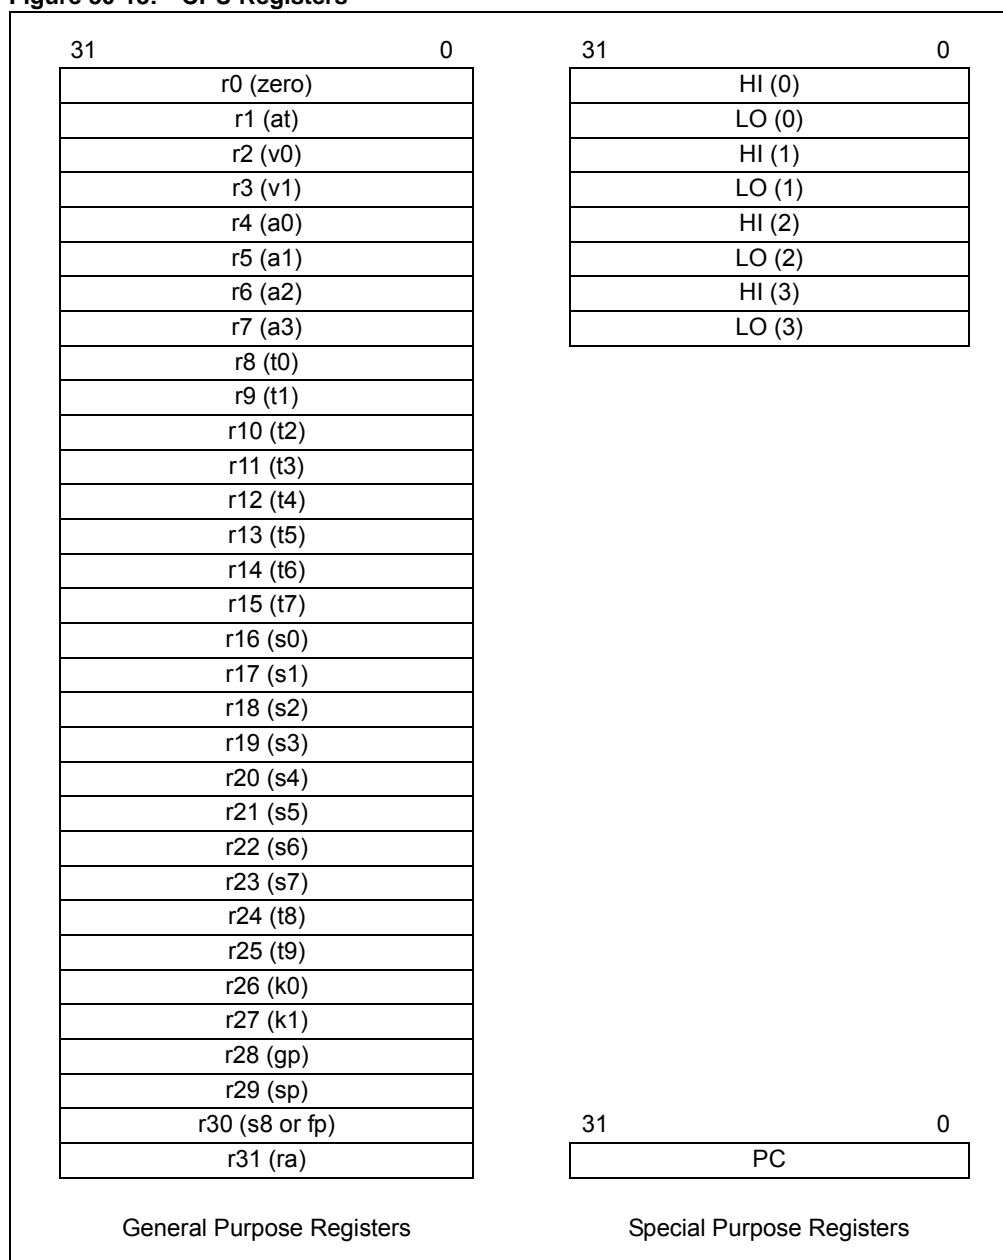


Table 50-6: microMIPS 16-bit Instruction Register Usage

16-bit Register Encoding	32-bit MIPS Register Encoding	Symbolic Name	Description
0	16/0	s0/zero	General-purpose register
1	17	s1	General-purpose register
2	2	v0	General-purpose register
3	3	v1	General-purpose register
4	4	a0	General-purpose register
5	5	a1	General-purpose register
6	6	a2	General-purpose register
7	7	a3	General-purpose register
N/A	28	gp	microMIPS implicitly referenced General-pointer register
N/A	29	sp	microMIPS implicitly referenced Stack pointer register
N/A	31	ra	microMIPS implicitly referenced Return address register

Table 50-7: microMIPS Special Registers

Symbolic Name	Purpose
PC	Program counter. The PC-relative instructions can access this register as an operand.
HI	Contains high-order word of multiply or divide result.
LO	Contains low-order word of multiply or divide result.

50.11.3 How to Implement Stack/MIPS Calling Conventions

The PIC32 CPU does not have hardware stacks. Instead, the processor relies on software to provide this functionality. Since the hardware does not perform stack operations itself, a convention must exist for all software within a system to use the same mechanism. For example, a stack can grow either toward lower addresses, or grow toward higher addresses. If one piece of software assumes that the stack grows toward a lower address, and calls a routine that assumes that the stack grows toward a higher address, the stack would become corrupted.

Using a system-wide calling convention prevents this problem from occurring. The Microchip MPLAB® XC32 C/C++ Compiler assumes the stack grows toward lower addresses.

50.11.4 Processor Modes

There are two operational modes and one special mode of execution in the PIC32 family CPUs: User mode, Kernel mode and Debug mode. The processor starts execution in Kernel mode, and if desired, can stay in Kernel mode for normal operation. User mode is an optional mode that allows a system designer to partition code between privileged and unprivileged software. Debug mode is normally only used by a debugger or monitor.

One of the main differences between the modes of operation is the memory addresses that software is allowed to access. Peripherals are not accessible in User mode. [Figure 50-14](#) shows the different memory maps for each mode. For more information on the processor's memory map, refer to **Section 3. "Memory Organization"** (DS60001115) of the *"PIC32 Family Reference Manual"*.

Figure 50-14: CPU Modes

Virtual Address	User Mode	Kernel Mode	Debug Mode
0xFFFF_FFFF			kseg3
0xFF40_0000			dseg
0xFF3F_FFFF		kseg3	kseg3
0xFF20_0000			
0xFF1F_FFFF		kseg2	kseg2
0xE000_0000			
0xDFFF_FFFF			
0xC000_0000		kseg1	kseg1
0xBFFF_FFFF			
0xA000_0000		kseg0	kseg0
0x9FFF_FFFF			
0x8000_0000			
0x7FFF_FFFF			
	useg	kuseg	kuseg
0x0000_0000			

50.11.4.1 KERNEL MODE

To access many of the hardware resources, the processor must be operating in Kernel mode. Kernel mode gives software access to the entire address space of the processor as well as access to privileged instructions.

The processor operates in Kernel mode when the DM bit in the Debug register is '0' and the Status register contains one, or more, of the following values:

- UM = 0
- ERL = 1
- EXL = 1

When a non-debug exception is detected, EXL or ERL will be set and the processor will enter Kernel mode. At the end of the exception handler routine, an Exception Return (ERET) instruction is generally executed. The ERET instruction jumps to the Exception PC (EPC or ErrorPC depending on the exception), clears ERL, and clears EXL if ERL = 0.

If UM = 1 the processor will return to User mode after returning from the exception when ERL and EXL are cleared back to '0'.

50.11.4.2 USER MODE

When executing in User mode, software is restricted to use a subset of the processor's resources. In many cases it is desirable to keep application-level code running in User mode where if an error occurs it can be contained and not be allowed to affect the Kernel mode code.

Applications can access Kernel mode functions through controlled interfaces such as the SYSCALL mechanism.

As seen in [Figure 50-14](#), User mode software has access to the USEG memory area.

To operate in User mode, the Status register must contain each the following bit values:

- UM = 1
- EXL = 0
- ERL = 0

50.11.4.3 DEBUG MODE

Debug mode is a special mode of the processor normally only used by debuggers and system monitors. Debug mode is entered through a debug exception and has access to all Kernel mode resources as well as special hardware resources used to debug applications.

The processor is in Debug mode when the DM bit in the Debug register is '1'.

Debug mode is normally exited by executing a DERET instruction from the debug handler.

50.12 FLOATING POINT UNIT (FPU)

PIC32 devices with the M-Class core contain a Floating Point Unit (FPU) that implements the MIPS Release 5 Instruction Set Architecture for floating point computation.

Note: This module is not available on all devices. Refer to the “CPU” chapter in the specific device data sheet to determine availability.

50.12.1 Features

Some of the most important features of this module include:

- The PIC32 implementation supports the “*IEEE Standard for Binary Floating-Point Arithmetic*” (ANSI/IEEE 754 Standard) for single and double precision data formats. See [50.12.6.5 “IEEE 754-1985 Standard”](#) for more information.
- Full 64-bit operation is implemented in both the register file and functional units. The FPU has 32 64-bit floating point registers used for all of the floating point operations.
- A 32-bit Floating Point Control Register controls the operation of the FPU, and monitors condition codes and exception conditions
- The performance of the unit is optimized for single precision formats. Most instructions have one FPU cycle throughput and four FPU cycle latency.
- The FPU implements compound multiply-add (MADD) and multiply-sub (MSUB) instructions with intermediate rounding after the multiply function. The result is guaranteed to be the same as executing a MUL followed by an ADD/SUB instruction, but the instruction latency, instruction fetch, dispatch bandwidth, and the total number of register accesses is improved.
- IEEE denormalized input operands and results are supported by hardware for some instructions. A fast flush-to-zero mode is provided to optimize performance for IEEE denormalized results. The fast flush-to-zero mode has to be enabled through the FPU control registers, and use of this mode is recommended for best performance when denormalized results are generated.
- Additional arithmetic operations not specified by IEEE 754 Standard (for example, reciprocal and reciprocal square root) are specified by the MIPS[®] architecture (see **Note**) and are implemented by the FPU. To achieve low latency counts, these instructions satisfy more relaxed precision requirements.

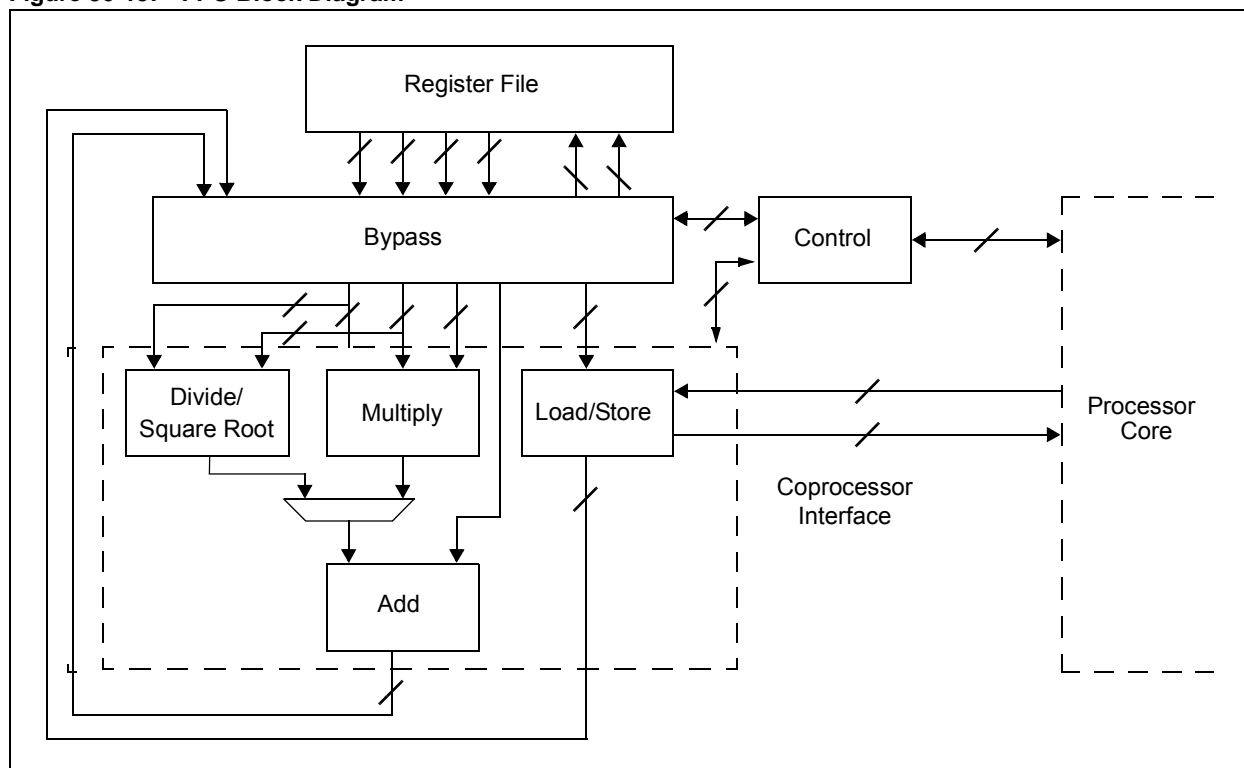
Note: Refer to the Imagination Technologies Ltd. website, www.imgtec.com, for information on the MIPS architecture.

- The MIPS FPU architecture is designed such that a combination of hardware and software can be used to implement the architecture. The PIC32 FPU can operate on numbers within a specific range (the IEEE normalized numbers), but it relies on a software handler to operate on numbers not handled by the FPU hardware (the IEEE denormalized numbers).
- The FPU has a separate pipeline for floating point instruction execution. This pipeline operates in parallel with the integer core pipeline and does not stall when the integer pipeline stalls. This allows long-running FPU operations, such as divide or square root, to be partially masked by system stalls and/or other integer unit instructions.
- The FPU access is provided through Coprocessor 1. Like the main processor core, Coprocessor 1 is programmed and operated using a Load/Store instruction set. The processor core communicates with Coprocessor 1 using a dedicated coprocessor interface. The FPU functions as an autonomous unit. The hardware is completely interlocked such that, when writing software, the programmer does not have to worry about inserting delay slots after loads and between dependent instructions.
- Arithmetic instructions are always dispatched and completed in order, but loads and stores can complete out of order. The exception model is ‘precise’ at all times.

Refer to [50.14 “Coprocessor 1 \(CP1\) Registers”](#) for information on the related FPU registers.

[Figure 50-15](#) shows a block diagram of the PIC32 FPU.

Figure 50-15: FPU Block Diagram



50.12.2 FPU data formats

The FPU supports the single-precision and double-precision floating point data types as specified by the IEEE 754 Standard.

In addition, fixed point data types are supported: signed integers that are provided by the MIPS architecture.

50.12.2.1 FLOATING POINT FORMATS

The PIC32 FPU supports the following two floating point formats:

- a 32-bit single-precision floating point (type S, shown in [Figure 50-16](#))
- a 64-bit double-precision floating point (type D, shown in [Figure 50-17](#))

The single and double floating point data types are composed of three fields: sign, exponent, fraction.

They represent numeric values as well as the following special entities:

- Two infinities: $+\infty$ and $-\infty$
- Signaling non-numbers (SNaNs)
- Quiet non-numbers (QNaNs)
- Numbers of the form: $(-1)^s 2^E b_0.b_1 b_2..b_p-1$, where:
 - $s = 0$ or 1
 - $E =$ any integer between E_{min} and E_{max} , inclusive
 - $b_i = 0$ or 1 (the high bit, b_0 , is to the left of the binary point)
 - p is the signed-magnitude precision

The sizes for the single and double precision numbers supported by the architecture are listed in [Table 50-8](#).

Table 50-8: Parameters of Floating Point Data Types

Parameter	Single	Double
Bits of mantissa precision, p	24	53
Maximum exponent, E_max	+127	+1023
Minimum exponent, E_min	-126	-1022
Exponent bias	+127	+1023
Bits in exponent field, e	8	11
Representation of b0 integer bit	Hidden	Hidden
Bits in fraction field, f	23	52
Total format width in bits	32	64
Magnitude of largest representable number	3.4028234664e+38	1.7976931349e+308
Magnitude of smallest normalized representable number	1.1754943508e-38	2.2250738585e-308

Figure 50-16: Single-Precision Floating Point Format (S)

S	Exponent <0:7>	Fraction <0:22>
---	----------------	-----------------

Figure 50-17: Double-Precision Floating Point Format (D)

S	Exponent <0:10>	Fraction <0:51>
---	-----------------	-----------------

The fields in the [Figure 50-16](#) and [Figure 50-17](#) are:

- 1-bit sign, S
- Biased exponent, e = E + bias
- Binary fraction, f = .b1 b2..bp-1 (the b0 bit is hidden; it is not recorded)

Values are encoded in the specified format using the unbiased exponent, fraction, and sign values listed in [Table 50-9](#).

The high-order bit of the Fraction field, identified as b1, has also special importance for NaNs.

Table 50-9: Single or Double Floating Point Data Type Encoding

Unbiased E	f	s	b1	Value V	Type of Value	Value of Typical Single Bit Pattern (see Note 1)	Value of Typical Double Bit Pattern
E_max+1	≠ 0		1	SNaN	Signaling NaN (FCSR _{NAN2008} = 0)	0x7FFFFFFF	0x7FFFFFFF FFFFFFFF
			0	QNaN	Quiet NaN (FCSR _{NAN2008} = 0)	0x7FBFFFFFF	0x7FF7FFFF FFFFFFFF
E_max+1	≠ 0		1	QNaN	Quiet NaN (FCSR _{NAN2008} = 1)	0x7FFFFFFF	0x7FFFFFFF FFFFFFFF
			0	SNaN	Signaling NaN (FCSR _{NAN2008} = 1)	0x7FBFFFFFF	0x7FF7FFFF FFFFFFFF
E_max+1	0	1		-∞	Minus infinity	0xFF800000	0xFFF00000 00000000
		0		+∞	Plus infinity	0x7F800000	0x7FF00000 00000000
E_max to E_min		1		-(2 ^E)(1.f)	Negative normalized number	0x80800000 through 0xFF7FFFFF	0x80100000 00000000 through 0xFFEFFFFF FFFFFFFF
		0		+(2 ^E)(1.f)	Positive normalized number	0x00800000 through 0x7F7FFFFF	0x00100000 00000000 through 0x7FEFFFFF FFFFFFFF
E_min-1	≠ 0	1		-(2 ^{E_min})(0.f)	Negative denormalized number	0x807FFFFF	0x800FFFFF FFFFFFFF
		0		+(2 ^{E_min})(0.f)	Positive denormalized number	0x007FFFFF	0x000FFFFF FFFFFFFF
E_min-1	0	1		-0	Negative zero	0x80000000	0x80000000 00000000
		0		+0	Positive zero	0x00000000	0x00000000 00000000

Note 1: The “Typical” nature of the bit patterns for the NaN and denormalized values reflects the fact that the sign might have either value (NaN) and that the fraction field might have any non-zero value (both). As such, the bit patterns shown are one value in a class of potential values that represent these special values.

50.12.2.1.1 Normalized and Denormalized Numbers

There is just one encoding for each nonzero numerical value that could be represented as a single or double data type. This is called the normalized form.

The high-order bit of the p-bit mantissa, which lies to the left of the binary point, is “hidden,” and not recorded in the Fraction field. The value of this bit can be determined by looking at the value of the exponent:

- When the unbiased exponent is in the range E_min to E_max, inclusive, the number is normalized and the hidden bit must be ‘1’.
- If the numeric value cannot be normalized because the exponent is less than E_min, the representation is denormalized, the encoded number has an exponent of E_min – 1, and the hidden bit has the value ‘0’.

Please note that plus and minus zero are special cases that are not regarded as denormalized values.

50.12.2.1.2 Infinity and NaNs

A floating point operation under certain conditions, such as not using initialized variables, violations of mathematical rules, or results that cannot be represented, can signal exception conditions that are part of the IEEE 754 Standard.

Usually a program will take a trap when an exception condition is encountered. However a different approach is possible where a computation that encounters any of these conditions proceeds without trapping, but generates a result indicating that an exceptional condition arose during the computation. To allow this behavior each floating point format defines representations for plus infinity (+∞), minus infinity (-∞), quiet non-numbers (QNaN), and signaling non-numbers (SNaN) as required by the IEEE 754 Standard. See Table 50-9 for these values.

50.12.2.2 INFINITY ARITHMETIC

Infinity represents a number with magnitude too large to be represented in the given format. During a computation it represents a magnitude overflow. A correctly signed $+\infty$ or $-\infty$ will be generated as the default result in division by zero operations and some cases of overflow as described in [50.12.5 “Floating Point Exceptions Overview”](#).

When created as a default result, ∞ can become an operand in a subsequent operation. The ordering is such that $-\infty < (\text{every finite number}) < +\infty$. Arithmetic with ∞ is the limiting case of real arithmetic with operands of arbitrarily large magnitude, when such limits exist. In these cases, arithmetic on ∞ is regarded as exact, and exception conditions do not arise. The out-of-range indication represented by ∞ is propagated through subsequent computations.

For some cases, there is no meaningful limiting case in real arithmetic for operands of ∞ . These cases raise the Invalid Operation exception condition as described in [50.12.3 “General Floating Point Registers”](#).

50.12.2.2.1 Signaling Non-Number (SNaN)

SNaN operands cause an Invalid Operation exception for arithmetic operations. SNaNs are useful values to put in uninitialized variables. A SNaN is never produced as a result value.

The MIPS architecture makes the formatted operand move instructions (MOV.fmt, MOVT.fmt, MOVF.fmt, MOVN.fmt, MOVZ.fmt) non-arithmetic; they do not signal IEEE 754 Standard exceptions.

50.12.2.2.2 Quiet Non-Number (QNaN)

QNaNs provide diagnostic information propagated from invalid or unavailable data and results. This propagation requires that the information contained in a QNaN be preserved through arithmetic operations and floating point format conversions.

Arithmetic operations with QNaN operands do not signal an exception. When a floating point result is to be delivered, a QNaN operand causes an arithmetic operation to supply a QNaN result. When possible, this QNaN result is one of the operand QNaN values.

QNaNs have similar effects to SNaNs on operations that do not deliver a floating point result (i.e., comparison operations).

When certain invalid operations not involving QNaN operands are performed and the trap is not enabled, a new QNaN value is created. [Table 50-10](#) shows the QNaN value generated when no input operand QNaN value can be copied. The values listed for the fixed point formats are the values supplied to satisfy the IEEE 754 Standard when a QNaN or infinite floating point value is converted to fixed point. There is no other feature of the architecture that detects or utilizes these “integer QNaN” values.

Table 50-10: Value Supplied When a New QNaN is Created

Format	New QNaN Value (FCSR _{NAN2008} = 0)	New QNaN Value (FCSR _{NAN2008} = 1)
Single floating point	0x7FBF FFFF	0x7FFF FFFF
Double floating point	0x7FF7 FFFF FFFF FFFF	0x7FFF FFFF FFFF FFFF
Word fixed point	0x7FFF FFFF	0x7FFF FFFF
Long word fixed point	0x7FFF FFFF FFFF FFFF	0x7FFF FFFF FFFF FFFF

50.12.2.3 FIXED POINT FORMATS

The PIC32 FPU provides two fixed point data types which are the signed integers that are provided by the MIPS architecture:

- 32-bit Word Fixed Point Format (type W)
- 64-bit Long Word Fixed Point Format (type L)

The fixed point values are held in 2's complement format, which is used for signed integers in the CPU. Unsigned fixed point data types are not provided by the architecture; application software can synthesize computations for unsigned integers from the existing instructions and data types.

50.12.3 General Floating Point Registers

The FPU is a 64b floating point unit with 64-bit Floating Point General Registers (FPRs) but a 32b register mode for MIPS backwards compatibility is also supported. The FR bit (CP0<26>) in the CP0 Status register determines which mode is selected:

- When the FR bit is a '1', the FPU is in FR64 mode and the 64b register model is used, which defines 32 64-bit registers with all formats supported in a register.
- When the FR bit is a '0', the FPU is in FR32 mode and the 32b register model is used, which defines 32 32-bit registers with double format values stored in even-odd pairs of registers; The register file can also be viewed as having 16 64-bit registers. When configured this way, there are several restrictions for double operations:
 - Any double operations which specify an odd register as a source or destination will cause a Reserved Instruction exception
 - MTHC1/MFHC1 instructions which access an odd FPU register will signal a Reserved Instruction exception.

50.12.3.1 FPR REGISTERS AND FORMATTED OPERAND LAYOUT

The FPU instructions that operate on formatted operand values specify the Floating Point Register (FPR) that holds the value. Operands that are only 32 bits wide (W and S formats) use only half the space in an FPR. See [Figure 50-18](#) and [Figure 50-19](#) for the FPR organization and the way that operand data is stored in them.

Figure 50-18: Single Floating Point or Word Fixed Point Operand in an FPR

FPR<32:63> - Undefined/Unused	FPR<0:31> - Data Word
-------------------------------	-----------------------

Figure 50-19: Double Floating Point or Long Word Fixed Point Operand in an FPR

FPR<0:63> - Data Double Word/Long Word
--

50.12.3.2 FORMATS OF VALUES USED IN FPR REGISTERS

Unlike the CPU, the FPU neither interprets the binary encoding of source operands nor produces a binary encoding of results for every operation. The value held in a FPR has a format, or type, and it can be used only by instructions that operate on that format. The format of a value in an FPR is one of the valid numeric formats: single or double floating point, and word or long fixed point. Otherwise the value is either uninterpreted or unknown.

The value in an FPR is always set when a value is written to the register as follows:

- When a data transfer instruction writes binary data into an FPR (a load destination of LWC1, LDC1, or MTC1 instructions), the FPR receives a binary value that is uninterpreted.
- A computational or FP register move instruction that produces a result of type *fnt* puts a value of type *fnt* into the result register.

When an FPR with an uninterpreted value is used as a source operand by an instruction that requires a value of format *fnt*, the binary contents are interpreted as an encoded value in format *fnt*, and the value in the FPR changes to a value of format *fnt*. The binary contents cannot be reinterpreted in a different format.

If an FPR contains a value of format *fnt*, a computational instruction must not use the FPR as a source operand of a different format. If this case occurs, the value in the register becomes unknown, and the result of the instruction is also a value that is unknown. Using an FPR containing an unknown value as a source operand produces a result that has an unknown value.

The format of the value in the FPR is unchanged when it is read by a data transfer instruction (a store i.e. source operand of SWC1, SDC1, or MFC1 instructions). A data transfer instruction produces a binary encoding of the value contained in the FPR. If the value in the FPR is unknown, the encoded binary value produced by the operation is not defined.

50.12.3.3 32-BIT AND 64-BIT BINARY DATA TRANSFER

The data transfer instructions move words and double words between the FPU FPRs and the system.

The operations of the word and double word load and move-to instructions are shown in [Figure 50-20](#) and [Figure 50-21](#), respectively.

The store and move-from instructions do the reverse, reading data from the location that the corresponding load or move-to instruction had written.

Figure 50-20: FPU Word Load and Move-to Operations

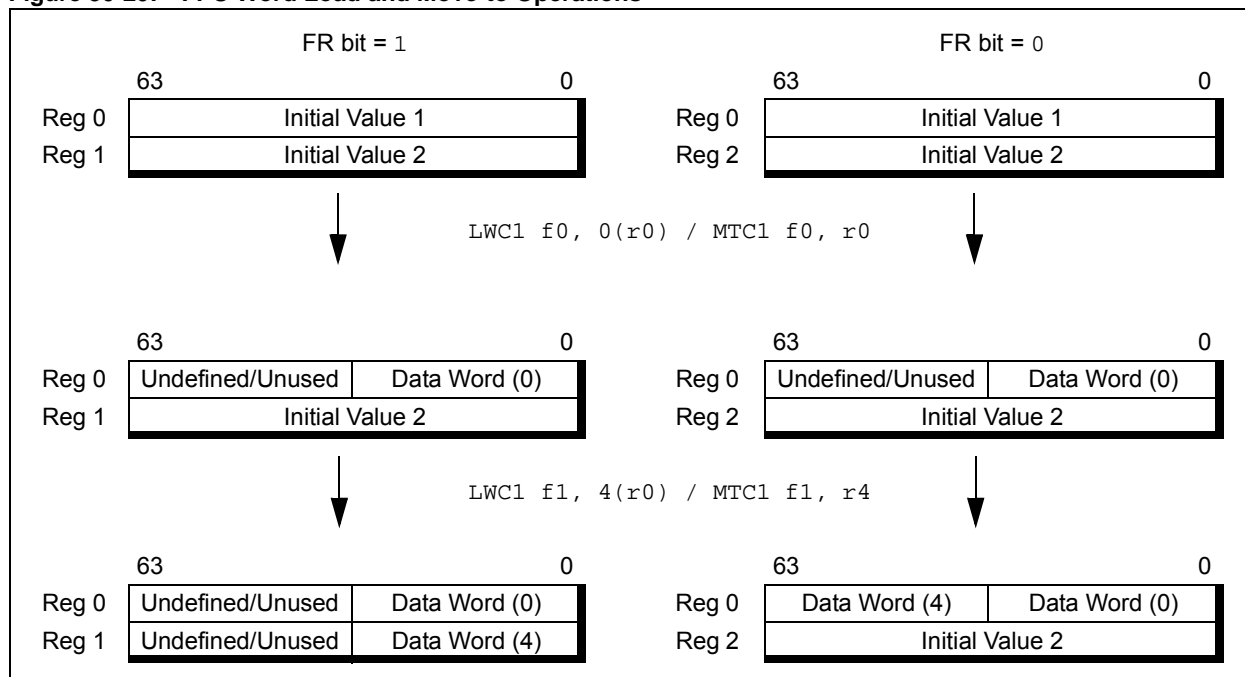
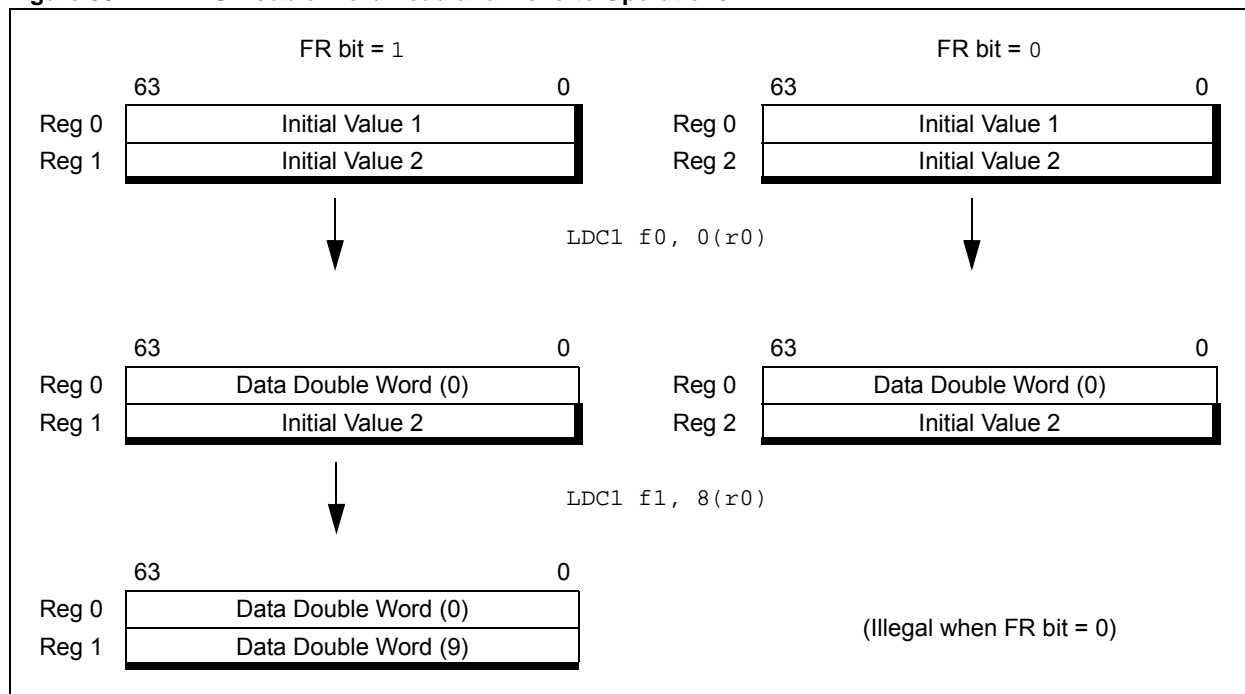


Figure 50-21: FPU Double Word Load and Move-to Operations



50.12.4 Floating Point Instruction Overview

The FPU instructions are divided into these categories:

- [FPU Data Transfer Instructions](#)
- [FPU Arithmetic Instructions](#)
- [FPU Conversion Instructions](#)
- [FPU Formatted Operand-Value Move Instructions](#)
- [FPU Conditional Branch Instructions](#)
- [FPU Miscellaneous Instructions](#)

The instructions are described in detail in **Chapter 14, “M5150 Processor Core Instructions”** of the “MIPS32® M5150 Processor Core Family Software User’s Manual”. This document is available for download by registered users from the Imagination Technologies Ltd. website (www.imgtec.com).

50.12.4.1 FPU DATA TRANSFER INSTRUCTIONS

The FPU has two register sets: Coprocessor General Registers (FPRs) and Coprocessor Control Registers (FCRs). The FPU has a load/store architecture: all computations are done on data held in coprocessor general registers.

The control registers are used to control FPU operation. Data is transferred between registers and the rest of the system with dedicated load, store, and move instructions. The transferred data is treated as unformatted binary data. No format conversions are performed, and therefore no IEEE floating point exceptions can occur.

Table 50-11: FPU Data Transfer Instructions

Transfer Direction			Transferred Data
FPU general register	<->	Memory	Word/double word load/store
FPU general register	<->	CPU general register	Word move
FPU control register	<->	CPU general register	Word move

All coprocessor loads and stores operate on naturally aligned data items. An attempt to load or store to an address that is not naturally aligned for the data item causes an Address Error exception. The address of a word or double word is the smallest byte address in the object. For the PIC32 architecture this is the least-significant byte.

50.12.4.2 FPU DATA TRANSFER INSTRUCTIONS ADDRESSING

The FPU has loads and stores using the same register + offset addressing as that used by the CPU. Moreover, for the FPU only, there are load and store instructions using register + register addressing.

Table [Table 50-12](#) and [Table 50-13](#) list the FPU data transfer instructions.

Table 50-12: FPU Load and Store Instructions

Mnemonic	Instruction	Addressing Mode
LDC1	Load Double word to Floating Point	Register + offset
LWC1	Load Word to Floating Point	Register + offset
SDC1	Store Double word from Floating Point	Register + offset
SWC1	Store Word from Floating Point	Register + offset
LDXC1	Load Double word Indexed to Floating Point	Register + Register
LUXC1	Load Double word Indexed Unaligned to Floating Point	Register + Register
LWXC1	Load Word Indexed to Floating Point	Register + Register
SDXC1	Store Double word Indexed from Floating Point	Register + Register
SUXC1	Store Double word Indexed Unaligned from Floating Point	Register + Register
SWXC1	Store Word Indexed from Floating Point	Register + Register

Table 50-13: FPU Move To and From Instructions

Mnemonic	Instruction
CFC1	Move Control Word From Floating Point
CTC1	Move Control Word To Floating Point
MFC1	Move Word From Floating Point
MFHC1	Move Word From High Half of Floating Point
MTC1	Move Word To Floating Point
MTHC1	Move Word to High Half of Floating Point

50.12.4.3 FPU ARITHMETIC INSTRUCTIONS

Arithmetic instructions operate on formatted data values. The results of most floating point arithmetic operations meet the IEEE 754 Standard for accuracy. A result is identical to an infinite-precision result that has been rounded to the specified format using the current rounding mode. The rounded result differs from the exact result by less than one Unit in the Least-significant Place (ULP).

In general, the arithmetic instructions take an Unimplemented Operation exception for denormalized numbers, except for the *ABS*, *C*, and *NEG* instructions, which can handle denormalized numbers. The FS, FO, and FN bits in the CP1 FCSR register can override this behavior as described in [50.14.6 “Floating Point Operation of the FS/FO/FN Bits”](#).

[Table 50-14](#) lists the FPU IEEE compliant arithmetic operations.

Table 50-14: FPU IEEE Arithmetic Instructions

Mnemonic	Instruction
ABS.fmt	Floating Point Absolute Value
ADD.fmt	Floating Point Add
C.cond.fmt	Floating Point Compare
DIV.fmt	Floating Point Divide
MUL.fmt	Floating Point Multiply
NEG.fmt	Floating Point Negate
SQRT.fmt	Floating Point Square Root
SUB.fmt	Floating Point Subtract
RECIP.fmt	Floating Point Reciprocal Approximation. See Note 1 .
RSQRT.fmt	Floating Point Reciprocal Square Root Approximation. See Note 2 .

Note 1: This low latency operation might be less accurate than the IEEE specification. The result of the RECIP differs from the exact reciprocal by no more than one Unit in the Least-significant Place (ULP).

2: This low latency operation might be less accurate than the IEEE specification. The result of the RSQRT differs from the exact reciprocal square root by no more than two ULPs.

Four compound-operation instructions perform variations of multiply-accumulate operations: multiply two operands, accumulate the result to a third operand, and produce a result. The product is rounded according to the current rounding mode prior to the accumulation. This model meets the IEEE accuracy specification; the result is numerically identical to an equivalent computation using multiply, add, subtract, or negate instructions.

The compound-operation instructions are listed in [Table 50-15](#).

Table 50-15: FPU Multiply-Accumulate Arithmetic Instructions

Mnemonic	Instruction
MADD.fmt	Floating Point Multiply Add
MSUB.fmt	Floating Point Multiply Subtract
NMADD.fmt	Floating Point Negative Multiply Add
NMSUB.fmt	Floating Point Negative Multiply Subtract

50.12.4.4 FPU CONVERSION INSTRUCTIONS

These instructions perform conversions between floating point and fixed point data types. Each instruction converts values from a number of operand formats to a particular result format. Some conversion instructions use the rounding mode specified in the Floating Control/Status register (FCSR), while others specify the rounding mode directly.

In general, the conversion instructions only take an Unimplemented Operation exception for denormalized numbers.

The FS and FN bits in the CP1 FCSR register can override this behavior as described in [50.14.6 “Floating Point Operation of the FS/FO/FN Bits”](#).

[Table 50-16](#) and [Table 50-17](#) list the FPU conversion instructions according to their rounding mode.

Table 50-16: FPU Conversion Operations Using the FCSR Rounding Mode Instructions

Mnemonic	Instruction
CVT.D.fmt	Floating Point Convert to Double Floating Point
CVT.L.fmt	Floating Point Convert to Long Fixed Point
CVT.S.fmt	Floating Point Convert to Single Floating Point
CVT.W.fmt	Floating Point Convert to Word Fixed Point

Table 50-17: FPU Conversion Operations Using a Directed Rounding Mode Instructions

Mnemonic	Instruction
CEIL.L.fmt	Floating Point Ceiling to Long Fixed Point
CEIL.W.fmt	Floating Point Ceiling to Word Fixed Point
FLOOR.L.fmt	Floating Point Floor to Long Fixed Point
FLOOR.W.fmt	Floating Point Floor to Word Fixed Point
ROUND.L.fmt	Floating Point Round to Long Fixed Point
ROUND.W.fmt	Floating Point Round to Word Fixed Point
TRUNC.L.fmt	Floating Point Truncate to Long Fixed Point
TRUNC.W.fmt	Floating Point Truncate to Word Fixed Point

50.12.4.5 FPU FORMATTED OPERAND-VALUE MOVE INSTRUCTIONS

These instructions move formatted operand values among FPU general registers. A particular operand type must be moved by the instruction that handles that type. There are three kinds of move instructions:

- Unconditional move
- Conditional move that tests an FPU true/false condition code
- Conditional move that tests a CPU general-purpose register against zero

Conditional move instructions operate in a way that might be unexpected. They always force the value in the destination register to become a value of the format specified in the instruction. If the destination register does not contain an operand of the specified format before the conditional move is executed, the contents become undefined. For more information, see the individual descriptions of the conditional move instructions in “MIPS32[®] Architecture Reference Manual, Volume II” and “microMIPS32[™] Architecture Reference Manual, Volume II”. These documents are available for download by registered users from the Imagination Technologies Ltd. website (www.imgtec.com).

Table [Table 50-18](#) lists the formatted operand-value move instructions.

Table 50-18: FPU Formatted Operand Move Instructions

Mnemonic	Instruction
MOV.fmt	Floating Point Move
MOV.F.fmt	Floating Point Move Conditional on FP False
MOV.T.fmt	Floating Point Move Conditional on FP True
MOV.N.fmt	Floating Point Move Conditional on Nonzero
MOV.Z.fmt	Floating Point Move Conditional on Zero

50.12.4.6 FPU CONDITIONAL BRANCH INSTRUCTIONS

The FPU has PC-relative conditional branch instructions that test condition codes set by FPU compare instructions (C.cond.fmt).

All branches have an architectural delay of one instruction. When a branch is taken, the instruction immediately following the branch instruction is said to be in the branch delay slot. It is executed before the branch to the target instruction takes place.

Conditional branches come in two versions, depending upon how they handle an instruction in the delay slot when the branch is not taken and execution falls through:

- Branch instructions execute the instruction in the delay slot.
- Branch likely instructions do not execute the instruction in the delay slot if the branch is not taken (they are said to nullify the instruction in the delay slot).

Note: Although the Branch Likely instructions are included, software is strongly encouraged to avoid the use of the Branch Likely instructions, as they will be removed from a future revision of the MIPS architecture.

The MIPS architecture defines eight condition codes for use in compare and branch instructions. For backward compatibility with previous revisions of the ISA, condition code bit 0 and condition code bits 1 through 7 are in discontinuous fields in the FCSR.

[Table 50-19](#) lists the conditional branch (branch and branch likely) FPU instructions.

Table 50-19: FPU Conditional Branch Instructions

Mnemonic	Instruction
BC1F	Branch on FP False
BC1T	Branch on FP True
BC1FL	Branch on FP False Likely. See Note 1 .
BC1FT	Branch on FP True Likely. See Note 1 .

Note 1: Deprecated FPU instruction. Software should avoid their use.

50.12.4.7 FPU MISCELLANEOUS INSTRUCTIONS

The MIPS32 architecture defines various miscellaneous instructions that conditionally move one CPU general register to another, based on an FPU condition code.

Table 50-20: CPU Conditional Move on FPU True/False Instructions

Mnemonic	Instruction
MOVN	Move Conditional on FP False
MOVZ	Move Conditional on FP True

50.12.5 Floating Point Exceptions Overview

There are five exception conditions defined by the IEEE 754 Standard:

- Invalid Operation Exception
- Division By Zero Exception
- Underflow Exception
- Overflow Exception
- Inexact Exception

There is also a MIPS-specific exception condition, the Unimplemented Operation Exception, that is used to signal a need for software emulation of an instruction.

Normally an IEEE arithmetic operation can cause only one exception condition. The only case in which two exceptions can occur simultaneously are Inexact With Overflow and Inexact With Underflow.

At the program's control an IEEE exception condition can either cause a trap or not cause a trap. The IEEE 754 Standard specifies the result to be delivered if no trap is taken. The FPU will supply these results whenever the exception condition does not result in a trap. The default action taken depends on the type of exception condition, and in the case of the Overflow and Underflow, the current rounding mode.

[Table 50-21](#) summarizes the default results supplied by the FPU.

FPU exceptions are implemented in the PIC32 FPU architecture with the Cause, Enables, and Flags fields of the FCSR. The flag bits implement IEEE exception status flags and the cause and enable bits control exception trapping.

Each field has a bit for each of the five IEEE exception conditions. The Cause field has an additional exception bit, Unimplemented Operation, that could be used to trap for software emulation assistance. If an exception type is enabled through the Enables field of the FCSR, the FPU is operating in precise exception mode for this type of exception.

50.12.5.1 FLOATING POINT PRECISE EXCEPTION MODE

In precise exception mode, a trap occurs before the instruction that causes the trap or any following instruction can complete and write its results. So the software trap handler can resume execution of the interrupted instruction stream after handling the exception, if desired.

The Cause field reports per-bit instruction exception conditions. The cause bits are written during each floating point arithmetic operation to show any exception conditions that arise during the operation. A cause bit is set to '1' if its corresponding exception condition arises; otherwise, it is cleared to '0'.

A floating point trap is generated any time both a cause bit and its corresponding enable bit are set. This case occurs either during the execution of a floating point operation or when moving a value into the FCSR. There is no enable bit for Unimplemented Operations: this exception always generates a trap.

In a trap handler, exception conditions that arise during any trapped floating point operations are reported in the Cause field. Before returning from a floating point interrupt or exception, or before setting cause bits with a move to the FCSR, software first must clear the enabled cause bits by executing a move to the FCSR to prevent the trap from being erroneously retaken.

If a floating point operation sets only non-enabled cause bits, no trap occurs and the default result defined by the IEEE 754 Standard is stored (see [Table 50-21](#)). When a floating point operation does not trap, the program can monitor the exception conditions by reading the Cause field.

The Flags field is a cumulative report of IEEE exception conditions that arise as instructions complete; instructions that trap do not update the flag bits. The flag bits are set to '1' if the corresponding IEEE exception is raised, otherwise the bits are unchanged. There is no flag bit for the Unimplemented Operation exception. The flag bits are never cleared as a side effect of floating point operations, but they can be set or cleared by the software by moving a new value into the FCSR.

Table 50-21: FPU Supplied Results for Not Trapped Exceptions

Bit Name	Description	Default Action
V	Invalid Operation	Supplies a quiet NaN.
Z	Divide by Zero	Supplies a properly signed infinity.
U	Underflow	Depends on the rounding mode as shown below: (RN): Supplies a zero with the sign of the exact result. (RZ): Supplies a zero with the sign of the exact result. (RP): For positive underflow values, supplies $2^{E_{\min}}$ (MinNorm). For negative underflow values, supplies a positive zero. (RM): For positive underflow values, supplies a negative zero. For negative underflow values, supplies a negative $2^{E_{\min}}$ (MinNorm). Note: this behavior is only valid if the FCSR FN bit is cleared.
I	Inexact	Supplies a rounded result. If caused by an overflow without the overflow trap enabled, supplies the overflowed result. If caused by an underflow without the underflow trap enabled, supplies the underflowed result.
O	Overflow	Depends on the rounding mode, as follows: <ul style="list-style-type: none"> • (RN): Supplies a infinity with the sign of the exact result. • (RZ): Supplies the format's largest finite number with the sign of the exact result. • (RP): For positive overflow values, supplies positive infinity. For negative overflow values, supplies the format's most negative finite number. • (RM): For positive overflow values, supplies the format's largest finite number. For negative overflow values, supplies minus infinity.

50.12.5.2 FLOATING POINT INVALID OPERATION EXCEPTION

An Invalid Operation exception is signaled when one or both of the operands are invalid for the operation to be performed. When the exception condition occurs without a precise trap, the result is a quiet NaN.

The following operations are invalid:

- One or both operands are a signaling NaN (except for the non-arithmetic MOV.fmt, MOVT.fmt, MOVF.fmt, MOVN.fmt, and MOVZ.fmt instructions).
- Addition or subtraction: magnitude subtraction of infinities, such as $(+\infty) + (-\infty)$ or $(-\infty) - (-\infty)$.
- Multiplication: $0 \times \infty$, with any signs.
- Division: $0/0$ or ∞/∞ , with any signs.
- Square root: An operand of less than 0 (-0 is a valid operand value).
- Conversion of a floating point number to a fixed point format when either an overflow or an operand value of infinity or NaN precludes a faithful representation in that format.
- Some comparison operations in which one or both of the operands is a QNaN value.

50.12.5.3 FLOATING POINT DIVISION BY ZERO EXCEPTION

The divide operation signals a Division By Zero exception if the divisor is zero and the dividend is a finite nonzero number. When no precise trap occurs, the result is a correctly signed infinity. Divisions (0/0 and $\infty/0$) do not cause the Division By Zero exception. The result of (0/0) is an Invalid Operation exception. The result of ($\infty/0$) is a correctly signed infinity.

50.12.5.4 FLOATING POINT UNDERFLOW EXCEPTION

There are two related events that contribute to underflow:

- **Tininess:** The creation of a tiny, nonzero result between $\pm 2^{E_{\min}}$ which, because it is tiny, might cause some other exception later such as overflow on division. The IEEE 754 Standard allows choices in detecting tininess events. The PIC32/MIPS architecture specifies that tininess be detected after rounding, when a nonzero result computed as though the exponent range were unbounded would lie strictly between $\pm 2^{E_{\min}}$.
- **Loss of accuracy:** The extraordinary loss of accuracy occurs during the approximation of such tiny numbers by denormalized numbers. The IEEE 754 Standard allows choices in detecting loss of accuracy events. The PIC32/MIPS architecture specifies that loss of accuracy be detected as inexact result, when the delivered result differs from what would have been computed if both the exponent range and precision were unbounded.

The way that an underflow is signaled depends on whether or not underflow traps are enabled:

- When an underflow trap is not enabled, underflow is signaled only when both tininess and loss of accuracy have been detected. The delivered result might be zero, denormalized, or $\pm 2^{E_{\min}}$.
- When an underflow trap is enabled (through the FCSR Enables field), underflow is signaled when tininess is detected regardless of loss of accuracy.

50.12.5.5 FLOATING POINT OVERFLOW EXCEPTION

An Overflow exception is signaled when the magnitude of a rounded floating point result (if the exponent range is unbounded) is larger than the destination format's largest finite number.

When no precise trap occurs, the result is determined by the rounding mode and the sign of the intermediate result.

50.12.5.6 FLOATING POINT INEXACT EXCEPTION

An Inexact exception is signaled when one of the following occurs:

- The rounded result of an operation is not exact.
- The rounded result of an operation overflows without an overflow trap.
- When a denormal operand is flushed to zero.

50.12.5.7 FLOATING POINT UNIMPLEMENTED OPERATION EXCEPTION

The Unimplemented Operation exception is a MIPS-defined exception that provides software emulation support. This exception is not IEEE-compliant.

The PIC32/MIPS architecture is designed so that a combination of hardware and software can implement the FPU functionality. Operations not fully supported in hardware cause an Unimplemented Operation exception, allowing software to perform the operation.

There is no enable bit for this condition. It will always causes a trap (but the condition is effectively masked for all operations when $FS = 1$). After the appropriate emulation or other operation is done in a software exception handler, the original instruction stream can be continued.

An Unimplemented Operation exception is taken when denormalized operands or tiny results are encountered for instructions not supporting denormalized numbers and where such are not handed by the FS/FO/FN bits.

50.12.6 Floating Point Pipeline and Performance

This section describes the structure and operation of the FPU pipeline.

50.12.6.1 FPU PIPELINE OVERVIEW

The FPU has a seven stage pipeline to which the integer pipeline dispatches instructions:

- Decode, register read and unpack (FR stage)
- Multiply tree - double pumped for double (M1 stage)
- Multiply complete (M2 stage)
- Addition first step (A1 stage)
- Addition second and final step (A2 stage)
- Packing to IEEE format (FP stage)
- Register writeback (FW stage)

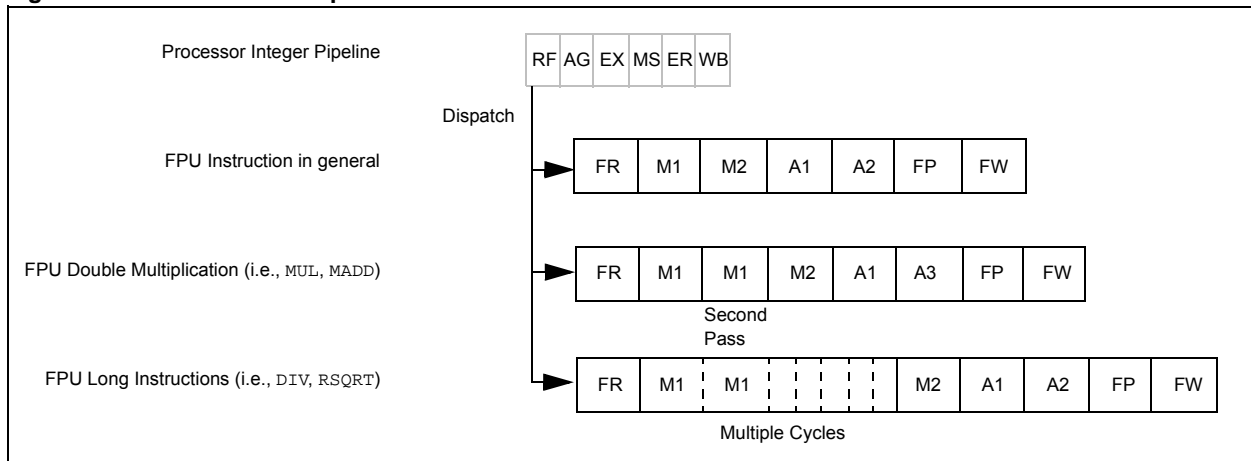
The FPU implements a bypass mechanism that allows the result of an operation to be forwarded directly to the instruction that needs it without having to write the result to the FPU register and then read it back.

The FPU pipeline runs in parallel with the PIC32 core integer pipeline. The FPU is built to run at the same frequency as the PIC32 core.

The FPU pipeline is optimized for single-precision instructions, such that the basic multiply, ADD/SUB, and MADD/MSUB instructions can be performed with single-cycle throughput and low latency. Executing double-precision multiply and MADD/MSUB instructions requires a second pass through the M1 stage to generate all 64 bits of the product.

Executing long latency instructions, such as DIV and RSQRT, extends the M1 stage. [Figure 50-22](#) shows the FPU pipeline.

Figure 50-22: PIC32 FPU Pipeline



Stage 1: FPU Pipeline: FR Stage – Decode, Register Read, and Unpack

The FR stage has the following functionality:

- The dispatched instruction is decoded for register accesses.
- Data is read from the register file.
- The operands are unpacked into an internal format.

Stage 2: FPU Pipeline: M1 Stage – Multiply Tree

The M1 stage has the following functionality:

- A single-cycle multiply array is provided for single-precision data format multiplication, and two cycles are provided for double-precision data format multiplication
- The long instructions, such as divide and square root, iterate for several cycles in this stage.
- Sum of exponents is calculated.

Stage 3: FPU Pipeline: M2 Stage – Multiply Complete

The M2 stage has the following functionality:

- Multiplication is complete when the carry-save encoded product is compressed into binary
- Rounding is performed
- Exponent difference for addition path is calculated

Stage 4: FPU Pipeline: A1 Stage – Addition First Step

This stage performs the first step of the addition.

Stage 5: FPU Pipeline: A2 Stage - Addition Second and Final Step

This stage performs the second and final step of the addition.

Stage 6: FPU Pipeline: FP Stage – Result Pack

The FP stage has the following functionality:

- The result coming from the data path is packed into the IEEE 754 Standard format for the FPR register file
- Overflow and underflow exceptional conditions are resolved

Stage 7: FPU Pipeline: FW Stage – Register Write

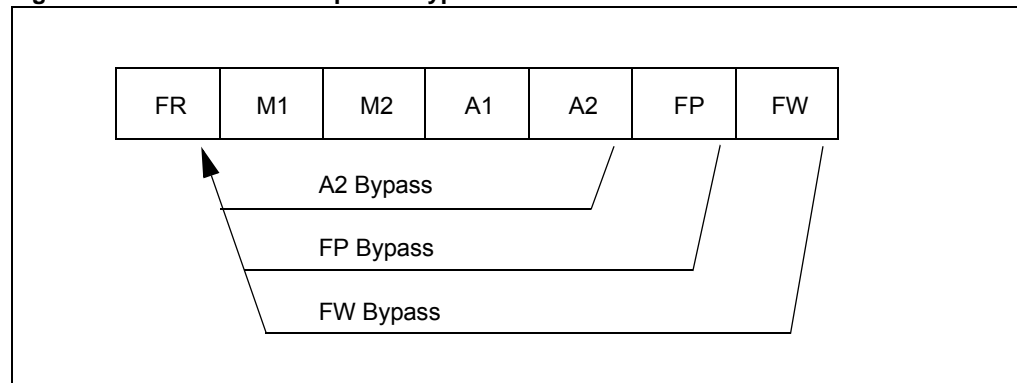
The result is written to the FPR register file.

50.12.6.2 FPU BYPASSING

The FPU pipeline implements extensive bypassing so that the results do not need to be written into the register file and read back before they can be used, but can be forwarded directly to an instruction already in the pipe.

Some bypassing is disabled when operating in 32-bit register file mode (the FR bit in the CP0 Status<26> register is '0'), due to the paired even-odd 32-bit registers that provide 64-bit registers.

Figure 50-23: PIC32 FPU Pipeline Bypass Paths



50.12.6.3 FPU REPEAT RATE AND LATENCY

Table 50-22 shows the repeat rate and latency for the FPU instructions.

Note that cycles related to floating point operations are listed in terms of FPU clocks.

Table 50-22: FPU Latency and Repeat Rate

Op code (see Note 1)	Latency (FPU Cycles)	Repeat Rate (FPU Cycles)
ABS.[S,D], NEG.[S,D], ADD.[S,D], SUB.[S,D], MUL.S, MADD.S, MSUB.S, NMADD.S, MSUB.S	4	1
MUL.D, MADD.D, MSUB.D, NMADD.D, NMSUB.D	5	2
RECIP.S	13	10
RECIP.D	25	21
RSQRT.S	17	14
RSQRT.D	35	31
DIV.S, SQRT.S	17	14
DIV.D, SQRT.D	32	29
C.cond.[S,D] to MOVF.fmt and MOVT.fmt instruction/ MOVT, MOVN, BCL instruction	One-half	1
CVT.D.S, CVT.[S,D].[W,L]	4	1
CVT.S.D	6	1
CVT.[W,L].[S,D], CEIL.[W,L].[S,D], FLOOR.[W,L].[S,D], ROUND.[W,L].[S,D], TRUNC.[W,L].[S,D]	5	1
MOV.[S,D], MOVF.[S,D], MOVN.[S,D], MOVT.[S,D], MOVZ.[S,D]	4	1
LWC1, LDC1, LDXC1, LUXC1, LWXC1	3	1
MTC1, MFC1	2	1

Legend: S = Single; D = Double; W = Word; L = Long Word

50.12.6.4 FLOATING POINT 2008 FPU SUPPORT

The PIC32 FPU implements the following status/control bits to provide greater compatibility with the IEEE 754 Standard Floating Point released in 2008:

- The Has2008 bit (FIR<23>) will always read as '1' to signify that 2008 FPU is implemented.
- The MAC2008 bit (FCSR<20>) will always read as '0' to signify that Fused Multiply Add operation is not implemented.
- The ABS2008 bit (FCSR<19>) is always set to '1' which makes ABS and NEG instructions non-arithmetic instructions. All floating point exceptions will be disabled.
- The NAN2008 bit (FCSR<18>) is always set to '1' to show Quiet and signaling NaN encodings recommended by the IEEE 754-2008 Standard. In addition, the following behaviors are implemented:
 - In the case of one or more QNaN operands (no SNaN), the QNaN operand is propagated from one of the input register operands (in order of priority): *fs*, *ft*, and *fr* (see the following **Note**).
 - When SNaN is used as an input, and exceptions are disabled, QNaN is the expected output
 - The QNaN output will not be a fixed value. To comply with IEEE, an input NaN should produce a NaN with the payload of the input NaN if representable in the destination format, where the payload is defined as the Mantissa field less its most-significant bit.
 - If ABS2008 = 1 and MAC2008 = 0 (as it always is in PIC32), the sign of NMADD and NMSUB do not flip the sign of any QNaN input, and the sign is retained and propagated to the output.
 - When a NaN is an input, the output will be one of the input NaNs with as much of the mantissa preserved as possible.
 - SNaN inputs have higher priority than QNaN inputs and then *fs* has higher priority than *ft* which has higher priority than *fr* register (see the following **Note**).
 - The sign of the selected NaN input is preserved. If the input that is selected for the output is already a QNaN, the entire mantissa is preserved. However, if the input that is selected for the output is a SNaN, the most significant bit of the SNaN mantissa is complemented to convert the SNaN into a QNaN. If this conversion to a QNaN would result in an infinity, the next most significant bit of the mantissa is set.
 - For CVT.s.d, the NaN mantissa most significant bits are preserved. For CVT.d.s, the NaN mantissa is padded with '0's in the least significant bits.
 - For multiply-add, if both *fs/ft* and *fr* registers are QNaNs, the multiply produces a QNaN based upon *fs/ft*, and this QNaN has priority over *fr* in the add operation. However, if both *fs/ft* and *fr* registers are SNaNs and the invalid trap is not enabled, the multiply generates a QNaN based upon *fs/ft*, which is then added to the signaling *fr* register and the signaling *fr* has priority
 - When a NaN is needed for output but there is no NaN input, a positive QNaN is created that has all other mantissa bits set.

Note: Registers *fs*, *ft*, and *fr* are floating point registers that occur in the multiply-add floating point operations. For example: MADD.D fd, fr, fs, ft.

50.12.6.5 IEEE 754-1985 STANDARD

The IEEE 754-1985 Standard, "IEEE Standard for Binary Floating-Point Arithmetic" defines the following:

- Floating Point data types
- The basic arithmetic, comparison, and conversion operations
- A computational model

The IEEE 754-1985 Standard does not define specific processing resources nor does it define an instruction set. For additional information about the IEEE 754-1985 standard, visit the IEEE Web page at <http://stds.bbs.ieee.org/>.

50.13 COPROCESSOR 0 (CP0) REGISTERS

The PIC32 uses a special register interface to communicate status and control information between system software and the CPU. This interface is called Coprocessor 0, or CP0. The features of the CPU that are visible through Coprocessor 0 are:

- Translation Lookaside Buffer (TLB)
- Core timer
- Interrupt and exception control
- Virtual memory configuration
- Shadow register set control
- Processor identification
- Debugger control
- Performance counters

System software accesses the registers in CP0 using coprocessor instructions such as MFC0 and MTC0. [Table 50-23](#) describes the CP0 registers found on PIC32 devices.

Table 50-23: CP0 Registers

Register Number	Register Name	Function
0	Index	Index into the TLB array (MPU only).
1	Random	Randomly generated index into the TLB array (MPU only).
2	EntryLo0	Low-order portion of the TLB entry for even-numbered virtual pages (MPU only).
3	EntryLo1	Low-order portion of the TLB entry for odd-numbered virtual pages (MPU only).
4	Context/ UserLocal	Pointer to the page table entry in memory (MPU only). User information that can be written by privileged software and read via the RDHWR instruction.
5	PageMask/ PageGrain	PageMask controls the variable page sizes in TLB entries. PageGrain enables support of 1 KB pages in the TLB (MPU only).
6	Wired	Controls the number of fixed (i.e., wired) TLB entries (MPU only).
7	HWREna	Enables access via the RDHWR instruction to selected hardware registers in Non-privileged mode.
8	BadVAddr	Reports the address for the most recent address-related exception.
	BadInstr	Reports the instruction that caused the most recent exception.
	BadInstrP	Reports the branch instruction if a delay slot caused the most recent exception.
9	Count	Processor cycle count.
10	EntryHi	High-order portion of the TLB entry (MPU only).
11	Compare	Core timer interrupt control.
12	Status	Processor status and control.
	IntCtl	Interrupt control of vector spacing.
	SRSCtl	Shadow register set control.
	SRSMap	Shadow register mapping control.
	View_IPL	Allows the Priority Level to be read/written without extracting or inserting that bit from/to the Status register.
	SRSMP2	Contains two 4-bit fields that provide the mapping from a vector number to the shadow set number to use when servicing such an interrupt.
13	Cause	Describes the cause of the last exception.
	NestedExc	Contains the error and exception level status bit values that existed prior to the current exception.
	View_RIPL	Enables read access to the RIPL bit that is available in the Cause register.
14	EPC	Program counter at last exception.
	NestedEPC	Contains the exception program counter that existed prior to the current exception.

Table 50-23: CP0 Registers (Continued)

Register Number	Register Name	Function
15	PRID	Processor identification and revision
	Ebase	Exception base address of exception vectors.
	CDMMBase	Common device memory map base.
16	Config	Configuration register.
	Config1	Configuration register 1.
	Config2	Configuration register 2.
	Config3	Configuration register 3.
	Config4	Configuration register 4.
	Config5	Configuration register 5.
	Config7	Configuration register 7.
17	LLAddr	Load link address (MPU only).
18	WatchLo	Low-order watchpoint address (MPU only).
19	WatchHi	High-order watchpoint address (MPU only).
20-22	Reserved	Reserved in the PIC32 core.
23	Debug	EJTAG debug register.
	TraceControl	EJTAG trace control.
	TraceControl2	EJTAG trace control 2.
	UserTraceData1	EJTAG user trace data 1 register.
	TraceBPC	EJTAG trace breakpoint register.
	Debug2	Debug control/exception status 1.
24	DEPC	Program counter at last debug exception.
	UserTraceData2	EJTAG user trace data 2 register.
25	PerfCtl0	Performance counter 0 control.
	PerfCnt0	Performance counter 0.
	PerfCtl1	Performance counter 1 control.
	PerfCnt1	Performance counter 1.
26	ErrCtl	Software test enable of way-select and data RAM arrays for I-Cache and D-Cache (MPU only).
27	CacheErr	Records information about Cache/SPRAM parity errors.
28	TagLo/DataLo	Low-order portion of cache tag interface (MPU only).
29	Reserved	Reserved in the PIC32 core.
30	ErrorEPC	Program counter at last error.
31	DeSAVE	Debug handler scratchpad register.
	KScratch1	Scratch register for Kernel mode.
	KScratch2	Scratch register for Kernel mode.

50.13.1 Index Register (CP0 Register 0, Select 0) (MPU only)

The Index register is a 32-bit read/write register that contains the index used to access the TLB for TLBP, TLBR, and TLBW instructions.

Register 50-1: Index; TLB Index Register; CP0 Register 0, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-x	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	P	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	—	—	—	Index<4:0> ⁽¹⁾				

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **P**: Probe Failure Detect bit

1 = The previous TLBP instruction failed to find a match in the TLB

0 = The previous TLBP instruction found a match in the TLB

bit 30-5 **Unimplemented**: Read as '0'

bit 4-0 **Index<4:0>**: Index to TLB Entry Affected by the TLBR and TLBW Instructions bits⁽¹⁾

11111 = TLB Entry 31

•
•
•

00000 = TLB Entry 0

Note 1: Depending on the configuration of the MMU, not all bits may be used. The number of TLB entries supported by the MMU can be read from the MMU Size<5:0> field of the Config1 CP0 register.

50.13.2 Random Register (CP0 Register 1, Select 0) (MPU only)

The Random register is a read-only register whose value is used to index the TLB during a TLBWR instruction.

The value of the register varies between an upper and lower bound as follows:

- A lower bound is set by the number of TLB entries reserved for exclusive use by the operating system (the contents of the Wired register). The entry indexed by the Wired register is the first entry available to be written by a TLB Write Random operation.
- An upper bound is set by the total number of TLB entries minus 1

The Random register is decremented by one almost every clock, wrapping after the value in the Wired register is reached. To enhance the level of randomness and reduce the possibility of a live lock condition, an LFSR register is used that prevents the decrement pseudo-randomly.

The processor initializes the Random register to the upper bound on a Reset exception and when the Wired register is written.

Register 50-2: Random; Random Field Register; CP0 Register 1, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	R-1	R-1	R-1	R-1
	—	—	—	—	Random<3:0>			

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-4 **Unimplemented:** Read as '0'

bit 3-0 **Random<3:0>:** TLB Random Index bits

50.13.3 EntryLo0 Register (CP0 Register 2, Select 0) and EntryLo1 Register (CP0 Register 3, Select 0) (MPU only)

The pair of EntryLo registers act as the interface between the TLB and the TLBR, TLBWI, and TLBWR instructions. EntryLo0 holds the entries for even pages and EntryLo1 holds the entries for odd pages.

The contents of the EntryLo0 and EntryLo1 registers are undefined after an address error, TLB invalid, TLB modified, or TLB refill exception.

Register 50-3: EntryLo0; Even Page TLB Entries Register; CP0 Register 2, Select 0 and EntryLo1; Odd Page TLB Entries Register; CP0 Register 3, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-0	R/W-0	U-0	U-0	U-0	U-0	R/W-x	R/W-x
	RI	XI	—	—	—	—	PFN<19:18>	
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	PFN<17:10>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	PFN<9:2>							
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	PFN<1:0>		C<2:0>			D	V	G

Legend:

R = Readable bit
-n = Value at POR

W = Writable bit
'1' = Bit is set

U = Unimplemented bit, read as '0'
'0' = Bit is cleared
x = Bit is unknown

- bit 31 **RI:** Read Inhibit bit
If this bit is set, an attempt to read data from the page causes a TLB Invalid exception, even if the V (Valid) bit is set. The RI bit is enabled only if the RIE bit of the PageGrain register is set. If the RIE bit of PageGrain is not set, the RI bit of EntryLo0/EntryLo1 is a reserved '0' bit as per the MIPS32 specification.
- bit 30 **XI:** Execute Inhibit bit
If this bit is set, an attempt to fetch from the page causes a TLB Invalid exception, even if the V (Valid) bit is set. The XI bit is enabled only if the XIE bit of the PageGrain register is set. If the XIE bit of PageGrain is not set, the XI bit of EntryLo0/EntryLo1 is a reserved '0' bit as per the MIPS32 specification.
- bit 29-26 **Unimplemented:** Read as '0'
- bit 25-6 **PFN<19:0>:** Page Frame Number bits
Contributes to the definition of the high-order bits of the physical address.
The PFN<19:0> bits correspond to bits <31:12> of the physical address.
- bit 5-3 **C<2:0>:** Coherency Page Attribute bits
111 = Reserved
110 = Reserved
101 = Reserved
100 = Reserved
011 = Cacheable, non-coherent, write-back, write allocate
010 = Uncached
001 = Cacheable, non-coherent, write-through, write allocate
000 = Cacheable, non-coherent, write-through, no write allocate
- bit 2 **D:** Dirty (write-enable) bit
1 = Stores to the page are permitted
0 = Stores to the page cause a TLB modified exception
- bit 1 **V:** Valid bit
1 = Accesses to the page are permitted
0 = Accesses to the page cause a TLB invalid exception
- bit 0 **G:** Global bit
On a TLB write, the logical AND of the G bits in both the EntryLo0 and EntryLo1 register becomes the G bit in the TLB entry. If the TLB entry G bit is a '1', the ASID comparisons are ignored during TLB matches. On a read from a TLB entry, the G bits of both EntryLo0 and EntryLo1 reflect the state of the TLB G bit.

50.13.4 Context Register (CP0 Register 4, Select 0) (MPU Only)

The Context register is a read/write register containing a pointer to an entry in the page table entry (PTE) array. This array is an operating system data structure that stores virtual-to-physical translations. During a TLB miss, the operating system loads the TLB with the missing translation from the PTE array. The Context register duplicates some of the information provided in the BadVAddr register but is organized in such a way that the operating system can directly reference an 8-byte page table entry (PTE) in memory.

Register 50-4: Context: Context Register; CP0 Register 4, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	PTEBase<8:1>							
23:16	R/W-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	PTEBase<0>	BadVPN2<19:13>						
15:8	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadVPN2<12:5>							
7:0	R-x	R-x	R-x	R-x	R-x	U-0	U-0	U-0
	BadVPN2<4:0>					—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-23 **PTEBase<8:0>**: Context Register PTE Array Pointer bits

These bits are for use by the operating system and are normally written with a value that allows the operating system to use the Context register as a pointer into the current PTE array in memory.

bit 22-4 **BadVPN2<19:0>**: TLB Hardware Miss Status bits

These bits contain the value of bits VA<31:13> of the virtual address that was missed.

bit 3-0 **Unimplemented**: Read as '0'

50.13.5 UserLocal Register (CP0 Register 4, Select 2)

The UserLocal register is a read-write register that is not interpreted by hardware and is conditionally readable through the RDHWR instruction.

Privileged software may write this register with arbitrary information and make it accessible to unprivileged software through register 29 (ULR) of the RDHWR instruction. To do so, the URL bit (HWREna<29>) must be set to a '1' to enable unprivileged access to the register.

In some operating environments, the UserLocal register contains a pointer to a thread-specific storage block that is obtained through the RDHWR register.

Register 50-5: UserLocal: User Local Register; CP0 Register 4, Select 2

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	USERLOCAL<31:24>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	USERLOCAL<23:16>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	USERLOCAL<15:8>							
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	USERLOCAL<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **USERLOCAL<31:0>**: User Local bits

50.13.6 PageMask Register (CP0 Register 5, Select 0) (MPU only)

The PageMask register is a read/write register used for reading from and writing to the TLB. It holds a comparison mask that sets the variable page size for each TLB entry, as shown in Table 50-24.

Register 50-6: PageMask; TLB Variable Page Size Register; CP0 Register 5, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	—	—	—	Mask<15:11>				
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	Mask<10:3>							
15:8	R/W-x	R/W-x	R/W-x	U-0	U-0	U-0	U-0	U-0
	Mask<2:0>			—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-29 **Unimplemented:** Read as '0'

bit 28-13 **Mask<15:0>:** Virtual Address Mask bits

When this bit is a '1', this indicates that the corresponding bit of the virtual address should not participate in the TLB match.

bit 12-0 **Unimplemented:** Read as '0'

Table 50-24: Values for the Mask bits of the PageMask Register

Page Size	Register Bit Location															
	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13
4 KB	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
16 KB	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1
64 KB	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1
256 KB	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1
1 MB	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1
4 MB	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
16 MB	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	1
64 MB	0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1
256 MB	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1

50.13.7 PageGrain Register (CP0 Register 5, Select 1) (MPU only)

The PageGrain register is used on the PIC32 device to enable or disable the read and execute inhibit bits in the EntryLo0 and EntryLo1 registers.

Register 50-7: PageGrain; TLB Page Grain Enable Register; CP0 Register 5, Select 1

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-0	R/W-0	U-0	U-0	R/W-0	U-0	U-0	U-0
	RIE	XIE	—	—	IEC	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **RIE:** Read Inhibit Enable bit

1 = RI bit of EntryLo0 and EntryLo1 registers is enabled

0 = RI bit of EntryLo0 and EntryLo1 registers is disabled and is not writable

bit 30 **XIE:** Execute Inhibit Enable bit

1 = XI bit of EntryLo0 and EntryLo1 registers is enabled

0 = XI bit of EntryLo0 and EntryLo1 registers is disabled and is not writable

bit 29-28 **Unimplemented:** Must be written as '0'; returns '0' on a read

bit 27 **IEC:** Enable Read-Inhibit and Execute-Inhibit Exception Codes bit

1 = Read-Inhibit exceptions use the TLBRI exception code. Execute-Inhibit exceptions use the TLBXI exception code

0 = Read-Inhibit and Execute-Inhibit exceptions both use the TLBL exception code

bit 26-0 **Unimplemented:** Must be written as '0'; returns '0' on a read

50.13.8 Wired Register (CP0 Register 6, Select 0) (MPU Only)

The Wired register is a read/write register that specifies the boundary between the wired and random entries in the TLB. The width of the Wired field is calculated in the same manner as that described for the Index register. Wired entries are fixed, non-replaceable entries that are not overwritten by a TLBWR instruction. Wired entries can be overwritten by a TLBWI instruction.

The Wired register is reset to zero by a Reset exception. Writing the Wired register causes the Random register to reset to its upper bound.

Register 50-8: Wired; TLB Boundary Entries Register; CP0 Register 6, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	—	—	—	Wired<4:0> ⁽¹⁾				

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-5 **Unimplemented:** Read as '0'

bit 4-0 **Wired<4:0>:** TLB Wired Boundary bits⁽¹⁾

11111 = Entry 31 is random, entries 0-30 are wired

.

.

.

01111 = Entry 15 is random, entries 0-14 are wired

.

.

.

00111 = Entries 7 and above are random, below 7 are wired

.

.

.

00000 = All 16 entries are random

Note 1: Depending on the configuration of the MMU, not all bits may be used. The number of TLB entries supported by the MMU can be read from the MMU Size<5:0> field of the Config1 CP0 register.

50.13.9 HWREna Register (CP0 Register 7, Select 0)

The HWREna register contains a bit mask that determines which hardware registers are accessible through the RDHWR instruction.

Register 50-9: HWREna: Hardware Accessibility Register; CP0 Register 7, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	R/W-0	U-0	U-0	U-0	U-0	U-0
	—	—	ULR	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	R/W-0	R/W-0	R/W-0	R/W-0
	—	—	—	—	MASK<3:0>			

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-30 **Unimplemented:** Read as '0'

bit 29 **ULR:** User Local Register bit

1 = Enable unprivileged access to UserLocal register

0 = Disable unprivileged access to UserLocal register

This bit provides read access to the Coprocessor 0 UserLocal register.

bit 28-4 **Unimplemented:** Read as '0'

bit 3-0 **MASK<3:0>:** Bit Mask bits

1 = Access is enabled to corresponding hardware register

0 = Access is disabled

Each of these bits enables access by the RDHWR instruction to a particular hardware register (which may not be an actual register). See the RDHWR instruction for a list of valid hardware registers.

50.13.10 BadVAddr Register (CP0 Register 8, Select 0)

BadVAddr is a read-only register that captures the most recent virtual address that caused an address error exception. Address errors are caused by executing load, store, or fetch operations from unaligned addresses, or by trying to access Kernel mode addresses from User mode.

For devices with the MPU core, the BadVAddr register will also capture the most recent virtual address that caused a TLB refill, TLB invalid, or TLB modified exception.

BadVAddr does not capture address information for bus errors, because they are not addressing errors.

Register 50-10: BadVAddr: Bad Virtual Address Register; CP0 Register 8, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadVAddr<31:24>							
23:16	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadVAddr<23:16>							
15:8	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadVAddr<15:8>							
7:0	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadVAddr<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **BadVAddr<31:0>**: Bad Virtual Address bits

Captures the virtual address that caused the most recent address error exception.

50.13.11 BadInstr Register (CP0 Register 8, Select 1) (M-Class only)

The BadInstr register is an optional read-only register that captures the most recent instruction that caused one of the following exceptions:

- Instruction Validity: Coprocessor Unusable, Reserved Instruction
- Execution Exception: Integer Overflow, Trap, System Call, Breakpoint, Floating-point, Coprocessor 2 exception
- Addressing: Address Error, TLB Refill, TLB Invalid, TLB Read Inhibit, TLB Execute Inhibit, TLB Modified

The BadInstr register is provided to allow acceleration of instruction emulation. The BadInstr register is only set by exceptions that are synchronous to an instruction. The BadInstr register is not set by Interrupts or by NMI, Machine check, Bus Error, Cache Error, Watch, or EJTAG exceptions.

When a synchronous exception occurs for which there is no valid instruction word (for example TLB Refill - Instruction Fetch), the value stored in BadInstr is unpredictable.

Presence of the BadInstr register is indicated by the Config3BI bit. The BadInstr register is instantiated per-VPE in an MT ASE processor.

Register 50-11: BadInstr: Bad Instruction Word Register; CP0 Register 8, Select 1

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
BadInstr<31:24>								
23:16	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
BadInstr<23:16>								
15:8	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
BadInstr<15:8>								
7:0	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
BadInstr<7:0>								

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **BadInstr<31:0>**: Bad (Faulting) Instruction Word bits

Instruction words smaller than 32 bits are placed in bits 15:0, with bits 31:16 containing '0'.

50.13.12 BadInstrP Register (CP0 Register 8, Select 2) (M-Class only)

The BadInstrP register is an optional register that is used in conjunction with the BadInstr register. The BadInstrP register contains the prior branch instruction when the faulting instruction is in a branch delay slot.

The BadInstrP register is updated for these exceptions:

- Instruction Validity: Coprocessor Unusable, Reserved Instruction
- Execution Exception: Integer Overflow, Trap, System Call, Breakpoint, Floating-point, Coprocessor 2 exception
- Addressing: Address Error, TLB Refill, TLB Invalid, TLB Read Inhibit, TLB Execute Inhibit, TLB Modified

The BadInstrP register is provided to allow acceleration of instruction emulation. The BadInstrP register is only set by exceptions that are synchronous to an instruction. The BadInstrP register is not set by Interrupts or by NMI, Machine check, Bus Error, Cache Error, Watch, or EJTAG exceptions. When a synchronous exception occurs, and the faulting instruction is not in a branch delay slot, then the value stored in BadInstrP is unpredictable.

Presence of the BadInstrP register is indicated by the Config3BP bit. The BadInstrP register is instantiated per-VPE in an MT ASE processor.

Register 50-12: BadInstrP: Bad Prior Branch Instruction Register; CP0 Register 8, Select 2

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadInstrP<31:24>							
23:16	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadInstrP<23:16>							
15:8	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadInstrP<15:8>							
7:0	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	BadInstrP<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **BadInstrP<31:0>**: Bad Prior Branch Instruction bits

Instruction words smaller than 32 bits are placed in bits 15:0, with bits 31:16 containing '0'.

50.13.13 Count Register (CP0 Register 9, Select 0)

The Count register acts as a timer, incrementing at a constant rate, whether or not an instruction is executed, retired, or any forward progress is made through the pipeline. The counter increments every other clock, if the DC bit in the Cause register is '0'.

Count can be written for functional or diagnostic purposes, including at Reset or to synchronize processors.

By writing the COUNTDM bit in the Debug register, it is possible to control whether Count continues to increment while the processor is in Debug mode.

Register 50-13: Count: Interval Counter Register; CP0 Register 9, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COUNT<31:24>								
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COUNT<23:16>								
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COUNT<15:8>								
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COUNT<7:0>								

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **COUNT<31:0>**: Interval Counter bits

This value is incremented every other clock cycle.

50.13.14 EntryHi Register (CP0 Register 10, Select 0) (MPU only)

The EntryHi register contains the virtual address match information used for TLB read, write, and access operations.

A TLB exception (TLB Refill, TLB Invalid, or TLB Modified) causes bits VA31...13 of the virtual address to be written into the VPN2 field of the EntryHi register. A `TLBR` instruction writes the EntryHi register with the corresponding fields from the selected TLB entry. The ASID field is written by software with the current address space identifier value and is used during the TLB comparison process to determine TLB match.

Because the ASID field is overwritten by a `TLBR` instruction, software must save and restore the value of ASID around use of the `TLBR`. This is especially important in TLB Invalid and TLB Modified exceptions, and in other memory management software.

The VPN2 field of the EntryHi register is not defined after an address error exception, and may be modified by hardware during the address error exception sequence. Software writes of the EntryHi register (via `MTC0`) do not cause the implicit write of address-related fields in the `BadVAddr` or `Context` registers.

Register 50-14: EntryHi: TLB Address Match Register; CP0 Register 10, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	VPN2<18:11>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	VPN2<10:3>							
15:8	R/W-x	R/W-x	R/W-x	U-0	U-0	U-0	U-0	U-0
	VPN2<2:0>			—	—	—	—	—
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	ASID<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-13 **VPN2<18:11>**: Virtual Page Number bits

These bits are written by hardware on a TLB exception or on a TLB read, and are written by software before a TLB write.

bit 12-8 **Unimplemented**: Read as '0'

bit 7-0 **ASID<7:0>**: Address Space Identifier bits

These bits are written by hardware on a TLB read and by software to establish the current ASID value for a TLB write and against which TLB references match each entry's TLB ASID field.

50.13.15 Compare Register (CP0 Register 11, Select 0)

The Compare register acts in conjunction with the Count register to implement a timer and timer interrupt function. Compare maintains a stable value and does not change on its own.

When the value of Count equals the value of Compare, the CPU asserts an interrupt signal to the system interrupt controller. This signal will remain asserted until Compare is written.

Register 50-15: Compare: Interval Count Compare Register; CP0 Register 11, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COMPARE<31:24>								
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COMPARE<23:16>								
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COMPARE<15:8>								
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
COMPARE<7:0>								

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **COMPARE<31:0>**: Interval Count Compare Value bits

50.13.16 Status Register (CP0 Register 12, Select 0)

The read/write Status register contains the operating mode, interrupt enabling, and the diagnostic states of the processor. The bits of this register combine to create operating modes for the processor.

50.13.16.1 INTERRUPT ENABLE

Interrupts are enabled when all of the following conditions are true:

- IE = 1
- EXL = 0
- ERL = 0
- DM = 0

If these conditions are met, the settings of the IPL bits enable the interrupts.

50.13.16.2 OPERATING MODES

If the DM bit in the Debug register is '1', the processor is in Debug mode; otherwise, the processor is in either Kernel mode or User mode. The CPU Status register bit settings shown in [Table 50-25](#) determine User or Kernel mode:

Table 50-25: CPU Status Register Bits That Determine Processor Mode

Mode	Bit/Setting		
User (requires <i>all</i> of the following bits and values)	UM = 1	EXL = 0	ERL = 0
Kernel (requires <i>one</i> or more of the following bit values)	UM = 0	EXL = 1	ERL = 1

Note: The Status register CU0 bit (Status<28>) controls Coprocessor 0 accessibility. If Coprocessor 0 is unusable, an instruction that accesses it generates an exception.

Register 50-16: Status: Status Register; CP0 Register 12, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	R/W-0	R/W-x	R/W-0	R/W-0	R/W-x	R/W-0
	—	—	CU1 ⁽²⁾	CU0	RP	FR ⁽²⁾	RE	MX ⁽³⁾
23:16	U-0	R/W-1	R/W-0, HS, CS	R/W-1	R/W-0	R/W-x	U-0	R/W-0
	—	BEV	TS ⁽¹⁾	SR	NMI	IPL<7>	—	IPL<6>
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	U-0	U-0
	IPL<5:0>						—	—
7:0	U-0	U-0	U-0	R/W-x	U-0	R/W-x	R/W-x	R/W-x
	—	—	—	UM	—	ERL	EXL	IE

Legend: HS = Set by hardware CS = Cleared by software
R = Readable bit W = Writable bit U = Unimplemented bit, read as '0'
-n = Value at POR '1' = Bit is set '0' = Bit is cleared x = Bit is unknown

bit 31-30 **Unimplemented:** Read as '0'

- Note 1:** This bit is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.
- 2:** This bit is only available on devices with the M-Class core. Refer to the specific device data sheet for availability.
- 3:** DSP ASE is not available on all devices. Refer to the “CPU” chapter in the specific device data sheet for availability.

PIC32 Family Reference Manual

Register 50-16: Status: Status Register; CP0 Register 12, Select 0 (Continued)

- bit 29 **CU1:** Coprocessor 1 Usable bit⁽²⁾
This bit allows access to Coprocessor 1. This bit can only be written if the FPU is configured. This bit will read as '0' if the FPU is not present.
1 = Access is allowed
0 = Access is not allowed
- bit 28 **CU0:** Coprocessor 0 Usable bit
This bit controls access to Coprocessor 0.
1 = Access is allowed
0 = Access is not allowed
Coprocessor 0 is always usable when the processor is running in Kernel mode, independent of the state of the CU0 bit.
- bit 27 **RP:** Reduced Power bit
1 = Enables Reduced Power mode
0 = Disables Reduced Power mode
- bit 26 **FR:** Floating Point Register Mode⁽²⁾
1 = Floating-point registers can contain any data type
0 = Floating-point registers can contain any 32-bit data type. 64-bit data types are stored in even-odd pairs of registers
Note: This bit must be ignored on write and read as '0' for CPUs with no FPU.
- bit 25 **RE:** Reverse-endian Memory Reference Enable bit
Used to enable reverse-endian memory references while the processor is running in User mode
1 = User mode uses reversed endianness
0 = User mode uses configured endianness
Debug, Kernel, or Supervisor mode references are not affected by the state of this bit.
- bit 24 **MX:** MIPS DSP Resource Enable bit⁽³⁾
This bit must be set prior to executing any DSP ASE instruction. An attempt to execute a DSP ASE instruction while this bit is cleared will result in a DSP State Disabled exception.
1 = Access is enabled
0 = Access is disabled
- bit 23 **Unimplemented:** Read as '0'
- bit 22 **BEV:** Bootstrap Exception Vector Control bit
Controls the location of exception vectors.
1 = Bootstrap
0 = Normal
- bit 21 **TS:** TLB Shutdown Control bit⁽¹⁾
Indicates that the TLB has detected a match on multiple entries. This bit is also set if a TLBWI or TLBWR instruction is issued that would cause a TLB shutdown condition if allowed to complete. A machine check exception is also issued.
1 = TLB shutdown event
0 = No TLB shutdown event
Software can only write a '0' to this bit to clear it and cannot force a '0' to '1' transition.
- bit 20 **SR:** Soft Reset bit
Indicates that the entry through the Reset exception vector was due to a Soft Reset.
1 = Soft Reset; this bit is always set for any type of reset on the PIC32 core
0 = Not used on PIC32
Software can only write a '0' to this bit to clear it and cannot force a '0' to '1' transition.

- Note 1:** This bit is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.
- 2:** This bit is only available on devices with the M-Class core. Refer to the specific device data sheet for availability.
- 3:** DSP ASE is not available on all devices. Refer to the “CPU” chapter in the specific device data sheet for availability.

Register 50-16: Status: Status Register; CP0 Register 12, Select 0 (Continued)

- bit 19 **NMI:** Non-Maskable Interrupt bit
Indicates that the entry through the reset exception vector was due to a NMI.
1 = NMI
0 = Not NMI (Soft Reset or Reset)
Software can only write a '0' to this bit to clear it and cannot force a '0' to '1' transition.
- bit 18 **IPL<7>:** Interrupt Priority Level bits
This field is the encoded (0-256) value of the current IPL. An interrupt will be signaled only if the requested IPL is higher than this value.
- bit 17 **Unimplemented:** Read as '0'
- bit 16-10 **IPL<6:0>:** Interrupt Priority Level bits
This field is the encoded (0-256) value of the current IPL. An interrupt will be signaled only if the requested IPL is higher than this value.
- bit 9-5 **Unimplemented:** Read as '0'
- bit 4 **UM:** User Mode bit
This bit denotes the base operating mode of the processor. On the encoding of this bit is:
1 = Base mode is User mode
0 = Base mode in Kernel mode
The processor can also be in Kernel mode if ERL or EXL is set, regardless of the state of the UM bit.
- bit 3 **Unimplemented:** Read as '0'
- bit 2 **ERL:** Error Level bit
Set by the processor when a Reset, Soft Reset, NMI or Cache Error exception are taken.
1 = Error level
0 = Normal level
When ERL is set:
- Processor is running in Kernel mode
 - Interrupts are disabled
 - **ERET** instruction will use the return address held in the ErrorEPC register instead of the EPC register
 - Lower 2^{29} bytes of kuseg are treated as an unmapped and uncached region. This allows main memory to be accessed in the presence of cache errors. The operation of the processor is undefined if the ERL bit is set while the processor is executing instructions from kuseg.
- bit 1 **EXL:** Exception Level bit
Set by the processor when any exception other than Reset, Soft Reset, or NMI exceptions is taken.
1 = Exception level
0 = Normal level
When EXL is set:
- Processor is running in Kernel mode
 - Interrupts are disabled
- EPC, BD, and SRSCtl will not be updated if another exception is taken.
- bit 0 **IE:** Interrupt Enable bit
Acts as the master enable for software and hardware interrupts:
1 = Interrupts are enabled
0 = Interrupts are disabled
This bit may be modified separately via the **DI** and **EI** instructions

- Note 1:** This bit is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.
- 2:** This bit is only available on devices with the M-Class core. Refer to the specific device data sheet for availability.
- 3:** DSP ASE is not available on all devices. Refer to the “CPU” chapter in the specific device data sheet for availability.

50.13.17 IntCtl: Interrupt Control Register (CP0 Register 12, Select 1)

The IntCtl register controls the vector spacing of the PIC32 architecture.

Register 50-17: IntCtl: Interrupt Control Register; CP0 Register 12, Select 1

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-1	R/W-1
	—	PF	ICE	STKDEC<4:0>				
15:8	R/W-0	R/W-0	R/W-0	U-0	U-0	U-0	R/W-0	R/W-0
	CLREXL	APE	USESTK	—	—	—	VS<4:3>	
7:0	R/W-0	R/W-0	R/W-0	U-0	U-0	U-0	U-0	U-0
	VS<2:0>			—	—	—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-23 **Unimplemented:** Read as '0'

bit 22 **PF:** Vector Prefetching Enable bit

1 = Vector Prefetching is Enabled

0 = Vector Prefetching is Disabled

bit 21 **ICE:** Interrupt Chaining Enable bit

1 = Interrupt chaining is Enabled

0 = Interrupt chaining is Disabled

bit 20-16 **STKDEC<4:0>:** Stack Pointer Decrement bits

For the Auto-Prologue feature, this is the number of 4-byte words that are decremented from the stack pointer value.

31-4 = Specifies the number of words to be decremented

3-0 = Decrement 3 words (12 bytes)

bit 15 **CLREXL:** Clear KSU/ERL/EXL bit

For the Auto-Prologue feature and IRET instruction, this bit, if set, during Auto-Prologue and IRET interrupt chaining, clears the KSU/ERL/EXL bits.

1 = Bits are cleared by these operations

0 = Bits are not cleared by these operations

bit 14 **APE:** Auto-Prologue Enable bit

1 = Auto-Prologue is enabled

0 = Auto-Prologue is disabled

bit 13 **USEKSTK:** Stack Use bit

Chooses which Stack to use during Interrupt Auto-Prologue.

1 = Use r29 of the Current SRS at the beginning of IAP

Used for environments where there are separate User mode and Kernel mode stacks. In this case, r29 of the SRS used during IAP must be preinitialized by software to hold the Kernel mode stack pointer.

0 = Copy r29 of the Previous SRS to the Current SRS at the beginning of IAP

Used for Bare-Iron environments with only one stack.

bit 12-10 **Unimplemented:** Read as '0'

Register 50-17: IntCtl: Interrupt Control Register; CP0 Register 12, Select 1 (Continued)

bit 9-5 **VS<4:0>**: Vector Spacing bits

These bits specify the spacing between each interrupt vector.

Encoding	Spacing Between Vectors (hex)	Spacing Between Vectors (decimal)
0x00	0x000	0
0x01	0x020	32
0x02	0x040	64
0x04	0x080	128
0x08	0x100	256
0x10	0x200	512

All other values are reserved. The operation of the processor is undefined if a reserved value is written to these bits.

bit 4-0 **Unimplemented**: Read as '0'

PIC32 Family Reference Manual

50.13.18 SRSCtl Register (CP0 Register 12, Select 2)

The SRSCtl register controls the operation of GPR shadow sets in the processor.

Table 50-26: Sources for New CSS on an Exception or Interrupt

Exception Type	Bit Source	Condition	Comment
Exception	ESS	All	—
Non-Vectored Interrupt	ESS	IV bit = 0 (Cause<23>)	Treat as exception
Vectored EIC Interrupt	EICSS	IV bit = 1 (Cause<23>) and, VEIC bit = 1 (Config3<6>)	Source is external interrupt controller.

Register 50-18: SRSCtl: Shadow Register Set Register; CP0 Register 12, Select 2

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	R-0	R-1	R-1	R-1	U-0	U-0
	—	—	HSS<3:0>				—	—
23:16	U-0	U-0	R-x	R-x	R-x	R-x	U-0	U-0
	—	—	EICSS<3:0>				—	—
15:8	R/W-0	R/W-0	R/W-0	R/W-0	U-0	U-0	R/W-0	R/W-0
	ESS<3:0>				—	—	PSS<3:2>	
7:0	R/W-0	R/W-0	U-0	U-0	R-0	R-0	R-0	R-0
	PSS<1:0>		—	—	CSS<3:0>			

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-30 **Unimplemented:** Read as '0'

bit 29-26 **HSS<3:0>:** High Shadow Set bits

This bit contains the highest shadow set number that is implemented by this processor. A value of '0000' in these bits indicates that only the normal GPRs are implemented.

1111 = Reserved

0111 = Eight shadow sets are present

0110 = Reserved

0101 = Reserved

0100 = Reserved

0011 = Four shadow sets are present

0010 = Reserved

0001 = Two shadow sets are present

0000 = One shadow set (normal GPR set) is present

The value in this bit also represents the highest value that can be written to the ESS<3:0>, EICSS<3:0>, PSS<3:2>, and CSS<3:0> bits of this register, or to any of the bits of the SRSSMap register. The operation of the processor is undefined if a value larger than the one in this bit is written to any of these other bits.

bit 25-22 **Unimplemented:** Read as '0'

bit 21-18 **EICSS<3:0>:** External Interrupt Controller Shadow Set bits

EIC Interrupt mode shadow set. This bit is loaded from the external interrupt controller for each interrupt request and is used in place of the SRSSMap register to select the current shadow set for the interrupt.

bit 17-16 **Unimplemented:** Read as '0'

bit 15-12 **ESS<3:0>:** Exception Shadow Set bits

This bit specifies the shadow set to use on entry to Kernel mode caused by any exception other than a vectored interrupt. The operation of the processor is undefined if software writes a value into this bit that is greater than the value in the HSS<3:0> bits.

bit 11-10 **Unimplemented:** Read as '0'

Register 50-18: SRSCtl: Shadow Register Set Register; CP0 Register 12, Select 2 (Continued)

bit 9-6 **PSS<3:0>**: Previous Shadow Set bits

Since GPR shadow registers are implemented, this bit is copied from the CSS bit when an exception or interrupt occurs. An `ERET` instruction copies this value back into the CSS bit if the BEV bit (`Status<22>`) = 0.

This bit is not updated on any exception which sets the ERL bit (`Status<2>`) to '1' (i.e., Reset, Soft Reset, NMI, cache error), an entry into EJTAG Debug mode, or any exception or interrupt that occurs with the EXL bit (`Status<1>`) = 1, or BEV = 1. This bit is not updated on an exception that occurs while ERL = 1.

The operation of the processor is undefined if software writes a value into this bit that is greater than the value in the HSS<3:0> bits.

bit 5-4 **Unimplemented**: Read as '0'

bit 3-0 **CSS<3:0>**: Current Shadow Set bits

Since GPR shadow registers are implemented, this bit is the number of the current GPR set. This bit is updated with a new value on any interrupt or exception, and restored from the PSS bit on an `ERET`.

[Table 50-26](#) describes the various sources from which the CSS<3:0> bits are updated on an exception or interrupt.

This bit is not updated on any exception which sets the ERL bit (`Status<2>`) to '1' (i.e., Reset, Soft Reset, NMI, cache error), an entry into EJTAG Debug mode, or any exception or interrupt that occurs with EXL bit (`Status<1>`) = 1, or BEV = 1. Neither is it updated on an `ERET` with ERL = 1 or BEV = 1. This bit is not updated on an exception that occurs while ERL = 1.

The value of the CSS<3:0> bits can be changed directly by software only by writing the PSS<3:0> bits and executing an `ERET` instruction.

50.13.19 SRSSMap: Register (CP0 Register 12, Select 3)

The SRSSMap register contains eight 4-bit fields that provide the mapping from a vector number to the shadow set number to use when servicing such an interrupt. The values from this register are not used for a non-interrupt exception, or a non-vectored interrupt (IV bit = 0, Cause<23> or VS<4:0> bit = 0, IntCtl<9:5>). In such cases, the shadow set number comes from the ESS<3:0> bits (SRSCtl<15:12>).

If the HSS<3:0> bits (SRSCtl<29:26>) are '0', the results of a software read or write of this register are unpredictable.

The operation of the processor is undefined if a value is written to any bit in this register that is greater than the value of the HSS<3:0> bits.

The SRSSMap register contains the shadow register set numbers for vector numbers 7-0. The same shadow set number can be established for multiple interrupt vectors, creating a many-to-one mapping from a vector to a single shadow register set number.

Register 50-19: SRSSMap: Shadow Register Set Map Register; CP0 Register 12, Select 3

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	SSV7<3:0>				SSV6<3:0>			
23:16	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	SSV5<3:0>				SSV4<3:0>			
15:8	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	SSV3<3:0>				SSV2<3:0>			
7:0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	SSV1<3:0>				SSV0<3:0>			

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

- bit 31-28 **SSV7<3:0>**: Shadow Set Vector 7 bits
Shadow register set number for Vector Number 7.
- bit 27-24 **SSV6<3:0>**: Shadow Set Vector 6 bits
Shadow register set number for Vector Number 6.
- bit 23-20 **SSV5<3:0>**: Shadow Set Vector 5 bits
Shadow register set number for Vector Number 5.
- bit 19-16 **SSV4<3:0>**: Shadow Set Vector 4 bits
Shadow register set number for Vector Number 4.
- bit 15-12 **SSV3<3:0>**: Shadow Set Vector 3 bits
Shadow register set number for Vector Number 3.
- bit 11-8 **SSV2<3:0>**: Shadow Set Vector 2 bits
Shadow register set number for Vector Number 2.
- bit 7-4 **SSV1<3:0>**: Shadow Set Vector 1 bits
Shadow register set number for Vector Number 1.
- bit 3-0 **SSV0<3:0>**: Shadow Set Vector 0 bit
Shadow register set number for Vector Number 0.

50.13.20 View_IPL Register (CP0 Register 12, Select 4)

This register gives read and write access to the IPL<7:0> bits that are also available in the Status register. The use of this register allows the Priority Level to be read/written without extracting/inserting that field from/to the Status register.

The IPL field is located in non-contiguous bits within the Status register. All of the IPL bits are presented as a contiguous field within this register.

Register 50-20: View_IPL: View Interrupt Priority Level Register; CP0 Register 12, Select 4

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	R/W-0	R/W-0
	—	—	—	—	—	—	IPL<7:6>	
7:0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	U-0	U-0
	IPL<5:0>						—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-10 **Unimplemented:** Read as '0'

bit 9-2 **IPL<7:0>:** These bits contain the encoded (0...256) value of the current IPL.

bit 1-0 **Unimplemented:** Read as '0'

50.13.21 SRSMAP2 Register (CP0 Register 12, Select 5)

The SRSMAP2 register contains two 4-bit bits that provide the mapping from an vector number to the shadow set number to use when servicing such an interrupt. The values from this register are not used for a non-interrupt exception, or a non-vectored interrupt (IV bit (Cause<23>) = 0 or the VS<4:0> bits (IntCtl<9:5>) = 0). In such cases, the shadow set number comes from the ESS<3:0> bits (SRSCtl<15:12>).

If the HSS<3:0> bits (SRSCtl<9:6>) are '0', the results of a software read or write of this register are unpredictable.

The operation of the processor is undefined if a value is written to any bit in this register that is greater than the value of the HSS<3:0> bits.

The SRSMAP2 register contains the shadow register set numbers for vector numbers 9:8. The same shadow set number can be established for multiple interrupt vectors, creating a many-to-one mapping from a vector to a single shadow register set number.

Register 50-21: SRSMAP2: Shadow Register Set Map 2 Register; CP0 Register 12, Select 5

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	SSV9<3:0>				SSV8<3:0>			

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-8 **Unimplemented:** Read as '0'

bit 7-4 **SSV9<3:0>:** Shadow Set Vector 9 bits

Shadow register set number for Vector Number 9.

bit 3-0 **SSV8<3:0>:** Shadow Set Vector 8 bits

Shadow register set number for Vector Number 8.

50.13.22 Cause Register (CP0 Register 13, Select 0)

The Cause register primarily describes the cause of the most recent exception. In addition, bits also control software interrupt requests and the vector through which interrupts are dispatched. With the exception of the IP1, IP0, DC, and IV bits, all bits in the Cause register are read-only.

These are the causes of the exceptions, such as `void __attribute__((nomips16)) _general_exception_handler(void)`!

Table 50-27: Cause Register EXCCODE<4:0> Bits

Exception Code Value		Mnemonic	Description
Decimal	Hex		
0	0x00	Int	Interrupt
1	0x01	MOD ⁽¹⁾	TLB modified exception
2	0x02	TLBL ⁽¹⁾	TLB exception (load or instruction fetch)
3	0x03	TLBS ⁽¹⁾	TLB exception (store)
4	0x04	AdEL	Address error exception (load or instruction fetch)
5	0x05	AdES	Address error exception (store)
6	0x06	IBE	Bus error exception (instruction fetch)
7	0x07	DBE	Bus error exception (data reference: load or store)
8	0x08	Sys	Syscall exception
9	0x09	Bp	Breakpoint exception
10	0x0a	RI	Reserved instruction exception
11	0x0B	CPU	Coprocessor Unusable exception
12	0x0C	Ov	Arithmetic Overflow exception
13	0x0D	Tr	Trap exception
14	0x0E	—	Reserved
15	0x0F	FPE ⁽²⁾	Floating Point exception
16-18	0x10-0x0x12		Reserved
19	0x13	TLBRI ⁽¹⁾	TLB read-inhibit
20	0x14	TLBEI ⁽¹⁾	TLB execute-inhibit
21-22	0x15-0x16	—	Reserved
23	0x17	WATCH ⁽¹⁾	Reference to WatchHi/WatchLo address
24	0x18	MCheck ⁽¹⁾	Machine check
25	0x19	—	Reserved
26	0x1A	DSPDis	DSP ASE state disabled exception ⁽³⁾
27-31	0x1B-0x1F	—	Reserved

Note 1: This feature is only available on PIC32 devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

2: This feature is only available on PIC32 devices with the M-Class core. Refer to the “FPU” chapter in the specific device data sheet to determine availability.

3: DSP ASE is not available on all devices. Please consult the “CPU” chapter of the specific device data sheet to determine availability

PIC32 Family Reference Manual

Register 50-22: Cause: **Exception Cause Register; CP0 Register 13, Select 0**

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-x	R-x	R-x	R-x	R/W-0	R-0	R-x	R-x
	BD	TI	CE<1:0>		DC	PCI	IC	AP
23:16	R/W-x	R/W-x	R-x	U-0	U-0	U-0	R-x	R-x
	IV	WP ⁽¹⁾	FDCI	—	—	—	RIPL<7:6>	
15:8	R-x	R-x	R-x	R-x	R-x	R-x	R/W-x	R/W-x
	RIPL<5:0>						IP1	IP0
7:0	U-0	R-x	R-x	R-x	R-x	R-x	U-0	U-0
	—	EXCCODE<4:0>					—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **BD:** Branch Delay bit

Indicates whether the last exception taken occurred in a branch delay slot:

1 = In delay slot

0 = Not in delay slot

The processor updates BD only if the EXL bit (Status<1>) was '0' when the exception occurred.

bit 30 **TI:** Timer Interrupt bit

Timer Interrupt. This bit denotes whether a timer interrupt is pending (analogous to the IP bits for other interrupt types):

1 = Timer interrupt is pending

0 = No timer interrupt is pending

bit 29-28 **CE<1:0>:** Coprocessor Exception bits

Coprocessor unit number referenced when a Coprocessor Unusable exception is taken. This bit is loaded by hardware on every exception, but is unpredictable for all exceptions except for Coprocessor Unusable.

11 = Reserved

10 = Reserved

01 = Reserved

00 = Coprocessor 0

bit 27 **DC:** Disable Count Register bit

In some power-sensitive applications, the Count register is not used and can be stopped to avoid unnecessary toggling.

1 = Disable counting of Count register

0 = Enable counting of Count register

bit 26 **PCI:** Performance Counter Interrupt bit

1 = Performance counter interrupt is pending

0 = No performance counter interrupt is pending

bit 25 **IC:** Interrupt Chaining bit

Indicates if Interrupt chaining occurred on the last IRET instruction.

1 = Interrupt Chaining occurred during last IRET instruction

0 = Interrupt Chaining did not happen on last IRET instruction

bit 24 **AP:** Interrupt Auto-Prologue Exception bit

Indicates whether an exception occurred during Interrupt Auto-Prologue.

1 = Exception occurred during Auto-Prologue operation

0 = Exception did not occur during Auto-Prologue operation

Note 1: This bit is only available on PIC32 devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

Register 50-22: Cause: Exception Cause Register; CP0 Register 13, Select 0 (Continued)

bit 23 **IV:** Interrupt Vector bit

Indicates whether an interrupt exception uses the general exception vector or a special interrupt vector

1 = Use the special interrupt vector (0x200)

0 = Use the general exception vector (0x180)

If the IV bit (Cause<23>) is '1' and the BEV bit (Status<22>) is '0', the special interrupt vector represents the base of the vectored interrupt table.

bit 22 **WP:** Watch Exception Pending bit⁽¹⁾

This bit indicates that a watch exception was deferred because the status bits EXL or ERL were set at the time the watch exception was detected.

1 = Watch exception pending

0 = Watch exception not pending

When set, this bit indicates a pending watch exception, and causes the exception to be initiated once the Status EXL and ERL bits are both zero. The software must clear this bit as part of the watch exception handler to prevent a watch exception loop.

Note: Software should not write a '1' to this bit when its value is '0', thereby causing a 0-to-1 transition. If such a transition is caused by software, the results are unpredictable.

bit 21 **FDCI:** Fast Debug Channel Interrupt bit

This bit indicates that a FDC interrupt is pending

1 = Fast Debug Channel interrupt is pending

0 = Fast Debug Channel interrupt is not pending

bit 20-18 **Unimplemented:** Read as '0'

bit 17-10 **RIPL<7:0>:** Requested Interrupt Priority Level bits

This bit is the encoded (255-0) value of the requested interrupt. A value of '0' indicates that no interrupt is requested.

bit 9-8 **IP<1:0>:** Software Interrupt Request Control bits

Controls the request for software interrupts

1 = Request software interrupt

0 = No interrupt requested

These bits are exported to the system interrupt controller for prioritization in EIC Interrupt mode with other interrupt sources.

bit 7 **Unimplemented:** Read as '0'

bit 6-2 **EXCCODE<4:0>:** Exception Code bits

See Table 50-27 for the list of Exception codes.

bit 1-0 **Unimplemented:** Read as '0'

Note 1: This bit is only available on PIC32 devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

50.13.23 View_RIPL Register (CP0 Register 13, Select 4)

This register gives read access to the RIPL bit that is also available in the Cause register. The use of this register allows the Requested Priority Level to be read without extracting that bit from the Cause register.

Register 50-23: View_RIPL: View Requested Priority Level Register; CP0 Register 13, Select 4

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	R-0	R-0
	—	—	—	—	—	—	RIPL<7:6>	
7:0	R-0	R-0	R-0	R-0	R-0	R-0	R/W-0	R/W-0
	RIPL<5:0>						IP<1:0>	

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-10 **Unimplemented:** Read as '0'

bit 9-2 **RIPL<7:0>:** Requested Interrupt Priority Level bits

If EIC Interrupt mode is enabled, this bit indicates the encoded (0...255) value of the current Requested Priority Level of the pending interrupt.

bit 1-0 **IP<1:0>:** Software Interrupt Pending bits

If EIC Interrupt mode is not enabled, controls which SW interrupts are pending.

50.13.24 NestedExc Register (CP0 Register 13, Select 5)

The NestedExc register is a read-only register that contains the values of Status<1> (EXL) and Status<2> (ERL) prior to acceptance of the current exception.

Register 50-24: NestedExc: Nested Exception Register; CP0 Register 13, Select 5

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	R/W-x	R/W-x	U-0
	—	—	—	—	—	NERL	NEXL	—

Legend:	r = Reserved bit	
R = Readable bit	W = Writable bit	U = Unimplemented bit, read as '0'
-n = Value at POR	'1' = Bit is set	'0' = Bit is cleared x = Bit is unknown

bit 31-3 **Unimplemented:** Read as '0'

bit 2 **NERL:** Nested Error Level bit

This bit contains the value of the ERL bit prior to acceptance of the current exception. This bit is updated by all exceptions that would set either the EXL bit or the ERL bit in the status register. This bit is not updated by Debug exceptions.

bit 1 **NEXL:** Nested Exception Level bit

This bit contains the value of the EXL bit prior to acceptance of current exception. This bit is updated by exceptions that would update the exception program counter if the EXL bit is not set (MCheck, interrupt, Address Error, all TLB exceptions, Bus Error, CopUnusable, Reserved Instruction, Overflow, Trap, Syscall, FPU, etc.). For these exception types, this register field is updated regardless of the value of StatusEXL. This bit is not updated by exception types that update ErrorEPC (Reset, Soft Reset, NMI, and Cache Error). In addition, this bit is not updated by Debug exceptions.

bit 0 **Unimplemented:** Read as '0'

50.13.25 EPC Register (CP0 Register 14, Select 0)

The Exception Program Counter (EPC) is a read/write register that contains the address at which processing resumes after an exception has been serviced. All bits of the EPC register are significant and are writable.

For synchronous (precise) exceptions, the EPC contains one of the following:

- The virtual address of the instruction that was the direct cause of the exception
- The virtual address of the immediately preceding **BRANCH** or **JUMP** instruction, when the exception causing instruction is in a branch delay slot and the Branch Delay bit in the Cause register is set

On new exceptions, the processor does not write to the EPC register when the EXL bit in the Status register is set; however, the register can still be written through the **MTC0** instruction.

Since the PIC32 family implements MIPS16e® or microMIPS ASE, a read of the EPC register (via **MFC0**) returns the following value in the destination GPR:

$$\text{GPR[rt]} \leftarrow \text{ExceptionPC}_{31..1} \parallel \text{ISAMode}_0$$

That is, the upper 31 bits of the exception PC are combined with the lower bit of the ISA<1:0> bits (Config3<15:14>) and are written to the GPR.

Similarly, a write to the EPC register (via **MTC0**) takes the value from the GPR and distributes that value to the exception PC and the ISA<1:0> bits (Config3<15:14>), as follows:

$$\begin{aligned} \text{ExceptionPC} &\leftarrow \text{GPR[rt]}_{31..1} \parallel 0 \\ \text{ISAMode} &\leftarrow 2\#0 \parallel \text{GPR[rt]}_0 \end{aligned}$$

That is, the upper 31 bits of the GPR are written to the upper 31 bits of the exception PC, and the lower bit of the exception PC is cleared. The upper bit of the ISA<1:0> bits (Config3<15:14>) is cleared and the lower bit is loaded from the lower bit of the GPR.

Register 50-25: EPC: Exception Program Counter Register; CP0 Register 14, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<31:24>								
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<23:16>								
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<15:8>								
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<7:0>								

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **EPC<31:0>**: Exception Program Counter bits

50.13.26 NestedEPC Register (CP0 Register 14, Select 2)

The NestedEPC register is a read/write register with the same behavior as the EPC register, with the following exceptions:

- The NestedEPC register ignores the value of Status<1> (EXL) and is therefore updated on the occurrence of any exception, including nested exceptions
- The NestedEPC register is not used by the `ERET/DERET/IERET` instructions. To return to the address stored in NestedEPC, software must copy the value of the NestedEPC register to the EPC register.

Register 50-26: NestedEPC: Nested Exception Program Counter Register; CP0 Register 14, Select 2

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<31:24>								
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<23:16>								
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<15:8>								
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
EPC<7:0>								

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **EPC<31:0>**: Nested Exception Program Counter bits

These bits are updated by exceptions that would update the exception program counter if the EXL bit is not set (MCheck, Interrupt, Address Error, all TLB exceptions, Bus Error, CopUnusable, Reserved Instruction, Overflow, Trap, Syscall, FPU, etc.). For these exception types, this register field is updated regardless of the value of EXL. These bits are not updated by exception types that update the exception program counter (Reset, Soft Reset, NMI, and Cache Error). In addition, these bits are not updated by Debug exceptions.

50.13.27 PRID Register (CP0 Register 15, Select 0)

The Processor Identification (PRID) register is a 32-bit read-only register that contains information identifying the manufacturer, manufacturer options, processor identification, and revision level of the processor.

Register 50-27: PRID: Processor Identification Register; CP0 Register 15, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	R-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	R-0	R-0	R-0	R-0	R-0	R-0	R-0	R-1
	COMPANYID<23:16>							
15:8	R-1	R-0	R-0	R-1	R-1	R-1	R-1	R-0
	PROCESSORID<15:8>							
7:0	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	MAJORREV<2:0>			MINORREV<2:0>			PATCHREV<1:0>	

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-24 **Unimplemented:** Read as '0'

bit 23-16 **COMPANYID<23:16>:** Company Identification bits

In PIC32 devices, these bits contain a value of '1' to indicate Imagination Technologies Ltd. as the processor manufacturer/designer.

bit 15-8 **PROCESSORID<15:8>:** Processor Identification bits

These bits allow software to distinguish between the various types of Imagination Technologies Ltd. processors.

0x9E = Microprocessor core

bit 7-5 **MAJORREV<2:0>:** Processor Major Revision Identification bits

These bits allow software to distinguish between one revision and another of the same processor type. This number is increased on major revisions of the processor core.

bit 4-2 **MINORREV<2:0>:** Processor Minor Revision Identification bits

This number is increased on each incremental revision of the processor and reset on each new major revision.

bit 1-0 **PATCHREV<1:0>:** Processor Patch Level Identification bits

If a patch is made to modify an older revision of the processor, the value of these bits will be incremented.

50.13.28 Ebase Register (CP0 Register 15, Select 1)

The Ebase register is a read/write register containing the base address of the exception vectors used when the BEV bit (Status<22>) equals '0', and a read-only CPU number value that may be used by software to distinguish different processors in a multi-processor system.

The Ebase register provides the ability for software to identify the specific processor within a multi-processor system, and allows the exception vectors for each processor to be different, especially in systems composed of heterogeneous processors. Bits 31-12 of the Ebase register are concatenated with zeros to form the base of the exception vectors when the BEV bit is '0'. The exception vector base address comes from the fixed defaults when the BEV bit is '1', or for any EJTAG Debug exception. The Reset state of bits 31-12 of the Ebase register initialize the exception base register to 0x80000000.

Bits 31 and 30 of the Ebase Register are fixed with the value 2#10 to force the exception base address to be in the kseg0 or kseg1 unmapped virtual address segments.

If the value of the exception base register is to be changed, this must be done with the BEV bit equal to '1'. The operation of the processor is undefined if the Ebase<17:0> bits are written with a different value when the BEV bit (Status<22>) is '0'.

Combining bits 31-20 of the Ebase register allows the base address of the exception vectors to be placed at any 4 KB page boundary.

Register 50-28: Ebase: Exception Base Register; CP0 Register 15, Select 1

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-1	U-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	—	—	EBASE<17:12>					
23:16	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	EBASE<11:4>							
15:8	R/W-0	R/W-0	R/W-0	R/W-0	U-0	U-0	R-0	R-0
	EBASE<3:0>				—	—	CPUNUM<9:8>	
7:0	R-0	R-0	R-0	R-0	R-0	R-0	R-0	R-0
	CPUNUM<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **Unimplemented:** Read as '1'

bit 30 **Unimplemented:** Read as '0'

bit 29-12 **EBASE<17:0>:** Exception Vector Base Address bits

In conjunction with bits 31-30, these bits specify the base address of the exception vectors when the BEV bit (Status<22>) is '0'.

bit 11-10 **Unimplemented:** Read as '0'

bit 9-0 **CPUNUM<9:0>:** CPU Number bits

These bits specify the number of CPUs in a multi-processor system and can be used by software to distinguish a particular processor from others. In a single processor system, this value is set to '0'.

50.13.29 CDMMBase Register (CP0 Register 15, Select 2)

The 36-bit physical base address for the Common Device Memory Map (CDMM) is defined by this register. Since the PIC32 is a 32-bit device, the upper four bits of the address are zero. The CDMM is a region of physical address space that is reserved for mapping I/O device Configuration registers within a MIPS processor. The CDMM helps aggregate various device mappings into one area, preventing fragmentation of the memory address space. On PIC32 devices, the CDMM contains the Fast Debug Channel (FDC) control registers.

Register 50-29: CDMMBase: Common Device Memory Map Base Register; CP0 Register 15, Select 2

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	CDMMUA<20:13>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	CDMMUA<12:5>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-0	R-0	R-0
	CDMMUA<4:0>					EN	CI	CDMMSize<8>
7:0	R-0	R-0	R-0	R-0	R-0	R-0	R-1	R-0
	CDMMSize<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-11 **CDMMUA<20:0>**: CDMM Upper Address bits

Bits 35-15 of the base physical address of the memory mapped registers. Unimplemented bits ignore writes and return '0' on reads.

bit 10 **EN**: CDMM Enable bit

1 = CDMM region is enabled

0 = CDMM region is disabled

bit 9 **CI**: Device Register Block Reserved bit

1 = The first 64-byte Device Register Block of the CDMM is reserved for additional registers that manage CDMM region behavior and are not IO device registers

0 = The first 64-byte Device Register Block of the CDMM is not reserved

bit 8-0 **CDMMSize<8:0>**: CDMM Size bits

These bits indicate the number of 64-byte Device Register Blocks instantiated in the CPU core.

11111111 = 512 Device Register Blocks

.

.

.

00000000 = 1 Device Register Block

50.13.30 Config Register (CP0 Register 16, Select 0)

The Config register specifies various configuration and capabilities information. Most of the fields in the Config register are initialized by hardware during the Reset exception process, or are constant.

Note: The Core Configuration register bit fields and default values may vary between devices. Refer to the “**Core Configuration**” chapter in the specific device data sheet for the core configuration of a specific device.

Register 50-30: Config: Configuration Register; CP0 Register 16, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	r-1 —	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R-x
		K23<2:0> ⁽¹⁾			KU<2:0> ⁽¹⁾			ISP
23:16	R-x DSP	R-x UDI	R-x SB	R-x MDU	U-0 —	R-x MM<1:0> ⁽²⁾	R-x	R-x DS ⁽¹⁾ BM ⁽²⁾
15:8	R-0 BE	R-0	R-0	R-0	R-0	R-1	R-0	R-x
		AT<1:0>		AR<2:0>			MT<2:1>	
7:0	R-x MT<0>	U-0 —	U-0 —	U-0 —	U-0 —	R/W-x	R/W-x	R/W-x K0<2:0>

Legend:	r = Reserved bit
R = Readable bit	W = Writable bit
-n = Value at POR	‘1’ = Bit is set
	U = Unimplemented bit, read as ‘0’
	‘0’ = Bit is cleared
	x = Bit is unknown

bit 31 **Reserved:** This bit is hardwired to ‘1’ to indicate the presence of the Config1 register.

bit 30-28 **K23<2:0>**: kseg2 and kseg3 bits⁽¹⁾

These bits control the cacheability of the kseg2 and kseg3 address segments. Refer to [Table 50-28](#) for the bit encoding.

bit 27-25 **KU<2:0>**: kuseg and useg bits⁽¹⁾

These bits control the cacheability of the kuseg and useg address segments. Refer to [Table 50-28](#) for the bit encoding.

bit 24 **ISP:** Instruction Scratchpad RAM bit

This bit indicates whether instruction scratchpad RAM is implemented.

1 = Instruction scratchpad RAM is implemented

0 = Instruction scratchpad RAM is not implemented

bit 23 **DSP:** Data Scratchpad RAM bit

This bit indicates whether data scratchpad RAM is implemented.

1 = Data scratchpad RAM is implemented

0 = Data scratchpad RAM is not implemented

Note 1: This bit is only available on devices with the MCU Microprocessor core. On devices with the MPU Microprocessor core, this bit is reserved, and is always read as ‘0’. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.

2: This bit is only available on devices with the MPU Microprocessor core. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.

PIC32 Family Reference Manual

Register 50-30: Config: Configuration Register; CP0 Register 16, Select 0 (Continued)

- bit 22 **UDI**: User-defined bit
This bit indicates that CorExtend User-Defined Instructions have been implemented.
1 = User-defined instructions are implemented
0 = No user-defined instructions are implemented
- bit 21 **SB**: SimpleBE bit
This bit indicates whether SimpleBE Bus mode is enabled.
1 = Only simple byte enables allowed on internal bus interface
0 = No reserved byte enables on internal bus interface
This bit is hardwired to '1' to indicate only simple byte enables allowed on internal bus interface.
- bit 20 **MDU**: Multiply/Divide Unit bit
1 = Iterative, area-efficient MDU
0 = Fast, high-performance MDU
This bit is hardwired to '0' to indicate the fast, high-performance MDU.
- bit 19 **Unimplemented**: Read as '0'
- bit 18-17 **MM<1:0>**: Merge Mode bits⁽²⁾
11 = Reserved
10 = Merging is allowed
01 = Reserved
00 = Merging is not allowed
- bit 16 **DS**: Dual SRAM bit⁽¹⁾
1 = Dual instruction/data SRAM internal bus interfaces
0 = Unified instruction/data SRAM internal bus interface
BM: Burst Mode bit⁽²⁾
This bit is hardwired to a '0' to indicate burst order is sequential.
- bit 15 **BE**: Big-Endian bit
Indicates the endian mode in which the processor is running, PIC32 is always little-endian.
1 = Big-endian
0 = Little-endian
- bit 14-13 **AT<1:0>**: Architecture Type bits
Architecture type implemented by the processor. This bit is always '00' to indicate the MIPS32 architecture.
- bit 12-10 **AR<2:0>**: Architecture Revision Level bits
Architecture revision level. This bit is always '001' to indicate MIPS32 Release 2.
111 = Reserved
110 = Reserved
101 = Reserved
100 = Reserved
011 = Reserved
010 = Reserved
001 = Release 2
000 = Release 1

Note 1: This bit is only available on devices with the MCU Microprocessor core. On devices with the MPU Microprocessor core, this bit is reserved, and is always read as '0'. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.

2: This bit is only available on devices with the MPU Microprocessor core. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.

Register 50-30: Config: Configuration Register; CP0 Register 16, Select 0 (Continued)

bit 9-7 **MT<2:0>**: MMU Type bits

PIC32 devices with the MCU Microprocessor core use a fixed-mapping MMU.

PIC32 devices with the MPU Microprocessor core use a TLB-based MMU.

111 = Reserved

110 = Reserved

101 = Reserved

100 = Reserved

011 = Fixed mapping

010 = Reserved

001 = Standard TLB

000 = Reserved

bit 6-3 **Unimplemented**: Read as '0'

bit 2-0 **K0<2:0>**: Kseg0 bits

Kseg0 coherency algorithm. Refer to [Table 50-28](#) for the bit encoding.

- Note 1:** This bit is only available on devices with the MCU Microprocessor core. On devices with the MPU Microprocessor core, this bit is reserved, and is always read as '0'. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.
- 2:** This bit is only available on devices with the MPU Microprocessor core. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.

Table 50-28: Cache Coherency Attributes

K0<2:0> Value	Cache Coherency Attribute
011	Cacheable, non-coherent, write-back, write allocate
010	Uncached
001	Cacheable, non-coherent, write-through, write allocate
000	Cacheable, non-coherent, write-through, no write allocate

50.13.31 Config1 Register (CP0 Register 16, Select 1)

The Config1 register is an adjunct to the Config register and encodes additional information about capabilities present on the core. All fields in the Config1 register are read-only.

Note: The Core Configuration register bit fields and default values may vary between devices. Refer to the “**Core Configuration**” chapter in the specific device data sheet for the core configuration of a specific device.

Register 50-31: Config1: Configuration Register 1; CP0 Register 16, Select 1

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	r-1	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	—	MMU Size<5:0> ⁽¹⁾						IS<2>
23:16	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	IS<1:0> ⁽¹⁾		IL<2:0> ⁽¹⁾			IA<2:0> ⁽¹⁾		
15:8	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	DS<2:0> ⁽¹⁾			DL<2:0> ⁽¹⁾			DA<2:1> ⁽¹⁾	
7:0	R-x	U-0	U-0	R-x	R-x	R-x	R-x	R-x
	DA<0>	—	—	PC	WR	CA	EP	FP

Legend:

r = Reserved bit
R = Readable bit
W = Writable bit
U = Unimplemented bit, read as '0'
-n = Value at POR
'1' = Bit is set
'0' = Bit is cleared
x = Bit is unknown

bit 31 **Reserved:** This bit is hardwired to a '1' to indicate the presence of the Config2 register.

bit 30-25 **MMU Size<5:0>**: Contains the number of TLB entries minus 1⁽¹⁾

11111 = 32 TLB entries

•
•
•

00000 = No TLB entries

bit 24-22 **IS<2:0>**: Instruction Cache Sets bits⁽¹⁾

Contains the number of instruction cache sets per way.

0x0 = 64

0x1 = 128

0x2 = 256

0x3 = 512

0x4 = 1024

0x5 = Reserved

0x6 = Reserved

0x7 = Reserved

bit 21-19 **IL<2:0>**: Instruction-Cache Line bits⁽¹⁾

Contains the instruction cache line size.

0x0 = No instruction cache is present

0x3 = 16 bytes

All other values are Reserved.

Note 1: For the PIC32 devices with the MCU Microprocessor core, these bits are reserved and always read as '0'. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.

Register 50-31: Config1: Configuration Register 1; CP0 Register 16, Select 1 (Continued)

- bit 18-16 **IA<2:0>**: Instruction-Cache Associativity bits⁽¹⁾
 Contains the level of instruction cache associativity.
 0x0 = Direct-mapped
 0x1 = 2-way
 0x2 = 3-way
 0x3 = 4-way
 0x4 = Reserved
 0x5 = Reserved
 0x6 = Reserved
 0x7 = Reserved
- bit 15-13 **DS<2:0>**: Data-Cache Sets bits⁽¹⁾
 Contains the number of data cache sets per way.
 0x0 = 64
 0x1 = 128
 0x2 = 256
 0x3 = 512
 0x4 = 1024
 0x5 = Reserved
 0x6 = Reserved
 0x7 = Reserved
- bit 12-10 **DL<2:0>**: Data-Cache Line bits⁽¹⁾
 Contains the data cache line size.
 0x0 = No instruction cache is present
 0x3 = 16 bytes
 All other values are Reserved.
- bit 9-7 **DA<2:0>**: Data-Cache Associativity bits⁽¹⁾
 Contains the type of set associativity for the data cache.
 0x0 = Direct-mapped
 0x1 = 2-way
 0x2 = 3-way
 0x3 = 4-way
 0x4 = Reserved
 0x5 = Reserved
 0x6 = Reserved
 0x7 = Reserved
- bit 6-5 **Unimplemented**: Read as '0'
- bit 4 **PC**: Performance Counter bit
 Performance Counter registers implemented.
 1 = The processor core contains Performance Counters
 0 = The processor core does not contain Performance Counters
- bit 3 **WR**: Watch Register Presence bit
 1 = No Watch registers are present
 0 = One or more Watch registers is present
- bit 2 **CA**: Code Compression Implemented bit
 1 = MIPS16e is implemented
 0 = No MIPS16e present
- bit 1 **EP**: EJTAG Present bit
 1 = EJTAG unit is implemented
 0 = EJTAG unit is not implemented
- bit 0 **FP**: Floating Point Unit bit
 1 = Floating Point Unit is implemented
 0 = Floating Point Unit is not implemented

Note 1: For the PIC32 devices with the MCU Microprocessor core, these bits are reserved and always read as '0'. Refer to the “CPU” chapter in the specific device data sheet to determine availability.

50.13.32 Config2 (CP0 Register 16, Select 2)

The Config2 register is an adjunct to the Config register and is reserved to encode additional capabilities information. Config2 is allocated for showing the configuration of level 2/3 caches. These bits are reset to '0' because L2/L3 caches are not supported by the PIC32 core. All bits in the Config2 register are read-only.

Note: The Core Configuration register bit fields and default values may vary between devices. Refer to the “**Core Configuration**” chapter in the specific device data sheet for the core configuration of a specific device.

Register 50-32: Config2: Configuration Register 2; CP0 Register 16, Select 2

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	r-1	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—

Legend:

R = Readable bit

-n = Value at POR

r = Reserved bit

W = Writable bit

'1' = Bit is set

U = Unimplemented bit, read as '0'

'0' = Bit is cleared

x = Bit is unknown

bit 31 **Reserved:** This bit is hardwired to a '1' to indicate the presence of the Config3 register.

bit 30-0 **Unimplemented:** Read as '0'

50.13.33 Config3 Register (CP0 Register 16, Select 3)

The Config3 register encodes additional capabilities. All fields in the Config3 register are read-only.

Note: The Core Configuration register bit fields and default values may vary between devices. Refer to the “Core Configuration” chapter in the specific device data sheet for the core configuration of a specific device.

Register 50-33: Config3: Configuration Register 3; CP0 Register 16, Select 3

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	r-1 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —
23:16	U-0 —	R-0 IPLW<1:0>	R-x	R-0	R-0	R-0	R-x	R/W-y ISAONEXC
15:8	R-y ISA<1:0>	R-y	R-x ULRI	R-x RXI ⁽¹⁾	R-x DSP2P ⁽²⁾	R-x DSPP ⁽²⁾	U-0 —	R-x ITL
7:0	U-0 —	R-x VEIC	R-x VINT	R-x SP ⁽¹⁾	R-x CDDM	U-0 —	U-0 —	R-x TL

Legend:

U = Unimplemented bit, read as ‘0’ y = Value set from BOOTISA Configuration bit (CFG0<6>) on a POR
R = Readable bit W = Writable bit U = Unimplemented bit, read as ‘0’
-n = Value at POR ‘1’ = Bit is set ‘0’ = Bit is cleared x = Bit is unknown

bit 31 **Reserved:** This bit is hardwired to a ‘1’ to indicate the presence of a Config4 register

bit 30-23 **Unimplemented:** Read as ‘0’

bit 22-21 **IPLW<1:0>:** Width of the Status IPL and Cause RIPL bits

11 = Reserved

10 = Reserved

01 = IPL and RIPL bits are 8-bits in width

00 = IPL and RIPL bits are 6-bits in width

If the IPL field is 8-bits in width, bits 18 and 16 of the Status register are used as the Most Significant bit and second Most Significant bit, respectively, of that bit.

If the RIPL field is 8-bits in width, bits 17 and 16 of the Cause register are used as the Most Significant bit and second Most Significant bit, respectively, of that bit.

Note: The PIC32 device core uses 8-bit IPL and RIPL fields, so these bits are set to ‘01’.

bit 20-18 **MMAR<2:0>:** microMIPS Architecture Revision level bits

111 = Reserved

110 = Reserved

101 = Reserved

100 = Reserved

011 = Reserved

010 = Reserved

001 = Reserved

000 = Release 1

bit 17 **MCU:** MIPS MCU ASE Implemented bit

1 = MCU ASE is implemented

0 = MCU ASE is not implemented

bit 16 **ISAONEXC:** ISA on Exception bit

Reflects the ISA used when vectoring to an exception. Affects exceptions whose vectors are offsets from EBASE. The reset value of this bit is determined by the BOOTISA Configuration bit (CFG0<6>).

1 = microMIPS is used on entrance to an exception vector

0 = MIPS32 ISA is used on entrance to an exception vector

Note 1: This bit is only available on devices with the MPU microprocessor core. Refer to the “CPU” chapter in the specific device data sheet for availability.

2: This bit is not available on all devices. Please consult the “CPU” chapter of the specific device data sheet to determine availability

PIC32 Family Reference Manual

Register 50-33: Config3: Configuration Register 3; CP0 Register 16, Select 3 (Continued)

- bit 15-14 **ISA<1:0>**: Indicates Instruction Set Availability
The reset value of this bit is determined by the BOOTISA Configuration bit (CFG0<6>).
11 = Both MIPS32 and microMIPS are implemented; microMIPS is used when coming out of reset
10 = Both MIPS32 and microMIPS are implemented; MIPS32 ISA used when coming out of reset
01 = Only microMIPS is implemented
00 = Only MIPS32 is implemented
- bit 13 **ULRI**: UserLocal register implemented bit
This bit indicates whether the UserLocal coprocessor 0 register is implemented.
1 = UserLocal register is implemented
0 = UserLocal register is not implemented
- bit 12 **RXI**: RIE and XIE Implemented in PageGrain bit⁽¹⁾
1 = RIE and XIE bits are implemented
0 = RIE and XIE bits are not implemented
- bit 11 **DSP2P**: MIPS DSP ASE Revision 2 Presence bit⁽²⁾
1 = DSP Revision 2 is present
0 = DSP Revision 2 is not present
- bit 10 **DSPP**: MIPS DSP ASE Presence bit⁽²⁾
1 = DSP is present
0 = DSP is not present
- bit 9 **Unimplemented**: Read as '0'
- bit 8 **ITL**: iFlowtrace[®] Hardware bit
This bit indicates that iFlowtrace[®] hardware is present
1 = The iFlowtrace[®] is implemented in the core
0 = The iFlowtrace[®] is not implemented in the core
- bit 7 **Unimplemented**: Read as '0'
- bit 6 **VEIC**: External Vector Interrupt Controller bit
This bit indicates whether support for an external interrupt controller is implemented.
1 = Support for EIC Interrupt mode is implemented
0 = Support for EIC Interrupt mode is not implemented
PIC32 devices internally implement a MIPS "external interrupt controller"; therefore, this bit reads '1'.
- bit 5 **VINT**: Vector Interrupt bit
This bit indicates whether vectored interrupts are implemented.
1 = Vector interrupts are implemented
0 = Vector interrupts are not implemented
On the PIC32 core, this bit is always a '1' since vectored interrupts are implemented.
- bit 4 **SP**: Small Page bit⁽¹⁾
This bit indicates whether support for small pages (1 KB) is implemented.
1 = Small page support is implemented
0 = Small page support is not implemented
- bit 3 **CDMM**: Common Device Memory Map bit
This bit indicates whether support for the Common Device Memory Map is implemented.
1 = CDMM is implemented
0 = CDMM is not implemented
- bit 2-1 **Unimplemented**: Read as '0'
- bit 0 **TL**: Trace Logic bit
This bit indicates whether trace logic is implemented.
1 = Trace logic is implemented
0 = Trace logic is not implemented

- Note 1:** This bit is only available on devices with the MPU microprocessor core. Refer to the "CPU" chapter in the specific device data sheet for availability.
- 2:** This bit is not available on all devices. Please consult the "CPU" chapter of the specific device data sheet to determine availability

50.13.34 Config4 Register (CP0 Register 16, Select 4)

The Config4 register is an adjunct to the Config register and encodes additional information about capabilities present on the core. On PIC32 devices, the Config4 register indicates the presence of the Config5 register. All fields in the Config4 register are read-only.

Note: The Core Configuration register bit fields and default values may vary between devices. Refer to the “**Core Configuration**” chapter in the specific device data sheet for the core configuration of a specific device.

Register 50-34: Config4: Configuration Register 4; CP0 Register 16, Select 4

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	r-1	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—

Legend:

R = Readable bit	r = Reserved	W = Writable bit	U = Unimplemented bit, read as '0'
-n = Value at POR	'1' = Bit is set	'0' = Bit is cleared	x = Bit is unknown

bit 31 **Reserved:** This bit is hardwired to a '1' to indicate the presence of the Config5 register.

bit 30-0 **Unimplemented:** Read as '0'

50.13.35 Config5 Register (CP0 Register 16, Select 5)

The Config5 register is an adjunct to the Config register and encodes additional information about capabilities present on the core. The Config5 register indicates whether or not the Nested Fault feature is implemented. All fields in the Config5 register are read-only.

Note: The Core Configuration register bit fields and default values may vary between devices. Refer to the “**Core Configuration**” chapter in the specific device data sheet for the core configuration of a specific device.

Register 50-35: Config5: Configuration Register 5; CP0 Register 16, Select 5

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	R/W-0	U-0	R-1
	—	—	—	—	—	UFR ⁽¹⁾	—	NF

Legend:

R = Readable bit
 -n = Value at POR

r = Reserved
 W = Writable bit
 '1' = Bit is set

U = Unimplemented bit, read as '0'
 '0' = Bit is cleared
 x = Bit is unknown

bit 31-3 **Unimplemented:** Read as '0'

bit 2 **UFR:** User Mode Access bit⁽¹⁾

This feature allows user mode access to Status_{FR} with CTC1 and CFC1 instructions allowed.

1 = User mode Status_{FR} instructions are allowed

0 = User mode Status_{FR} instructions are not allowed

bit 1 **Unimplemented:** Read as '0'

bit 0 **NF:** Nested Fault bit

1 = Nested Fault feature is implemented

0 = Nested Fault feature is not implemented

Note 1: This bit is only available on devices with the M-Class core. Refer to the “**CPU**” chapter in the specific device data sheet to determine availability.

50.13.36 Config7 Register (CP0 Register 16, Select 7)

The Config7 register contains implementation specific configuration information. A number of these bits are writable to disable certain performance enhancing features within the microprocessor core.

Note: The Core Configuration register bit fields and default values may vary between devices. Refer to the “**Core Configuration**” chapter in the specific device data sheet for the core configuration of a specific device.

Register 50-36: Config7: Configuration Register 7; CP0 Register 16, Select 7

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-1	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	WII	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as ‘0’

-n = Value at POR

‘1’ = Bit is set

‘0’ = Bit is cleared

x = Bit is unknown

bit 31 **WII:** Wait IE Ignore bit

Indicates that this processor will allow an interrupt to unblock a WAIT instruction, even if IE is preventing the interrupt from being taken. This avoids problems using the WAIT instruction for ‘bottom half’ interrupt servicing. This bit always reads ‘1’ for devices with the microprocessor core.

bit 30-0 **Unimplemented:** Read as ‘0’

50.13.37 LLAddr Register (CP0 Register 17, Select 0) (MPU only)

The LLAddr register contains the physical address read by the most recent Load Linked (LL) instruction. This register is for diagnostic purposes only, and serves no function during normal operation.

Register 50-37: LLAddr: Load Linked Address Register; CP0 Register 17, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	R-x	R-x	R-x	R/x
	—	—	—	—	PAddr<27:24>			
23:16	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	PAddr< 23:16>							
15:8	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	PAddr<15:8>							
7:0	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	PAddr<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-28 **Unimplemented:** Read as '0'

bit 27-0 **PAddr<27:0>:** Load Linked Instruction Physical Address Encode bits

These bits encode the physical address read by the most recent Load Linked instruction.

50.13.38 WatchLo Register (CP0 Register 18, Select 0-3) (MPU only)

The WatchLo and WatchHi registers together provide the interface to a watchpoint debug facility that initiates a watch exception if an instruction or data access matches the address specified in the registers. As such, they duplicate some functions of the EJTAG debug solution. Watch exceptions are taken only if the EXL and ERL bits are both '0' in the Status register. If either bit is a '1', the WP bit is set in the Cause register, and the watch exception is deferred until both the EXL and ERL bits are '0'.

The WatchLo register specifies the base virtual address and the type of reference (instruction fetch, load, store) to match.

Register 50-38: WatchLo: Watchdog Debug Low Register; CP0 Register 18, Select 0-3

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	VAddr<28:21>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	VAddr< 20:13>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	VAddr<12:5>							
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-0	R/W-0	R/W-0
	VAddr<4:0>					I	R	W

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-3 **VAddr<28:0>**: Virtual Address Match bits

This field specifies the virtual address to match. This is a double word address, because bits <2:0> are used to control the type of match.

bit 2 **I**: Instruction Fetch Watch Address Exception Enable bit

1 = Watch exceptions are enabled for instruction fetches that match the address

0 = Watch exceptions are not enabled

bit 1 **R**: Instruction Fetch Watch Load Exception Enable bit

1 = Watch exceptions are enabled for loads that match the address

0 = Watch exceptions are not enabled

bit 0 **W**: Instruction Fetch Watch Stores Exception Enable bit

1 = Watch exceptions are enabled for stores that match the address

0 = Watch exceptions are not enabled

50.13.39 WatchHi Register (CP0 Register 19, Select 0-3) (MPU only)

The WatchLo and WatchHi registers together provide the interface to a watchpoint debug facility that initiates a watch exception if an instruction or data access matches the address specified in the registers. As such, they duplicate some functions of the EJTAG debug solution. Watch exceptions are taken only if the EXL and ERL bits are zero in the Status register. If either bit is a '1', the WP bit is set in the Cause register, and the watch exception is deferred until both the EXL and ERL bits are '0'.

The WatchHi register contains information that qualifies the virtual address specified in the WatchLo register: an ASID, a Global (G) bit, and an optional address mask. If the G bit is '1', any virtual address reference that matches the specified address will cause a watch exception. If the G bit is a '0', only those virtual address references for which the ASID value in the WatchHi register matches the ASID value in the EntryHi register cause a watch exception. The optional mask field provides address masking to qualify the address specified in WatchLo.

Register 50-39: WatchHi: Watchdog Debug High Register; CP0 Register 19, Select 0-3

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-x	R/W-x	U-0	U-0	U-0	U-0	U-0	U-0
	M	G	—	—	—	—	—	—
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	ASID<7:0>							
15:8	U-0	U-0	U-0	U-0	R/W-x	R/W-x	R/W-x	R/W-x
	—	—	—	—	Mask<8:5>			
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	W-x, HS, CS	W-x, HS, CS	W-x, HS, CS
	Mask<4:0>					I	R	W

Legend:	HS = Set in Hardware	CS = Cleared by Software
R = Readable bit	W = Writable bit	U = Unimplemented bit, read as '0'
-n = Value at POR	'1' = Bit is set	'0' = Bit is cleared
		x = Bit is unknown

- bit 31 **M:** Watch Register Pairs Detect bit
Indicates whether additional Watch register pairs beyond this one are present or not
- bit 30 **G:** Global bit
1 = Any address that matches that specified in the WatchLo register causes a watch exception
0 = ASID<7:0> must match the ASID bits of the EntryHi register to cause a watch exception
- bit 29-24 **Unimplemented:** Read as '0'
- bit 23-16 **ASID<7:0>:** Address Space ID Watch Exception Match bits
ASID value which is required to match that in the EntryHi register if the G bit is zero in the WatchHi register.
- bit 15-12 **Unimplemented:** Read as '0'
- bit 11-3 **Mask<8:0>:** Virtual Address Match Mask bits
Bit mask that qualifies the address in the WatchLo register. Any bit in this field that is a set inhibits the corresponding address bit from participating in the address match.
- bit 2 **I:** Instruction Fetch Condition Match bit
This bit is set by hardware when an instruction fetch condition matches the values in this watch register pair. When set, the bit remains set until cleared by software, which is accomplished by writing a '1'.
- bit 1 **R:** Load Condition Match bit
This bit is set by hardware when a load condition matches the values in this watch register pair. When set, the bit remains set until cleared by software, which is accomplished by writing a '1'.
- bit 0 **W:** Store Condition Match bit
This bit is set by hardware when a store condition matches the values in this watch register pair. When set, the bit remains set until cleared by software, which is accomplished by writing a '1'

50.13.40 Debug Register (CP0 Register 23, Select 0)

The Debug register is used to control the debug exception and provide information about the cause of the debug exception and when re-entering at the debug exception vector due to a normal exception in Debug mode. The read-only information bits are updated every time the debug exception is taken or when a normal exception is taken when already in Debug mode.

Only the DM bit and the VER<2:0> bits are valid when read from non-Debug mode; the values of all other bits and fields are unpredictable. Operation of the processor is undefined if the Debug register is written from non-Debug mode.

Some of the bits and fields are only updated on debug exceptions and/or exceptions in Debug mode, as follows:

- DSS, DBP, DDBL, DDBS, DIB, DINT are updated on both debug exceptions and on exceptions in debug modes
- DEXCCODE<4:0> are updated on exceptions in Debug mode, and are undefined after a debug exception
- HALT and DOZE are updated on a debug exception, and are undefined after an exception in Debug mode
- DBD is updated on both debug and on exceptions in debug modes

All bits are undefined when read from Normal mode, except VER<2:0> and DM.

Register 50-40: Debug: Debug Exception Register; CP0 Register 23, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-0	R-0	R-0	R/W-0	R-x	R-x	R/W-1	R/W-0
	DBD	DM	NODCR	LSNM	DOZE	HALT	COUNTDM	IBUSEP
23:16	U-0	U-0	R/W-0	R/W-0	R-0	R-0	R-1	R-0
	—	—	DBUSEP	IEXI	DDBSIMPR	DDBLIMPR	VER<2:1>	
15:8	R-1	R-x	R-x	R-x	R-x	R-x	R-0	R/W-0
	VER<0>	DEXCCODE<4:0>					NOSST	SST
7:0	U-0	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	—	DIBIMPR	DINT	DIB	DDBS	DDBL	DBP	DSS

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **DBD:** Branch Delay Debug Exception bit

Indicates whether the last debug exception or exception in Debug mode, occurred in a branch delay slot:

1 = In delay slot

0 = Not in delay slot

bit 30 **DM:** Debug Mode bit

Indicates that the processor is operating in Debug mode:

1 = Processor is operating in Debug mode

0 = Processor is operating in non-Debug mode

bit 29 **NODCR:** Debug Control Register bit

Indicates whether the dseg memory segment is present and the Debug Control Register is accessible:

1 = No dseg present

0 = dseg is present

bit 28 **LSNM:** Load Store Access Control bit

Controls access of load/store between dseg and main memory:

1 = Load/stores in dseg address range goes to main memory

0 = Load/stores in dseg address range goes to dseg

PIC32 Family Reference Manual

Register 50-40: Debug: Debug Exception Register; CP0 Register 23, Select 0 (Continued)

- bit 27 **DOZE**: Low-Power Mode Debug Exception bit
Indicates that the processor was in any kind of low-power mode when a debug exception occurred.
1 = Processor was in a low-power mode when a debug exception occurred
0 = Processor was not in a low-power mode when a debug exception occurred
- bit 26 **HALT**: System Bus Clock Stop bit
Indicates that the internal system bus clock was stopped when the debug exception occurred.
1 = Internal system bus clock running
0 = Internal system bus clock stopped
- bit 25 **COUNTDM**: Count Register Behavior bit
Indicates the Count register behavior in Debug mode.
1 = Count register is running in Debug mode
0 = Count register stopped in Debug mode
- bit 24 **IBUSEP**: Instruction Fetch Bus Error Exception Pending bit
Set when an instruction fetch bus error event occurs or if a '1' is written to the bit by software. Cleared when a Bus Error exception on instruction fetch is taken by the processor, and by Reset. If IBUSEP is set when IEXI is cleared, a Bus Error exception on instruction fetch is taken by the processor, and IBUSEP is cleared.
- bit 23-22 **Unimplemented**: Read as '0'
- bit 21 **DBUSEP**: Data Access Bus Error Exception Pending bit
Covers imprecise bus errors on data access, similar to behavior of IBUSEP for imprecise bus errors on an instruction fetch.
- bit 20 **IEXI**: Imprecise Error Exception Inhibit Control bit
Controls exceptions taken due to imprecise error indications. Set when the processor takes a debug exception or exception in Debug mode. Cleared by execution of the `DERET` instruction; otherwise modifiable by Debug mode software. When IEXI is set, the imprecise error exception from a bus error on an instruction fetch or data access, cache error, or machine check is inhibited and deferred until the bit is cleared.
- bit 19 **DDBSIMPR**: Debug Data Break Store Exception bit
Indicates that an imprecise Debug Data Break Store exception was taken. All data breaks are precise on the PIC32 core, so this bit will always read as '0'.
- bit 18 **DDBLIMPR**: Debug Data Break Load Exception bit
Indicates that an imprecise Debug Data Break Load exception was taken. All data breaks are precise on the PIC32 core, so this bit will always read as '0'.
- bit 17-15 **VER<2:0>**: EJTAG Version bit
Contains the EJTAG version number.
- bit 14-10 **DEXCCODE<4:0>**: Latest Exception in Debug Mode bit
Indicates the cause of the latest exception in Debug mode. The bit is encoded as the EXCCODE<4:0> bits in the Cause register for those normal exceptions that may occur in Debug mode. Value is undefined after a debug exception.
- bit 9 **NOSST**: Single-step Feature Control bit
Indicates whether the single-step feature controllable by the SST bit is available in this implementation.
1 = No single-step feature is available
0 = Single-step feature is available
- bit 8 **SST**: Debug Single-step Control bit
Controls if debug single-step exception is enabled.
1 = Debug single step exception is enabled
0 = No debug single-step exception is enabled
- bit 7 **Unimplemented**: Read as '0'
- bit 6 **DIBIMPR**: Imprecise Debug Instruction Break Exception bit
Indicates that an imprecise debug instruction break exception occurred (due to a complex breakpoint). Cleared on exception in Debug mode.

Register 50-40: Debug: Debug Exception Register; CP0 Register 23, Select 0 (Continued)

- bit 5 **DINT:** Debug Interrupt Exception bit
Indicates that a debug interrupt exception occurred. Cleared on exception in Debug mode.
1 = Debug interrupt exception
0 = No debug interrupt exception
- bit 4 **DIB:** Debug Instruction Break Exception bit
Indicates that a debug instruction break exception occurred. Cleared on exception in Debug mode.
1 = Debug instruction exception
0 = No debug instruction exception
- bit 3 **DBBS:** Debug Data Break Exception on Store bit
Indicates that a debug data break exception occurred on a store. Cleared on exception in Debug mode.
1 = Debug instruction exception on a store
0 = No debug data exception on a store
- bit 2 **DBDL:** Debug Data Break Exception on Load bit
Indicates that a debug data break exception occurred on a load. Cleared on exception in Debug mode.
1 = Debug instruction exception on a load
0 = No debug data exception on a load
- bit 1 **DBP:** Debug Software Breakpoint Exception bit
Indicates that a debug software breakpoint exception occurred. Cleared on exception in Debug mode.
1 = Debug software breakpoint exception
0 = No debug software breakpoint exception
- bit 0 **DSS:** Debug Single-step Exception bit
Indicates that a debug single-step exception occurred. Cleared on exception in Debug mode.
1 = Debug single-step exception
0 = No debug single-step exception

50.13.41 TraceControl Register (CP0 Register 23, Select 1)

The TraceControl register enables software trace control and is only implemented on devices with the EJTAG trace capability.

Register 50-41: TraceControl: Trace Control Register; CP0 Register 23, Select 1

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-0 TS	R/W-x UT	U-0 —	U-0 —	R/W-x TB	R/W-x IO	R/W-x D ⁽¹⁾	R/W-x E ⁽¹⁾
23:16	R/W-x K ⁽¹⁾	U-0 —	R/W-x U ⁽¹⁾	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	ASID_M<2:0>			ASID<7:3>				
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-0
	ASID<2:0>			G	MODE<2:0>			ON

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **TS:** Trace Select bit

This bit is used to select between the hardware and the software trace control bits.

1 = Selects the trace control bits

0 = Selects the external hardware trace block signals

bit 30 **UT:** User Type Select bit

This bit is used to indicate the type of user-triggered trace record.

1 = User type 2

0 = User type 1

bit 29-28 **Unimplemented:** Read as '0'

bit 27 **TB:** Trace Branch bit

1 = Trace the PC value for all taken branches

0 = Trace the PC value for branch targets with unpredictable static addresses

bit 26 **IO:** Inhibit Overflow bit

This signal is used to indicate to the core trace logic that slow but complete tracing is desired.

1 = Inhibit FIFO overflow or discard of trace data

0 = Allow FIFO overflow or discard of trace data

bit 25 **D:** Debug Mode Trace Enable bit⁽¹⁾

1 = Enable tracing in Debug mode

0 = Disable tracing in Debug mode

bit 24 **E:** Exception Mode Trace Enable bit⁽¹⁾

1 = Enable tracing in Exception mode

0 = Disable tracing in Exception mode

bit 23 **K:** Kernel Mode Trace Enable bit⁽¹⁾

1 = Enable tracing in Kernel mode

0 = Disable tracing in Kernel mode

bit 22 **Unimplemented:** Read as '0'

Note 1: The ON bit must be set to '1' to enable tracing.

Register 50-41: TraceControl: Trace Control Register; CP0 Register 23, Select 1 (Continued)

- bit 21 **U:** User Mode Trace Enable bit⁽¹⁾
1 = Enable tracing in User mode
0 = Disable tracing in User mode
- bit 20-13 **ASID_M<7:0>:** Address Space ID Mask Value bits
A '1' in any bit in this field inhibits the corresponding ASID bit from participating in the match.
A value of '0' in this field compares all bits of ASID.
- bit 12-5 **ASID<7:0>:** Address Space ID Value bits
These bits must match the Address Space ID when the G bit is set to '0'.
These bits are ignored when the G bit is set to '1'.
- bit 4 **G:** Global bit
1 = Tracing is to be enabled for all processes, provided that other enabling functions are also true
0 = Tracing is not enabled
- bit 3-1 **MODE<2:0>:** Trace Mode Control bits
111 = Trace PC and both load and store address and data
110 = Trace PC and store address and data
101 = Trace PC and load address and data
100 = Trace PC and load data
011 = Trace PC and both load and store addresses
010 = Trace PC and store address
001 = Trace PC and load address
000 = Trace PC
- bit 0 **ON:** Master Trace Enable bit
1 = Tracing is enabled when another trace enable bit is set to '1'
0 = Tracing is disabled

Note 1: The ON bit must be set to '1' to enable tracing.

50.13.42 TraceControl2 Register (CP0 Register 23, Select 2)

The TraceControl2 register provides additional control and status information. Note that some fields in the TraceControl2 register are read-only, but have a reset state of “Undefined”. This is because these values are loaded from the Trace Control Block (TCB). As such, these fields in the TraceControl2 register will not have valid values until the TCB asserts these values.

Register 50-42: TraceControl2: Trace Control Register 2; CP0 Register 23, Select 2

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	R-1	R-0	R-0	R-1	R-x	R-x	R-x
	—	VALIDMODES<1:0>		TBI	TBU	SYP<2:0>		

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as ‘0’

-n = Value at POR

‘1’ = Bit is set

‘0’ = Bit is cleared

x = Bit is unknown

bit 31-7 **Unimplemented:** Read as ‘0’

bit 6-5 **VALIDMODES<1:0>:** Valid Trace Mode Select bits

11 = Reserved

10 = PC, load and store address, and load and store data

01 = PC and load and store address tracing only

00 = PC tracing only

bit 4 **TBI:** Trace Buffers Implemented bit

1 = On-chip and off-chip trace buffers are implemented by the TCB

0 = Only one trace buffer is implemented

bit 3 **TBU:** Trace Buffers Used bit

1 = Trace data is being sent to an off-chip trace buffer

0 = Trace data is being sent to an on-chip trace buffer

bit 2-0 **SYP<2:0>:** Synchronization Period bits

The “On-chip” column value is used when the trace data is being written to an on-chip trace buffer (e.g., TBU bit = 0). Conversely, the “Off-chip” column is used when the trace data is being written to an off-chip trace buffer (e.g., TBU bit = 1).

Bit Setting	On-chip	Off-chip
111 =	2^2	2^7
110 =	2^3	2^8
101 =	2^4	2^9
100 =	2^5	2^{10}
011 =	2^6	2^{11}
010 =	2^7	2^{12}
001 =	2^8	2^{13}
000 =	2^9	2^{14}

50.13.43 UserTraceData1 Register (CP0 Register 23, Select 3)

A software write to any bits in the UserTraceData1 register will trigger a trace record to be written indicating a type 1 user format. The type is based on the UT bit in the TraceControl register. This register cannot be written in consecutive cycles. The trace output data is unpredictable if this register is written in consecutive cycles.

This register is only implemented on devices with the EJTAG trace capability.

Register 50-43: UserTraceData1: User Trace Data Register 1; CP0 Register 23, Select 3

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<31:24>							
23:16	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<23:16>							
15:8	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<15:8>							
7:0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **DATA<31:0>**: Software Readable/Writable Data bits

When written, this register triggers a user format trace record out of the PDtrace interface that transmits the Data bit to trace memory.

50.13.44 TraceBPC Register (CP0 Register 23, Select 4)

This register is used to control start and stop of tracing using an EJTAG Hardware breakpoint. The Hardware breakpoint would then be set as a trigger source and optionally also as a Debug exception breakpoint.

This register is only implemented on devices with both Hardware breakpoints and the EJTAG trace capability.

Register 50-44: TraceBPC: Trace Breakpoint Control Register; CP0 Register 23, Select 4

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-0 DE	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —
23:16	R/W-0 DBPO7 ⁽¹⁾	R/W-0 DBPO6 ⁽¹⁾	R/W-0 DBPO5 ⁽¹⁾	R/W-0 DBPO4 ⁽¹⁾	R/W-0 DBPO3 ⁽¹⁾	R/W-0 DBPO2 ⁽¹⁾	R/W-0 DBPO1	R/W-0 DBPO0
15:8	R/W-0 IE	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —
7:0	R/W-0 IBPO7 ⁽¹⁾	R/W-0 IBPO6 ⁽¹⁾	R/W-0 IBPO5	R/W-0 IBPO4	R/W-0 IBPO3	R/W-0 IBPO2	R/W-0 IBPO1	R/W-0 IBPO0

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **DE:** EJTAG Data Breakpoint Trigger Select bit

1 = Enable trigger signals from data breakpoints

0 = Disable trigger signals from data breakpoints

bit 30-24 **Unimplemented:** Read as '0'

bit 23 **DBPO7:** Data Breakpoint 7 bit⁽¹⁾

1 = Enable corresponding data instruction breakpoint trigger to start tracing

0 = Disable tracing with the trigger signal

bit 22 **DBPO6:** Data Breakpoint 6 bit⁽¹⁾

1 = Enable corresponding data instruction breakpoint trigger to start tracing

0 = Disable tracing with the trigger signal

bit 21 **DBPO5:** Data Breakpoint 5 bit⁽¹⁾

1 = Enable corresponding data instruction breakpoint trigger to start tracing

0 = Disable tracing with the trigger signal

bit 20 **DBPO4:** Data Breakpoint 4 bit⁽¹⁾

1 = Enable corresponding data instruction breakpoint trigger to start tracing

0 = Disable tracing with the trigger signal

bit 19 **DBPO3:** Data Breakpoint 3 bit⁽¹⁾

1 = Enable corresponding data instruction breakpoint trigger to start tracing

0 = Disable tracing with the trigger signal

bit 18 **DBPO2:** Data Breakpoint 2 bit⁽¹⁾

1 = Enable corresponding data instruction breakpoint trigger to start tracing

0 = Disable tracing with the trigger signal

bit 17 **DBPO1:** Data Breakpoint 1 bit

1 = Enable corresponding data instruction breakpoint trigger to start tracing

0 = Disable tracing with the trigger signal

Note 1: This bit is only available on PIC32 devices with the MPU core.

Register 50-44: TraceBPC: Trace Breakpoint Control Register; CP0 Register 23, Select 4 (Continued)

- bit 16 **DBPO0:** Data Breakpoint 0 bit
1 = Enable corresponding data instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 15 **IE:** EJTAG Instruction Breakpoint Select bit
1 = Enable trigger signals from instruction breakpoints
0 = Disable trigger signals from instruction breakpoints
- bit 14-8 **Unimplemented:** Read as '0'
- bit 5 **IBPO7:** Instruction Breakpoint 7 bit⁽¹⁾
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 4 **IBPO6:** Instruction Breakpoint 6 bit⁽¹⁾
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 5 **IBPO5:** Instruction Breakpoint 5 bit
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 4 **IBPO4:** Instruction Breakpoint 4 bit
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 3 **IBPO3:** Instruction Breakpoint 3 bit
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 2 **IBPO2:** Instruction Breakpoint 2 bit
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 1 **IBPO1:** Instruction Breakpoint 1 bit
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal
- bit 0 **IBPO0:** Instruction Breakpoint 0 bit
1 = Enable corresponding instruction breakpoint trigger to start tracing
0 = Disable tracing with the trigger signal

Note 1: This bit is only available on PIC32 devices with the MPU core.

50.13.45 Debug2 Register (CP0 Register 23, Select 6)

This register holds additional information about Complex Breakpoint exceptions. This register is only implemented if complex hardware breakpoints are present.

Register 50-45: Debug2: Debug Breakpoint Exceptions Register; CP0 Register 23, Select 6

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	r-1	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	R-x	R-x	R-x	R-x
	—	—	—	—	PRM	DQ	TUP	PACO

Legend:	r = Reserved bit
R = Readable bit	W = Writable bit
-n = Value at POR	'1' = Bit is set
	U = Unimplemented bit, read as '0'
	'0' = Bit is cleared
	x = Bit is unknown

- bit 31 **Reserved:** Read as '1'
- bit 30-4 **Unimplemented:** Read as '0'
- bit 3 **PRM:** Primed bit
Indicates whether a complex breakpoint with an active priming condition was seen on the last debug exception.
- bit 2 **DQ:** Data Qualified bit
Indicates whether a complex breakpoint with an active data qualifier was seen on the last debug exception.
- bit 1 **TUP:** Tuple Breakpoint bit
Indicates whether a tuple breakpoint was seen on the last debug exception.
- bit 0 **PACO:** Pass Counter bit
Indicates whether a complex breakpoint with an active pass counter was seen on the last debug exception.

50.13.46 DEPC Register (CP0 Register 24, Select 0)

The Debug Exception Program Counter (DEPC) register is a read/write register that contains the address at which processing resumes after a debug exception or Debug mode exception has been serviced.

For synchronous (precise) debug and Debug mode exceptions, the DEPC register contains either:

- The virtual address of the instruction that was the direct cause of the debug exception, or
- The virtual address of the immediately preceding branch or jump instruction, when the debug exception causing instruction is in a branch delay slot, and the Debug Branch Delay (DBD) bit in the Debug register is set.

For asynchronous debug exceptions (debug interrupt), the DEPC register contains the virtual address of the instruction where execution should resume after the debug handler code is executed.

Since the PIC32 family implements the MIPS16e or microMIPS ASE, a read of the DEPC register (via MFC0) returns the following value in the destination GPR:

$$\text{GPR}[\text{rt}] = \text{DebugExceptionPC}_{31..1} \parallel \text{ISAMode}_0$$

That is, the upper 31 bits of the debug exception PC are combined with the lower bit of the ISA<1:0> bits (Config3<15:14>) and are written to the GPR.

Similarly, a write to the DEPC register (via MTC0) takes the value from the GPR and distributes that value to the debug exception PC and the ISA<1:0> bits (Config3<15:14>), as follows:

$$\begin{aligned} \text{DebugExceptionPC} &= \text{GPR}[\text{rt}]_{31..1} \parallel 0 \\ \text{ISAMode} &= 2\#0 \parallel \text{GPR}[\text{rt}]_0 \end{aligned}$$

That is, the upper 31 bits of the GPR are written to the upper 31 bits of the debug exception PC, and the lower bit of the debug exception PC is cleared. The upper bit of the ISA<1:0> bits (Config3<15:14>) is cleared and the lower bit is loaded from the lower bit of the GPR.

Register 50-46: DEPC: Debug Exception Program Counter Register; CP0 Register 24, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DEPC<31:24>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DEPC<23:16>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DEPC<15:8>							
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DEPC<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **DEPC<31:0>**: Debug Exception Program Counter bits

The DEPC register is updated with the virtual address of the instruction that caused the debug exception. If the instruction is in the branch delay slot, the virtual address of the immediately preceding branch or jump instruction is placed in this register.

Execution of the `DERET` instruction causes a jump to the address in the DEPC register.

50.13.47 UserTraceData2 (CP0 Register 24, Select 3)

A software write to any bits in the UserTraceData2 register will trigger a trace record to be written indicating a type 2 user format. The type is based on the UT bit in the TraceControl register. This register cannot be written in consecutive cycles. The trace output data is unpredictable if this register is written in consecutive cycles.

This register is only implemented on devices with the EJTAG trace capability.

Register 50-47: UserTraceData2: User Trace Data Register 2; CP0 Register 24, Select 3

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<31:24>							
23:16	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<23:16>							
15:8	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<15:8>							
7:0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	DATA<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **DATA<31:0>**: Software Readable/Writable Data bits

When written, this register triggers a user format trace record out of the PDtrace interface that transmits the DATA bits to trace memory.

50.13.48 PerfCtlx Register (CP0 Register 25, Select 0/3)

The microprocessor core defines two performance counters, PerfCnt0 and PerfCnt1 (see [Register 50-49](#)), and two associated control registers, PerfCtl0 and PerfCtl1, which are mapped to CP0 register 25. The select bit of the MTC0/MFC0 instructions are used to select the specific register accessed by the instruction, as shown in [Table 50-29](#).

Table 50-29: Performance Counter Register Selects

Select<2:0>	Register
0	Register 0 Control
1	Register 0 Count
2	Register 1 Control
3	Register 1 Count

Each counter is a 32-bit read/write register and is incremented by one each time the countable event, specified in its associated control register, occurs. Each counter can independently count one type of event at a time.

Bit 31 of each of the counters are ANDed with an interrupt enable bit, IE, of their respective control register to determine if a performance counter interrupt should be signaled. The two values are then ORed together to create the Performance Counter Interrupt output. This signals an interrupt to the core. Counting is not affected by the interrupt indication. This output is cleared when the counter wraps to zero, and may be cleared in software by writing a value with bit 31 = 0 to the Performance Counter Count registers.

Note: The performance counter registers are connected to a clock that is stopped when the processor is in Sleep mode. Most events would not be active during that time, but others would be, notably the cycle count. This behavior should be considered when analyzing measurements taken on a system.

Register 50-48: PerfCtlx: Performance Counter Control Register; CP0 Register 25, Select 0/3)

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	r-x	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	R/W-x	R/W-x	R/W-x	R/W-x
	—	—	—	—	EVENT<6:3>			
7:0	R/W-x	R/W-x	R/W-x	R/W-0	R-0	U-0	R/W-x	R/W-x
	EVENT<2:0>			IE	U	—	K	EXL

Legend:

r = Reserved bit
R = Readable bit
-n = Value at POR
W = Writable bit
'1' = Bit is set
U = Unimplemented bit, read as '0'
'0' = Bit is cleared
x = Bit is unknown

bit 31 **Reserved:** Read as '1' for PerfCtl0, and '0' for PerfCtl1

bit 30-12 **Unimplemented:** Read as '0'

bit 11-5 **EVENT<6:0>:**

Counter event enabled for this counter. Possible events are listed in [Table 50-30](#).

bit 4 **IE:** Counter Interrupt Enable bit

This bit masks bit 31 of the associated count register from the interrupt exception request output.

bit 3 **U:** Count in User Mode bit

When this bit is set, the specified event is counted in User mode.

PIC32 Family Reference Manual

Register 50-48: PerfCtlx: Performance Counter Control Register; CP0 Register 25, Select 0/3) (Continued)

- bit 2 **Unimplemented:** Read as '0'
- bit 1 **K:** Count in Kernel Mode bit
When this bit is set, count the event in Kernel mode when EXL and ERL are both '0'.
- bit 0 **EXL:** Count when EXL bit
When this bit is set, count the event when EXL = 1 and ERL = 0.

Table 50-30 provides the events countable with two performance counters. The mode column indicates whether the event counting is influenced by the mode bits (U, K, EXL). The operation of a counter is unpredictable for events that are specified as Reserved.

Performance counters never count in Debug mode or when ERL = 1.

Table 50-30: Performance Countable Events

Event Number	Counter 0	Mode	Counter 1	Mode
0	Cycles	No	Cycles	No
1	Instructions completed	Yes	Instructions completed	Yes
2	Branch instructions	Yes	Reserved	N/A
3	JR r31 (return) instructions	Yes	Reserved	N/A
4	JR (not r31) instructions	Yes	Reserved	N/A
5	ITLB accesses ⁽¹⁾	Yes	ITLB misses ⁽¹⁾	Yes
6	DTLB accesses ⁽¹⁾	Yes	DTLB misses ⁽¹⁾	Yes
7	JTLB instruction accesses ⁽¹⁾	Yes	JTLB instruction misses ⁽¹⁾	Yes
8	JTLB data accesses ⁽¹⁾	Yes	JTLB data misses ⁽¹⁾	Yes
9	Instruction cache accesses ⁽¹⁾	Yes	Instruction cache misses ⁽¹⁾	Yes
10	Data cache accesses ⁽¹⁾	Yes	Data cache write-backs ⁽¹⁾	Yes
11	Data cache misses ⁽¹⁾	Yes	Data cache misses ⁽¹⁾	Yes
12	Reserved	N/A	Reserved	N/A
13	Reserved	N/A	Reserved	N/A
14	integer instructions completed	Yes	Reserved	N/A
15	loads completed	Yes	Stores completed	Yes
16	J/JAL completed	Yes	microMIPS instructions completed	Yes
17	no-ops completed	Yes	Integer multiply/divide completed	Yes
18	Stall cycles	No	Reserved	N/A
19	SC instructions completed	Yes	SC instructions failed	Yes
20	Prefetch instructions completed	Yes	Prefetch instructions completed with cache hit ⁽¹⁾	Yes
21	Reserved	N/A	Reserved	N/A
22	Reserved	N/A	Reserved	N/A
23	Exceptions taken	Yes	Reserved	N/A
24	Reserved	N/A	Reserved	N/A
25	IFU stall cycles ⁽¹⁾	No	ALU stall cycles	No
26	Reserved	N/A	Reserved	N/A
27	Reserved	N/A	Reserved	N/A
28	Reserved	N/A	Implementation-specific CP2 event	Yes
29	Reserved	N/A	Reserved	N/A
30	Reserved	N/A	Reserved	N/A

Note 1: This event is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

Table 50-30: Performance Countable Events (Continued)

Event Number	Counter 0	Mode	Counter 1	Mode
31	Reserved	N/A	Reserved	N/A
32	Reserved	N/A	Reserved	N/A
33	Uncached loads	N/A	Uncached stores	N/A
34	Reserved	N/A	Reserved	N/A
35	CP2 arithmetic instructions completed	Yes	CP2 To/From instructions completed	Yes
36	Reserved	N/A	Reserved	N/A
37	I-Cache miss stall cycles ⁽¹⁾	Yes	D-Cache miss stall cycles ⁽¹⁾	Yes
38	Reserved	N/A	Reserved	N/A
39	D-Cache miss cycles ⁽¹⁾	No	Reserved	N/A
40	Uncached stall cycles	Yes	Reserved	N/A
41	MDU stall cycles	Yes	Reserved	N/A
42	CP2 stall cycles	Yes	Reserved	N/A
43	Reserved	N/A	Reserved	N/A
44	CACHE instruction stall cycles ⁽¹⁾	No	Reserved	N/A
45	Load to Use stall cycles	Yes	Reserved	N/A
46	Other interlock stall cycles	Yes	Reserved	N/A
47	Reserved	N/A	Reserved	N/A
48	Reserved	N/A	Reserved	N/A
49	EJTAG Instruction Triggerpoints	Yes	EJTAG Data Triggerpoints	Yes
50 -51	Reserved	N/A	Reserved	N/A
52	LFQ < one-fourth full ⁽¹⁾	No	LFQ one-fourth to one-half full ⁽¹⁾	No
53	LFQ > one-half full ⁽¹⁾	No	LFQ full pipeline stall cycles ⁽¹⁾	No
54	WBB < one-fourth full ⁽¹⁾	No	WBB one-fourth to one-half full ⁽¹⁾	No
55	WBB > one-half full ⁽¹⁾	No	WBB full pipeline stall cycles ⁽¹⁾	No
56-63	Reserved	N/A	Reserved	N/A

Note 1: This event is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

PIC32 Family Reference Manual

Table 50-31: Event Description

Event Name	Counter	Event Number	Description
Cycles	0/1	0	Total number of cycles. The performance counters are clocked by the top-level gated clock. If the core is built with that clock gate is present, none of the counters will increment while the clock is stopped (e.g., due to a WAIT instruction).
Instruction Completion			
The following events indicate completion of various types of instructions			
Instructions	0/1	1	Total number of instructions completed.
Branch instructions	0	2	Counts all branch instructions that completed.
JR R31 (return) instructions	0	3	Counts all JR R31 instructions that completed.
JR (not R31)	0	4	Counts all JR rxx (not r31) and JALR instructions (indirect jumps).
Integer instructions	0	14	Non-floating point instructions.
Loads	0	15	Includes both integer and coprocessor loads.
Stores	1	15	Includes both integer and coprocessor stores.
J/JAL	0	16	Direct Jump (And Link) instruction.
microMIPS	1	16	All microMIPS instructions.
no-ops	0	17	This includes all instructions that normally write to a GPR, but where the destination register was set to r0.
Integer Multiply/Divide	1	17	Counts all Integer Multiply/Divide instructions (MULxx, DIVx, MADDx, MSUBx).
SC	0	19	Counts conditional stores regardless of whether they succeeded.
PREF	0	20	Note that this only counts PREFs that are actually attempted. PREFs to uncached addresses or ones with translation errors are not counted
Uncached loads ⁽¹⁾	0	33	Include both uncached and uncached accelerated CCAs.
Uncached stores ⁽¹⁾	1	33	
Instruction Execution Events			
ITLB accesses ⁽¹⁾	0	5	Counts ITLB accesses that are due to fetches showing up in IF stage of the pipe and do not use fixed mapping or are not in unmapped space. If an address is fetched twice down the pipe (as in the case of a cache miss), that instruction will count 2 ITLB accesses. Also, because each fetch returns 2 instructions, there is one access marked per double word.
ITLB misses ⁽¹⁾	1	5	Counts all misses in ITLB except ones that are on the back of another miss. We cannot process back to back misses and thus those are ignored for this purpose. Also ignored if there is some form of address error.
DTLB accesses ⁽¹⁾	0	6	Counts DTLB access including those in unmapped address spaces.
DTLB misses ⁽¹⁾	1	6	Counts DTLB misses. Back to back misses that result in only one DTLB entry getting refilled are counted as a single miss.
JTLB instruction accesses ⁽¹⁾	0	7	Instruction JTLB accesses are counted exactly the same as ITLB misses.
JTLB instruction misses ⁽¹⁾	1	7	Counts instruction JTLB accesses that result in no match or a match on an invalid translation.

Note 1: This event is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

Table 50-31: Event Description (Continued)

Event Name	Counter	Event Number	Description
JTLB data accesses ⁽¹⁾	0	8	Data JTLB accesses.
JTLB data misses ⁽¹⁾	1	8	Counts data JTLB accesses that result in no match or a match on an invalid translation.
I-Cache accesses ⁽¹⁾	0	9	Counts every time the instruction cache is accessed. All replays, wasted fetches etc. are counted. For example, following a branch, even the prediction is taken, the fall-through access is counted.
I-Cache misses ⁽¹⁾	1	9	Counts all instruction cache misses that result in a bus request.
D-Cache accesses ⁽¹⁾	0	10	Counts cached loads and stores.
D-Cache writebacks ⁽¹⁾	1	10	Counts cache lines written back to memory due to replacement or cache ops.
D-Cache misses ⁽¹⁾	0/1	11	Counts loads and stores that miss in the cache.
SC instructions failed	1	19	SC instruction that did not update memory. Note: While this event and the SC instruction count event can be configured to count in specific operating modes, the timing of the events is much different, and the observed operating mode could change between them, causing some inaccuracy in the measured ratio.
PREF completed with cache hit ⁽¹⁾	1	20	Counts PREF instructions that hit in the cache.
Exceptions Taken	0	23	Any type of exception taken.
EJTAG instruction triggers	0	49	Number of times an EJTAG Instruction Trigger Point condition matched.
EJTAG data triggers	1	49	Number of times an EJTAG Data Trigger Point condition matched.
Pipeline			
Cache fixup ⁽¹⁾	0	24	Counts cycles where the DCC is in a fixup and cannot accept a new instruction from the ALU. Fixups are replays within the DCC that occur when an instruction needs to reaccess the cache or the DTLB.
General Stalls			
IFU stall cycles ⁽¹⁾	0	25	Counts the number of cycles in which the fetch unit is not providing a valid instruction to the ALU.
ALU stall cycles	1	25	Counts the number of cycles in which the ALU pipeline cannot advance.
Stall cycles	0	18	Counts the total number of cycles in which no instructions are issued by SRAM to ALU (the RF stage does not advance). This includes both of the previous two events. However, this is different from the sum of them, because cycles when both stalls are active will only be counted once.
Specific Stalls			
These events will count the number of cycles lost due to this. This will include bubbles introduced by replays within the pipe. If multiple stall sources are active simultaneously, the counters for each of the active events will be incremented.			
I-Cache miss stall cycles ⁽¹⁾	0	37	Cycles when ICC stalls because an I-Cache miss caused the ICC not to have any runnable instructions. Ignores the stalls due to ITLB misses as well as the 4 cycles following a redirect.

Note 1: This event is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

PIC32 Family Reference Manual

Table 50-31: Event Description (Continued)

Event Name	Counter	Event Number	Description
D-Cache miss stall cycles ⁽¹⁾	1	37	Counts all cycles in which the integer pipeline waits on a Load to return data due to a D-Cache miss.
D-Cache miss cycle cycles ⁽¹⁾	0	39	D-Cache miss is outstanding, but not necessarily stalling the pipeline. The difference between this and D-Cache miss stall cycles can show the gain from non-blocking cache misses.
Uncached stall cycles	0	40	Cycles in which the processor is stalled on an uncached fetch, load, or store.
MDU stall cycles	0	41	Counts all cycles in which the integer pipeline waits on MDU return data.
CACHE instruction stall cycles ⁽¹⁾	0	44	Counts all cycles in which pipeline is stalled due to CACHE instructions. Includes cycles in which CACHE instructions themselves are stalled in the ALU, and cycles in which CACHE instructions cause subsequent instructions to be stalled.
Load to Use stall cycles	0	45	Counts all cycles in which the integer pipeline waits on Load return data.
Other interlocks stall cycles	0	46	Counts all cycles in which the integer pipeline waits on return data from MFC0 and RDHWR instructions.
LFB full pipeline stall cycles	1	53	Cycles in which the pipeline is stalled because the Load Fill Buffer (LFB) in the DCC is full.
Write-through buffer full stall cycles ⁽¹⁾	1	55	Cycles in which the pipeline is stalled because the write-through buffer in the BIU is full.

Note 1: This event is only available on devices with the MPU core. Refer to the “CPU” chapter in the specific device data sheet for availability.

50.13.49 PerfCntx Register (CP0 Register 25, Select 1/3)

The microprocessor core defines two performance counters, PerfCnt0 and PerfCnt1, and two associated control registers, PerfCtl0 and PerfCtl1 (see [Register 50-48](#)), which are mapped to CP0 register 25. The select bit of the MTC0/MFC0 instructions are used to select the specific register accessed by the instruction, as shown in [Table 50-29](#).

The performance counter resets to a low-power state, in which none of the counters will start counting events until software has enabled event counting, using an MTC0 instruction to the Performance Counter Control Registers.

Register 50-49: PerfCntx: Performance Counter Count Register; CP0 Register 25, Select 1/3 ('x' = 0 or 1)

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	COUNTER<31:24>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	COUNTER<23:16>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	COUNTER<15:8>							
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	COUNTER <7:0>							
Legend:								
R = Readable bit			W = Writable bit			U = Unimplemented bit, read as '0'		
-n = Value at POR			'1' = Bit is set			'0' = Bit is cleared		
						x = Bit is unknown		

bit 31-0 **COUNTER<31:0>**: Event Counter bits
Counter for events enabled for this counter.

50.13.50 ErrCtl Register (CP0 Register 26, Select 0) (MPU only)

The ErrCtl register provides for software testing of the way-selection and RAM.

The way-selection RAM test mode is enabled by setting the WST bit. It modifies the functionality of the CACHE Index Load Tag and Index Store Tag operations so that they modify the way-selection RAM and leave the Tag RAMs untouched. When this bit is set, the lower six bits of the PA field in the TagLo register are used as the source and destination for Index Load Tag and Index Store Tag CACHE operations.

The WST bit also enables the data RAM test mode. When this bit is set, the Index Store Data CACHE instruction is enabled. This CACHE operation writes the contents of the DataLo register to the word in the data array that is indicated by the index and byte address.

Register 50-50: ErrCtl: Parity Protection Control Register; CP0 Register 26, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	R/W-0	U-0	U-0	U-0	U-0	U-0
	—	—	WST	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-28 **Unimplemented:** Read as '0'

bit 29 **WST:** Way-Select/Tag Array bit

Indicates whether the tag array or the way-select array should be read/written on Index Load/Store Tag CACHE instructions. This bit also enables the Index Store Data CACHE instruction, which writes the contents of DataLo to the data array.

1 = Way-select array is read/written on Index Load/Store tag CACHE instruction

0 = Tag array is read/written on Index Load/Store tag CACHE instruction

bit 27-0 **Unimplemented:** Read as '0'

50.13.51 CacheErr Register (CP0 Register 27, Select 0) (M-Class only)

The CacheErr register provides an interface with the cache error-detection logic. When a caches/SPRAM Parity Error exception is signaled, the fields of this register are set accordingly.

Register 50-51: CacheErr: Cache Error Register; CP0 Register 27, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R-0	R-0	R-0	R-0	R-0	U-0	R-0	R-0
	ER	EC	ED	ET	ES	—	EB	EF
23:16	R-0	R-0	R-0	R-0	R/W-0	R-0	R-0	R-0
	SP	EW	Way<1:0>		Index<19:16>			
15:8	R-0	R-0	R-0	R-0	R-0	R-0	R-0	R-0
	Index<15:8>							
7:0	R-0	R-0	R-0	R-0	R-0	R-0	R-0	R-0
	Index<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31 **ER:** Error Reference bit

Indicates the type of reference that encountered an error.

1 = Data reference is encountered

0 = Instruction reference is encountered

bit 30 **EC:** Cache Level Error Detected bit

Indicates the cache level at which the error was detected.

1 = Non-primary cache level is detected

0 = Primary cache level is detected

bit 29 **ED:** Error Data bit

Indicates a data RAM error.

1 = Data RAM error is detected

0 = No data RAM error is detected

bit 28 **ET:** Error Tag bit

Indicates a tag RAM error.

1 = Tag RAM error is detected

0 = No tag RAM error is detected

bit 27 **ES:** Error Source bit

Indicates whether error was caused by internal processor or external snoop request.

1 = Error on external request

0 = Error on internal request

bit 26 **Unimplemented:** Read as '0'

bit 25 **EB:** Error Both bit

Indicates that a data caches/SPRAM parity error occurred in addition to an instruction caches/SPRAM parity error. In the case of an additional data caches/SPRAM parity error, the remainder of the bits in this register are set according to the instruction caches/SPRAM parity error.

1 = Additional data caches/SPRAM parity error

0 = No additional data caches/SPRAM parity error

Register 50-51: CacheErr: Cache Error Register; CP0 Register 27, Select 0 (Continued)

- bit 24 **EF:** Error Fatal bit
Indicates that a fatal cache error has occurred.
There are a few situations in which software will not be able to get all information about a cache error from the CacheErr register. These situations are fatal, because software cannot determine which memory locations have been affected by the error. To enable software to detect these cases, the EF bit has been added to the CacheErr register.
The following cases are indicated as fatal cache errors by the EF bit:
- Dirty parity error in dirty victim (dirty bit cleared)
 - Tag parity error in dirty victim
 - Data parity error in dirty victim
 - WB store miss and EW error at the requested index
- In addition, simultaneous instruction and data cache errors as indicated by the EB bit will cause information about the data cache error to be unavailable. However, that situation is not indicated by the EF bit.
- bit 23 **SP:** Scratchpad bit
1 = Scratchpad RAM error is detected
0 = No Scratchpad RAM error is detected
- bit 22 **EW:** Error Way bit
1 = Way selection RAM error is detected
0 = No way selection RAM is detected
- bit 21-20 **Way<1:0>:** Way bits
Specifies the cache way in which the error was detected. It is not valid if a Tag RAM error is detected (ET = 1) or Scratchpad RAM error is detected (SP = 1).
- bit 19-0 **Index<19:0>:** Index bits
Specifies the cache or Scratchpad RAM index of the double word in which the error was detected. The way of the faulty cache is written by hardware in the Way field. Software must combine the Way and Index read in this register with cache configuration information in the Config1 register in order to obtain an index which can be used in an indexed Cache instruction to access the faulty cache data or tag. Note that Index is aligned as a byte index, so it does not need to be shifted by software before it is used in an indexed Cache instruction. Index bits 4-3 are undefined upon tag RAM errors, and Index bits above the MSB actually used for cache indexing will also be undefined. Bits 19-16 are only used for errors in the Scratchpad RAM.

50.13.52 TagLo Register (CP0 Register 28, Select 0) When WST = 0 (ErrCtl<29>) (MPU only)

The TagLo register acts as the interface to the cache tag array. The Index Store Tag and Index Load Tag operations of the CACHE instruction use the TagLo register as the source of tag information.

When the WST bit of the ErrCtl register is asserted, this register becomes the interface to the way-selection RAM. In this mode, the fields are redefined to give appropriate access the contents of the WS array instead of the Tag array.

Note: [Register 50-53](#) describes the fields in the TagLo register when WST = 1.

Register 50-52: TagLo: Cache Tag Array Interface Register; CP0 Register 28, Select 0 (When WST = 0)

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	PA<21:14>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	PA<13:6>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	U-0	U-0
	PA<5:0>						—	—
7:0	R/W-x	R/W-x	R/W-x	U-0	U-0	U-0	U-0	U-0
	V	D	L	—	—	—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-10 **PA<21:0>**: Cache Line Physical Address bits

These bits contain the physical address of the cache line. PA<21> corresponds to bit 31 of the physical address and PA<0> corresponds to bit 10 of the physical address.

bit 9-8 **Unimplemented**: Read as '0'

bit 7 **V**: Valid Cache Line Status bit

This bit indicates whether the cache line is valid.

bit 6 **D**: Dirty Cache Line Status bit

This bit indicates whether the cache line is dirty. This bit is only set if bit 7 (V) is also set.

bit 5 **L**: Cache Tag Lock Status bit

This bit specifies the lock bit for the cache tag. When this bit is set, and bit 7 (V) is set, the corresponding cache line will not be replaced by the cache replacement algorithm.

bit 4-0 **Unimplemented**: Read as '0'

50.13.53 TagLo Register (CP0 Register 28, Select 0) When WST = 1 (ErrCtl<29>) (MPU only)

The TagLo register acts as the interface to the cache tag array. The Index Store Tag and Index Load Tag operations of the CACHE instruction use the TagLo register as the source of tag information.

When the WST bit of the ErrCtl register is asserted, this register becomes the interface to the way-selection RAM. In this mode, the fields are redefined to give appropriate access the contents of the WS array instead of the Tag array.

Note: [Register 50-52](#) describes the fields in the TagLo register when WST = 0.

Register 50-53: TagLo: Cache Tag Array Interface Register; CP0 Register 28, Select 0 (When WST = 1)

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	R/W-x	R/W-x	R/W-x	R/W-x
	—	—	—	—	WSD<3:0>			
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	U-0	U-0
	WSLRU<5:0>						—	—
7:0	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-20 **Unimplemented:** Read as '0'

bit 19-16 **WSD<3:0>:** Way-Select Dirty Status bits

These bits contain the value read from the WS array after a CACHE Index Load WS operation. It is used to store into the WS array during CACHE Index Store WS operations.

bit 15-10 **WSLRU<5:0>:** Way-Select Least Recently Used bits

These bits contain the value read from or to be stored to the WS array.

bit 9-0 **Unimplemented:** Read as '0'.

50.13.54 DataLo Register (CP0 Register 28, Select 1) (MPU only)

The DataLo register is a register that acts as the interface to the cache data array and is intended for diagnostic operations only. The Index Load Tag operation of the `CACHE` instruction reads the corresponding data values into the DataLo register. If the WST bit in the ErrCtl register is set, the contents of DataLo can be written to the cache data array by doing an Index Store Data `CACHE` instruction.

Register 50-54: DataLo: Cache Data Array Interface Register; CP0 Register 28, Select 1

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DATA<31:24>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DATA<23:16>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DATA<15:8>							
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	DATA<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **DATA<31:0>**: Low-order Data Read from Cache Data Array bits

50.13.55 ErrorEPC (CP0 Register 30, Select 0)

The ErrorEPC register is a read/write register, similar to the EPC register, except that ErrorEPC is used on error exceptions. All bits of the ErrorEPC register are significant and must be writable. It is also used to store the program counter on Reset, Soft Reset, and non-maskable interrupt (NMI) exceptions.

The ErrorEPC register contains the virtual address at which instruction processing can resume after servicing an error. This address can be:

- The virtual address of the instruction that caused the exception
- The virtual address of the immediately preceding branch or jump instruction when the error causing instruction is in a branch delay slot

Unlike the EPC register, there is no corresponding branch delay slot indication for the ErrorEPC register.

Since the PIC32 family implements the MIPS16e® or microMIPS ASE, a read of the ErrorEPC register (via MFC0) returns the following value in the destination GPR:

```
GPR[rt] = ErrorExceptionPC31..1 || ISAMode0
```

That is, the upper 31 bits of the error exception PC are combined with the lower bit of the ISA<1:0> bits (Config3<15:14>) and are written to the GPR.

Similarly, a write to the ErrorEPC register (via MTC0) takes the value from the GPR and distributes that value to the error exception PC and the ISA<1:0> bits (Config3<15:14>), as follows:

```
ErrorExceptionPC = GPR[rt]31..1 || 0
ISAMode = 2#0 || GPR[rt]0
```

That is, the upper 31 bits of the GPR are written to the upper 31 bits of the error exception PC, and the lower bit of the error exception PC is cleared. The upper bit of the ISA<1:0> bits (Config3<15:14>) is cleared and the lower bit is loaded from the lower bit of the GPR.

Register 50-55: ErrorEPC: Error Exception Program Counter Register; CP0 Register 30, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
ErrorEPC<31:24>								
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
ErrorEPC<23:16>								
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
ErrorEPC<15:8>								
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
ErrorEPC<7:0>								

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **ErrorEPC<31:0>**: Error Exception Program Counter bits

50.13.56 DeSAVE Register (CP0 Register 31, Select 0)

The DeSAVE register is a read/write register that functions as a simple memory location. This register is used by the debug exception handler to save one of the GPRs that is then used to save the rest of the context to a predetermined memory area (such as in the EJTAG Probe). This register allows the safe debugging of exception handlers and other types of code where the existence of a valid stack for context saving cannot be assumed.

Register 50-56: DeSAVE: Debug Exception Save Register; CP0 Register 31, Select 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
DESAVE<31:24>								
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
DESAVE<23:16>								
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
DESAVE<15:8>								
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
DESAVE<7:0>								

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **DESAVE<31:0>**: Debug Exception Save bits

Scratch Pad register used by Debug Exception code.

50.13.57 KScratch n Registers (CP0 Register 31, Select 2-3) (M-Class only)

The KScratch n registers are optional read/write registers available for scratchpad storage by kernel-mode software. These registers are 32 bits in width for 32-bit processors and 64 bits for 64-bit processors.

The existence of these registers is indicated by the KScrExist field in the Config4 register. The KScrExist field specifies which of the selects are populated with a kernel scratch register.

Debug-mode software should not use these registers; instead, use the DeSave register. If EJTAG is implemented, select 0 should not be used for a KScratch register. Select 1 is being reserved for future debug use and should not be used for a KScratch register.

Register 50-57: KScratch n : Kernel Mode Scratchpad Registers; CP0 Register 31, Select 2-3

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	Data<31:24>							
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	Data<23:16>							
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	Data<15:8>							
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	Data<7:0>							

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as '0'

-n = Value at POR

'1' = Bit is set

'0' = Bit is cleared

x = Bit is unknown

bit 31-0 **Data<31:0>**: Scratchpad Data Saved by Kernel Software bits

50.14 COPROCESSOR 1 (CP1) REGISTERS

There are five Coprocessor 1 (CP1) registers, also referred to as FPU Control Registers (FCRs), that are used to control the FPU. These registers are 32 bits wide: FIR, FCCR, FEXR, FENR, FCSR. Three of these registers: FCCR, FEXR, and FENR will select subsets of the FCSR, the floating point Control/Status register.

Note: The Coprocessor 1 (CP1) registers are only available in devices with the M-Class core. Refer to the “CPU” chapter of the specific device data sheet to determine availability.

CP1 control registers are summarized in [Table 50-32](#).

Table 50-32: Coprocessor 1 Register Summary

Register Number	Register Name	Function
0	FIR	Floating Point Implementation register. Contains information that identifies the FPU.
25	FCCR	Floating Point Condition Codes register.
26	FEXR	Floating Point Exceptions register.
28	FENR	Floating Point Enables register.
31	FCSR	Floating Point Control and Status register.

50.14.1 Floating Point Register (FIR, CP1 Control Register 0) (M-Class only)

The Floating Point Implementation Register (FIR) is a 32-bit read-only register that contains information identifying the capabilities of the FPU, the Floating Point processor identification, and the revision level of the FPU.

Register 50-58: FIR: Floating Point Implementation Register; CP1 Control Register 0

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	R-1	U-0	U-0	U-0	R-1
	—	—	—	UFRP	—	—	—	FC
23:16	R-1	R-1	R-1	R-1	R-1	R-1	R-1	R-1
	Has2008	F64	L	W	3D	PS	D	S
15:8	R-1	R-1	R-1	R-0	R-0	R-1	R-1	R-1
	ProcessorID<7:0>							
7:0	R-x	R-x	R-x	R-x	R-x	R-x	R-x	R-x
	Revision<7:0>							

Legend:	W = Writable bit	U = Unimplemented bit, read as '0'
R = Readable bit	HC = Hardware Set	HS = Hardware Cleared
-n = Value at POR	'1' = Bit is set	'0' = Bit is cleared
		x = Bit is unknown

bit 31-29 **Unimplemented:** Read as '0'

bit 28 **UFRP:** User Mode Status_{FR} Switching Instruction Support Status bit
1 = User mode Status_{FR} switching instructions are supported. The UFR bit in the [Config5 Register \(CP0 Register 16, Select 5\)](#) specifies whether user access mode instructions are enabled.

bit 27-25 **Unimplemented:** Read as '0'

bit 24 **FC:** Full Convert Ranges Implementation bit
1 = This bit is always '1' to indicate that full convert ranges are implemented. This means that all numbers can be converted to another type by the FPU (If FS bit in FCSR is not set the Unimplemented Operation exception can still occur on denormal operands though).

bit 23 **Has2008:** IEEE-754-2008 Implemented bit
1 = This bit is always set to '1' to indicate that the MAC2008, ABS2008, NAN2008 bits within the FCSR register exist (see [Register 50-62](#))

bit 22 **F64:** 64-bit FPU bit
1 = This bit is always '1' to indicate that this is a 64-bit FPU.

bit 21 **L:** Long Fixed point data type implementation bit
1 = This bit is always '1' to indicate that long fixed point data types are implemented.

bit 20 **W:** Word Fixed point data type implementation bit
1 = This bit is always '1' to indicate that word fixed point data types are implemented.

bit 19 **3D:** MIPS-3D ASE implementation bit
0 = This bit is always '0' to indicate that MIPS-3D is not implemented

bit 18 **PS:** Paired-Single implementation bit
0 = This bit is always '0' to indicate that paired-single floating point data types are not implemented.

bit 17 **D:** Double-Precision implementation bit
1 = This bit is always '1' to indicate that double-precision floating point data types are implemented.

bit 16 **S:** Single-Precision implementation bit
1 = This bit is always '1' to indicate that single-precision floating point data types are implemented.

bit 15-8 **ProcessorID<7:0>:** Processor ID bits
This value matches the corresponding field of the PRID Register (CP0 Register 15, Select 0).

bit 7-0 **Revision<7:0>:** FPU Revision number bits
Specifies the revision number of the FPU. This field allows software to distinguish between different revisions of the same floating point processor type.

50.14.2 Floating Point Condition Codes Register (FCCR, CP1 Control Register 25) (M-Class only)

The Floating Point Condition Codes Register (FCCR) provides an alternative way to read and write the floating point condition code values that also appear in the FCSR register. Unlike the FCSR, all eight FCC bits are contiguous in the FCCR register.

Register 50-59: FCC: Floating Point Condition Codes Register; CP1 Control Register 25

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	FCC<7:0>							

Legend:

R = Readable bit

-n = Value at POR

W = Writable bit

HC = Hardware Set

'1' = Bit is set

U = Unimplemented bit, read as '0'

HS = Hardware Cleared

'0' = Bit is cleared

r = Reserved

x = Bit is unknown

bit 31-8 **Unimplemented:** Read as '0'

bit 7-0 **FCC<7:0>:** FPU Floating Point Condition Code bits

Refer to [Register 50-62](#): "Floating Point Control and Status Register (FCSR, CP1 Control Register 31)".

50.14.3 Floating Point Exceptions Register (FEXR, CP1 Control Register 26) (M-Class only)

The Floating Point Exceptions Register (FEXR) provides an alternative way to read and write the Cause and Flags bits that also appear in the FCSR register.

Register 50-60: FEXR: Floating Point Exceptions Register; CP1 Control Register 26

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —
23:16	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	U-0 —	R/W-x Cause<5:4>	R/W-x —
15:8	R/W-x Cause<3:0>	R/W-x —	R/W-x —	R/W-x —	U-0 —	U-0 —	U-0 —	U-0 —
7:0	U-0 —	R/W-x Flags<4:0>	R/W-x —	R/W-x —	R/W-x —	R/W-x —	U-0 —	U-0 —

Legend:	W = Writable bit	U = Unimplemented bit, read as '0'
R = Readable bit	HC = Hardware Set	HS = Hardware Cleared
-n = Value at POR	'1' = Bit is set	'0' = Bit is cleared
		r = Reserved
		x = Bit is unknown

bit 31-18 **Unimplemented:** Read as '0'

bit 17-12 **Cause<5:0>:** Floating Point Cause bits

The Cause bits are E, V, Z, O, U, I. See [Table 50-33](#) and [Register 50-62](#) for the definitions of these bits.

bit 11-7 **Unimplemented:** Read as '0'

bit 6-2 **Flags<4:0>:** Floating Point Flag bits

The Flag bits are V, Z, O, U, I. See [Table 50-33](#) and [Register 50-62](#) for the definitions of these bits.

bit 1-0 **Unimplemented:** Read as '0'

Table 50-33: Cause, Enables, and Flags Definitions

Bit Name	Description
E	Unimplemented Operation (this bit exists only in the Cause field).
V	Invalid Operation.
Z	Divide-by-Zero.
O	Overflow.
U	Underflow.
I	Inexact.

50.14.4 Floating Point Enables Register (FENR, CP1 Control Register 28) (M-Class only)

The Floating Point Enables Register (FENR) provides an alternative way to read and write the Enables, FS, and RM bits that also appear in the FCSR register.

Register 50-61: FENR: Floating Point Enables Register; CP1 Control Register 28

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
23:16	U-0	U-0	U-0	U-0	U-0	U-0	U-0	U-0
	—	—	—	—	—	—	—	—
15:8	U-0	U-0	U-0	U-0	R/W-0	R/W-0	R/W-0	R/W-0
	—	—	—	—	Enables<4:1>			
7:0	R/W-0	U-0	U-0	U-0	U-0	R/W-0	R/W-0	R/W-0
	Enables<0>	—	—	—	—	FS	RM<1:0>	

Legend:	W = Writable bit	U = Unimplemented bit, read as '0'
R = Readable bit	HC = Hardware Set	HS = Hardware Cleared
-n = Value at POR	'1' = Bit is set	'0' = Bit is cleared
		r = Reserved
		x = Bit is unknown

bit 31-12 **Unimplemented:** Read as '0'

bit 11-7 **Enables<4:0>:** Floating Point Enable bits

The Enable bits are V, Z, O, U, I. See [Table 50-33](#) and [Register 50-62](#) for the definitions of these bits.

bit 6-3 **Unimplemented:** Read as '0'

bit 2 **FS:** Flush-to-Zero bit

1 = Flush-to-Zero is enabled

0 = Flush-to-Zero is disabled

See [Register 50-62](#) for details on the Flush-to-Zero operation.

bit 1-0 **RM<1:0>:** Rounding Mode bits

These bits indicate the rounding mode used for most floating point operations.

See [Table 50-34](#) and [Register 50-62](#) for details on Rounding mode.

50.14.5 Floating Point Control and Status Register (FCSR, CP1 Control Register 31) (M-Class only)

The Floating Point Control and Status Register (FCSR) controls the operation of the Floating Point Unit.

It allows control and status report of the floating point unit:

- Selects the default rounding mode for FPU arithmetic operations
- Enables traps of FPU exception conditions
- Controls denormalized number handling options
- Reports any IEEE exceptions that occur during the execution of the floating point instructions
- Reports any IEEE exceptions that cumulatively arose in completed floating point instructions
- Indicates the condition code result of floating point compare instructions

Access to the FCSR is not privileged; it can be read or written by any program that has access to the FPU (via the coprocessor enables in the Status Register (CP0 Register 12, Select 0).

Register 50-62: FCSR: Floating Point Control and Status Register; CP1 Control Register 31

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	FCC<7:1>							FS
23:16	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	FCC<0>	FO	FN	MAC2008	ABS2008	NAN2008	Causes<5:4>	
15:8	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	Causes<3:0>				Enables<4:1>			
7:0	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x	R/W-x
	Enables<0>	Flags<4:0>					RM<1:0>	

Legend:	W = Writable bit	U = Unimplemented bit, read as '0'
R = Readable bit	HC = Hardware Set	HS = Hardware Cleared
-n = Value at POR	'1' = Bit is set	'0' = Bit is cleared
		x = Bit is unknown

bit 31-25 **FCC<7:1>**: Floating Point Condition Codes bits

These bits record the result of floating point compares and are tested for floating point conditional branches and conditional moves.

Use of the FCC bits is specified in the compare, branch, or conditional move instruction.

bit 24 **FS**: Flush-to-Zero bit

This bit controls the handling of denormalized operands and tiny results. See [50.14.6 "Floating Point Operation of the FS/FO/FN Bits"](#) for details regarding the operation of these bits.

bit 23 **FCC<0>**: Floating Point Condition Codes bit

This bit controls the handling of denormalized operands and tiny results. See [50.14.6 "Floating Point Operation of the FS/FO/FN Bits"](#) for details regarding the operation of this bit.

bit 22 **FO**: Flush Override bit

See [50.14.6 "Floating Point Operation of the FS/FO/FN Bits"](#) for details regarding the operation of this bit.

bit 21 **FN**: Flush to Nearest bit

See [50.14.6 "Floating Point Operation of the FS/FO/FN Bits"](#) for details regarding the operation of this bit

Register 50-62: FCSR: Floating Point Control and Status Register; CP1 Control Register 31 (Continued)

bit 20 **MAC2008:** Fused Multiply-Add Mode bit

0 = Unfused multiply-add

The IEEE 754-2008 Standard fused multiply-add operation multiplies and adds with unbounded range and precision, rounding only once to the destination format.

The fused multiply-add is not supported in the PIC32 core.

The PIC32 FPU implements the unfused multiply-add, which rounds the intermediary multiplication result to the destination format.

This field applies to the `MADD.fmt`, `NMADD.fmt`, `MSUB.fmt`, and `NMSUB.fmt` instructions.

bit 19 **ABS2008:** ABS and NEG instructions compliant bit

1 = ABS & NEG accept QNaN input without trapping.

The IEEE 754-2008 standard requires that the ABS and NEG functions accept QNaN inputs without trapping.

This bit is always set in the PIC32 core to indicate support for the IEEE 754-2008 standard.

bit 18 **NAN2008:** Quiet and signaling NaN encodings bit

1 = IEEE 754-2008 NaN encoding

As recommended by the IEEE 754-2008 Standard a quiet NaN is encoded with the first bit of the fraction being 1 and a signaling NaN is encoded with the first bit of the fraction field being 0.

This bit is always set to in the PIC32 core to indicate support for the IEEE 754-2008 Standard encoding.

bit 17-12 **Cause<3:0>:** Cause bits

1 = Corresponding exception condition has occurred during the execution of an instruction

0 = Corresponding exception condition has not occurred during the execution of an instruction

These bits indicate the exception conditions that occur during execution of an FPU arithmetic instruction.

By reading the registers, the exception condition caused by the preceding FPU arithmetic instruction can be determined.

See [Table 50-33](#) for the definitions of the cause bits

bit 11-7 **Enables<4:0>:** Enable trap bits

1 = The trap occurs when the corresponding cause bit is set

0 = The trap does not occur when the corresponding cause bit is set

These bits control whether or not a trap is taken when an floating point exception condition occurs for any of the five conditions (V, Z, O, U, I). The trap occurs when both an enable bit and its corresponding cause bit are set either during an FPU arithmetic operation or by moving a value to the FCSR or one of its alternative representations.

The Cause bit E has no corresponding enable bit; the MIPS architecture defines non-IEEE Unimplemented Operation exceptions as always enabled.

See [Table 50-33](#) for the definitions of the enable bits.

bit 6-2 **Flags<4:0>:** Exception Flag bits

1 = Corresponding exception condition has occurred for completed FPU instructions

0 = Corresponding exception condition has not occurred for completed FPU instructions

This field shows any exception conditions that have occurred for completed instructions since the flag was last reset by software. When an FPU arithmetic operation raises an IEEE exception condition that does not result in a Floating Point Exception (the enable bit was off), the corresponding bit(s) in the Flags field are set, while the others remain unchanged.

Arithmetic operations that result in a Floating Point Exception (the enable bit was on) do not update the Flags field.

Note: Hardware never resets this field; software must explicitly reset this field.

See [Table 50-33](#) for the definitions of the flag bits.

bit 1-0 **RM<1:0>:** Rounding Mode bits

This field indicates the rounding mode used for most floating point operations (some operations use a specific rounding mode).

See [Table 50-34](#) for details about Rounding mode.

Table 50-34: Rounding Mode Definitions

RM Field Encoding	RM Field Description
0	RN - Round to Nearest Rounds the result to the nearest representable value. When two representable values are equally near, the result is rounded to the value whose least significant bit is zero (even).
1	RZ - Round Toward Zero Rounds the result to the value closest to but not greater in magnitude than the result.
2	RP - Round Towards Plus Infinity Rounds the result to the value closest to but not less than the result.
3	RM - Round Towards Minus Infinity Rounds the result to the value closest to but not greater than the result.

50.14.6 Floating Point Operation of the FS/FO/FN Bits

The FS, FO, and FN bits in the CP1 FCSR register control handling of denormalized operands and tiny results (i.e. nonzero result between $\pm 2^{E_{\min}}$), whereby the FPU can handle these cases directly instead of relying on the much slower software handler. The trade-off is a loss of IEEE compliance and accuracy (except for use of the FO bit), because a minimal normalized or zero result is provided by the FPU instead of the more accurate denormalized result that a software handler would give. The benefit is a significantly improved performance and precision.

Use of the FS, FO, and FN bits affects handling of denormalized floating point numbers and tiny results for the instructions listed below:

Table 50-35: FS/FO/FN Bits Floating Point Operations

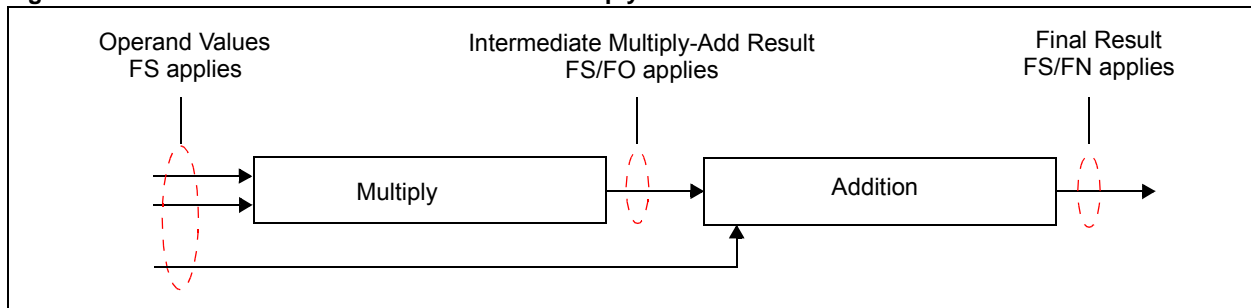
FS/FO/FN Bits	Affected Instructions
FS and FN	ADD, CEIL, CVT, DIV, FLOOR, MADD, MSUB, MUL, NMADD, NMSUB, RECIP, ROUND, RSQRT, SQRT, TRUNC, SUB, ABS, C.cond, and NEG. See Note 1 .
FO	MADD, MSUB, NMADD, and NMSUB

Note 1: For ABS, C.cond, and NEG, denormal input operands or tiny results do not result in Unimplemented exceptions when FS = 0. Regardless, flushing to zero is implemented when FS = 1, such that these operations return the same result as an equivalent sequence of arithmetic FPU operations.

Instructions not listed in [Table 50-35](#) do not cause Unimplemented Operation exceptions on denormalized numbers in operands or results.

[Figure 50-24](#) illustrates how the FS, FO, and FN bits control handling of denormalized numbers. For instructions that are not multiply or add types (such as DIV), only the FS and FN bits apply.

Figure 50-24: FPU FS/FO/FN Bits Influence on Multiply and Addition Results



50.14.6.1 FPU FLUSH-TO-ZERO BIT

When the Flush-To-Zero (FS) bit is set, denormal input operands are flushed to zero. Tiny results are flushed to either zero or the applied format's smallest normalized number (MinNorm) depending on the rounding mode settings. [Table 50-36](#) lists the flushing behavior for tiny results.

Table 50-36: Zero Flushing for Tiny Results

Rounding Mode	Negative Tiny Result	Positive Tiny Result
RN (RM = 0)	-0	+0
RZ(RM = 1)	-0	+0
RP (RM = 2)	-0	+MinNorm
RM (RM = 3)	-MinNorm	+0

The flushing of results is based on an intermediate result computed by rounding the mantissa using an unbounded exponent range; that is, tiny numbers are not normalized into the supported exponent range by shifting in leading zeros prior to rounding.

Handling of denormalized operand values and tiny results depends on the FS bit setting as shown in [Table 50-37](#).

Table 50-37: Handling of Denormalized Operand Values and Tiny Results Based on FS Bit Setting

FS Bit Setting	Handling of Denormalized Operand Values
0	An Unimplemented Operation exception is taken
1	Instead of causing an Unimplemented Operation exception, operands are flushed to zero, and tiny results are forced to zero or MinNorm.

50.14.6.2 FPU FLUSH OVERRIDE BIT

When the Flush Override (FO) bit is set, a tiny intermediate result of any multiply-add type instruction is not flushed according to the FS bit. The intermediate result is maintained in an internal normalized format to improve accuracy.

FO only applies to the intermediate result of a multiply-add type instruction.

Handling of tiny intermediate results depends on the FO and FS bits as shown in [Table 50-38](#).

Table 50-38: Handling of Tiny Intermediate Result Based on the FO and FS Bit Settings


FO Bit Setting	FS Bit Setting	Handling of Tiny Result Values
0	0	An Unimplemented Operation exception is taken.
0	1	The intermediate result is forced to the value that would have been delivered for an untrapped underflow (see Table 50-21 : "FPU Supplied Results for Not Trapped Exceptions") instead of causing an Unimplemented Operation exception.
1	x	The intermediate result is kept in an internal format, which can be perceived as having the usual mantissa precision but with unlimited exponent precision and without forcing to a specific value or taking an exception.

50.14.6.3 FPU FLUSH-TO-NEAREST BIT

When the Flush-to-Nearest (FN) bit is set and the rounding mode is Round to Nearest (RN), a tiny final result is flushed to zero or 2E_min (MinNorm). If a tiny number is strictly below MinNorm/2, the result is flushed to zero; otherwise, it is flushed to MinNorm (see [Figure 50-25](#)). The flushed result has the same sign as the result prior to flushing. Note that the FN bit takes precedence over the FS bit.

For all rounding modes other than Round to Nearest (RN), setting the FN bit causes final results to be flushed to zero or MinNorm as if the FS bit was set.

Figure 50-25: FPU Flushing to Nearest when Rounding Mode is Round to Nearest



Handling of tiny final results depends on the FN and FS bits, as shown in [Table 50-39](#).

Table 50-39: FPU Handling of Tiny Final Result Based on FN and FS Bit Settings

FN Bit Setting	FS Bit Setting	Handling of Tiny Result Values
0	0	An Unimplemented Operation exception is taken.
0	1	Final result is forced to the value that would have been delivered for an untrapped underflow (see Table 50-21 : "FPU Supplied Results for Not Trapped Exceptions") rather than causing an Unimplemented Operation exception.
1	x	Final result is rounded to either zero or $2^{E_{\min}}$ (MinNorm), whichever is closest when in Round to Nearest (RN) rounding mode. For other rounding modes, a final result is given as if FS was set to 1.

50.14.6.4 FPU RECOMMENDED FS/FO/FN BIT SETTINGS

[Table 50-40](#) summarizes the recommended settings for the FPU FS/FO/FN bits.

Table 50-40: Recommended settings for the FPU FS/FO/FN bits

FS Bit Setting	FO Bit Setting	FN Bit Setting	Remarks
0	0	0	IEEE-compliant mode. Low performance on denormal operands and tiny results
1	0	0	Regular embedded applications. High performance on denormal operands and tiny results.
1	1	1	Highest accuracy and performance configuration. Note: in this mode MADD might return a different result other than the equivalent MUL and ADD operation sequence.

50.14.6.5 FPU FCSR CAUSE BIT UPDATE FLOW

50.14.6.5.1 FPU Exceptions Triggered by CTC1 instruction

Regardless of the targeted control register, the co-processor CTC1 instruction causes the Enables and Cause fields of the FCSR to be inspected to determine if an exception is to be thrown.

50.14.6.6 FPU COMPUTATIONS GENERIC FLOW

Computations are performed in two steps:

1. Compute rounded mantissa with unbound exponent range.
2. Flush to default result if the result from Step 1 is overflow or tiny (no flushing happens on denormalized results for instructions supporting denormalized results, such as MOV).

The Cause field is updated after each of these two steps. Any enabled exceptions detected in these two steps cause a trap, and no further updates to the Cause field are done by subsequent steps.

Step 1 can set Cause bits I, U, O, Z, V, and E.

E has priority over V; V has priority over Z; and Z has priority over U and O. Thus when E, V, or Z is set in Step 1, no other cause bits can be set. However, note that I and V both can be set if a denormalized operand was flushed (FS = 1). I, U, and O can be set alone or in pairs (IU or IO). U and O never can be set simultaneously in Step 1. U and O are set if the computed unbounded exponent is outside the exponent range supported by the normalized IEEE format.

Step 2 can set I if a default result is generated.

50.14.6.7 FPU MULTIPLY-ADD FLOW

For multiply-add type instructions the computation is extended with two more steps:

1. Compute rounded mantissa with unbound exponent range for the multiply.
2. Flush to default result if the result from Step 1 is overflow or tiny (no flushing happens on tiny results if FO = 1).
3. Compute rounded mantissa with unbounded exponent range for the addition.
4. Flush to default result if the result from Step 3 is overflow or tiny.

The Cause field is updated after each of these four steps. Any enabled exceptions detected in these four steps cause a trap, and no further updates to the Cause field are done by subsequent steps.

Step 1 and Step 3 can set a cause bit as described for Step 1 in [50.14.6.6 “FPU Computations Generic Flow”](#).

Step 2 and Step 4 can set I if a default result is generated.

Although U and O can never both be set in Step 1 or Step 3, both U and O might be set after the multiply-add has executed in Step 3 because U might be set in Step 1 and O might be set in Step 3.

50.14.6.8 FPU CAUSE UPDATE FLOW FOR INPUT OPERANDS

Denormalized input operands to Step 1 or Step 3 always set Cause bit I when FS = 1. For example, SNaN+DeNorm will set I (and V) provided that Step 3 was reached (in case of a multiply-add type instruction).

Conditions directly related to the input operand (for example, I/E set due to DeNorm, V set due to SNaN and QNaN propagation) are detected in the step where the operand is logically used. For example, for multiply-add type instructions, exceptional conditions caused by the input operand *fr* are detected in Step 3.

Note: The “*fr*” register is a FPU register that occurs in the MADD.fmt, NMADD.fmt, MSUB.fmt, and NMSUB.fmt FPU instructions.

50.14.6.9 FPU CAUSE UPDATE FLOW FOR UNIMPLEMENTED OPERATIONS

Note that the Cause bit, E, is special; it clears any Cause updates done in previous steps. For example, if Step 3 caused E to be set, any I, U, or O Cause update done in Step 1 or Step 2 is cleared. Only E is set in the Cause field when an Unimplemented Operation trap is taken.

50.15 microMIPS EXECUTION

microMIPS minimizes the code footprint of applications and therefore reduces the cost of memory, which is particularly high for embedded memory. Simultaneously, the high performance of MIPS cores is maintained. Using this technology, best results can be achieved without the need to spend time to profile the application. The smaller code footprint typically leads to reduced power consumption per executed task because of the smaller number of memory accesses.

The MIPS32 Release 3.0 Architecture supports both the MIPS32 instruction set and microMIPS, the enhanced MIPS32 instruction set.

50.16 MCU ASE EXTENSION

The MCU ASE extends the microMIPS and MIPS32 Architecture with a set of new features designed for the microcontroller market.

The MCU ASE contains enhancements in several distinct areas: interrupt delivery and interrupt latency.

50.16.1 Interrupt Delivery

The MCU ASE extends the number of interrupt hardware inputs from 63 to 255 (External Interrupt Controller (EIC) mode), with separate priority and vector generation.

50.16.2 Interrupt Latency Reduction

The MCU ASE includes a package of extensions to microMIPS and MIPS32 that decrease the latency of the processor's response to a signaled interrupt.

50.16.2.1 INTERRUPT VECTOR PREFETCHING

Normally on MIPS architecture processors, when an interrupt or exception is signaled, execution pipelines must be flushed before the interrupt/exception handler is fetched. This is necessary to avoid mixing the contexts of the interrupted/faulting program and the exception handler. The MCU ASE introduces a hardware mechanism in which the interrupt exception vector is prefetched whenever the interrupt input signals change. The prefetch memory transaction occurs in parallel with the pipeline flush and exception prioritization. This decreases the overall latency of the execution of the interrupt handler's first instruction.

50.16.2.2 AUTOMATED INTERRUPT PROLOGUE

The use of Shadow Register Sets avoids the software steps of having to save general-purpose registers before handling an interrupt.

The MCU ASE adds additional hardware logic that automatically saves some of the CP0 state in the stack and automatically updates some of the CP0 registers in preparation for interrupt handling.

50.16.2.3 AUTOMATED INTERRUPT EPILOGUE

A mirror to the Automated Prologue, this feature automates the restoration of some of the CP0 registers from the stack and the preparation of some of the CP0 registers for returning to Non-Exception mode. This feature is implemented within the `IRET` instruction, which is introduced in this ASE.

50.16.2.4 INTERRUPT CHAINING

An optional feature of the Automated Interrupt Epilogue, this feature allows handling a second interrupt after a primary interrupt is handled, without returning to Non-Exception mode (and the related pipeline flushes that would normally be necessary).

50.17 MIPS DSP ASE EXTENSION

Note: DSP ASE is not available on all devices. Refer to the “CPU” chapter in the specific device data sheet to determine availability.

The MIPS DSP Application-Specific Extension Revision 2 is an extension to the MIPS32 architecture. This extension comprises new integer instructions and states that include new HI/LO accumulator register pairs and a DSP control register. This extension is crucial in a wide range of DSP, multimedia, and DSP-like algorithms covering Audio and Video processing applications. The extension supports native fractional format data type operations, register Single Instruction Multiple Data (SIMD) operations, such as add, subtract, multiple, and shift. In addition, the extension includes the following features that are essential in making DSP algorithms computationally efficient:

- Support for multiplication of complex operands
- Variable bit insertion and extraction
- Implementation and use of virtual circular buffers
- Arithmetic saturation and overflow handling support
- Zero cycle overhead saturation and rounding operations

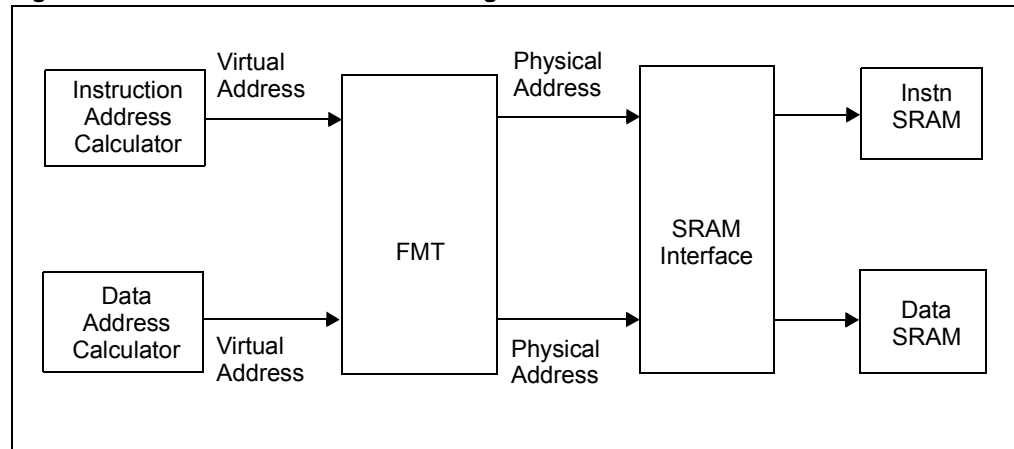
Refer to “MIPS32[®] Architecture for Programmers Volume IV-e: The MIPS[®] DSP Application-Specific Extension to the MIPS32[®] Architecture” – MD00374 for information on this extension. This document is available for download from the Imagination Technologies Ltd. website at: <http://www.imgtec.com/mips/architectures/dsp.asp>.

50.18 MEMORY MODEL (MCU ONLY)

Virtual addresses used by software are converted to physical addresses by the memory management unit (MMU) before being sent to the CPU busses. PIC32 devices based on the MCU Microprocessor core use a fixed mapping for this conversion.

For more information regarding the system memory model, refer to **Section 3. “Memory Organization”** (DS60001115) of the “PIC32 Family Reference Manual”.

Figure 50-26: Address Translation During SRAM Access



50.18.1 Cacheability

The CPU uses the virtual address of an instruction fetch, load or store to determine whether to access the cache or not. Memory accesses within kseg0, or useg/kuseg can be cached, while accesses within kseg1 are non-cacheable. The CPU uses the CCA bits in the Config register to determine the cacheability of a memory segment. A memory access is cacheable if its corresponding CCA = 011₂. For more information on cache operation, refer to **Section 4. “Prefetch Cache Module”** (DS60001119) of the “PIC32 Family Reference Manual”.

50.18.1.1 LITTLE ENDIAN BYTE ORDERING

On CPUs that address memory with byte resolution, there is a convention for multi-byte data items that specify the order of high-order to low-order bytes. Big-endian byte-ordering is where the lowest address has the MSB. Little-endian ordering is where the lowest address has the LSB of a multi-byte datum. The PIC32 CPU supports little-endian byte ordering.

Figure 50-27: Big-Endian Byte Ordering


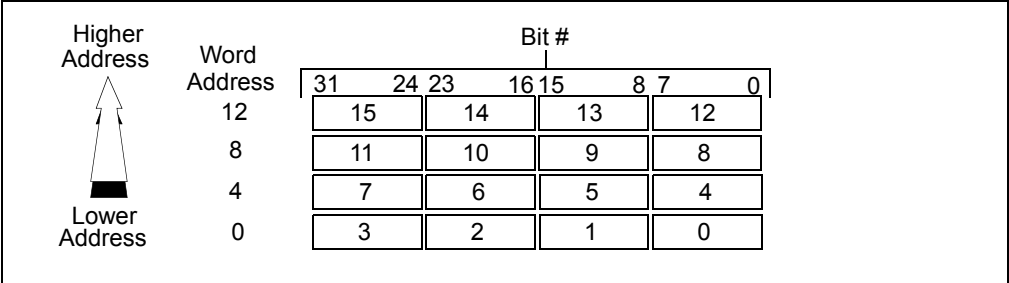


Figure 50-28: Little-Endian Byte Ordering



50.19 MEMORY MANAGEMENT (MPU ONLY)

PIC32 devices with the MPU core include a Memory Management Unit (MMU) that uses a Translation Lookaside Buffer (TLB) to translate a virtual page address to a physical page address. This feature is used by operating systems to manage multiple tasks running in the same virtual memory, by mapping them into separate physical memory locations. The MMU can also provide protection of physical memory areas and define the cache protocol. PIC32 devices with the MPU core support page sizes from 4 KB to 1 MB.

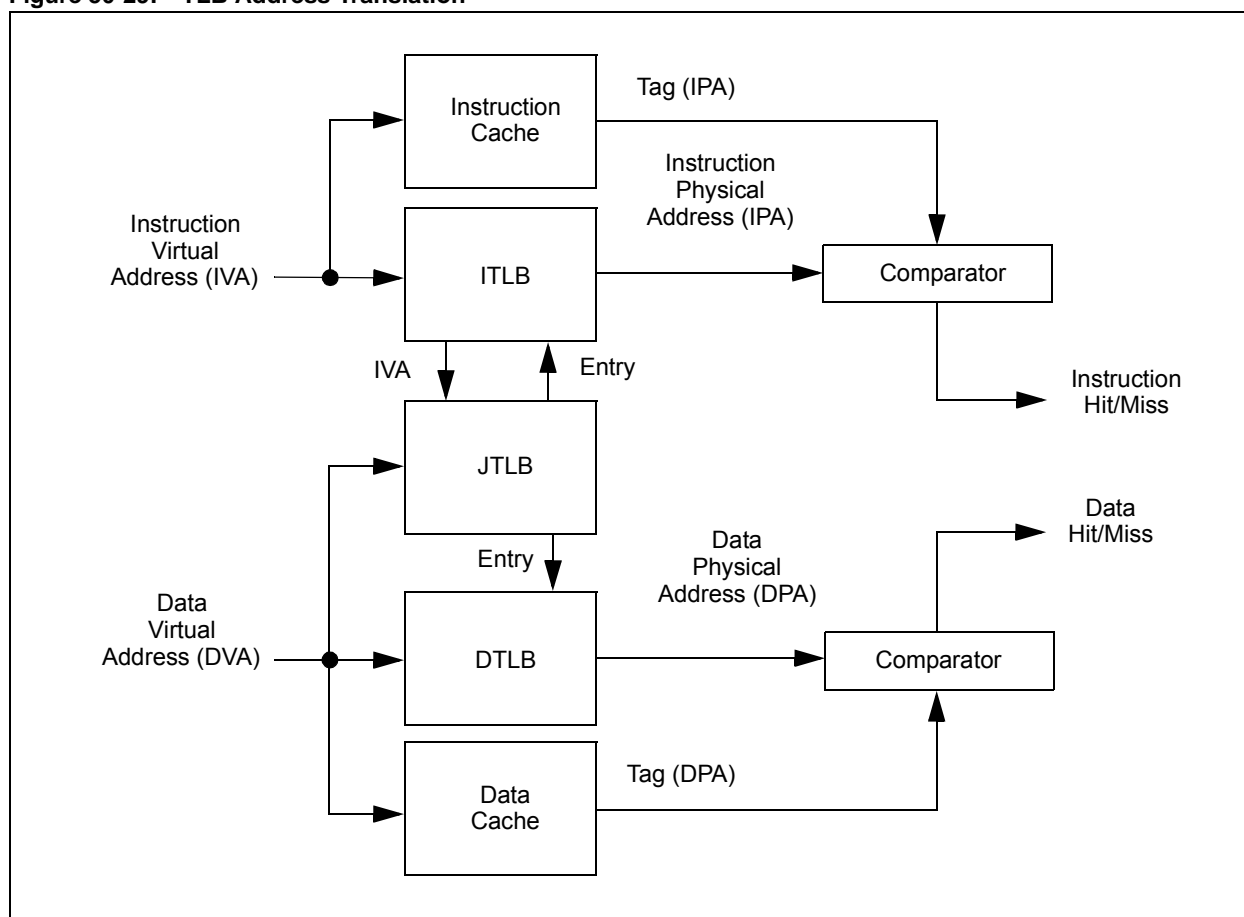
The core of the MMU is the TLB. The TLB contains three address translation buffers:

- 16 dual-entry fully associative Joint TLB (JTLB)
- 4-entry Instruction micro-TLB (ITLB)
- 4-entry Data micro-TLB (DTLB)

When a page address is translated, the ITLB or DTLB is accessed first. If the translation is not found in the micro-TLB, the JTLB is accessed. If the translation is not found in the JTLB, an exception is taken.

Figure 50-29 shows how the TLB interacts with cache accesses in PIC32 devices with the MPU core.

Figure 50-29: TLB Address Translation



50.19.1 Virtual Memory and Modes of Operation

All PIC32 devices support three modes of operation:

- User mode
- Kernel mode
- Debug mode

The core enters Kernel mode at reset, and when an exception is recognized. While in Kernel mode, the software has access to the entire 4 GB address space, as well as the CP0 registers.

User mode access is restricted to the first 2 GB of the address space (0x00000000 through 0x7FFFFFFF), and can be excluded from accessing CP0 functions. Accessing a virtual address above 0x7FFFFFFF in User mode will cause an exception.

Debug mode is entered on a debug exception. While in Debug mode, the software has access to all Kernel mode addresses and functions, as well as debug segment dseg, which overlays part of the kernel segment kseg3.

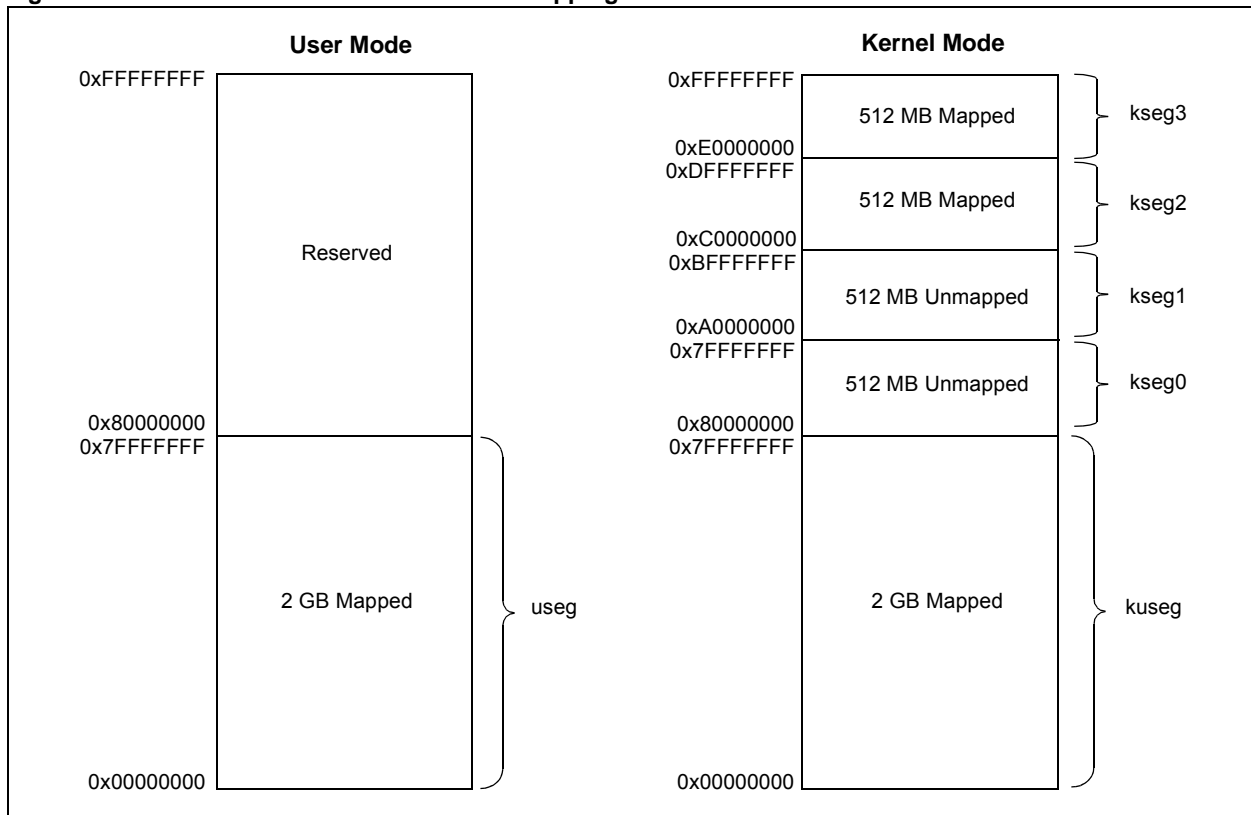
The virtual memory segments are different depending on the mode of operation. See [Figure 50-14](#) in [50.11.4 “Processor Modes”](#) for the PIC32 virtual memory map.

50.19.1.1 MAPPED AND UNMAPPED SEGMENTS

Memory segments that use the MMU to translate virtual page address into physical memory locations are considered to be mapped. Those that do not use the MMU are considered unmapped. Unmapped segments have a fixed translation from virtual to physical addresses. At reset, it is important to execute from unmapped memory until the TLB is programmed to perform address translation.

Except for kseg0, unmapped segments are always uncached. The cacheability of kseg0 is set in the K0 field of the CP0 register. The cacheability of mapped segments is set in the K23 and KU fields of the CP0 register. [Figure 50-30](#) shows the mapped and unmapped virtual memory areas for Kernel mode and User mode. When operating in Debug mode, the dseg area is unmapped. The mapping for all other areas is the same as Kernel mode.

Figure 50-30: User and Kernel Mode Virtual Mapping



50.19.2 Translation Lookaside Buffer (TLB)

The TLB consists of one joint and two micro address translation buffers:

- 16 dual-entry fully associative Joint TLB (JTLB)
- 4-entry fully associative Instruction micro-TLB (ITLB)
- 4-entry fully associative Data micro-TLB (DTLB)

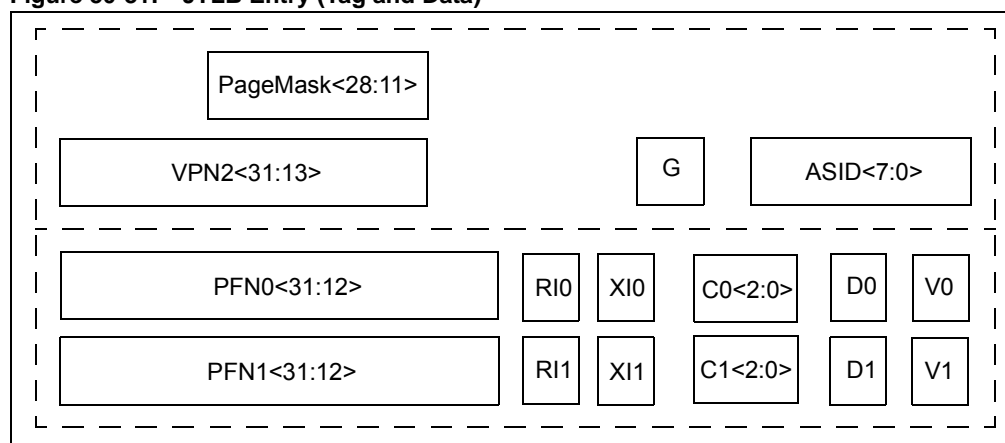
50.19.2.1 JOINT TLB

The JTLB maps 32 virtual page addresses to their corresponding physical addresses. The purpose of the TLB is to translate virtual page addresses along with their corresponding Address Space ID (ASID) into a physical memory address. Operating systems typically assign an ASID to each user program or process. The TLB helps the operating system separate each user process into its own physical memory space.

The virtual to physical translation is performed by comparing the upper bits of the virtual address and the ASID bits against each of the JTLB tag entries. The JTLB is referred to as a “Joint” TLB because it is used to translate both instruction and data virtual page addresses.

The JTLB is organized as pairs of even and odd entries containing address translations of pages ranging from 4 KB to 1 MB in size. Each virtual tag entry corresponds to two physical data entries, an even page entry and an odd page entry. The highest order virtual address bit not participating in the tag comparison is used to determine which of the two data entries is used. [Figure 50-31](#) shows the contents of one the dual entries in the JTLB.

Figure 50-31: JTLB Entry (Tag and Data)



PIC32 Family Reference Manual

Table 50-41 describes each field of a JTLB entry.

Table 50-41: TLB Tag Entry Fields

Bit Name	Description																		
PageMask<28:11>	<div>Page Mask Value. The Page Mask defines the page size by masking the appropriate VPN2 bits. It also determines which address bit is used to make the even-odd page (PFN0-PFN1) determination.</div> <table><tr><th>PageMask</th><th>Page Size</th><th>Even/Odd Bank Select Bit</th></tr><tr><td>00 0000 0000 0000 0011</td><td>4 KB</td><td>VAddr<12></td></tr><tr><td>00 0000 0000 0000 1111</td><td>16 KB</td><td>VAddr<14></td></tr><tr><td>00 0000 0000 0011 1111</td><td>64 KB</td><td>VAddr<16></td></tr><tr><td>00 0000 0000 1111 1111</td><td>256 KB</td><td>VAddr<18></td></tr><tr><td>00 0000 0011 1111 1111</td><td>1 MB</td><td>VAddr<20></td></tr></table>	PageMask	Page Size	Even/Odd Bank Select Bit	00 0000 0000 0000 0011	4 KB	VAddr<12>	00 0000 0000 0000 1111	16 KB	VAddr<14>	00 0000 0000 0011 1111	64 KB	VAddr<16>	00 0000 0000 1111 1111	256 KB	VAddr<18>	00 0000 0011 1111 1111	1 MB	VAddr<20>
PageMask	Page Size	Even/Odd Bank Select Bit																	
00 0000 0000 0000 0011	4 KB	VAddr<12>																	
00 0000 0000 0000 1111	16 KB	VAddr<14>																	
00 0000 0000 0011 1111	64 KB	VAddr<16>																	
00 0000 0000 1111 1111	256 KB	VAddr<18>																	
00 0000 0011 1111 1111	1 MB	VAddr<20>																	
VPN2<31:13>	Virtual Page Number divided by two. This field contains the upper bits of the virtual page number divided by two. Because it represents a pair of TLB pages, it is divided by two. Bits <31:25> are always included in the TLB lookup comparison. Bits <24:13> are included depending on the page size, defined by PageMask.																		
G	Global Bit. When set, indicates that this entry is global to all address spaces, and therefore excludes the ASID in the TLB lookup comparison.																		
ASID<7:0>	Address Space Identifier. Identifies which process or thread this TLB entry is associated with.																		
PFN0<31:12> PFN1<31:12>	Physical Frame Number. Defines the upper bits of the physical address. For page sizes larger than 4KB, only a subset of these bits is actually used.																		
C0<2:0> C1<2:0>	<div>Cacheability. Indicates the cacheability attributes and determines whether the page should be placed in the cache or not.</div> <table><tr><th>C<2:0></th><th>Coherency Attribute</th></tr><tr><td>000</td><td>Cacheable, non-coherent, write-through, no write-allocate</td></tr><tr><td>001</td><td>Cacheable, non-coherent, write-through, write-allocate</td></tr><tr><td>010</td><td>Uncached</td></tr><tr><td>011</td><td>Cacheable, non-coherent, write-back, write-allocate</td></tr><tr><td>100-111</td><td>Reserved</td></tr></table>	C<2:0>	Coherency Attribute	000	Cacheable, non-coherent, write-through, no write-allocate	001	Cacheable, non-coherent, write-through, write-allocate	010	Uncached	011	Cacheable, non-coherent, write-back, write-allocate	100-111	Reserved						
C<2:0>	Coherency Attribute																		
000	Cacheable, non-coherent, write-through, no write-allocate																		
001	Cacheable, non-coherent, write-through, write-allocate																		
010	Uncached																		
011	Cacheable, non-coherent, write-back, write-allocate																		
100-111	Reserved																		
R10 R11	Read Inhibit bit. Indicates that the page is read protected. If this bit is set, and the IEC bit of the PageGrain register is set, any attempt to read from the page will result in a TLB Read Inhibit exception.																		
X10 X11	Execute Inhibit bit. Indicates that the page is execution protected. If this bit is set, and the IEC bit of the PageGrain register is set, any attempt to fetch an instruction from the page will result in a TLB Execute Inhibit exception.																		
D0 D1	Dirty bit. Indicates that the page has been written and/or is writable. If this bit is set, writes to the page are allowed. If this bit is not set, writes to the page will result in a TLB Modified exception.																		
V0 V1	Valid bit. Indicates that the TLB entry is valid. If this bit is set, accesses to the page are permitted. If the bit is not set, accesses to the page will result in a TLB Invalid exception.																		

A table entry is filled with a TLBWI or TLBWR instruction. Before executing either instruction, the following CP0 registers must be updated with the information to be written to the TLB entry:

- PageMask is set in the PageMask register (CP0 Register 5, Select 0).
- VPN2, VPN2X, and ASID are set in the EntryHi register (CP0 Register 10, Select 0)
- PFN0, C0, D0, V0, and G are set in the EntryLo0 register (CP0 Register 2, Select 0)
- PFN1, C1, D1, V1, and G are set in the EntryLo1 register (CP0 Register 3, Select 0)

Note: The global bit is part of the EntryLo0 and EntryLo1 registers. The value written to the G bit in the JTLB is the logical AND of the G bits in the EntryLo0 and EntryLo1 registers.

50.19.2.2 micro-TLBs (ITLB AND DTLB)

The ITLB is a small (i.e., micro) TLB dedicated to the instruction stream. The DTLB is a small TLB that provides a faster translation for Load/Store addresses than is possible with the JTLB. Both are 4-entry, fully associative, and managed entirely by hardware.

Instruction fetch address translation is handled first by the ITLB. If the fetch misses the ITLB, the JTLB is accessed in the following clock cycle. If successful, the entry is copied into the ITLB. The ITLB is then accessed again, and the address is successfully translated.

Data translations access the ITLB and the JTLB in parallel. If there is a DTLB miss and a JTLB hit, the DTLB can be reloaded in the same clock cycle. The DTLB is then accessed on the next clock cycle.

50.19.3 Virtual to Physical Address Translation


Essentially, the TLB acts a cache for page table entries. When a virtual address is converted to a physical address, the Virtual Page Number (VPN) of the address (the upper bits of the address) is compared with the virtual page numbers stored in the TLB. A TLB hit occurs when the following conditions are met:

- The VPN of the virtual address matches a VPN stored in the TLB and:
 - The Global (G) bit of both the even and odd pages of the TLB entry are set, or
 - The ASID field of the virtual address is the same as the ASID field of the TLB entry

If these two conditions are not met, a TLB miss exception is taken by the processor and software is allowed to refill the TLB from a page table of virtual/physical addresses in memory.

Figure 50-32 shows the translation of a virtual address into a physical address.

Figure 50-32: Virtual to Physical Address Translation



Each JTLB entry contains a tag and two data fields. On a TLB hit, the upper bits of the virtual address are replaced with the Page Frame Number (PFN) stored in the corresponding data field. On a TLB miss, an exception is taken and software refills the TLB from a page table stored in memory.

50.19.4 TLB Entry Replacement

On a normal TLB miss, JTLB entries are replaced by a random replacement algorithm. The CPU also provides the ability to lock a programmable number of mappings into the TLB via the CP0 Wired register (see [Register 50-8](#) for details).

The JTLB supports pages of different sizes ranging from 4 KB to 1 MB in powers of four. The page size can be configured on a per-entry basis by loading the CP0 PageMask register with the desired page size prior to writing a new entry. A common use for this feature is to map a large memory block (such as a frame buffer) with a single TLB entry.

50.19.5 TLB Instructions

[Table 50-42](#) lists the TLB-related instructions supported by the PIC32. See [50.21.5.5 “TLB Instructions \(MPU Only\)”](#) for details on these instructions.

Table 50-42: TLB Instructions

Instruction	Description
TLBP	Translation Lookaside Buffer Probe
TLBR	Translation Lookaside Buffer Read
TLBWI	Translation Lookaside Buffer Write Index
TLBWR	Translation Lookaside Buffer Write Random

50.20 L1 CACHES (MPU ONLY)

PIC32 devices with the MPU microprocessor core have separate instruction and data caches. The use of separate caches allows instruction and data references to happen simultaneously. Both caches are Virtually Index, Physically-tagged (VIPT), allowing cache access to occur in parallel with virtual-to-physical address translation.

50.20.1 Cache Configuration

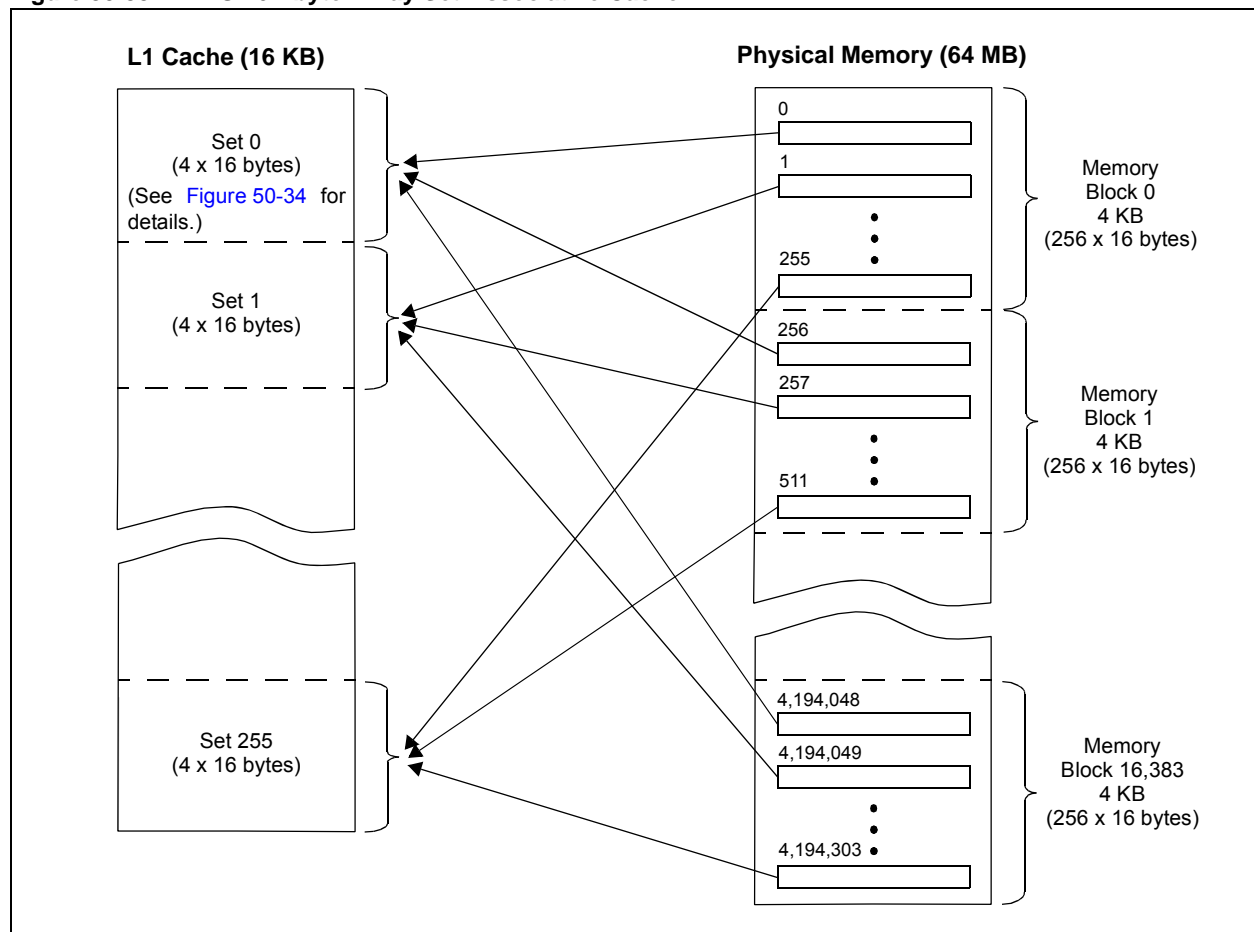
The instruction and data caches may have different sizes and associativities. For example, the instruction cache may be 16 KB in size with 4-way set associativity, while the data cache may be 4 KB with 2-way set associativity. A set-associative cache is a compromise between direct-mapped and fully-associative caches.

In a direct-mapped cache, any location in memory can only be mapped to a single location in the cache. A direct-mapped cache is simple to implement and address, but can be inflexible and lead to thrashing when two heavily used memory locations share the same mapping.

A fully-associative cache allows any location in memory to map to any location in cache. This minimizes collisions, but is expensive to implement.

In a set-associative cache, the cache is divided into groups of lines known as sets. The number of lines per set is known as the associativity. Each memory location is mapped to a set, and may be cached in any line within the set. For example, in a 16 KB, 16 bytes per line, 4-way set-associative cache, the cache is divided into 256 sets of four lines each. As shown, in [Figure 50-33](#), any cacheable (16-byte) memory location is mapped to a single set, and may be mapped to any of four lines within the set. In a system with 64 MB of cacheable memory, each set is shared by 16,384 (256×16 -bytes) blocks of memory.

Figure 50-33: MPU 16 Kbyte 4-way Set-Associative Cache



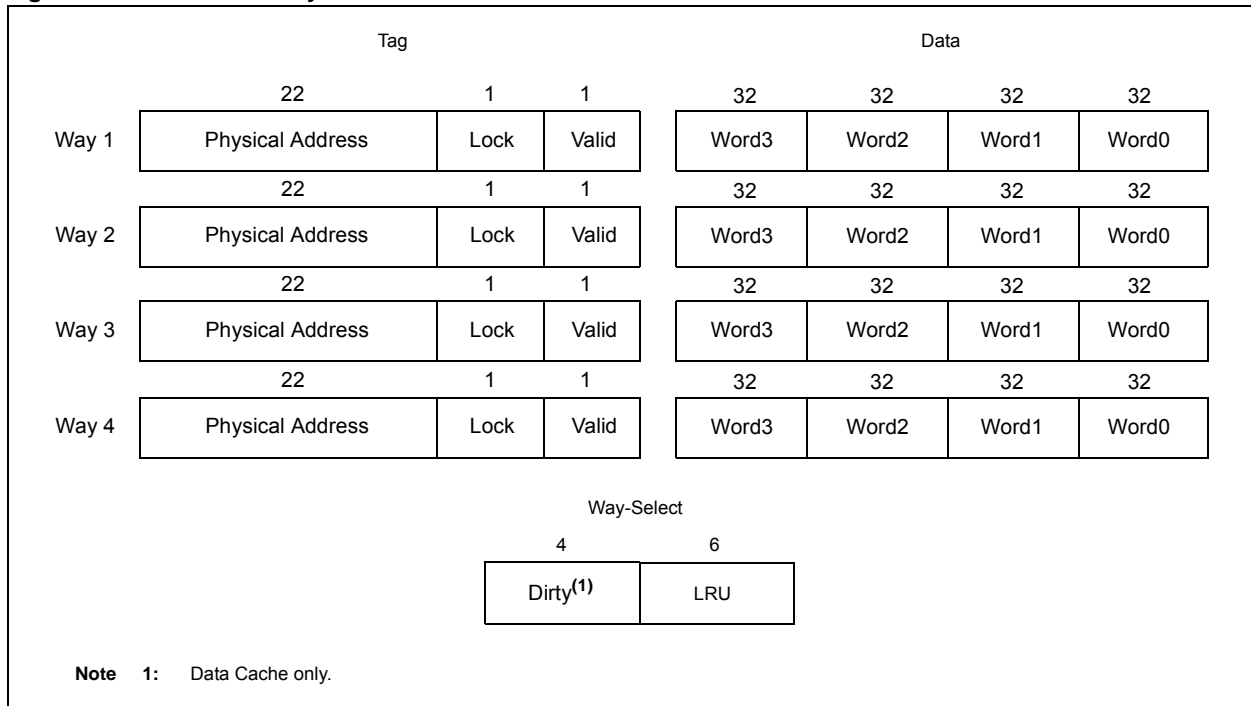
50.20.2 Cache Organization

Each cache line is organized into three arrays:

- Tag
- Data
- Way-select

The tag array holds the physical address of the cached memory location, and the data array holds the cached (16-byte) instruction or data. The tag and data arrays hold four lines of information per set, corresponding to the 4-way set associativity of the cache. The way-select array holds Least Recently Used (LRU) bits that are decoded to select the way to be replaced, according to a LRU algorithm. In the data cache, this array also holds the dirty bits, which indicate whether or not the data needs to be written back to memory before being replaced in the cache.

Figure 50-34: Cache Array Formats



Definitions for [Figure 50-34](#) are as follows:

- **Physical Address:** The upper 22 bits of the physical address (bits <31:10>)
- **Lock:** This bit is set or cleared using the CACHE instruction. When set, the cache line is “locked” (i.e., it cannot be selected for replacement on a cache miss).
- **Valid:** Indicates whether or not the data in the cache line is valid
- **Word 0-3:** Instruction or data words from memory. Each cache line stores four 32-bit words.
- **Dirty:** Present only in the data cache array, there is one dirty bit for each way. The dirty bit is set when data in the cache line is modified. This lets the CPU know to write the cache line back to memory before it is replaced on a cache fill.
- **LRU:** These bits indicate the way to be replaced according to a Least Recently Used (LRU) algorithm

50.20.3 Cacheability Attributes

The cacheability attributes of kseg0 can be set with the CP0 Config1 register. The cacheability options are:

- **Uncached:** Addresses in an uncached memory area are read from main memory, and not from cache. Writes to uncached memory areas are written directly to main memory, without changing the cache contents.
- **Write-back with write allocation:** Loads and instruction fetches first search the cache, reading main memory only if the desired data does not reside in the cache. On data store operations, the cache is first searched to see if the target address is resident in the cache. If it is resident, the cache contents are updated, but main memory is not written. If the cache lookup misses on a store, main memory is read to bring the line into the cache and merge it with the new store data. Therefore, the allocation policy on a cache miss is read- or write-allocate. Data stores will update the appropriate dirty bit in the way-select array to indicate that the line contains modified data. When a line with dirty data is displaced from the cache, it is written back to memory.
- **Write-through with no write allocation:** Loads and instruction fetches first search the cache, reading main memory only if the desired data does not reside in the cache. On data store operations, the cache is first searched to see if the target address is resident in the cache. If it is resident, the cache contents are updated, and main memory is also written. If the cache lookup misses on a store, only main memory is written. Therefore, the allocation policy on a cache miss is read-allocate only.
- **Write-through with write allocation:** Loads and instruction fetches first search the cache, reading main memory only if the desired data does not reside in the cache. On data store operations, the cache is first searched to see if the target address is resident in the cache. If it is resident, the cache contents are updated, and main memory is also written. If the cache lookup misses on a store, main memory is read to bring the line into the cache and merge it with the new store data. In addition, the store data is also written to main memory. Therefore, the allocation policy on a cache miss is read- or write-allocate.

50.20.4 Cache Replacement Policy

The cache replacement policy refers to how the cache determines which way within a set to fill on a cache miss. In a normal cache miss, the lock and LRU tag bits are used to determine the way that is filled. If all ways are valid, any locked ways will not be replaced. If all ways are locked, fill data will not fill into the cache, and Write-back stores turn into Write-through, Write-allocate stores. If the way being replaced is dirty, the 16-byte line will be written back to memory before the fill takes place.

The LRU field in the way-select array is updated in the following ways:

- On a cache hit, the associated way is updated to be the most recently used
- On a cache fill, the filled way is updated to be the most recently used
- On a CACHE instruction, the update of the LRU bits depends on the operation:
 - Index (Writeback) Invalidate: Least recently used
 - Index Load Tag: No update
 - Index Store Tag (CP0 ErrCtl<WST> = 0): Most recently used if CP0 TagLo<V> = 1
Least recently used if CP0 TagLo<V> = 0
 - Index Store Tag (CP0 ErrCtl<WST> = 1): Updated with contents of CP0 TagLo<LRU>
 - Index Store Data: No update
 - Hit Invalidate: Least recently used if a hit is generated; otherwise, no update
 - Fill: Most recently used
 - Hit Writeback: No update
 - Fetch and Lock: Most recently used

50.20.5 Cache Instruction

The contents of the tag, data and way-select arrays can be manipulated by the user with the `CACHE` instruction. The user may fill, lock and invalidate individual cache lines by setting or clearing bits in the tag and data arrays. See [50.21.7 “MPU Cache Instruction”](#) for details on how to use the `CACHE` instruction.

50.20.6 Cache Coherency

Because a cache holds a copy of memory-resident data, it is possible for another bus master to modify cached memory locations, rendering the cached data stale. Likewise, the CPU may update the cache contents, rendering the corresponding memory data stale until it is written back. The PIC32 CPU has no hardware support for maintaining coherency between cache and memory, so it must be handled by software. Most operating systems manage cache coherency issues automatically. Programs running without the benefit of an operating system must manage cache coherency internally.

In Write-through mode, all data writes are written to main memory. In Write-back mode, data writes only go to the cache; the cache may contain the only valid copy of the data until it is written to main memory by a line refill or a `CACHE` instruction.

50.20.7 Cache Initialization

The PIC32 L1 cache tag and data arrays power up to an unknown state and are not affected by reset. Therefore, the caches must be initialized before use. This is typically done by the boot code by writing all-zero values to the tag array. Note that the boot code runs from uncached memory (kseg1).

50.21 CPU INSTRUCTIONS

CPU instructions are organized into the following functional groups:

- Load and store
- Computational
- Jump and branch
- Miscellaneous
- Coprocessor

Each instruction is 32 bits long.

50.21.1 CPU Load and Store Instructions

MIPS processors use a load/store architecture; all operations are performed on operands held in processor registers and main memory is accessed only through load and store instructions.

50.21.1.1 TYPES OF LOADS AND STORES

There are several different types of load and store instructions, each designed for a different purpose:

- Transferring variously-sized fields (for example, LB, SW)
- Trading transferred data as signed or unsigned integers (for example, LHU)
- Accessing unaligned fields (for example, LWR, SWL)
- Atomic memory update (read-modify-write: for instance, LL/SC)

50.21.1.2 LIST OF CPU LOAD AND STORE INSTRUCTIONS

The following data sizes (as defined in the AccessLength field) are transferred by CPU load and store instructions:

- Byte
- Half-word
- Word

Signed and unsigned integers of different sizes are supported by loads that either sign-extend or zero-extend the data loaded into the register.

Unaligned words and double words can be loaded or stored in just two instructions by using a pair of special instructions. For loads a `LWL` instruction is paired with a `LWR` instruction. The load instructions read the left-side or right-side bytes (left or right side of register) from an aligned word and merge them into the correct bytes of the destination register.

50.21.1.3 LOADS AND STORES USED FOR ATOMIC UPDATES

The paired instructions, Load Linked and Store Conditional, can be used to perform an atomic read-modify-write of word or double word cached memory locations. These instructions are used in carefully coded sequences to provide one of several synchronization primitives, including test-and-set, bit-level locks, semaphores, and sequencers and event counts.

50.21.1.4 COPROCESSOR LOADS AND STORES

If a particular coprocessor is not enabled, loads and stores to that processor cannot execute and the attempted load or store causes a Coprocessor Unusable exception. Enabling a coprocessor is a privileged operation provided by the System Control Coprocessor, CP0.

50.21.2 Computational Instructions

Two's complement arithmetic is performed on integers represented in 2's complement notation. These are signed versions of the following operations:

- Add
- Subtract
- Multiply
- Divide

The add and subtract operations labelled “unsigned” are actually modulo arithmetic without overflow detection.

There are also unsigned versions of multiply and divide, as well as a full complement of shift and logical operations. Logical operations are not sensitive to the width of the register.

MIPS32 provides 32-bit integers and 32-bit arithmetic.

50.21.2.1 SHIFT INSTRUCTIONS

The ISA defines two types of shift instructions:

- Those that take a fixed shift amount from a 5-bit field in the instruction word (for instance, SLL, SRL)
- Those that take a shift amount from the low-order bits of a general register (for instance, SRAV, SRLV)

50.21.2.2 MULTIPLY AND DIVIDE INSTRUCTIONS

The multiply instruction performs 32-bit by 32-bit multiplication and creates either 64-bit or 32-bit results. Divide instructions divide a 64-bit value by a 32-bit value and create 32-bit results. With one exception, they deliver their results into the HI and LO special registers. The MUL instruction delivers the lower half of the result directly to a GPR.

- Multiply produces a full-width product twice the width of the input operands; the low half is loaded into LO and the high half is loaded into HI
- Multiply-Add and Multiply-Subtract produce a full-width product twice the width of the input operations and adds or subtracts the product from the concatenated value of HI and LO. The low half of the addition is loaded into LO and the high half is loaded into HI.
- Divide produces a quotient that is loaded into LO and a remainder that is loaded into HI

The results are accessed by instructions that transfer data between HI/LO and the general registers.

50.21.3 Jump and Branch Instructions

50.21.3.1 TYPES OF JUMP AND BRANCH INSTRUCTIONS DEFINED BY THE ISA

The architecture defines the following jump and branch instructions:

- PC-relative conditional branch
- PC-region unconditional jump
- Absolute (register) unconditional jump
- A set of procedure calls that record a return link address in a general register

50.21.3.2 BRANCH DELAYS AND THE BRANCH DELAY SLOT

All branches have an architectural delay of one instruction. The instruction immediately following a branch is said to be in the “branch delay slot”. If a branch or jump instruction is placed in the branch delay slot, the operation of both instructions is undefined.

By convention, if an exception or interrupt prevents the completion of an instruction in the branch delay slot, the instruction stream is continued by re-executing the branch instruction. To permit this, branches must be restartable; procedure calls may not use the register in which the return link is stored (usually GPR 31) to determine the branch target address.

50.21.3.3 BRANCH AND BRANCH LIKELY

There are two versions of conditional branches; they differ in the manner in which they handle the instruction in the delay slot when the branch is not taken and execution falls through.

- Branch instructions execute the instruction in the delay slot
- Branch likely instructions do not execute the instruction in the delay slot if the branch is not taken (they are said to nullify the instruction in the delay slot)

Note: Although the Branch Likely instructions are included in this specification, software is strongly encouraged to avoid the use of the Branch Likely instructions, as they will be removed from a future revision of the MIPS architecture.

50.21.4 microMIPS Instructions

The microMIPS ISA introduces several new instructions, including the ability to load or store multiple 16-bit or 32-bit words from/to memory. Refer to the “MIPS® Architecture for Programmers Volume II-B: The microMIPS32™ Instruction Set” - MD00582, for the full listing and complete description of the microMIPS ISA. This document is available from the Imagination Technologies Ltd. website at: <http://www.imgtec.com/mips/architectures/mips32.asp>.

50.21.5 Miscellaneous Instructions

50.21.5.1 INSTRUCTION SERIALIZATION (SYNC AND SYNCI)

In normal operation, the order in which load and store memory accesses appear to a viewer *outside* the executing processor (for instance, in a multiprocessor system) is not specified by the architecture.

The SYNC instruction can be used to create a point in the executing instruction stream at which the relative order of some loads and stores can be determined: loads and stores executed before the SYNC are completed before loads and stores after the SYNC can start.

The SYNCI instruction synchronizes the processor caches with previous writes or other modifications to the instruction stream.

50.21.5.2 EXCEPTION INSTRUCTIONS

Exception instructions transfer control to a software exception handler in the kernel. There are two types of exceptions, conditional and unconditional. These are caused by the following instructions: `syscall`, `trap`, and `break`.

Trap instructions cause conditional exceptions based upon the result of a comparison. System call and breakpoint instructions cause unconditional exceptions.

50.21.5.3 CONDITIONAL MOVE INSTRUCTIONS

MIPS32 includes instructions to conditionally move one CPU general register to another, based on the value in a third general register.

50.21.5.4 NOP INSTRUCTIONS

The `NOP` instruction is actually encoded as an all-zero instruction. MIPS processors special-case this encoding as performing no operation, and optimize execution of the instruction. In addition, the `SSNOP` instruction takes up one issue cycle on any processor, including super-scalar implementations of the architecture.

50.21.5.5 TLB INSTRUCTIONS (MPU ONLY)

The `TLBR`, `TLBWI` and `TLBWR` instructions can be used to read or write TLB entries. The Index register is loaded with the VPN of a TLB entry, and the tag contents are read into the `EntryLo0`, `EntryLo1` and `PageMask` registers. TLB entries are written in a similar manner using the `TLBWI` and `TLBWR` instructions.

50.21.6 Coprocessor Instructions

50.21.6.1 WHAT COPROCESSORS DO

Coprocessors are alternate execution units, with register files separate from the CPU. In abstraction, the MIPS architecture provides for up to four coprocessor units, numbered 0 to 3. Each level of the ISA defines a number of these coprocessors. Coprocessor 0 is always used for system control and coprocessor 1 and 3 are used for the floating point unit. Coprocessor 2 is reserved for implementation-specific use. PIC32 devices only implement Coprocessor 0.

A coprocessor may have two different register sets:

- Coprocessor general registers
- Coprocessor control registers

Each set contains up to 32 registers. Coprocessor computational instructions may use the registers in either set.

50.21.6.2 SYSTEM CONTROL COPROCESSOR 0 (CP0)

The system controller for all MIPS processors is implemented as coprocessor 0 (CP0), the System Control Coprocessor. It provides the processor control, memory management, and exception handling functions.

50.21.6.3 COPROCESSOR LOAD AND STORE INSTRUCTIONS

Explicit load and store instructions are not defined for CP0; the move to and from coprocessor instructions must be used to write and read the CP0 registers.

50.21.7 MPU Cache Instruction

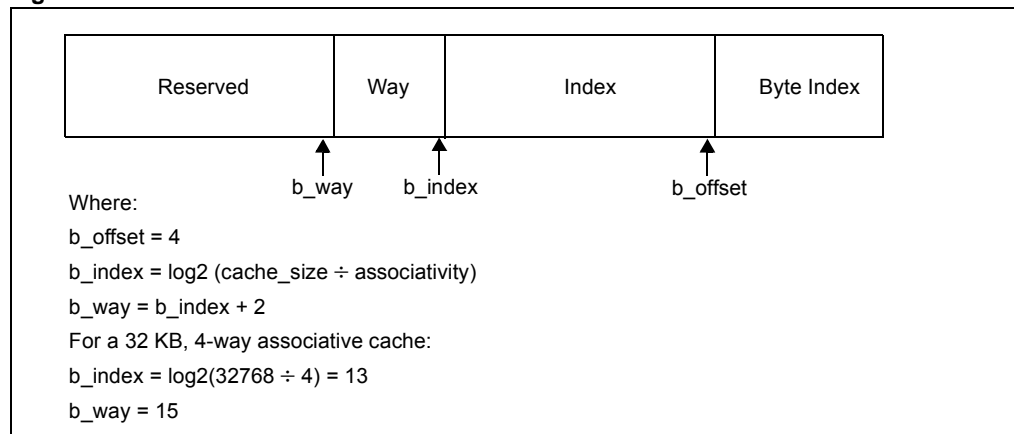
The `CACHE` instruction can be used to read or write the data, tag or way-select arrays. It takes two arguments: a 5-bit operation field, and a 16-bit offset field. The format is similar to a load and store instruction, with the offset field consisting of a base register plus 16-bit signed immediate offset:

```
cache OP, offset(base)
```

The 16-bit offset is sign-extended and added to the contents of the base register to form an effective address. The effective address can be one of two types:


- **Address:** The effective address is the virtual memory location of some instruction or data, and is processed just like a normal cached access.
- **Index:** The effective address encodes the cache virtual index, the byte location within the cache line, and the way. The exact size and boundary of each value depends on the cache size and configuration. In general, an index effective address has the following format, as shown in [Figure 50-35](#).

Figure 50-35: Cache Effective Address



The operation field encodes the cache on which to perform the operation and the operation itself, as shown in Figure 50-36.

Figure 50-36: Cache Operation



The lower two bits of the operation field specify the cache on which to perform the operation, as shown in Table 50-43.

Table 50-43: Cache Instruction OP bits Encoding (Cache)

OP<1:0> bits	Cache
`0b00	Instruction Cache
`0b01	Data Cache
`0b10	Reserved
`0b11	Reserved

The remaining bits of the operation field specify the operation to perform, as shown in Table 50-44.

Table 50-44: Cache Instruction OP bits Encoding

OP<4:2> bits	Cache	Name	Type	Operation
`0b000	Instruction	Index Invalidate	Index	Set the state of the cache block at the specified address to invalid. Commonly used to initialize the instruction cache at startup.
	Data			Set the state of the cache block at the specified address to invalid. If invalid and dirty, write back to memory first. Should NOT be used to initialize the data cache at startup.
`0b001	—	—	—	Reserved.
`0b010	Both	Index Store Tag	Index	Sets the cache tag using the values from the TagLo register. Commonly used to initialize the data cache at startup by setting TagLo to zero.
`0b011	—	—	—	Reserved.
`0b100	Both	Hit Invalidate	Address	If the cache block contains the specified address, set to invalid. Do not write back if dirty.
`0b101	Instruction	Fill	Address	Fill the cache from the specified address. Similar to a cache miss.
	Data			If the cache block contains the specified address, set to invalid. Write back if dirty. Recommended way to invalidate a cache line in a running cache.
`0b110	—	—	—	Reserved.
`0b111	—	—	—	Reserved.

PIC32 Family Reference Manual

The CACHE Index Load Tag and Index Store Tag instructions can be used to read and write the LRU bits in the instruction or data cache way-select arrays by setting the WST bit in the ErrCtl CP0 register. Valid LRU field values are shown in [Table 50-45](#).

Table 50-45: LRU bit Way Selection Encoding

Selection Order	LRU<5:0> bits	Selection Order	LRU<5:0> bits	Selection Order	LRU<5:0> bits	Selection Order	LRU<5:0> bits
0123	000000	1023	000100	2013	100010	3012	011001
0132	000001	1032	000101	2031	110010	3021	011011
0213	000010	1203	100100	2103	100110	3102	011101
0231	010010	1230	101100	2130	101110	3120	111101
0312	010001	1302	001101	2301	111010	3201	111011
0321	010011	1320	101101	2310	111110	3210	111111

The order is indicated by listing the least-recently used way to the left and the most-recently used way to the right. Note that not all values are valid.

50.22 MIPS DSP ASE INSTRUCTIONS

Note: DSP ASE is not available on all devices. Refer to the “CPU” chapter in the specific device data sheet to determine availability.

The MIPS DSP ASE Revision 2 instructions can be classified into the following subclasses:

- Arithmetic
- GPR-based shift
- Multiply
- Bit manipulation
- Compare-Pick
- DSP Control Access
- Indexed-Load
- Branch

The MIPS DSP ASE adds four new registers. The software is required to recognize the presence of the MIPS DSP ASE and to include these additional registers in context save and restore operations.

Three additional HI/LO registers are available to create a total of four accumulator registers. Many common DSP computations involve accumulation (e.g., convolution). MIPS DSP ASE Revision 2 instructions that target the accumulators use two bits to specify the destination accumulator, with the zero value referring to the original accumulator of the MIPS architecture.

A new control register, DSPControl, is used to hold extra state bits needed for efficient support of the new instructions. [Register 50-63](#) illustrates the bits in this register.

50.22.1 DSPControl Register

The DSPControl register is used to hold extra state bits needed for efficient support of the MIPS DSP ASE Revision 2 instruction set.

Note: This register is not available on all devices. Refer to the “CPU” chapter in the specific device data sheet to determine availability.

Register 50-63: DSPControl: MIPS DSP ASE Control Register

Bit Range	Bit 31/23/15/7	Bit 30/22/14/6	Bit 29/21/13/5	Bit 28/20/12/4	Bit 27/19/11/3	Bit 26/18/10/2	Bit 25/17/9/1	Bit 24/16/8/0
31:24	U-0	U-0	U-0	U-0	R/W-0	R/W-0	R/W-0	R/W-0
	—	—	—	—	CCOND<3:0>			
23:16	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	OUFLAG<7:0>							
15:8	U-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	—	EFI	C	SCOUNT<5:1>				
7:0	R/W-0	U-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0	R/W-0
	SCOUNT<0>	—	POS<5:0>					

Legend:

R = Readable bit

W = Writable bit

U = Unimplemented bit, read as ‘0’

-n = Value at POR

‘1’ = Bit is set

‘0’ = Bit is cleared

x = Bit is unknown

bit 31-28 **Unimplemented:** Read as ‘0’

bit 27-24 **CCOND<3:0>:** Condition Code bits

These bits are set by vector comparison instructions and are used as source selectors by the group of **PICK** instructions. The vector element size determines the number of bits set by a comparison (1, 2, or 4). Bits that are not set after the comparison are unpredictable.

PIC32 Family Reference Manual

Register 50-63: DSPControl: MIPS DSP ASE Control Register (Continued)

bit 23-16 **OUFLAG<7:0>**: Overflow/Underflow Indication bits

The bits of the overflow flag OUFLAG<7:0> are set by the instructions listed in the following table. These bits are sticky and can be reset only by an explicit write to these bits in the register (using the WRDSP instruction).

Bit Number	Description
23	This bit is set when the destination is accumulator (HI-LO pair) zero, and an operation overflow or underflow occurs. These instructions are: DPAQ_S, DPAQ_SA, DPSQ_S, DPSQ_SA, DPAQX_S, DPAQX_SA, DPSQX_S, DPSQX_SA, MAQ_S, MAQ_SA and MULSAQ_S.
23	Same instructions as above, when the destination is accumulator (HI-LO pair) one.
21	Same instructions as above, when the destination is accumulator (HI-LO pair) two.
20	Same instructions as above, when the destination is accumulator (HI-LO pair) three.
19	Instructions that set this bit on an overflow/underflow: ABSQ_S, ADDQ, ADDQ_S, ADDU, ADDU_S, ADDWC, SUBQ, SUBQ_S, SUBU and SUBU_S.
18	Instructions that set this bit on an overflow/underflow: MUL, MUL_S, MULEQ_S, MULEU_S, MULQ_RS, and MULQ_S.
17	Instructions that set this bit on an overflow/underflow: PRECRQ_RS, SHLL, SHLL_S, SHLLV, and SHLLV_S.
16	Instructions that set this bit on an overflow/underflow: EXTR, EXTR_S, EXTR_RS, EXTRV, and EXTRV_RS.

bit 15 **Unimplemented**: Read as '0'

bit 14 **EFI**: Extract Fail Indicator bit

This bit is set to '1' when one of the extraction instructions (EXTP, EXTPV, EXTPDP, or EXTPDPV) fails. A failure occurs when there are insufficient bits to extract (i.e., when the value of the POS<5:0> bits is less than the size argument specified in the instruction).

bit 13 **C**: Carry bit

This bit is set and used by a special add instruction to implement a 64-bit addition across two GPRs in a microMIPS32 implementation. Instruction ADDSC sets the bit and instruction ADDWC uses this bit.

bit 12-7 **SCOUNT<5:0>**: Size Count bit

This bit is used by the INSV instruction to specify the size of the bit field to be inserted.

bit 6 **Unimplemented**: Read as '0'

bit 5-0 **POS<5:0>**: Insert/Extract Position bits

These bits are used by the variable insert instruction, INSV, to specify the position to insert bits. It is also used to indicate the extract position for the EXTP, EXTPV, EXTPDP, or EXTPDPV instructions.

50.23 CPU INITIALIZATION

Software is required to initialize the following parts of the device after a Reset event.

50.23.1 General Purpose Registers

The CPU register file powers up in an unknown state with the exception of r0 which is always '0'. Initializing the rest of the register file is not required for proper operation in hardware. However, depending on the software environment, several registers may need to be initialized. Some of these are:

- sp – Stack pointer
- gp – Global pointer
- fp – Frame pointer

50.23.2 Coprocessor 0 State

Miscellaneous CP0 states need to be initialized prior to leaving the boot code. There are various exceptions that are blocked by ERL = 1 or EXL = 1, and which are not cleared by Reset. These can be cleared to avoid taking spurious exceptions when leaving the boot code.

Table 50-46: CPU Initialization

CP0 Register	Action
Cause	WP (Watch Pending), SW0/1 (Software Interrupts) should be cleared.
Config	Typically, the K0, KU and K23 fields should be set to the desired Cache Coherency Algorithm (CCA) value prior to accessing the corresponding memory regions.
Count ⁽¹⁾	Should be set to a known value if Timer Interrupts are used.
Compare ⁽¹⁾	Should be set to a known value if Timer Interrupts are used. The write to Compare will also clear any pending Timer Interrupts (Thus, Count should be set before Compare to avoid any unexpected interrupts).
Status	Desired state of the device should be set.
Other CP0 state	Other registers should be written before they are read. Some registers are not explicitly writable, and are only updated as a by-product of instruction execution or a taken exception. Uninitialized bits should be masked off after reading these registers.

Note 1: When the Count register is equal to the Compare register a timer interrupt is signaled. There is a mask bit in the interrupt controller to disable passing this interrupt to the CPU if desired.

50.23.3 System Bus

The System Bus should be initialized before switching to User mode or before executing from DRM. The values written to the System Bus are based on the memory layout of the application to be run.

50.23.4 Caches (MPU Only)

The cache tag and data arrays power-up to an unknown state and are not affected by a reset. Every tag in the cache arrays should be initialized to an invalid state using the `CACHE` instruction (Index Invalidate function).

50.24 EFFECTS OF A RESET

50.24.1 Master Clear Reset

The PIC32 core is not fully initialized by a hardware Reset. Only a minimal subset of the processor state is cleared. This is enough to bring the core up while running in unmapped and uncached code space. All other processor state can then be initialized by software. Power-up Reset brings the device into a known state. A Soft Reset can be forced by asserting the Master Clear (MCLR) pin. This distinction is made for compatibility with other MIPS processors. In practice, both resets are handled identically.

50.24.1.1 COPROCESSOR 0 STATE

Much of the hardware initialization occurs in Coprocessor 0, which are described in [Table 50-47](#).

Table 50-47: Bits Cleared or Set by Reset

Register Name	Bit Name	Cleared or Set	Value	Cleared or Set By
Status	BEV	Set	1	Reset or Soft Reset
	TS	Cleared	0	Reset or Soft Reset
	SR	Set	1	Soft Reset
	SR	Cleared	0	Reset
	NMI	Cleared	0	Reset or Soft Reset
	ERL	Set	1	Reset or Soft Reset
	RP	Cleared	0	Reset or Soft Reset
All Configuration Registers: Config Config1 Config2 Config3 Config4 Config5 Config7	Note: The reset state of the CP0 Configuration registers may vary between devices. Refer to the “CPU” chapter in the specific device data sheet for details.			
Debug	DM	Cleared	0	Reset or Soft Reset ⁽¹⁾
	LSNM	Cleared	0	Reset or Soft Reset
	IBUSEP	Cleared	0	Reset or Soft Reset
	IEXI	Cleared	0	Reset or Soft Reset
	SSt	Cleared	0	Reset or Soft Reset

Note 1: Unless EJTAGBOOT option is used to boot into Debug mode.

50.24.1.2 BUS STATE MACHINES

All pending bus transactions are aborted and the state machines in the SRAM interface unit are reset when a Reset or Soft Reset exception is taken.

50.24.2 Fetch Address

Upon Reset/Soft Reset, unless the EJTAGBOOT option is used, the fetch is directed to VA 0xB-FC00000 (PA 0x1FC00000). This address is in kseg1, which is unmapped and uncached.

50.24.3 WDT Reset

The status of the CPU registers after a WDT event depends on the operational mode of the CPU prior to the WDT event.

If the device was not in Sleep a WDT event will force registers to a Reset value.

50.25 RELATED APPLICATION NOTES

This section lists application notes that are related to this section of the manual. These application notes may not be written specifically for the PIC32 device family, but the concepts are pertinent and could be used with modification and possible limitations. The current application notes related to the PIC32 CPU for Devices with MIPS32® microAptiv™ and M-Class Cores include the following:

Title	Application Note #
No related application notes at this time.	N/A

Note: Please visit the Microchip Web site (www.microchip.com) for additional application notes and code examples for the PIC32 family of devices.

50.26 REVISION HISTORY

Revision A (April 2013)

This is the initial released version of the document.

Revision B (July 2015)

This revision includes the following updates:

- The document title was changed to: “CPU for Devices with MIPS32® microAptiv™ and M-Class Cores”
- M-Class core related updates were implemented throughout the document
- All references to MIPS Technologies were changed to Imagination Technologies Ltd.
- The Select number for the Debug2 register was changed from 5 to 6 (see [Register 50-45](#))
- The following registers were added:
 - [BadInstr](#): Bad Instruction Word Register; CP0 Register 8, Select 1 ([Register 50-11](#))
 - [BadInstrP](#): Bad Prior Branch Instruction Register; CP0 Register 8, Select 2 ([Register 50-12](#))
 - [CacheErr](#): Cache Error Register; CP0 Register 27, Select 0 ([Register 50-51](#))
 - [KScratchn](#): Kernel Mode Scratchpad Registers; CP0 Register 31, Select 2-3 ([Register 50-57](#))
- [50.5.2 “Architecture Release 5”](#) was added
- [50.12 “Floating Point Unit \(FPU\)”](#) was added
- [50.14 “Coprocessor 1 \(CP1\) Registers”](#) was added, which includes the following registers:
 - [FIR](#): Floating Point Implementation Register; CP1 Control Register 0 ([Register 50-58](#))
 - [FCC](#): Floating Point Condition Codes Register; CP1 Control Register 25 ([Register 50-59](#))
 - [FEXR](#): Floating Point Exceptions Register; CP1 Control Register 26 ([Register 50-60](#))
 - [FENR](#): Floating Point Enables Register; CP1 Control Register 28 ([Register 50-61](#))
 - [FCSR](#): Floating Point Control and Status Register; CP1 Control Register 31 ([Register 50-62](#))
- In addition, minor updates to text and formatting were incorporated throughout the document

Note the following details of the code protection feature on Microchip devices:

- Microchip products meet the specification contained in their particular Microchip Data Sheet.
- Microchip believes that its family of products is one of the most secure families of its kind on the market today, when used in the intended manner and under normal conditions.
- There are dishonest and possibly illegal methods used to breach the code protection feature. All of these methods, to our knowledge, require using the Microchip products in a manner outside the operating specifications contained in Microchip's Data Sheets. Most likely, the person doing so is engaged in theft of intellectual property.
- Microchip is willing to work with the customer who is concerned about the integrity of their code.
- Neither Microchip nor any other semiconductor manufacturer can guarantee the security of their code. Code protection does not mean that we are guaranteeing the product as "unbreakable."

Code protection is constantly evolving. We at Microchip are committed to continuously improving the code protection features of our products. Attempts to break Microchip's code protection feature may be a violation of the Digital Millennium Copyright Act. If such acts allow unauthorized access to your software or other copyrighted work, you may have a right to sue for relief under that Act.

Information contained in this publication regarding device applications and the like is provided only for your convenience and may be superseded by updates. It is your responsibility to ensure that your application meets with your specifications. MICROCHIP MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WHETHER EXPRESS OR IMPLIED, WRITTEN OR ORAL, STATUTORY OR OTHERWISE, RELATED TO THE INFORMATION, INCLUDING BUT NOT LIMITED TO ITS CONDITION, QUALITY, PERFORMANCE, MERCHANTABILITY OR FITNESS FOR PURPOSE. Microchip disclaims all liability arising from this information and its use. Use of Microchip devices in life support and/or safety applications is entirely at the buyer's risk, and the buyer agrees to defend, indemnify and hold harmless Microchip from any and all damages, claims, suits, or expenses resulting from such use. No licenses are conveyed, implicitly or otherwise, under any Microchip intellectual property rights unless otherwise stated.

Trademarks

The Microchip name and logo, the Microchip logo, dsPIC, FlashFlex, flexPWR, JukeBlox, KEELOQ, KEELOQ logo, Klear, LANCheck, MediaLB, MOST, MOST logo, MPLAB, OptoLyzer, PIC, PICSTART, PIC³² logo, RightTouch, SpyNIC, SST, SST Logo, SuperFlash and UNI/O are registered trademarks of Microchip Technology Incorporated in the U.S.A. and other countries.

The Embedded Control Solutions Company and mTouch are registered trademarks of Microchip Technology Incorporated in the U.S.A.

Analog-for-the-Digital Age, BodyCom, chipKIT, chipKIT logo, CodeGuard, dsPICDEM, dsPICDEM.net, ECAN, In-Circuit Serial Programming, ICSP, Inter-Chip Connectivity, KlearNet, KlearNet logo, MiWi, MPASM, MPF, MPLAB Certified logo, MPLIB, MPLINK, MultiTRAK, NetDetach, Omniscient Code Generation, PICDEM, PICDEM.net, PICkit, PICtail, RightTouch logo, REAL ICE, SQI, Serial Quad I/O, Total Endurance, TSHARC, USBCheck, VariSense, ViewSpan, WiperLock, Wireless DNA, and ZENA are trademarks of Microchip Technology Incorporated in the U.S.A. and other countries.

SQTP is a service mark of Microchip Technology Incorporated in the U.S.A.

Silicon Storage Technology is a registered trademark of Microchip Technology Inc. in other countries.

GestIC is a registered trademark of Microchip Technology Germany II GmbH & Co. KG, a subsidiary of Microchip Technology Inc., in other countries.

All other trademarks mentioned herein are property of their respective companies.

© 2013-2015, Microchip Technology Incorporated, Printed in the U.S.A., All Rights Reserved.

ISBN: 978-1-63277-631-0

QUALITY MANAGEMENT SYSTEM
CERTIFIED BY DNV
= ISO/TS 16949 =

Microchip received ISO/TS-16949:2009 certification for its worldwide headquarters, design and wafer fabrication facilities in Chandler and Tempe, Arizona; Gresham, Oregon and design centers in California and India. The Company's quality system processes and procedures are for its PIC® MCUs and dsPIC® DSCs, KEELOQ® code hopping devices, Serial EEPROMs, microperipherals, nonvolatile memory and analog products. In addition, Microchip's quality system for the design and manufacture of development systems is ISO 9001:2000 certified.

Worldwide Sales and Service

AMERICAS

Corporate Office
2355 West Chandler Blvd.
Chandler, AZ 85224-6199
Tel: 480-792-7200
Fax: 480-792-7277
Technical Support:
<http://www.microchip.com/support>
Web Address:
www.microchip.com

Atlanta
Duluth, GA
Tel: 678-957-9614
Fax: 678-957-1455

Austin, TX
Tel: 512-257-3370

Boston
Westborough, MA
Tel: 774-760-0087
Fax: 774-760-0088

Chicago
Itasca, IL
Tel: 630-285-0071
Fax: 630-285-0075

Cleveland
Independence, OH
Tel: 216-447-0464
Fax: 216-447-0643

Dallas
Addison, TX
Tel: 972-818-7423
Fax: 972-818-2924

Detroit
Novi, MI
Tel: 248-848-4000

Houston, TX
Tel: 281-894-5983

Indianapolis
Noblesville, IN
Tel: 317-773-8323
Fax: 317-773-5453

Los Angeles
Mission Viejo, CA
Tel: 949-462-9523
Fax: 949-462-9608

New York, NY
Tel: 631-435-6000

San Jose, CA
Tel: 408-735-9110

Canada - Toronto
Tel: 905-673-0699
Fax: 905-673-6509

ASIA/PACIFIC

Asia Pacific Office
Suites 3707-14, 37th Floor
Tower 6, The Gateway
Harbour City, Kowloon

Hong Kong
Tel: 852-2943-5100
Fax: 852-2401-3431

Australia - Sydney
Tel: 61-2-9868-6733
Fax: 61-2-9868-6755

China - Beijing
Tel: 86-10-8569-7000
Fax: 86-10-8528-2104

China - Chengdu
Tel: 86-28-8665-5511
Fax: 86-28-8665-7889

China - Chongqing
Tel: 86-23-8980-9588
Fax: 86-23-8980-9500

China - Dongguan
Tel: 86-769-8702-9880

China - Hangzhou
Tel: 86-571-8792-8115
Fax: 86-571-8792-8116

China - Hong Kong SAR
Tel: 852-2943-5100
Fax: 852-2401-3431

China - Nanjing
Tel: 86-25-8473-2460
Fax: 86-25-8473-2470

China - Qingdao
Tel: 86-532-8502-7355
Fax: 86-532-8502-7205

China - Shanghai
Tel: 86-21-5407-5533
Fax: 86-21-5407-5066

China - Shenyang
Tel: 86-24-2334-2829
Fax: 86-24-2334-2393

China - Shenzhen
Tel: 86-755-8864-2200
Fax: 86-755-8203-1760

China - Wuhan
Tel: 86-27-5980-5300
Fax: 86-27-5980-5118

China - Xian
Tel: 86-29-8833-7252
Fax: 86-29-8833-7256

ASIA/PACIFIC

China - Xiamen
Tel: 86-592-2388138
Fax: 86-592-2388130

China - Zhuhai
Tel: 86-756-3210040
Fax: 86-756-3210049

India - Bangalore
Tel: 91-80-3090-4444
Fax: 91-80-3090-4123

India - New Delhi
Tel: 91-11-4160-8631
Fax: 91-11-4160-8632

India - Pune
Tel: 91-20-3019-1500

Japan - Osaka
Tel: 81-6-6152-7160
Fax: 81-6-6152-9310

Japan - Tokyo
Tel: 81-3-6880-3770
Fax: 81-3-6880-3771

Korea - Daegu
Tel: 82-53-744-4301
Fax: 82-53-744-4302

Korea - Seoul
Tel: 82-2-554-7200
Fax: 82-2-558-5932 or
82-2-558-5934

Malaysia - Kuala Lumpur
Tel: 60-3-6201-9857
Fax: 60-3-6201-9859

Malaysia - Penang
Tel: 60-4-227-8870
Fax: 60-4-227-4068

Philippines - Manila
Tel: 63-2-634-9065
Fax: 63-2-634-9069

Singapore
Tel: 65-6334-8870
Fax: 65-6334-8850

Taiwan - Hsin Chu
Tel: 886-3-5778-366
Fax: 886-3-5770-955

Taiwan - Kaohsiung
Tel: 886-7-213-7828

Taiwan - Taipei
Tel: 886-2-2508-8600
Fax: 886-2-2508-0102

Thailand - Bangkok
Tel: 66-2-694-1351
Fax: 66-2-694-1350

EUROPE

Austria - Wels
Tel: 43-7242-2244-39
Fax: 43-7242-2244-393

Denmark - Copenhagen
Tel: 45-4450-2828
Fax: 45-4485-2829

France - Paris
Tel: 33-1-69-53-63-20
Fax: 33-1-69-30-90-79

Germany - Dusseldorf
Tel: 49-2129-3766400

Germany - Karlsruhe
Tel: 49-721-625370

Germany - Munich
Tel: 49-89-627-144-0
Fax: 49-89-627-144-44

Italy - Milan
Tel: 39-0331-742611
Fax: 39-0331-466781

Italy - Venice
Tel: 39-049-7625286

Netherlands - Druenen
Tel: 31-416-690399
Fax: 31-416-690340

Poland - Warsaw
Tel: 48-22-3325737

Spain - Madrid
Tel: 34-91-708-08-90
Fax: 34-91-708-08-91

Sweden - Stockholm
Tel: 46-8-5090-4654

UK - Wokingham
Tel: 44-118-921-5800
Fax: 44-118-921-5820